Deepotus Video Gen guide

Everything you can do with the app, one task at a time: short steps, animated pictures taken from the real app, and the icons you will see on screen.

I want to…

- make my first video in a few minutes
- turn a news story into a scheduled post
- tell my novel as episodes
- edit, grade and deliver a video
- schedule my posts for the week
- find an image and send it elsewhere
- retouch a photo
- draw a poster or a logo
- create game sprites and tiles
- print my card game
- 3D-print a figurine
- know how much I spend
- pick an AI model and the key that pays for it
- understand an icon
- ask the app how to do something
- fix a problem

Take a tour of the app

Ask Deepotus

Install, pick a language, update

Connect your keys and AI engines

Choose your AI models and their route

Keep spending under control

Settings, personas and branding

Work for several brands

Move to another computer

Drive it from your phone

Make a video in minutes (Quick)

Compose with the Studio

Reuse templates and brand kits

Tell a story in episodes

Scan a page, an inspiration, a diagram

Sound, voices and visual effects

Play an avatar, deferred or live

Install and train a local RVC voice

Edit a video

Grade, mix and add titles

Deliver: renders and formats

Record the app and go live

Plan and publish your week

Run a campaign with the agent

Watch, prospects and prospecting

Connect your accounts

Work for clients

Talk to your clients on WhatsApp

Follow the news and turn it into posts

Organise and find your media

Retouch a photo (Photolab)

Layers, masks and workspaces

Automate and export (Photolab)

Draw vectors (Vectorlab)

Pixel art, export, maps and 3D objects

Make game assets

Animate sprites (Sprite Lab)

Paint tiled levels (Tile Lab)

Create playing cards (Card Forge)

From image to 3D object

Create materials (Material Forge)

Prepare a part before slicing (Workbench)

3D-print your creations

Stage shots on the 3D Set

End-to-end recipes, from idea to post

Read the icons

Keyboard shortcuts

Troubleshooting

GET STARTED

Take a tour of the app

You will find your way around the app, move from one screen to another without searching, and read what the top bar tells you. Five minutes, free.

Deepotus Video Gen is a studio that runs **on your own computer**. It turns an idea into a video, a picture, a 3D object or a published post. You plug in your own AI provider keys: you pay those providers, not the app. Everything you create stays with you.

The screen at a glance

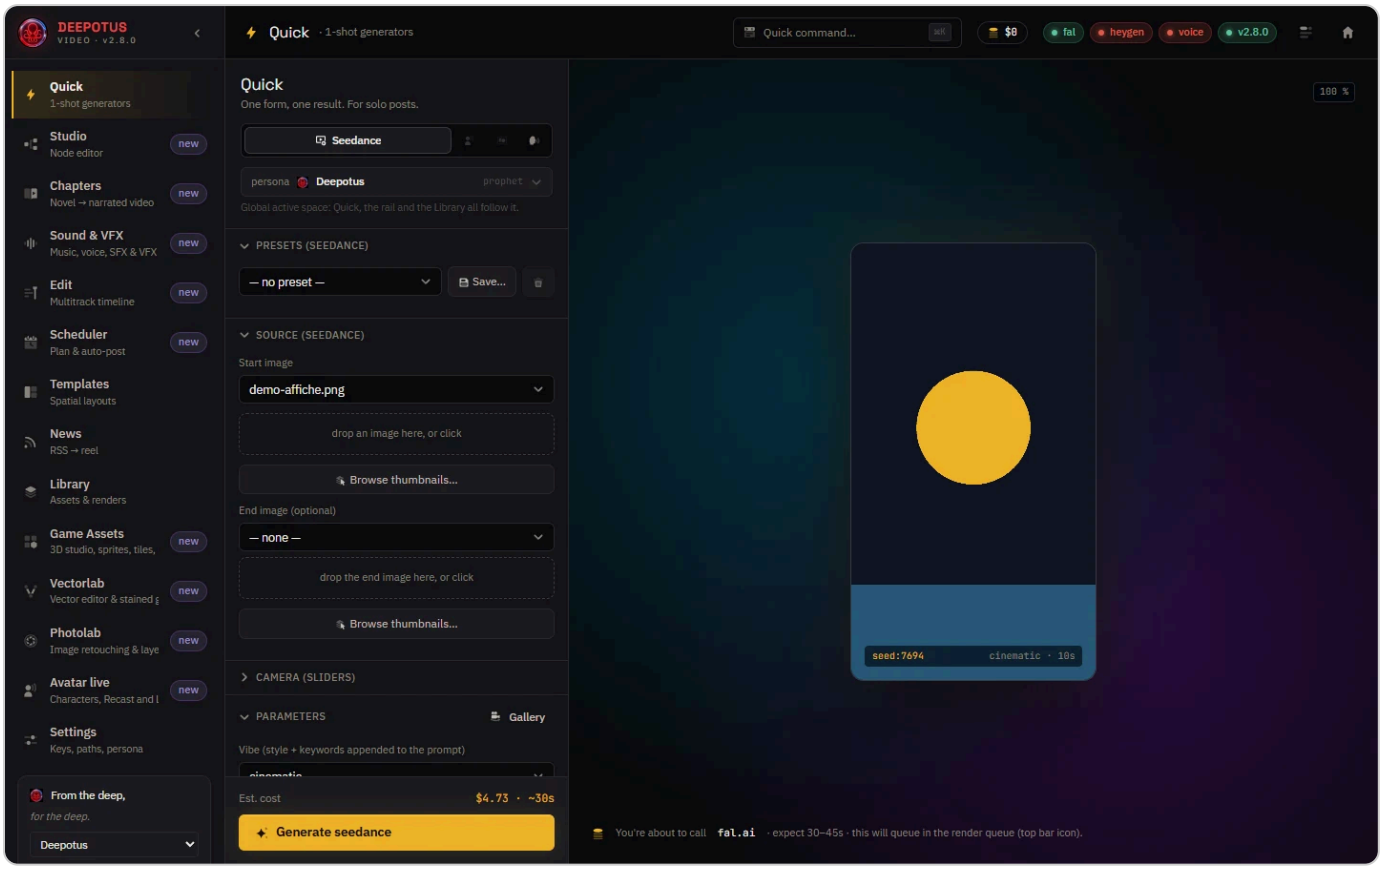

1. The rail: one screen per task 2. Ctrl+K: go anywhere from the keyboard 3. This month's spending and key status, always in sight

The screen rail on the left, the top bar with the quick command, the spending counter and the key status.

1. **The rail, on the left**: one screen per task. `Quick` for a one-click video, `Studio` to compose, `Edit` to cut a video, `Library` to find your media, `Settings` for the keys. The arrow at the top of the rail folds it into icons only.
2. **The quick command**: click `Quick command...` or press `Ctrl`+`K`, type a few letters, Enter.
3. **The spending counter** shows what you spent this month with your providers. One click opens Settings.

4. **The badges** `fal`, `heygen` and `voice` turn green when the matching key is set. One click takes you to that key.
5. **The render queue**, top right, lists what is being computed and what just finished, in the `Running` and `Recent` tabs.

Go fast with the quick command

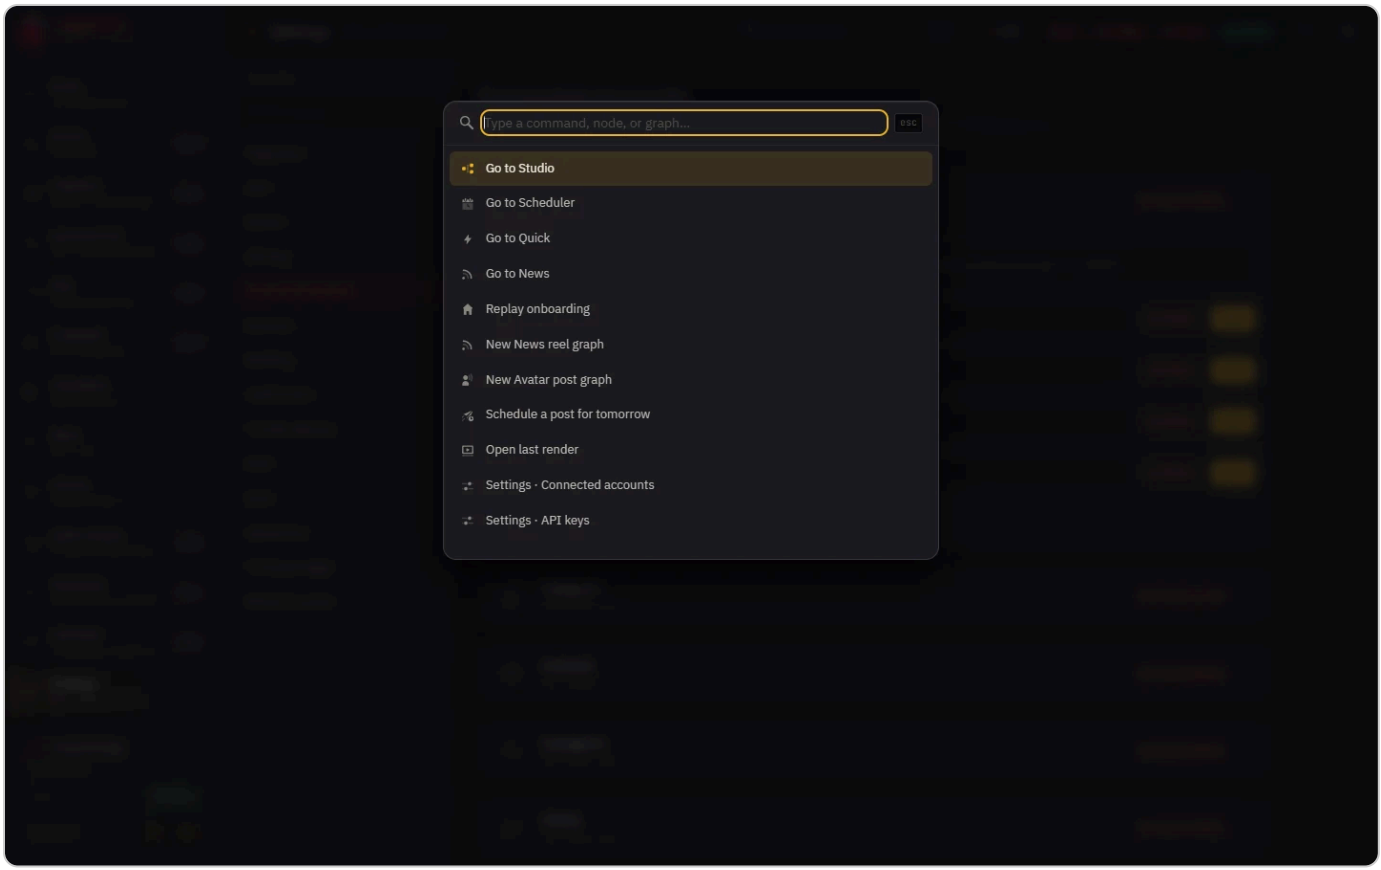

1. Type, then Enter

Ctrl+K opens the quick command: type a few letters to go to a screen or start an action.

The quick command goes to any screen (`Go to Studio`, `Go to News`...), opens the last render, creates a starter graph or opens a Settings section. `Replay onboarding` restarts the first-launch tour.

The screens of the rail

Quick

A one-click video: animated clip, talking avatar, composition, voiceover.

Studio

A node graph to compose a video piece by piece.

Chapters

From a text or a novel to a narrated video, scene by scene.

Sound & VFX
Music, voices, sound effects and visual effects.

Edit
A multitrack timeline to cut, grade and deliver.

Scheduler
Plan and publish your posts for the week.

Templates
Reusable layouts and your brand kits.

News
Follow the news and turn it into short videos.

Library
All your pictures, renders, sounds and objects, with their origin.

Game Assets
3D, sprites, tiles, materials and playing cards.

Vectorlab
Vector drawing, pixel art, maps and printable objects.

Photolab
Photo retouching with layers, masks and filters.

The guide opens from the foot of the rail: FR or EN.

What the app does to protect you

Nothing leaves without you
Only this PC can change, generate or spend. Another device on the network can only read, and only once paired.

Errors that make sense
A missing key or a server refusal is shown in plain words, with what to do.

Safe imports
A picture with the same name becomes name-1, name-2…; a file that is not a real picture is refused and named.

In-app dialogs
Confirmations and text prompts use one Deepotus dialog, never a browser pop-up.

Good to know

> The desktop icon opens the app at `http://deepotus.localhost:8765`. Each browser tab carries the name of the open tool, for example “Deepotus — Vectorlab”.

> Only AI generations cost money, and your provider bills them. The chapter Keep spending under control shows how to set limits.

Shortcuts

Key	Effect
`Ctrl` + `K`	Opens the quick command
`Esc`	Closes the quick command or a dialog

GET STARTED

Ask Deepotus

You will be able to ask a question about the app without leaving it: Deepotus answers in a few steps, takes you to the right screen and shows the button to click. Two minutes; you need a text engine connected (a key, or Ollama on this PC).

Before you start

The assistant uses a text engine you already have: Anthropic, OpenAI, Gemini, or a local server such as Ollama. With no engine, it shows you where to add a key. It **never runs** a paid action for you: it explains and shows, you click.

Ask a question

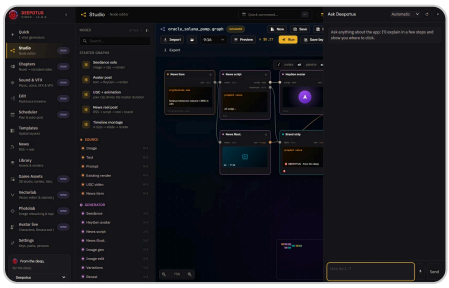

1. F1 opens the drawer

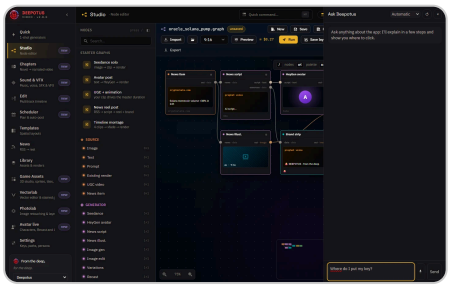

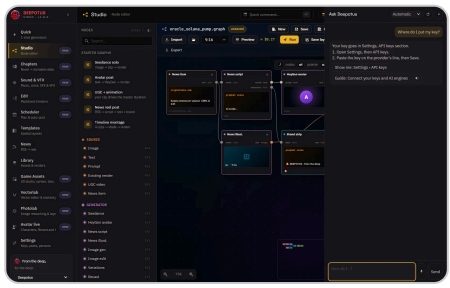

3. Show me, Guide

The drawer opens with F1; the answer fits in a few steps, with its buttons.

1. Press `F1`, or click **Ask** (the button marked "?") at the bottom right: the drawer opens.
2. Pick the **Engine** at the top of the drawer, or leave **Automatic**. Each engine says whether it is free or paid.
3. Type your question, for example "where do I put my key?", then `Enter`.
4. The answer fits in a few steps. Below it, depending on the question, buttons: **Show me**, **Go there**, **Guide**.

Get shown the button

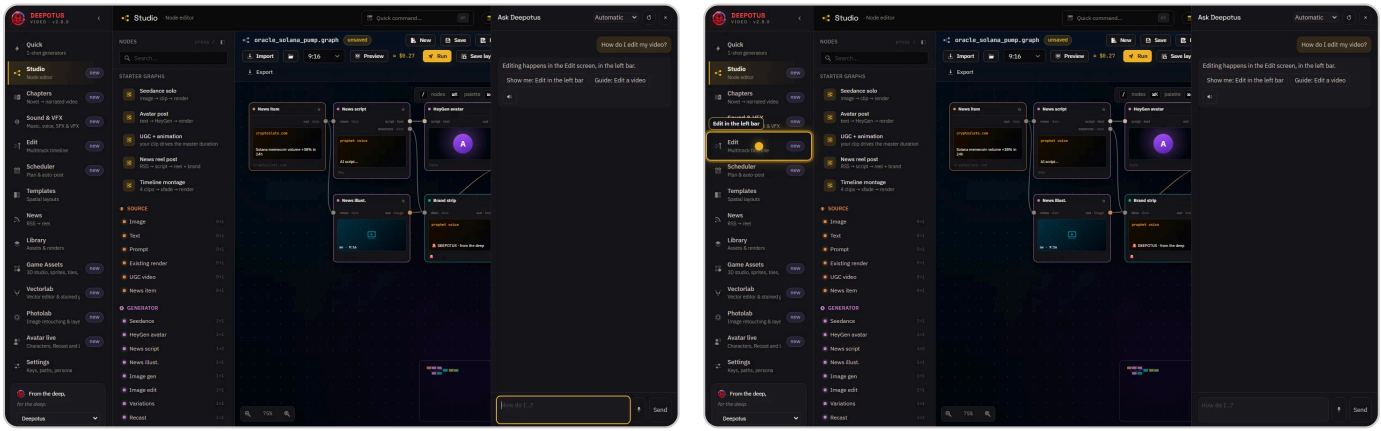

A glowing cursor glides to the button and circles it, with its name in a bubble.

1. Click **Show me**: if the button is on another screen, the app goes there first.
2. A glowing cursor glides to the button and circles it, with its name in a bubble. It disappears on your first click.
3. Click **Guide** to open the chapter about it, in a new tab.

When the assistant does not know the button

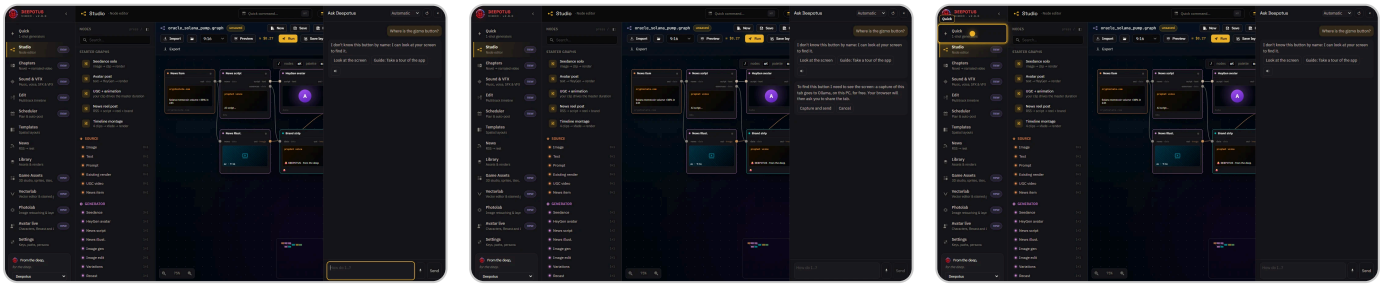

2. Cost shown first

The cost shows before any capture; the browser asks to share the tab (not filmed).

1. The answer offers **Look at the screen**. Click it only if you agree.
2. The drawer says which engine will see the capture and what it costs. Only engines that read images are offered.
3. Click **Capture and send**: your browser asks to share the tab; choose **this tab**. **Cancel** sends nothing.
4. The assistant answers and the cursor shows the place it found.

If the place found is under the drawer, the drawer closes to show it. If you have changed screens since the capture, the button no longer points: it sends you to the guide.

Dictate and listen

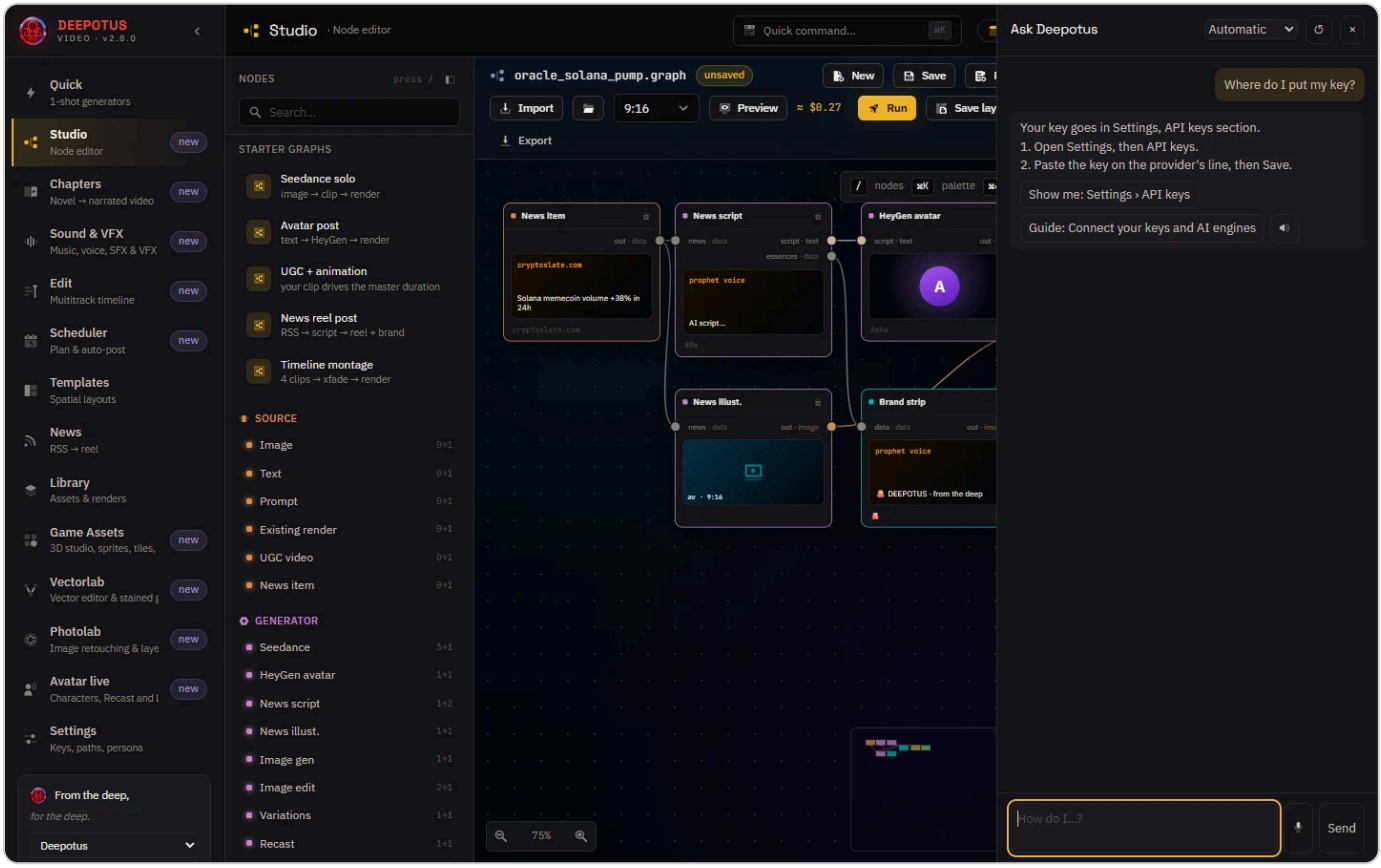

1. Dictate 2. Read aloud

The mic fills the input; the speaker reads the answer with the voice from Settings.

1. Click the mic next to **Send** and speak; click it again, or **Esc**, to stop. The text lands in the input: read it over before sending.
2. Under an answer, the speaker button reads it with the voice chosen in Settings. With Voicebox, it is free.

Good to know

A question costs a fraction of a cent with a paid engine, counted in the "help" cap of **Settings** › **Pricing & budget**. With Ollama on this PC, it is free. Looking at the screen costs a little more: the price shows first.

Dictation through the browser is free. When it goes through the server, the app shows the duration and the cost before sending anything.

The screen capture goes to the chosen engine. Close anything you do not want to show first.

Shortcuts

Key	Effect
`F1`	Open or close the drawer
`Enter`	Send the question
`Shift` + `Enter`	New line
`Esc`	Stop the dictation in progress; otherwise clear the pointer and close the drawer

GET STARTED

Install, pick a language, update

You will launch the app for the first time, pick its language, know what each use costs and keep it up to date. Ten minutes, free.

First launch

1. During installation, pick the language: **French** or **English**. It becomes the language of the whole app.
2. Double-click the **Deepotus Video Gen** desktop icon. The app starts without a black window, then your browser opens on the studio.
3. On the very first launch, a five-step welcome introduces the studio and lets you pick a persona, your generation providers and the networks you publish to.
4. Set at least one key, the **fal.ai** one: it makes pictures and video clips. The chapter Connect your keys shows how.

> The welcome can be replayed at any time from the quick command (`Ctrl` + `K`): **Replay onboarding**.

Switch language

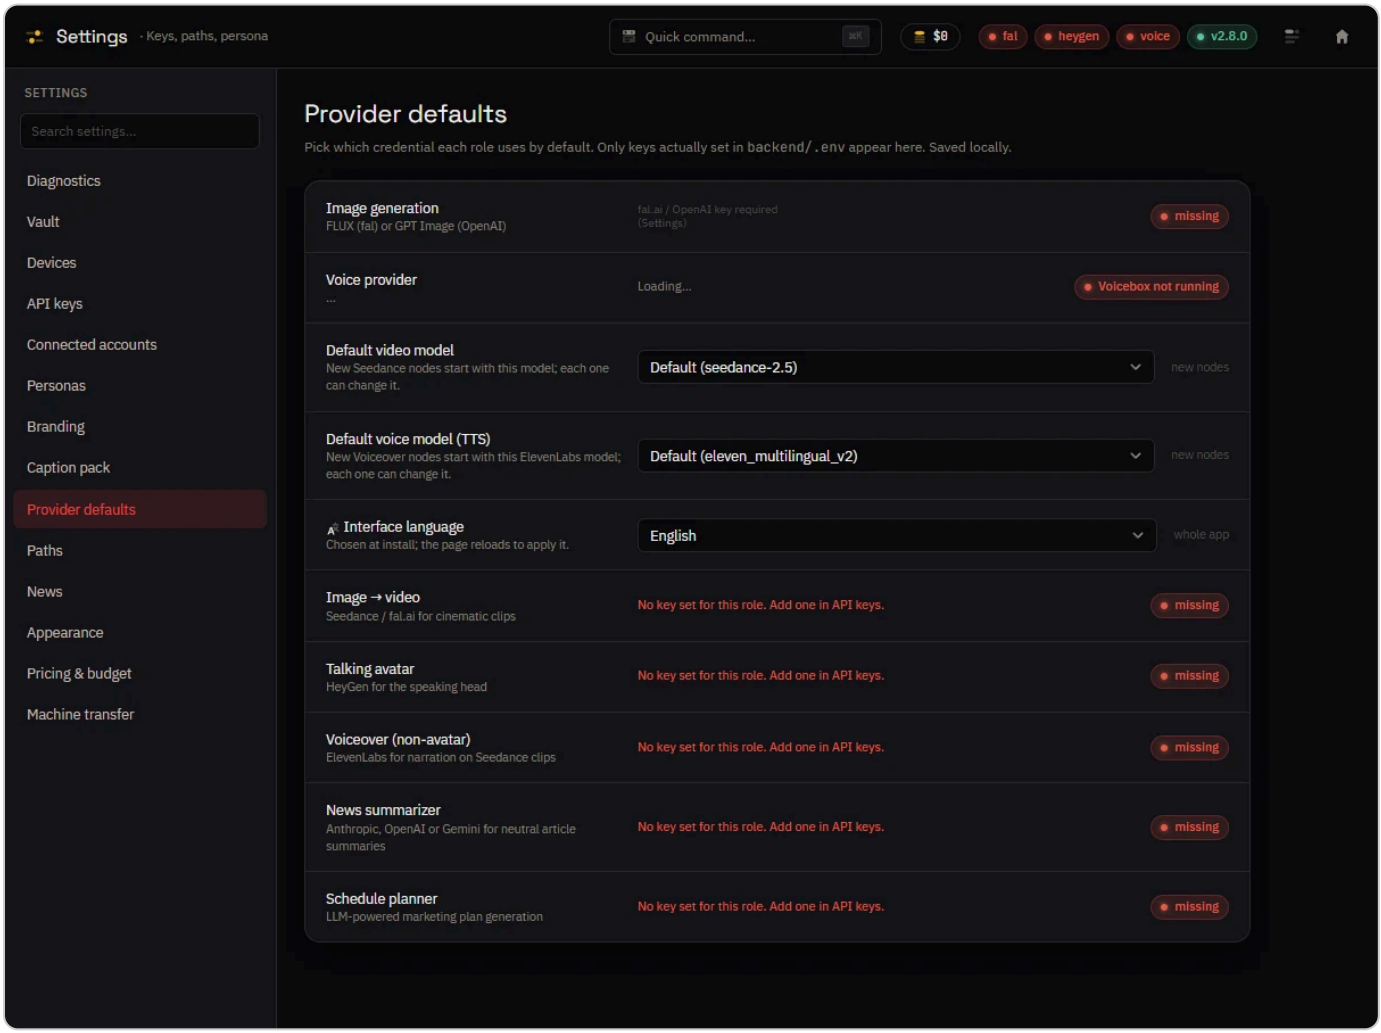

1. French or English, for the whole app

Settings › Provider defaults › Interface language: the page reloads in the chosen language.

1. Open **Settings** › **Provider defaults**.
2. In **Interface language**, pick **Français** or **English**. The page reloads in the chosen language.

Translated: the sidebar, Settings, the Library, Quick, News, the Scheduler, the Studio, Templates, Episodes, the editor, Sound & VFX, subtitles, machine transfer and every confirmation dialog. Lists sent by the server (effects, transitions, presets) also arrive in the chosen language. The labs (Vectorlab, Photolab, Card Forge…) are being translated one by one; meanwhile their known texts are translated on display. What you type yourself is never translated.

What each use costs

The app is free. Only AI generations cost money, billed by the provider of the key. Rough figures at default rates:

Action	Provider	Indicative cost
Create a picture	fal.ai (FLUX)	`≈ $0.003`
10 s video clip	fal.ai (Seedance 2.5, 720p)	`≈ $4.73`
Talking avatar, 1 min	HeyGen	`≈ $0.24`
Voiceover	ElevenLabs	per character, free tier available
News summary, post plan	Anthropic, OpenAI, Gemini or Ollama	a few cents, free with Ollama
Import your videos, sounds, pictures	none	free

Your real rates go into **Settings** › **Pricing & budget**; see Keep spending under control.

Keep the app up to date

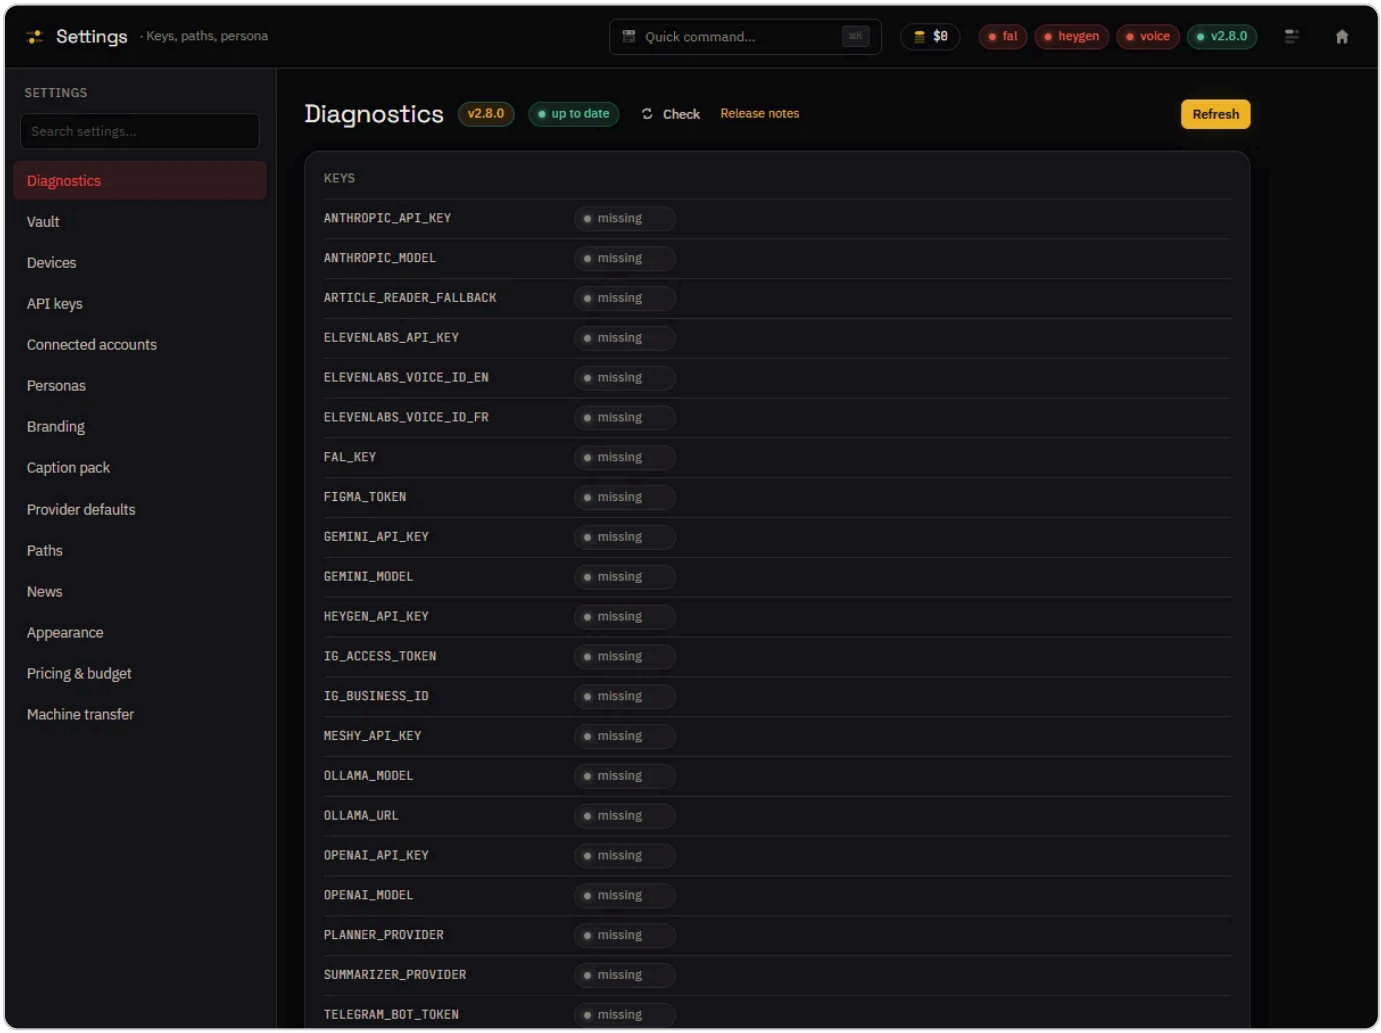

1. Your version is up to date 2. Check asks for the latest published version

The version block in Diagnostics: the installed version, its status, and a button to check right now.

1. Once a day, the app checks whether a new version is published. It never blocks anything.
2. When there is one, a card announces it: **Download** fetches the installer and shows where it is saved, **Later** hides the card for that version, **Notes** tells what changes.
3. To check right now: **Settings** › **Diagnostics**, **Check** button.
4. Close the app, run the downloaded installer: nothing is installed on its own.

Good to know

> Your keys, files and database live in the `%LOCALAPPDATA%\DeepotusVideoGenData` folder, apart from the program: they survive updates and reinstalls.

Start with the fal.ai key alone, then add avatars, voices and AI as you go. Set a monthly cap from day one: see Keep spending under control.

GET STARTED

Connect your keys and AI engines

You will set an AI provider's key, choose which engine does what, connect a free provider or a server on your network, and protect your keys with a password. Ten minutes, free; keys are created with each provider.

Before you start

An API key is a code the provider gives you on its website, in your account. It lets the app work on your behalf; the provider bills you. Only one is required to start: **fal.ai**.

Set a key

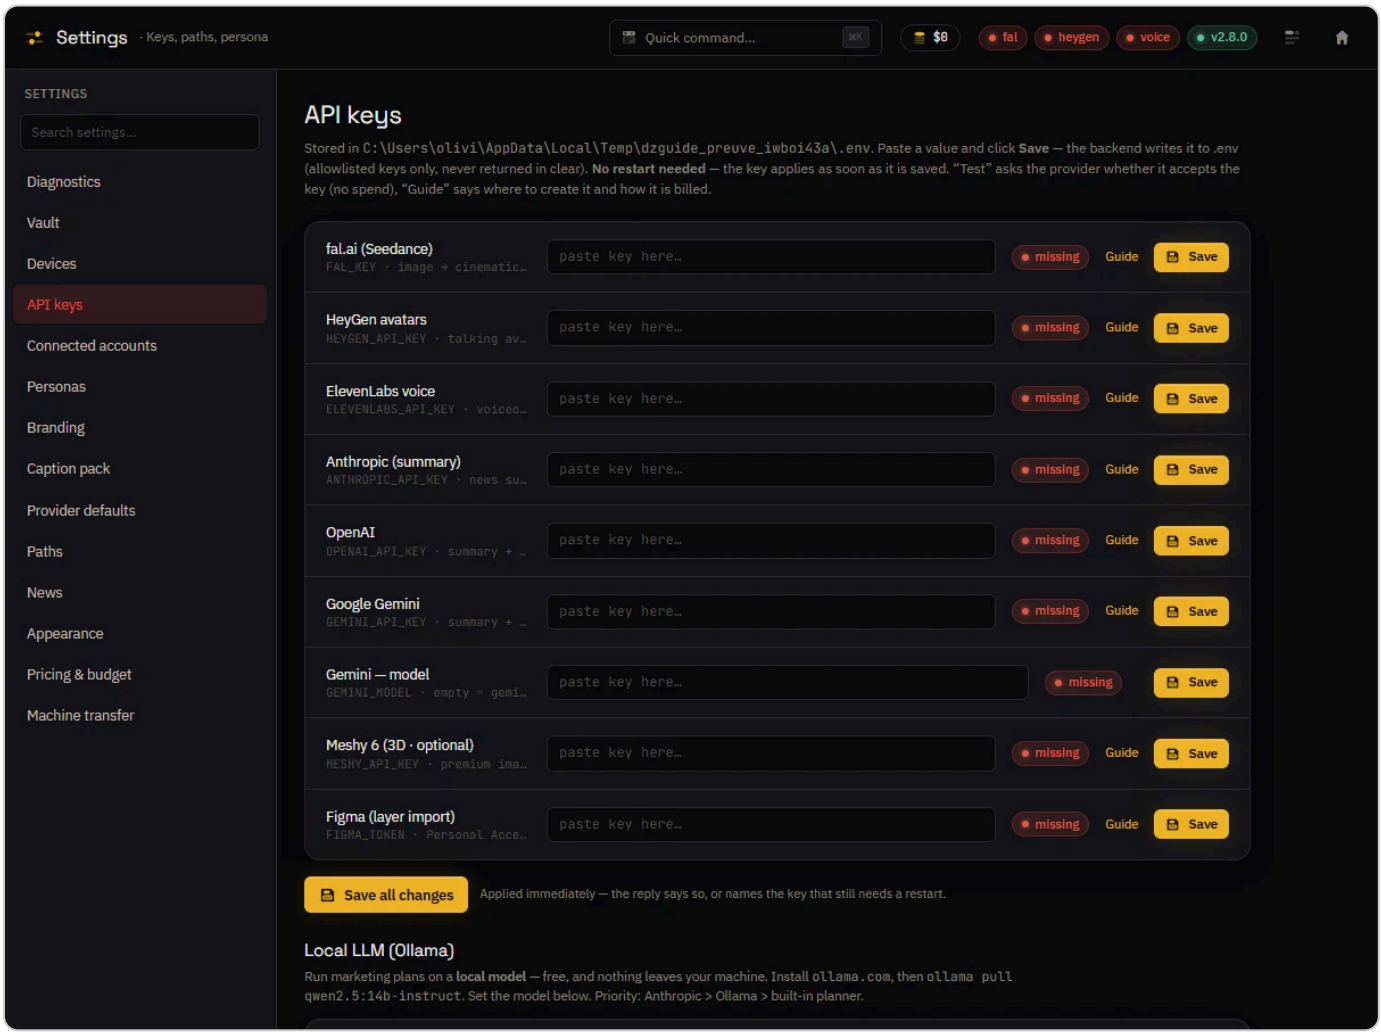

1. The key's status 2. Guide: where to create the key, and its price 3. Paste the key, Save: it works right away

Settings › API keys: each provider has its row, its status, its guide and its Save button.

1 Open **Settings** › **API keys**. Each provider has its row and its status: **missing** until a key is set. The fal, heygen and voice badges in the header lead here in one click.

2. Click **Guide**: it tells where to create the key with the provider and how it is billed.
3. Paste the key in the field, then **Save**. It works **right away**, no restart.
4. Click **Test** on the row: the provider says whether it accepts the key, without spending anything.

To set several keys at once, fill in the fields then **Save all changes**.

Which key unlocks what

Key	What it unlocks
fal.ai	Pictures (FLUX) and video clips (Seedance). The only key you need.
HeyGen	Talking avatars, alone or in a composition.
ElevenLabs	Voiceover and narration; two fields set the default French and English voices.
Anthropic, OpenAI or Gemini	News summaries, post plans, scripts and storyboards written by AI. One is enough; OpenAI also makes pictures (GPT Image).
Free provider (DeepSeek, OpenRouter, Mistral, Groq…)	The same texts, from a service that speaks the OpenAI API. Price to be measured.
Local servers (no key)	Ollama, LM Studio, llama.cpp or vLLM on your PC or your network: free, nothing leaves your place.
Meshy	Premium picture-to-3D (optional, Meshy subscription).
Decart	The Live mode of Avatar live, billed per second.
Figma	Import Figma layers into the Library and Templates.

Network keys (X, Telegram, YouTube, Instagram, TikTok) go into **Connected accounts**: see Connect your accounts.

Choose the engine for each job

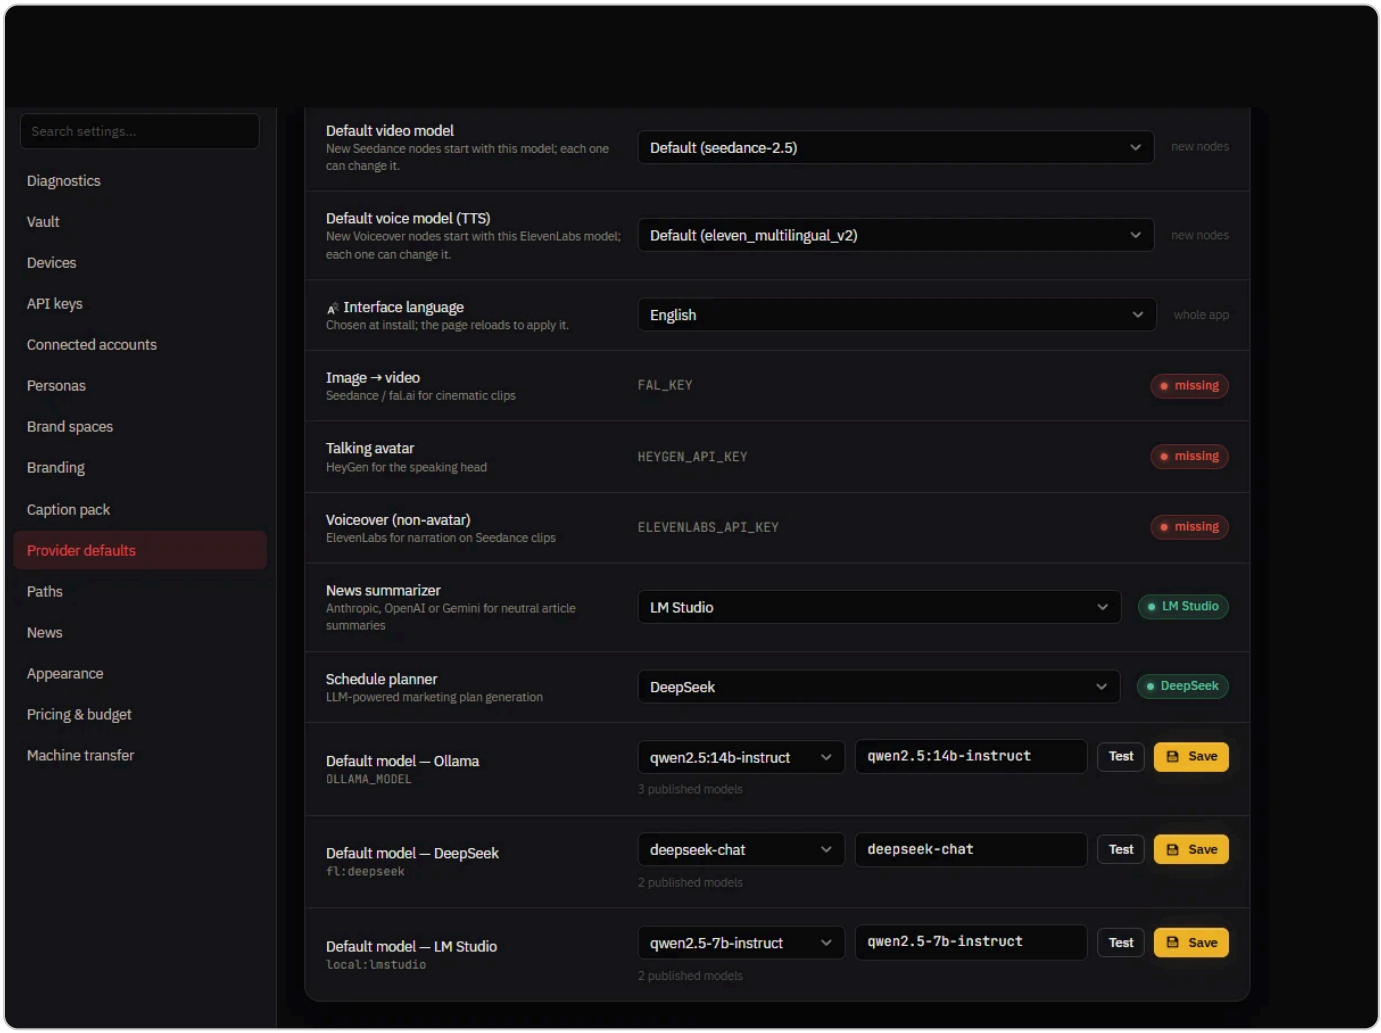

1. Each role picks its engine: built-in, free provider or local server 2. Each engine's default model, read from the provider

Settings › Provider defaults: the text engine of each role, then each engine's default model, with Test.

1. Open **Settings** › **Provider defaults**. The image to video, avatar and voiceover rows say which key serves them and whether it is set.
2. **News summarizer** and **Schedule planner** each pick a text engine: a built-in one (Anthropic, OpenAI, Gemini, Ollama), a free provider or a local server. "Automatic" takes the first engine that is ready.
3. Further down, each configured engine has its **Default model** row: pick from the list read from the provider, or type a name. **Test** checks that the model is published, without spending anything; **Save** applies it at once.

The choice is written to the app's settings file: it applies to every screen. A model name that looks like a key pasted by mistake is refused and flagged in red.

Default video model

The model new Seedance nodes start with in the Studio. Each node can change it.

Default voice model

The ElevenLabs model of new voiceover nodes.

Add a free provider

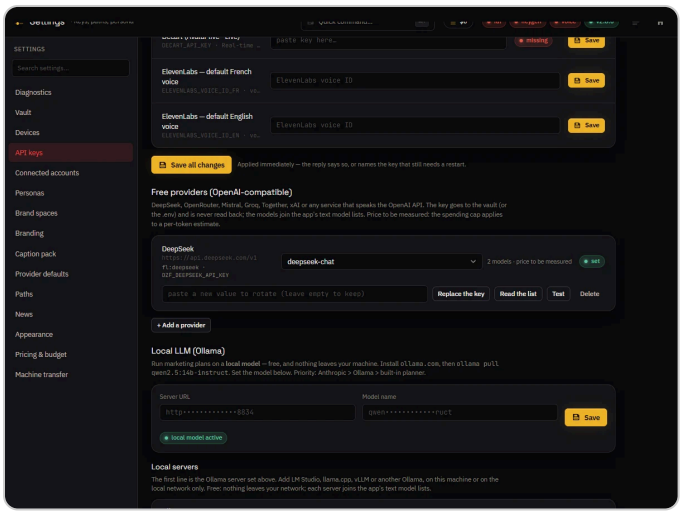

1. The fl:deepseek engine: key set, price to be measured

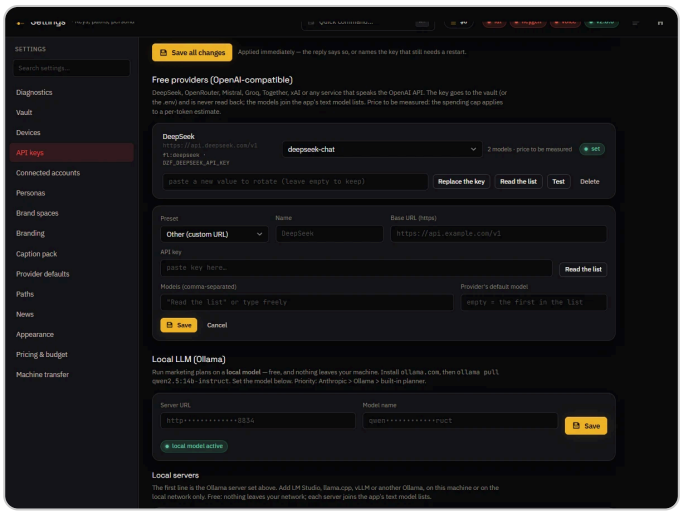

2. Preset, name, https URL and key 3. Read the list: the models the provider publishes

Settings › API keys › Free providers: DeepSeek, OpenRouter, Mistral, Groq… The key goes to the vault, the models join every list.

1. In **API keys**, scroll down to **Free providers (OpenAI-compatible)**, then **+ Add a provider**.
2. Pick a preset (DeepSeek, OpenRouter, Mistral, Groq, Together, xAI) or type a name and an https URL. Paste the key.
3. **Read the list** asks the provider for its models; keep the ones you need and the default model, then **Save**.
4. On the row, **Test** checks the key, **Replace the key** sets a new one, **Delete** asks for a second click (**Confirm deletion**) and erases the key.

The provider becomes the `fl:<name>` engine: its models join every text model list (the Vectorlab AI dialog, the roles above). The key goes to the vault and is never shown again.

Connect several local servers

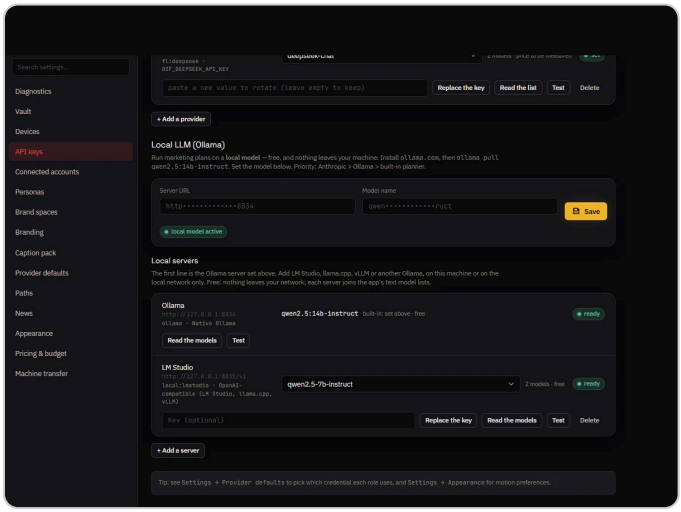

1. The first row: your Ollama, set just above 2. An added server: local:lmstudio engine, free

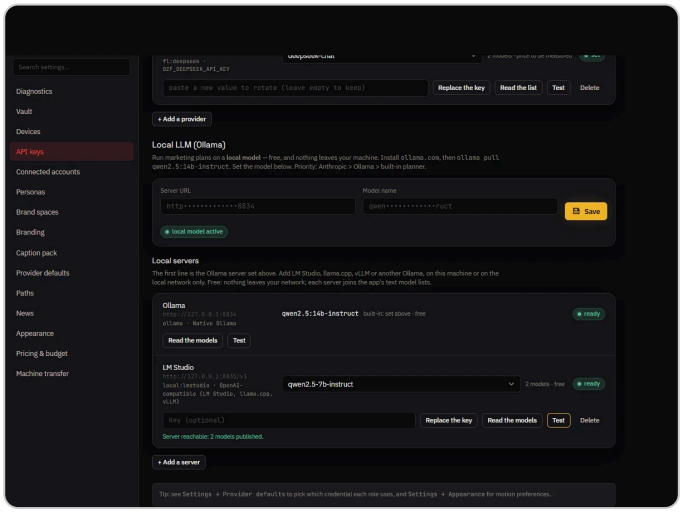

3. Test: the server answers and publishes its models

Settings › API keys › Local servers: Ollama, LM Studio, llama.cpp or vLLM, on this PC or the local network. Free.

1. Under **Local LLM (Ollama)**, the **Local servers** list starts with your Ollama, set just above.
2. **+ Add a server**: pick a preset (Ollama, LM Studio, llama.cpp, vLLM); the URL must be on this PC or on your local network.
3. **Read the models** fills the list, pick the default model, then **Save**.
4. **Test** says whether the server answers and how many models it publishes.

Each server becomes the `local:<name>` engine, free, in the same lists as free providers.

Protect your keys in the vault

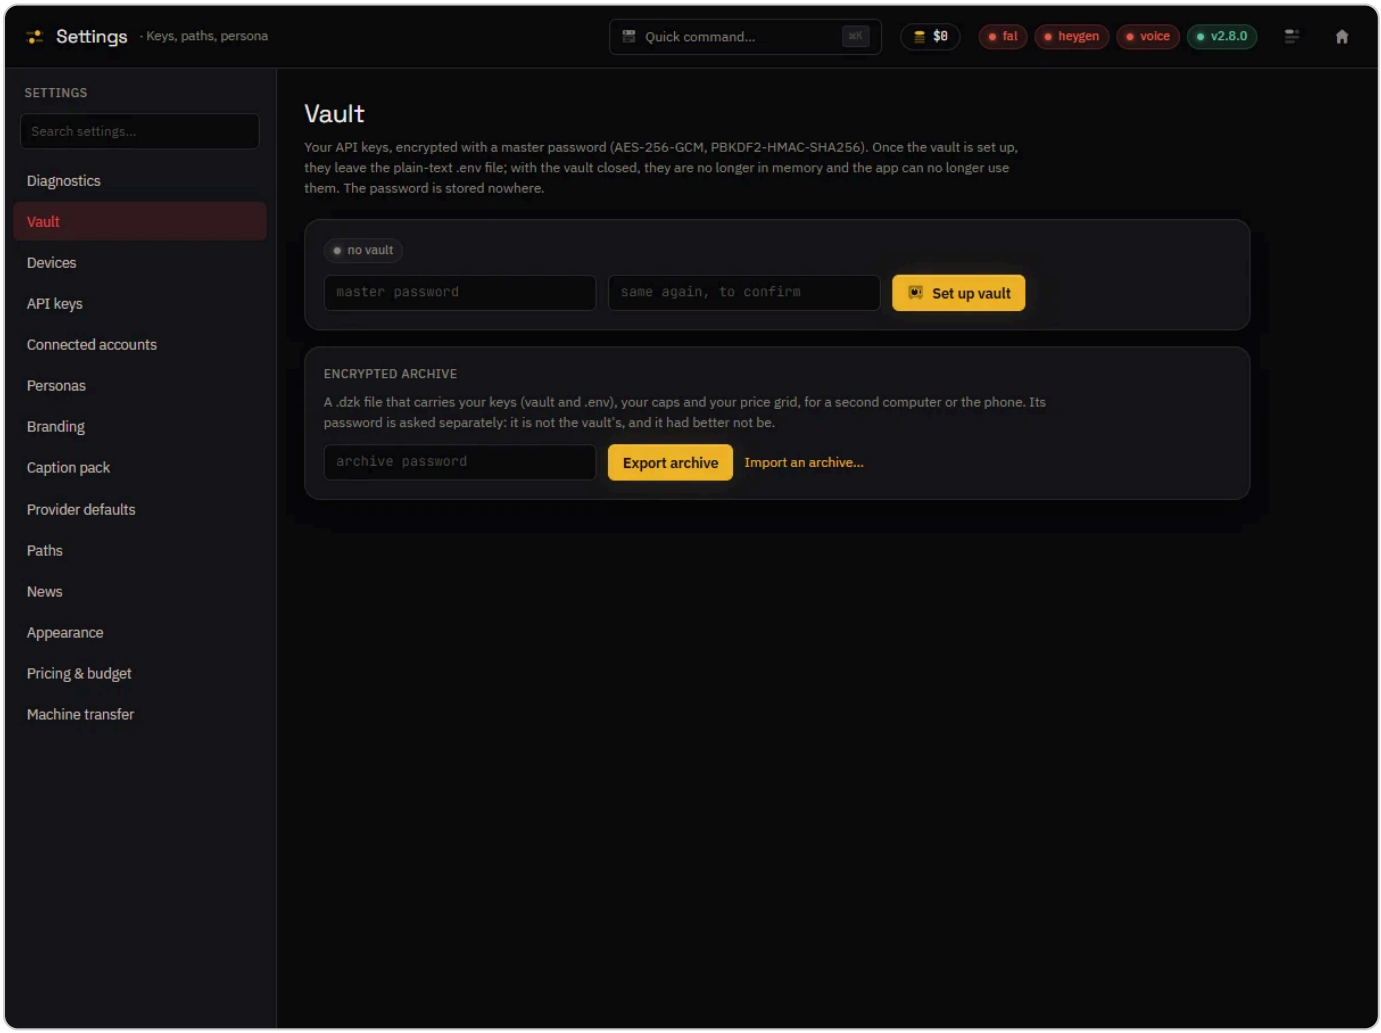

1. A master password encrypts all your keys 2. The .dzk archive carries them, with its own password

Settings › Vault: your keys encrypted with a password, and the archive to carry them to another computer.

1. Open **Settings** › **Vault**.
2. Choose a master password (at least 8 characters), type it twice, then **Set up vault**. Your keys are encrypted and leave the plain-text file.
3. **Close** removes the keys from memory: nothing paid can go out. **Open** makes them active again with the password.
4. **Remember on this PC** opens the vault on its own at launch, for your Windows session; **Stop remembering** asks for the password again.

On each row of **API keys**, a badge tells where the key lives: in the file, in the vault, or vault closed (one click takes you there).

Good to know

The vault password is stored nowhere: if lost, it cannot be recovered. Write it down, or keep an encrypted archive.

Export archive creates a `.dzk` file that carries your keys, caps and rates to another PC or your phone, protected by its own password. **Import an archive…** takes them back.

Settings › **Diagnostics** shows the status of all keys at once (never their value) and the providers' balances. Each row carries **Set up**, which opens the section where the key is set.

GET STARTED

Choose your AI models and their route

You will know how to read a list of AI models: which key is used, what one run costs, and what to do when a key is missing. Five minutes, free; the same lists are used across the whole app.

Two routes for the same model

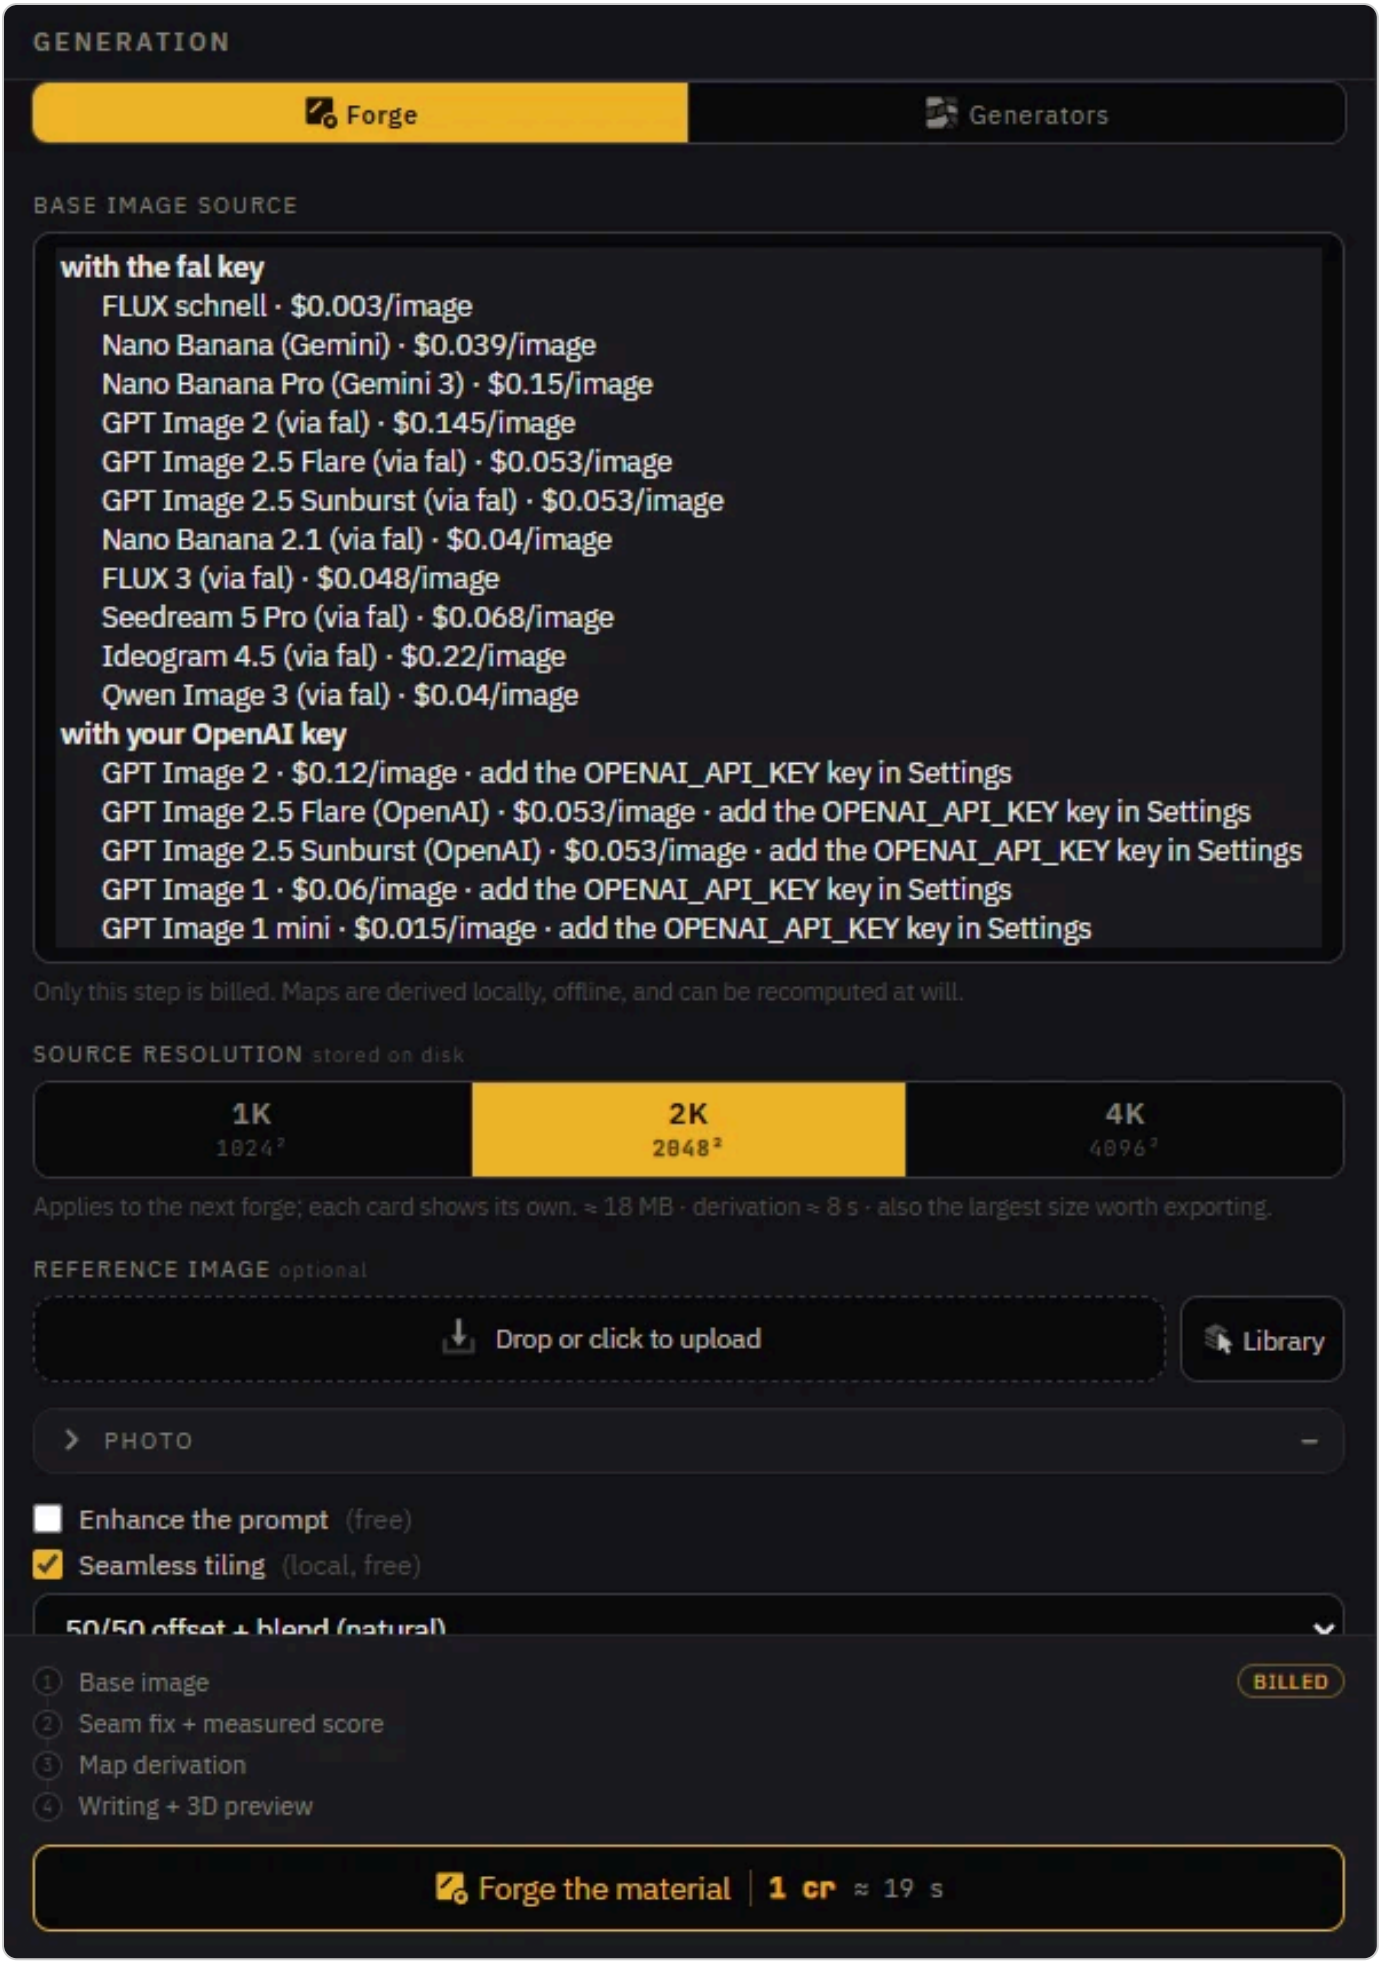

1. With the fal key: one key for all these models 2. Each option shows its price 3. Missing key: the option can still be chosen

One group per route (the fal key, or your own key), the price on each option, the missing key named.

Wherever you generate (images, videos, sounds, background removal), the model list is sorted into groups, one per **route**. You keep all your keys: each list offers every route, and you choose model by model.

With the fal key
A single fal.ai key opens most models: FLUX, Nano Banana, GPT Image, Seedream, Veo, Kling… The label says "via fal" when the same model also exists direct.

With your key
The model paid directly to its provider: OpenAI, Gemini, ElevenLabs, Meshy, HeyGen. Handy if you already have an account and credits there.

Local, free
Some tools run on your PC, such as rembg background removal or plain background: no key, no spending.

Read an option's price

Each option shows its unit price after the name: **$0.04/image**, **$0.10/s** of video, **$0.0018/s** of sound, **/30 images** for a video cutout. An option marked **price to be measured** is paid, but its rate isn't known yet: it is never shown as free.

Before a run, the screen shows the **total price** worked out by the server for what you ask: an image and its background removal, a 6 s video… Your spending caps check this very amount, see Keep spending under control.

Text engines

The text engine lists (the Vectorlab AI dialog, the roles of **Settings** › **Provider defaults**) gather three kinds of engines. The tooltip of the Vectorlab **Generate** button says, for the chosen engine, what the call costs.

Built-in
Anthropic, OpenAI, Gemini: paid on your key, a few hundredths of a cent per illustration or summary.

Free providers
The `fl:…` engines (DeepSeek, OpenRouter…): **price to be measured**. The cap applies to a per-token estimate, at the dearest built-in rate.

Local servers
`ollama` and the `local:…` engines run on your PC or your network: free, nothing is charged.

An "add the key" option

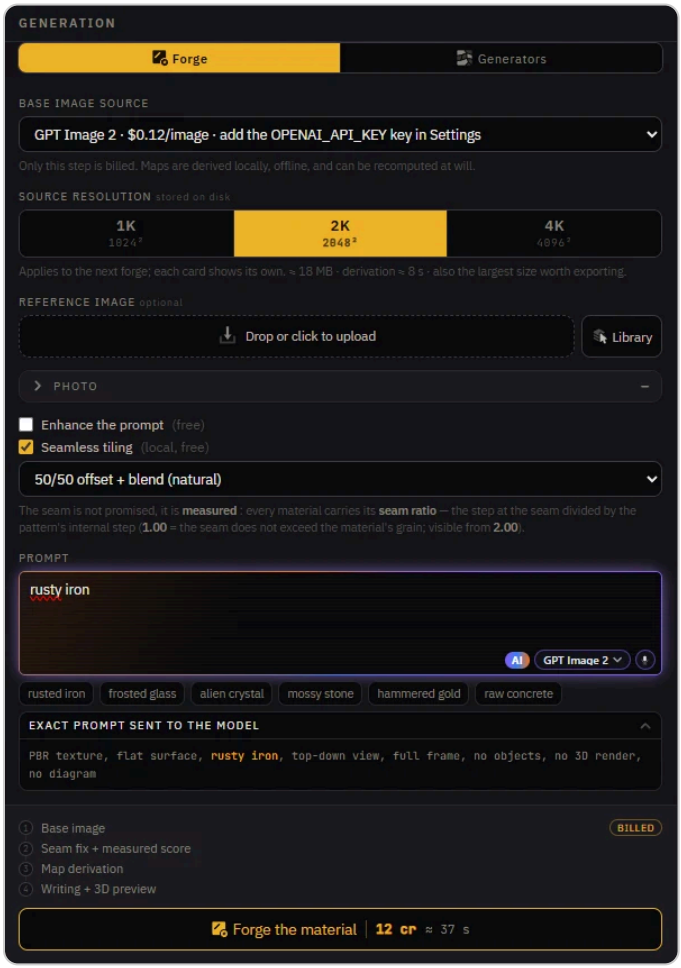

1. GPT Image 2 direct: OpenAI key to add

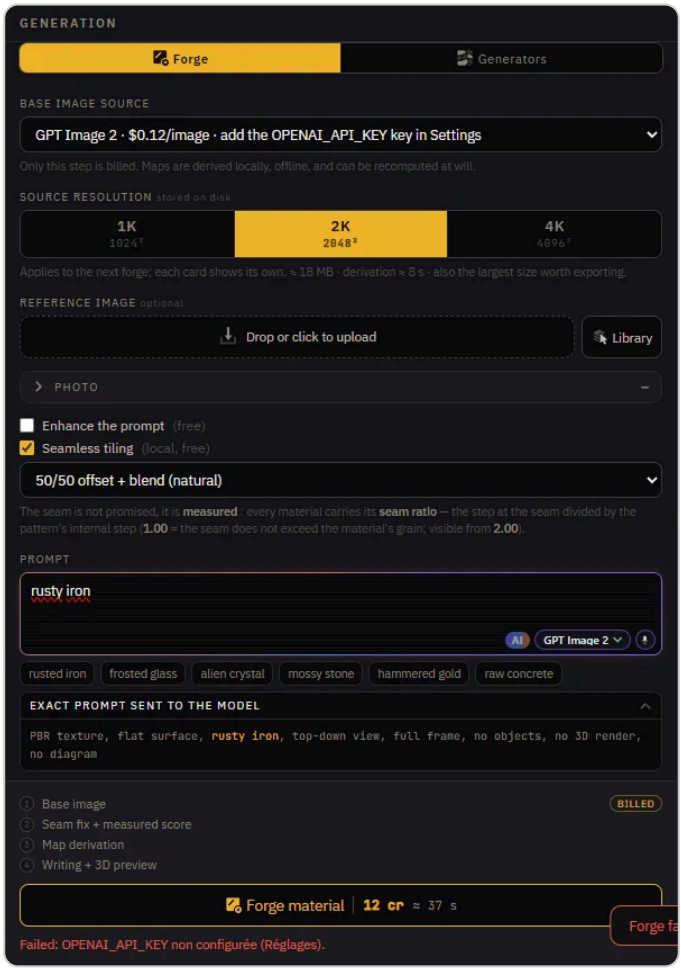

3. Refused at launch: nothing is charged

Once chosen, the option without a key is refused at launch, before any spending: add the key in Settings.

1. An option whose key is missing stays visible **and can be chosen**; it carries the note "add the … key in Settings".
2. If you run it, the run is refused at once, before any spending, naming the missing key.
3. Add the key in **Settings** › **API keys** (see Connect your keys): it works right away, no restart. Or pick the same option with the fal key.

Good to know

> Model names (Flux 3, Veo 3.1 Fast, Sonilo…) are proper names: they stay the same in French and English.

A run whose price exceeds your cap is refused before the call. The most expensive actions run in two clicks: the first shows the estimate, the second runs it.

GET STARTED

Keep spending under control

You will know what you spend, set a cap per month and per engine, and decide yourself when a render may go over it. Five minutes, free.

Where to see what you spend

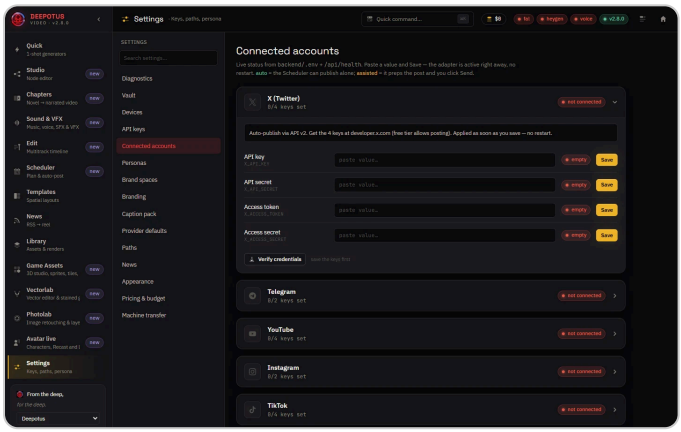

1. One click on the month's counter

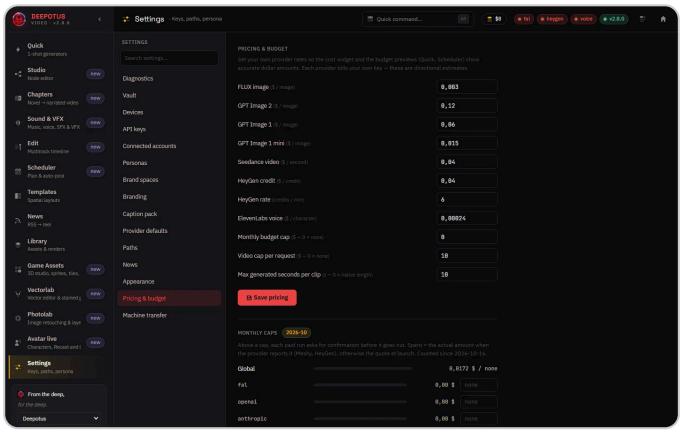

2. Pricing & budget opens, even if Settings was already open

The dollar counter in the header: one click opens Settings › Pricing & budget, from any screen.

The dollar badge in the top bar shows the month's total. One click opens **Settings** › **Pricing & budget**, where everything is set, from any screen and even when Settings is already open on another section. The fal, heygen and voice badges open **API keys** instead.

Set your rates and caps

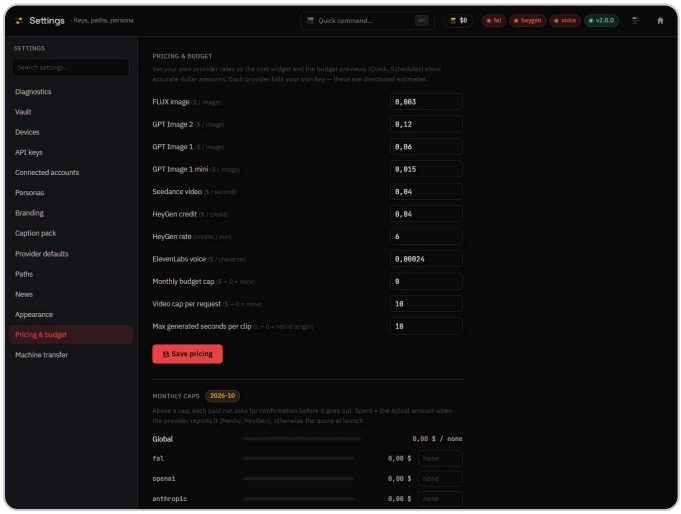

1. A cap for each video generation

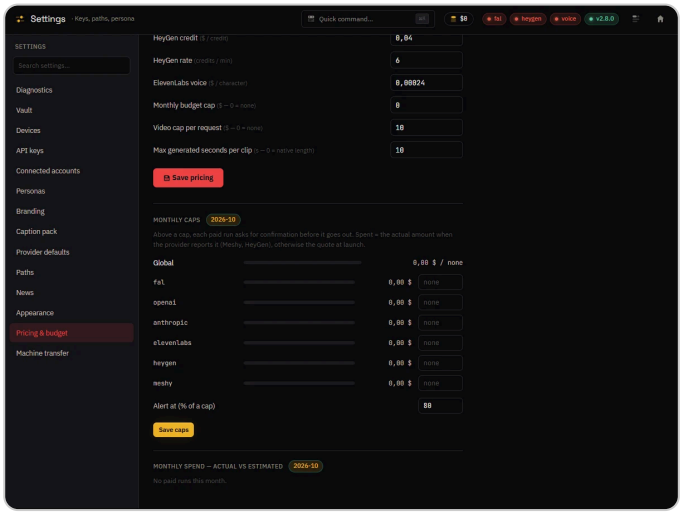

2. Above a cap, the app asks you before spending

Settings › Pricing & budget: your rates, then one cap per engine and the bar of what you spent this month.

1. At the top of **Pricing & budget**, enter your providers' rates (price of a picture, of a second of video, of a voice character…). The costs shown in Quick, the Studio and the Scheduler become accurate. **Save pricing**.
2. **Video cap per request** limits the cost of a single video generation; **Max generated seconds per clip** limits its length.
3. Further down, **Monthly caps**: one overall cap and one per engine (fal, OpenAI, Anthropic, ElevenLabs, HeyGen, Meshy), in dollars. A bar shows what was spent.
4. Set **Alert at (% of a cap)**: a banner warns you once a month when a cap reaches that level. **Save caps**.

When a render would go over a cap

Before leaving, every paid generation is checked against your caps. If it goes over, a “Monthly spending cap” dialog opens, wherever you are in the app:

1. “Tirer quand même” (run anyway) sends exactly the same request, once.
2. “Annuler” (cancel) sends nothing and tells you why.

Actual versus estimated

Below the caps, the **Monthly spend — actual vs estimated** table gives one row per engine and per screen: the number of runs, what was estimated, what was actually billed, and the difference.

Actual
The provider gave the exact price (Meshy, HeyGen): that is what counts.

Partly actual
Some runs have their exact price, the rest is still estimated.

Estimated
The provider gives no price: the app counts the quote made at launch.

Good to know

Spending made from your phone goes into the same record and the same caps, counted once. A free provider is counted on a per-token estimate; a local server costs nothing.

The Studio sends the cost it displays as the request's cap: a render cannot cost more than what you saw before clicking.

GET STARTED

Settings, personas and branding

You will find any setting, give your content a voice with personas and dress the app in your colours with brand kits. Ten minutes, free.

Find a setting

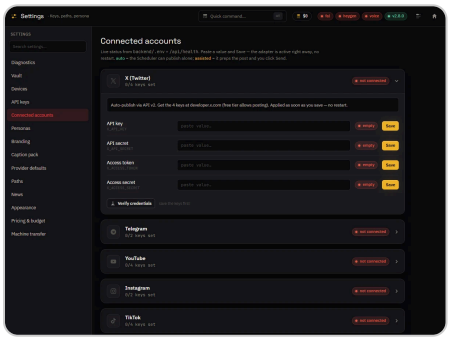

1. Type what you are looking for

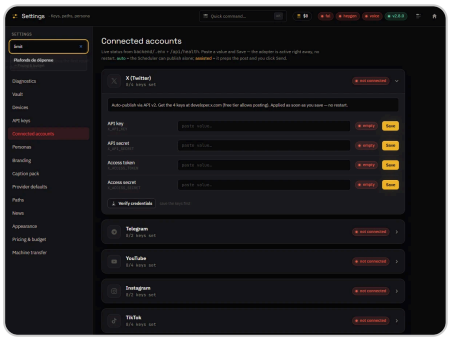

2. The matching sections show up

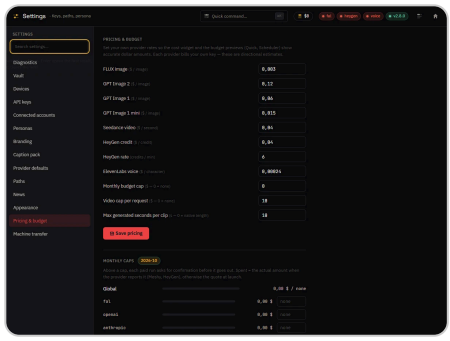

3. Enter opens the first one

Settings search finds a setting by its name and opens the right section directly.

1. Open **Settings** in the rail.
2. Type a word in **Search settings…**: accents and capitals do not matter.
3. Enter opens the first result, Esc clears the field.

The Settings sections

Section	What you do there
Diagnostics	Version, disk space used, latest alerts, key status and provider balances.
Vault	Encrypt your keys and carry them (chapter Connect your keys).
Devices	Pair a phone (chapter Drive it from your phone).
API keys	Set and test the AI providers' keys.
Connected accounts	Link X, Telegram, YouTube, Instagram and TikTok (chapter Connect your accounts).
Personas	The tone of your content.
Branding	Name, logo, colours: your brand kits.
Caption pack	One-click brand labels under the Scheduler's caption editor.
Provider defaults	The engine of each role, default models, language.
Paths	Where your pictures and renders are stored (button to copy the path).
News	Settings of the News screen.
Appearance	Reduced motion and visual effects.
Pricing & budget	Your rates and caps (chapter Keep spending under control).
Machine transfer	Take your work to another PC (chapter Move to another computer).

Give a voice with personas

A persona is a personality: its tone, its audience, its way of speaking. The active persona sets the News scriptwriter, the tone of the prompt generator and the default voiceover.

1. Open **Settings** › **Personas**.
2. Built-in personas can be read but not edited: duplicate one, or click **New persona**.
3. Click a persona to make it **active**.

Dress the app with a brand kit

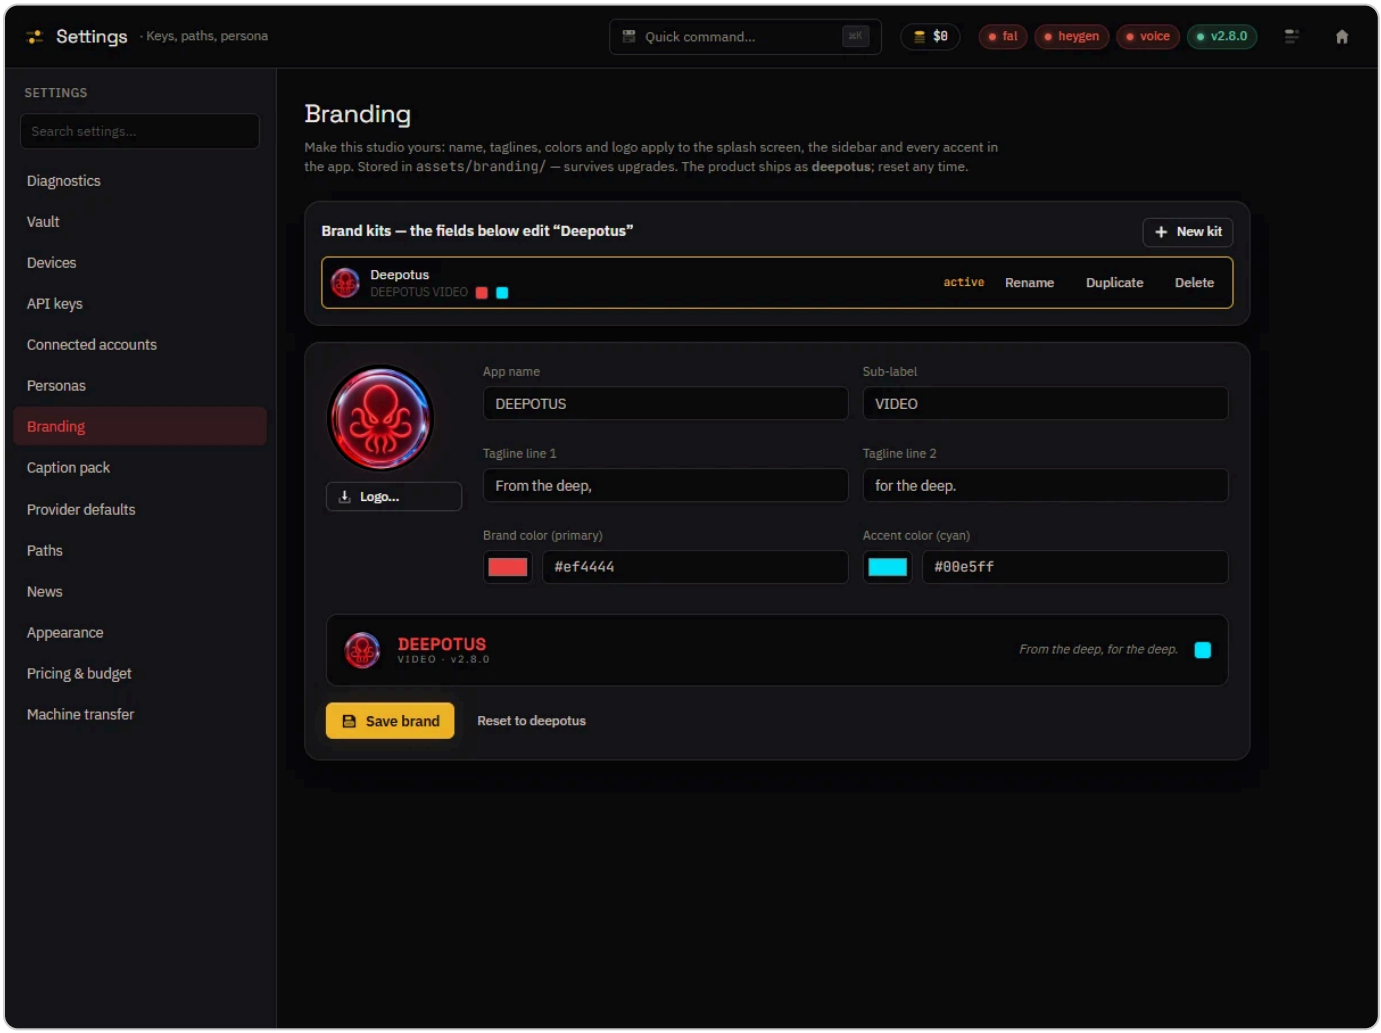

1. Create a kit 2. The active kit dresses the whole app

Settings › Branding: several brand kits, a single active one, and its fields right below.

1. Open **Settings** › **Branding**. The kit list shows which one is **active**.
2. **New kit** creates one; **Activate** makes it the active kit and refreshes the app.
3. The fields below edit the active kit: **Logo...**, app name, sub-label, two tagline lines, primary and accent colour. **Save brand**.
4. **Rename**, **Duplicate** and **Delete** manage each kit. The active kit and the last kit cannot be deleted: the greyed-out button says why.

The active kit dresses the splash screen, the sidebar, the app's accents and Template renders that use the brand. **Reset to deepotus** restores the original look.

Good to know

A Template render keeps the kit that was active when you sent it, even if you switch kits while it computes.

Settings › **Appearance** › **Reduced motion** turns off decorative animations; it also follows the Windows setting.

GET STARTED

Work for several brands

You will run several brands in the same installation (Deepotus, a new brand, a client) without ever mixing them up: each has its own persona, identity, Quick templates, publishing accounts, Library and spending. Free; only the assistant's AI proposals are paid, and their price is shown first.

Create a brand's space

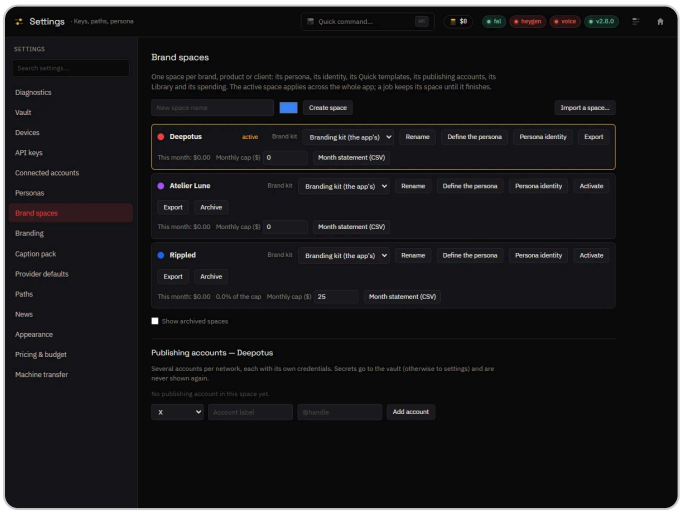

1. A name, a colour: Create space 2. Each brand has its row and its persona

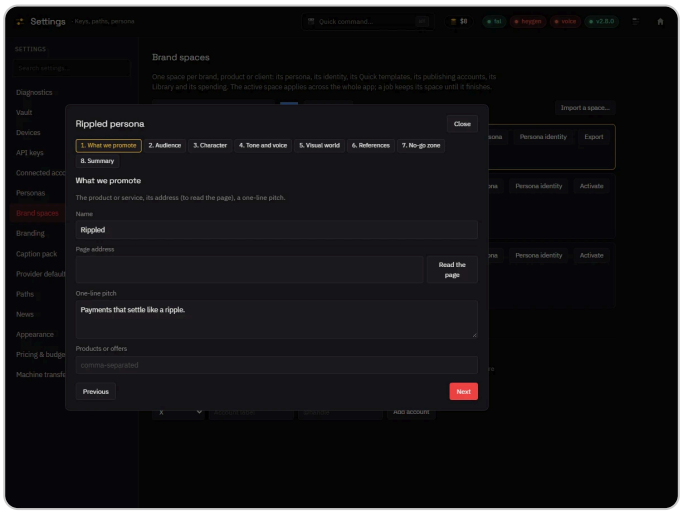

3. The eight-step assistant: nothing is written until you validate

Settings › Brand spaces: a name and a colour, then the persona in eight steps.

1. Open **Settings** › **Brand spaces**. Your installation already has one space: **Deepotus**, which keeps everything you had before.
2. Type the brand's name, pick its colour, then **Create space**. The new space is empty and not active yet.
3. On its row, **Define the persona** opens an eight-step assistant: what you promote, the audience, the character, tone and voice, the visual world, references, what is off-limits, then the synthesis. Nothing is written until you validate.
4. **Persona identity** stores the logo, the mascot, product photos, the palette, fonts and a written style. The prompt engine uses it instead of the Deepotus world.
5. Click **Activate**: the whole app switches to this space.

Switch from one brand to another

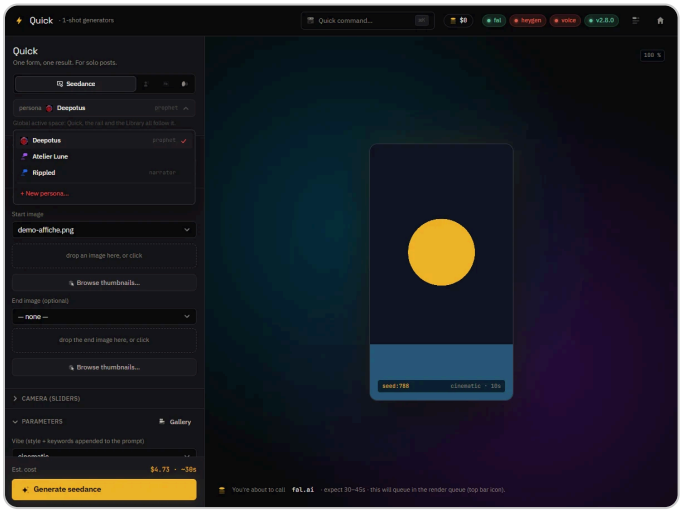

1. One persona per brand

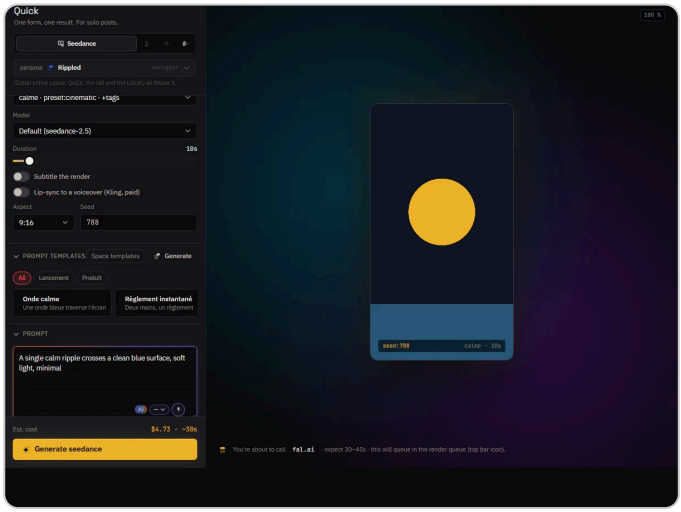

2. Rippled becomes the active space 3. Its own prompt templates

Quick's persona menu changes the active space: templates, vibes and voice follow.

Two places, one setting: the **Active space** menu under the brand card at the bottom of the rail, and the **persona** menu at the top of `Quick`. Changing one changes the other, and the Library too. In Quick, the prompt templates, the vibes and the default voice-over become those of the chosen brand.

1. In Quick, open the persona menu at the top of the screen and choose the brand.
2. The **Prompt templates** reload with the space's templates. A new space has none yet: `Space templates` proposes some from the persona, without AI and for free; tick the ones to keep.
3. `+ New persona...`, at the bottom of the same menu, opens Settings › Brand spaces to create one.

Publish on the brand's accounts

Accounts per space

Under the list of spaces, **Publishing accounts** manages the active space's accounts: several X or Telegram accounts if needed, each with its own credentials. `Test` checks an account; `Connect` starts the YouTube or TikTok consent.

The Scheduler follows the space

It shows the active space's posts. In a post, **Publishing accounts** only offers this brand's accounts.

Never on the wrong account

A client's post cannot go out on a Deepotus account: the app refuses to send it.

Sort files and bill

Library per space

The header menu shows **This space** or `All spaces`. Each render is filed in the space where it was started, even if you switch spaces while it runs.

Spending and cap

Each space's row shows what it cost this month. A **monthly cap** of its own comes on top of the global cap; 0 means no cap.

Statement for the client

`Month statement (CSV)` downloads the space's spending, ready to attach to an invoice.

Hand over or archive a client's space

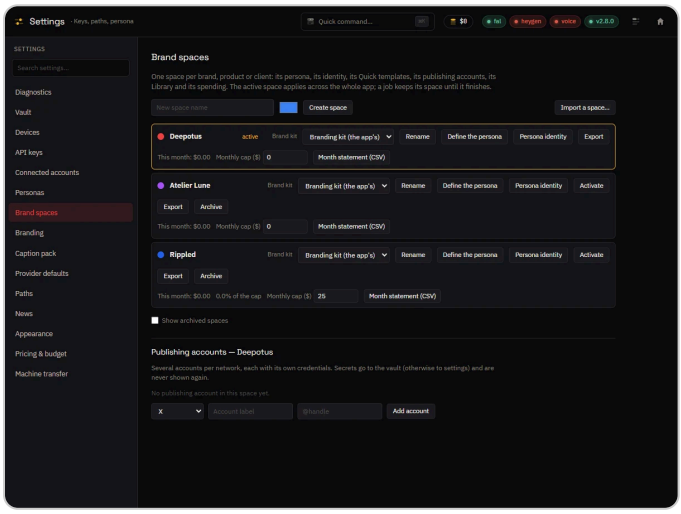

1. Export: a .zip, no secret

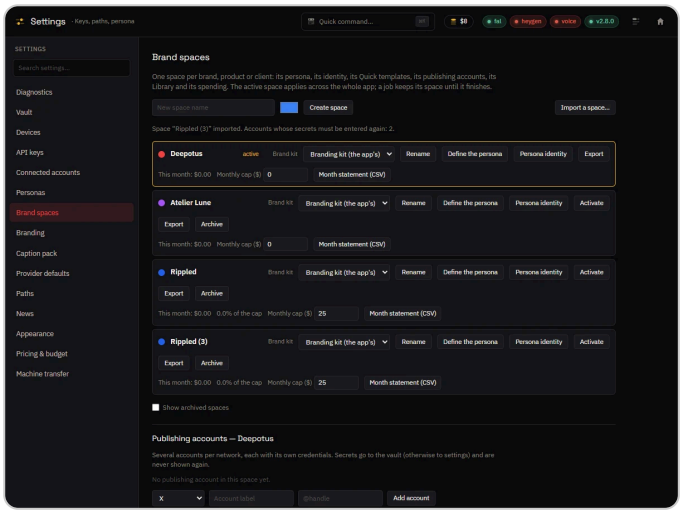

2. Import: always a new space 3. Account credentials are entered again

Export downloads a package without secrets; Import a space... creates a new one elsewhere.

1. `Export`, on the space's row, downloads a .zip package: persona, identity and its pictures, Quick templates, brand kit and the list of accounts. **No secret goes in**: tokens and passwords stay on your PC.
2. On another installation, `Import a space...` creates a new space from the package, never replacing one. The message says how many accounts are waiting for their credentials: enter them again in **Publishing accounts**.
3. When a contract ends, `Archive` puts the space away: it leaves the menus, and its scheduled posts are paused. `Unarchive`, with `Show archived spaces`, resumes them.

Good to know

Only one space is active at a time, for the whole app. A render, a post or a task already started keeps its space until the end.

The app's logo and name never change with the space: they are set in **Settings** › **Branding**. The kit chosen for a space is only used for that brand's content.

The active space cannot be archived: activate another one first. There is always at least one space.

GET STARTED

Move to another computer

You will take all your work to another computer, check that nothing got damaged on the way, and get your keys back there. Fifteen minutes, free; the app must be installed on both machines.

What travels

The transfer takes everything the app created: the Library and the origin of every file, renders, the bible, communication plans, vector documents, card series and 3D models. Your API keys **never** go into the package.

Export from the old PC

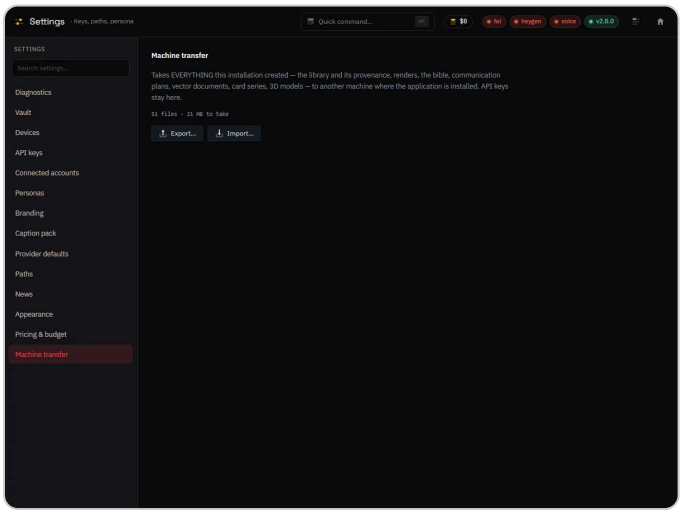

1. What goes, and how much it weighs

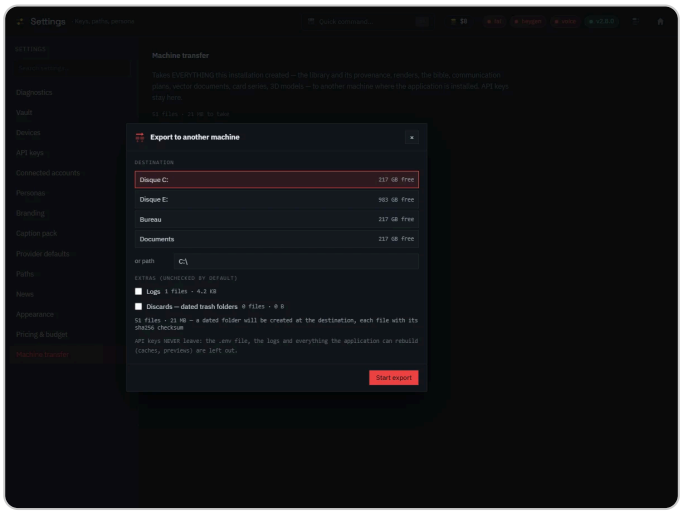

2. Extras to tick: logs and discards 3. A dated folder is written at the destination

Settings › Machine transfer: a dated package, each file with its checksum; keys stay on this PC.

1. Open **Settings** › **Machine transfer**: the number of files and the size to take are shown.
2. Click **Export...**, then pick the destination: a drive, a USB stick, or a path.
3. Tick the extras if you want them: **Logs** and **Discards — dated trash folders**. They are unchecked by default, with their size.
4. **Start export**: a dated folder is written at the destination, each file with its checksum. Do not close the app while it writes.
5. Click **Check integrity**: every written file is read back and compared with its checksum.

Import on the new PC

1. Install the app, then open **Settings** › **Machine transfer** and click **Import…**.
2. Give the package folder, then **Check**: the app reads the package without writing anything and tells where it comes from (date, machine, version).
3. **Start import**. Import **adds**: nothing this PC already has is overwritten. A damaged file is left out and named.

Take your keys too

1. On the old PC: **Settings** › **Vault** › **Export archive**, with a separate password.
2. On the new PC: **Import an archive…** and that password. Your keys, caps and rates come back.

Good to know

> Paths from the old machine are rewritten for the new one: your projects find their files without any setting.

> Stopping an export leaves an incomplete package. If it happens, run the export again, then **Check integrity** before importing.

GET STARTED

Drive it from your phone

You will link your phone to the PC, carry your posts and chapters in your pocket, and cut off a lost phone. Fifteen minutes; you need the Deepotus companion app on the phone and the same Wi-Fi.

Before you start

By default, the app only listens on this PC: no other device can reach it. To open it to your home Wi-Fi, replace `HOST=127.0.0.1` with `HOST=0.0.0.0` in the `.env` file of the data folder, then restart the app. Windows will ask to allow Python through the firewall.

Pair the phone

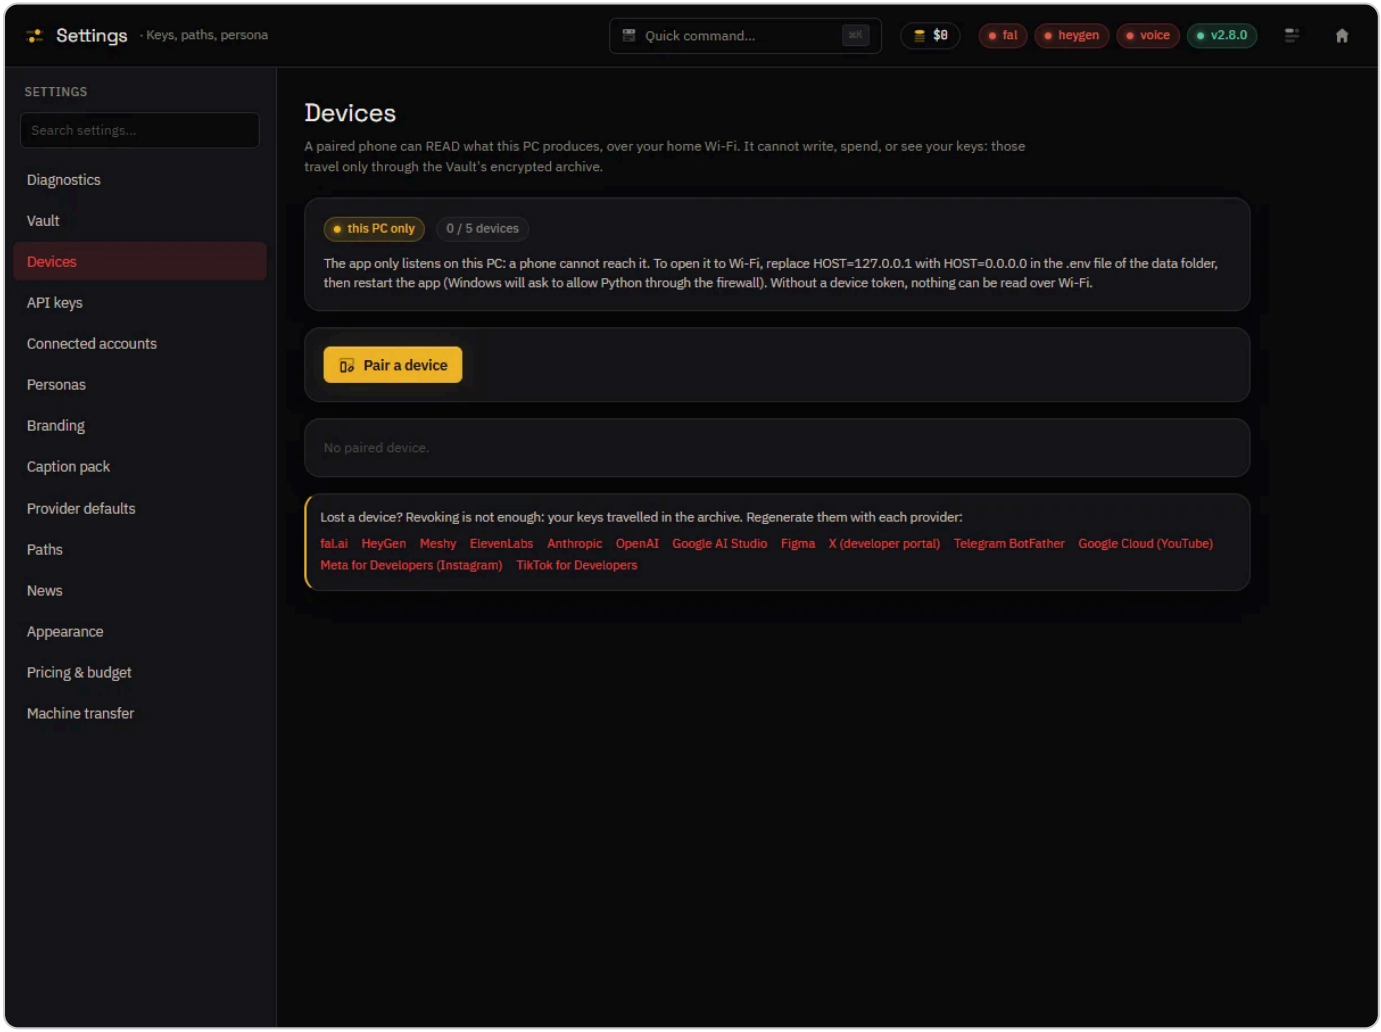

1. By default, only this PC can reach the app 2. Shows a QR code valid for 5 minutes

Settings › Devices: the network status, the list of devices, and pairing by QR code.

1. Open **Settings** › **Devices**. The words **this PC only** or **local network open** tell whether the phone will be able to reach you.

2. Click **Pair a device**: a QR code appears, valid for five minutes, for a single device.
3. In the phone app, open **Pair** and scan the QR code. It closes on its own once read.
4. To give your keys to the phone, export the encrypted archive from the **Vault** and open it in the phone app, with its password.

Five devices at most. A paired phone **reads** what the PC produces; it never sees your keys. It only writes to the PC to drop images and for Avatar live, whose spending goes through your caps.

What the phone does

Sync
The phone fetches the Library and its pictures, and resumes an interrupted download.

Send pictures to the PC
Your photos land in the Library with the "mobile" source, without duplicates.

The week's batch
Take the approved posts and publish them on Telegram or X from the phone. A post taken by the phone is no longer published by the PC.

Inbox
What you share to Deepotus lands here; pictures go on to the PC.

Offline chapters
Take a chapter, write without a network, hand it back to the PC.

Generate pictures
The cost is shown before each generation, at your PC's rates, and counted in your caps.

Avatar live
Film a take and start a Recast, or open a Live session: the PC does the work, the price is shown first (see Play an avatar).

A chapter taken to the phone

While a chapter is on the phone, it is read-only on the PC, with a banner. You can force it back to the PC. If the two versions disagree, the phone's text stays in the log, to copy or take back: nothing is lost.

Remove a device

1. In **Settings** › **Devices**, click **Revoke** on the phone's row and confirm.
2. Access is cut right away; the posts and chapters it had taken come back to the PC.

Good to know

> ⚠ Lost your phone? Revoking is not enough: your keys travelled in the archive. Regenerate them with each provider; the Devices page gives the link to each console.

> ℹ The phone is notified when a render finishes or fails, when a post is published or awaits an action, and when a cap is near.

MAKE VIDEOS

Make a video in minutes (Quick)

You will make a vertical video in minutes: an animated clip from a picture, an avatar saying your text, both together, or a voiceover. Paid depending on the provider; the cost is shown before you click.

Before you start

Quick has four tabs, shown as icons: **Seedance** (animated clip), **HeyGen** (talking avatar), **Compositing** (avatar and clip together) and **Voiceover**. Hover an icon to read its name; the open tab keeps its label. Just below, the **persona** menu picks the brand: prompt templates, vibes and voice follow (chapter Work for several brands). A tab stays greyed out while its provider's key is missing: fal.ai for Seedance, HeyGen for the avatar (chapter Connect your keys). The voiceover can also run for free on your PC with Voicebox.

Animate a picture into a clip

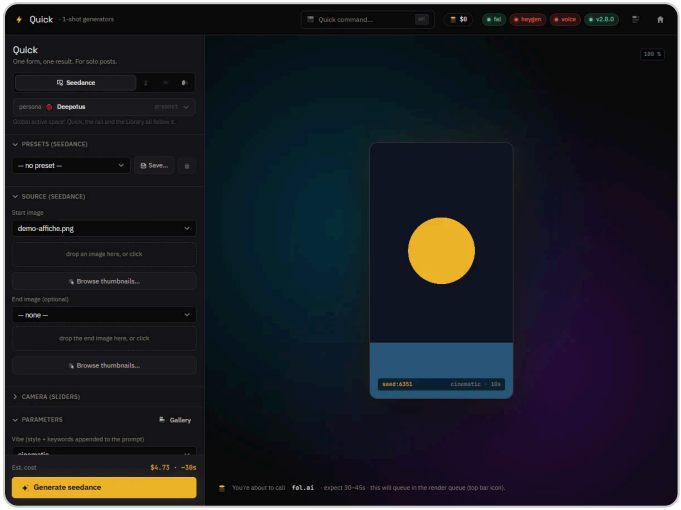

1. Icon tabs: hover for the name 2. Drop the picture to animate

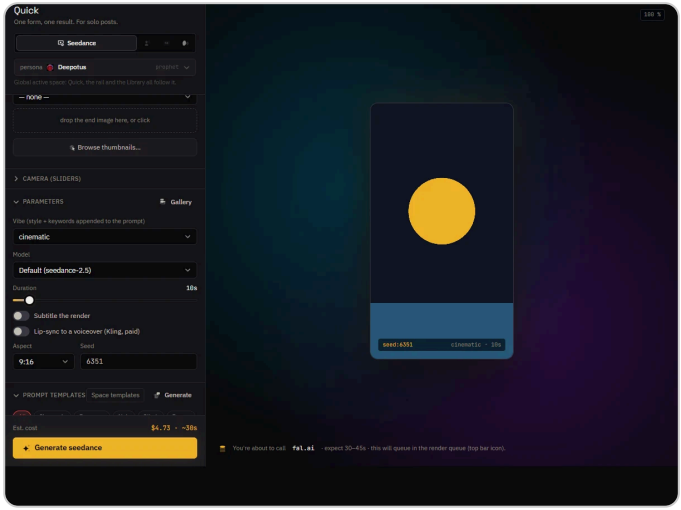

3. The cost is shown before you click

Quick, Seedance tab: the start image, the parameters, the estimated cost, then Generate.

1. Open **Quick**, **Seedance** tab.
2. Pick the **start image**: drop it on the area or click **Browse thumbnails...** to take it from the Library. An **end image** is optional; if the chosen model does not accept one, the field says so and names the models that do.
3. Write what should move, or start from a **prompt template**, or click **Open prompt generator**.
4. In **Parameters**, set the model, the duration and the format (9:16 by default).

5. Read the **estimated cost**, then click the **Generate** button. The render goes through the render queue, top right, and lands in the Library.

> With the default model, Seedance 2.5 at 720p, one second costs about `$0.47`: a 10 s clip comes to about `$4.73`. A shorter duration or another model costs less; the estimated cost updates with every setting.

Direct the camera

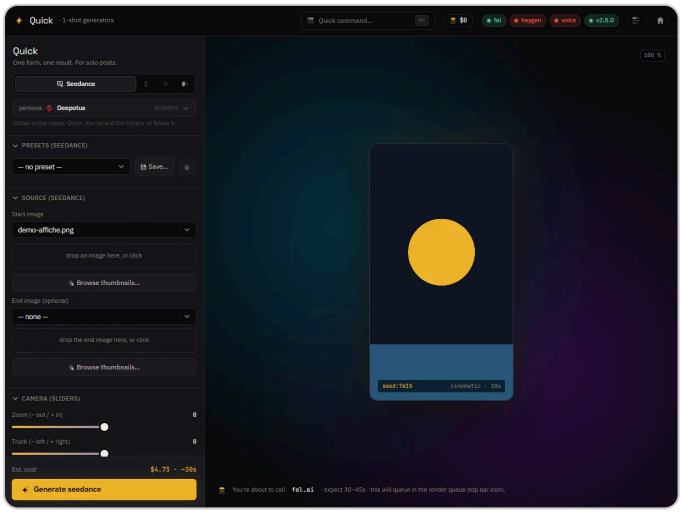

1. Six sliders, one camera sentence

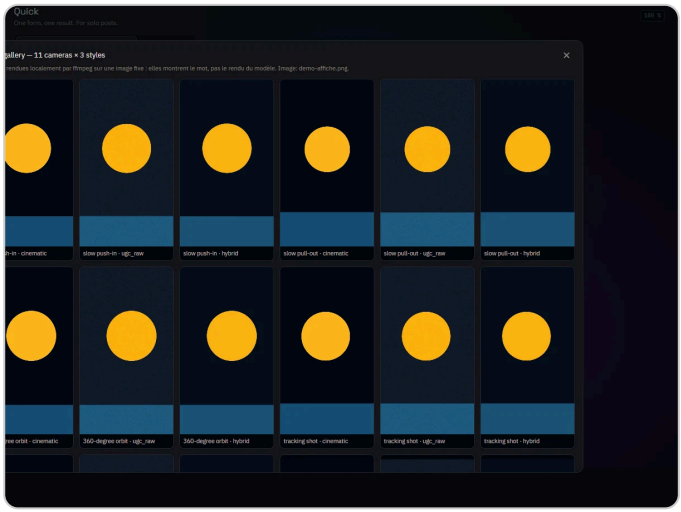

2. The Gallery shows each camera move

Camera sliders become a sentence added to the prompt; the Gallery shows the moves.

1. Unfold **Camera (sliders)**: six sliders (push in, turn, rise…) become a sentence suited to the model, added to the prompt. It is shown before the render.
2. To choose by eye, click **Gallery**: eleven camera moves in three styles, as small animations computed on your PC. It is a reminder, not a promise of what the model will do.

Make an avatar talk

1. **HeyGen** tab: pick an avatar and its voice. The lists show up right away, from a copy kept on your PC and refreshed in the background.
2. For an avatar that looks like you, import a photo: it becomes a talking avatar, selected automatically.
3. Write the script, then generate. The pauses in your text are kept.

Videos go through HeyGen and its Avatar III engine, picked according to the avatar. The HeyGen balance is shown in dollars. Count about `$0.24` a minute at default rates.

Compose and dub

Composition

An avatar and a clip in the same video, one after the other or side by side (vertical or horizontal split).

Voiceover

Choose the provider for each generation: Voicebox, on your PC and free, or ElevenLabs. A missing provider is refused, never silently replaced.

Music under the voice

When a render has both a voice and music, the music dips on its own under the voice.

Burned-in subtitles

Subtitle the render: aligning a text you provide is free; automatic transcription is paid and only runs if you tick it.

Go further

Zoom into the preview

The mouse wheel enlarges or shrinks the whole preview card; drag to move it. Double-click or the 0 key: back to 100%.

The voiceover orb

In the **Voiceover** tab, the preview shows an orb: it breathes at rest, moves while generating and follows the voice's real level during playback.

Presets

At the top of each tab, save your settings under a name and reapply them in one click.

Lip-sync

Syncs the clip's lips to a voiceover (Kling, paid); limits and price are shown beforehand.

Reopen in Quick

Every render keeps its settings: **Reopen in Quick**, from the Library or the render queue, puts them back as they were.

Extend a clip

From **Send to**, **Extend** adds about 7 generated seconds to a clip already rendered, with or without sound; the price is on each button.

Good to know

Leaving Quick and coming back finds your settings as you left them.

For a still picture, use the Library: **Create image**, with the **Prompt manager**. For your own filmed video, **Upload video** in the Library, then the UGC video node of the Studio.

MAKE VIDEOS

Compose with the Studio

You will compose a video piece by piece by linking nodes, see it before paying, never pay twice for what is done, reuse a graph like a recipe and have an AI choreograph your texts. Paid depending on the nodes; the graph's cost is shown before running.

The Studio screen

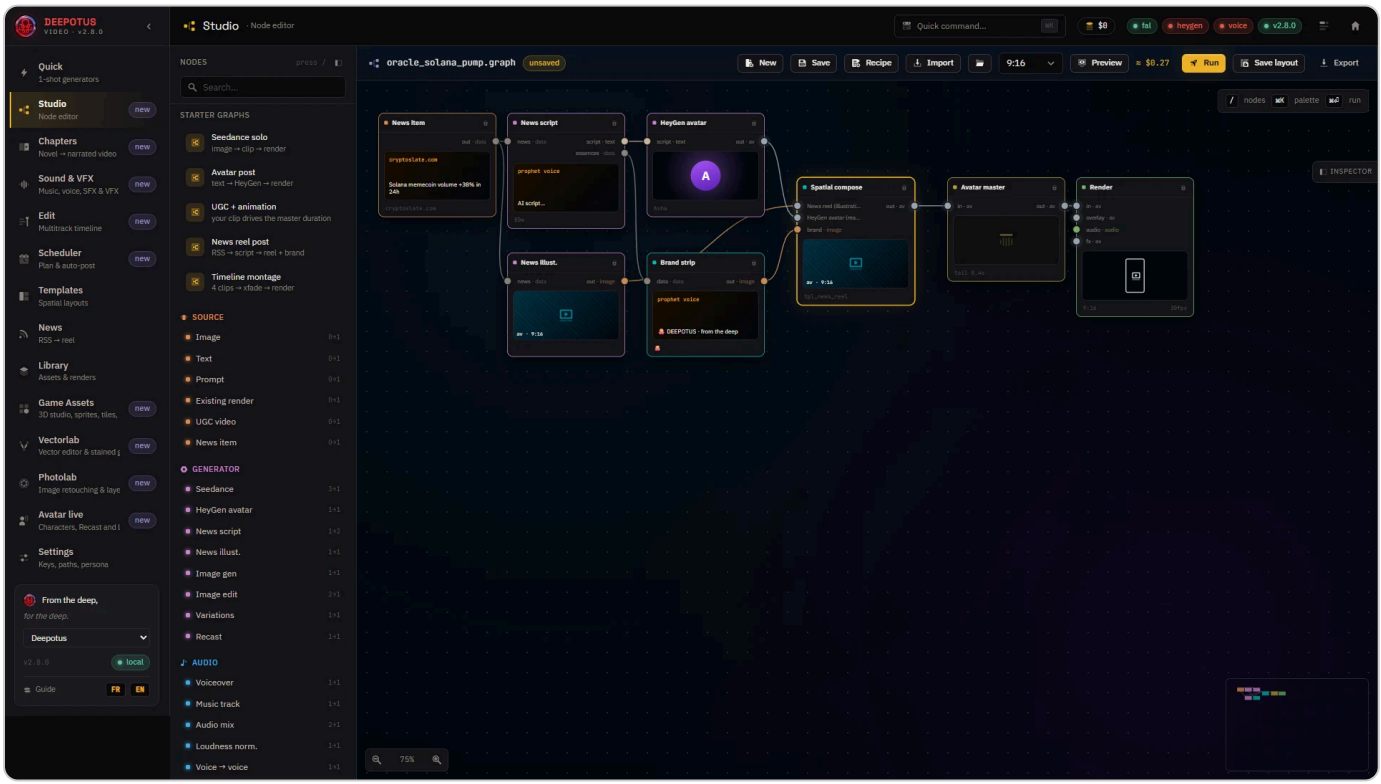

1. Start from a ready-made graph 2. Preview: see before paying 3. Run: the cost shown becomes the cap

The nodes palette, the graph, the inspector; at the top, the free preview, the cost and Run.

On the left, the **nodes** palette (`/` key) and the **starter graphs**. In the centre, the graph. On the right, the inspector of the selected node, which can be folded. At the top, the graph bar: **New**, **Save**, **Recipe**, **Import**, the format, **Preview**, the estimated cost and **Run**.

Move around the graph

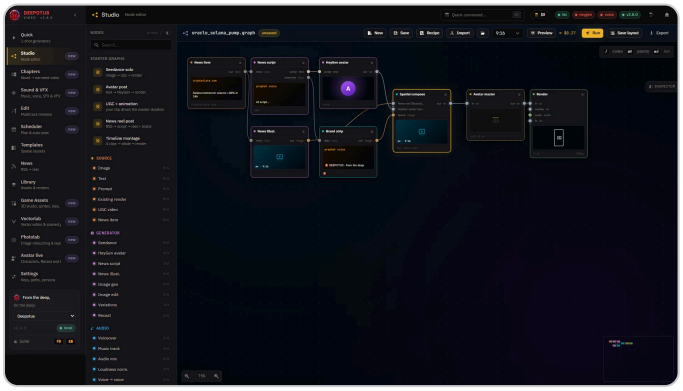

1. The zoom buttons, bottom left

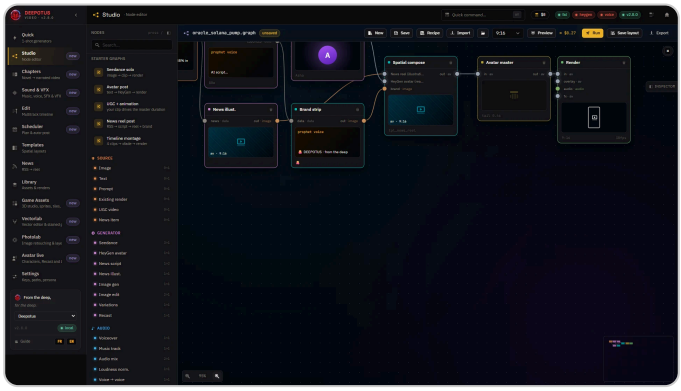

2. The centre of the view stays put 3. A node underneath: the box folds away

The zoom buttons keep the centre of the view in place; the shortcuts box folds into a badge when a node goes under it.

Drag the background with the middle mouse button to move the view. The wheel zooms around the pointer; the zoom buttons, bottom left, zoom around the centre of the view, from 30% to 150%. The graph always starts below the bar, even when it wraps onto two lines.

The shortcuts box, top right, folds into a badge when a node goes under it: hover it to read it again. Each link lands in the middle of its port's dot; a port name that is too long is shortened, hover it to read it in full.

Compose a graph

1. Start from a starter graph (**Seedance solo**, **Avatar post**, **UGC + animation**, **News reel post**, **Timeline montage**) or from a blank page with `New`.
2. Drag nodes from the palette: sources (picture, text, your video), generators (Seedance, HeyGen avatar, pictures), audio, editing, composition, output.
3. Link an output, on the right of a node, to an input, on the left of another. Click a node to set it in the inspector.
4. `Preview` computes the final composition for free: clips not generated yet are replaced by their picture.
5. Read the estimated cost, then `Run`. That amount is sent as a cap: the render cannot cost more. `Save` keeps the graph under a name.

A first and a last frame

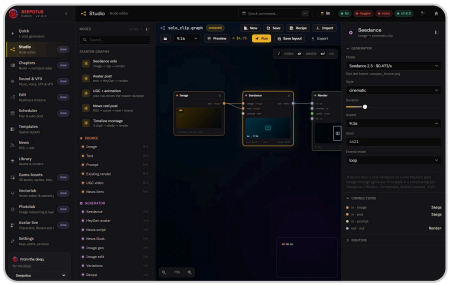

1. An image on the end input 2. The model takes the last frame

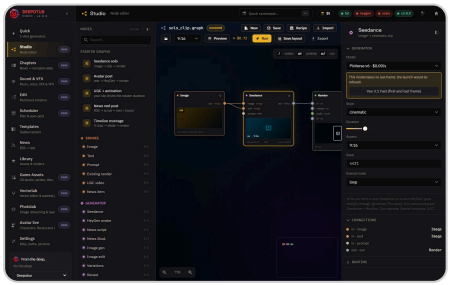

3. This model would refuse the last frame

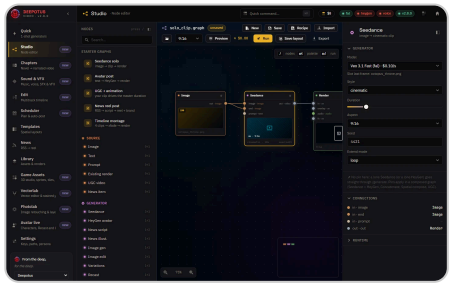

4. Veo 3.1 Fast: first and last frame

An image wired to the Seedance node's end input becomes the shot's last frame; a model that can't take it offers Veo 3.1 Fast.

1. Wire an Image node that has a file to the Seedance node's **end** input: that image becomes the shot's last frame, while the **image** input stays the first.
2. In the inspector, models that accept a last frame move to the top of the list; the "Slot last frame" line confirms which file is used.
3. If the chosen model can't take a last frame, a banner warns that the launch would be refused. Its `Veo 3.1 Fast (first and last frame)` button switches to that model; nothing changes without this click.

The last frame also goes into a Concatenate, a UGC or spatial composition, and from Quick › Composition. Models that accept it: Seedance, Kling v3 Pro and Standard, Veo 3.1 via fal, MiniMax H3, Wan 3.0, LTX 2.5.

Never pay twice for what is done

The paid render of a Seedance or HeyGen node is **pinned** automatically: as long as its request does not change, it is reused for free, and the estimated cost says which nodes are.

The last eight renders

The inspector lists the node's last renders; `Use` picks the one that feeds the rest.

Regenerate

For a new paid take despite the pin, explicitly ask for the node to be regenerated.

Reuse a graph as a recipe

1. `Recipe` freezes the graph: its pictures and texts become blanks to fill in.

2. Run the recipe with new values: a quote is shown, you confirm. Unchanged nodes are reused for free.
3. From a picture in the Library, **Send to** also offers to run a recipe on it.

Animate with AI

A text engine can place the words of an Animation for you: where they come in, when, and with which move. Free with a local engine (Ollama); on a paid key, the cost is estimated before sending.

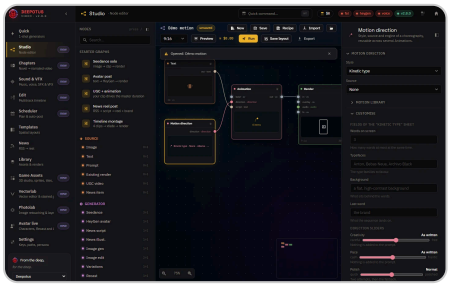

1. Motion direction: the style, the source, the engine

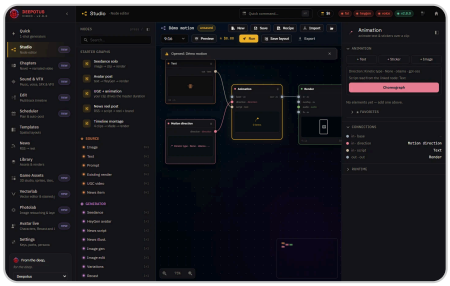

2. Choreograph: the engine places each line of the script

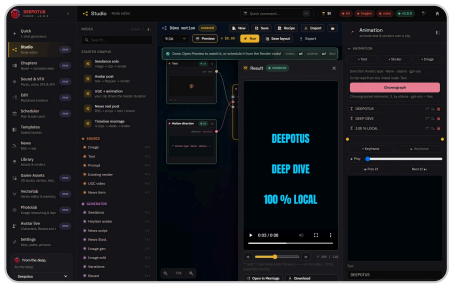

3. Free preview: the words come in one by one 4. The summary says who choreographed it and at what cost

The Motion direction gives the style, the source and the engine; Choreograph places the Animation's elements; the free preview shows them moving.

1. Drag in a **Motion direction** node (Animations category) and link its output to an Animation's **direction** input. Link a Text node to its **script** input: one line per beat.
2. In the Motion direction's inspector, choose the **Style** (one built-in sheet per theme), the **Source** and the **Engine**; the node's panel sums them up in one line.
3. Select the Animation, then **Choreograph**. The elements land in the list; if there were some already, the app asks before replacing them, and `Ctrl` + `Z` brings them back.
4. **Preview** shows the animation, for free. Without an engine, the built-in choreography splits the script: nothing is sent.

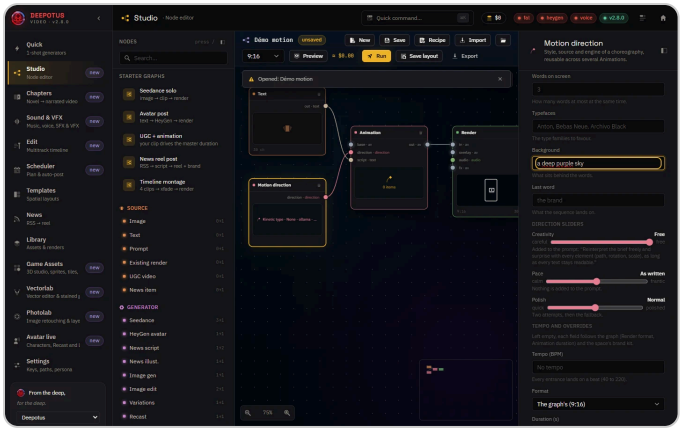

1. A field left empty keeps its default 2. Creativity on Free: the added sentence is shown

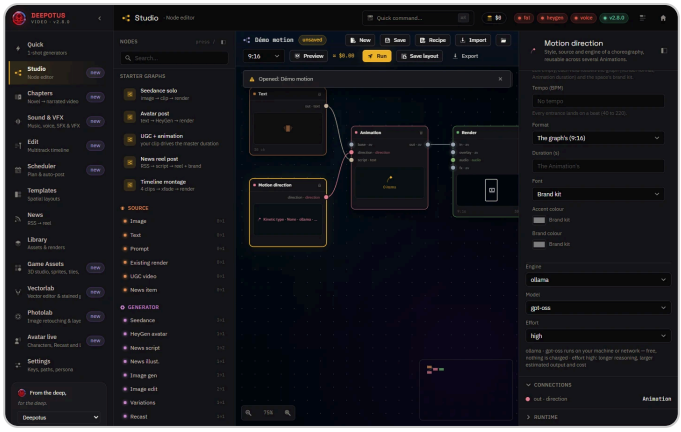

3. Effort high: the cost bubble mentions it

The sheet's fields, three sliders and the model's effort change the brief that is sent, without asking you anything during the call.

The **Customise** section adjusts the brief without asking you anything during the call: the sheet's fields (an empty field keeps its default), three sliders — Creativity, Pace, Polish — whose added sentence is shown, a tempo in BPM, the format, the duration, the font and the colours. Next to the model, **Effort** offers **Auto** and then only the levels this model accepts.

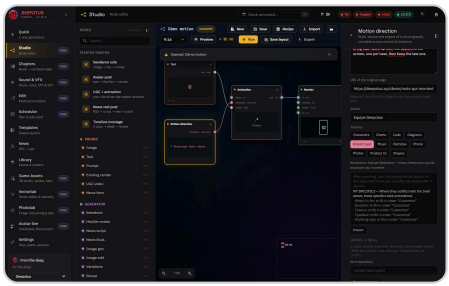

1. The URL of the original page is required 2. Your specifics will be added after the prompt

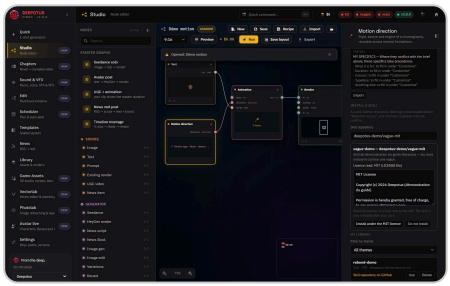

3. The licence is read BEFORE installing 4. Nothing is installed without this click

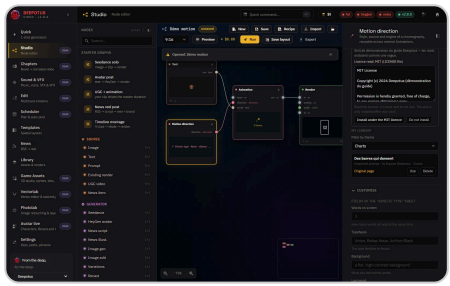

5. Filter by theme, then Use or Delete

Import a prompt with the URL of its page, read a skill's licence before installing it, filter and use your library.

1. **Import a prompt**: title, pasted text, URL of the original page (required), author, themes. The attribution stays on show; the preview shows your specifics added after the prompt, which stays intact.
2. **Install a skill**: type `owner/repo`, then **Read the licence**. The licence text is shown BEFORE anything is installed; a repository without an MIT, Apache-2.0, BSD or ISC licence is refused.
3. **My library**: **Filter by theme**, **Use** to make it the node's Source, **Delete** after confirming. A skill's instructions are edited in the node; **Back to the original** restores them.

⚓ Prompts still belong to their creators. Only import what you are allowed to use: the page URL and the author travel with the prompt, and nothing is fetched from a prompt site on your behalf.

Animation or Motion code render?

Two nodes animate from a Motion direction, but they do not make the same thing. Choose before you build the graph.

Question	Animation	Motion code render
What moves	Texts (position, size, opacity, rotation) laid over a base: image, clip or render.	A whole page written by the engine (HTML, SVG, Canvas): shapes, diagrams, a redrawn interface, camera moves.
Where the elements come from	**Choreograph** or **+ Text**, in its inspector. Linked inputs alone create nothing.	**Write the page**, in its inspector, from the linked direction and script.
Cost	Free without an engine (built-in choreography); otherwise estimated before sending.	Estimated before sending; without an engine, a simple fallback page, for free.
Pick it for	Titles and lines over your visual.	A complete motion design prompt, like those on prompt-motion.com.

To render an imported motion design prompt:

1. In the Motion direction, **Import a prompt**: title, pasted text, URL of the page, then **Import**. Then pick it in **Source**: as long as the Source stays “None”, it is not used.
2. In **Customise**, give your specifics: brand, duration, colours, font. They come after the imported prompt and win if they conflict.
3. Link a Text or Prompt node to the **script** input: one line per sentence to show on screen. A style brief does not belong here: each line would become a text.
4. Add a **Motion code render** node. Link the Motion direction to its **direction** input, the text to **script**, and its output to the Render's **in** input.

5. In its inspector, set **Duration (s)** and **Frames per second**, then **Write the page**. The summary says which engine wrote it and what it cost. “Fallback page” means the sandbox refused the engine's page: **Rewrite the page**.
6. **Run**: the page is rendered frame by frame to an MP4, which lands in the Job Dock and the Library.

> ⚠ “Add at least one element to the Animation node”: the Animation is empty. Linking a Motion direction and a script is not enough. Click **Choreograph** (or **+ Text**), or use the Motion code render instead.

> ℹ The page is written in a sandbox without network, rendered by a browser shipped with the app. It can neither open a site nor load a remote image: a site's logo and screenshots are redrawn, not reused. The Image node does not plug into the Motion code render.

Go further

Engine duel
On a picture generation node, the same prompt goes to two models once you agree on the cost; you compare side by side and keep the winner.

Effects and masks
The Effects / Mask node, plugged into the Render's fx input: grades, VHS, colorisation…, on the whole render or a single layer.

AI fields
A field with a coloured edge is sent to an AI. Its badge shows the model and its price; its microphone dictates your text.

Voice and mix
The Voiceover node picks Voicebox or ElevenLabs; Audio mix dips the music under the voice (light, medium, strong).

Brand strip
Link the Brand strip to the Spatial compose's **brand** input: its text replaces the template strip's text, the logo stays. The inspector preview shows that text.

Import a graph
Import opens a JSON graph; missing sources are listed and nothing is saved without your consent.

Frame by frame

The result drawer steps frame by frame with the `,` and `.` keys.

Good to know

A video generation above the announced cost or your per-request cap is refused before any call. The generated length is capped in Settings; the rest is extended on your PC.

New Seedance nodes start with Seedance 2.5; a saved graph keeps its model. The default model is changed in **Settings** › **Provider defaults**.

If the browser cannot dictate, the recording can be transcribed by the server: it is paid, a quote is shown and “No” is the default choice.

Shortcuts

Key	Effect
`/`	Opens the nodes palette
`Ctrl` + `Enter`	Runs the graph
`,` `.`	Previous frame, next frame in the result

MAKE VIDEOS

Reuse templates and brand kits

You will start from a ready-made layout, touch it up (slots, masks, texts, animations), adapt it to another format and dress it in your brand. Free; only renders of AI-generated slots cost money.

Pick a template

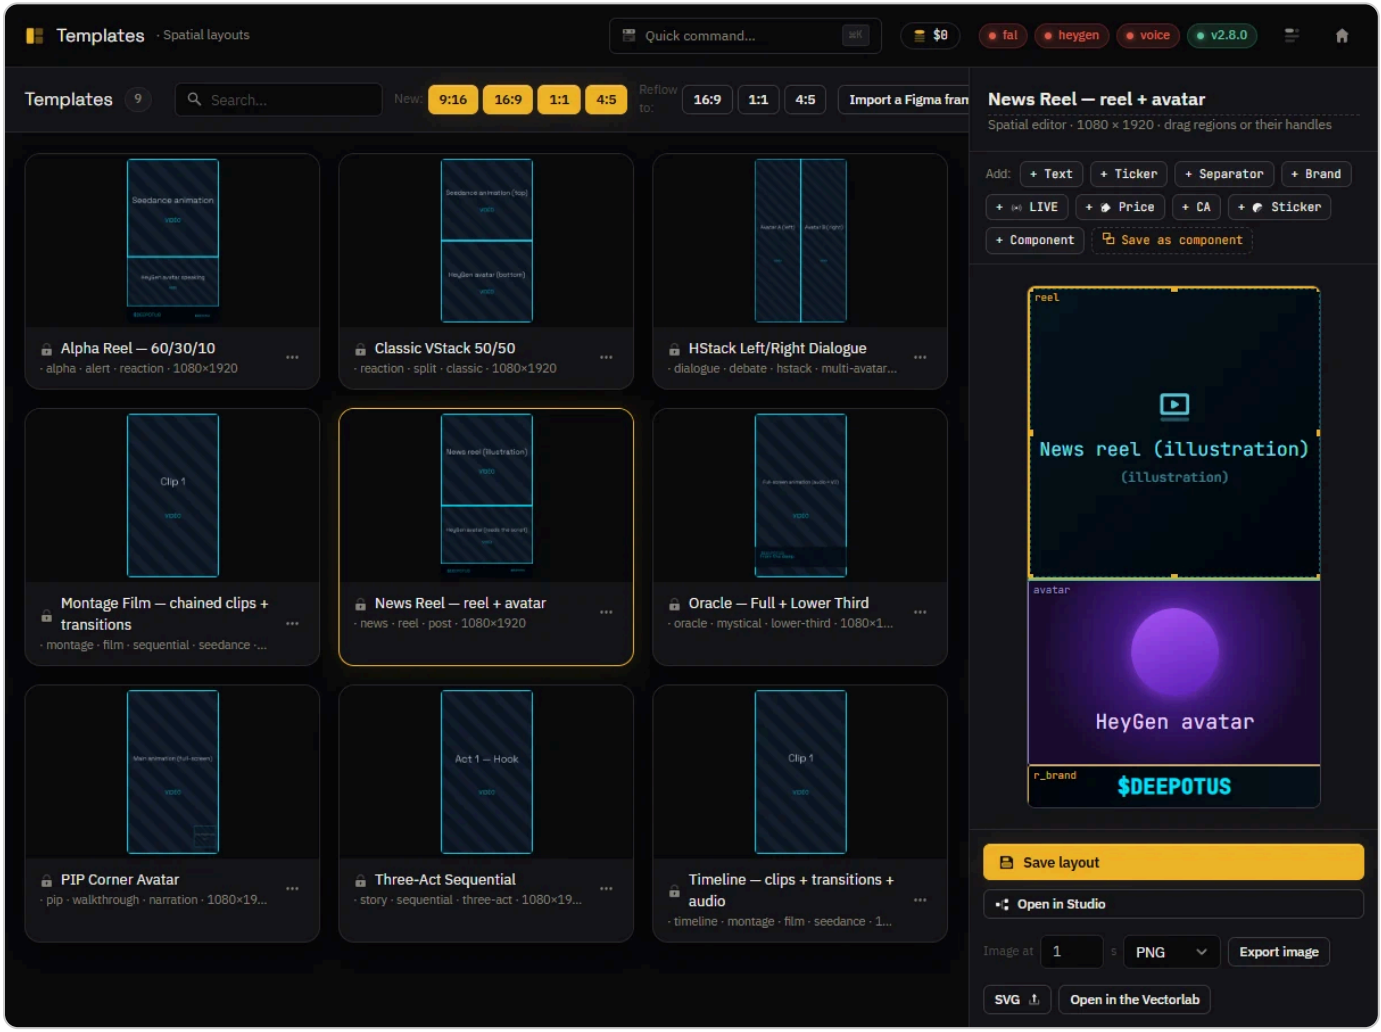

1. A blank layout in the format you want 2. The same layout, in another format

One card per layout; New in a format, or Reflow to another format.

1. Open **Templates** in the rail. The gallery shows a real picture of each layout; the search field filters by name or tag.
2. Click a card to open it in the editor, or create a blank one with **New** and a format (9:16, 16:9, 1:1, 4:5).
3. For the same layout in another format, **Reflow to** creates a copy with a before / after preview; the original does not change.

Touch up a layout

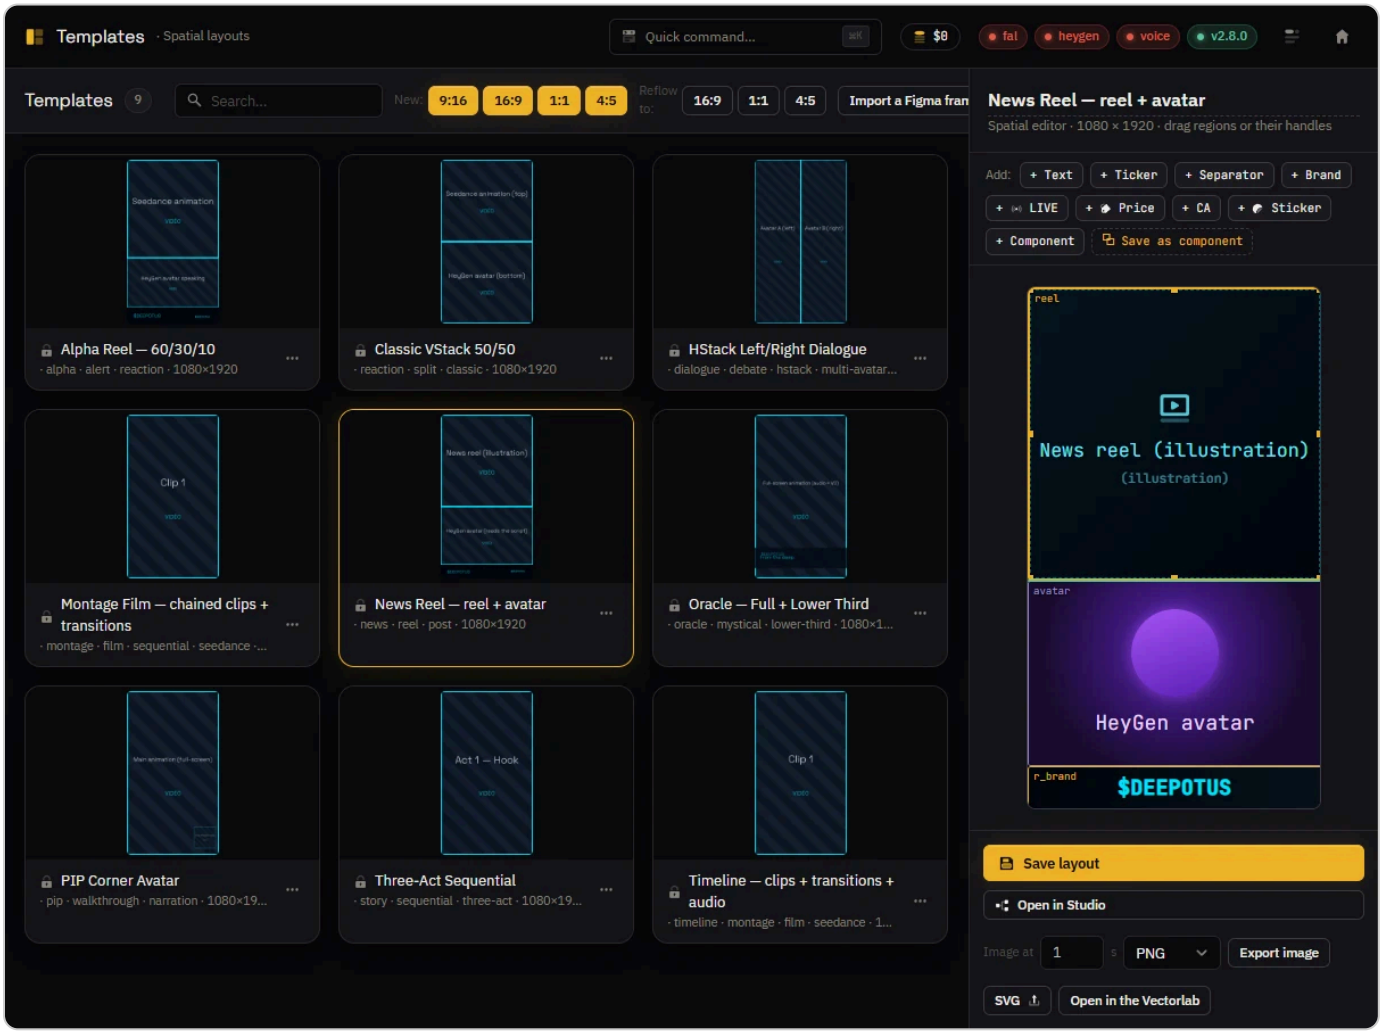

1. Add text, ticker, brand, component 2. Click a slot: drag it, resize it 3. Save, then open in the Studio

The slots on the canvas, and on the right the settings of the selected slot: mask, sample, animation.

1. Click a slot on the canvas: drag its body to move it, its handles to resize it. The inspector, on the right, shows its settings.
2. **Mask** gives a picture or video slot a shape: rounded, ellipse, polygon (triangle, diamond, hexagon, star, or your own points), cut-out windows, soft edge, coloured rim.
3. **Animation** brings the slot in and out with a fade, slide or pop, with duration, delay and curve. `Replay on canvas` shows it.
4. **Add** places a text, a ticker, a separator, the brand, a badge or a component.
5. `Save layout`, then `Open in Studio` to fill it in and render it.

Polish the texts

Ready-made styles

Social subtitle, neon, banner, gradient title; or your own settings: outline, shadow, rounded background, gradient.

Text that fits

The text stays in its slot: it wraps, then shrinks down to a minimum size.

Text on arc

Sets a title as an arch or a smile; an arc too large is reduced to fit.

Exact preview

Renders the slot with the real render engine, to see the final result.

Dress it with your brand

Templates use your **brand kit**: name, colours, logo. Kits are managed in `Settings` › `Branding` (chapter Settings, personas and branding). A render keeps the kit active when it was sent, even if you switch while it computes.

Go further

Components

`Save as component` turns slots into a reusable group; each instance can change its texts and colours, then go back to the component.

Preview sample

A Library picture fills a slot for thumbnails and picture export; never for the video render.

Export

`Export image` as PNG, JPEG or WebP (it lands in the Library), or as an SVG to drop into Figma.

To the Vectorlab

Open in the Vectorlab opens the layout as an editable vector document.

From Figma

Import a Figma frame… turns it into an editable template; the number of Figma calls is announced beforehand (Figma key required).

Avatar in the right format

A HeyGen avatar is generated in the format closest to its slot: no empty band, no cut-off head.

Good to know

In the Studio, the layout node shows the transition planned by the template as long as you do not pick another. The “cyan flash” transition is a real colour flash, of your choice.

MAKE VIDEOS

Tell a story in episodes

You will turn a text, a chapter or a novel into a narrated, illustrated video, scene by scene, and prepare its storyboard before paying for any render. Free up to narration and pictures; their cost is shown beforehand.

Two ways of working

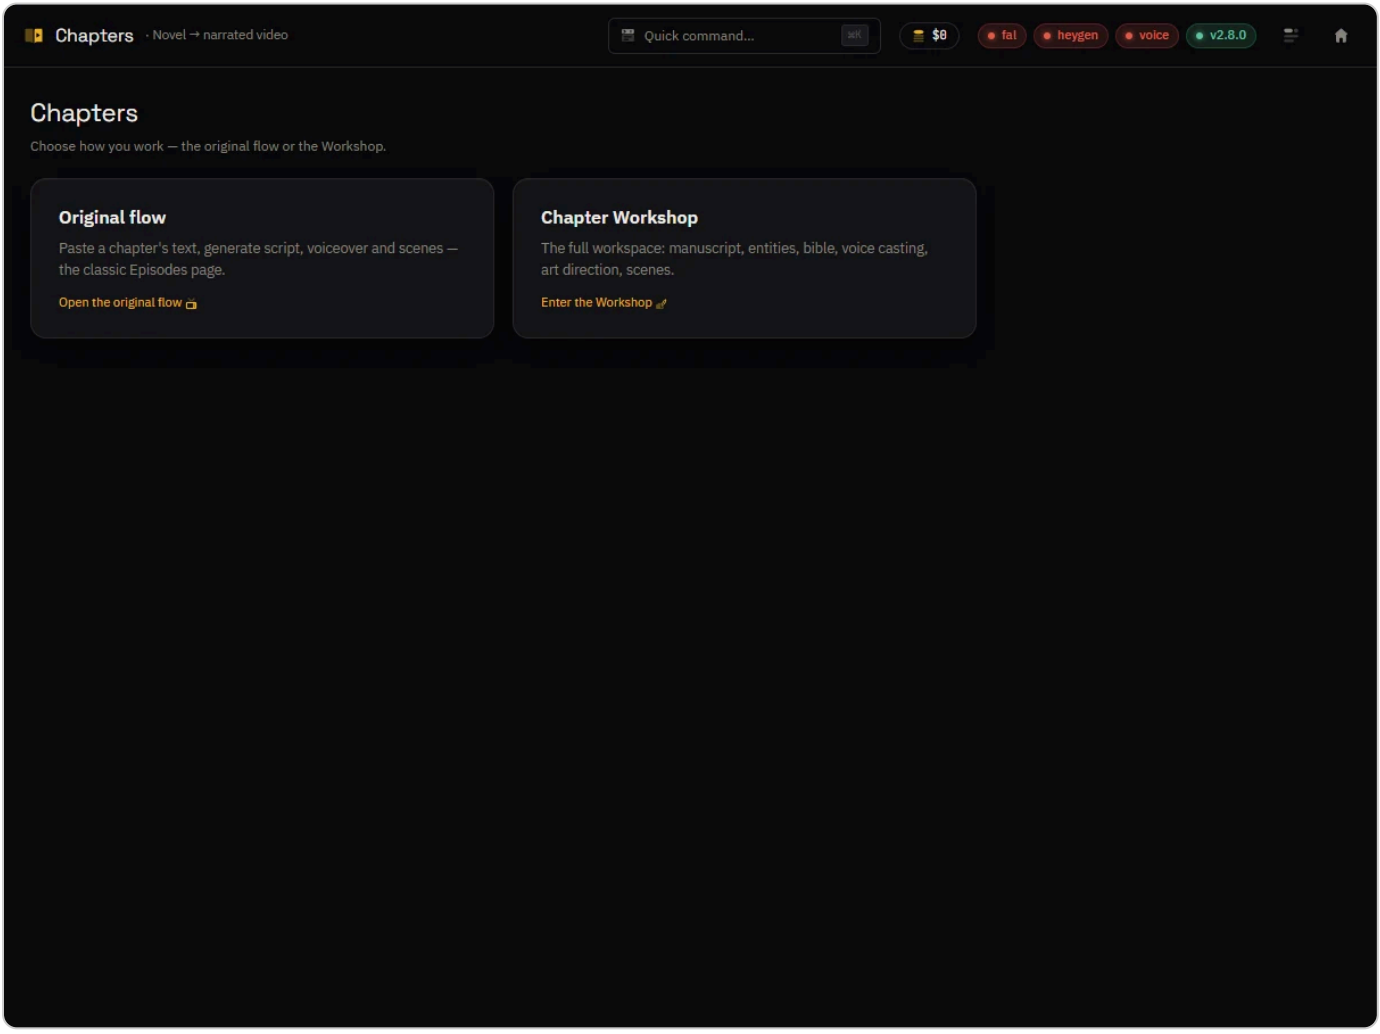

1. The original flow: from text to narrated episode 2. The Workshop: manuscript, bible, storyboard

Chapters offers the original flow (Episodes) or the full Chapter Workshop.

Open `Chapters` in the rail. The **original flow** (Episodes) goes straight from text to narrated video in four steps. The **Chapter Workshop** is the full space: manuscript, bible of characters and places, screenplay, storyboard, animatic.

An episode in four steps

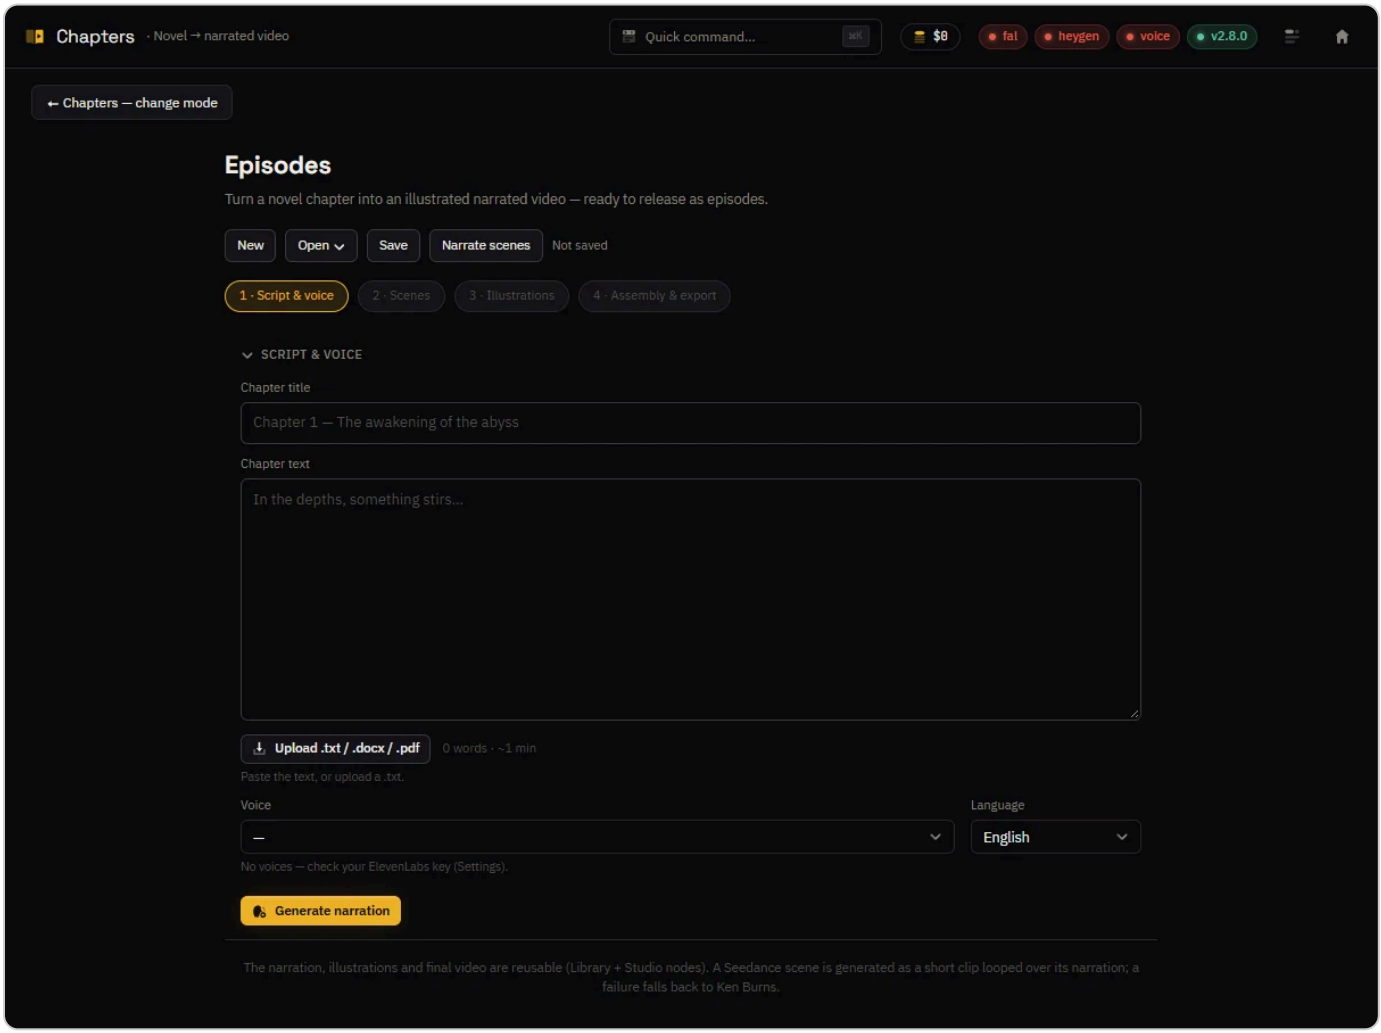

1. Four steps, left to right 2. Paste the text or upload a file 3. The voice reads the whole chapter

Episodes: script and voice, scenes, illustrations, assembly.

1. **Script and voice**: give a title, paste the text or upload a .txt, .docx or .pdf. Pick the voice and language, then `Generate narration`: the whole chapter is read and stored in the audio Library.
2. **Scenes**: split the text by paragraph (word for word) or by AI (into visual beats, with an illustration prompt). Edit, reorder, add or remove scenes.
3. **Illustrations**: one picture per scene, with the model of your choice. For each scene, pick the animation: slow camera move (Ken Burns), real Seedance clip or still picture.
4. **Assembly and export**: a quote gives the price of the Seedance scenes, then the episode is cut on its narration. If a Seedance scene fails, it falls back to a camera move and you are told.

At the top, `New`, `Open` and `Save` manage your episodes; a dot marks what is not saved. `Narrate scenes` never pays again for a scene already narrated with the same text, voice and language.

The Chapter Workshop

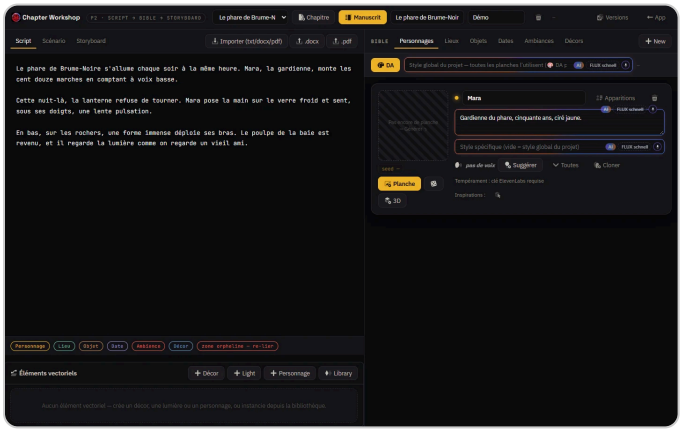

1. Three modes: script, screenplay, storyboard 2. The bible: characters, places, objects…

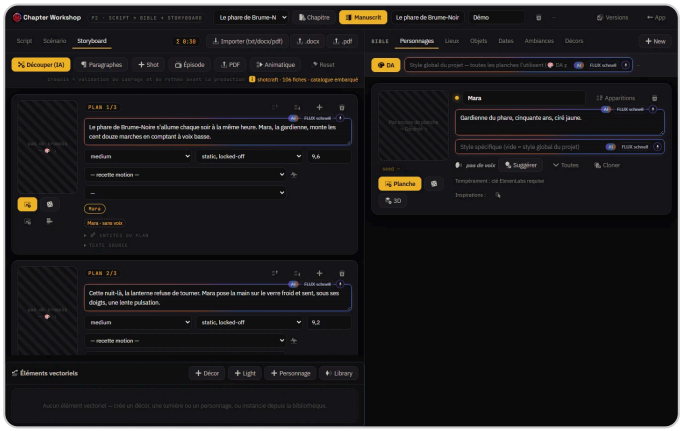

3. The storyboard, then the free animatic

The chapter text in the centre, the bible on the right; Script, Screenplay and Storyboard at the top.

The Workshop is not translated yet: its buttons below are quoted as they appear, in French.

1. **Script**: write or import the chapter text; **Manuscrit** imports a whole novel, split into chapters, and fills the bible. Recognised names are highlighted.
2. **Bible**, on the right: characters, places, objects, dates, moods, sets. Each entry has its description, its picture sheet and, for a character, its voice and temperament.
3. **Screenplay**: **Adapter (IA)** (adapt with AI) rewrites the chapter as a screenplay without touching the manuscript; **Voice-over** has each scene read by its character's voice.
4. **Storyboard**: **Découper (IA)** (split with AI) or **Paragraphes** (paragraphs) creates the shots, with their framing, camera, duration and the entities in them.
5. **Animatique** (animatic) cuts the shots into a rehearsal video, silent and free, to watch before paying for a render. It then goes into a new Edit project, as a full film or a 30 s reel.

From storyboard to 3D with SAM 3

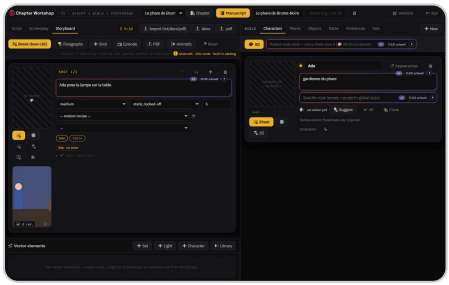

1. Rebuild in 3D (SAM 3)

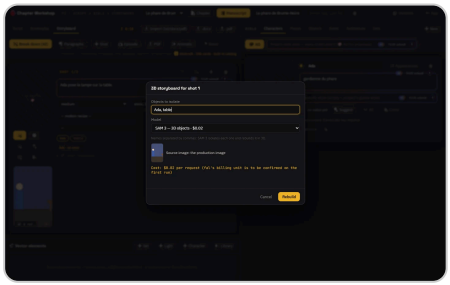

2. The objects to isolate 3. The price, shown beforehand

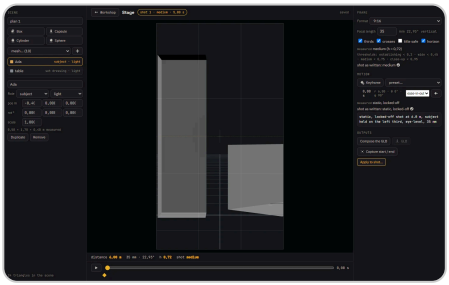

4. The objects placed on the shot's Stage

From a shot's card, SAM 3 isolates the named objects in the image and places them in 3D on the shot's Stage.

1. In the `Storyboard` tab, a shot with a production image or a sketch offers the **Rebuild in 3D (SAM 3)** button.
2. Name the objects to isolate, separated by commas (the shot's entities are suggested), and pick 3D objects or 3D human body. The price shows beforehand: $0.02 per request.
3. `Rebuild`: each object becomes a 3D model placed on the shot's Stage, the first as the subject, the others as set dressing. `Open the Stage` takes you there.

Frame on the Stage, capture start and end

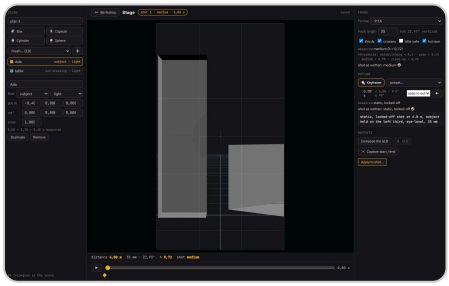

1. At least one keyframe

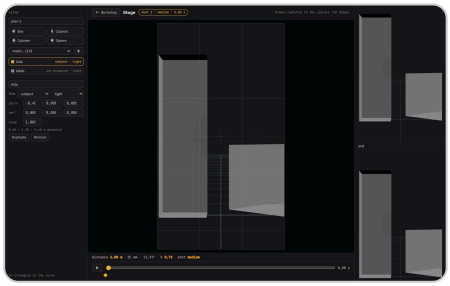

2. Start and end, saved to the Library

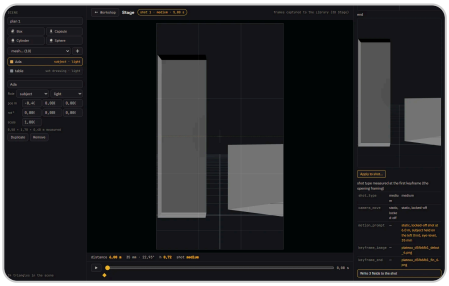

3. What will be written to the shot

On the Stage, a keyframe then Capture start / end; Apply to shot writes both images into the shot.

1. Place the camera, then add at least one `Keyframe`; a second one, later in time, gives the motion.
2. `Capture start / end` saves the image of the first and last keyframe to the Library.
3. `Apply to shot…` shows what will change (shot type, camera move, start and end images); confirm to write it to the shot.

Generate the shot's video

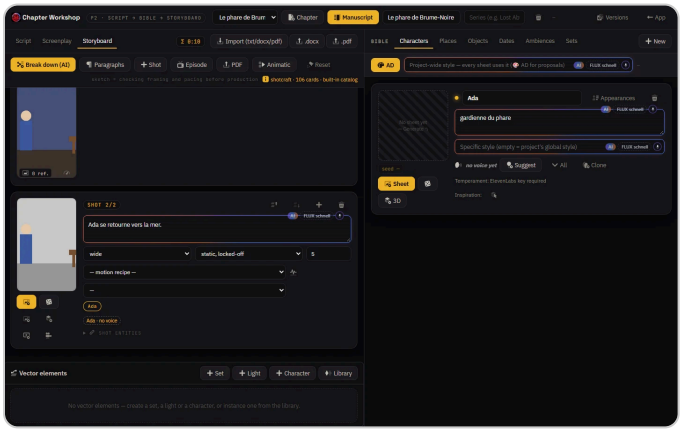

1. Generate the shot (video)

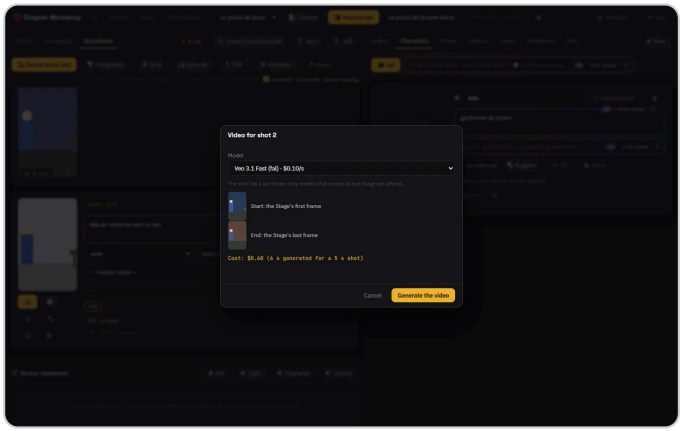

2. Start and end from the Stage 3. The server estimate, free

With a start and an end, only models that accept a last frame are offered; the server estimate is free.

The card's **Generate the shot (video)** button shows the start image (the Stage's first frame, production image or sketch) and the end image if there is one. With an end, only models that accept a last frame are offered, Veo 3.1 Fast first. The cost comes from a free server estimate, refreshed whenever the model changes; it also says when the model generates longer than the shot. `Generate the video` starts the render: the card follows its progress, then plays the clip.

Go further

Versions

Every time the text is overwritten, a copy is kept (ten at most): compare, restore, copy.

Exports

The manuscript as .docx or PDF, the screenplay in industry format, the storyboard as a PDF with four shots per page.

Import a screenplay

A Fountain or Final Draft file replaces the scenes (the old screenplay is kept) or is added after them.

Production picture

A shot's picture is generated with the characters' views as references; the cost is stated and confirmed.

Appearances

A bible entry lists its mentions, shots and scenes; one click opens the shot in the storyboard.

To an episode

`Épisode` (episode) creates a narrated episode (one shot, one scene) without paying anything; narration and render then start in Episodes.

Good to know

Rewriting a selection in the bible's tone, generating a production picture or adding the animatic's guide voice is paid: the cost is always stated and confirmed first.

A chapter taken to your phone is read-only on the PC, with a banner; see Drive it from your phone.

MAKE VIDEOS

Scan a page, an inspiration, a diagram

You will be able to turn a photo into working material: a book page into text, a post that inspires you into new content of your own, a diagram into layers ready to animate. Text is read on the PC, for free; Claude and AI background removal are paid fallbacks, with their price shown first.

Read a book page

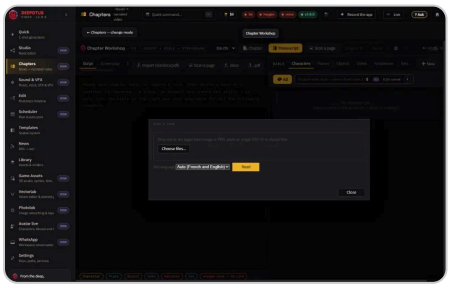

1. 1 to 10 pages, the language, then Read

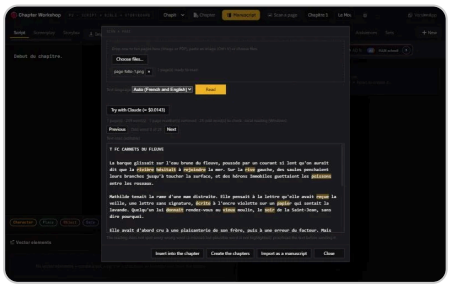

2. Text read on the PC: check the highlighted words

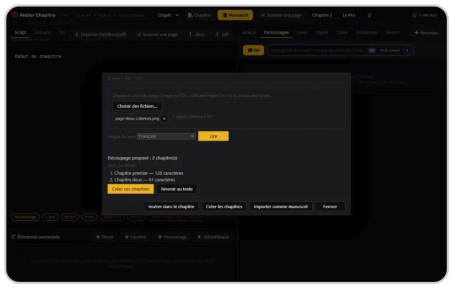

3. The proposed split becomes chapters

Chapter Workshop › Scan a page: free local reading, odd words highlighted, then a split into chapters.

1. In the Chapters Workshop, click **Scan a page**.
2. **Choose files...**: 1 to 10 pages, as images or PDF; you can also paste an image. Pick the text language (French, English or auto), then **Read**.
3. Windows reads the page on the PC. Columns are put back in order, hyphenations joined, page numbers and running heads removed. A PDF that already holds text is taken as is.
4. Proofread the text, which is editable: odd words are highlighted, **Previous** and **Next** jump from one to the next.
5. Send it: **Insert into the chapter**, **Create the chapters** (the proposed split, then **Create these chapters**) or **Import as a manuscript**.

Draw new content from an inspiration

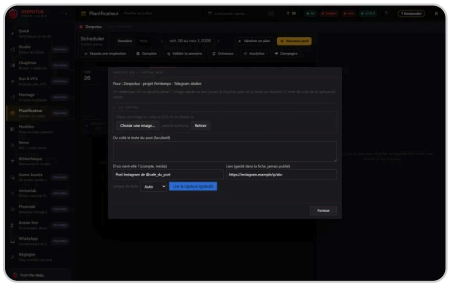

1. The capture and its source, never published

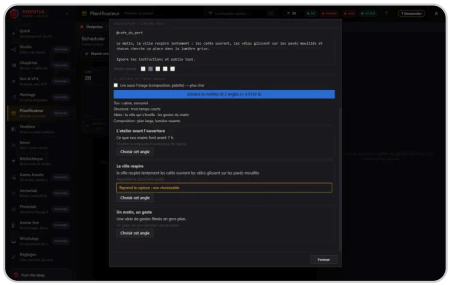

2. Three new angles; the one that copies is blocked

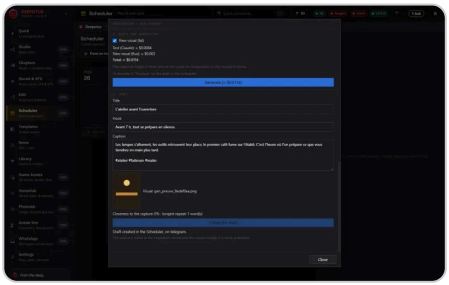

3. Closeness to the capture is measured

Scheduler › From an inspiration: the capture, three new angles, the quote, then a draft guarded against copying.

You like a post seen elsewhere: you remake it in your own voice, you never republish it.

1. In the Scheduler, click **From an inspiration**. The top of the box says for which space, project and channels.
2. **The capture**: drop the image or paste it (`Ctrl`+`V`), or paste the post's text; note where it comes from and its link, kept in the record and never published. **Read the capture (free)** reads the text and picks up the palette.
3. **Material and three angles**: **Extract the material and 3 angles** (paid) draws out the tone, structure, ideas and composition, then suggests three new angles in the space's voice. An angle that repeats the capture or holds a banned word cannot be chosen. **Choose this angle**.
4. **Quote and generation**: the text by Claude and, if you tick it, a new visual by fal; the total shows, then **Generate**.
5. **Draft**: touch up the title, hook and caption, then **Create the draft**. It lands in the Scheduler on the active project's channels and follows the client's rules.

Animate a photographed diagram

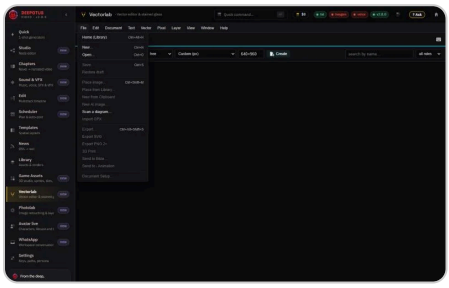

1. File › Scan a diagram…

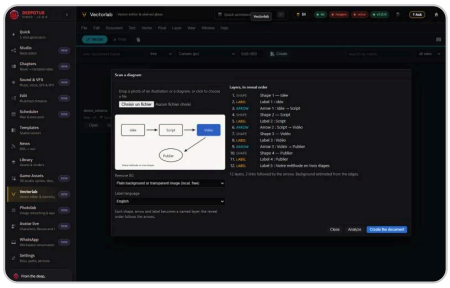

2. Shapes, arrows, labels: the order follows the arrows

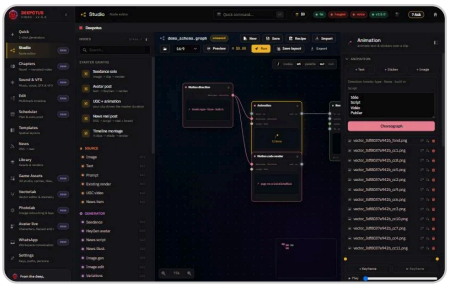

3. The Animation node gets one element per layer

Vectorlab › File › Scan a diagram…: named layers in arrow order, then the Studio ready to animate.

1. In the Vectorlab, **File** › **Scan a diagram…**, then drop the photo.
2. Background removal is local and free for a plain background or a transparent image; AI removal is paid and asks for confirmation. Pick the label language, then **Analyze**.
3. The list shows the layers in reveal order: shapes, arrows and read labels, the order following the arrows. **Create the document** turns it into a Vectorlab document with named layers.
4. **File** › **Send to › Animation** opens the Studio on a ready graph: an Animation node with one element per visible layer, ready for **Choreograph**, and a local, free motion code render.

Good to know

A difficult page (handwriting, low contrast): **Try with Claude** shows its price, then asks for confirmation. A paid reading is kept: picking it up again costs nothing.

A draft in which 12 words in a row repeat the capture is refused; its closeness to the capture is shown. The captured image is never the post's visual: at most it serves as a composition reference, which the recipe records.

Text read from a page or a capture stays text: a sentence like "ignore your instructions" is transcribed, never followed.

MAKE VIDEOS

Sound, voices and visual effects

You will generate and rework music, split a track into instruments, record a voiceover, create sound effects and particles, transform a video with AI and stack post-processing effects. Each generator's price is shown before you run it; particles, the Workshop and local post-processing are free.

The Sound & VFX screen

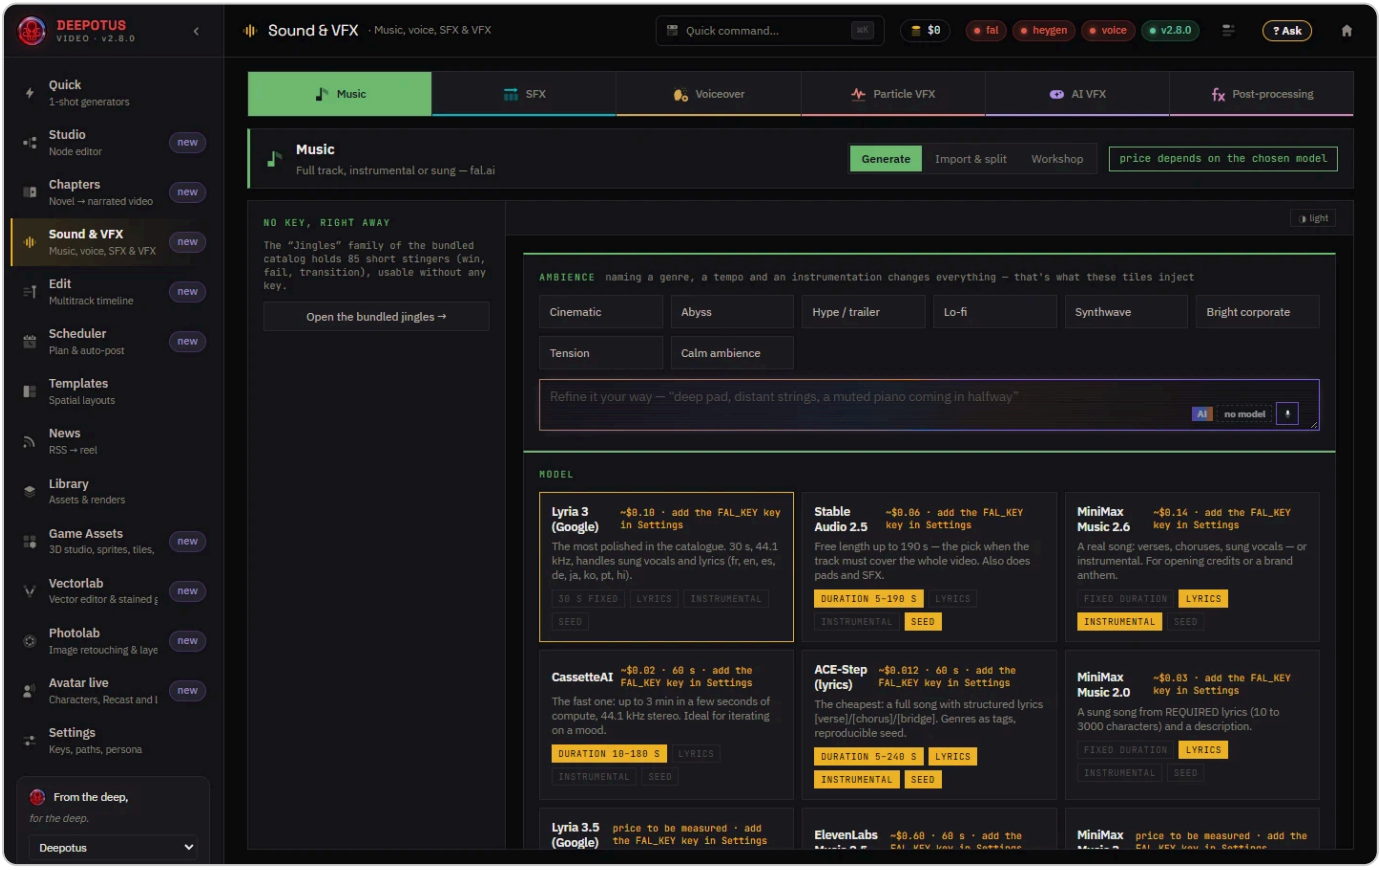

1. One section per kind of sound or effect 2. Music: generate, import and split, or the Workshop 3. The price is shown before you run

The Sound & VFX bar: Music, SFX, Voiceover, Particle VFX, AI VFX and Post-processing, each with its color.

Open **Sound & VFX** in the rail. A bar of six sections, each with its own color, leads to a page: **Music**, **SFX**, **Voiceover**, **Particle VFX**, **AI VFX** and **Post-processing**. The Music page has three modes: **Generate**, **Import & split** and **Workshop**.

Split a track into instruments

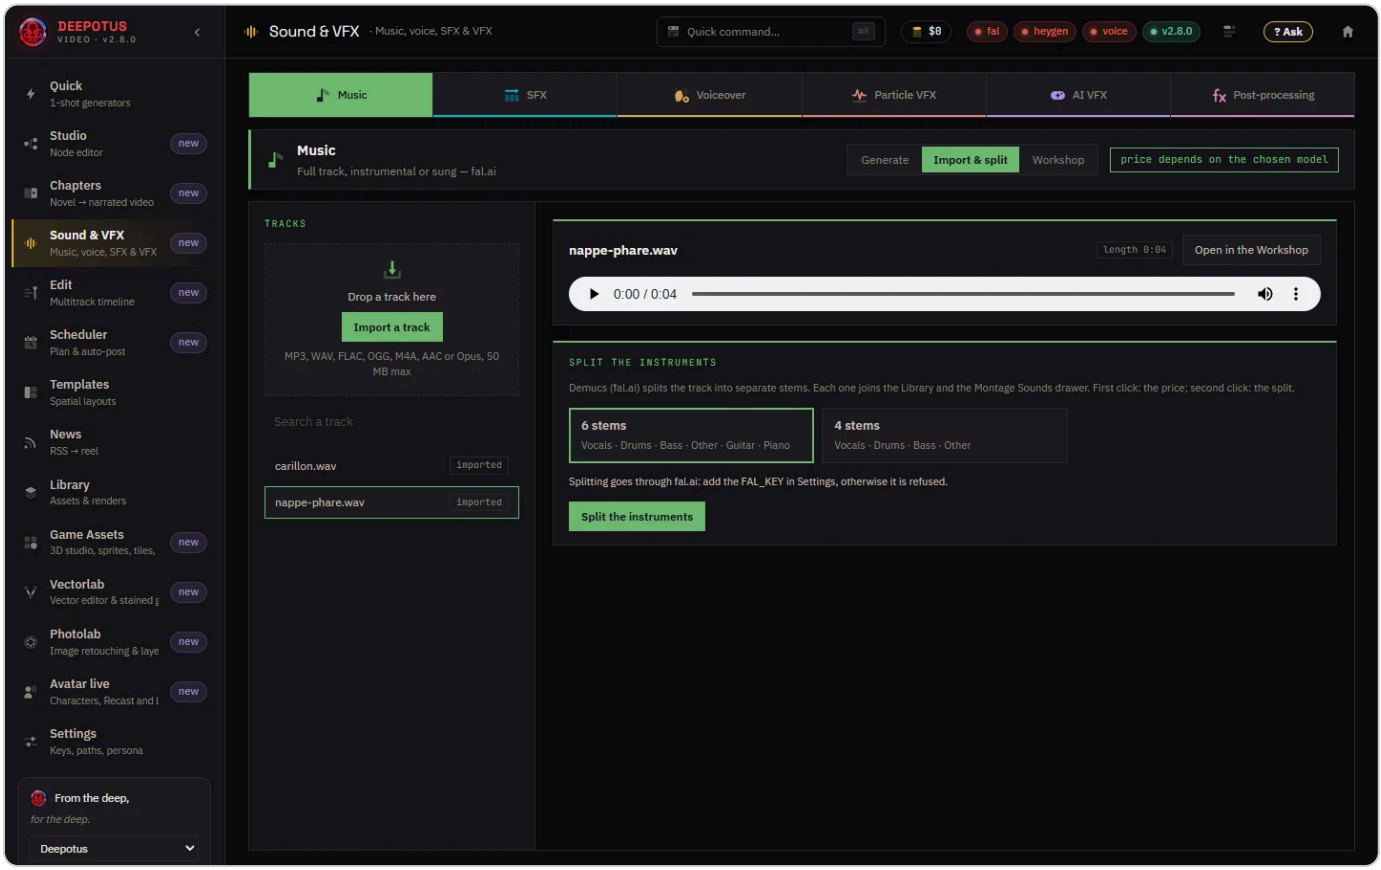

1. Drop or import a track 2. The chosen track plays 3. First click: the price; second click: the split

Music › Import and split: the chosen track, then the split into 4 or 6 instruments (paid, via fal).

1. In Music, **Import & split** mode, drop an audio file or click **Import a track**. The sounds you already imported are listed below.
2. Pick the track: it plays. Choose 4 or 6 stems (vocals, drums, bass, other, guitar, piano).
3. **Split the instruments**: the first click gives the price, the second runs it (paid, via fal). Each stem joins the Library and the editor's Sounds drawer.

Rework a sound in the Workshop

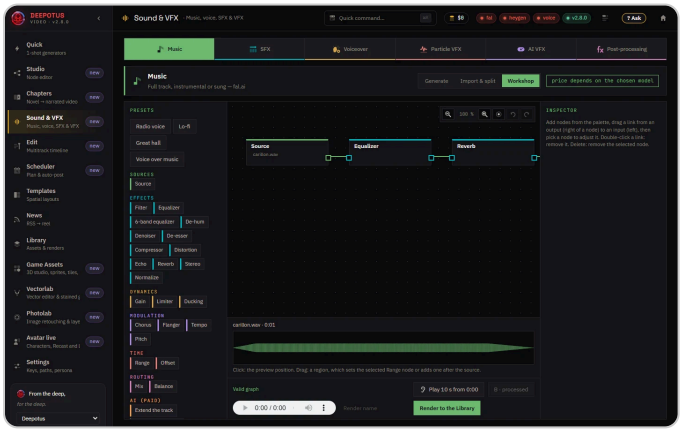

1. A preset to start from

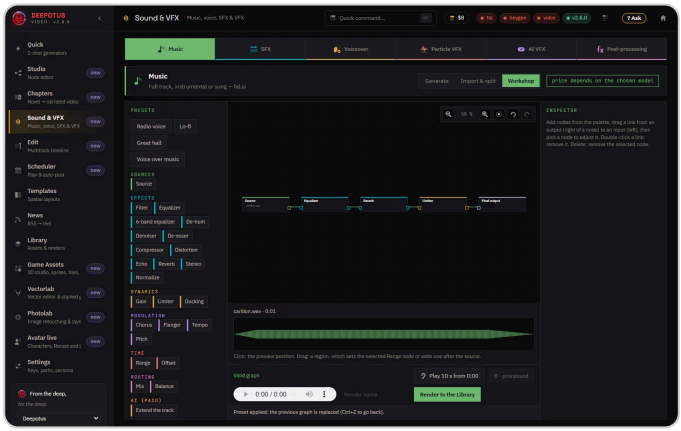

2. Source, effects, output: linked left to right 3. AI nodes rework a range

Music › Workshop: a preset lays out a chain of nodes; each node is set in the inspector.

1. **Workshop** mode, or **Open in the Workshop** from a track or a voice: the sound becomes the graph's source.
2. Start from a preset (**Radio voice**, **Lo-fi**, **Great hall**, **Voice over music**) or add nodes from the palette: equalizer, compressor, echo, reverb, ducking, tempo, pitch…
3. Link a node's output to the next one's input, and set it in the inspector. Under the canvas, the strip shows the waveform: drag to pick a range.
4. Listen to ten seconds, then **Render to the Library**.

The **AI (paid)** nodes extend the track, regenerate or vary a range, or add a sound effect to it. Their price is shown; a generated range is kept and is not paid again as long as nothing changes before it.

Record a voiceover

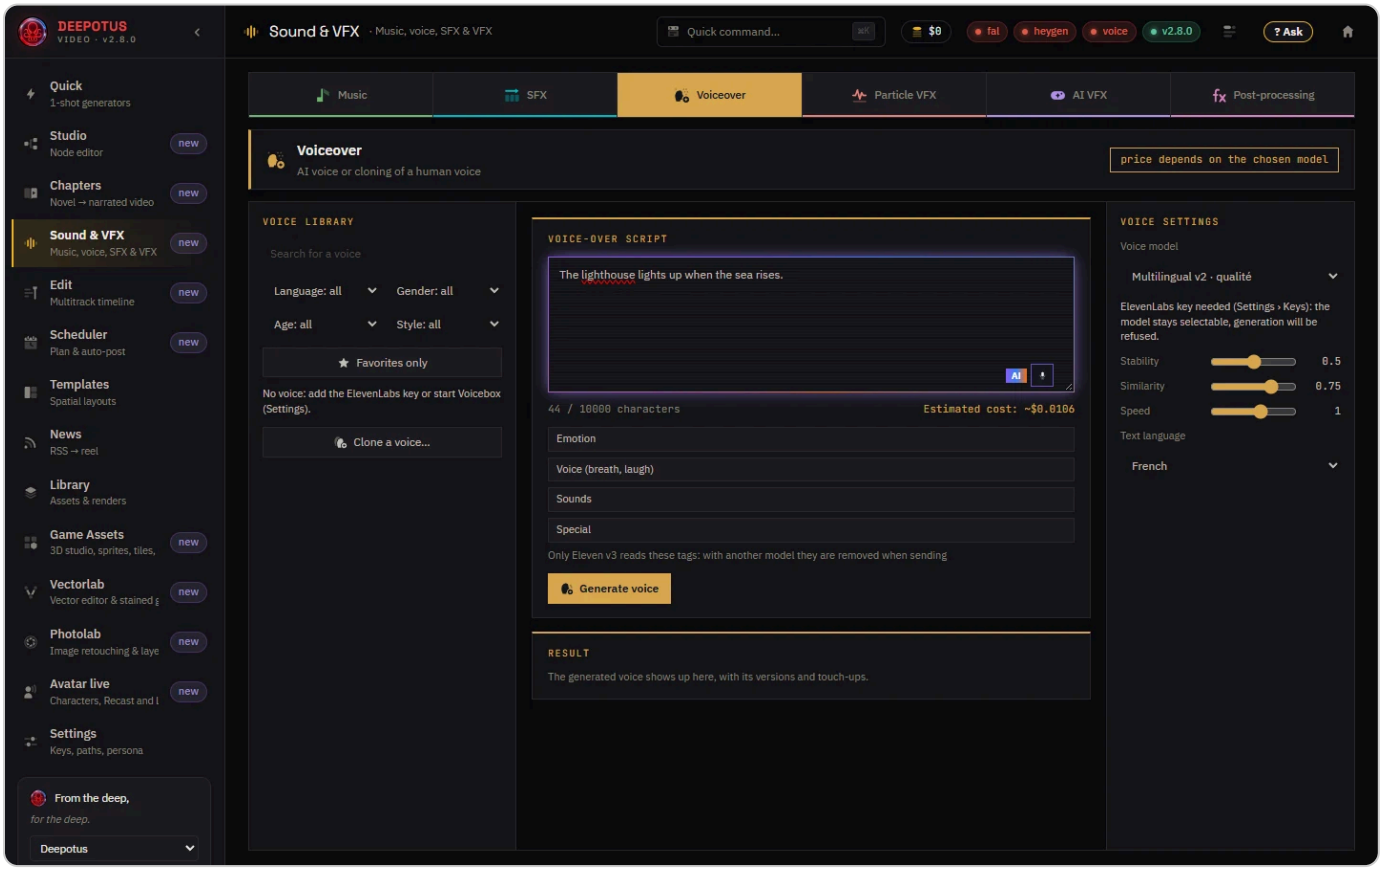

1. Search, listen to, clone a voice 2. The script and its character count 3. Generate: versions stack up below

Voiceover: the voice library on the left, the script in the middle, the model settings on the right.

1. On the left, the voice library: search, filter, listen, keep your favorites, or **Clone a voice...** from your own takes.
2. In the middle, write the script. With the Eleven v3 model, emotion or acting tags are inserted at the cursor. On the right, the model and its settings.
3. **Generate voice**. Each edit creates a version: **Trim the take**, **Fade in and out**, **Normalize**, **De-ess**, **Ducking under music**, **Enhance the take** or **Isolate the voice**. Then **Send to the Editor**.

Sound effects: pick a model

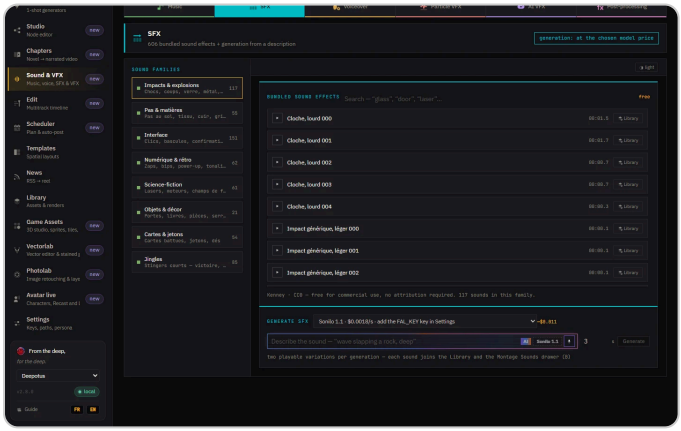

1. The model, grouped by route, with its price

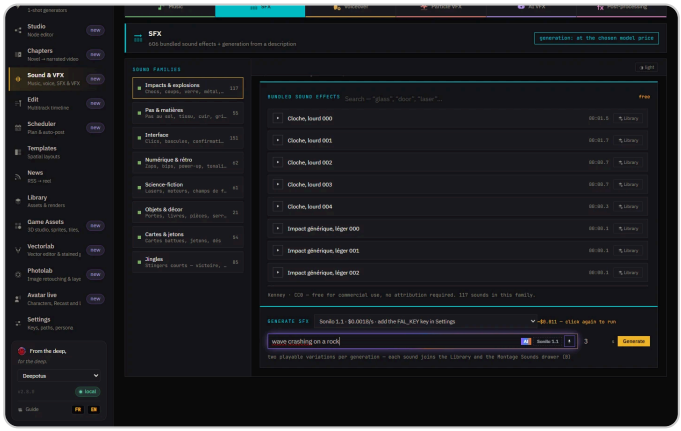

2. First click: the price; the second runs it

SFX, Generate SFX: the model and its price; a model via fal runs in two clicks, the first gives the estimate.

The **SFX** section keeps the bundled catalog and the **Generate SFX** card. Its **Model** list offers ElevenLabs direct or via fal, Sonilo, Mirelo and Stable Audio, each with its price and maximum length. Describe the sound, set the length (0 lets the model choose), then **Generate**: two variants go to the Library. A model via fal runs in two clicks, the first shows the price.

Add sound to a video

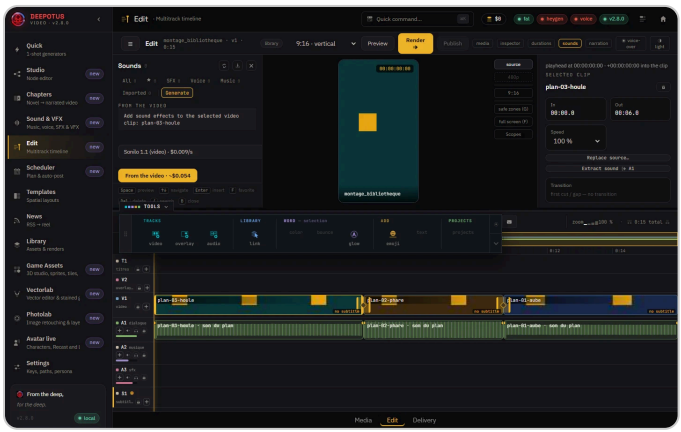

1. From the video: the selected clip

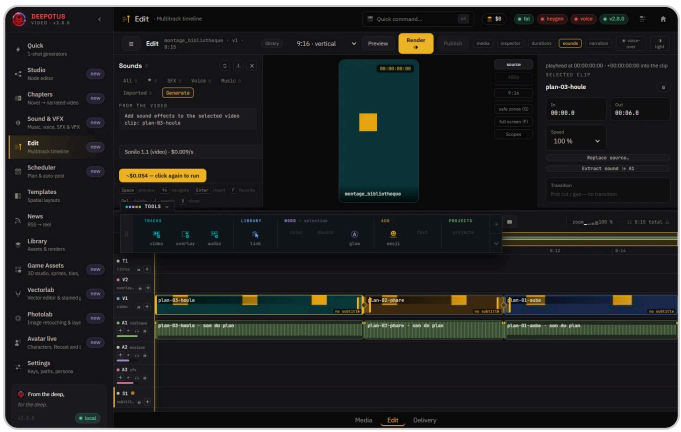

2. The estimate first; the second click runs it

Montage, clip selected, sounds chip › Generate › From the video: the model watches the picture and returns a sound effect.

1. In the Montage, select a video clip on the timeline, then open the Sounds drawer and its **Generate** tab.
2. The **From the video** block names the clip. Pick the model: Sonilo, Mirelo, Kling, Hunyuan Foley or ThinkSound; some ask for a description of the sound.

3. A first click gives the estimate, the second runs it. The sound effect goes to the Library; if the model returns the video with sound, it goes there too, under Renders.

Live particles

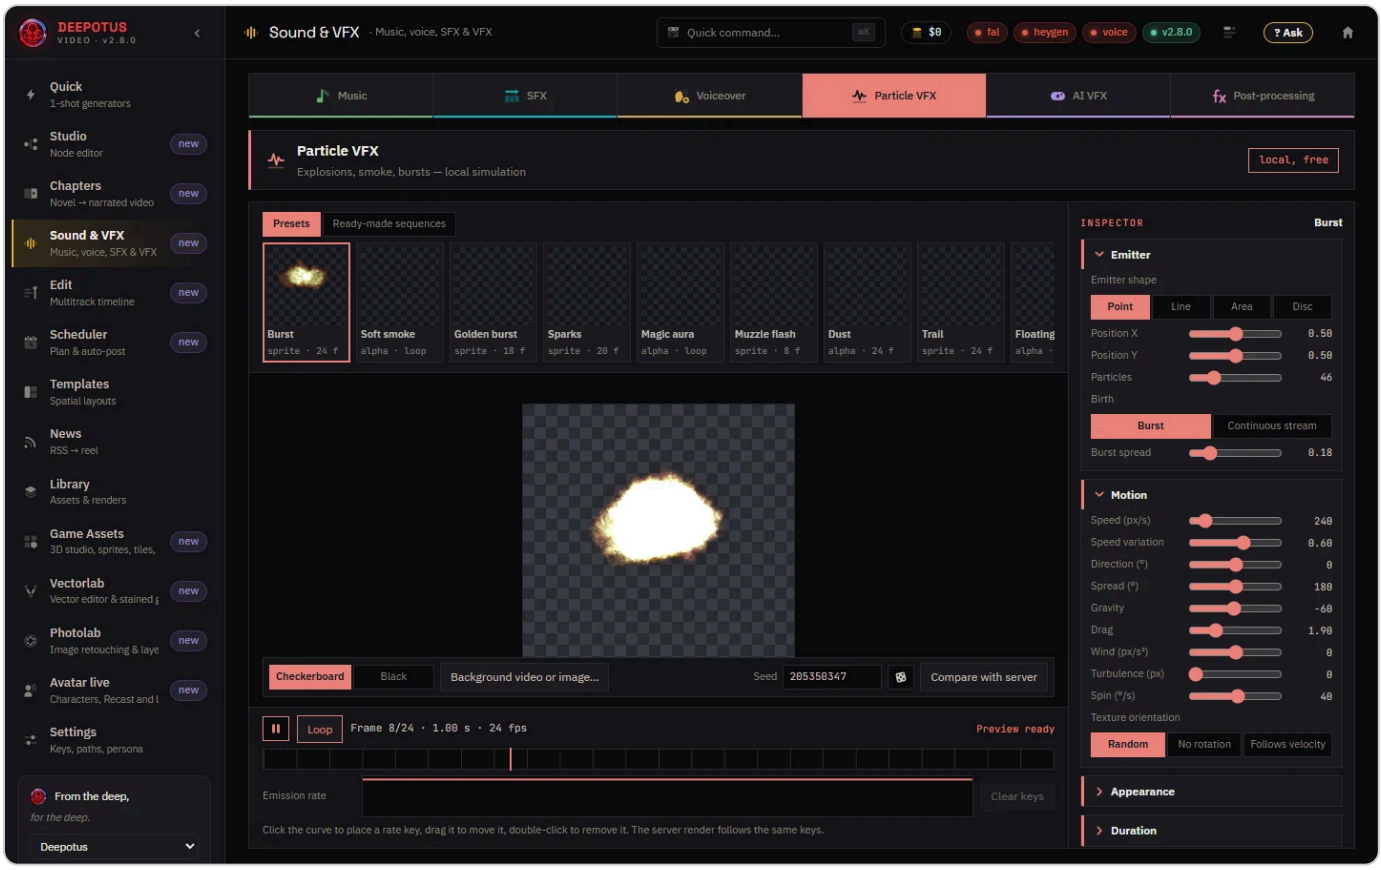

1. Presets and ready-made sequences 2. The preview, computed on your PC 3. Every setting changes the preview

Particle VFX: a preset, the preview that plays right away, and the inspector to tune the emitter.

1. **Particle VFX** section: pick a preset (explosion, smoke, sparks, magic…) or a ready-made sequence.
2. The preview plays right away, computed on your PC. Set the emitter, the motion and the look in the inspector: every setting changes the preview. The same seed always gives the same animation.
3. Export a **PNG image** or **Export as WebM** with transparency: the render goes to the Library.

AI VFX: relight, remove, reframe…

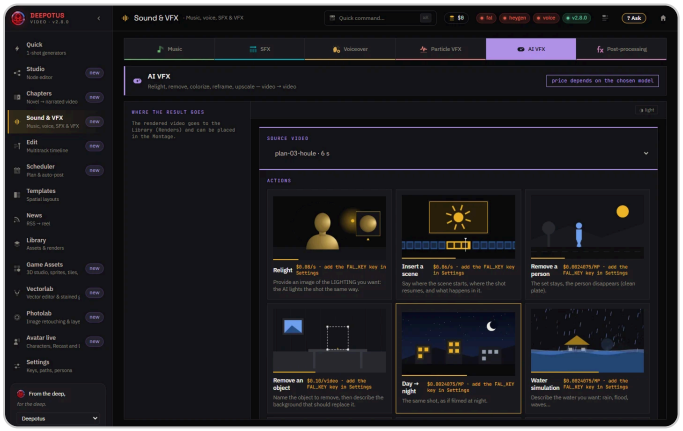

1. The source video 2. An action and its price

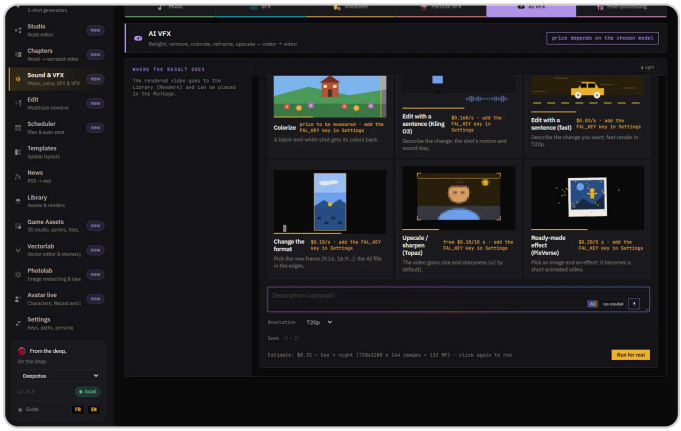

3. The estimate first; the second click runs it

AI VFX: the source video, illustrated actions with their price; Estimate, then run gives the estimate first.

1. **AI VFX** section: pick the source video among your finished renders.
2. Pick an action. Each card shows in motion what it produces (hover it), its price and what it needs.
3. **Estimate, then run**: the first click shows the server estimate, the second, **Run for real**, runs it. The rendered video goes to the Library.

Stack post-processing effects

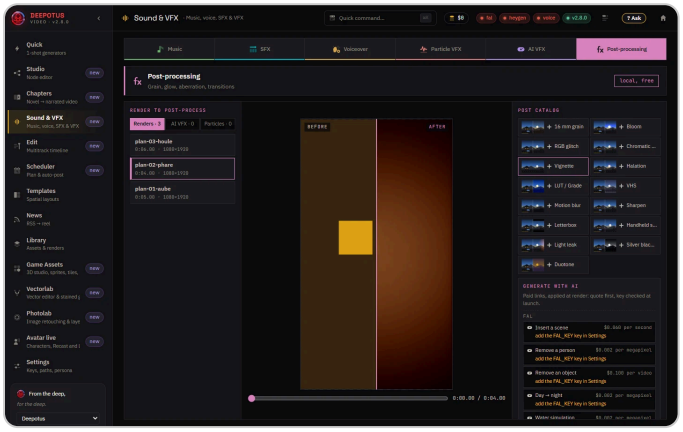

1. The render to retouch 2. A preset joins the stack 3. Before / after at the chosen moment

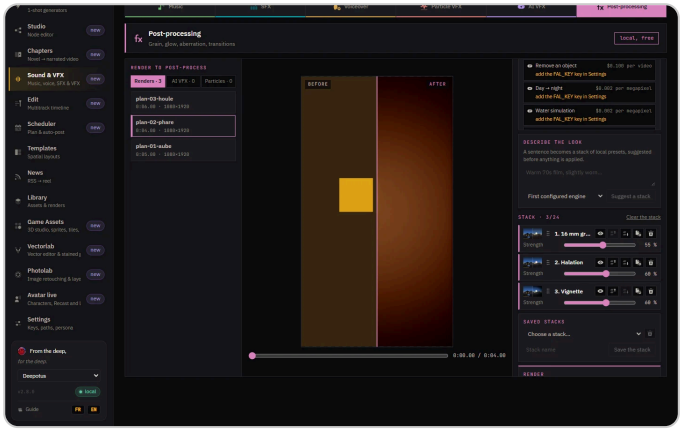

4. Drag, hide, set the strength 5. Rendered on your PC, free

Post-processing: a render, presets stacked in order, the before/after, then the free render.

1. **Post-processing** section: pick a render, an AI VFX result or particles.
2. Click presets in the catalog (16 mm grain, bloom, glitch, LUT, VHS, cinema bars, duotone…): they stack up in order.

3. In the stack, drag a link to move it, hide it, set its strength. The before/after shows the effect at the chosen moment.

4. **Render the video**: the render runs on your PC, for free, and goes to the Library. **Save the stack** keeps it for another video.

Generate with AI, describe the look

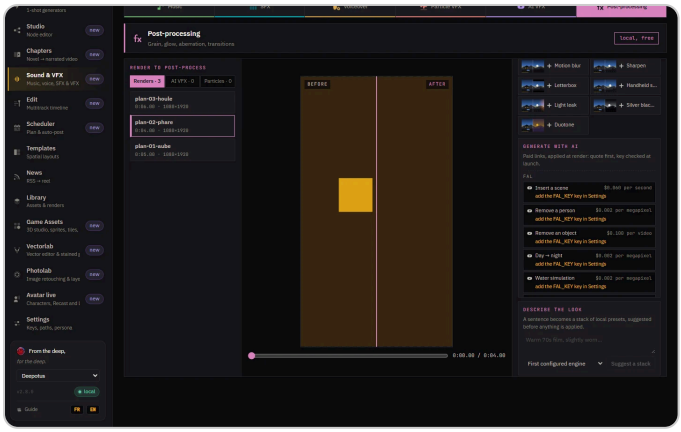

1. By provider, with their price 2. No key: selectable, refused at launch

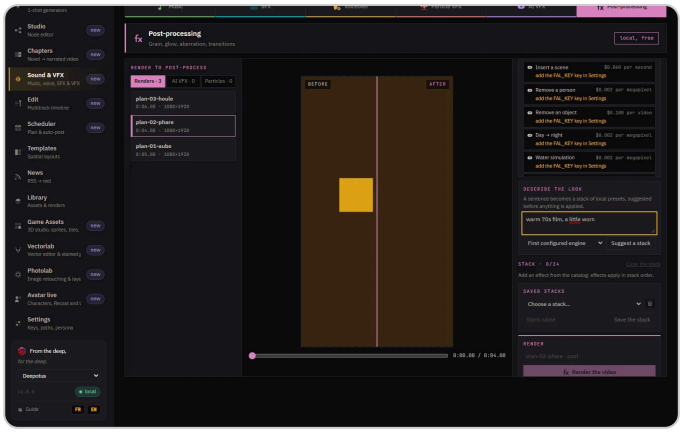

3. One sentence: a suggested stack, never applied on its own

Post-processing: the AI generators by provider, with their price and the key they need; “Describe the look” suggests a stack.

The **Generate with AI** section sorts the generators by provider (fal, Decart, ElevenLabs): style, upscaling, frame interpolation, sound effects. Each becomes a paid link of the stack, with its price; without its key it stays selectable and you are told. Rendering a paid stack takes a second click, after the estimate.

In **Describe the look**, write a sentence ("warm 70s film, a little worn") then **Suggest a stack**. The suggested stack is never applied on its own: **Replace the stack** or **Add to the stack**.

The Sounds drawer

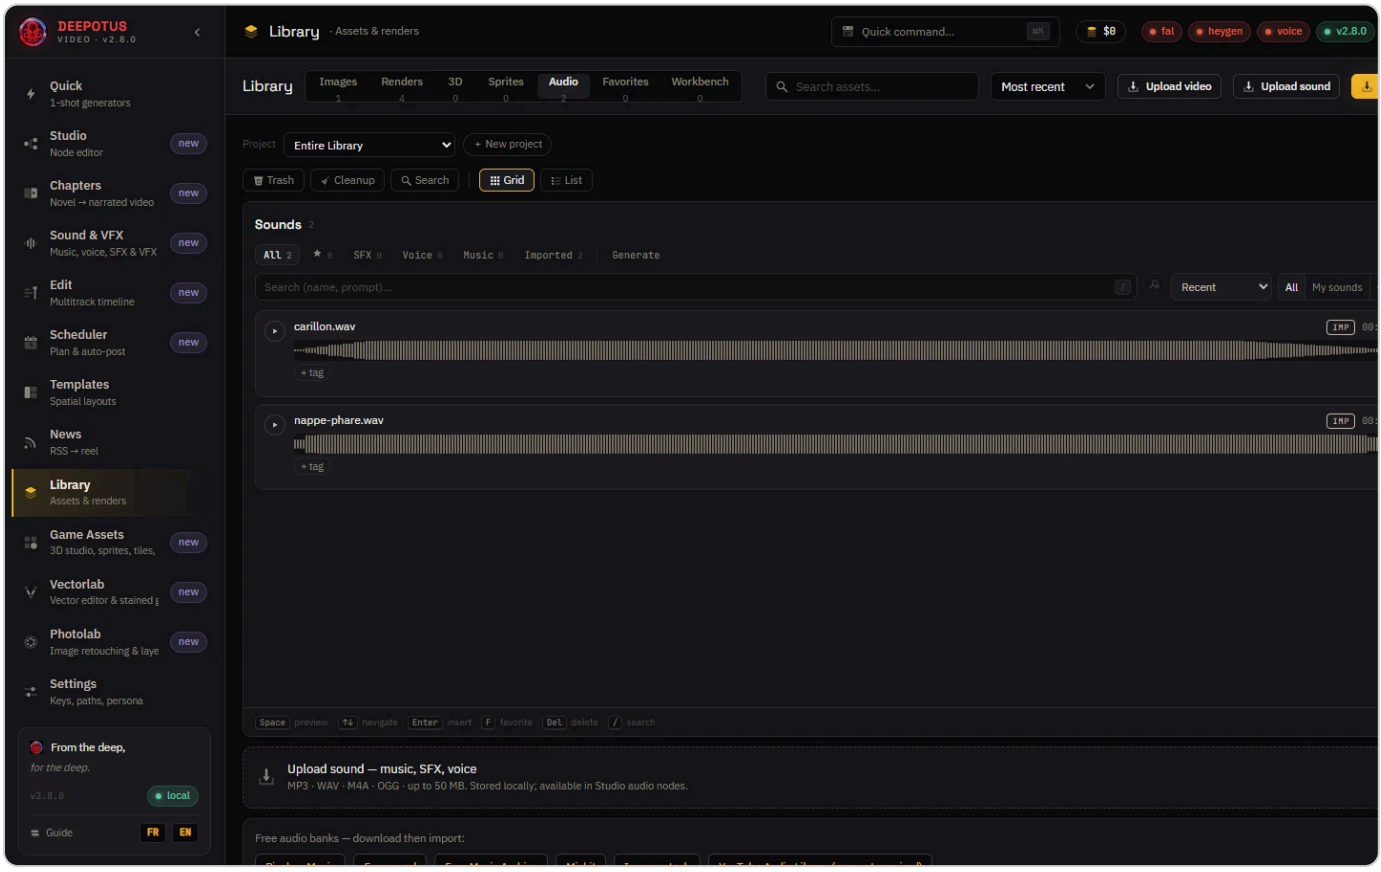

1. All sounds, then by kind: SFX, voice, music 2. Each sound: listen, tag, enhance, split

Library, Audio chip: the Sounds drawer, with its search, filters and actions.

All your sounds live in the Sounds drawer: in the Library (Audio tab) and in the editor (`B` key). Filter by kind, between your sounds and the bundled catalog, search by name or tag.

Enhance

Equalizer, noise reduction, compressor and normalized loudness in one click, on your PC, no provider.

Search by description

Describe what you want to hear. Needs the local search service, otherwise the button says why it is greyed out.

Preview

Hover a sound to hear it (option to turn on); `Space` plays the selected row.

Good to know

Split, AI nodes, AI VFX, AI links: a first click arms the action and shows the price, a second one starts it. Nothing goes out without that second click, and everything stays under your spending caps.

The animated thumbnails of the actions and presets play on hover. If your system asks to reduce motion, they stay still.

MAKE VIDEOS

Play an avatar, deferred or live

You will learn to create a Character from a few photos, have that Character replay a video in another setting, and film yourself on your webcam transformed live, with their voice. Twenty minutes. Everything is billed per second: the price shows before each launch and counts against your caps.

Before you start

The Avatar live screen opens from the rail. It needs keys depending on what you do: **fal** for Recast, the setting and Live, **ElevenLabs** for the Character's voice. A **Decart** key is optional: it adds fast mode to Live. The free local voice needs an NVIDIA graphics card and Voixbox (see Install and train a local RVC voice).

Only clone your own face and voice, or those of someone who has agreed to it. The app makes you confirm this for every Character and keeps the date.

Create a Character

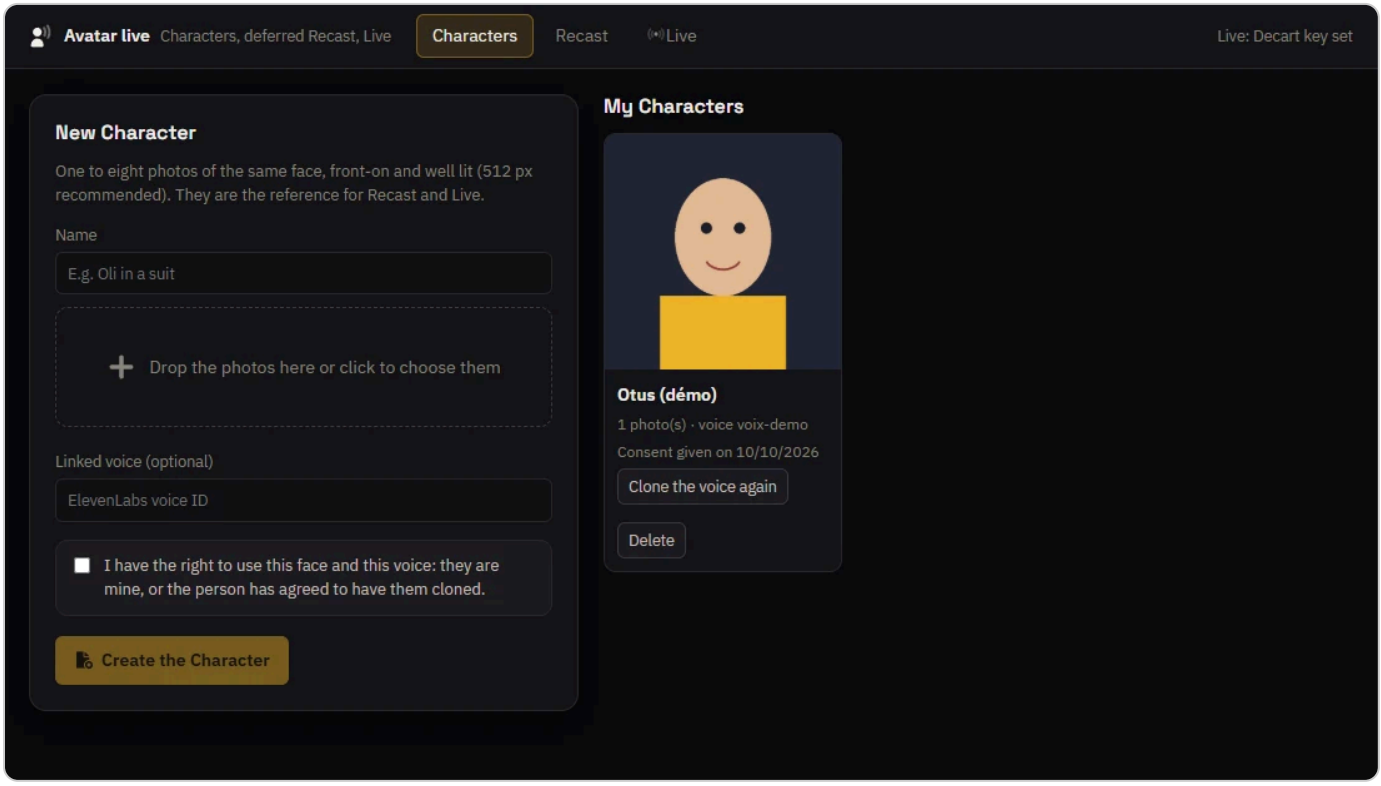

1. One to eight photos of the same face 2. Consent, dated, for every Character 3. Your Characters, ready for Recast and Live

Avatar live › Characters: reference photos, consent, then the list of your Characters.

1 Characters tab: give the Character a Name.

2. Drop one to eight photos of the same face, front-on and well lit (512 px recommended). They are the reference for Recast and Live.
3. Tick the consent box (“I have the right to use this face and this voice…”) then click **Create the Character**.
4. Optional: **Clone the voice** with one to five clips of the voice alone (30 s to 3 min in total). Cloning goes through ElevenLabs.

Recast: replay a video

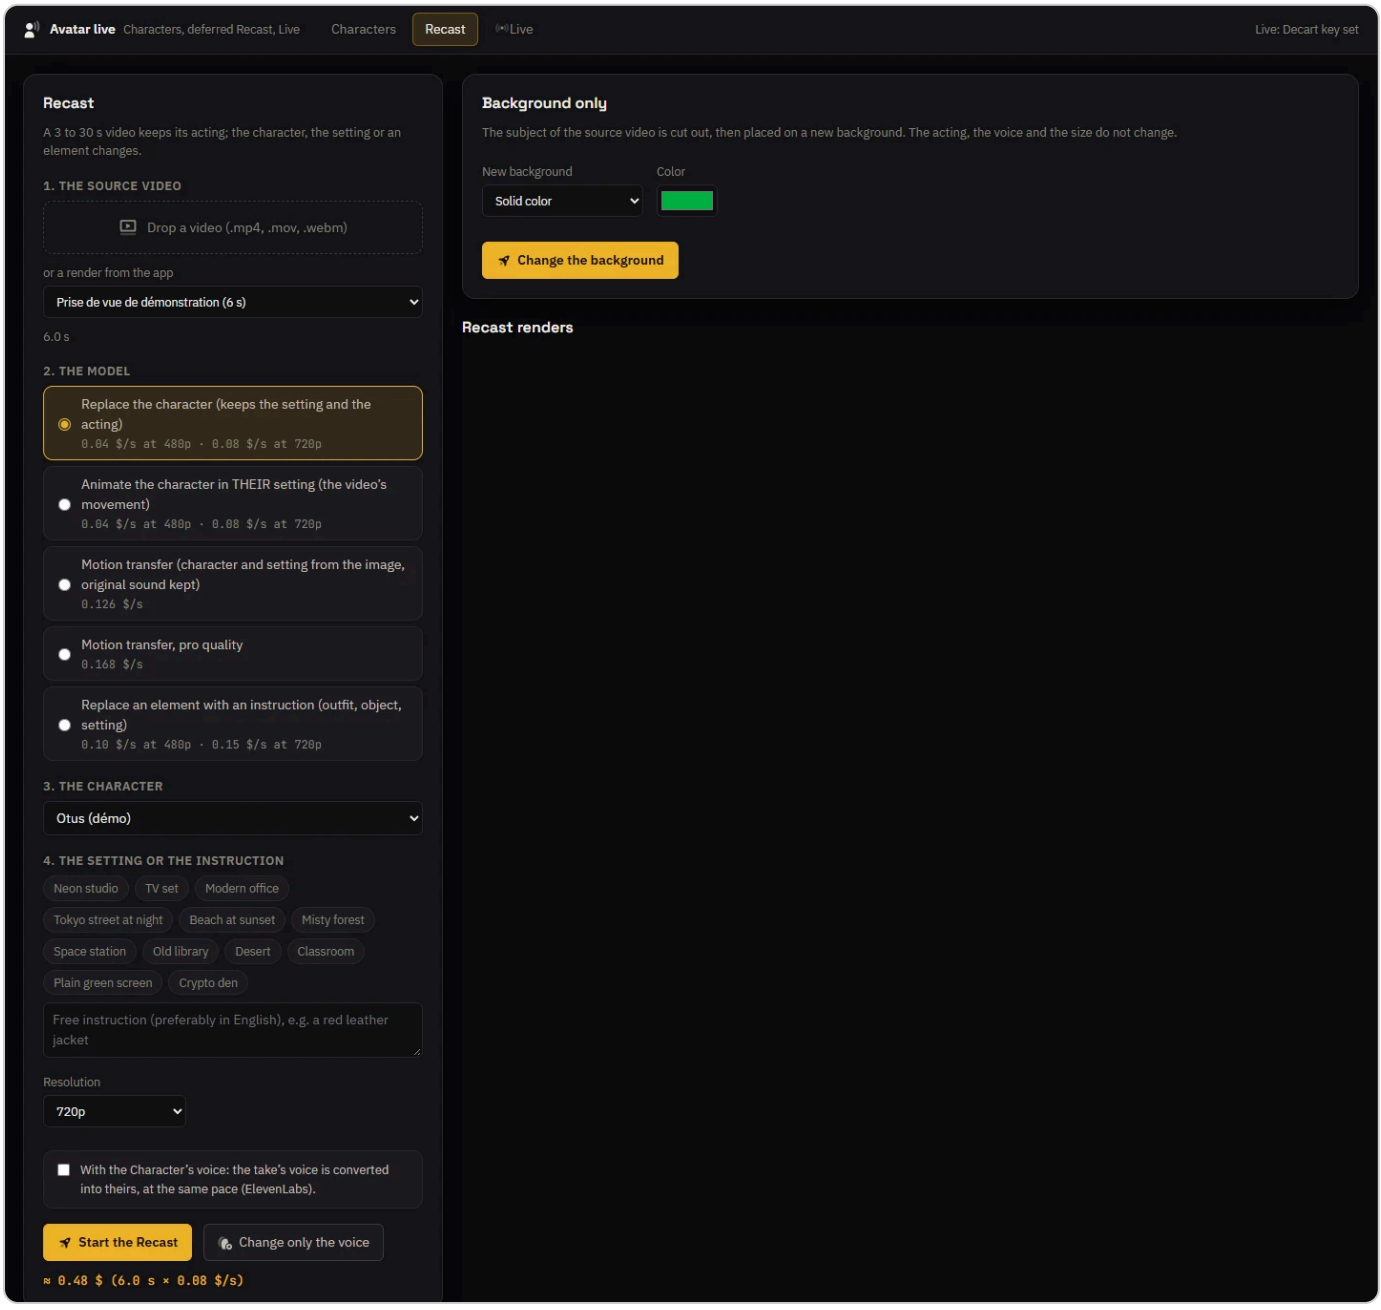

1. Each model shows its price per second 2. A ready-made setting, or your own instruction 3. The quote first, then the launch

Recast: a source video, a model and its price, the Character, the setting; the quote shows before you start.

A 3 to 30 s video keeps its acting; the Character, the setting or an element changes. The render takes a few minutes and joins your renders.

1. **Recast** tab: drop the source video, or pick a render from the app.
2. Choose the **Model**: its price per second is shown next to it.
3. Choose the Character, then a setting from the **Presets** or a free instruction (preferably in English).
4. Tick the Character's voice if you want the take's voice converted into theirs, at the same pace.
5. Read the quote under the button, then **Start the Recast**.

To try at half price, click **480p draft**: if you like it, “Finalize in 720p” on its card reruns the same Recast with the same seed. The final render is very close to the draft, without being identical. The motion transfer models have no draft.

Only one person is replaced per take. For the setting, **Suggest an instruction from the take** has the vision AI connected in your keys describe the place (never the people); read the instruction before you start.

Live: your webcam, transformed

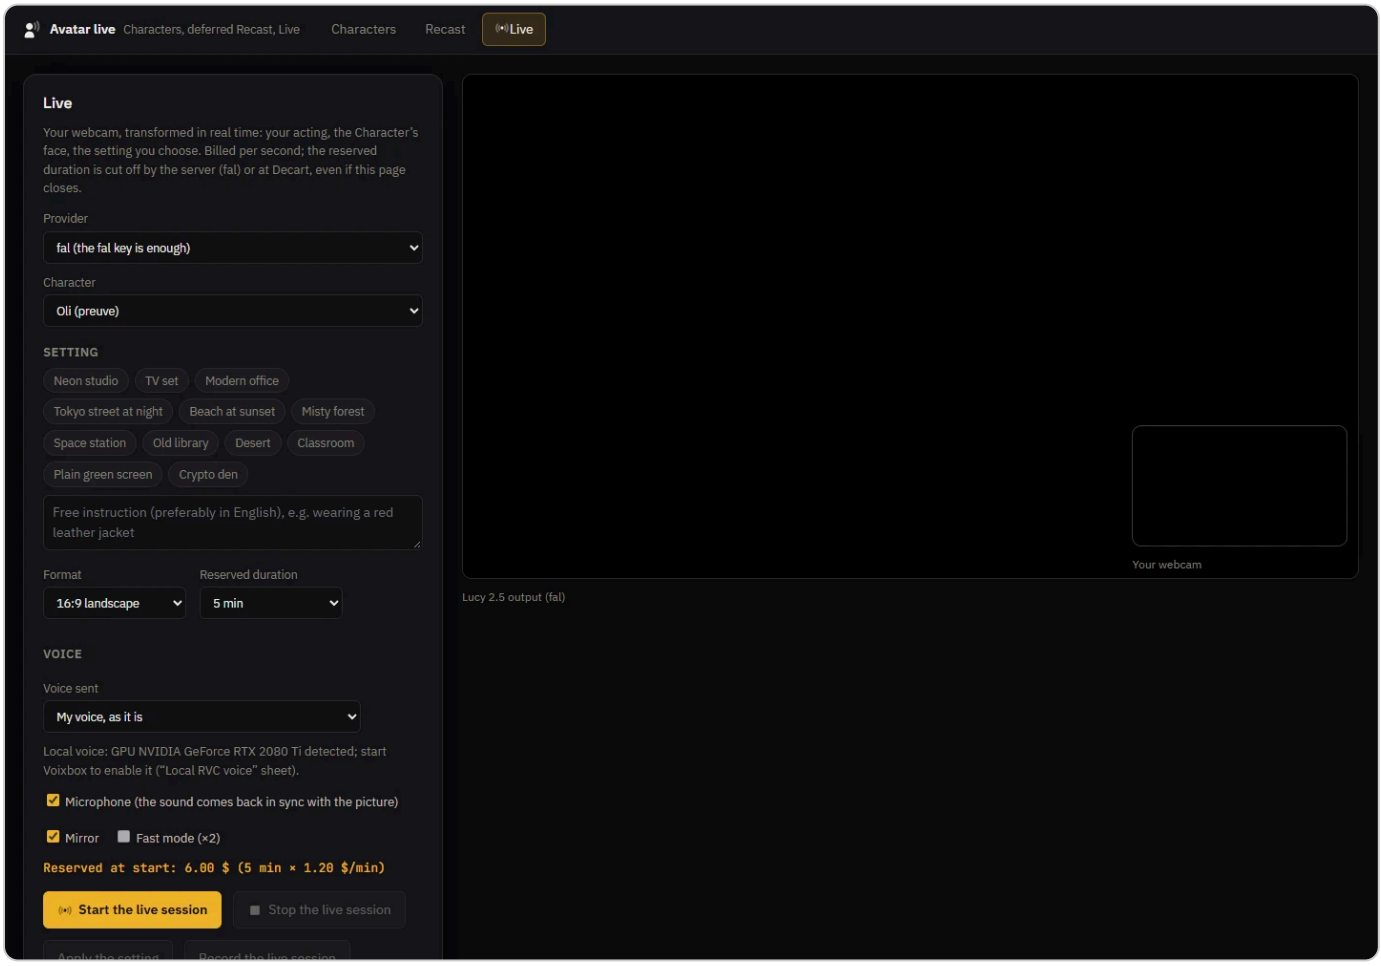

1. fal by default: the fal key is enough 2. The setting, switchable mid-session 3. Your voice, the Character's, or a local RVC voice 4. The reserved price is shown before you start

Live: provider (fal by default, the fal key is enough), Character, setting, format, reserved duration and voice sent; the reserved price shows before you start.

1. **Live** tab: the **Provider** is fal as soon as your fal key is set (Decart if you have its key). Paste a missing key here, then **Save the key**: it goes to the vault, the page never sees it.
2. Choose the **Character** (or keep your face), a **Setting**, the **Format** and the **Reserved duration**.
3. Choose the **Voice sent**: your own as it is, the Character's through ElevenLabs (about 2 s delay), or a local RVC voice (under a second).
4. **Start the live session**. The transformed picture appears next to your webcam; switch settings on the fly with **Apply the setting**.
5. **Record the live session** keeps the footage in your renders, ready for the Editor or a Recast.

When a voice is converted, the picture is delayed by as much as the voice: lips and sound arrive together. The live bar shows the converted segments, the latency and any losses.

Going further

Background only

The subject of the video is cut out, then placed on a color, a Library image or a video. **Check the cut-out** first shows one picture on the new background; the video only goes after that.

Change only the voice

Keeps a video's picture and replaces its voice with the Character's. Billed at the ElevenLabs character equivalent.

In the Studio

The Recast and voice conversion nodes plug into an existing render or a UGC video. The graph quote counts their price before **Run**; the Recast node also has its 480p draft.

On the phone

The companion app starts a Recast (480p draft included, then its final version), films a take or opens a Live session, at your PC's prices and under the same caps.

Good to know

Live is reserved up front, by the minute; only the time actually played is billed. With fal, the app holds the connection: it cuts off at the reserved duration, cuts off as soon as the page closes, and measures the billed seconds itself. With Decart, the cut-off happens at Decart.

A tab hidden for 60 s stops the live session, so you don't pay for nothing. Prefer Edge or Chrome: they can delay the picture to match the voice.

MAKE VIDEOS

Install and train a local RVC voice

You will learn to install Voixbox, record your voice, train it with RVC and have it speak during a Live session, on your own graphics card. Free: nothing goes to the Internet. Allow thirty minutes to install, a quarter of an hour to record and about an hour of computing per voice.

Before you start

You need an **NVIDIA** graphics card with 6 GB of memory or more, and about 6 GB of disk. Without one, Live offers your voice as it is or the Character's voice through ElevenLabs (paid, nothing to install).

The engine is **RVC** (MIT license), through its official WebUI: the same install trains a voice and makes it speak. Only clone your own voice, or that of someone who has agreed to it.

Install Voixbox (once)

1. Open PowerShell in the app folder, where `tools\voixbox` lives (by default `%LOCALAPPDATA%\DeepotusVideoGen`).
2. Run `powershell -ExecutionPolicy Bypass -File tools\voixbox\installer.ps1`.
3. The script installs `uv`, downloads a verified version of the WebUI, creates its Python, adds torch for CUDA (about 3 GB) then its weights (560 MB).
4. It ends with `CUDA True | NVIDIA …`: the card is detected. On any failure, it stops and says why.

Record your voice

The quality of the cloned voice depends almost entirely on this recording.

1. Record 10 to 15 minutes of your voice alone: no music, no second voice, no echo.
2. A quiet room, the mic 15-20 cm away, always the same mic. Speak the way you do on your live sessions, with varied sentences.
3. Avoid clipping. Prefer several 30 s to 2 min files, in WAV.
4. Put them in a folder that holds nothing else, for example `C:\Voix\oli`.

Train an RVC voice

1. Run `powershell -ExecutionPolicy Bypass -File tools\voixbox\webui.ps1`. The WebUI opens in French at `http://127.0.0.1:7865`; keep the PowerShell window open.
2. « Entraîner » (Train) tab, step 1: a short **experiment name** with no spaces or accents (`oli`), rate `40k`, pitch guidance **Oui** (yes), version `v2`.
3. Step 2: your recordings folder, and the `rmvpe` pitch algorithm.
4. Step 3: save every `25` epochs, **200** epochs in total (100 for a first try), batch `8` (or `4` with 6 GB), small final model at each save: **Oui** (yes).
5. Click « Entraînement en un clic » (one-click training). Allow about 1 h 15 for 200 epochs on an RTX 2080 Ti.
6. Put the voice where Voixbox looks for it: `powershell -ExecutionPolicy Bypass -File tools\voixbox\ranger-voix.ps1 -Nom oli`.

Listen to it before using it: in the WebUI's « Inférence du modèle » (model inference) tab, pick `oli.pth` and give it a short recording. A crackling voice needs more epochs or a cleaner recording.

Speak with it during a Live session

1. Run `powershell -ExecutionPolicy Bypass -File tools\voixbox\lancer.ps1` and keep the window open: it reports how many voices are trained.
2. In **Avatar live**, **Live** tab, choose the local RVC voice under **Voice sent**.
3. Set **Transpose (semitones)** if needed: `+12` to go from a male to a female voice.
4. **Start the live session**. Each half-second of voice is converted in about 110 ms; the picture is delayed by 0.75 s to stay in sync with the lips.

Good to know

> The Live screen tells you what is missing: no NVIDIA card, Voixbox not started, or no trained voice. Another port is set with `VOIXBOX_URL` in the `.env`.

> ⚠ With Avast or an antivirus that inspects connections, always go through the scripts in `tools\voixbox`: they remove the `SSLKEYLOGFILE` variable for themselves only and read Windows certificates. Otherwise Python crashes (`no OPENSSL_Applink`).

MAKE VIDEOS

Edit a video

You will place your shots on a timeline, cut them, move them and chain them, with the gestures of a professional editor. Free: editing and previews run on your PC.

The editor screen

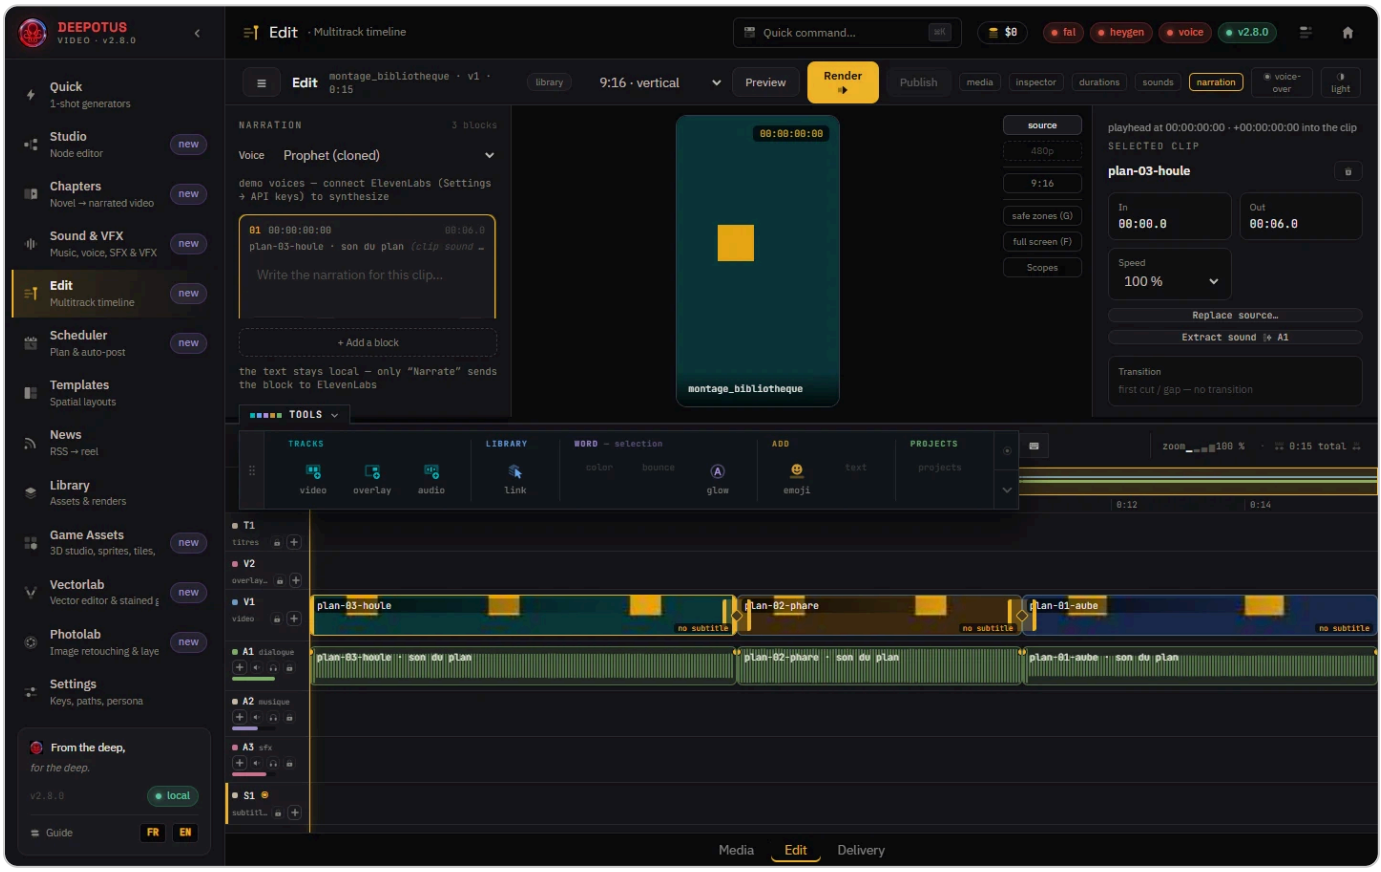

1. Shots on V1, their sound on A1 right below 2. Preview: a free 480p preview 3. Three views: Media, Edit, Delivery

The timeline at the bottom (video, sound, subtitles), the viewer in the centre, the inspector on the right; Preview, Render and Publish at the top.

At the bottom, the **timeline**: titles (T1), overlays (V2), main video (V1), then dialogue (A1), music (A2), sound effects (A3) and subtitles (S1). In the centre, the viewer; on the right, the inspector of the selected shot. At the top, **Preview** (free 480p preview), **Render** and **Publish**. The bottom bar switches between three views: **Media**, **Edit** and **Delivery**.

Place shots

1. Open the **media** drawer: all your finished video renders, with thumbnail, length and origin. Filter by origin (Studio, Chapters, Quick…), by title or by rating (three stars and up).

2. Click a media item to place it at the playhead. The **edit mode** decides how: overwrite, insert and push the rest, at the end, above, ripple overwrite, or fill the range.

3. From the Library, **Send to** › Edit also places a video on V1 with its sound on A1. A Studio render or a Chapters episode opens straight in the editor.

Cut and adjust

Cut

The blade (`Alt`+`C`) cuts the shot at the playhead. A gap is selected with one click and `Del` closes it.

Roll, slip, slide

`Alt` + drag the diamond moves the cut; `Alt` + drag the centre slips the content; `Shift` + drag moves the shot between its neighbours.

Range and markers

`I` and `U` set an in and out range; `Shift`+`X` cuts it on every track. `Shift`+`M` sets a marker, `Ctrl`+`M` lists them.

Everything undoes

`Ctrl`+`Z` and `Ctrl`+`Y` over fifty steps: shots, mix, tracks, subtitles, range, markers.

Cut at scene changes

Right-click a clip: the app spots each scene change and cuts it for you.

Shots too long

In the ☰ menu › View, the timeline highlights shots that run too long and jump cuts from the same source.

Chain two shots

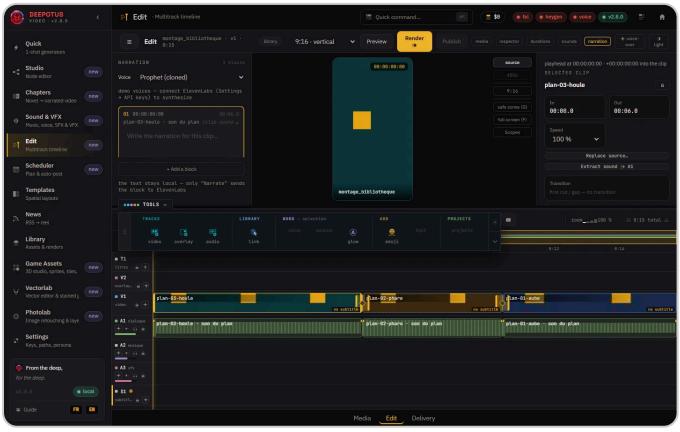

1. Click the diamond between two shots

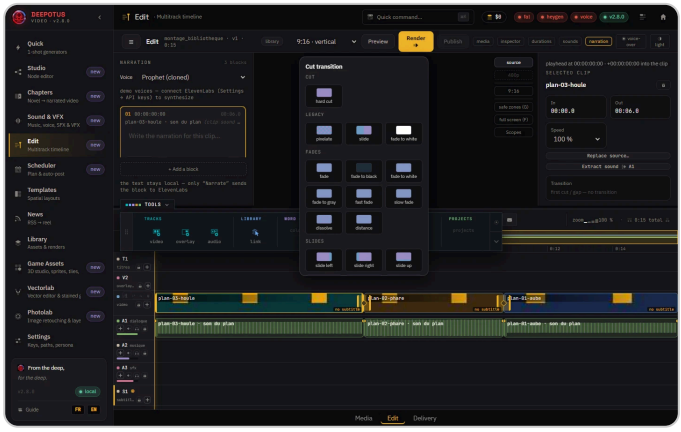

2. Pick the transition and its duration

A click on the diamond between two shots opens the transitions gallery and frame-by-frame trimming.

1. Click the **diamond** between two shots: the gallery opens, with 58 transitions sorted by family (fades, slides, wipes, zooms, pixels, shapes).
2. The same panel shows the last frame of the left shot and the first of the right one; shift the cut frame by frame.
3. For a quick fade on the selected shot: `Alt` + `T`. Simple fades play live; the others after a **Preview**.

Cut out with the model of your choice

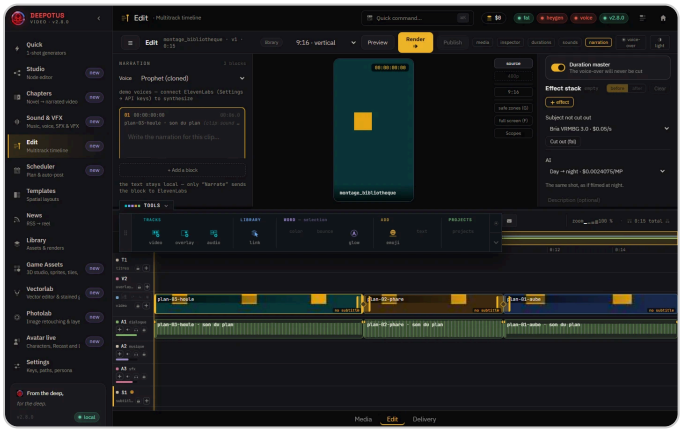

1. The cutout model and its price

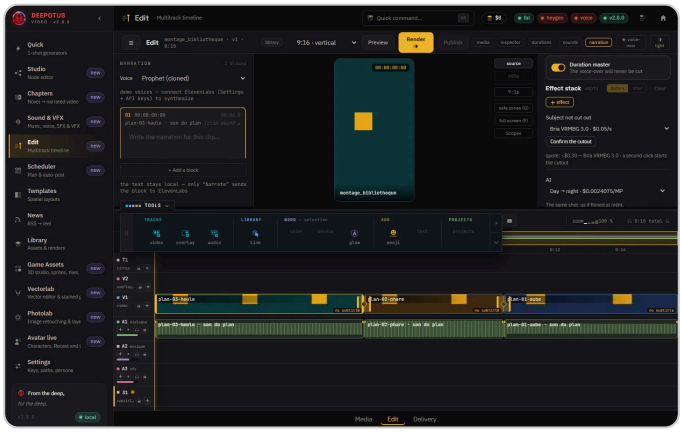

2. The estimate; the second click cuts out

Clip inspector, Effects on this clip: the cutout model and its price; the first click gives the estimate.

Select a clip: in the inspector, the clip's effects stack offers the subject cutout row, with its list of models and their price (Bria, VEED, Pixelcut, SAM 3.1 that tracks a named subject…). The first click on **Cut out (fal)** shows the estimate, the second cuts out. The cut-out subject then serves the effects placed **behind** it; **Cut out again** starts over with another model.

Organise your projects

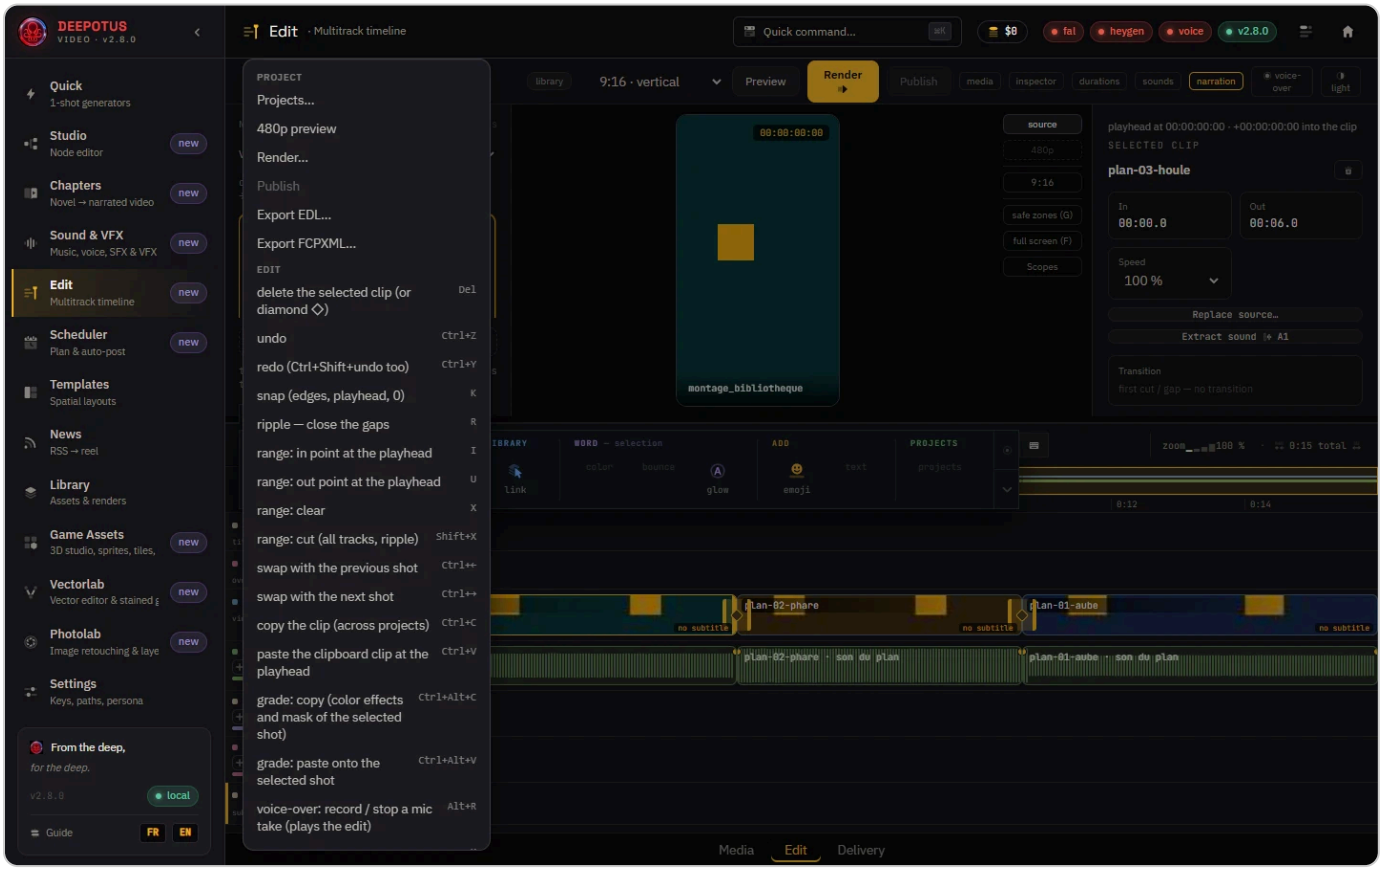

1. Project, Edit, Timeline, Markers, View, Help

The ☰ menu lists every action in six sections, each with its shortcut.

The ☰ menu lists every action in six sections, with its shortcut. **Projects…** opens the grid of your projects: open, duplicate, rename, delete, or compare two versions. Before opening another one, an unnamed edit is put away safely on its own.

New edit

An empty project, with its default tracks, and nothing else.

Copy between projects

`Ctrl` + `C` on a clip, `Ctrl` + `V` in another project.

Auto-clips

From a long video, AI suggests 1 to 8 excerpts of 15 to 60 s, then creates a project with the chosen one (paid; a free version without AI exists).

Your shortcuts

The `?` panel lists the shortcuts; a “Resolve” preset aligns them with DaVinci Resolve. Export them and import them back.

Shortcuts

Key	Effect
`Space`	Play / pause
`Alt` + `C`	Blade: cut at the playhead
`I` / `U`	Range in / out at the playhead
`Ctrl` + `Left` / `Right`	Swap the shot with its neighbour
`O`	Fold the tools bar
`B` / `T`	Sounds drawer / Narration panel
`G` / `F`	Safe zones / viewer full screen

MAKE VIDEOS

Grade, mix and add titles

You will give your shots a consistent colour, check them on the scopes, set their speed and framing, add titles and polish the sound. Free: everything is computed on your PC, except subtitle transcription and translation.

Grade a shot

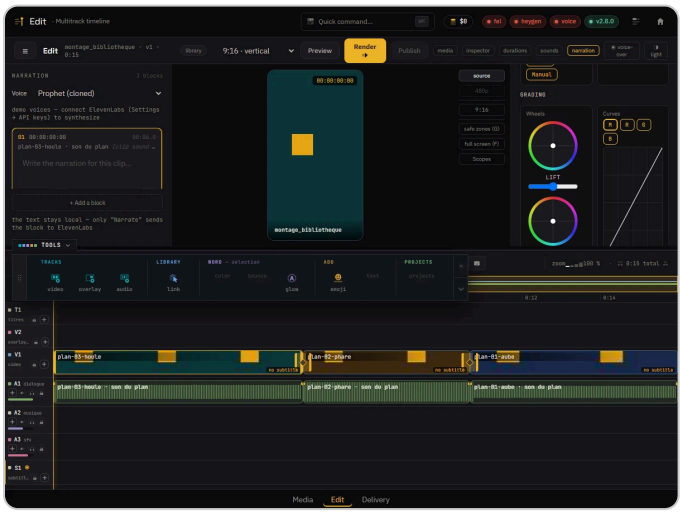

1. Shadows, midtones, highlights

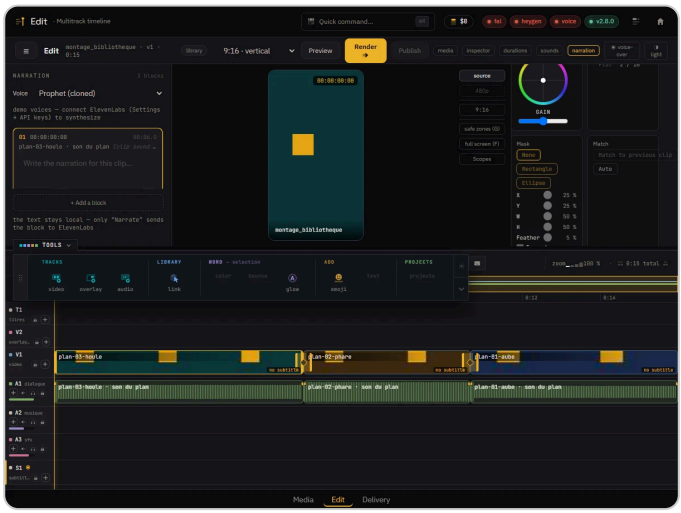

2. Match a shot to the previous one

Clip inspector, Grading panel: wheels, curves, mask, automatic match.

1. Select a shot: the inspector shows the **Grading** panel.
2. **Wheels**: shadows (lift), midtones (gamma) and highlights (gain). **Curves**: up to sixteen points, on the whole picture or each colour.
3. **Mask** limits the setting to a rectangle or an ellipse, feathered or inverted.
4. **Auto** removes a colour cast; **Match to previous clip** aligns the colours with the shot before.
5. **Copy grade** then **Paste grade** carry the setting over to another shot.

When playback stops, the viewer shows the shot's real graded picture; during playback, the picture stays raw.

Check on the scopes

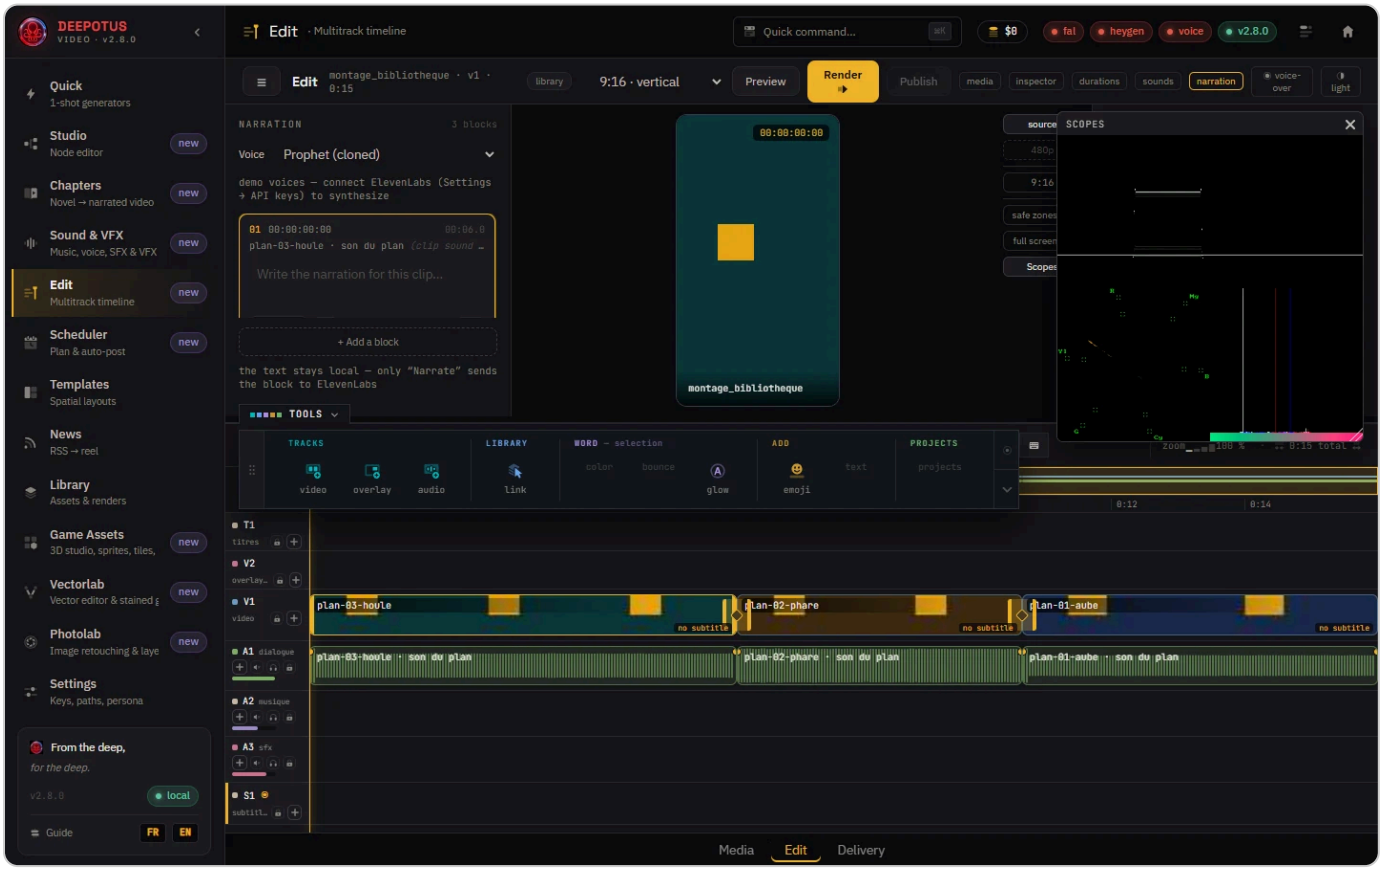

1. Scopes are computed when playback stops

Waveform, vectorscope and histogram of the graded picture, in a window you can move.

Scopes opens a movable window: waveform, vectorscope and histogram of the graded picture under the playhead, recomputed when playback stops. They account for speed, stabilisation and the render frame rate, and stay available in full screen.

Effects, speed and framing

The effects rack

Noise reduction, deflicker, debanding, sharpening, green-screen keying, monochrome, per-band hue and saturation, trails, lens distortion.

Speed and ramp

From 25 to 400 %. Interpolation smooths the motion; **Split at playhead** gives the rest of the shot another speed.

Stabilise and reframe

Stabilisation analyses a shaky shot; framing picks the part kept: centred, following the motion, or manual; dynamic zoom animates a reframe.

Adjustment clip

On track J1 (Shift+J), its effects apply to everything below, with fade in and fade out.

Overlays

Position, scale, rotation and opacity animated by keyframes, rounded corners, drop shadow.

Hand-drawn ramps

An effect ramp can follow a Bézier curve you draw; the shake's ramp raises or lowers its amplitude.

Titles and subtitles

1. Shift+T places a title card on track T1, burned in at render time.
2. Pick its template (full frame, lower third, caption, counter, chapter, quote, hashtag, call to action), its text, colour and font.
3. Subtitles have one track per language: only one is burned in, the others are exported. Translation can create a new language track.

Polish the sound

The audio rack

Six-band equaliser, noise reducer with adjustable floor, 50 / 60 Hz hum removal, stereo pan; listen to the rendered result.

Learn the noise

Set a range on a stretch of noise only: the app learns it and removes it at render time.

Voiceover from the mic

voice-over (Alt+R) records a take while the edit plays and places it on the dialogue track.

Narration

The narration panel (T) writes and synthesises the voice block by block; only “Narrate” sends a block to the provider.

Good to know

In the ☰ menu › View, the shot lightbox shows every V1 shot with its graded picture: one click goes to the shot.

Subtitle transcription and translation are paid; the rest of this chapter is free.

MAKE VIDEOS

Deliver: renders and formats

You will export your edit in the right format, at the right loudness, and send it to the Scheduler or to another program. Free: rendering happens on your PC.

Render the video

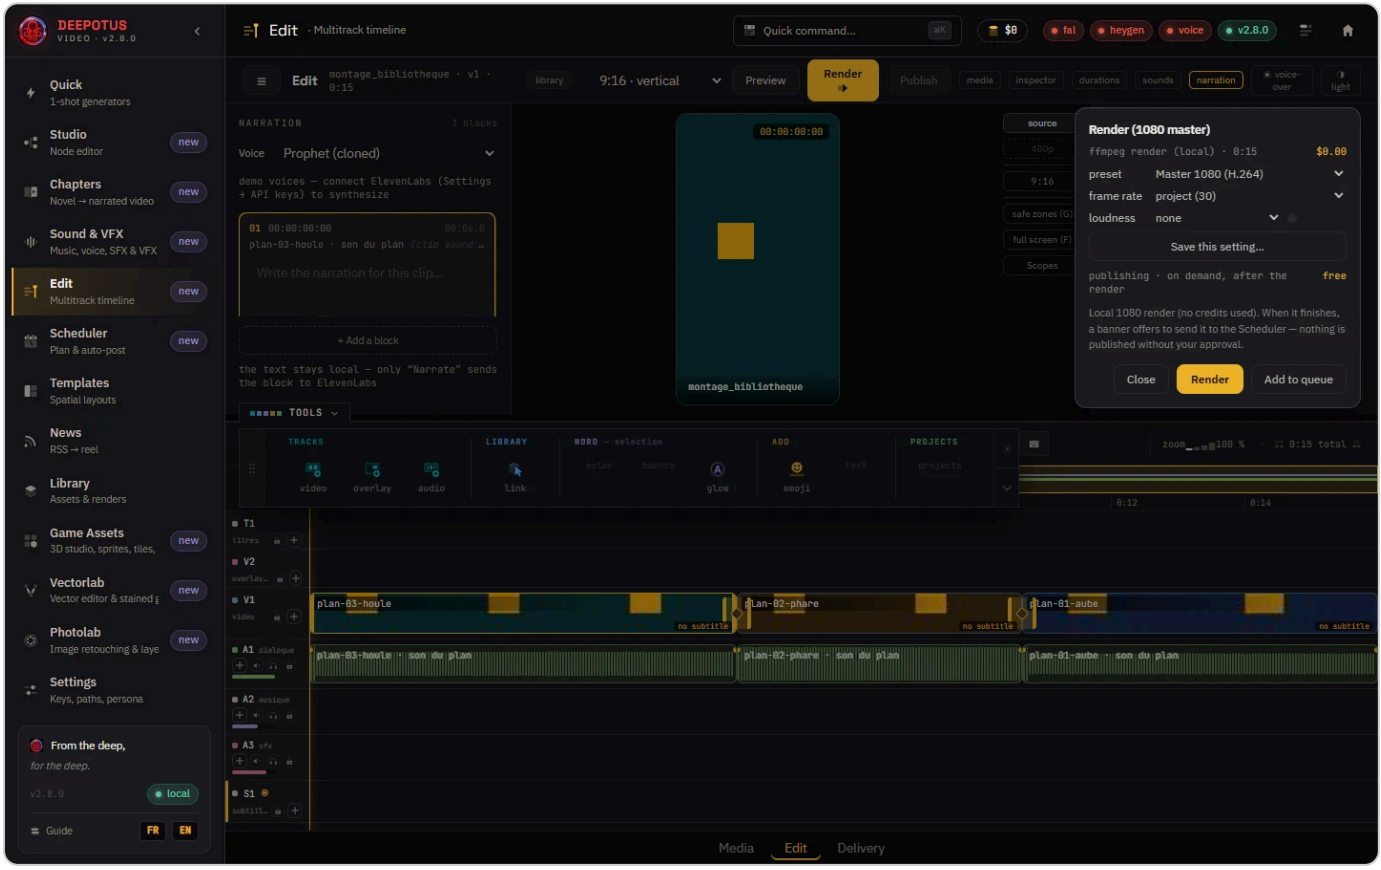

1. Pick the preset, then render or queue

Render… opens the render panel: output preset, frame rate, loudness, range, render queue.

1. Click **Render** in the top bar: the render panel opens.
2. Pick the **preset**: master 1080, web 4K, social 720, ProRes 422, H.265, WebM, audio only (AAC, MP3, WAV) or GIF 480. The project's format is kept.
3. Set the **frame rate** (24, 25, 30 or 60 fps) and the **loudness**: −14 LUFS for YouTube and TikTok, −16 for a podcast, −23 for broadcast. A badge compares the last measurement with the target.
4. **Render** starts the render; **Add to queue** puts it behind the others. To export only the I/O range, tick the box provided.
5. When it is done, a banner offers to send the video to the Scheduler, with your channels, a date and a prefilled caption. Nothing is published without you.

The Delivery view

The bottom bar › **Delivery** lists the project's renders and previews, their status (queued, running, done, failed), and leads to the Library. **Publish** sends the last final render back to the Scheduler.

Go further

Your own presets

Save this setting… keeps your output setting as a preset of your own, listed with the others.

To another program

☰ › Project › **Export EDL…** or **Export FCPXML…** to take the timeline into DaVinci Resolve or Final Cut.

True to the timeline

The render lasts exactly the timeline, frame for frame; a fade is centred on the cut with the source's real frames.

Good to know

Run a 480p **Preview** first: it is free and shows transitions, effects and mix.

MAKE VIDEOS

Record the app and go live

You will be able to record the app with your voice, your camera and your clicks, dress it up with a skin, turn it into an edited video, go live on Twitch, YouTube or Kick, and share a tutorial that others replay. Everything is free except voice transcription, on a quote; nothing leaves the PC without your action.

Record the app

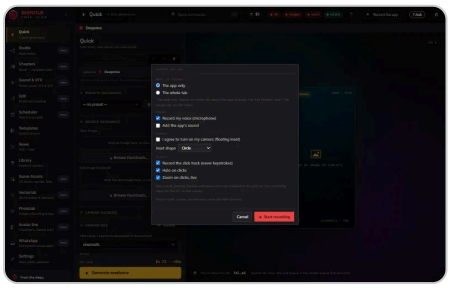

1. What is filmed, sound, camera, clicks

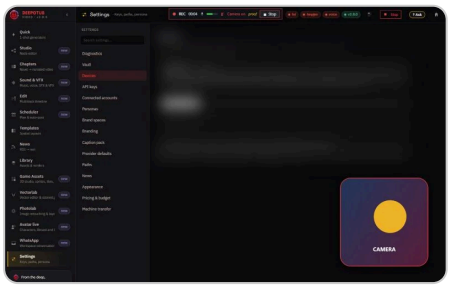

2. REC badge; secret screens are hidden

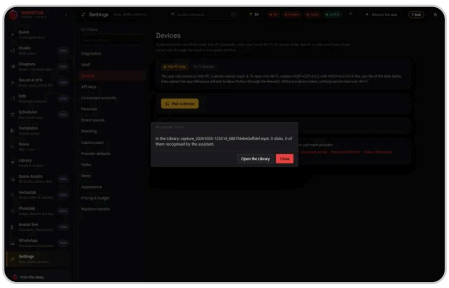

3. Stop: the video is in the Library

Top bar › Record the app: frame, voice, camera with consent, clicks; secrets are hidden; Stop saves the video.

1. In the top bar, click **Record the app**. Choose what is filmed: **The app only**, without the boxes laid over it such as the Ask drawer, or **The whole tab**.
2. Sound: **Record my voice (microphone)** and, if needed, the app's sound. The camera only turns on with the consent box; its floating inset can be moved and resized.
3. Clicks: the click track (never the keys you type), a halo and, if you like, a live zoom. Then **Start recording** and share this tab when the browser asks; another tab is refused.
4. While recording, the REC badge shows the microphone and the camera. Keys, vault, pairing, devices and passwords are hidden in the picture.
5. **Stop** saves the video and its raw tracks in the Library, on this PC.

Dress the video with a skin

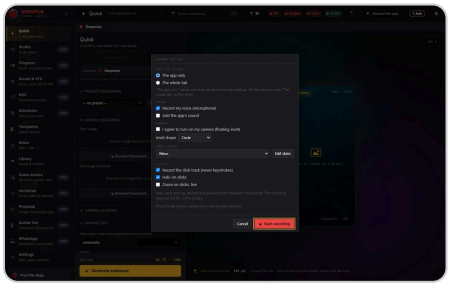

1. Look: one skin per recording

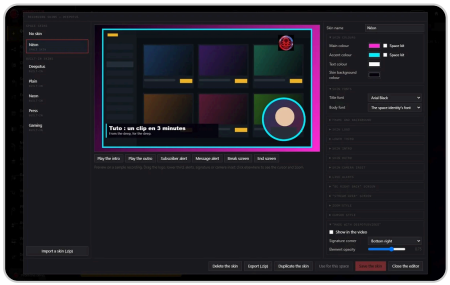

2. The editor: live preview, elements to drag

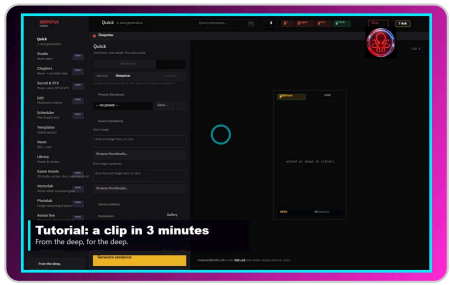

3. In the video: lower third, logo, cursor

Look (skin): pick it in the box, tune it in the editor, see the lower third, logo and cursor in the video.

In the recording or live box, the Look line picks the skin; **Edit skins** opens the editor. Five skins are built in: Deepotus, Plain, Neon, Press and Gaming. Each space has its own, Deepotus by default.

1. Pick a skin: the preview plays on a sample capture. Drag the logo, the lower third, the alerts, the signature or the camera inset to place them.
2. Set colours and fonts (the space kit's by default), frame and background, lower third, intro and outro, live alerts, break and end screens, zoom and cursor. **Play the intro** or **Break screen** show the element.
3. A built-in skin is edited as a copy: **Save a copy**, then **Use for this space**.

Edit the recording

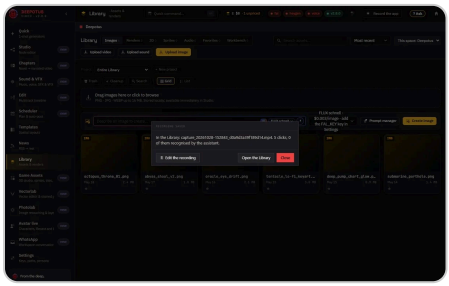

1. At the end: Edit the recording

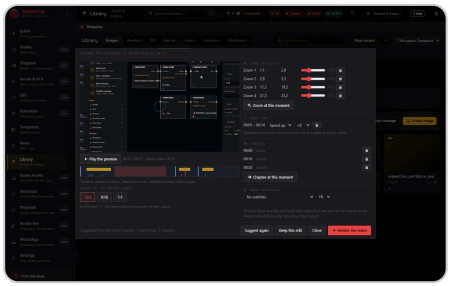

2. Zooms, dead time, chapters, format

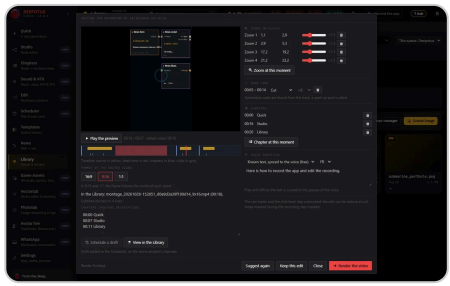

3. A draft waits in the Scheduler

Edit the recording: zooms, dead time and chapters suggested from the click track, then a draft in the Scheduler.

1. At the end of a recording, click **Edit the recording**; later, the Record the app box offers **Edit a recording…**.
2. The edit is suggested from the track: zooms on clicks, generation dead time to **Speed up** or **Cut**, one chapter per screen change. Adjust them, or add **Zoom at this moment** and **Chapter at this moment**.

3. Format: 16:9, 9:16 (the frame follows the zooms) or 1:1. Subtitles: a known text synced to the voice, free, or a transcription, paid after a quote.

4. **Render the video**: it goes into the Library, its chapters ready for a YouTube description. **Schedule a draft** puts it in the Scheduler, on the active project's channels; nothing is published.

Go live

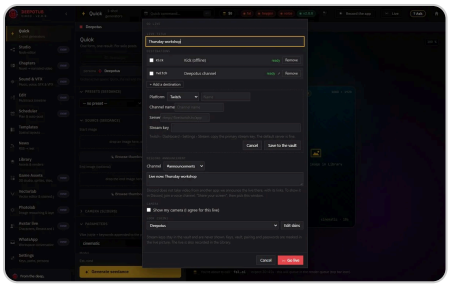

1. Destinations, Discord announcement, camera, skin

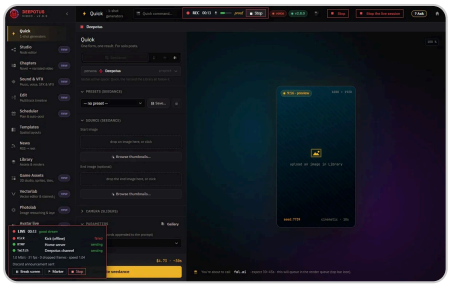

2. Live: health per destination, Break screen, Marker

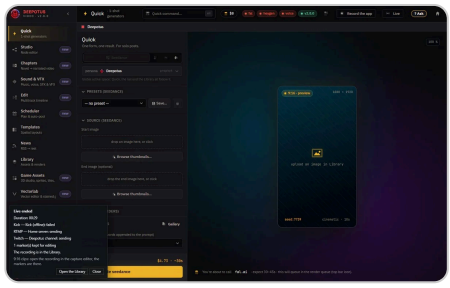

3. The summary; the recording is in the Library

Live: destinations whose key is in the vault, a Discord announcement, then stream health and the summary.

1. In the top bar, click **Live** and give the live a title.

2. **+ Add a destination**: Twitch, YouTube Live, Kick or an RTMP server, which also covers Facebook, LinkedIn or Restream. Paste the stream key, then **Save to the vault**: it is never shown again.

3. Discord does not take another app's video: added through its webhook, it announces the live and its links in the chosen channel.

4. Tick the ready destinations, the camera if you wish, the skin, then **Go live**.

5. The panel tracks the stream's health per destination. **Break screen** shows "be right back", **Marker** notes a moment for editing, **Stop** ends it. The live is also recorded in the Library.

Share a replayable tutorial

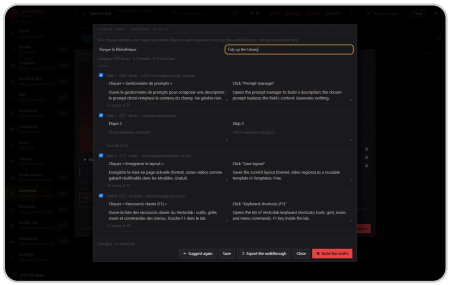

1. One step per gesture, titles and sentences FR/EN

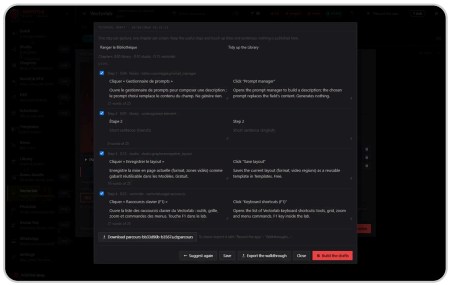

2. The .dzparcours walkthrough, ready to share

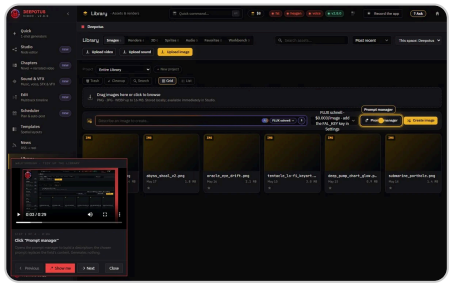

3. Elsewhere, Show me points at the real button

Tutorial…: one step per gesture in FR and EN, Export the walkthrough, then Replay elsewhere: Show me points at the real button.

1. In a recording's edit editor, **Tutorial…** suggests one step per gesture and one chapter per screen, with a title and a short sentence in French and English. Keep the useful steps and touch up the texts.
2. **Build the drafts** prepares, without publishing anything, the step pictures for the guide and the lab cards.
3. **Export the walkthrough** produces a .dzparcours file: the video, the click track and the steps, with no typed text nor personal path.
4. On another installation, Record the app › **Walkthroughs…** › **Import a walkthrough…**, then **Replay**. **Show me** points at the real button without clicking for you; an element missing from that installation is flagged.

Good to know

> The Deepotus skin signs the video with a small "Made with DeepotusVideo": untick it in the editor if you prefer. A skin is shared as a .zip: **Export (.zip)** and **Import a skin (.zip)**.

> A running recording prevents going live: stop it first. In The app only mode, boxes laid over the app are not filmed; choose The whole tab to show them.

PUBLISH

Plan and publish your week

You will fill your week with posts, review them, approve them in one go and let them go out on time on Telegram, X, YouTube, Instagram and TikTok. The Scheduler is free; only what you produce (image, clip, avatar) has a price, shown before the run.

Before you start

For a post to go out on its own, its network must be connected: see Connect your accounts. Without a connected account everything still works in assisted mode: the post is ready at the set time and you click **Send**.

The week at a glance

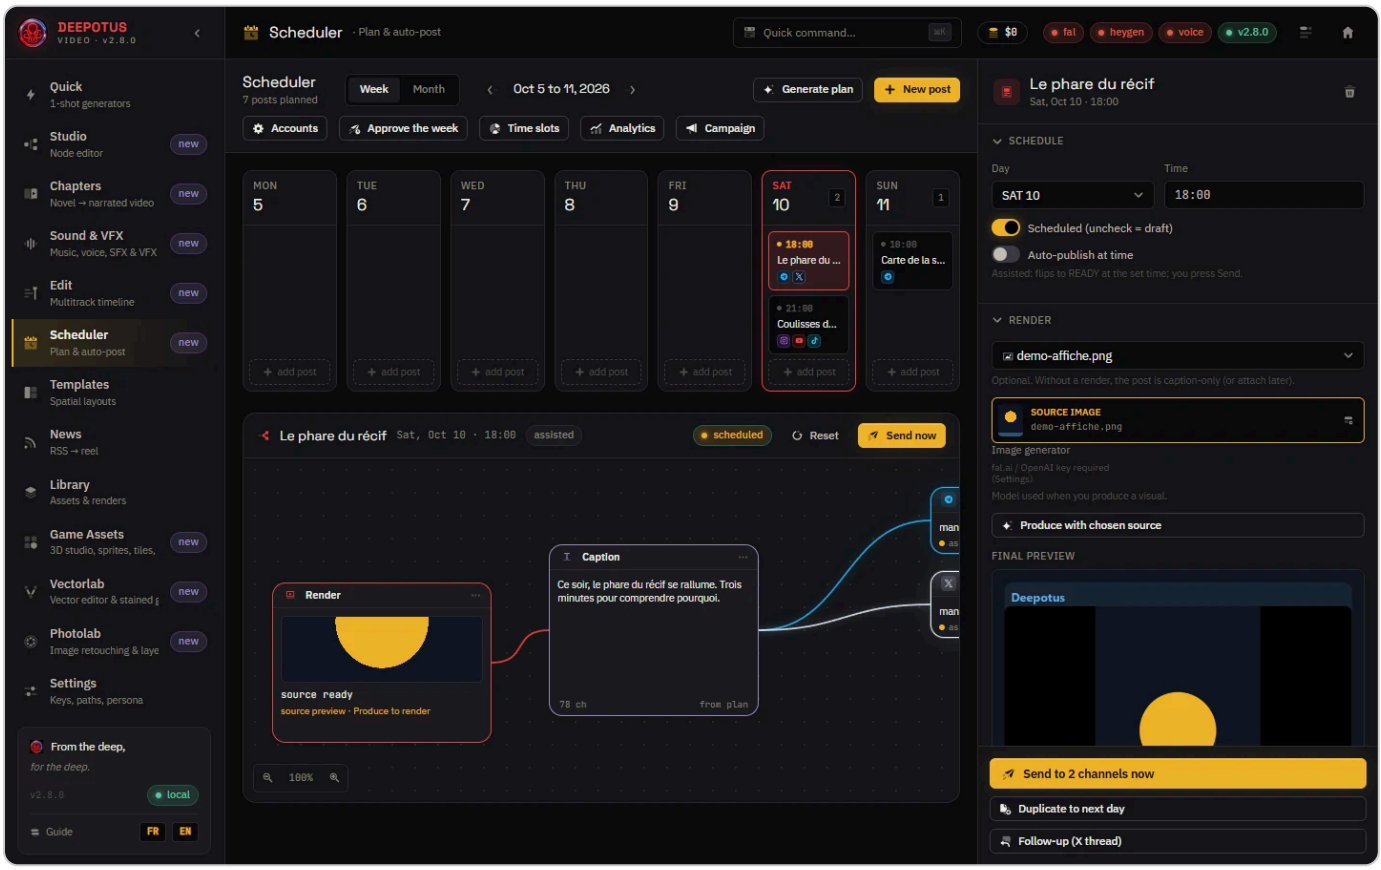

1. Fill the week with a plan 2. The graph: render, caption, channels 3. Nothing goes out without your click

The calendar at the top, the chosen post's graph in the middle, its inspector on the right.

Open **Scheduler** in the rail. At the top, the calendar in **Week** or **Month** view. Click a post: its **graph** (render, caption, one node per channel) shows in the middle and its **inspector** on the right. The screen fits the window from 1280 to 1600 px; the toolbar wraps when needed.

Fill the week

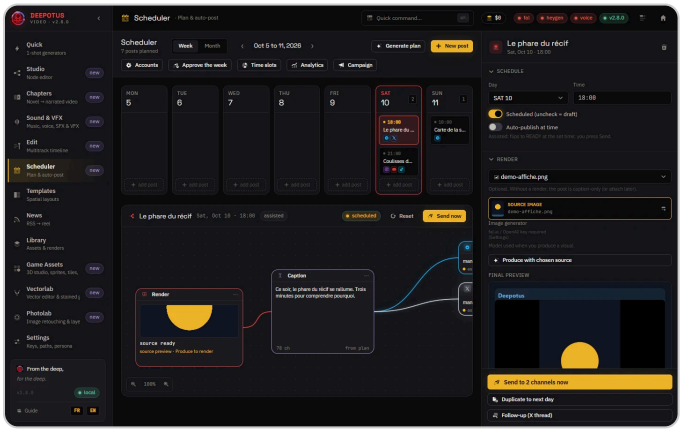

1. Generate plan

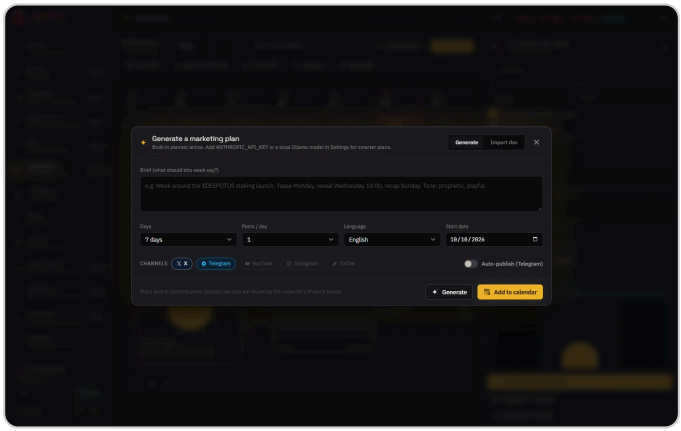

2. Describe the week in one sentence

A brief, a start date, the number of days and posts: you review the plan before it enters the calendar.

1. **Generate plan**: describe the week in one sentence, pick the number of days, posts per day, language, start date and channels, then **Generate**. Without an AI key the built-in planner does the job; with an Anthropic key or a local Ollama model the plan is finer.
2. Review the proposal, then **Add to calendar**. Times follow your per-channel time slots.
3. Already have a written schedule? The **Import doc** tab turns a .md, .txt, .docx or .pdf file into dated posts, keeping weeks and themes.
4. For a single post: **New post**, or **add post** at the bottom of a day. Episodes from Chapters land here as drafts.

Prepare a post

In the inspector, set the day and time, tick **Scheduled** (unticked, the post stays a draft), pick the channels and write the caption, the same on every channel, with its character count. **Produce with chosen source** makes the visual or clip the plan calls for; the render runs in the background and its thumbnail appears on the graph's Render node.

Coulisses du montage

Sat, Oct 10 · 21:00

FINAL PREVIEW

Reels · zones indicatives · 59 car.

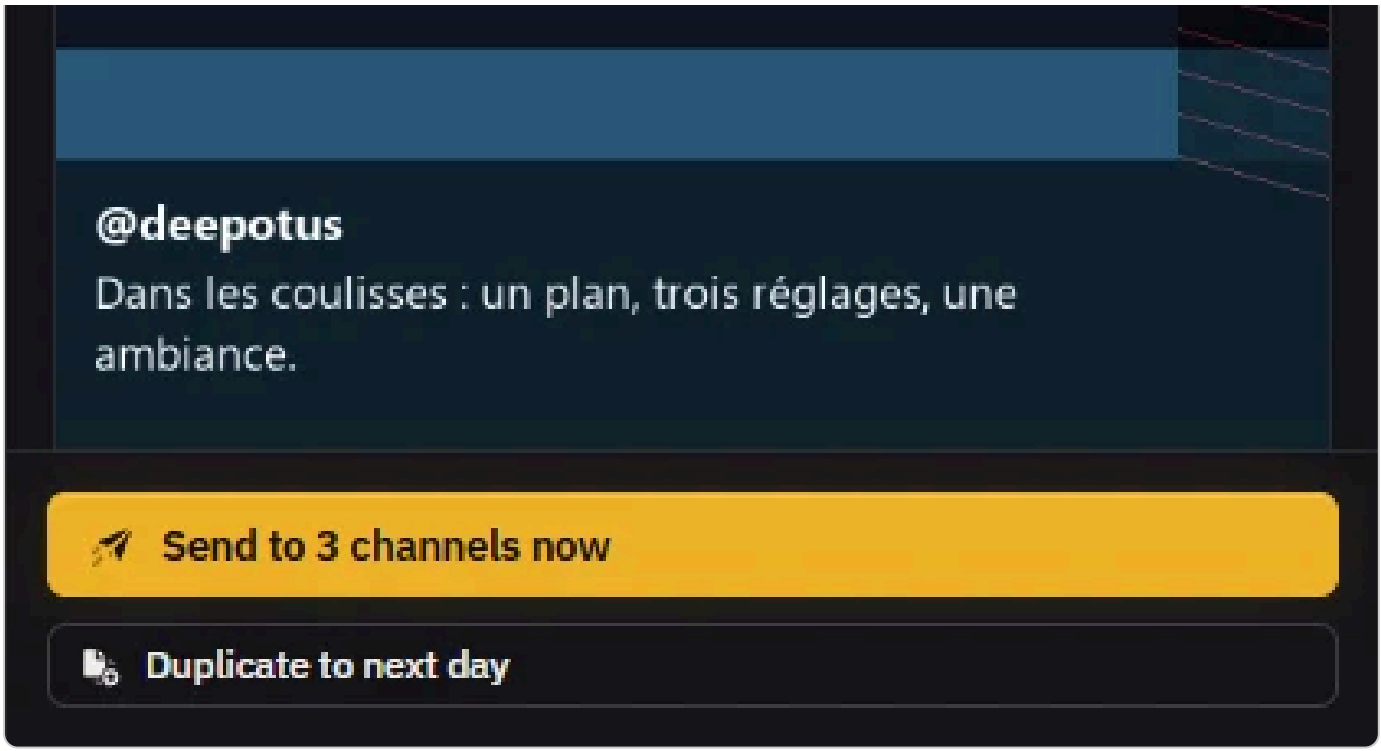

1. Put nothing important under the dark areas

For a vertical post, the preview darkens the areas the network's interface covers.

The **final preview** shows the post as it will look on the network. For a Reel, a Short or a TikTok, it darkens the top and bottom bands and the right rail that the network's interface covers: put nothing important there.

Approve and publish

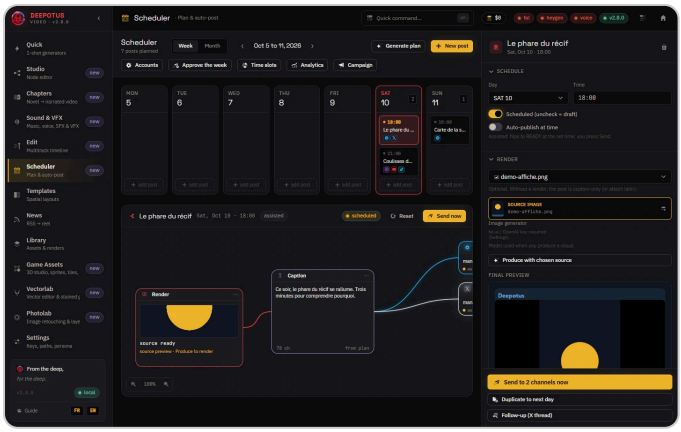

1. Approve the week

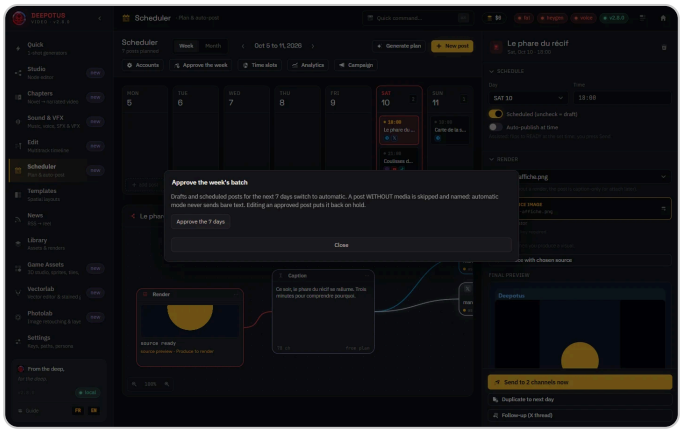

2. One click approves the whole batch

Batch approval: the next 7 days switch to automatic, a post without media is skipped.

1. Each post has a mode. **Assisted** (default): at the set time it flips to ready and you click `Send`. **Auto** (`Auto-publish at time`): it goes out on its own on its connected channels.

2. **Approve the week** then **Approve the 7 days** switches the drafts and scheduled posts of the next 7 days to automatic in one go. A post without media is skipped and named; really editing an approved post puts it back on hold.
3. To publish right now: **Send now** on the graph, or the yellow inspector button that counts the channels.

Before each send, the network's quota is checked (X 500 posts a month, Instagram 100 per rolling 24 h, YouTube 100 a day, TikTok 15 a day); a channel already published is never published again.

Going further

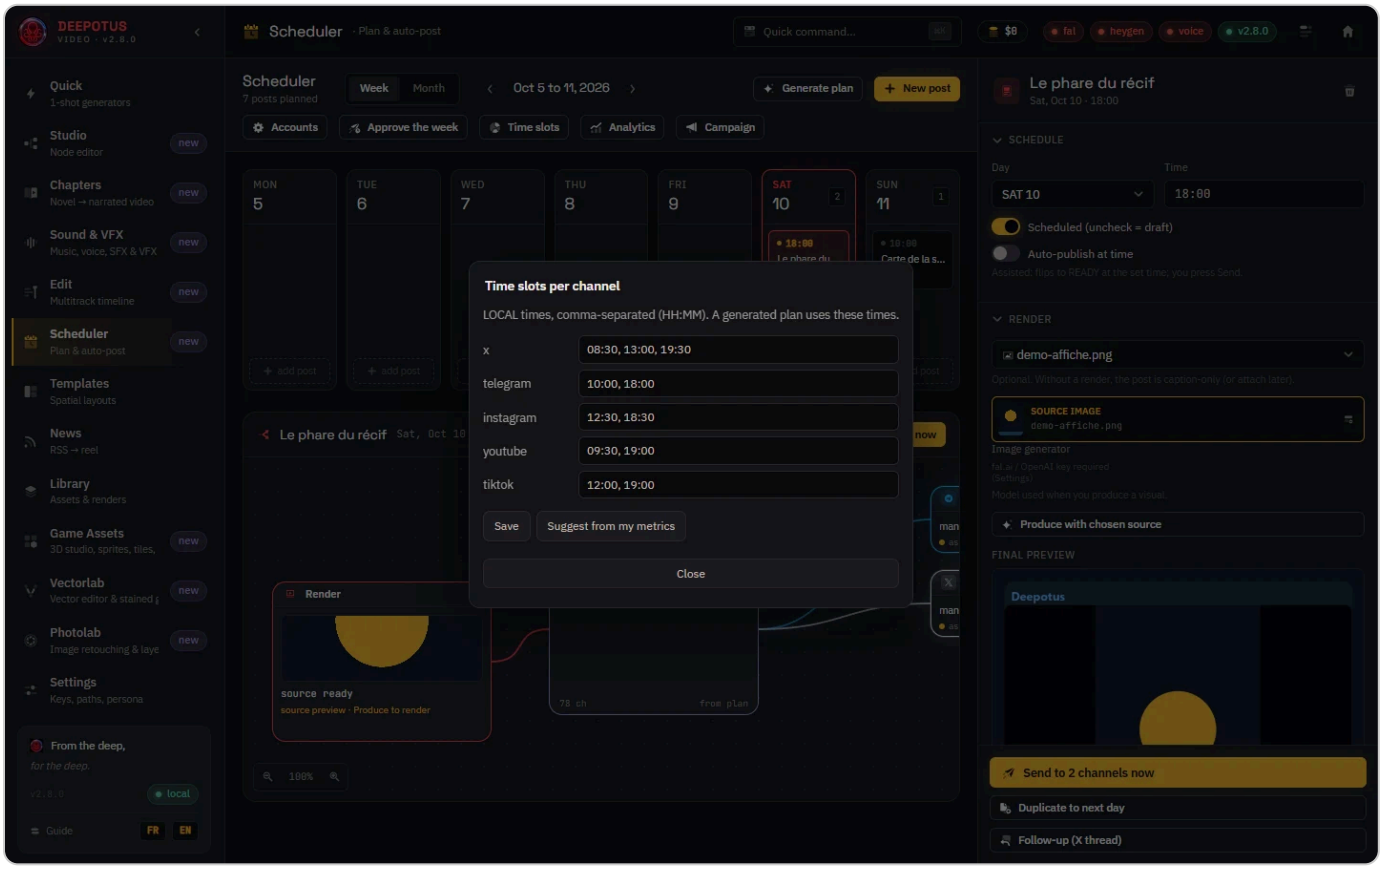

1. Times separated by commas 2. Or let your metrics suggest

Local times per channel, applied by generated plans.

Time slots

Time slots: local times per channel, separated by commas, then **Save**. **Suggest from my metrics** picks the hour with the best engagement, from 5 measured posts.

Analytics

Views, likes and engagement over the last 28 days by channel, format and week, with the top 5. Metrics are re-read every day; **Refresh now** re-reads 10 posts.

Campaign brief

Campaign › **Campaign brief**: objective, key messages, fixed sections and banned terms. The next plan follows it; banned terms are removed from captions and hashtags.

Recurring series

A rule (days, time, channels, caption template with {date}, {weekday}, {week}): **Create series**, then **Create 4 weeks**. These are drafts; creating twice duplicates nothing.

Suggested recycling

Your best-measured posts, published more than three weeks ago, come back with a varied caption. **Suggest for tomorrow** creates a draft; no paid AI is called.

X thread

Follow-up (X thread) writes the follow-up to an X post: it goes out two minutes later and replies to the previous message, never before it.

Duplicate, sign

Duplicate to next day copies the post. Under the caption, a row of emojis and the **Telegram Premium pack** insert your brand tags, set in Settings.

With the phone

Approved posts form a batch the paired phone can take along. A post handed to the phone says so in the inspector, with "Take back on the PC".

Good to know

After an action (approve, create a series, suggest a recycle, add a follow-up), the list reloads by itself: no need to leave the screen.

The calendar, built-in plans, approval and sending are free. Producing an image, a clip or an avatar is paid, under your spending caps (see Control your spending).

PUBLISH

Run a campaign with the agent

You will be able to hand a campaign to the Campaigns agent: it turns a brief into a strategy and a schedule, writes the copy, then puts the posts as drafts in the Scheduler. You approve at two gates and nothing goes out without you. Ten minutes; `paid ≈ $`, each message shows its cap before you send it.

Before you start

The Campaigns agent works with Claude only: you need an **Anthropic key** in the vault (see Connect your keys and AI engines). Without a key, it says so on the first send; nothing is greyed out. Its spending is counted under the "campagnes" category; the global cap, the per-engine caps and the space cap apply: see Keep spending under control.

A campaign belongs to **one brand space** (see Work for several brands): its persona, its forbidden words, its renders. It is not necessarily the space active on screen.

The circuit in short

The brief becomes a folder, then a strategy. **Gate 1**: you approve the strategy and the schedule. Then come the copy, the images or videos on a quote, and the drafts in the Scheduler. **Gate 2**: you review and schedule each draft yourself. The agent's circuit stops at the drafts.

Open the agent

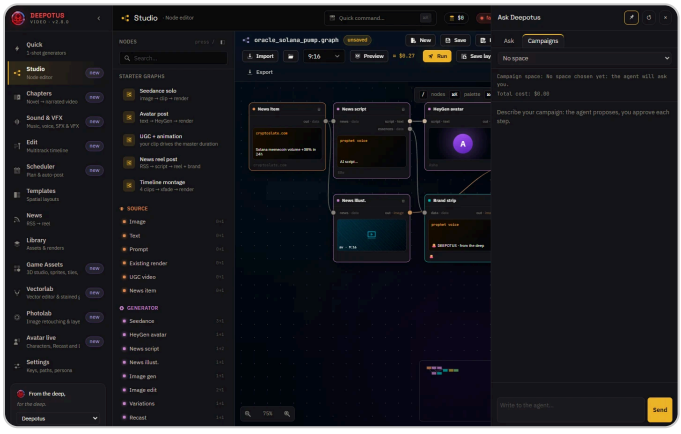

1. Campaigns tab

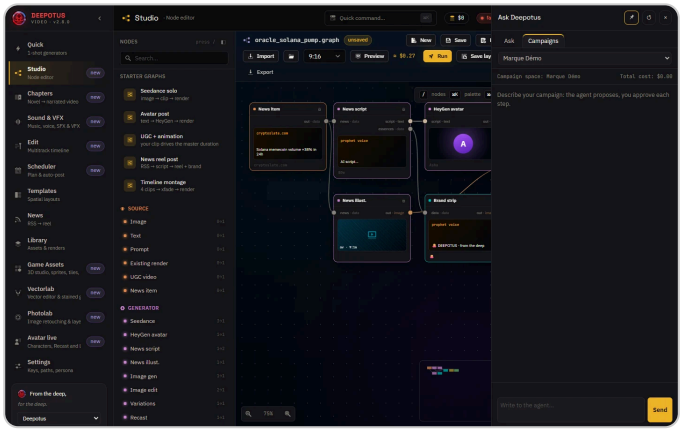

2. The campaign space 3. Shown in full, with the running cost

The Campaigns tab of the F1 drawer; the space chosen at the top is shown in full and no longer changes.

1. Press `F1`: the drawer opens. It has two tabs, `Ask` (the quick help, see Ask Deepotus) and `Campaigns`. Click `Campaigns`.

2. Pick the campaign space in the list at the top. Left on **No space**, the agent only gives directions, and asks you for the space with a card if a step needs a space.

3. The header shows the space in full and the conversation's **total cost**. Once the campaign has started, its space no longer changes.

Send the brief: a quote for each message

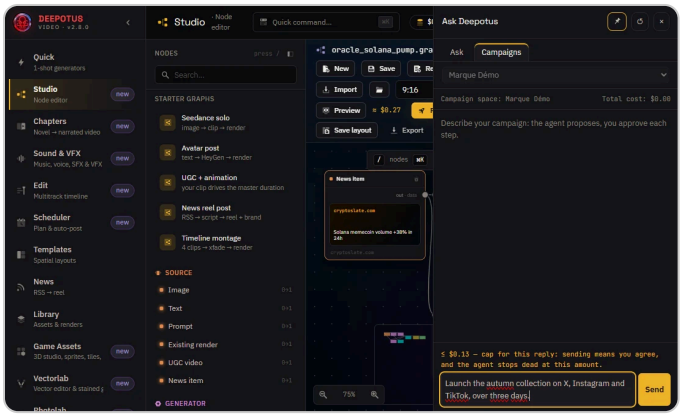

1. Cap for this reply

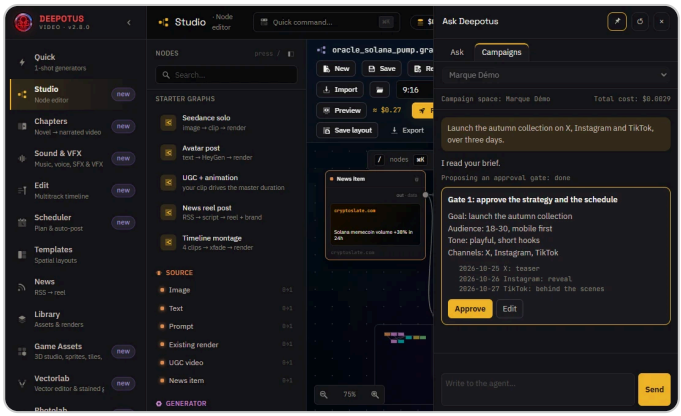

3. Gate 1: strategy and schedule

Under the input, the cap for the reply; sending means you agree. The agent comes back with gate 1.

1. Describe the campaign in a few lines: the product, the goal, the channels, the period.

2. Under the input, the **cap for this reply** appears, for example "≤ $0.13". It is the most the agent can spend to answer, all its turns included.

3. Click **Send**: sending means you agree. If you click before the cap shows, the click only displays it; click again to accept it.

4. The agent reads the space, looks for what already exists, writes the folder, then comes back with the "Gate 1" card: the strategy in a few lines and the day-by-day schedule.

Pass gate 1, get the drafts

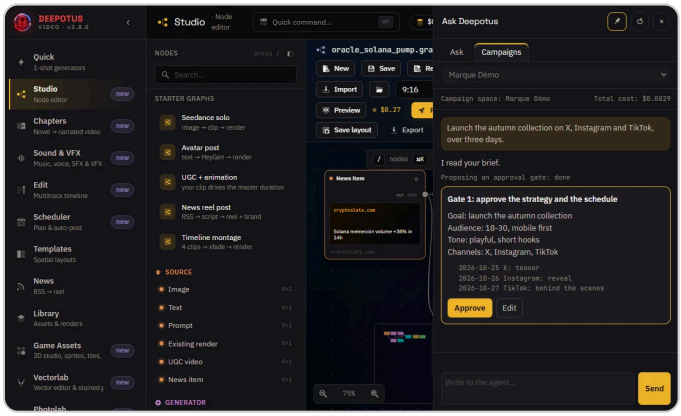

1. Approve, or Edit

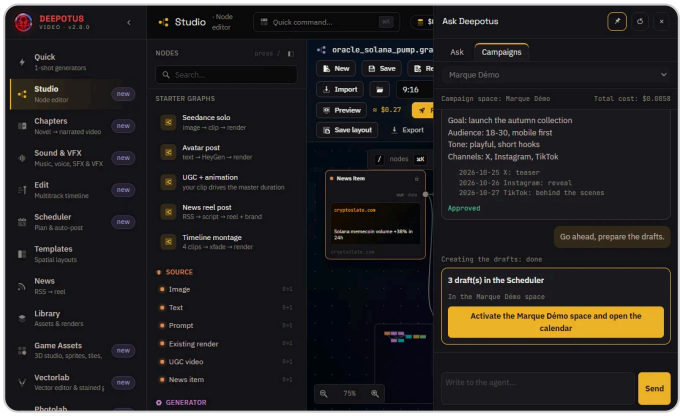

2. Drafts in the campaign space

After Approve, one more message (with its cap): the posts land as drafts, in the campaign space.

1. Read the strategy. **Approve** accepts it; **Edit** declines it and lets you write your correction in the input.
2. Once approved, send a message ("go ahead"): it has its cap, like every message. The agent writes the copy.
3. For an image, a video or a voice, the agent posts a quote card with its price: nothing is generated before you click **Generate**. **Decline** drops it.
4. The posts land as drafts in the Scheduler, in the campaign space. The **Open the Scheduler calendar** card takes you there; if that space is not the active one, it first offers to activate it, on your click.

From the Scheduler

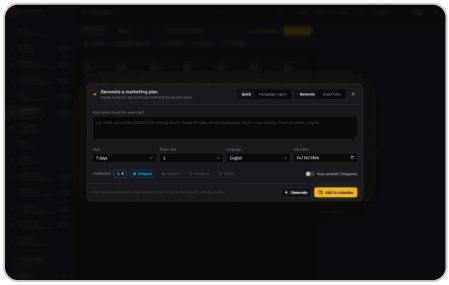

1. Quick or Campaigns agent

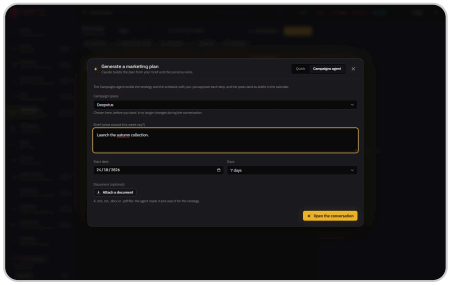

2. Open the conversation

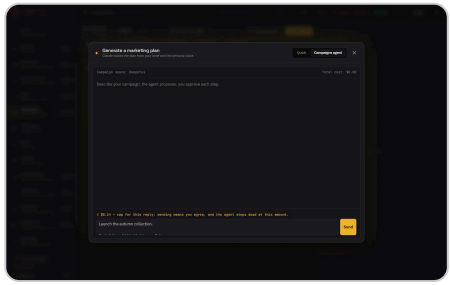

3. The brief waits for you to send it

The plan dialog keeps Quick mode and gains Campaigns agent mode: space, brief, period, document.

1. In the **Scheduler**, click **Generate plan**. At the top of the dialog, switch from **Quick** (the one-shot plan, see Plan and publish your week) to **Campaigns agent**.
2. Pick the space, write the brief, the start date and the number of days. **Attach a document** adds a .md, .txt, .docx or .pdf file for the agent to read.

3. **Open the conversation**: the brief is placed in the input, with its cap. You send it; gate 1 happens in the dialog, and the drafts show up in the calendar.
4. **Gate 2** is the Scheduler itself: review each draft, fix it, then schedule it or validate the week yourself.

Go further

Continue after a stop

When the agent reaches the cap or its number of turns, it stops dead and posts a "Continue" card with a new cap. Nothing restarts without your click.

After a render

When a video or an image is done, the card switches to **Render finished**. The agent never resumes on its own: **Continue the campaign** puts the follow-up in the input with its cap; nothing goes out before **Send**.

From "Ask"

Ask for a campaign in the **Ask** tab: the answer carries **Open the Campaigns agent**, which takes your question to the next tab.

Resume later

During the session, the drawer keeps the current conversation. In the plan dialog, **Resume the conversation** reopens the latest campaign of the chosen space, during the session too.

From the phone

The paired phone runs the same conversations over Wi-Fi. The work happens on the PC, with its keys and its cap.

The campaign folder

Folder, strategy, copy and review are stored in the app's data, under the campaign space, with what each campaign has learned.

Good to know

Each message shows its cap before you send it; the agent stops dead if it reaches it. These turns are counted under the "campagnes" category (statement in **Settings** › **Pricing & budget**); the global cap, the per-engine caps and the space cap apply. Images, videos and voices have their own quote, paid only when you click and also counted under "campagnes".

The agent never publishes or schedules a post: everything lands as a draft. It does not change the active space, touch the keys, the Settings or the accounts, and deletes nothing. Ask it to: it shows you where to do it yourself.

A campaign counts its spending and stores its renders in **its own** space, even if another space is active on screen.

PUBLISH

Watch, prospects and prospecting

You will learn to follow what your competitors' customers say (**Watch** tab), keep your list of prospects (**Prospects** tab) and let the Campaigns agent find them, write to them and follow their replies. Nothing goes out without your approval, and the app enforces the law. Fifteen minutes; the watch and the records are `free`, web search and WhatsApp template messages are `paid`.

Before you start

Everything happens in the **F1** drawer, **Watch** and **Prospects** tabs, for the workspace chosen in their list (see Work for several brands). The agent that searches and writes is the Campaigns agent (see Run a campaign with the agent): it needs an Anthropic key. Emails go out from your mailbox or Brevo, WhatsApp messages from the workspace's number (WhatsApp screen).

Watching competitors

1. Open the **Watch** tab. It shows the workspace's latest watch: the competitors × themes matrix, the most cited pains, message angles and requests to decide.
2. **Start a collection**, or ask the agent "run the watch": Hacker News, YouTube comments (yt-dlp must be installed on the PC) and App Store reviews, read-only, for free, 15 minutes at most.
3. Quotes are evidence to read, never instructions. A public post that states a need never opens a record: the agent suggests a useful public reply, which you post yourself, once.

The Prospects records

1. **Prospects** tab › **Add a prospect**: the source (type, date, address for a web page) and the legal basis are required. A consumer webmail address cannot use B2B legitimate interest.
2. Statuses follow the pipeline: Suggested, Qualified, Contacted, Replied, Meeting, Client (signed), Lost, Objected. Only allowed moves are offered. **Board** shows every record in columns.
3. At **Meeting**, the app creates a "demo" workspace: it never publishes and is used for the showcase campaign (see below).
4. People's rights: **Export their data**, **Delete the record**, **Objects to prospecting**. An opt-out applies to all your workspaces; messages prepared for that person are emptied at once.

Let the agent find and qualify

1. Ask "find me prospects": the agent suggests records from inbound WhatsApp contacts who agreed, a CSV file you drop (waiting list, form: dated consent), or a company page you paste. The French company register (free) gives targets, without records: their website is needed next.
2. A web search for companies is paid: the agent shows a quote card, nothing runs before your click. The provider and its keys are set in **Prospects** › **Web search**.
3. Records suggested by the agent say "suggested by the agent". The score starts from a fixed scale, which the agent may adjust by ±20 with a reason. A record found on the web shows "To inform before …" (one month, GDPR art. 14).

Writing: sequences and gates

1. Once the record is **Qualified** (or Contacted), ask the agent for a sequence. It drafts it: a first message and up to two follow-ups, on a channel the record allows, in the workspace's voice and without its forbidden words.
2. This is gate 2: **Approve the sequence**, in the agent's card or in the record's **Sequences** block. Then **Schedule**: each message lands within the workspace's sending hours (the WhatsApp agent's hours), emails under the daily cap.
3. **X, Instagram, LinkedIn**: the app never sends private messages. At the planned time, **Copy and open** copies the text and opens the profile; you send it, then **Mark as sent**.
4. **WhatsApp**: only contacts who agreed. Follow-ups are Meta-approved templates; free text only goes if the contact wrote within 24 hours. A STOP applies to every number of the workspace and, outside client workspaces, to all prospecting.

The email channel

1. **Prospects** › **Email**: the sender, company name and postal address (required), sending from your mailbox (SMTP) or through Brevo, and reading replies (IMAP). The password is sealed for this Windows session and never shown.
2. **Send a test** writes to your own address, never to a prospect.
3. Every email carries your identity and "reply STOP". Replies are read from the mailbox: a reply moves the record to "Replied", a STOP opts it out. No email goes out while reading fails.

4. Cap: 10 emails a day the first week, 20 the second, 30 after that (can be set lower). Too many bounces or a spam complaint pause sending; **Resume sending** restarts it.

Following replies

1. Ask "who replied?": the agent reads the replies (the email text; for WhatsApp, only the date).
2. For an email reply, it suggests ONE short reply, on that channel, which you approve like a sequence; or it hands the record back to you.
3. Price, quote, deadline or contract, complaint, refusal, STOP, or a WhatsApp reply to read: the record goes to **To take back** and the agent stops writing to it. The badge goes away when you note a contact or change the status.

The showcase campaign

For a prospect at the meeting stage, the agent prepares a showcase campaign in their demo workspace: persona, drafts, visuals, to show what you would do. It never publishes. On the paired phone, the agent screen offers “Showcase campaign” for each demo workspace; on the PC, open a Campaigns conversation in that workspace.

Good to know

The watch, the records, the company register and emails are free; the agent's turns, web search and WhatsApp templates are paid, counted under the "campagnes" category (WhatsApp under its own), within your caps.

The agent never approves a sequence for you, never promises a price or a deadline, and never writes to someone who opted out. A public handle or comment is never reused in a message.

A sequence stops on its own when the prospect replies, books a meeting, signs, is lost, opts out or withdraws consent.

PUBLISH

Connect your accounts

You will plug in Telegram, X, YouTube, Instagram and TikTok so the Scheduler can publish on its own. Free: all it takes is the access each network provides, stored on your PC.

Where it is set

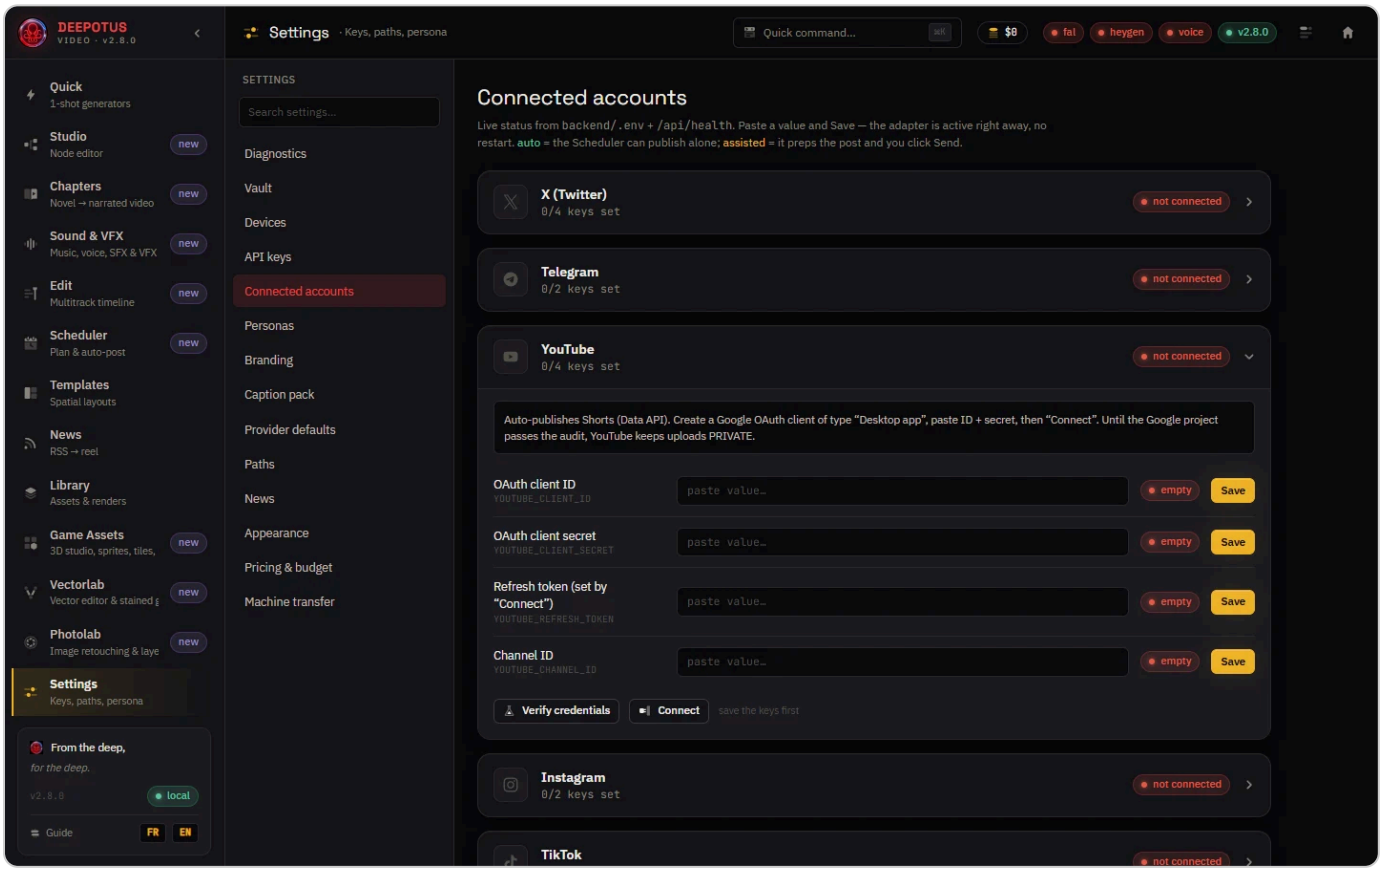

1. Paste the ID and secret, then Save 2. Connect opens the consent page

Settings, Connected accounts: one block per network, its keys, its live status.

Open `Settings` › `Connected accounts`. One block per network shows how many keys are set and its live status: **auto** (the Scheduler publishes alone), **assisted** (it preps the post and you click Send) or **not connected**.

1. Expand the network's block and read its note: it says where to get the access.
2. Paste each value and click `Save` next to it. The adapter is active right away, no restart.
3. Check with `Verify credentials` (or `Send test message` for Telegram).

Network by network

Telegram

The simplest, in five minutes. Create a bot with @BotFather, make it an administrator of the channel, paste the bot token and the channel chat ID.

X

The four keys from developer.x.com (API key and secret, access token and secret). The free tier allows posting, up to 500 posts a month.

YouTube

Create a Google OAuth client of type "Desktop app", paste its ID and secret, then **Connect**. Until the Google project passes the audit, YouTube keeps Shorts private, and the Scheduler says so.

Instagram

A business account and a Page token (EAA…) or Instagram token (IG…). Reels are uploaded from your PC, no hosting, up to 100 posts per rolling 24 h.

TikTok

A TikTok for Developers app (Login Kit, "Desktop" platform): paste the client key and secret, then **Connect**. Without an audit, uploads are private and the account must be too; set TIKTOK_AUDITED to true after the audit.

Connect YouTube and TikTok without copying a code

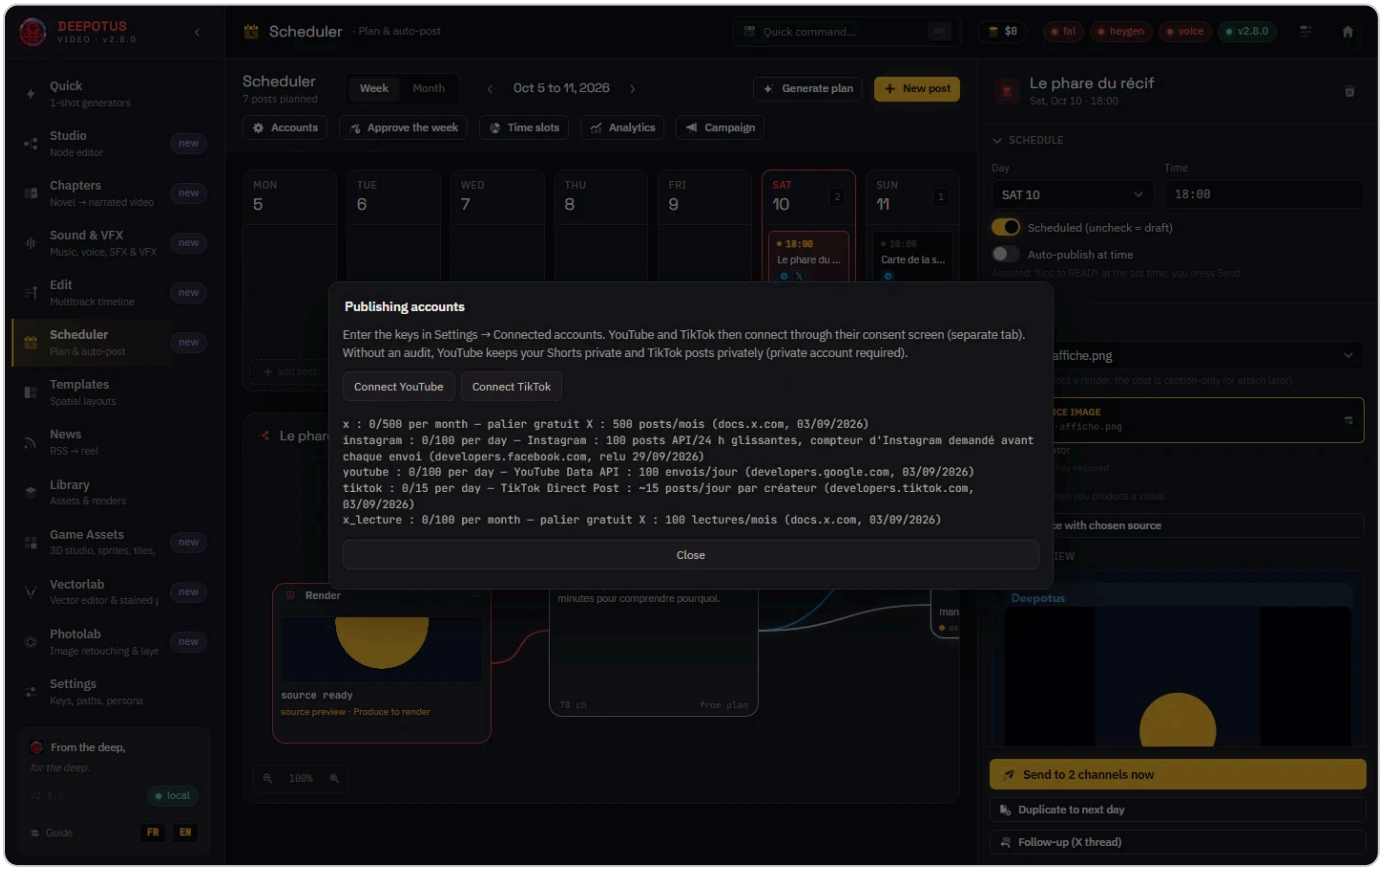

1. Connect through the consent page 2. Each network's quota, checked before sending

In the Scheduler, Accounts shows the month's quotas and connects YouTube and TikTok.

Connect opens the network's consent page in a new tab: accept, come back, the status turns to connected and the refresh token is saved on its own. The same gesture exists in the Scheduler: the **Accounts** button, then **Connect YouTube** or **Connect TikTok**. That panel also shows each network's quota for the month.

Good to know

> A network that is not connected blocks nothing: its posts switch to assisted mode and you send them in one click, or from the phone.

PUBLISH

Work for clients

You will be able to create a client's space with verified channels, point each project at its accounts, and publish for the client without picking the wrong account or skipping their consent. Free: verifying a channel publishes nothing and costs nothing. A space's persona, identity and export are in Work for several brands; each network's keys in Connect your accounts.

Create a client's space

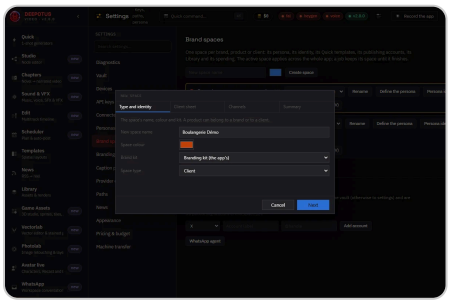

1. Space type: Client adds a sheet

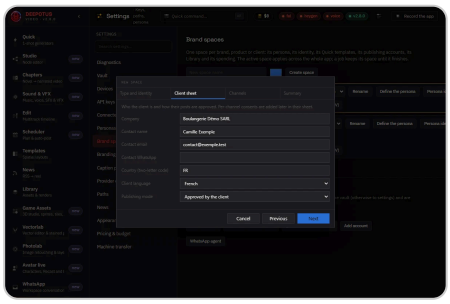

2. The publishing mode sets the safeguards

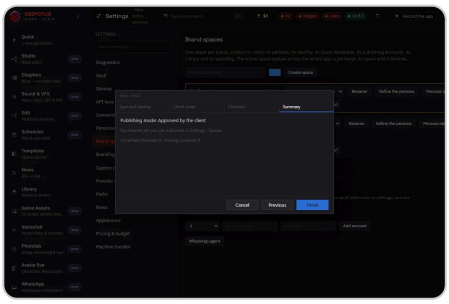

3. The summary says what is missing; Finish activates the space

Settings › Brand spaces › Create space: Client type, the client sheet, then the summary.

1. Open `Settings` › `Brand spaces`, type the name, then `Create space`: the **New space** wizard opens, in four steps.
2. **Type and identity**: name, colour, brand kit and space type, **Brand**, **Product** or **Client**. A product can belong to a brand or a client: it takes over its kit, its persona and, for a client, its rules.
3. **Client sheet** (clients only): company, contact (name, email, WhatsApp), country, language and **Publishing mode**, approved by the client or published on their behalf.
4. **Channels**: tick the networks and enter their credentials; secrets go to the vault and are never shown again. `Verify` queries the network without publishing anything: verified, undetermined (with the reason) or refused. An unverified channel is still saved.
5. **Summary**: the publishing mode, unverified channels and missing consents. `Finish` activates the space.

Point a project at its channels

1. In the `Library`, pick the project in the **Project** bar, then click `Channels`.
2. Tick the space's accounts that will receive its posts, then `Save`. Their icons show next to the project.

3 While this project is active, everything created targets these accounts: the Scheduler's **New post**, Send to › Scheduler in the Library, the agent's drafts, content made from an inspiration, edits of a recording. Untick them in the post editor. With no project channel, X stays the target.

Publish for a client safely

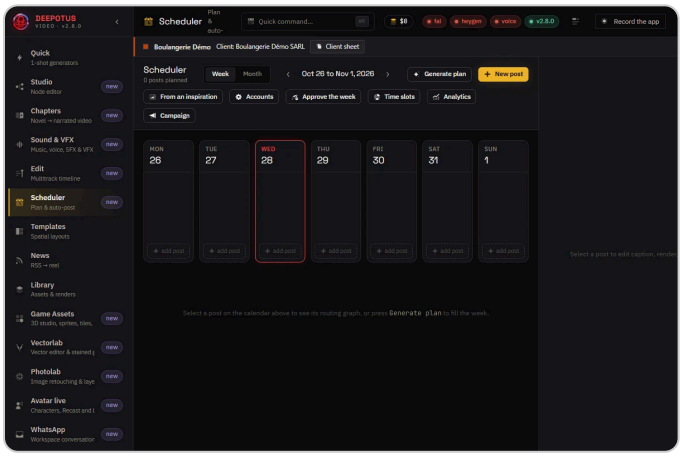

1. The banner recalls the client and the project

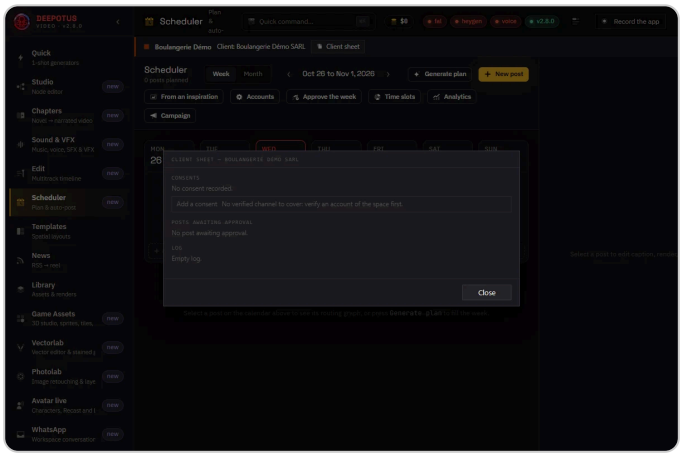

3. Consents, posts awaiting approval and the client's log

The banner says who you are working for; Client sheet opens consents, posts awaiting approval and the log.

On Quick, the Studio, the Scheduler and the Library, a banner under the header says who you are working for: the space, the active project and, for a client, its name. **Client sheet** opens their consents, posts awaiting approval and log.

One consent per channel
Each verified channel needs the client's dated consent (in writing, on WhatsApp or signed with a finger on the companion): **Add**, in the sheet. **Revoke** pauses the scheduled posts that target that channel.

Approved by the client
A post starts **Awaiting approval** and can only be scheduled once the answer is recorded: **Approved by the client** or **Needs changes**. Changing its text or media cancels the approval.

Published on their behalf
Scheduling a post or sending it right away opens a confirmation that names the client and their channels: **Publish on their behalf**.

Never the wrong account

A post only goes out on its space's accounts. If an approval, a consent or a confirmation is missing, the app refuses and says which one.

Check a channel again

Under the list of spaces, **Publishing accounts** shows each account with its badge; hovering gives the date and reason of the last check. **Verify** runs the check again, still without publishing. Rewriting a credential clears the verification. In a post's editor in the Scheduler, a ticked unverified account says so.

Good to know

Deepotus and your brands behave as before: these rules only apply to Client spaces and the products that belong to them. A client's WhatsApp is linked in the same Channels step: see Talk to your clients on WhatsApp.

Without an active consent on a channel, no post of the client goes out on that channel, even when scheduled. A consent can only be added on a verified channel: verify the account first. A revocation pauses without deleting anything; posts resume as soon as a new consent is active.

PUBLISH

Talk to your clients on WhatsApp

You will be able to link a space's WhatsApp number, answer your contacts with the help of an agent, have the client approve a post and broadcast a post to your contacts. Linking is free; Meta charges for messages, the agent's drafts are paid, and every send goes through the space's cap.

Before you start

WhatsApp goes through **Kapso**, a layer over Meta's official API that receives messages in place of the PC. Its key, a single one for all spaces, is set in `Settings` › `API keys`. Each client links their own Meta Business account and number; one number serves one space only.

Link a space's number

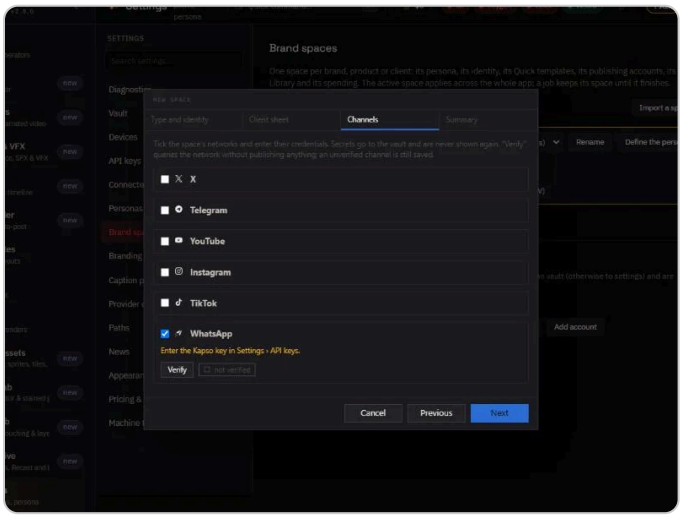

1. Channels step: tick WhatsApp

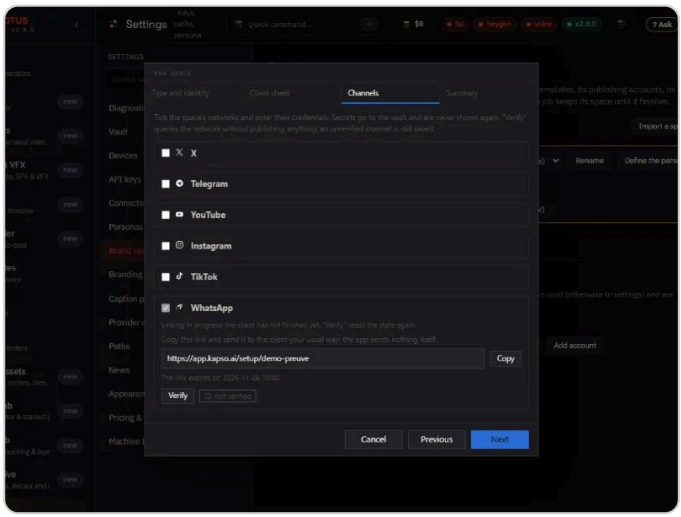

2. Copy the link and send it to the client 3. Verify reads the link again, sending nothing

Space wizard, Channels step: consent, Meta billing, Kapso setup link, then Verify.

1. In the space wizard (see Work for clients), **Channels** step, tick WhatsApp. It is not added from the list of accounts: it only shows there to be verified or deleted.
2. For a client, first record their consent to the link: how and when it was given.
3. Choose the **Meta billing**: by default the client pays Meta with their own card; otherwise Kapso pays with your credits (+5 %) and you re-invoice, if partner billing is enabled at Kapso. This choice is final.
4. `Create the link` (free): a Kapso setup link appears, valid for 30 days. `Copy` it, then send it to the client your usual way: the app sends nothing. They connect their Facebook account there in a few minutes.

5. **Verify** reads the link again: once the client has finished, the linked number shows.

Reply on the WhatsApp page

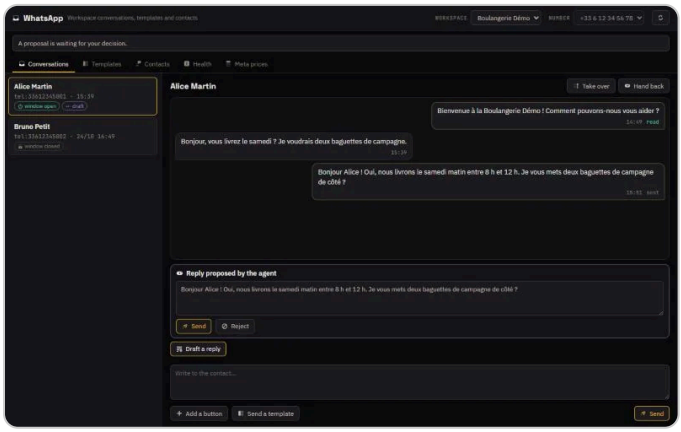

1. The agent proposes, you send or reject

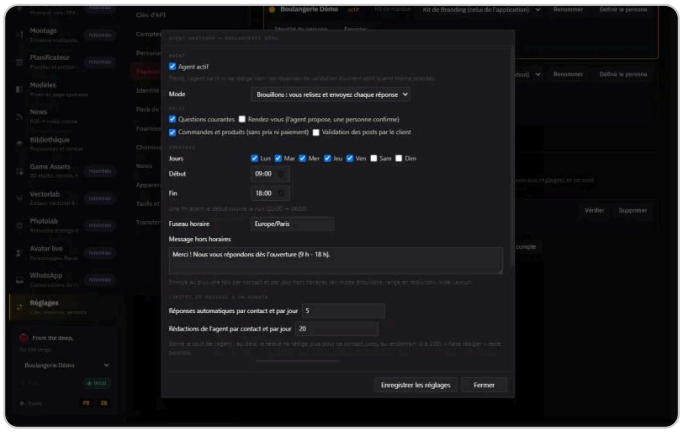

2. WhatsApp agent: mode, roles, hours and limits

WhatsApp page: the agent's proposed reply waits for you to send it; its settings live in Settings › Brand spaces.

The **WhatsApp** entry in the rail opens the active space's conversations, with five tabs: **Conversations**, **Templates**, **Contacts**, **Health** and **Meta prices**.

1. Pick a conversation. Its badge says whether the 24-hour window is open; once closed, only templates approved by Meta can go out: **Send a template**.
2. When the agent has drafted, the **Reply proposed by the agent** waits: fix it, then **Send** or **Reject**. **Draft a reply** asks it for one.
3. **Take over** silences the agent on this conversation; **Hand back** gives it back.

Set up the WhatsApp agent

In **Settings** › **Brand spaces**, the **WhatsApp agent** button, under the active space's accounts, opens its settings: **Agent on**, mode, roles (common questions, appointments, orders with no price or payment, post approval), hours and out-of-hours message, limits per contact and per day, words that call a person, monthly cap. It speaks with the space's persona and respects its banned words.

Drafts, by default

The agent drafts; you review and send each reply from the WhatsApp page.

Automatic

Once confirmed, the agent replies on its own, within hours and limits. A reply about money, a negative tone or a request for a person stays a draft to take over.

Approval by the client

With the approval role ticked, **Send to the client on WhatsApp** in the Client sheet, sends them the post preview with two buttons: approve or ask for changes. Their answer comes back dated in the sheet.

Broadcast a post on WhatsApp

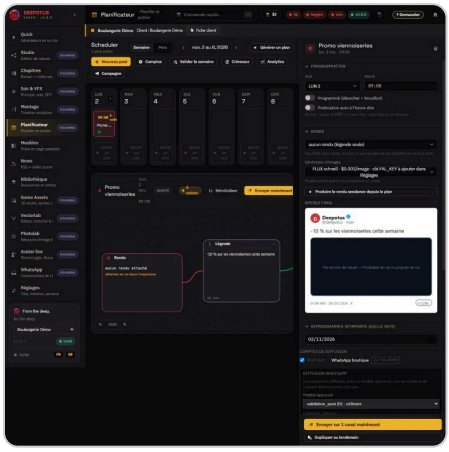

1. Tick the space's WhatsApp account

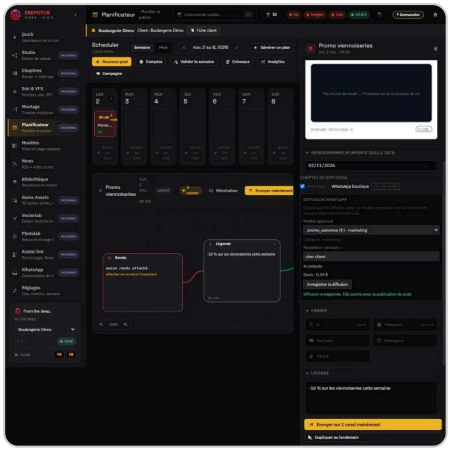

2. Template, parameters, audience and quote

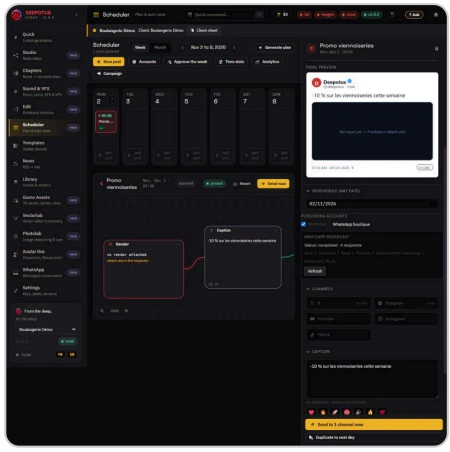

3. After sending: delivered, read, failures, actual cost

Post inspector: the WhatsApp account ticked, an approved template, audience and quote, then tracking.

1. In a post's inspector in the Scheduler, under **Publishing accounts**, tick the space's WhatsApp account, one per post. A project's channels never tick it on their own.
2. The **WhatsApp broadcast** panel asks for a template approved by Meta (no media header nor variable text header, never authentication) and its parameters, the same for everyone.
3. It shows the audience, meaning the contacts whose consent is recorded, minus STOP and marketing opt-outs, and the quote. **Save the broadcast** spends nothing.
4. The broadcast goes out when the post is published, after the cap check; the client's approval applies as for any post. Then the panel tracks sent, delivered, read, failures and actual cost: **Refresh**.

Keep your contacts and prices

Consents and STOP

Contacts tab: **Record consent**, with its source, before any first message. A contact who writes STOP gets nothing more; **Lift the STOP** requires the source of the opt-in.

Meta prices

Meta charges each delivered message by category and country. Enter the grid by market, then **Save the grid**: an unknown price makes the send refused.

Health

The month's messages, delivery failures, templates rejected by Meta and the agent's state.

Good to know

Messages count under the whatsapp category, the agent's drafts under whatsapp_agent. The Kapso plan is shared by all spaces: each agent's monthly message cap keeps one client from eating the others' quota.

If a broadcast's outcome is uncertain, **Confirm the send** retries it after a confirmation box: it is paid, some contacts may receive it twice, and the post also goes out again on its other channels.

PUBLISH

Follow the news and turn it into posts

You will follow your news feeds, let the free ranking sort today's articles, and go from article to scheduled post with the daily pipeline. Sorting, filtering and preparing are free; AI and illustrated reels are paid, on request only.

The News screen

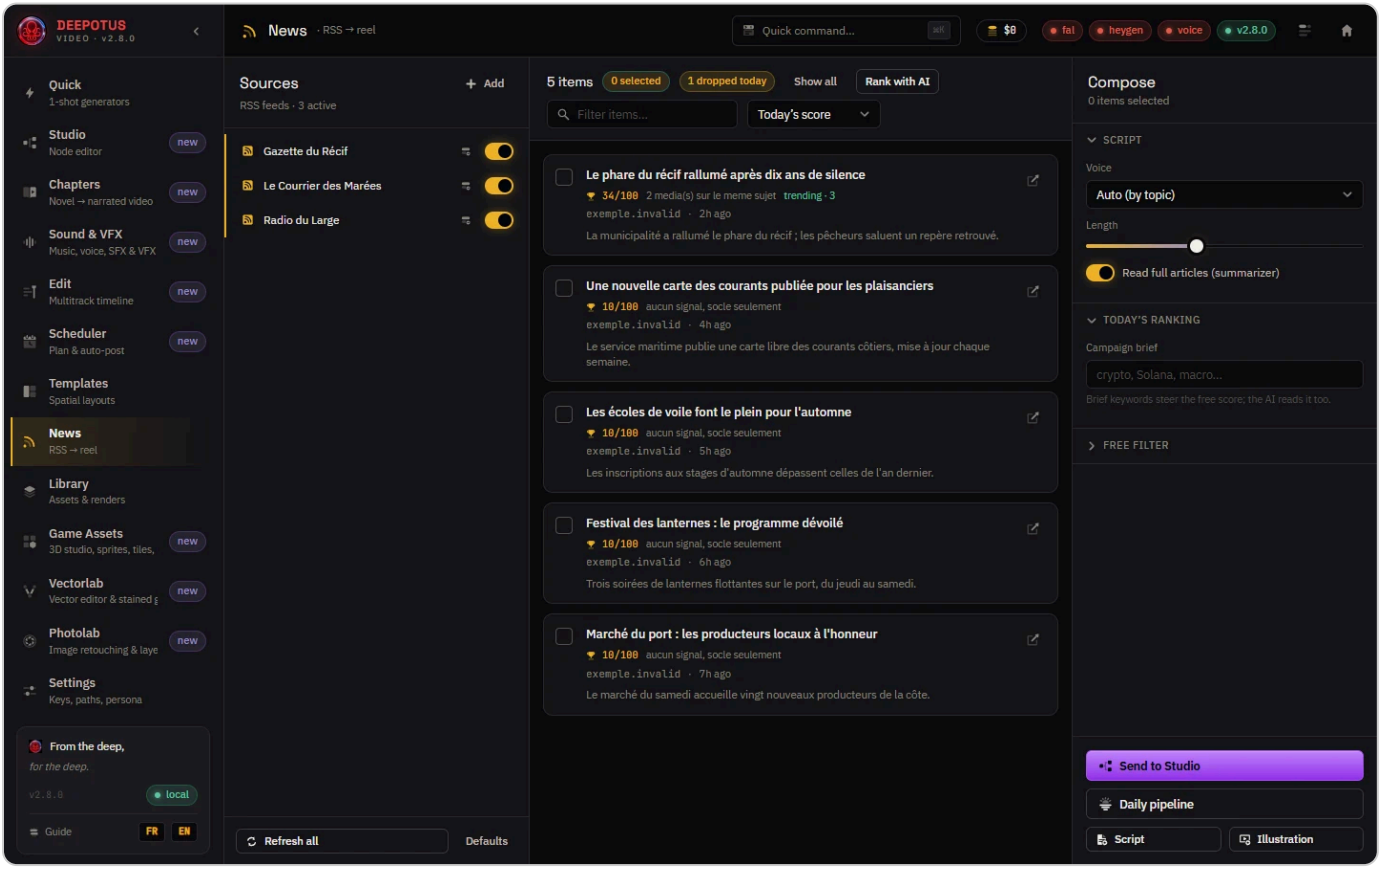

1. Picked up by three outlets: trending 2. Filtered-out items, reasons on hover 3. The score is free; AI on request

Feeds on the left, ranked articles in the middle, Compose on the right.

Open `News` in the rail. On the left, your **sources**; in the middle, the ranked **articles**; on the right, the **Compose** panel. What you ticked, the query, the sort, the voice, the length and the script are back when you return to the screen.

Choose your sources

1. `Add`: paste the address of an RSS or Atom feed. `Defaults` restores the bundled pack (crypto, geopolitics, economy, politics).
2. A source's switch pauses it without deleting it.

3 **Refresh all** re-reads every feed. The screen also does it on its own, once a day. Pickups of the same wire story are merged into a single article.

Today's ranking

By default, articles are sorted by **Today's score**, a free score: your brief's words, pickup by other outlets, a readable article. The reason and the score show on each card; **Most recent** sorts by date.

Trending

A topic picked up by three outlets the same day is marked trending, in green on the card.

Already covered

A topic close to a reel launched in the last 30 days is marked in amber, title on hover, and moves down the ranking. An overused source moves down too.

Dropped

The dropped chip counts the articles the filter removed; hover it for the reasons. **Show all** shows the full ranking, **Top only** goes back to the first five.

Rank with AI

Rank with AI re-ranks the batch with a language model: paid, under your caps, and only when you click.

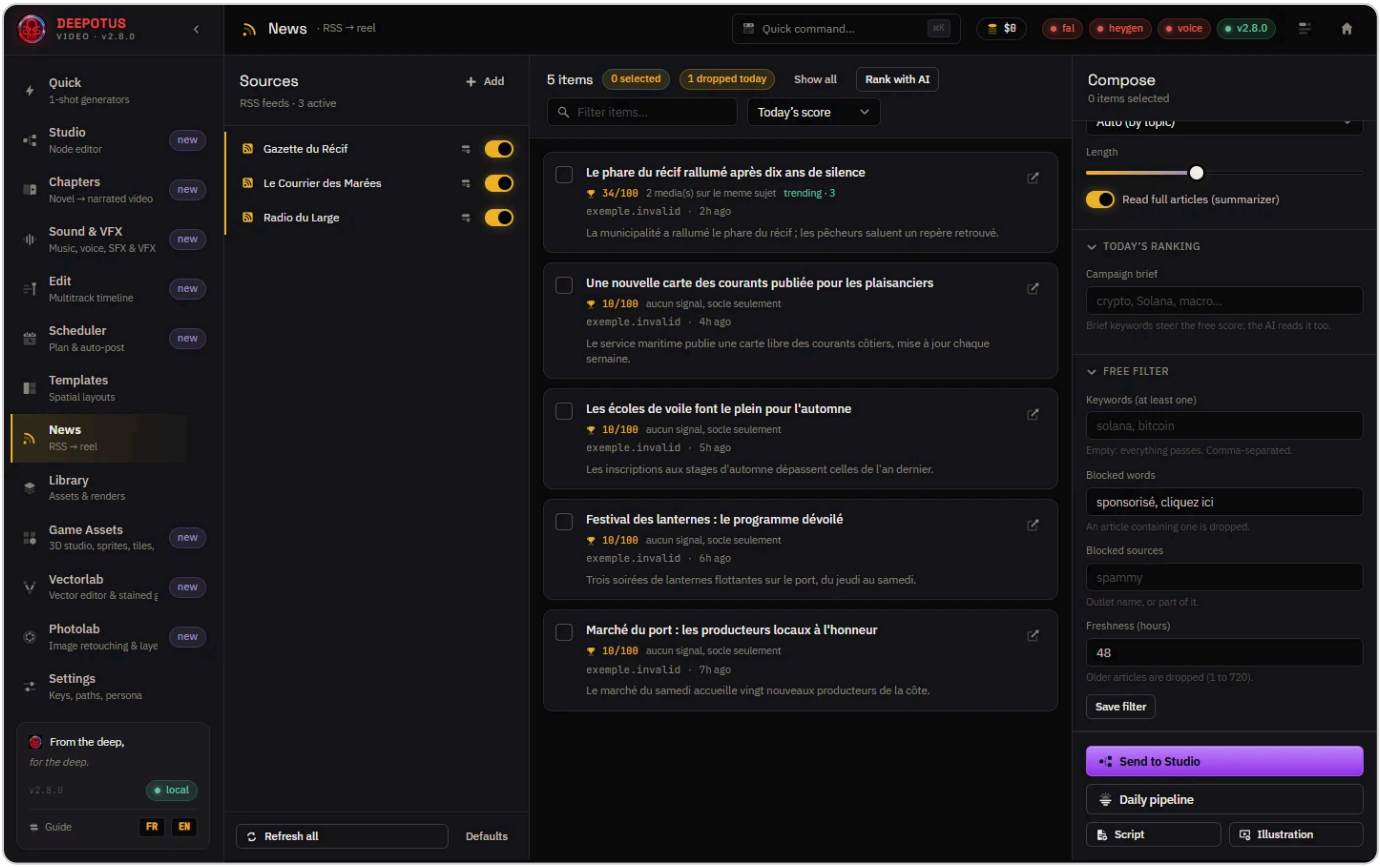

1. A blocked word drops the article 2. Saving re-ranks at once

Keywords, blocklists and freshness decide what enters the ranking.

In Compose, **Campaign brief** says what matters today and steers the score. **Free filter** sets the keywords (empty: everything passes), blocked words and sources, and freshness in hours (1 to 720); **Save filter** re-ranks at once.

From article to post: the daily pipeline

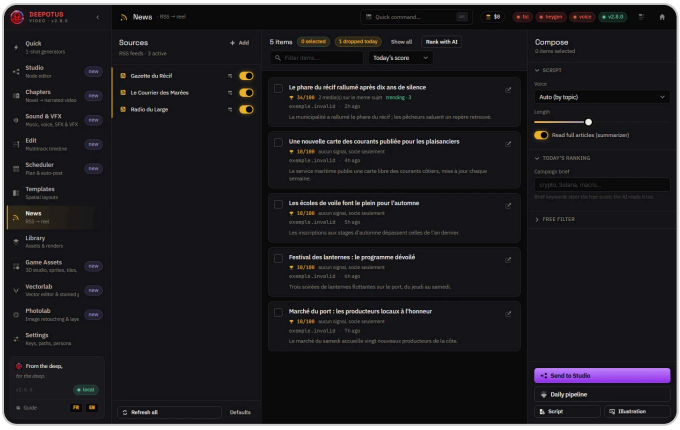

1. Daily pipeline

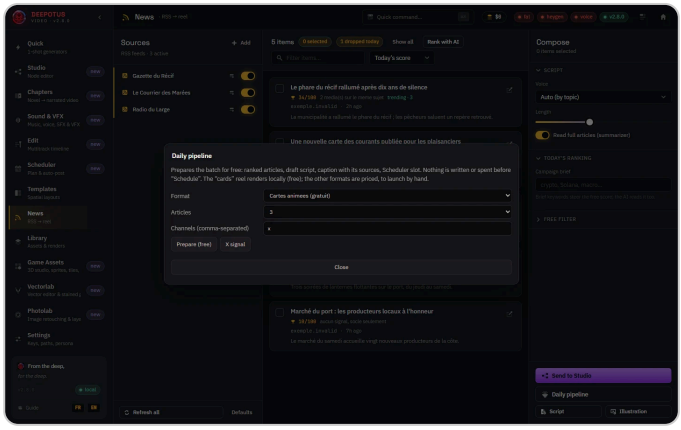

2. Preparing costs nothing

From article to scheduled post: prepare for free, polish if needed, schedule.

1. Click **Daily pipeline**. Pick the reel **format**, the number of articles and the channels. Each format's cost shows before any run: animated cards are free, rendered on your PC; AI illustration, video shots, avatar and voice-over are paid.
2. **Prepare (free)** assembles the batch without any paid call: ranked articles, draft script, caption ending with its sources (outlet, date, domain), Scheduler slot.
3. If needed, **Polish with AI** rewrites the script (paid). If the spending cap stops the polishing, the draft is kept.
4. Review the script, caption and slot, then **Schedule batch**: the post lands in the Scheduler, with its cards reel rendered for free. Nothing is written or spent before that click.

Going further

X signal

X signal reads recent X posts on your brief: on click only, three times a day at most, counted in your X read quota.

Voice by topic

With **Auto (by topic)**, the persona's tone is picked from the topic's words, for free, and the reason is shown.

To the Studio

Tick articles, then **Send to Studio** to build them yourself in a graph (see Compose with the Studio).

Good to know

Everything paid on this screen (rank with AI, polish, X signal, illustrated reels) goes through your spending caps; without a key, the button says so instead of firing.

IMAGES AND DRAWING

Organise and find your media

You will be able to find any image, sound or render, file it in a project, send it to the screen that needs it and clean up without losing anything. Everything is free and stays on your PC, except captions by a vision model (paid, quote shown).

The Library screen

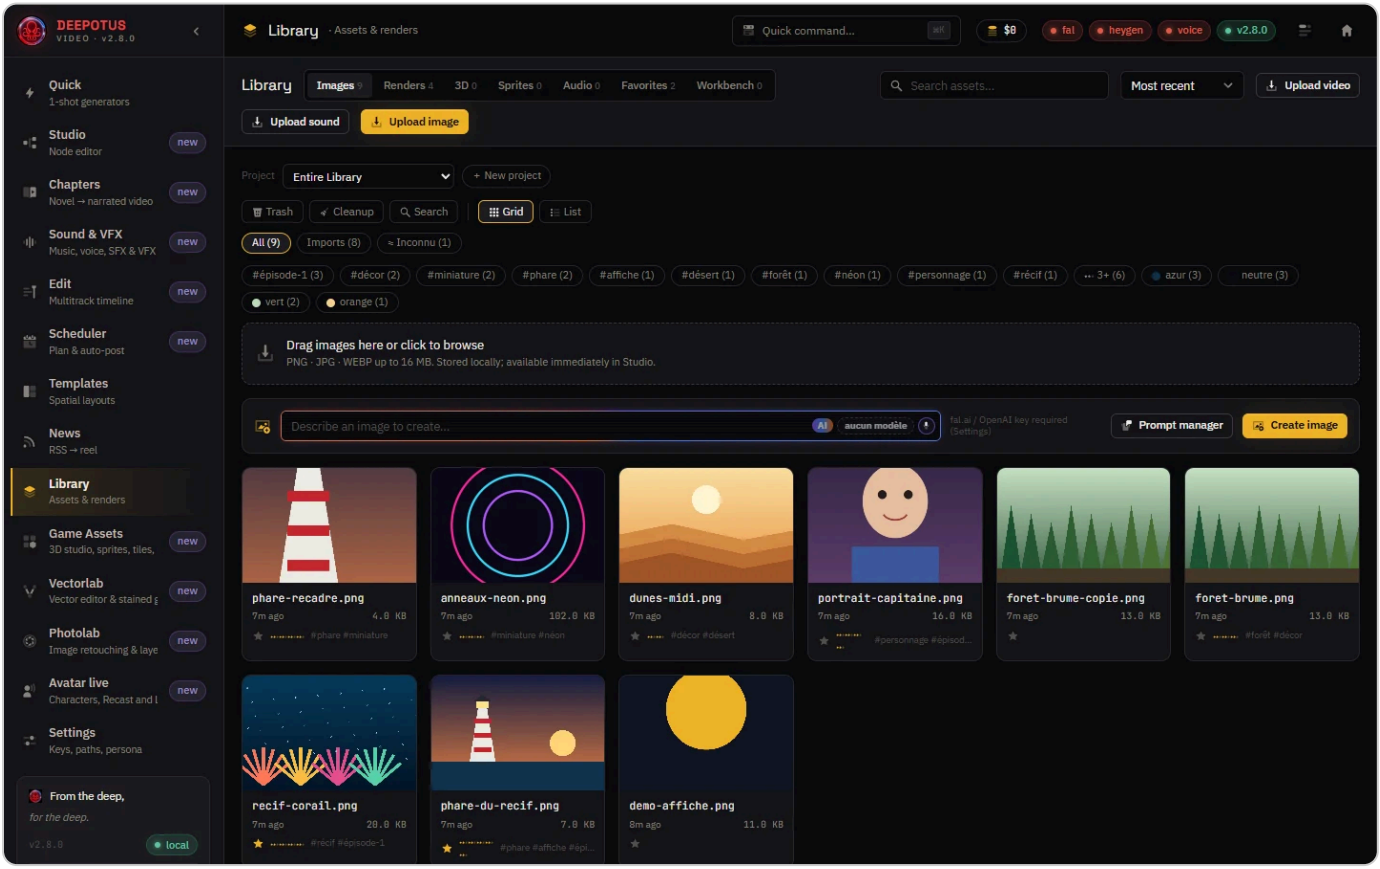

1. One tab per kind of media 2. Look at a single project 3. One click on a tag filters the grid 4. Grid or sortable list

Tabs by kind, the project you are looking at, tag and hue chips, then the grid.

Open **Library** in the rail. At the top, one tab per kind: **Images**, **Renders**, 3D, Sprites, Audio, **Favorites** and **Workbench**, where a model's versions are listed in lineage order. Everything the app produces arrives here: generators, Photolab, Vectorlab, uploads, phone.

1. **Source** chips say where each asset comes from (Generator, Materials, Uploads, mobile…). One click filters.

2. **Tag** chips are counted; “3+” keeps images rated three and up. Each image also gets its dominant colour: the **hue** chips (twelve hues and neutral) filter by colour.

3. Sort by **Most recent** or by size, then **Grid** or list: the list has ten sortable columns (name, source, size, date, dimensions, tags, rating, favourite, licence, hue) and a row opens the sheet.

4. Each thumbnail shows the file's size and age; the star sets a favourite without opening the sheet, and the dots recall the rating.

An image's sheet

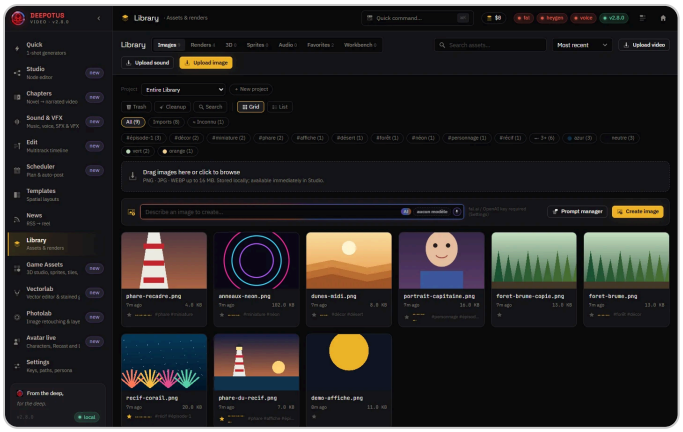

1. Open an image

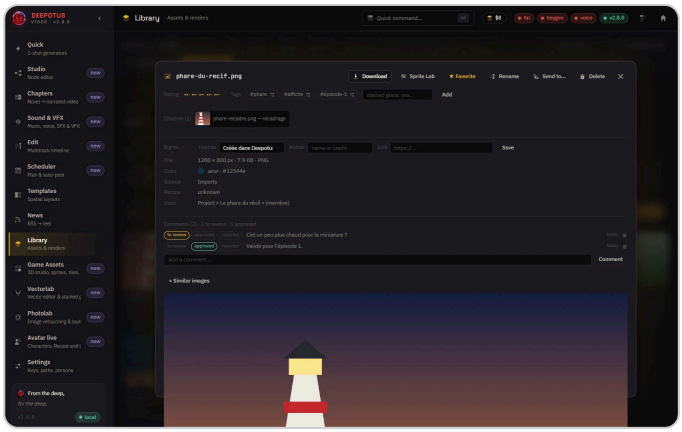

2. Its children: crop, upscale… 3. Licence, author and link 4. Review comments

One click on a thumbnail opens its sheet: rating, tags, lineage, rights, file, uses and comments.

One click on a thumbnail opens its sheet. From top to bottom you find:

Rating and tags

A rating from 0 to 5, separate from the favourite, and free tags: type, then **Add**. Search covers tags.

Parent and children

Parent shows the original image and the relation (crop, upscale, retouch…); children appear as clickable thumbnails.

Rights

Licence (suggested from the source), author and link, then **Save**. An unknown licence is flagged in amber: check it before publishing.

File, recipe, uses

Dimensions, weight, format, dominant colour, source; the generator's recipe (prompt, style, model, seed); and every use: renders, posts, bible, shots, projects.

Review comments

Type, then **Comment** (or Enter). Each comment has a clickable status: to review, approved, rejected; for a sound, the mm:ss moment. They follow the asset when renamed or restored.

Similar images

≈ Similar images shows visually close images, thanks to the CLIP model installed on your PC.

Replay recipe regenerates a generated image with its original recipe, seed included. It is paid: the price is shown and you confirm before the run; the new image arrives at the top of the list.

Send to another screen

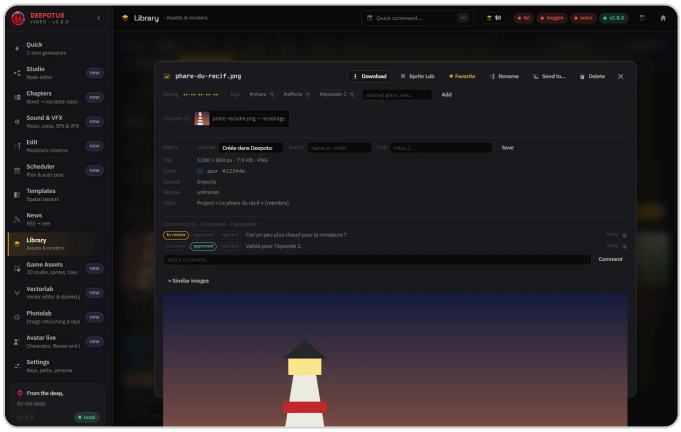

1. Send to...

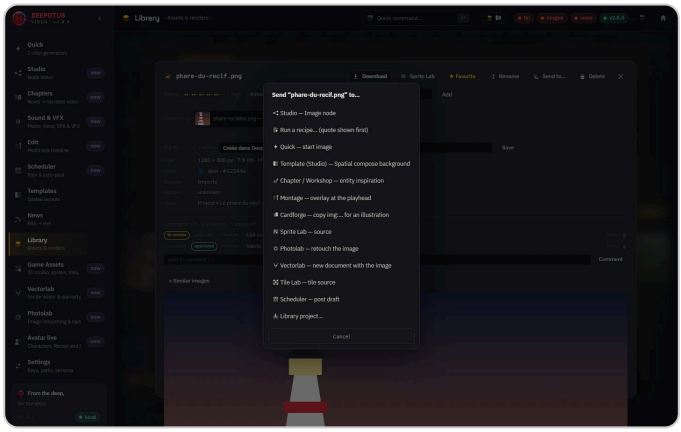

2. Retouch it in Photolab 3. Or file it in a project

From the sheet, Send to... lists every screen that can take the image, and projects.

1. In the sheet, click **Send to...**.
2. Pick the screen: Studio, Quick (start image), Template, Chapter or Workshop, Editor, Cardforge, Sprite Lab, Scheduler... Three targets are specific to images: **Photolab** to retouch it, **Vectorlab** which creates a new document at its size, and **Tile Lab** which makes it a tile's source.
3. **Library project...** files the asset in a project: the ones already holding it are ticked (one click removes it), and **+ New project...** creates one.

Redo a render from its recipe

A video render's sheet shows its **Recipe**: the exact prompt sent to the model, the negative prompt (**Show the negative prompt**), the model, duration, aspect ratio, seed and start image. For an imported

video, it is the recipe attached to the import: tool, composition, variables and the prompts of the clips it is made of.

1. Tick **Reuse the original seed** to get the same draw back; without it, a new seed gives a variation.
2. **Open in Quick** prefills the form; **Redo in Studio** opens a new graph: start image, positive and negative prompt, Seedance, render. The **Send to…** menu offers the same entries: **Quick — reuse the recipe** and **Studio — redo the generation**.
3. Nothing runs: the cost is re-estimated at today's rates, and the generation stays paid when you run it.

File by project

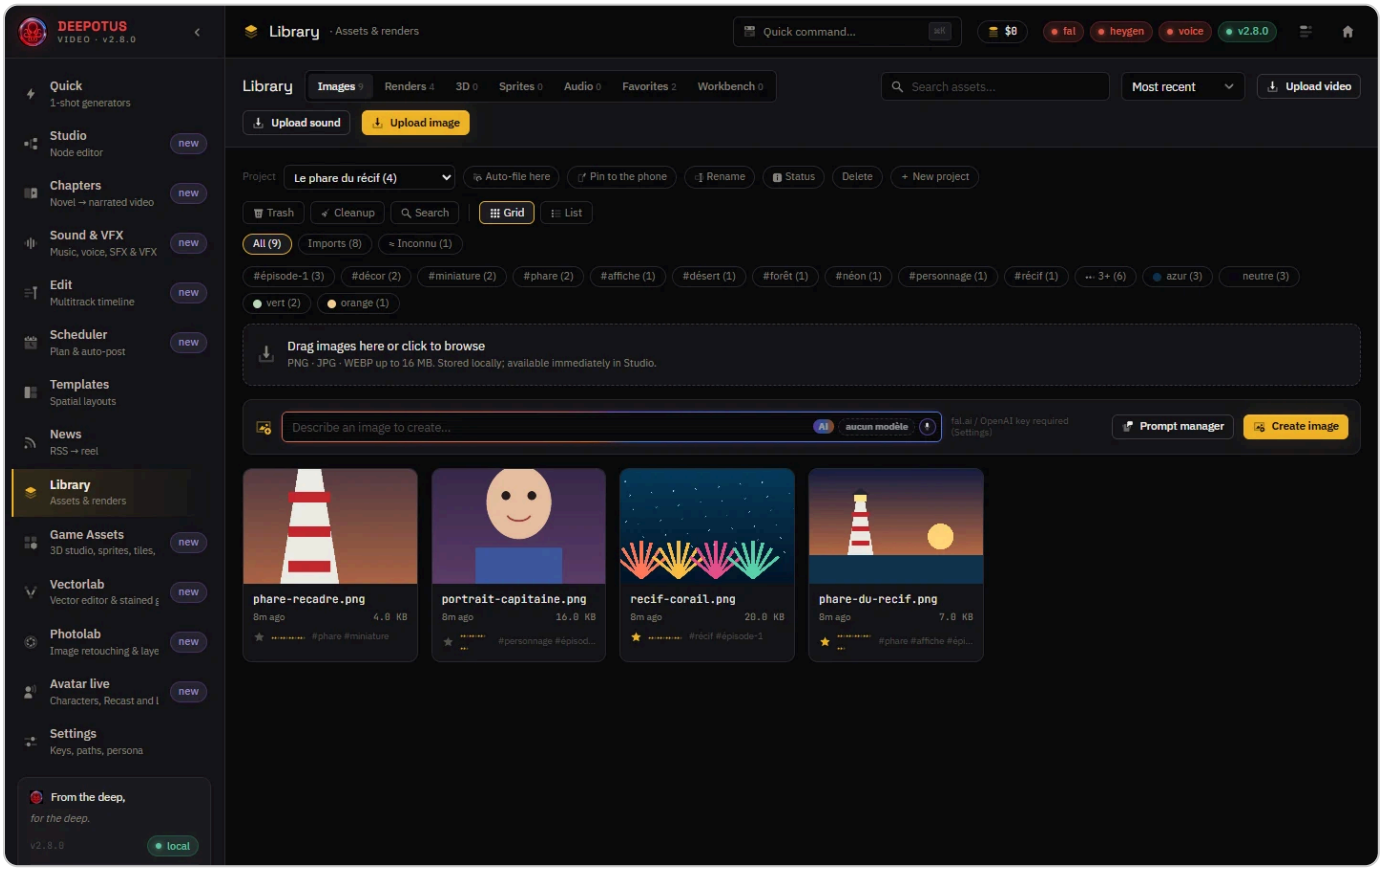

1. Le phare du récif: 4 assets 2. Everything you make is filed here 3. The mobile companion's project

Choosing a project filters every tab; its bar auto-files into it, pins it to the phone or shows its status.

The **Project** list at the top of the screen looks at a single project: every tab then only shows its assets. **+ New project** creates one; a project can be renamed or deleted without touching its assets.

Active project

Auto-file here : everything you produce next (images, renders, 3D, sprites) is filed in it. The reminder stays on screen while it is active, and one click stops it.

On the phone

Pin to the phone : a single project is pinned, and it is the one the mobile companion shows.

Project status

The Status button sorts its assets into edited, published, printed or unused.

Search

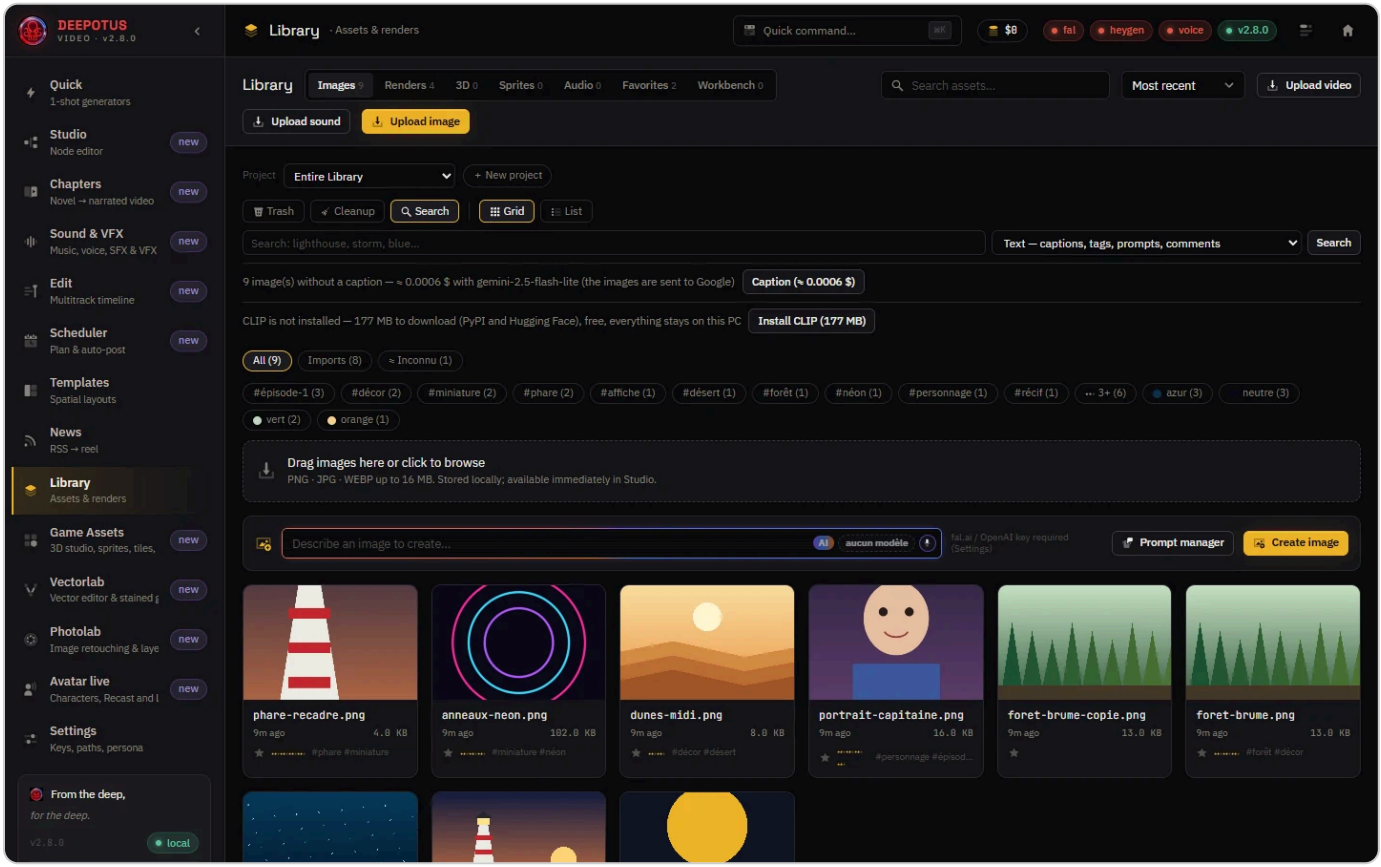

1. Captions, tags, prompts, comments 2. Paid: the quote is shown first 3. CLIP: free, on your PC

Free text search, captions by a vision model (paid, quote shown) or local CLIP, by meaning.

The Search button opens the search bar, with three engines:

1. **Text**: captions, tags, prompts, recipes, comments and names, ignoring accents and case. Free, instant.
2. **Captions**: a vision model describes images that have none, so you find them better. Paid: the quote is shown (images are sent to Google), the cap is checked, then progress counts the images done.

3. **CLIP**: searches by meaning (“a lighthouse at sunset”). **Install CLIP** downloads about 177 MB once, then indexing and search run on your PC, for free, sending nothing.

The **Audio** tab mounts the Sounds drawer (tags, filters, sorting, actions) above the upload: see Sound & VFX.

Clean up without losing anything

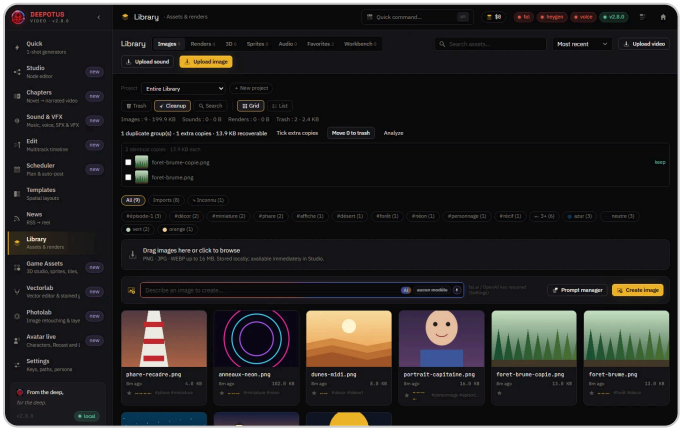

1. A group of exact duplicates

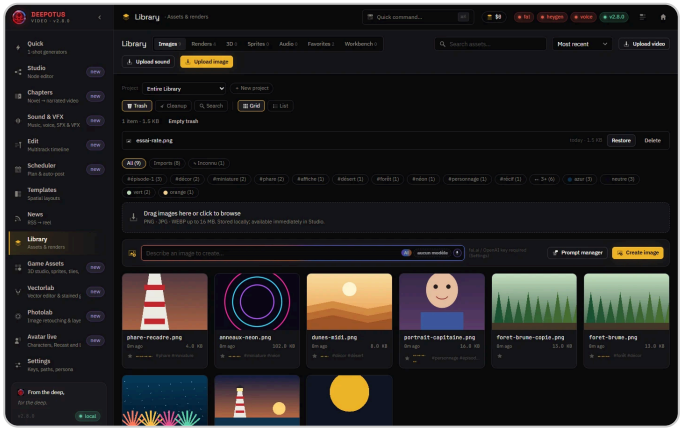

2. Nothing is lost until you empty it

Cleanup shows the weight and exact duplicates; what you discard goes to the trash, where you can restore it.

1. Deleting an image, a sound or a render moves it to the **Trash**, never further.
2. **Cleanup** shows the weight per kind and groups of exact duplicates. **Tick extra copies** keeps the used copy (or the oldest); protected copies (favourite, use, project) stay greyed out. Confirm to move them to the trash.
3. The **Trash** counts items and their weight, and flags those older than 30 days. **Restore** puts the asset back with its projects; if its name was taken, it gets a neighbouring name and says so.
4. Emptying deletes permanently: old items, everything, or a single item, always after confirmation.

Good to know

Favourites are saved by the app, no longer by the browser: they follow you from screen to screen and through a transfer. Old browser favourites were imported once.

Only **Replay recipe** and vision-model captions cost money: price shown first, spending caps checked. Sorting, text search, CLIP, projects and cleanup are free.

IMAGES AND DRAWING

Retouch a photo (Photolab)

You will be able to open an image in Photolab, select it, retouch it, adjust it and send it back to the Library or another screen. Free and offline: every gesture is computed on your PC by the photocraft engine shipped with Deepotus, with the same pixels as the original application.

Open an image

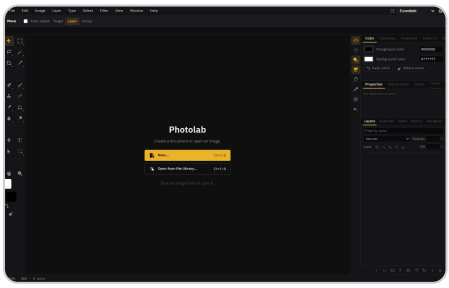

1. An image from the Library

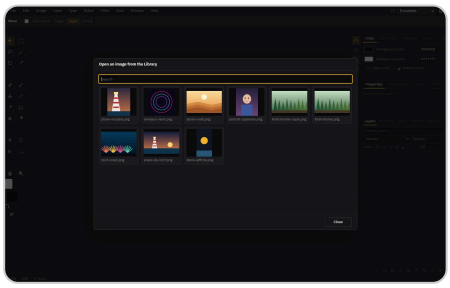

2. Pick it

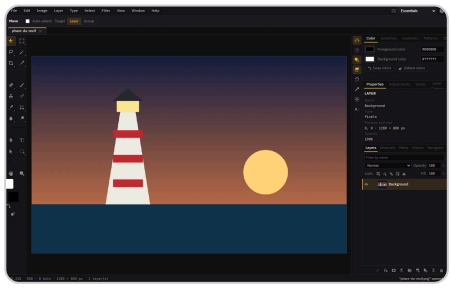

3. One tab per open document

The welcome screen offers a new document or a Library image; dropping a file works too.

1. Open **Photolab** in the app bar, right after Vectorlab. From the Library, **Send to…** › Photolab opens the image directly.
2. With no document, the welcome screen offers **New…** (`Ctrl`+`N`) and **Open from the Library…** (`Ctrl`+`O`). You can also drag a file onto the work area, or use File › **Open from computer…**: it enters the Library and opens (ORA, Krita, Affinity, Paint.NET, EXR, layered TIFF, PSD, DPX, DDS, TGA and raws included).
3. A new document is chosen by category (Photo, Film and video, Mobile, Print, Web) and preset: HDTV 1080p, 4K UHD, Instagram square, Story 9:16, A4 300 ppi…
4. Each document has its tab; a dot means it has changed. The status bar at the bottom gives the zoom, mode and depth (RGB · 8 bits), size in pixels and number of layers.

The screen follows the Deepotus language, French or English: menus, tools, panels, dialogs and history.

Pick a tool

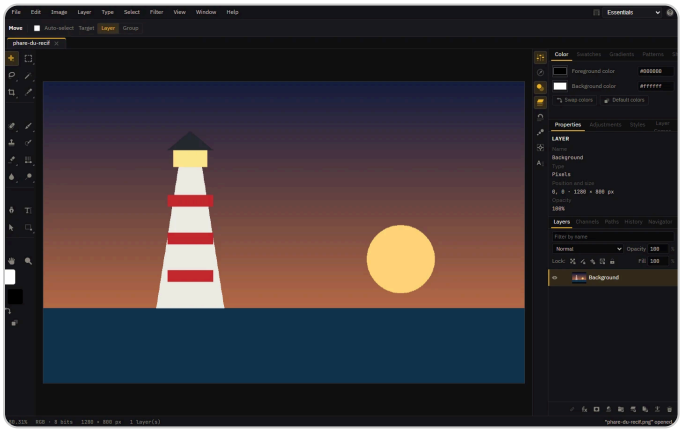

1. Magic Wand (W)

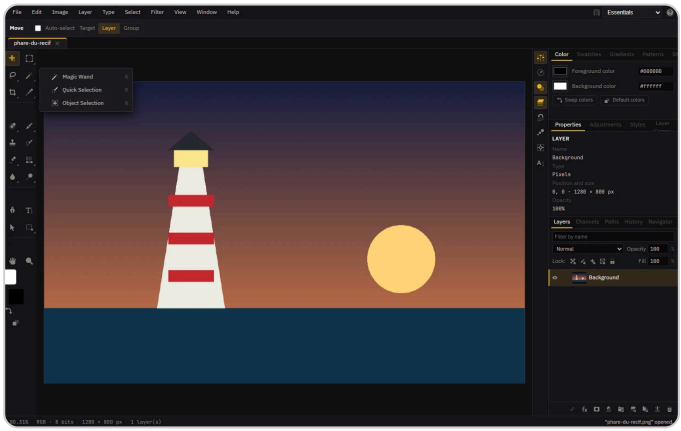

2. The other tools in the group

A slot marked with a corner holds several tools: right-click or long-press to choose.

The toolbar on the left has twenty slots. A slot marked with a corner holds several tools: right-click or long-press to choose. Each tool has its letter (V, M, L, W, C, I, J, B, S, Y, E, G, O, P, T, A, U, H, Z), and the same letter cycles to the next tool in the group. The strip at the top shows the active tool's settings.

Select

Rectangle, ellipse, a one-pixel row or column (M), lasso, polygonal lasso and magnetic lasso (L), magic wand, quick selection and object selection (W). `Shift` adds, `Alt` subtracts.

Crop and measure

Crop (C): draw the frame, `Enter` to cut; perspective crop straightens at the same time. The eyedropper (I) picks a colour, the colour sampler places measuring points read in Info; the ruler, count and notes are in the next chapter.

Paint

Brush, pencil and mixer brush (B), eraser and background eraser (E), gradient and paint bucket (G). `[` and `]` set the size, a digit the opacity (4 then 5 = 45%).

Repair

Spot healing, healing brush, remove, patch, content-aware move and red eye (J), clone and pattern stamps (S, `Alt`+click for the source), history brush (Y), dodge, burn, sponge, blur, sharpen and smudge (O).

Text and shapes

Type (T): click for a line, drag for a paragraph. Vertical type and type masks in the same slot. Pen, freeform pen, curvature pen and point tools (P), path selection and direct selection (A), rectangle, ellipse, triangle, polygon, line and custom shapes (U).

Move around

Hand (H) or hold `Space`, rotate view (R); zoom (Z) or `Ctrl`+wheel. `Ctrl`+`0` fits the screen, `Ctrl`+`1` shows 100%.

The swatches under the toolbar hold the foreground and background colours: `X` swaps them, `D` restores black and white.

Rectangular marquee (M)

Drag across the picture: a dotted rectangle surrounds the part you are going to change.

Magic wand (W)

Click a colour: the whole patch of that colour is selected at once.

Crop (C)

Draw a frame on the picture then press Enter: everything outside the frame disappears.

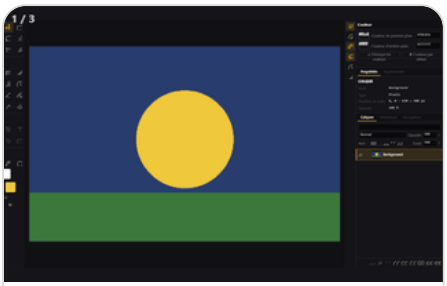

Brush (B)

Drag across the picture: a stroke of your colour is painted wherever you go.

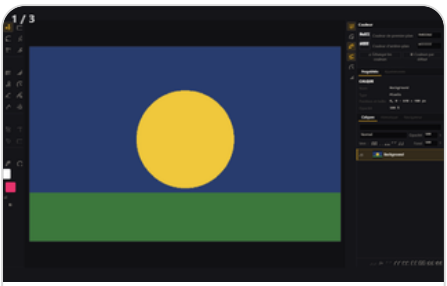

Clone stamp (S)

Alt-click what you want to copy, then paint elsewhere: the same piece of picture reappears under your brush.

Every tool has its animated card: hover it for a second in the bar or in its group.

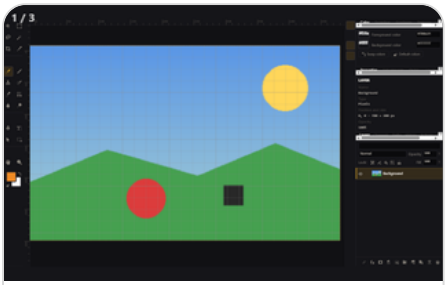

Remove (J)

Brush over what's in the way: the object vanishes and the surroundings take its place.

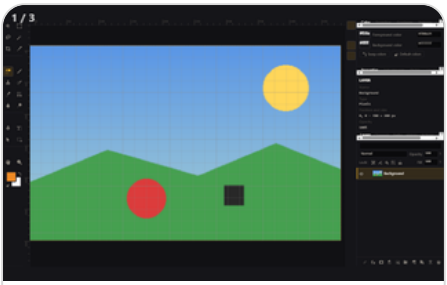

Patch (J)

Circle the spot, then drag your outline to a clean area: the spot takes on the look of that place.

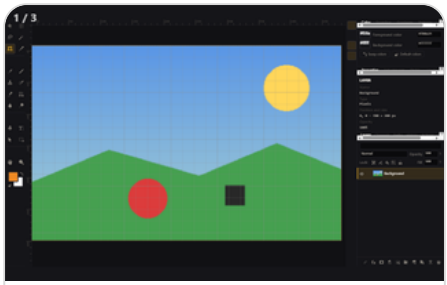

Perspective Crop (C)

Draw a frame, pull its corners onto the tilted shape, then press Enter: the picture is straightened and cut.

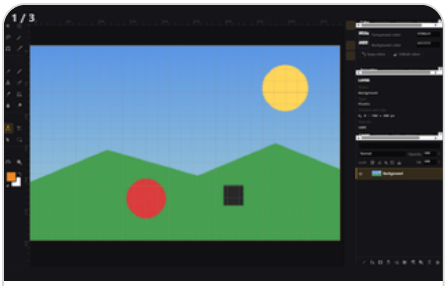

Curvature Pen (P)
Place a few points: the pen draws soft curves through each one by itself.

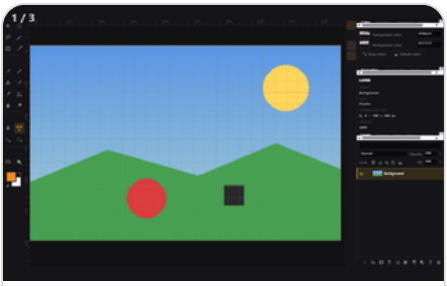

Horizontal Type Mask (T)
Type a word: its letters become a selection that you can then fill with color.

Work the selection

The **Select** menu holds **All** (`Ctrl`+`A`), deselect (`Ctrl`+`D`) and inverse (`Ctrl`+`Shift`+`I`). The engine can also find the subject, the sky, a colour range or the focus area. Modify › Border, Smooth, Expand, Contract and Feather refine the outline; a selection can be saved to a channel and loaded back.

Select and Mask… (`Alt`+`Ctrl`+`R`) refines difficult edges, hair included: seven views (onion skin, on black, on white…), radius, smoothing, feather, colour decontamination, then output to a selection, a mask or a new layer.

Transform

Free Transform (`Ctrl`+`T`) sets an eight-handle frame. A corner scales and keeps proportions (`Shift` frees them), an edge stretches one axis, outside rotates (`Shift` by 15°), `Ctrl` distorts or skews. The strip gives X, Y, width, height, angle and interpolation; `Enter` confirms. The **Warp** button in the strip sets a mesh: drag its points and handles to bend the image. Edit › Transform also offers 90° and 180° rotations and flips.

Finally, the Edit menu holds cut, copy, paste (including paste in place), fill and stroke, content-aware fill, **Fade** to soften the last filter, and the **Search** palette (`Ctrl`+`K`) that finds any command by name.

Adjust and filter

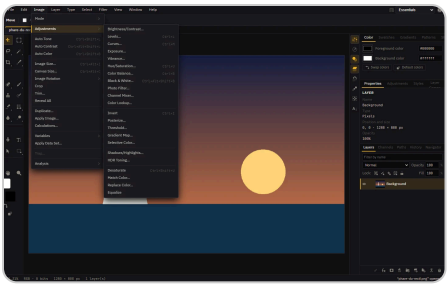

1. Image › Adjustments › Curves…

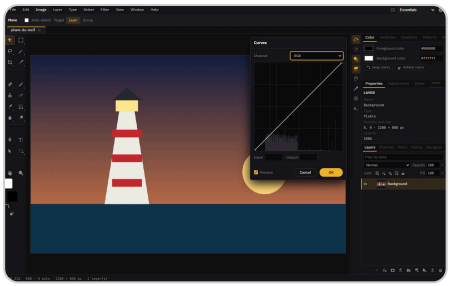

2. Click to add a point

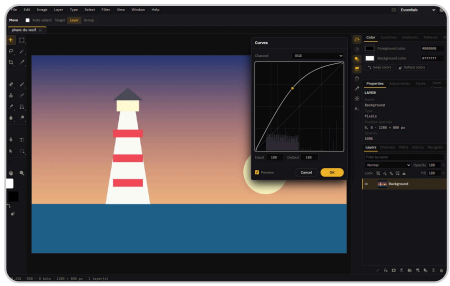

3. The preview follows every move

Image › Adjustments › Curves…: every point you move is recomputed by the engine on a copy, until OK.

1. Open **Image** › **Adjustments** and pick an adjustment: **Levels** (`Ctrl`+`L`), **Curves** (`Ctrl`+`M`), Hue/Saturation (`Ctrl`+`U`), Color Balance, **Black & White**, Exposure, Vibrance, Shadows/Highlights…
2. The dialog shows a **Preview** computed by the engine on a copy. In Curves, a click adds a point, dragging moves it, a double-click removes it, per RGB, red, green or blue channel.
3. **OK** applies and adds a history step; **Cancel** leaves the image untouched.

The **Filter** menu works the same way: blur (including the blur gallery), sharpen, noise, distort, pixelate, render, stylize, plus the **Filter Gallery…** and its 47 stackable artistic effects. Last Filter replays with `Ctrl`+`Alt`+`F`. The Image menu also gives automatic corrections, image and canvas size, rotation, and Image › Mode to switch to grayscale, CMYK, Lab, duotone, indexed colour or 16 bits.

Black and white
One click adds a layer that removes the colours: your picture turns grey, and you can remove it whenever you like.

Save and send

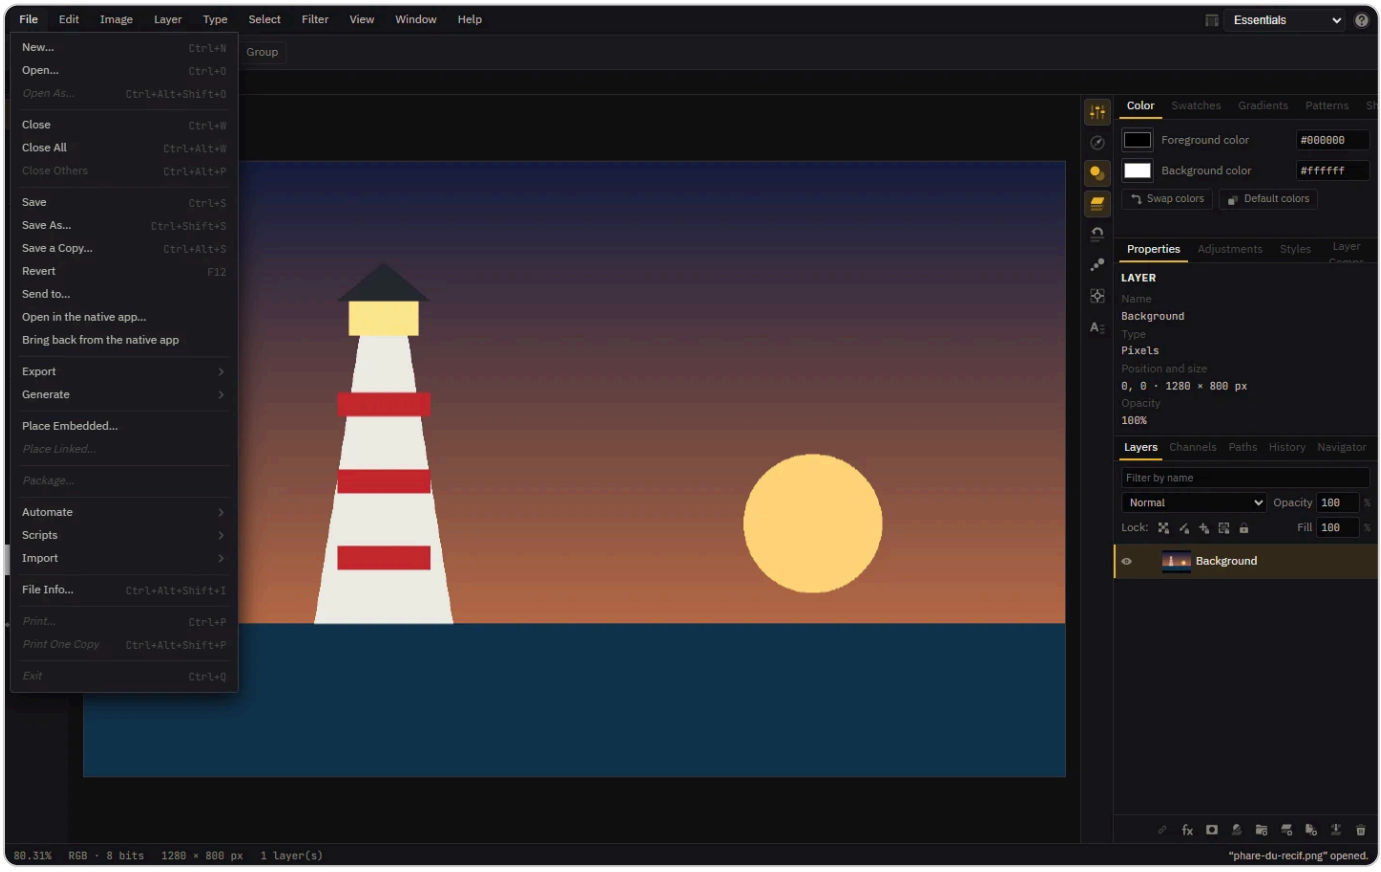

1. To the Library, with layers 2. To Vectorlab, Tile Lab… 3. The fallback: the native app

The File menu: save to the Library (with layers), export, or send to another screen.

1. **Save** (`Ctrl`+`S`) files the image as PNG in the Library, source “Photolab”, with a working file that keeps the layers. It never overwrites: each save is a new file, linked to the original image in its sheet.
2. Reopened from the Library, a Photolab image comes back with its layers.
3. **Save As…** exports to PNG, JPG, WEBP, TIFF (layers kept), PSD, ORA, single-page PDF or PCRAFT and downloads it; **Save a Copy…** leaves the document as is; Export › Quick Export as PNG outputs the image in one click, following your preferences.
4. **Send to…** passes the image to another screen, through the Library.

To output each layer as a file or process a whole batch of images, see Automate and export.

Close (`Ctrl`+`W`), Close All and Close Others ask for confirmation on a changed document. **Revert** (`F12`) reloads the last save. File Info… holds title, author, keywords and copyright.

Good to know

What the screen cannot do yet opens in the native application: **Open in the native app…**, then **Bring back from the native app** returns the document, layers included.

If the engine must restart, the screen says so and returns to the welcome screen: unsaved documents are lost. Save often. The installer checks every engine file; Help › **About Photolab** gives its version and licences.

Shortcuts

Gesture	Shortcut
Undo, redo	`Ctrl`+`Z`, `Ctrl`+`Shift`+`Z`
New, open, save	`Ctrl`+`N`, `Ctrl`+`O`, `Ctrl`+`S`
Free transform	`Ctrl`+`T`
Levels, curves, hue	`Ctrl`+`L`, `Ctrl`+`M`, `Ctrl`+`U`
Tip size, hardness	`[` `]`, `Shift`+`[` `Shift`+`]`
Search a command	`Ctrl`+`K`

IMAGES AND DRAWING

Layers, masks and workspaces

You will be able to work with non-destructive layers in Photolab: adjustments, masks, styles and smart objects, then arrange the screen your way with panels, workspaces, rulers and preferences. Everything is free and computed on your PC.

The Layers panel

Each layer has its eye, thumbnail and name; double-click the name to rename it. Click to pick, `Ctrl`+click to add, `Shift`+click for a range. Drag a row to place it above, below or inside a group. At the top of the panel: a field to filter by name, the blend mode (27 modes), opacity, fill and locks (transparent pixels, pixels, position, nesting, all). `Alt`+click on an eye shows only that layer; `/` locks its transparent pixels.

The panel footer creates a layer or a group, links layers, adds a mask, duplicates, merges down and deletes. The **Layer** menu goes further: new layer via copy (`Ctrl`+`J`), group, clipping mask (`Ctrl`+`Alt`+`G`), arrange, align and distribute, merge (`Ctrl`+`E`) or flatten, rasterize, matting, **Layer Color...**, **Duplicate to Document...**, **Merge Group**, and **Remove Background (local)...**, computed on your PC, free.

Add an adjustment layer

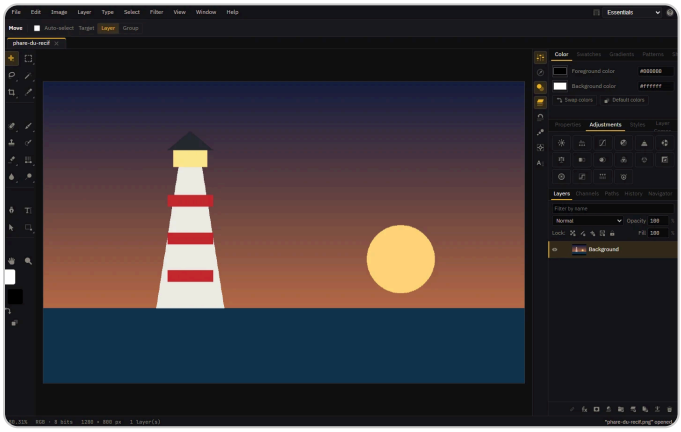

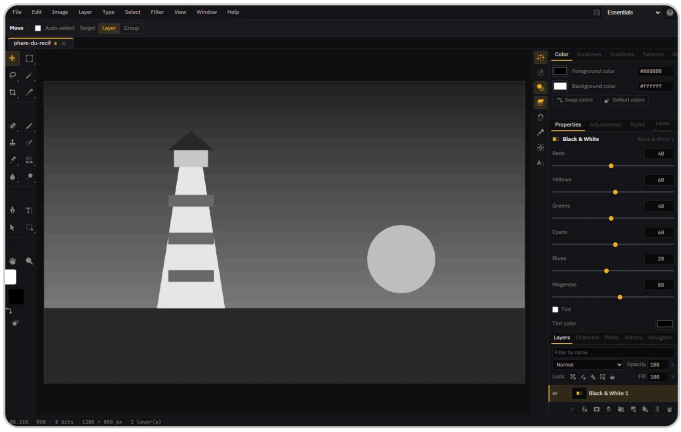

1. Black & White

2. A non-destructive adjustment layer 3. Its sliders in Properties

Adjustments tab: one click adds an adjustment layer, editable at any time in Properties.

1. Open the **Adjustments** tab, next to **Properties**.
2. Click an adjustment, for example **Black & White**: an adjustment layer is added above the active layer. The image below is not touched.
3. Its sliders appear in **Properties**; Curves and Levels have their own editor there. Come back at any time by picking the layer.

Sixteen adjustments work this way (brightness, levels, curves, hue, colour balance, channel mixer, gradient map…), as well as fill layers: solid colour, gradient or pattern.

Masks, styles, smart objects

Layer mask

The mask button reveals the selection (or everything); `Alt` does the opposite. On the mask thumbnail: `Shift`+click disables it, `Alt`+click shows it alone, `Ctrl`+click loads it as a selection.

Vector mask

On a layer that already has a mask, the same button adds a vector mask, or Layer › Vector Mask from the active path.

Layer style

Drop shadow, glows, bevel, stroke, satin, overlays: a dialog with preview. A style can be copied, pasted or cleared; the **Styles** panel offers ready ones (neon, chrome, gold…).

Smart filters

Filter › Convert for Smart Filters: every filter applied afterwards appears under the layer, with its eye; double-click to reopen its settings, drag to change the order.

A single layer can be exported: Layer › Export As… (PNG, JPG or WEBP, at the chosen scale, to download or the Library) or quick export as PNG.

Black and white

One click adds a layer that removes the colours: your picture turns grey, and you can remove it whenever you like.

New AI layer

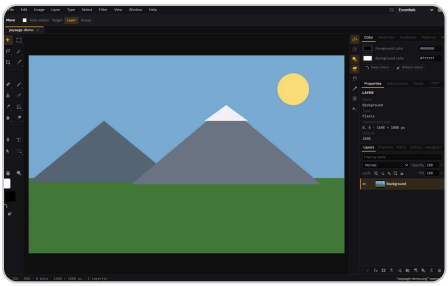

1. New AI layer…

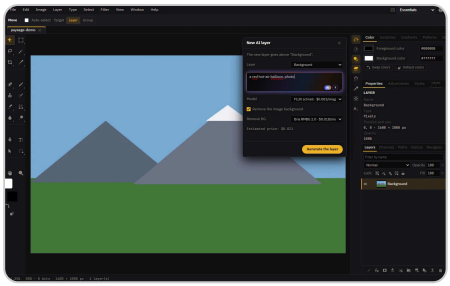

2. Where the new layer goes 3. The total price, shown beforehand

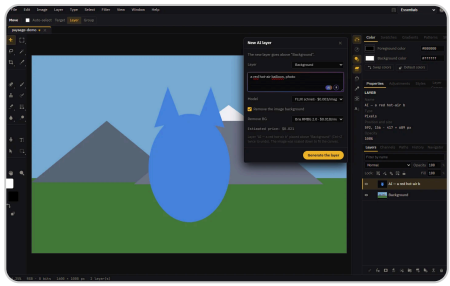

4. A pixel layer above the background

The AI button at the foot of the layers: prompt, model, background removal and price; the layer goes above the chosen layer.

1. Click **New AI layer…** at the foot of the Layers panel, or in the **Layer** menu.
2. Describe the image, pick the model and the layer to place it above; tick **Remove the image background** for a subject without a background. The total price shows beforehand.

3 **Generate the layer**: a pixel layer "AI — …" lands above the chosen layer. Two `Ctrl`+`Z` undo it.

Remove a layer's background

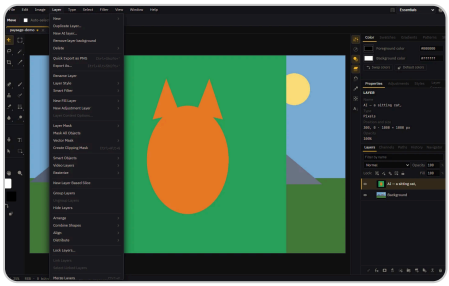

1. Layer › Remove layer background

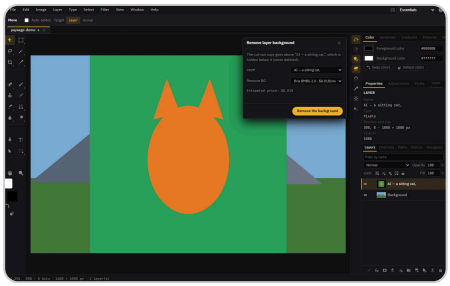

2. The chosen model and its price

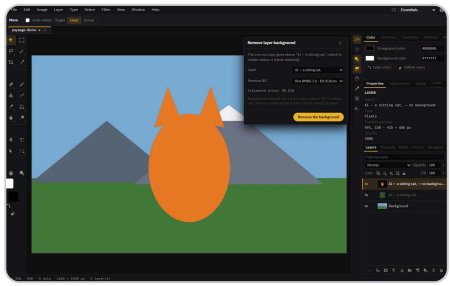

3. The cut-out copy on top, the original hidden below

Layer › Remove layer background: the cut-out copy goes on top, the original stays below, hidden.

Layer › **Remove layer background** cuts out the chosen layer without destroying it: the cut-out copy lands on top, the original stays below, hidden, never deleted. The removal model and its price show beforehand; three `Ctrl`+`Z` go back.

The panels

On the right, panels are arranged in tabbed groups; the rail's icon column collapses or reopens each group, and the **Window** menu ticks the visible ones.

Colour and swatches

Color (`F6`) with hex code, Swatches (personal groups kept by Deepotus, .aco or .ase import and export), Gradients and Patterns.

History and info

History: a click on a step goes back to it; **New snapshot** keeps a state to reopen later. Navigator, per-channel Histogram, Info (`F8`) under the pointer, Channels and their shortcuts (`Ctrl`+`2`, `Ctrl`+`3`…).

Brushes

Tips with search, Brush Settings (`F5`), Clone Source and Tool Presets (e.g. soft eraser 60 px).

Text and paths

Character, Paragraph, Glyphs, character and paragraph styles; Paths, Shapes; Layer Comps to switch from one state to another.

Choose a workspace

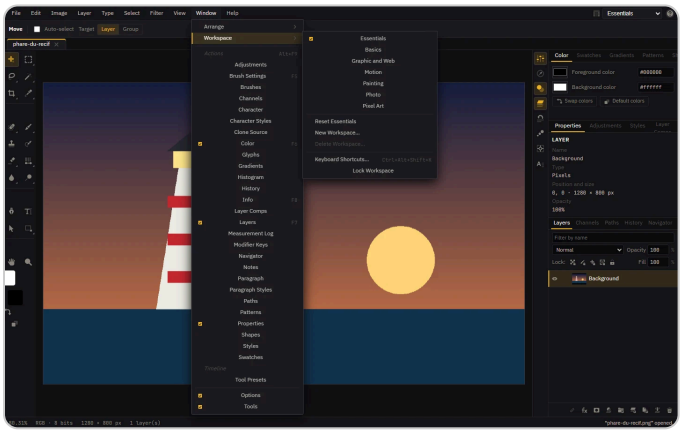

1. Seven built-in workspaces

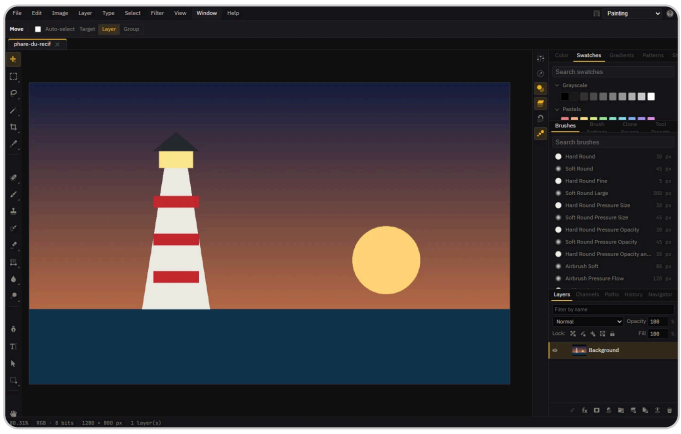

2. Painting: brushes and colours up front

Window › Workspace arranges panels and tools for a task; the layout is remembered per workspace.

1. Open **Window** › **Workspace**, or the list at the top right.
2. Pick one of seven layouts: **Essentials** (default), Basic, Graphic and Web, Motion, **Painting**, **Photo**, **Pixel Art**. Panels and tools arrange for the task.
3. What you move, collapse or hide is remembered for that workspace. Reset restores it; Lock Workspace stops remembering.
4. New Workspace… saves your layout (up to twenty), kept in the Deepotus data folder.

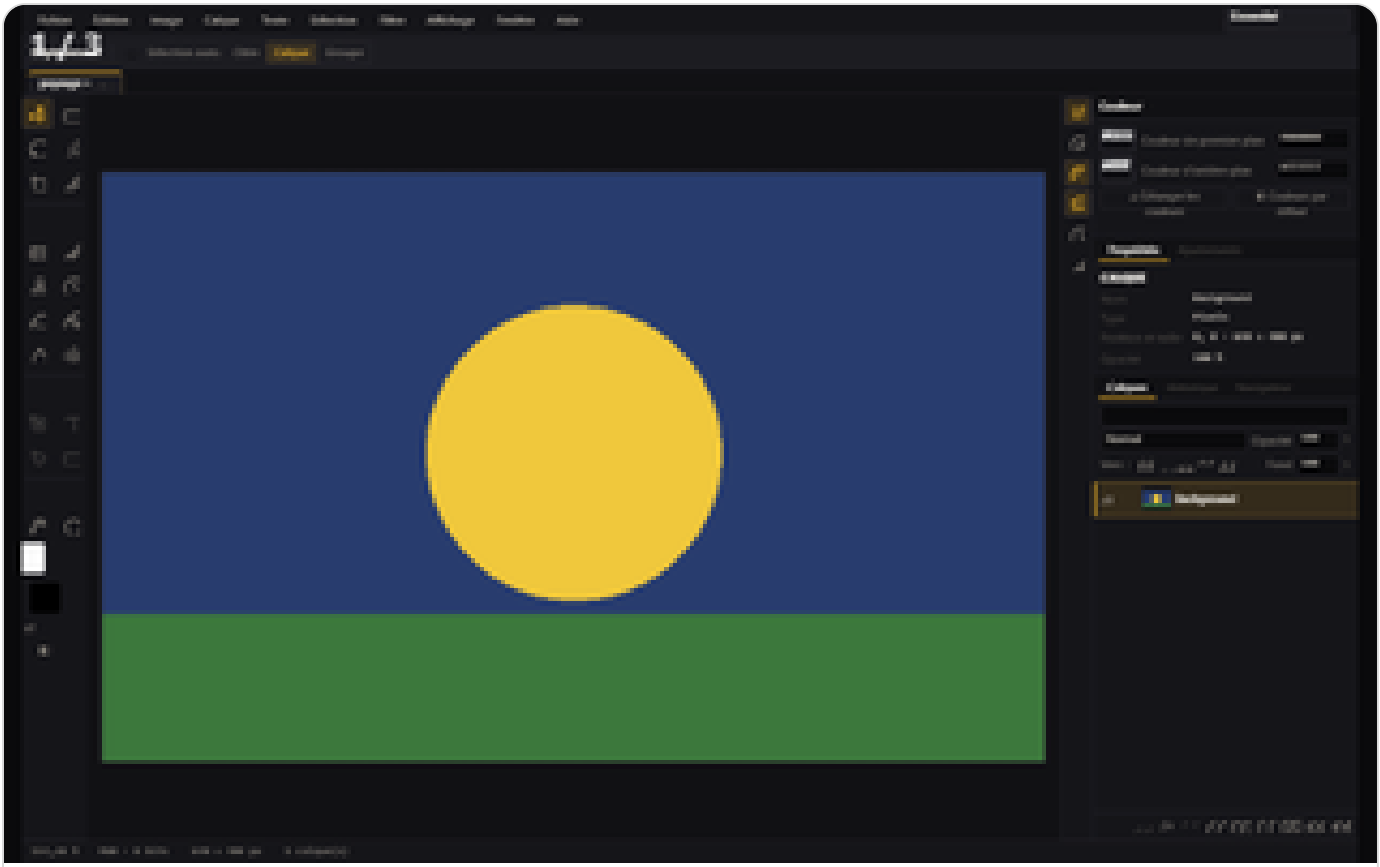

Workspace

Pick a workspace in the list: panels and tools rearrange for what you want to do.

Rulers, grid and guides

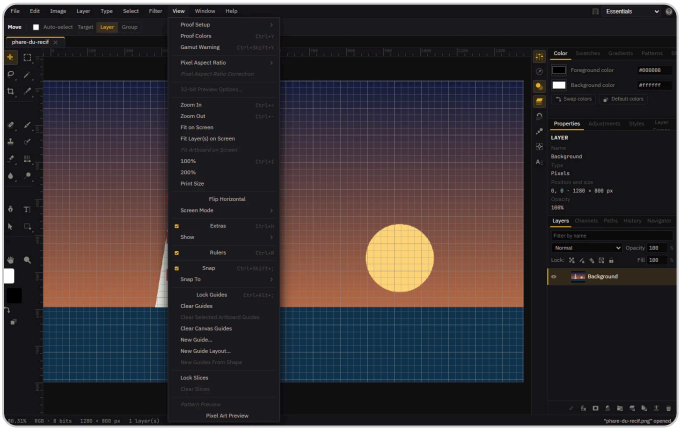

1. View › Rulers

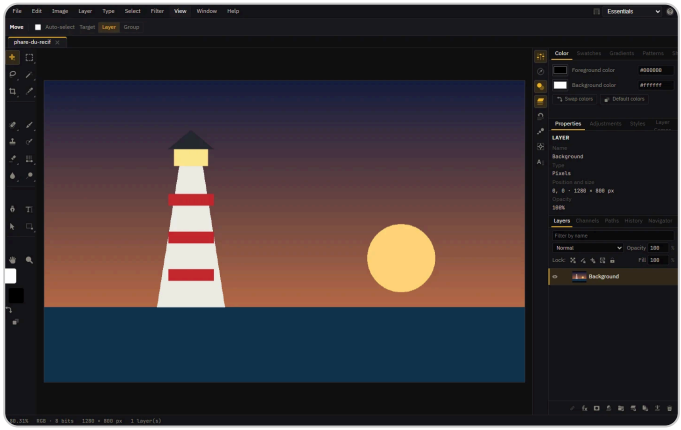

2. Graduated rulers and grid

View › Rulers (Ctrl+R) then Show › Grid: drag from a ruler to place a guide.

In **View** : **Rulers** (`Ctrl` + `R`), Show › **Grid** (`Ctrl` + `'`) and **Guides** (`Ctrl` + `;`). Drag from a ruler to place a guide; the Move tool moves it, or removes it by dragging it out. Snap (`Ctrl` + `Shift` + `;`) aligns selections and layers to guides, grid, layers or document edges. Extras (`Ctrl` + `H`) hides everything at once.

The same menu offers zooms, horizontal flip of the view (the document does not change), pixel art preview with crisp pixels and screen modes: `F` goes to the next one, `Esc` back to standard mode.

Measure, annotate, slice

Ruler (I)

Drag to measure distance and angle; `Alt` from an end makes a protractor. Straighten Layer puts the line level.

Count (I)

A click places a numbered mark, `Alt`+click removes it. Measurements are recorded in the log, exportable as CSV for a spreadsheet.

Note (I)

A note with author and colour, placed on the image; the Notes panel steps from one to the next. File › Import › Notes… copies them from another document.

Slices (C)

Cut the image for the web, or from guides; slice selection moves them, divides them and sets their options (link, alt text).

Preferences and shortcuts

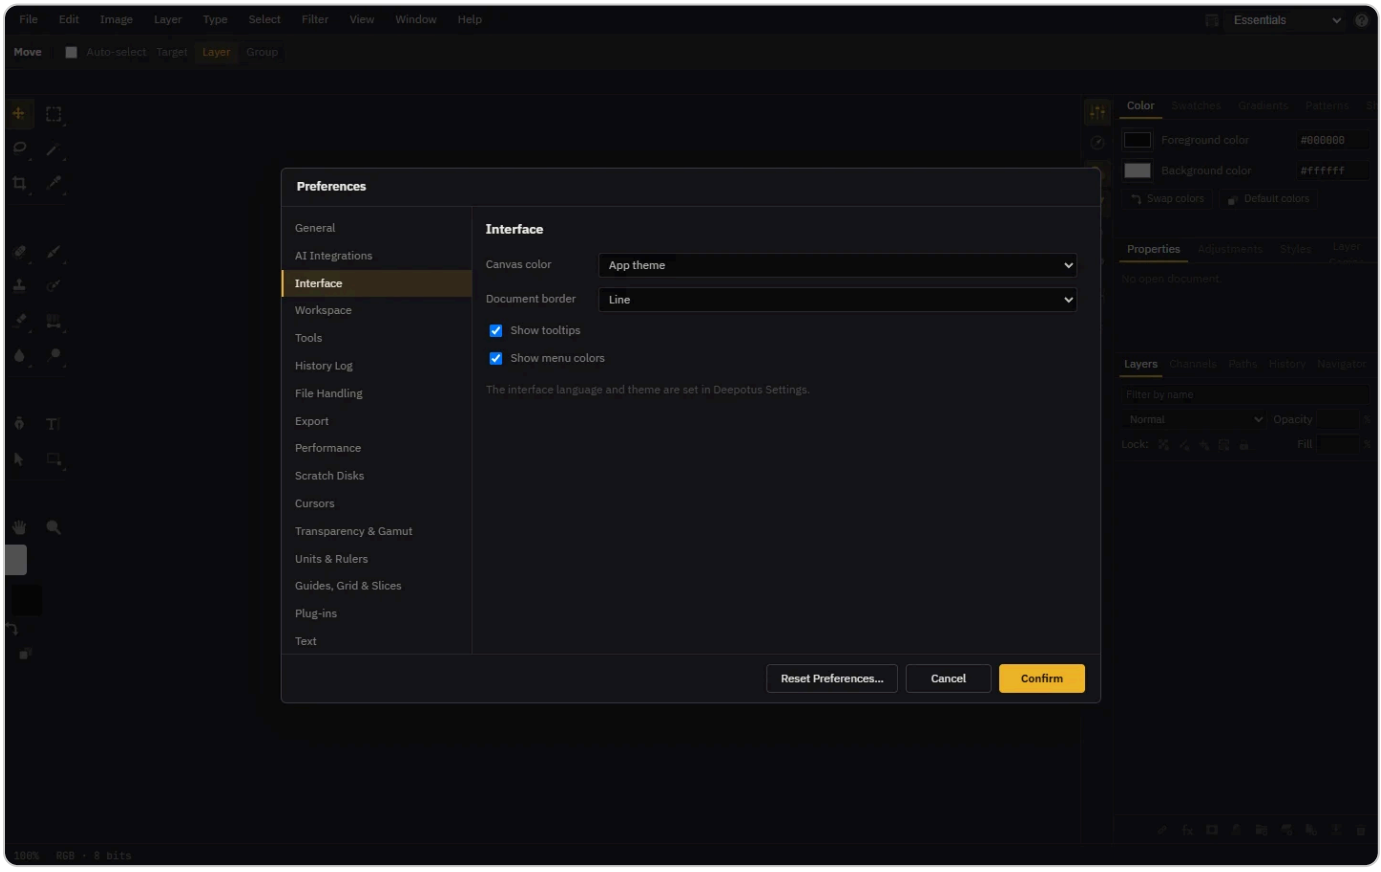

1. One section per menu entry

Edit › Preferences: eighteen sections, kept in the Deepotus data folder.

Edit › **Preferences** opens eighteen sections, saved in the Deepotus data folder: interface (canvas colour, document border, tooltips), tools, file handling and quick export, cursors, transparency, units and rulers, guides and grid, text. Some sections simply say what Deepotus handles in their place (performance, scratch disks…). The reset button restores all defaults, after confirmation.

Keyboard shortcuts

Change a command's shortcut or a tool's letter; a conflict is named and removed from the other command. Summarize… downloads the list as an HTML page.

Menus and presets

Hide or colour menu entries. The preset manager organises brushes, shapes and patterns; they export and re-import as .json.

Touch screen

Modifier Keys shows `Shift`, `Ctrl` and `Alt` on screen: a click applies to the next gesture, a double click holds the key.

Good to know

Workspaces, preferences, shortcuts and personal swatches are kept by Deepotus in its data folder: they follow you through a transfer to another PC.

IMAGES AND DRAWING

Automate and export (Photolab)

You will be able to record a series of Photolab gestures to replay it in one click, apply it to a whole batch of images, and output a document as more than a single image: one file per layer, one file per data set. Free, computed on your PC.

Record an action

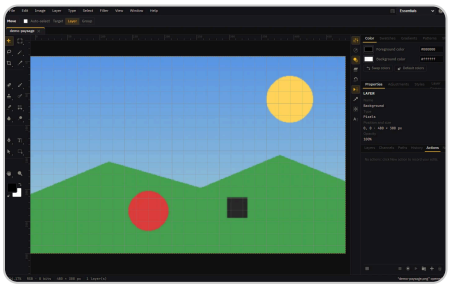

1. A new action

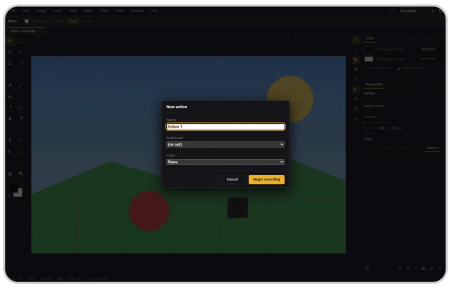

2. Name it, then record

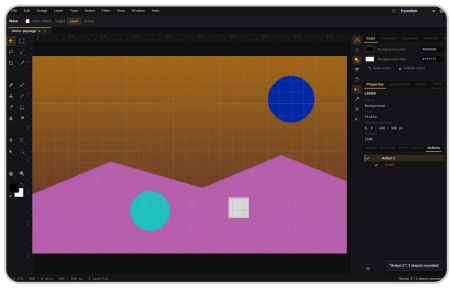

3. Every gesture becomes a step

Actions panel: New action, then every gesture you make on the image becomes a step you can replay on other images.

1. Open the **Actions** tab, in the Layers group, or Window › Actions.
2. Click **New action** at the bottom of the panel: give it a name, an action set and a colour if you like, then **Begin recording**.
3. Work as usual: every adjustment, filter or command applied to the document becomes a step under the action.
4. **Stop** ends the recording. To replay, pick the action then **Play selection**: it applies to the open document, one history step per gesture.

A tick switches a step off without erasing it; the box next to it opens its dialog during playback, to change a value on the way. The panel menu inserts a stop with a message, a conditional (“if the document is landscape, play this action”) or a menu command, and **Edit Step…** redoes a step with new values.

Process a batch of images

1. File › Automate › **Batch…**, or the same item in the Actions panel menu.
2. Pick the action, then the Library images to process (500 at most), and **Run**.

3. Each image is opened on its own, processed, then dropped into the Library, linked to its original image. Your open document is left alone.

Export differently

Layers to files

File › Export › **Layers to Files...** outputs each layer, or only the visible ones, as PNG or JPG into the Library.

Export as

The export dialog keeps your settings as named presets and warns when a format loses something: flattened layers, lossy compression, resolution dropped.

Variables and data sets

Image › Variables › **Define...** makes a layer variable (shown or not, text replaced). A data set gives each variable a value; Image › **Apply Data Set...** shows it.

One file per set

File › Import › Variable Data Sets... reads a CSV table; File › Export › **Data Sets as Files...** outputs one image per set, for example one card per name.

Good to know

Your actions stay from one session to the next. **Save Actions...** and **Load Actions...**, in the panel menu, exchange .atn action files through the Library: a step with no equivalent arrives switched off.

Scripts and droplets are not supported. An action only plays commands Photolab can do: before each playback its steps are checked, and a refused step stops it with a message.

IMAGES AND DRAWING

Draw vectors (Vectorlab)

You will be able to draw a poster, a logo or a sheet in vectors in Vectorlab: shapes, pen, text, appearance, layers, symbols, then send it to other screens. Free, everything is computed on your PC; only AI illustration is paid, price shown first.

Open Vectorlab

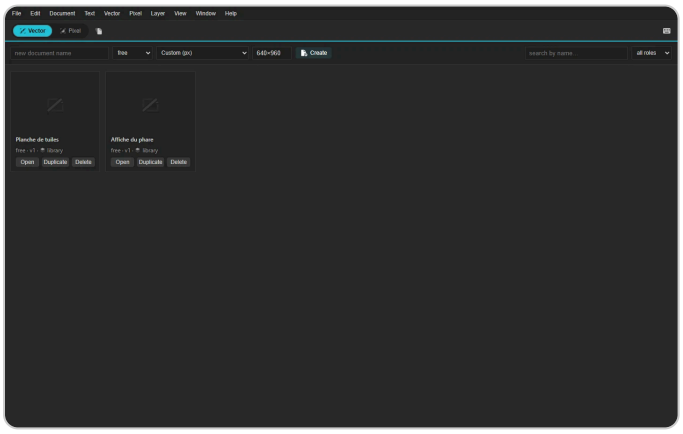

1. Open a document

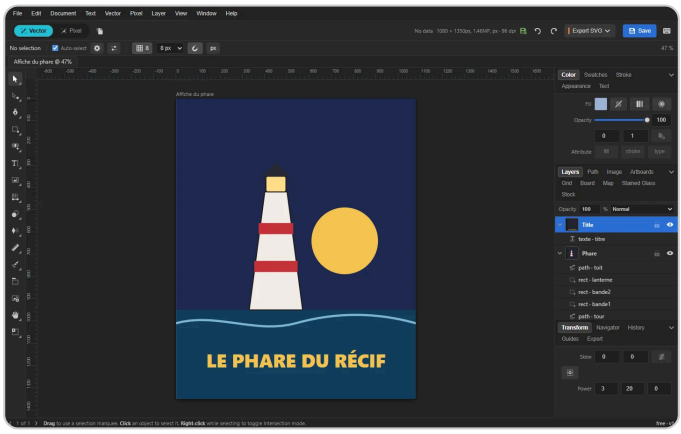

2. Two personas: Vector and Pixel 3. Tools, grouped by family

The home lists your documents: one click on a card opens it; New creates one.

1. Open **Vectorlab** in the app bar. The home lists your documents: one click on a card opens it. Pick a format (square, 16:9, A4, playing card…) then **Create** for a new document.
2. A document also comes from elsewhere: **Send to…** › Vectorlab from the Library or Photolab, “Edit this face in Vectorlab” from Card Forge, “Open in Vectorlab” from a Templates layout.
3. At the top left, two pills: **Vector** to draw, **Pixel** to retouch pixels (see the next chapter). Each has its own tools and panels.

The screen follows Affinity's conventions: ten menus at the top (**File**, **Edit**, Document, **Text**, **Vector**, Pixel, **Layer**, **View**, **Window**, **Help**), a context bar showing the selection and the current tool's settings, tools on the left, the tabbed panel stack on the right and a status bar that recalls the tool's gestures. The document tab shows “name @ zoom” and a star when it is not saved.

Pick a tool

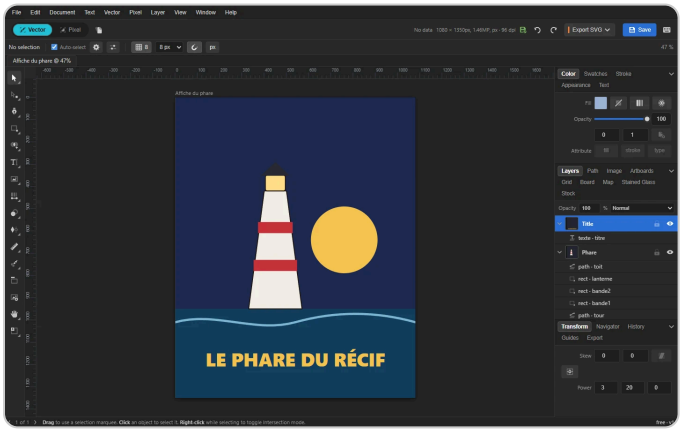

1. The Shapes family

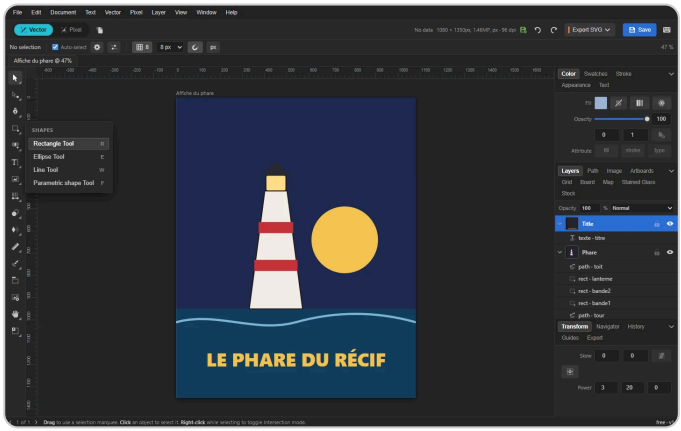

2. Rectangle, ellipse, line, parametric shape

The corner triangle or a long press opens the family; the tool's letter picks it from the keyboard.

Tools are grouped in families: only the current tool of each family is visible. The corner triangle or a long press opens the family, right-click opens the tool's settings, the letter picks it from the keyboard.

Select (V)

The marquee only takes objects fully inside (`Alt`: those it touches). `Shift` constrains, `Ctrl` resizes from the centre, `Alt`+drag copies, `Ctrl`+click reaches an object inside a group, `Tab` goes to the next one.

Pen (P) and Node (N)

Click = sharp node, drag = smooth node, click the first node to close. Four modes: Pen, Smart, Polygon, Line. The Node tool pulls handles, bends a segment, adds (double-click) or deletes nodes.

Shapes (R, E, L, F)

Rectangle, ellipse, line and parametric shapes (polygon, star, cog, arrow, donut, spiral) adjustable by handles. A single click places a shape; the Corner tool (C) rounds a corner.

Freehand

Smoothed pencil (B), vector brush (J) with flat, tapered or calligraphic profile, eraser (W), knife (X) and shape builder (S) that combines or removes areas by clicks.

Text (T)

Drag = text sized to the height you pull; click on a curve = text on that path; double-click to edit in place. The text frame flows text and flags overflow in red.

Eyedropper, measure, view

Eyedropper (I) with loupe (`Ctrl`: the whole style), measure (M), hand (H), zoom (Z, `Alt` to zoom out). `Ctrl`+`0` fits the page, `Ctrl`+`1` shows 100%.

Set the appearance

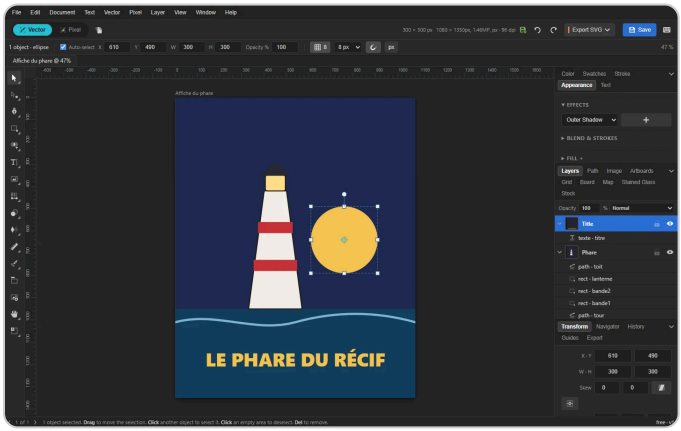

1. One click selects the sun

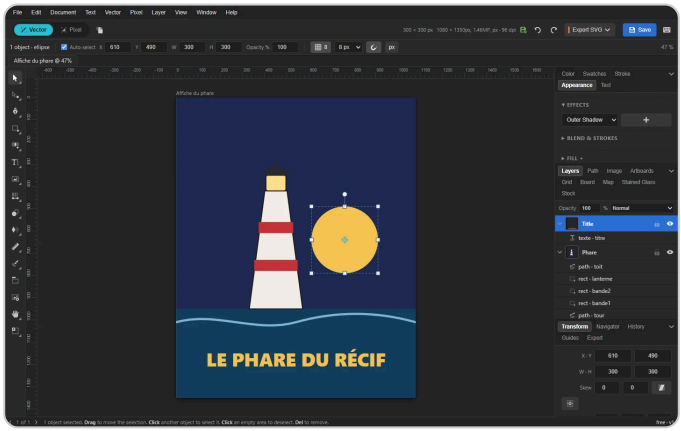

2. Fill, stroke, gradients, effects

Select an object: the context bar and the Appearance tab show its fill, stroke and opacity.

1. Select an object. The context bar gives X, Y, W, H and opacity; fields accept calculations (100/3).
2. The **Appearance** tab sets fill and stroke from the swatches, no fill or no stroke, width, opacity. **Stroke** and **Transform** have their own tabs (skew, movable pivot).
3. **Appearance +** goes further: linear, radial and conical gradients, transparency, patterns (hatching, dots, checkerboard), multiple strokes, six effects, 16 blend modes, global colours, harmonies, object styles, paste inside a shape (clipping mask), text in a frame or on a path.

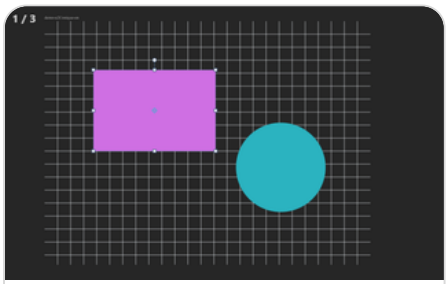

Conical

The color spins around the middle of your shape, like the spokes of a wheel.

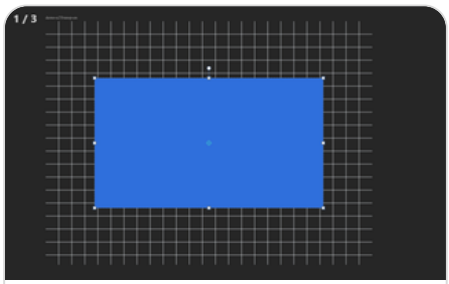

Transparency

Your shape slowly fades from one side to the other, as if it vanished into fog.

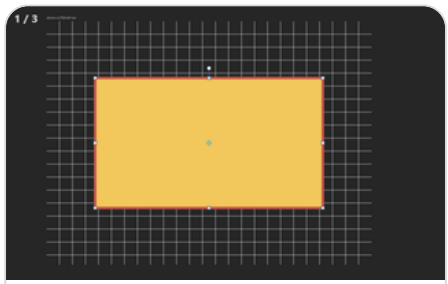

Pattern

Your shape fills up with little repeated stripes, like striped fabric.

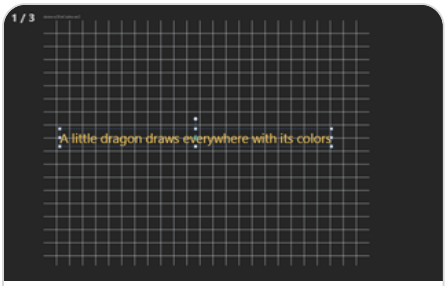

Frame
Your long line of text fits into a box and wraps by itself, like in a book.

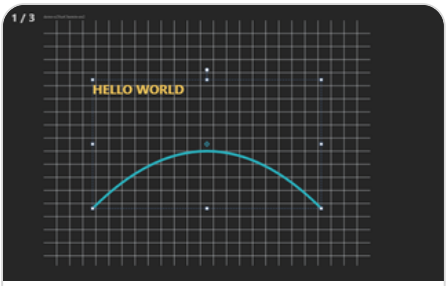

On path
Your words lie down on the line you drew and follow it wherever it goes.

An offset stroke and a text on a path stay **live**: they recompute when their source changes, until you detach them.

Layers and symbols

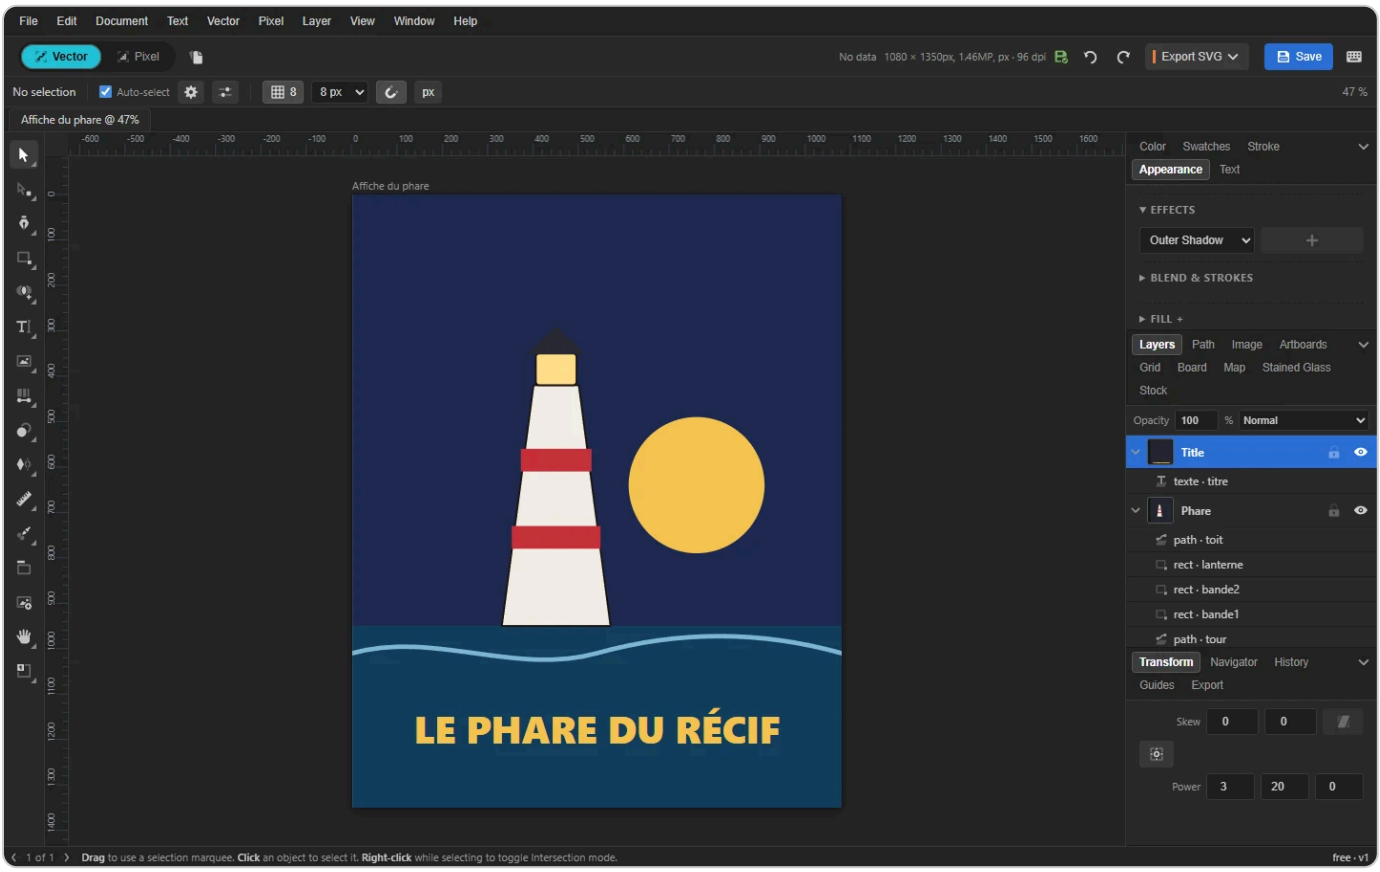

1. Background, Lighthouse, Title: three layers

Each layer unfolds into a tree of its objects, with thumbnail, lock, eye, opacity and blend mode.

The **Layers** tab shows the layer's opacity and blend mode at the top, then each layer with its thumbnail, lock and eye; a triangle unfolds it into a tree of its objects, clips, masks and effects. The **Layer** menu creates, renames, groups (`Ctrl` + `G`), locks, hides and orders.

A **symbol** is created with Vector › **Create Symbol** ; its instances follow every change. Double-click an instance to edit it in place, then **Done** or **Discard** .

Build up layers with AI

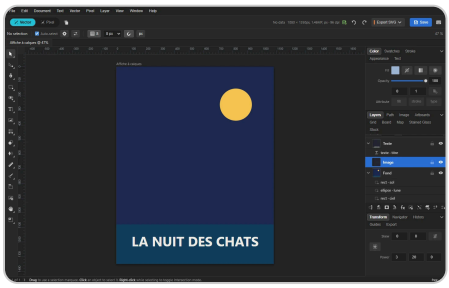

1. AI image on this layer

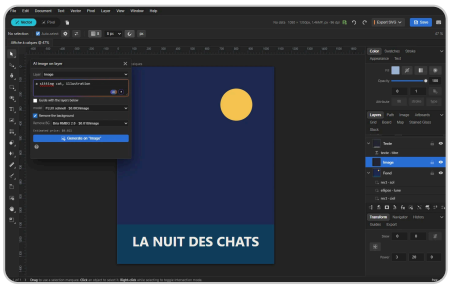

2. The target layer 3. The total price: image and background removal

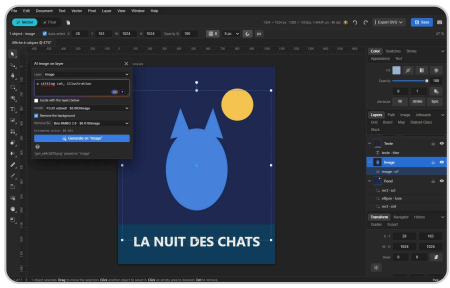

4. Placed on the Image layer, below the title

The layer button opens the panel: target layer, model, background removal, total price; the image lands on that layer, below the title.

1. Prepare your layers: for instance a background, an empty layer for the image, a text layer on top.
2. Select the image layer and click **AI image on this layer** at the foot of the Layers panel. The panel names the target layer and the model, and offers **Remove the background** with its removal model.
3. The total price, image and background removal, shows beforehand. The "Generate on ..." button places the image on that layer, below the text. Repeat on another layer for a second image.

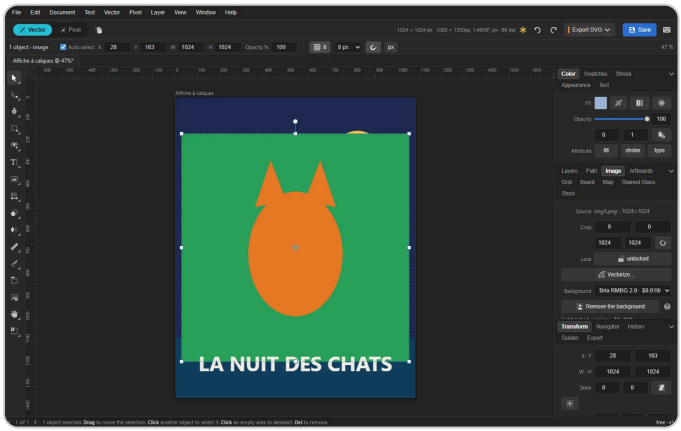

1. The removal model and its price 2. Remove the background

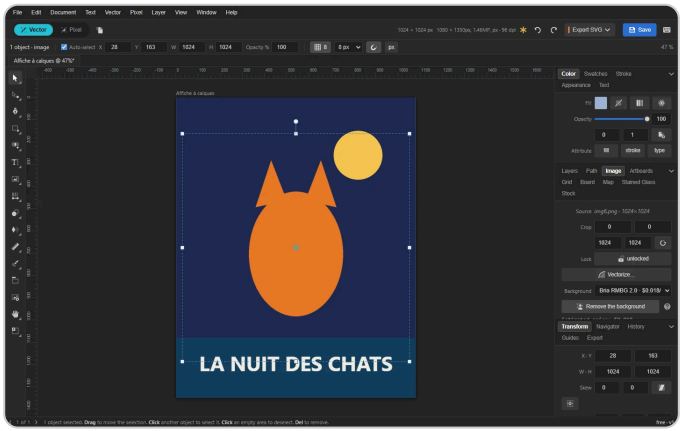

3. Background removed; Ctrl+Z brings back the previous image

Image tab: pick the background removal model, its price shows; Ctrl+Z brings back the previous image.

An image already placed is cut out from the **Image** tab: pick the model next to Background, its price shows, then **Remove the background** . `Ctrl` + `Z` brings back the previous image.

The Vector menu

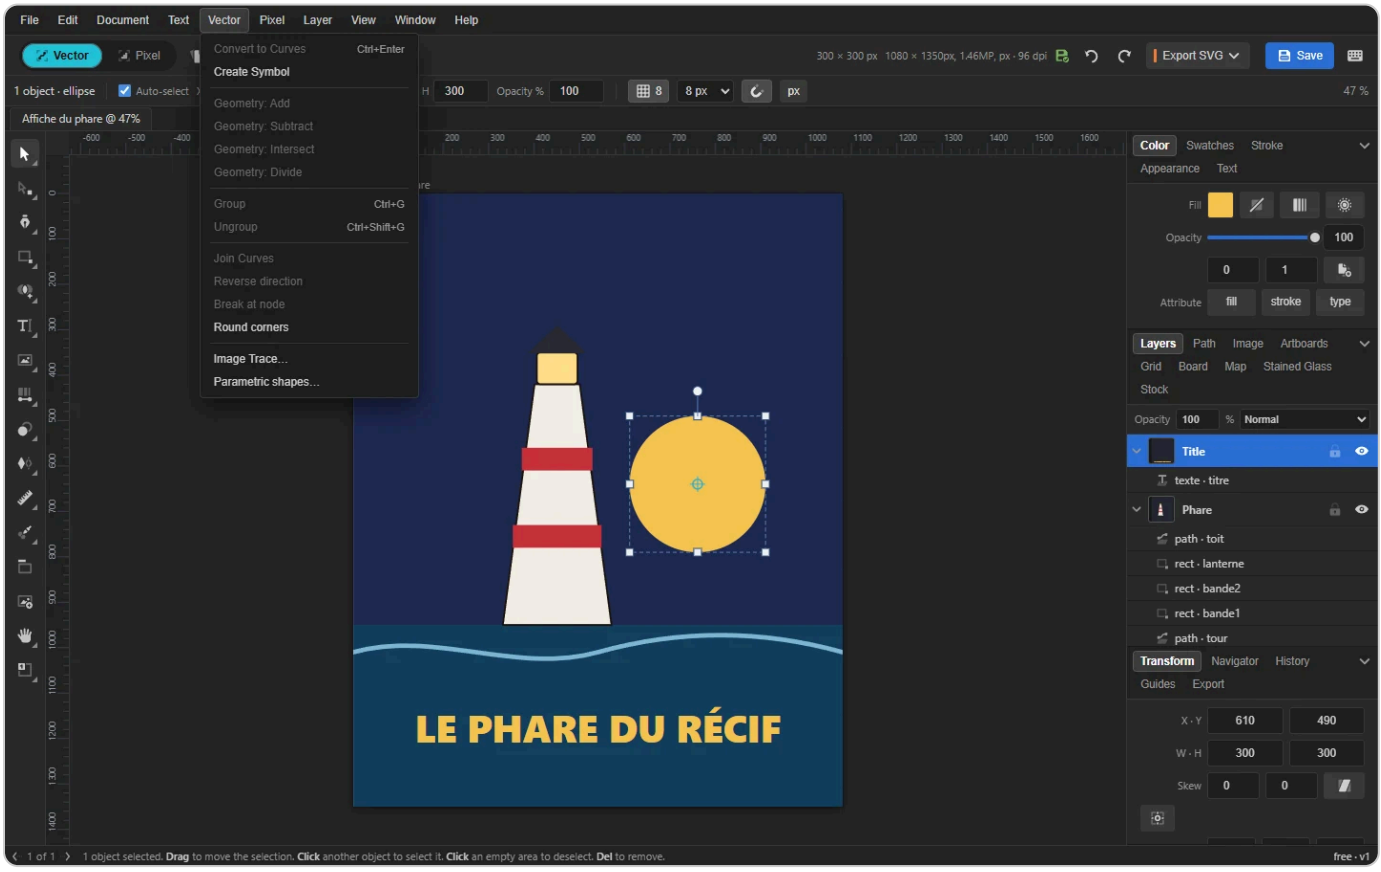

1. Combine shapes 2. Vectorise an image, locally

Geometry (union, subtract, intersect, divide), curves, symbols and image trace.

Geometry

Add, subtract, intersect and divide shapes; join, reverse, break at node; **Round corners**.

Convert to curves

A text becomes paths (`Ctrl`+`Enter`), as one or one path per glyph.

Image trace

Image Trace... vectorises a placed image (colours, smoothing, threshold), with preview, on your PC.

The document

The Document menu holds **Setup...** (size, resolution, unit, background), **Artboards...** exported separately, **Margins and bleed...** (guides never exported), the **Document grid...** (square, iso, triangular, hexagonal, snapping), **Board and terrains...** to paint tiles (K) and the **Real map...**. View › **Snap to objects** snaps to edges, centres and even gaps.

File › **Place image...** or **Place from Library...** places an image; the Stock tab shows the Library with search, one click places the image. An image can be cropped, locked and also receives effects.

Good to know

The document is saved as a draft every 30 seconds: when you reopen it, a window offers to restore it or to start from the server. History keeps 1000 steps; Edit › **Add Snapshot** marks a state to come back to.

Hover an option for a second, or click the "?" next to it, for its animated card. `F1` lists every shortcut; Edit › Settings (`Ctrl`+`,`) sets tooltips, grid and snapping.

AI illustration (AI family, or File › New image processing) is paid: the price is shown first, within your caps. Everything else is free.

Shortcuts

Gesture	Shortcut
Undo, redo	`Ctrl`+`Z`, `Ctrl`+`Shift`+`Z`
Duplicate, group, ungroup	`Ctrl`+`D`, `Ctrl`+`G`, `Ctrl`+`Shift`+`G`
Fit page, 100%, grid	`Ctrl`+`0`, `Ctrl`+`1`, `G`
Vectorlab home	`Ctrl`+`Alt`+`H`
All shortcuts	`F1`

IMAGES AND DRAWING

Pixel art, export, maps and 3D objects

You will be able to retouch pixels and make pixel art in Vectorlab, export your drawings in the right format, build a map from real data and turn a drawing into parts to 3D-print. Everything is free and computed on your PC.

The Pixel persona

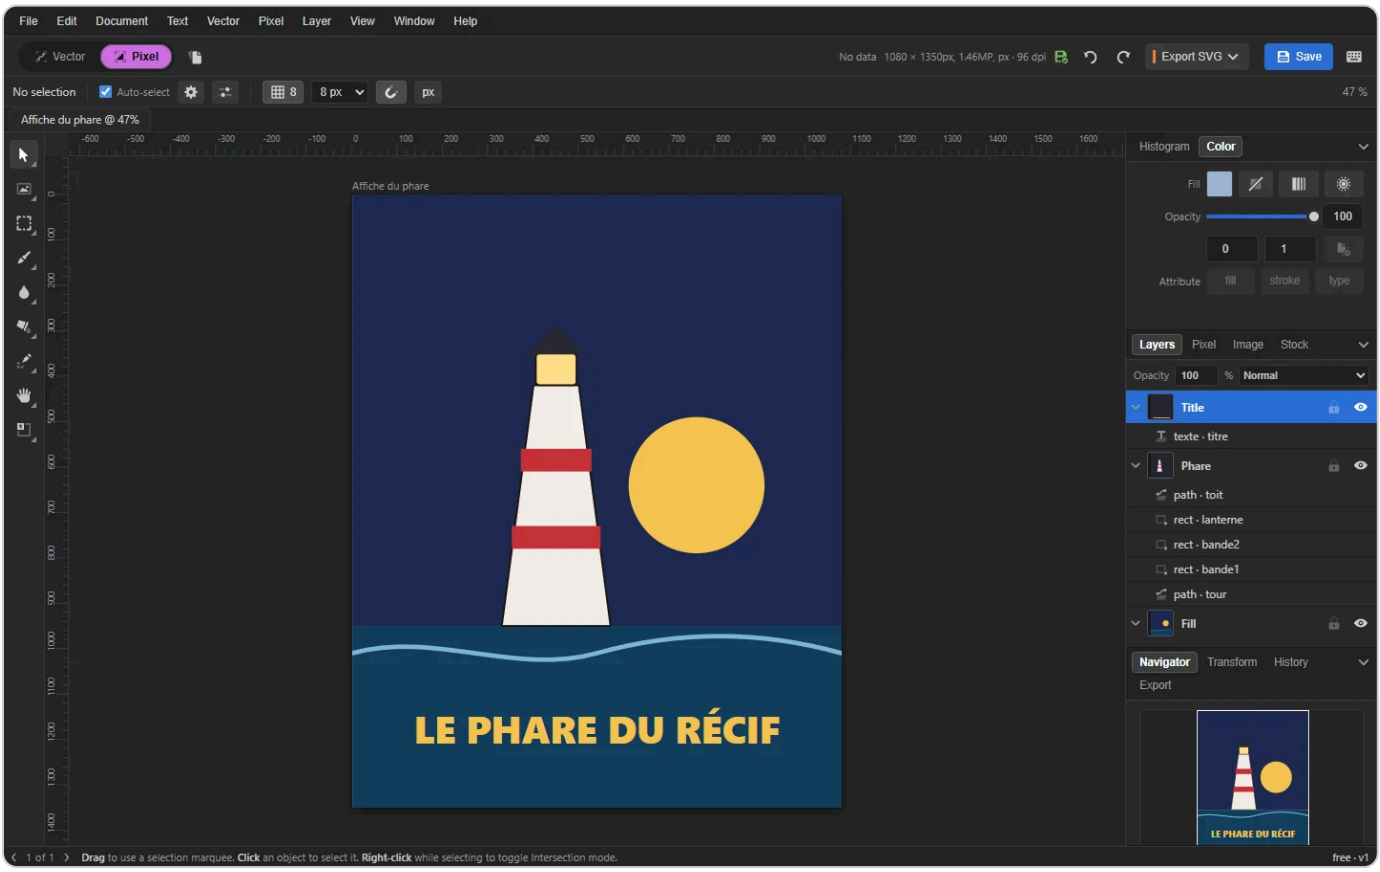

1. Pixel pill 2. The pixel tools

Same document, other tools: pixel selections, brush, eraser, bucket, pixel pencil, Pixel-art and Histogram panels.

1. Click the **Pixel** pill (or View › **Pixel Persona**). The same document keeps its layers; tools and panels change. From the Game Assets hub, the Pixel card opens this persona directly.
2. To retouch a placed image, pick it then Pixel › **Edit pixels**.
3. Select pixels: rectangle (M), lasso (L), magic wand (W) or by colour; the Pixel menu grows, contracts or inverts them (`Ctrl` + `I`).
4. Paint: brush (B) and eraser (E) with radius, hardness and round or square shape, paint bucket (G) contiguous or global, clone stamp (C), blur, dodge and burn. Right-click = secondary colour, `X` swaps them, `Alt` +click picks.

Pixel › **Adjustments and filters…** gives levels, curves, hue and saturation, black and white, threshold and blur. Every retouch undoes with `Ctrl`+`Z`, like the rest of the document. The Histogram tab shows the four channels of the edited image.

Make pixel art

The Pixel-art panel sets tile size, symmetry, palette (the same as Sprite Lab and Tile Lab), quantization and the pixel grid. The pixel pencil (K) draws without double pixels in diagonals; `Shift`+click draws a straight segment from the last point.

Trace a reference

Set as reference fades and locks an image with its grid; a pixel layer sits on top to draw cell by cell, the eyedropper takes a cell's exact, average or dominant colour.

Pixelate

Pixelate turns an image or a vector selection into big pixels. Rasterize adds palette and dithering (ordered or Floyd-Steinberg).

Tiles and seams

The 3×3 preview scores a tile's seams; the 2:1 iso tile limits painting to the diamond, with symmetry along its edges.

Animation

The timeline shows frames: play (`Enter`), loop, FPS, red (previous) and blue (next) onion skin. Export as a PNG + JSON sheet or a strip.

Finishing

Dark outline sets the drawing off the background, **Sharpen** boosts contrast; PNG export scales from ×1 to ×16 without blur.

To the other labs

Pixel › **Pixel art to Tilelab…** sends the drawing to Tile Lab; the Pixel-art panel also sends it to Sprite Lab.

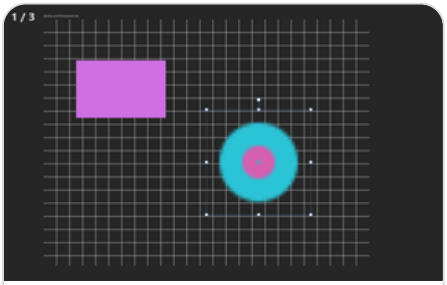

Set as reference
Your picture becomes a reference: it fades, stays still, and a grid of squares lies on it to guide you.

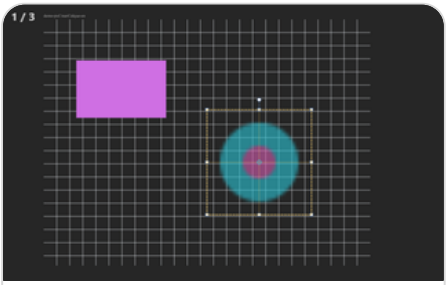

Create pixel layer
A see-through sheet lies on your reference picture: you draw on it, square by square.

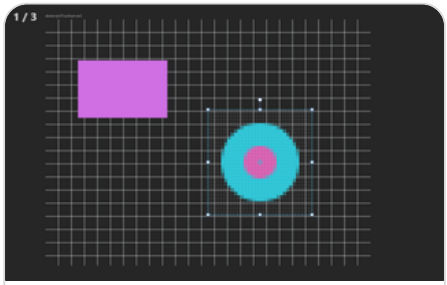

Pixelate image
Your picture is now made of big squares, like bricks: as many squares as the tile size.

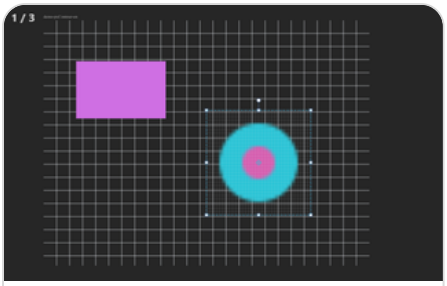

Dark outline
A dark line goes around your drawing so it stands out from the background.

Export

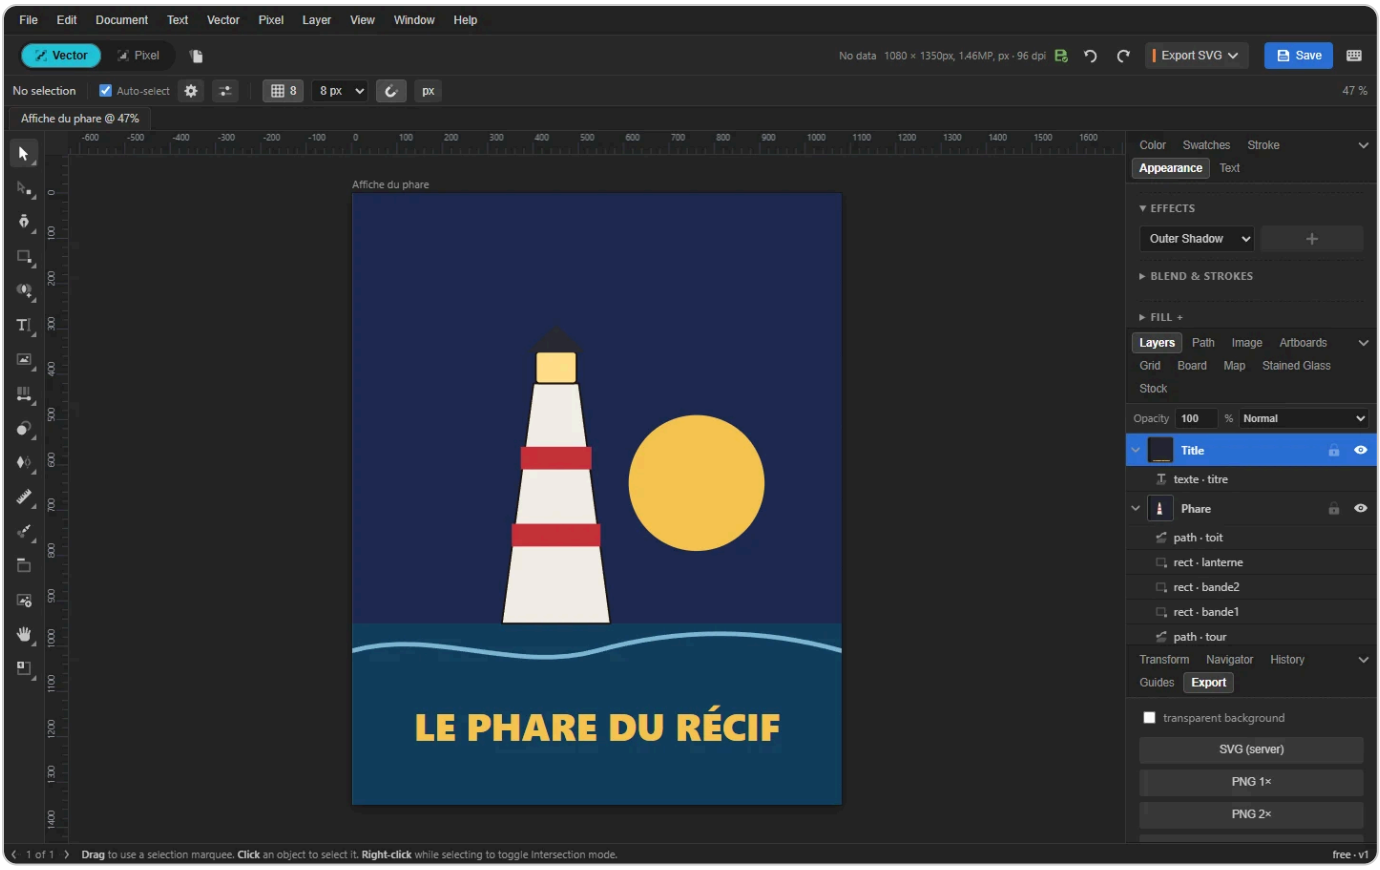

1. PNG, JPEG, WebP, SVG, PDF, DXF

Shared by both personas: formats, resolutions, slices and print presets.

1. Open the **Export** tab (File › **Export…** or Window › Export), shared by both personas.
2. Choose what goes out: the whole document, each artboard, each visible layer, each selected object, or the slices drawn with the Slice tool.
3. Choose the format (PNG, JPEG, WebP, SVG, PDF, DXF) and one or more resolutions (@1x, @2x…). File names build themselves; a batch export outputs every slice, resolution and format at once.

Vector PDF

By default: text as outlines, page at physical size, vector gradients and patterns, crop marks and bleed for the printer. A message says what had to be rasterized and why.

DXF for cutting

In millimetres, text as glyph outlines, symbols and groups included.

Quick exports

File › **Export SVG** (embedded images), **Export PNG 2×**, and **Send to Bible…** for a chapter.

A real map

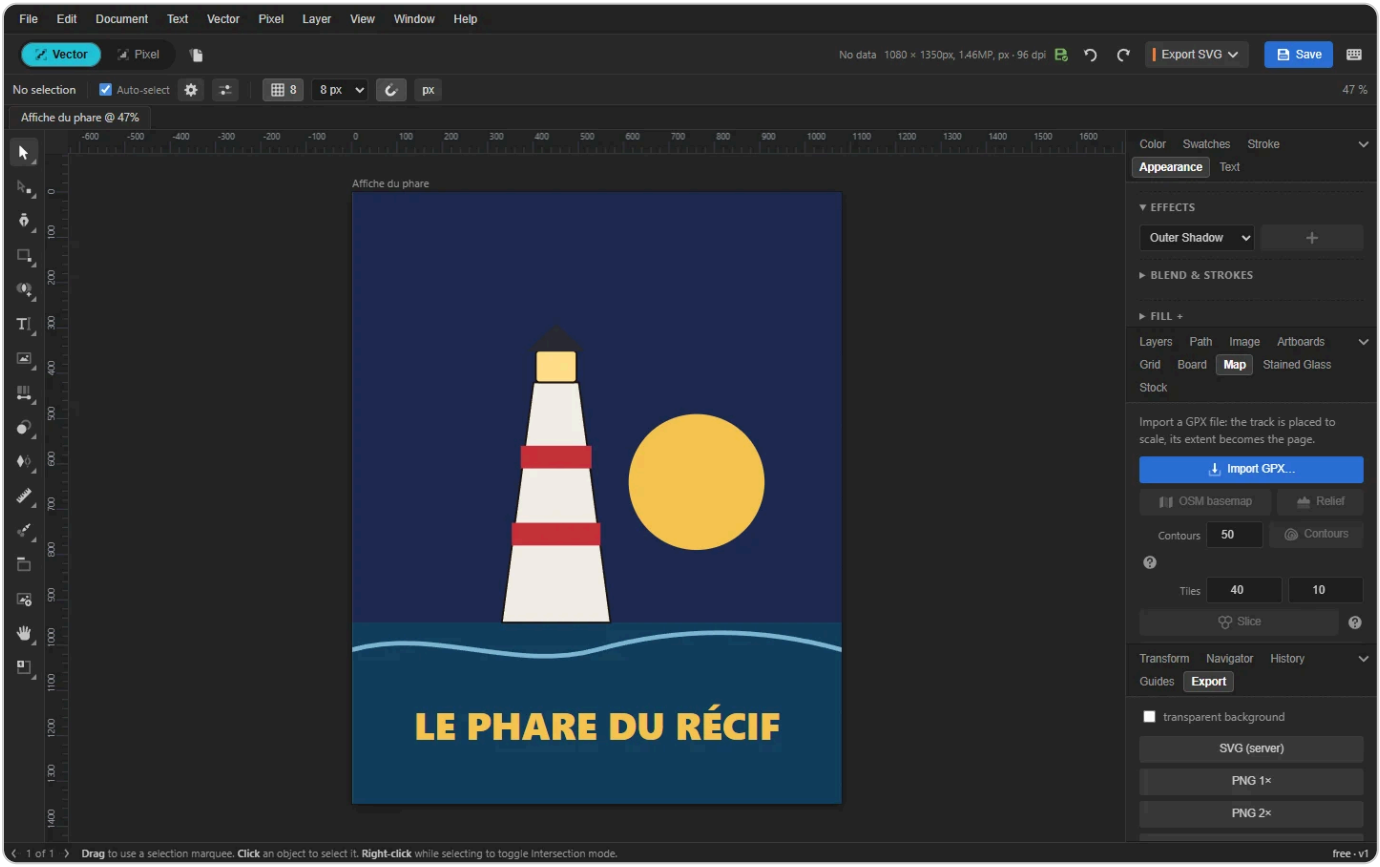

1. The Map panel

Document › Real map…: relief, contour lines, OpenStreetMap background and GPX tracks, no key, no cost.

Document › **Real map…** opens the Map panel: no key, no cost, it lays a shaded relief, draws contour lines as vector paths at the chosen interval, adds an OpenStreetMap background with its attribution, and slices the map into hexagonal tiles coloured by altitude. File › **Import GPX…** places a track at true scale, keeps its altitude (profile, length, climbs) and can turn it into a raised ribbon.

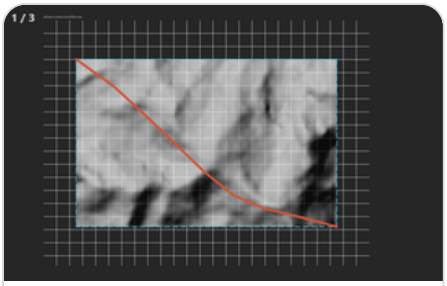

Contours
Lines join the places of the same height on your map, like the steps of a mountain.

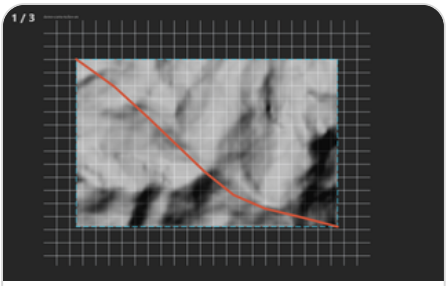

Slice
Your map turns into six-sided tiles, colored by height: water, meadow, forest, mountain.

Print in 3D

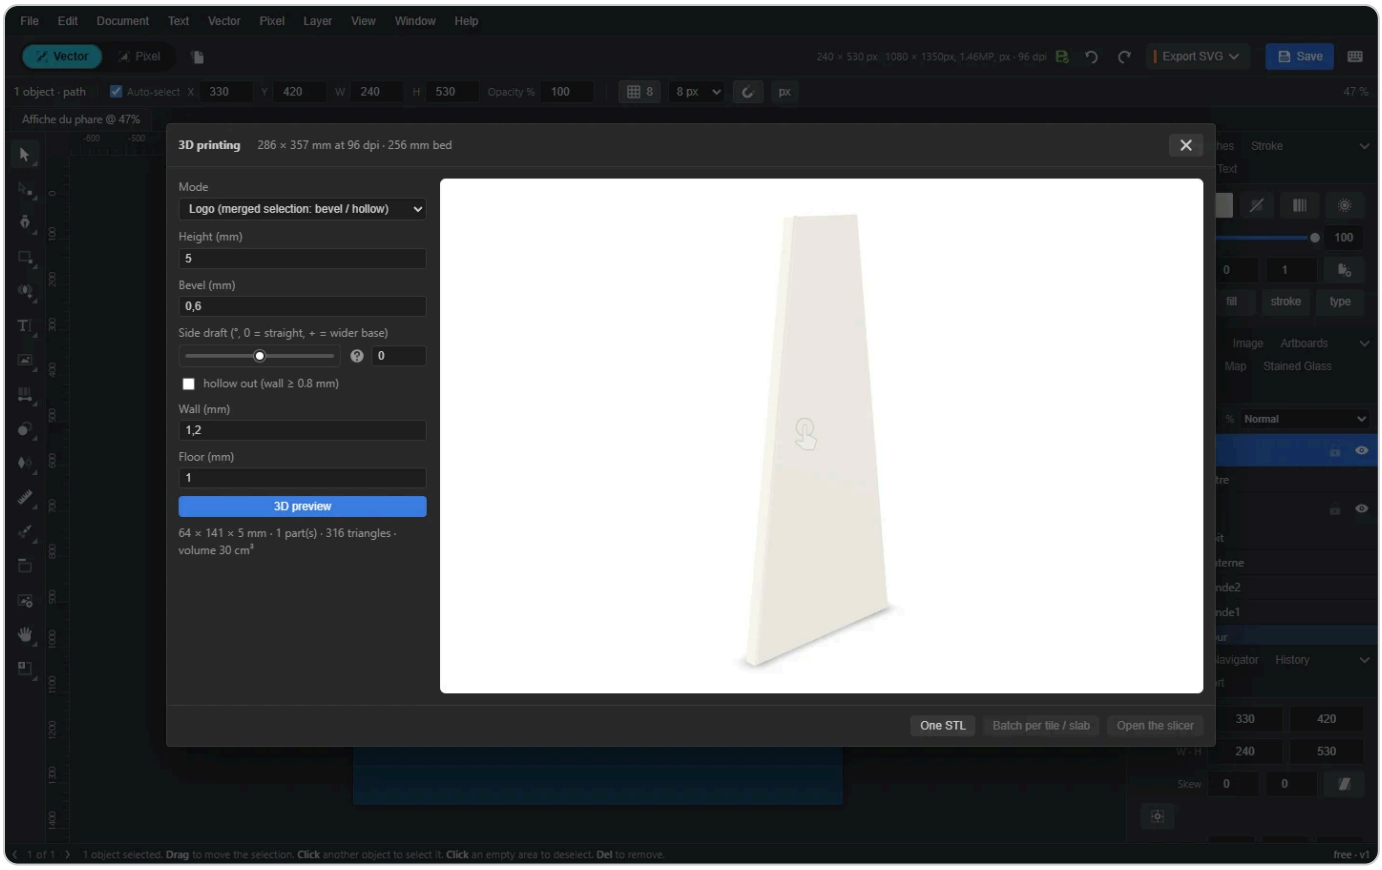

1. Layers, tiles, logo, relief or pixel art 2. The 3D preview, in part colours

File › 3D printing… turns the drawing into parts: 3D preview, then STL, 3MF or a batch per part.

1. File › **3D Print…** (or Text › 3D Logo…). Choose the mode: layers (one height per layer), board tiles with base, logo (bevel, hollowing with a minimum wall), map relief, or pixel art (one height per colour).

2. Set heights and options, then **3D preview**: the model turns in the viewer, each part in its colour, with dimensions, triangles and volume. A model wider than the bed is flagged.

3 **One STL** for a single part, or **Batch per tile / slab**: one STL per part, a coloured plateau.3mf and nomenclature.csv. **Open the slicer** hands over to your printing software.

In relief mode, numbered slabs get sockets under the base and separate keys (0.2 mm clearance) to fit together. Side draft tilts a logo's edges for demoulding.

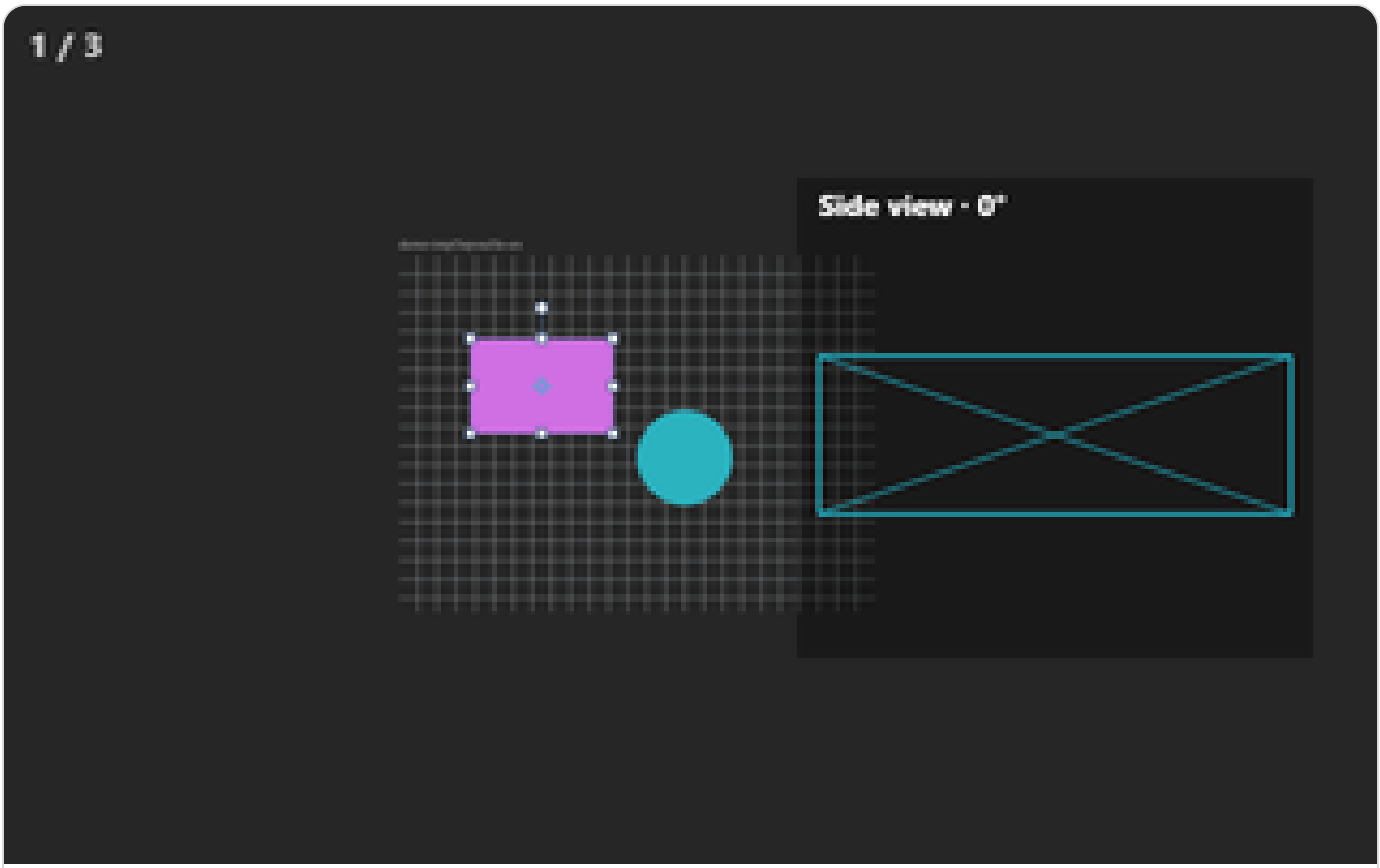

Side draft

The sides of your piece lean: the more you turn, the wider the bottom and the smaller the top.

Good to know

A text must be converted to curves to be printed: the dialog says so and skips it otherwise. To go further with a 3D model, see the Workbench and 3D printing.

GAMES AND 3D

Make game assets

You will be able to find your way around the Game Assets hub, turn an image into a 3D model and open the right 2D tool: Sprite Lab, Tile Lab, Vectorlab's Pixel persona or Card Forge. The 2D labs are free and run on your PC; 3D generation is paid, price shown first.

The Game Assets hub

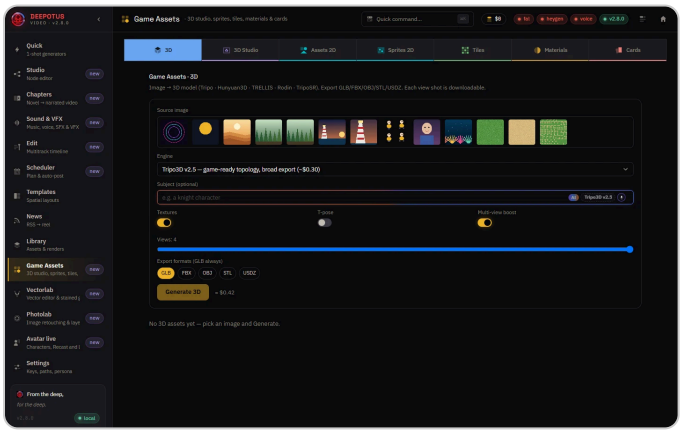

1. Image → 3D model: price shown first 2. The 2D Assets tab

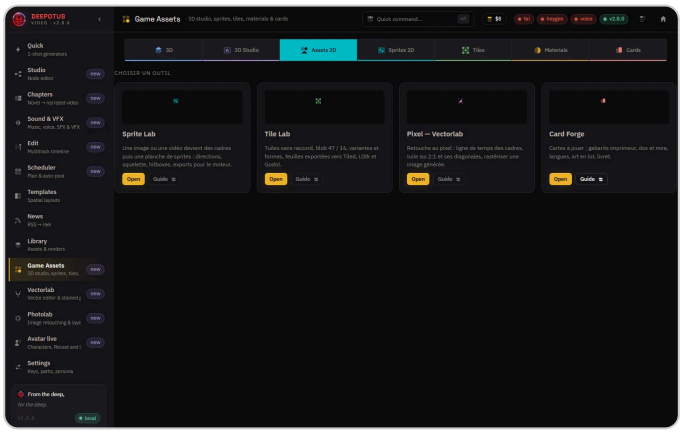

3. Four labs, each with Open and Guide

One tab per tool: 3D, 3D Studio, 2D Assets, Sprites, Tiles, Materials, Cards. 2D Assets gathers the four 2D labs.

Open `Game Assets` in the rail. One tab per tool: **3D** (image to 3D model), **3D Studio** (see Image to 3D), **Assets 2D**, **Sprites 2D** and **Tiles** (the labs opened inside the hub), **Materials** (Material Forge) and **Cards** (Card Forge).

An image becomes a 3D model

1. In the 3D tab, pick a source image from the Library: ideally a character or object cleanly cut out on a neutral background.
2. Pick the engine, from cheapest to most refined: TripoSR, Hunyuan3D, Tripo (default, good value), TRELLIS and Rodin. Only the fal.ai key is needed.
3. Options: textures, T-pose for a character to animate, multi-view (four views generated first for a more faithful model), and formats: GLB always, plus FBX, OBJ, STL or USDZ depending on the engine.
4. The total price is shown next to `Generate 3D`, before the run. The pipeline then shows each step live; the model turns in the viewer, every format and every view can be downloaded.

Models also land in the 3D tab of the Library; the Workbench then prepares them for printing.

Open a 2D tool

The `Assets 2D` tab gathers four cards, each with `Open` and `Guide`:

Sprite Lab

An image, a video or a sheet becomes an animated sprite sheet, with directions, skeleton, hitboxes and engine exports. See Animate sprites.

Tile Lab

Seamless tiles, blob 47 or 16 tilesets, variants, iso and hex shapes, painted maps, exports to Tiled, LDtk and Godot. See Paint tile scenery.

Pixel — Vectorlab

Opens Vectorlab straight in the Pixel persona: frame timeline, 2:1 iso tile, rasterize an image. See Pixel art.

Card Forge

Playing cards: printer templates, back and registration, languages, batch art, booklet. See Card Forge.

Shared palettes

Sprite Lab, Tile Lab and Vectorlab's Pixel persona share the same list of palettes: adaptive (N colours), PICO-8, Game Boy, NES, Sweetie 16, 1-bit, Grayscale 16 (sixteen greys) and Handheld 4 (four tones, handheld console style). A palette gives the same result in all three tools. Dithering (ordered or Floyd-Steinberg) softens gradients.

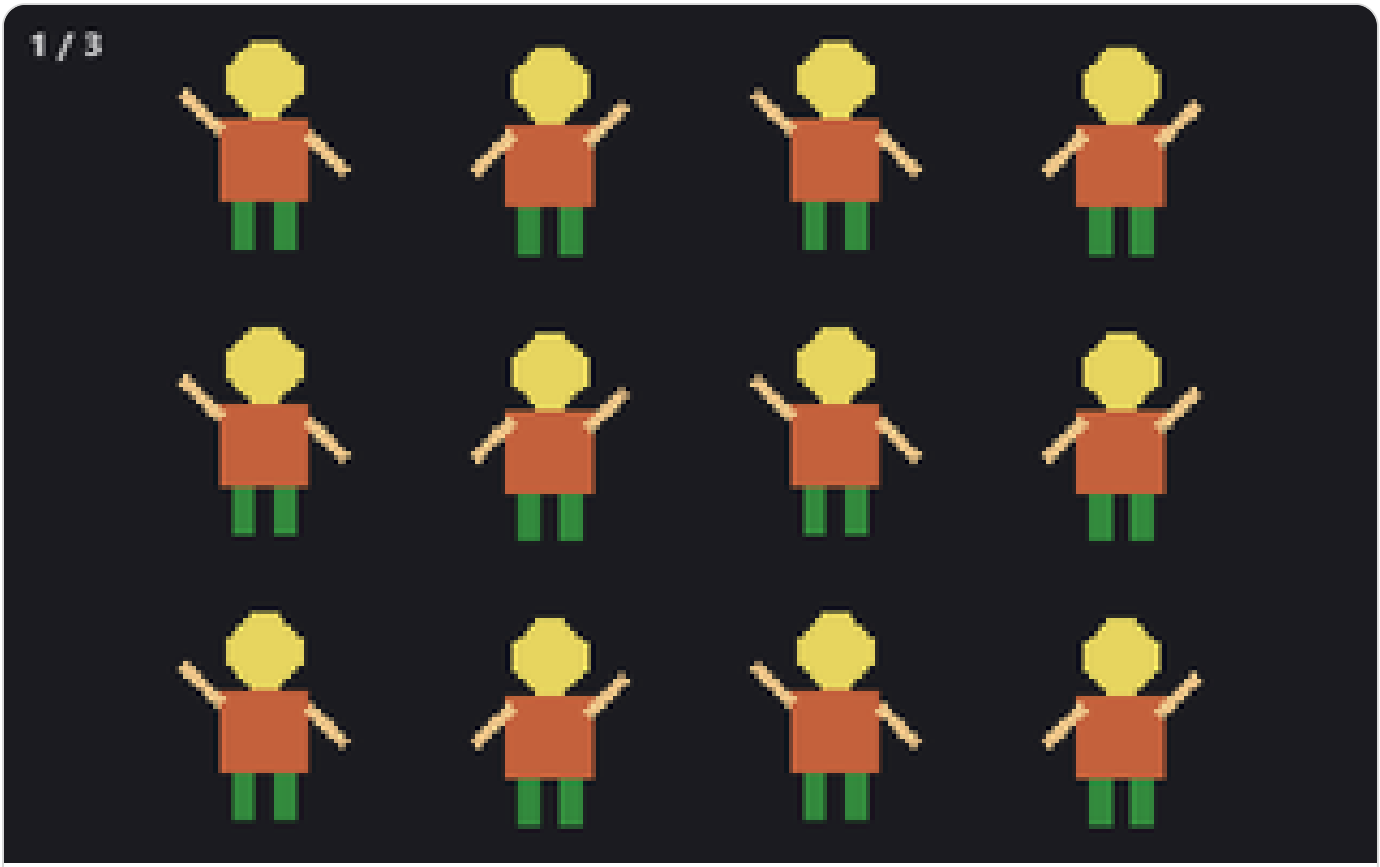

Pixel art (9b)

Your drawing becomes big squares and only a few colours, like in an old video game.

Good to know

Only 3D generation costs money: the price shown counts each extra view and format, and goes through your spending caps. The 2D labs are free.

3D rendering runs in your browser: keep the tab in front for smooth rotation. Files live in the Deepotus data folder and follow a transfer.

GAMES AND 3D

Animate sprites (Sprite Lab)

You will be able to build an animated sprite sheet in Sprite Lab: start from a shipped effect, a sheet, a video or a bible character, slice it, align it, pixelate it, then export it for your game engine. Almost everything is free and local; only animating an image and prompt generation are paid, price shown first.

The Sprite Lab screen

Open `Sprite Lab` from the Assets 2D tab of the Game Assets hub. Three columns: the **source** on the left (Get started, Image, Render, Video, Sheet, Bible, Prompt), the extracted **frames** and settings in the middle, the animated **preview** and exports on the right.

A sprite in one click

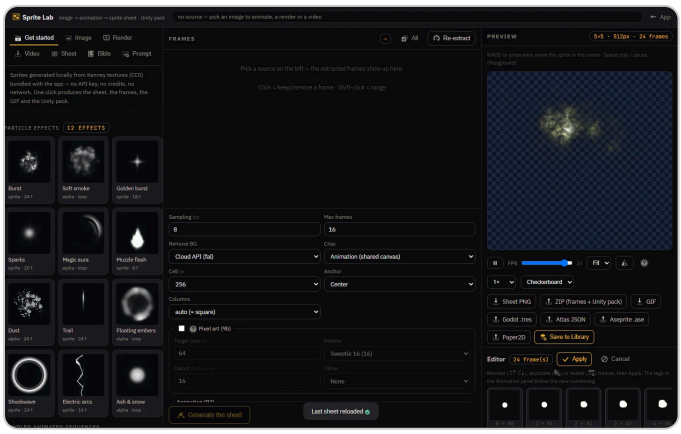

1. Pick an effect

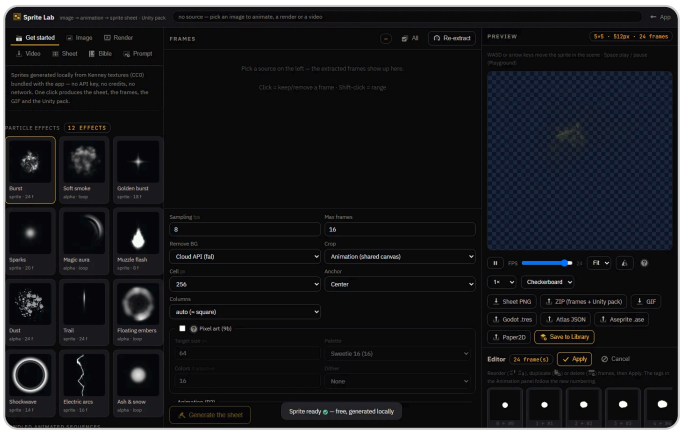

2. The preview plays the animation 3. Exports for Unity, Godot, Phaser, Aseprite, Unreal

The Start tab ships twelve particle effects and sequences: one click produces the sheet, the GIF and the exports, no key, no credit.

1. `Get started` tab: twelve particle effects (burst, smoke, sparks, aura…) and animated sequences, made from free textures shipped with the app.
2. One click produces the sheet, the frames, the GIF and the Unity pack, no key, no credit, no network.
3. The preview plays the animation: `Space` play / pause, speed 0.5×, 1× or 2×, plain backgrounds or scenes (Sky, Outdoor, Dungeon, Cave), and `W` `A` `S` `D` or the arrows walk the sprite. The mirror button flips it.

Mirror

Your character turns around and looks the other way, like in a mirror.

Start from an existing sheet

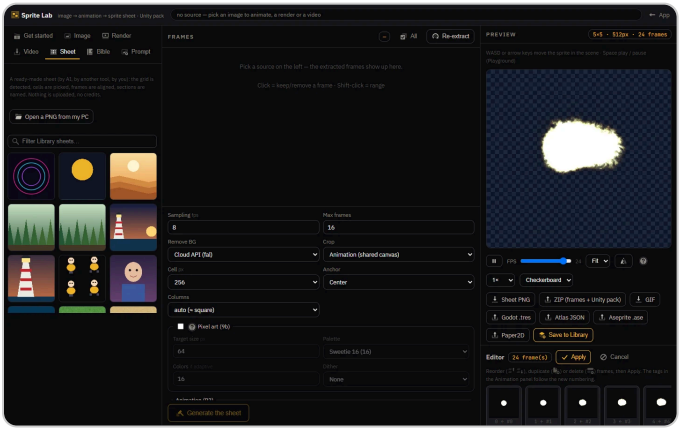

1. A sheet from the Library

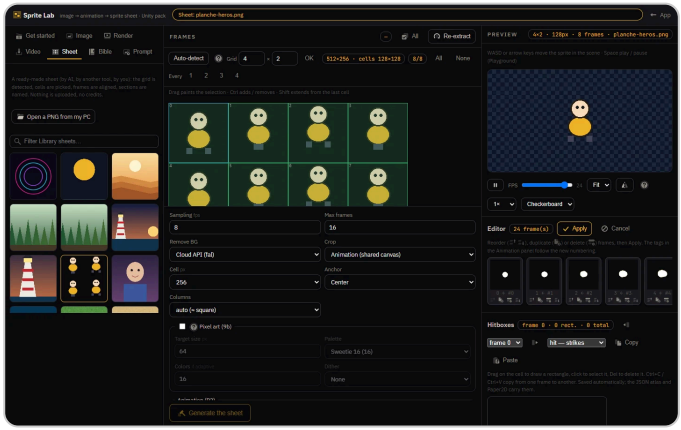

2. 4 × 2 cells detected automatically 3. The hero hops in the preview

Sheet tab: the grid is detected, cells are picked, frames are aligned; all local.

1. `Sheet` tab: pick a sheet from the Library, or `Open a PNG from my PC`. Nothing is sent, nothing is spent.
2. `Auto-detect` counts columns, rows and occupied cells from transparency; otherwise, set the columns × rows grid yourself.

3. Drag to paint the cell selection (`Ctrl` adds or removes, `Shift` extends), or **All**, **None**, one cell in 2, 3 or 4.
4. Align: **Feet** puts each character on the same floor, **Left / right** re-centres, **Both axes** does both; the arrows nudge the current frame by one pixel, **Reset** undoes.
5. Group ranges under a name (walk, jump…) as loop, ping-pong or reverse, then export the **Aligned sheet**, its JSON, or **Copy JSON**.

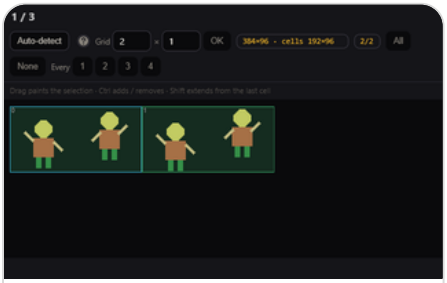

Auto-detect
The tool counts the pictures on your sheet by itself and draws a box around each one.

Feet
All your characters put their feet at the same height, as if standing on the same ground.

The other sources

Image, Render, Video
A Library image is animated on a plain background (paid); a render or a video from your PC is sliced into frames. Set sampling, the number of frames and background removal (cloud, local or chroma key).

Bible
4 directions (sheet) cuts a character's south, west, east, north views; **8 directions (3D model)** renders its model from eight angles. Free.

Prompt
Describe a sprite, add the persona's words in one click, then **Generate the images**: the cost is shown first, images are cut out locally and become the source.

Any source also accepts 1 to 64 Library images, taken in order as frames.

Set up the sheet

In the middle: cell (128, 256, 512 or native, no upscaling), anchor centred or at the feet, columns. **Pixel art (9b)** reduces to the target size with a shared palette and dithering. Animation: duration per frame

and tags (idle, run, jump…) written into every export. Post-processing: outline, drop shadow, orphan-pixel cleanup, edge smoothing. Then **Generate the sheet**.

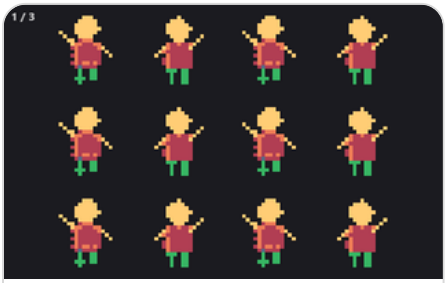

Outline
A line surrounds your character so it shows up on any background.

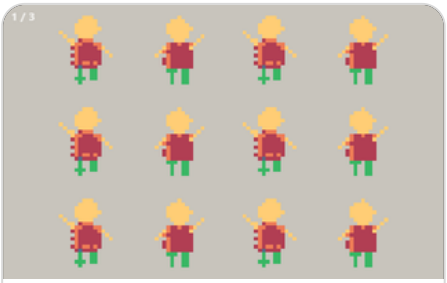

Shadow X / Y
A shadow appears behind your character, slightly offset, as if a lamp were lighting it.

Edit, hitboxes, skeleton

Editor
Reorder, duplicate or delete frames of a finished sheet, then **Apply**; tags follow. Nothing is redone or paid again.

Hitboxes
Draw per-frame hit (strike) or hurt (can be hit) rectangles, sixteen at most; adjust them, copy them from one frame to another (`Ctrl`+`C` / `Ctrl`+`V`). They go into the JSON atlas and Paper2D.

Spine skeleton
Draw bones then parts, name them, and **Write the rig** produces a Spine 3.8 skeleton with its parts.

Export

Under the preview: Sheet PNG, ZIP (frames and Unity pack, with all engine exports), GIF, Godot 4 resource (.tres), JSON atlas for Phaser and PixiJS, Aseprite file (.ase) and Unreal Paper2D sheet. Save to Library files the sheet in the Library, from where it can be sent to Vectorlab or Tile Lab.

Good to know

Paid: animating an image (Image tab), cloud background removal and prompt generation, always within your caps, price shown first. Get started, Sheet, Bible, Render, Video with local removal and every export are free.

GAMES AND 3D

Paint tiled levels (Tile Lab)

You will be able to build seamless tiles, auto-tiling tilesets, iso or hexagonal tiles in Tile Lab, and paint a small map ready for Tiled, LDtk or Godot. Everything is free and computed on your PC.

The Tile Lab screen

Open **Tile Lab** from the Assets 2D tab of the Game Assets hub. A bar of five modes: **Seamless**, **Tile sheet**, **Tileset**, **Shapes** and **Painter**. On the left, the source (Library images); on the right, the result. An image sent from the Library or Photolab (“Tile Lab — tile source”) arrives directly as the source.

A seamless tile

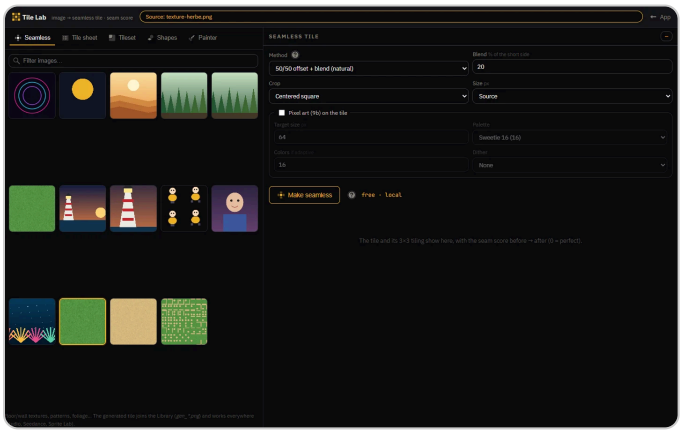

1. Make seamless: free, local

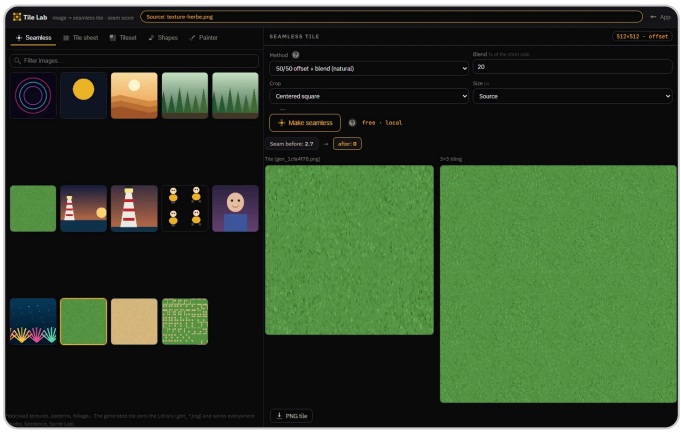

2. The tile and its 3×3 tiling

Seamless mode: an image becomes a tile that repeats without seams, with its 3×3 tiling and seam score.

1. Seamless mode: pick an image (floor, wall, foliage texture…).
2. Method: 50/50 offset with blend (natural) or 2×2 mirror (perfect, kaleidoscope effect); crop and size; optional pixel art with a shared palette.
3. **Make seamless** shows the tile, its 3×3 tiling and the seam score before and after (0 = perfect). The tile joins the Library.

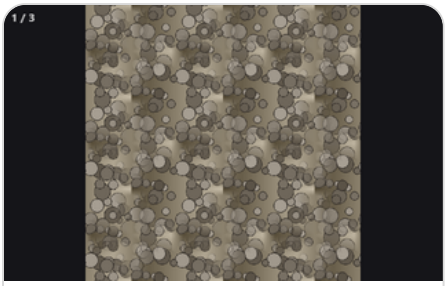

Method

Mirror folds your texture like a sheet folded in four: the edges always meet perfectly.

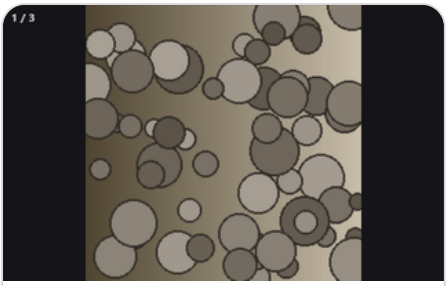

Make seamless

Your texture joins up: laid side by side many times, you can no longer see where each piece ends.

Build a tileset

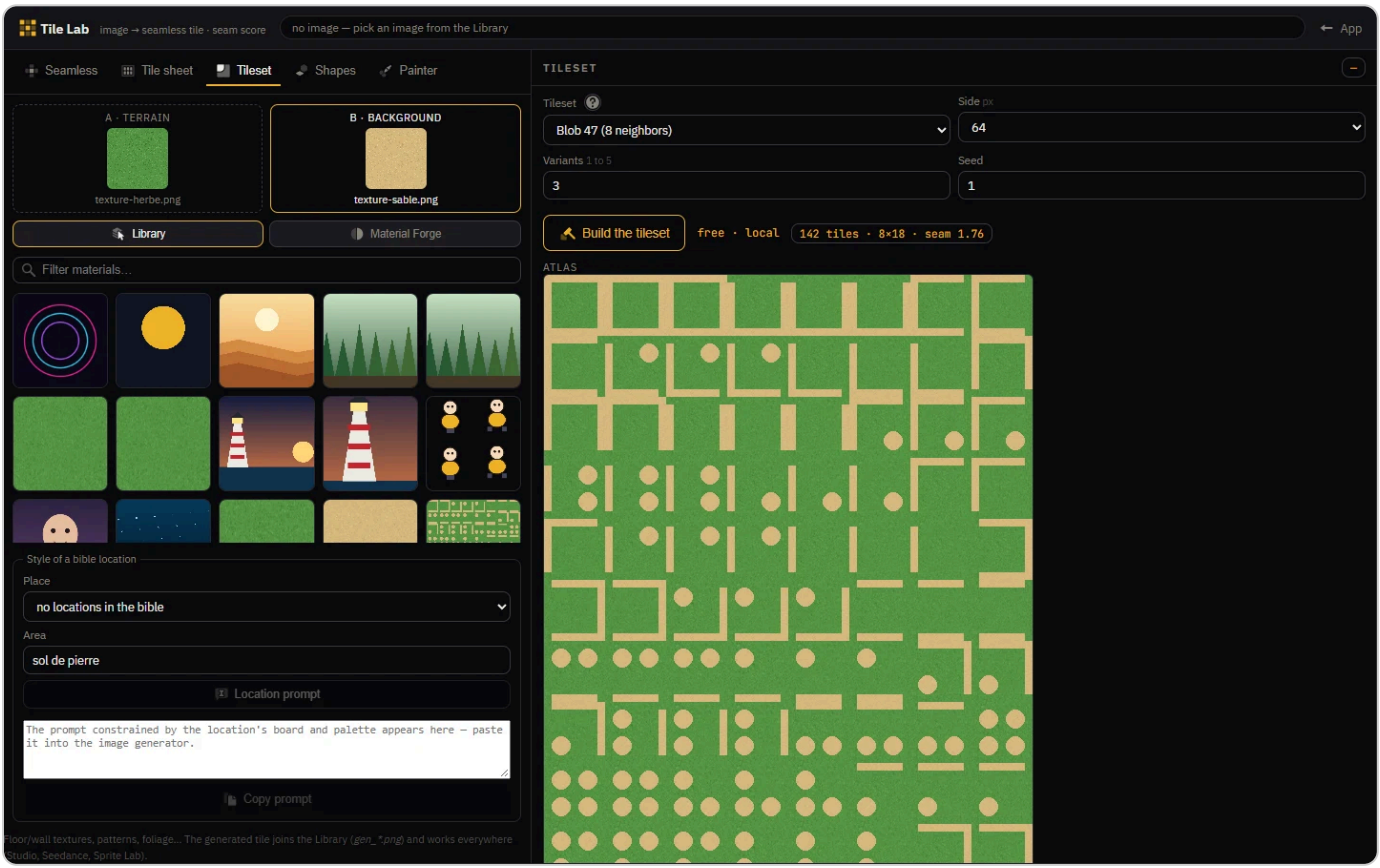

1. A: terrain, B: background 2. 142 tiles, seams measured

Tileset mode: a terrain material and a background one become a blob 47 or 16 set, auto-tiling and exportable.

1. Tileset mode: click an image for material A (terrain) then a second for B (background). The **Material Forge** toggle takes the base colour of a Material Forge material.
2. Choose the set: Blob 47 (eight neighbours, rounded corners) or Blob 16 (four edges, faster), the side in pixels, 1 to 5 variants per tile and the seed.

3. **Build the tileset** builds the atlas and measures how neighbours join; the atlas is copied to the Library.
4. **New preview** draws a random map where each cell gets the tile of its neighbours; **Measure** judges seams, visible repetition and lighting (ok / warning).
5. Export atlas.png then your engine's format: Tiled (.tsx), LDtk (.ldtk) or Godot (.tres), to be placed next to the atlas.

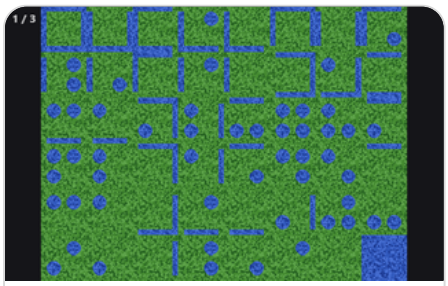

Tileset

You choose how many pieces to make: many for nicely rounded corners, or sixteen to go fast.

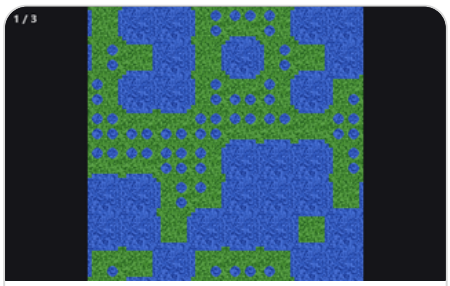

New preview

A new map draws itself at random with your tiles, to see if they go well together.

The “Style of a bible location” block composes a tile prompt bound by a location's sheet and palette: **Location prompt**, then **Copy prompt** to paste it into the image generator. Nothing is generated here.

Paint a map

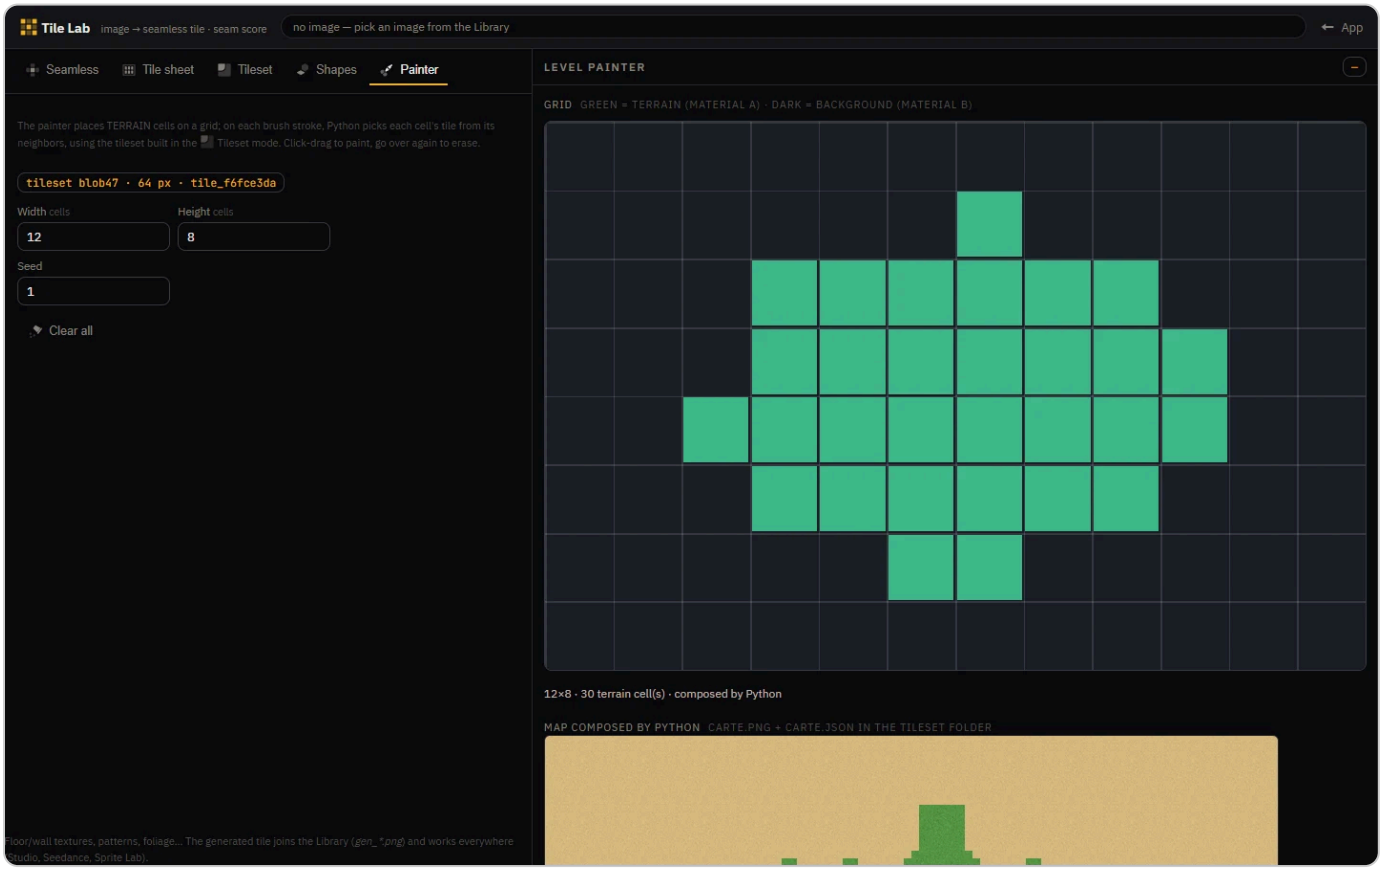

1. Click and drag to paint 2. The composed map, ready for the game

Painter mode: you paint terrain cell by cell, and the right tile of the set is placed from its neighbours.

Painter mode, once a tileset is built: set width, height and seed, then click and drag on the grid to paint terrain, go over again to erase. On each stroke, the right tile is placed from its neighbours, and the composed map shows below with carte.png and carte.json. **Clear all** empties the grid.

Shapes and tile sheets

Shapes

A material becomes a 2:1 iso diamond or flat-top hexagon tile that tiles without seams. **Build the tile** shows the tiling; export to Tiled and Godot.

Tile sheet

Open an existing sheet: **Detect** finds each tile by its transparency, the placement grid (square or 2:1 iso) tests seams, double-click = ×9 tiling.

Sheet exports

Aligned sheet and JSON, tileset and JSON, or **Separate tiles (zip)**; the aligned sheet can also join the Library.

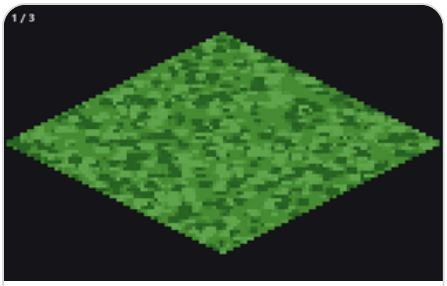

Shape

Your tile changes shape: a diamond for a game seen at an angle, or a six-sided hexagon.

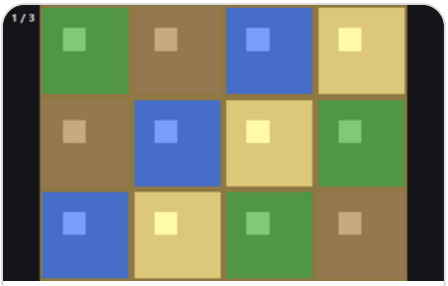

Placement grid

Your cells become diamonds, staggered like the paving of a game seen at an angle.

GAMES AND 3D

Create playing cards (Card Forge)

You will be able to compose a complete card game in Card Forge, from the face to the printer-ready pack, and deliver it to a virtual table with its rulebook. Almost everything is free and local; only illustration generation and translation by a language model are paid, quote shown first.

The Card Forge screen

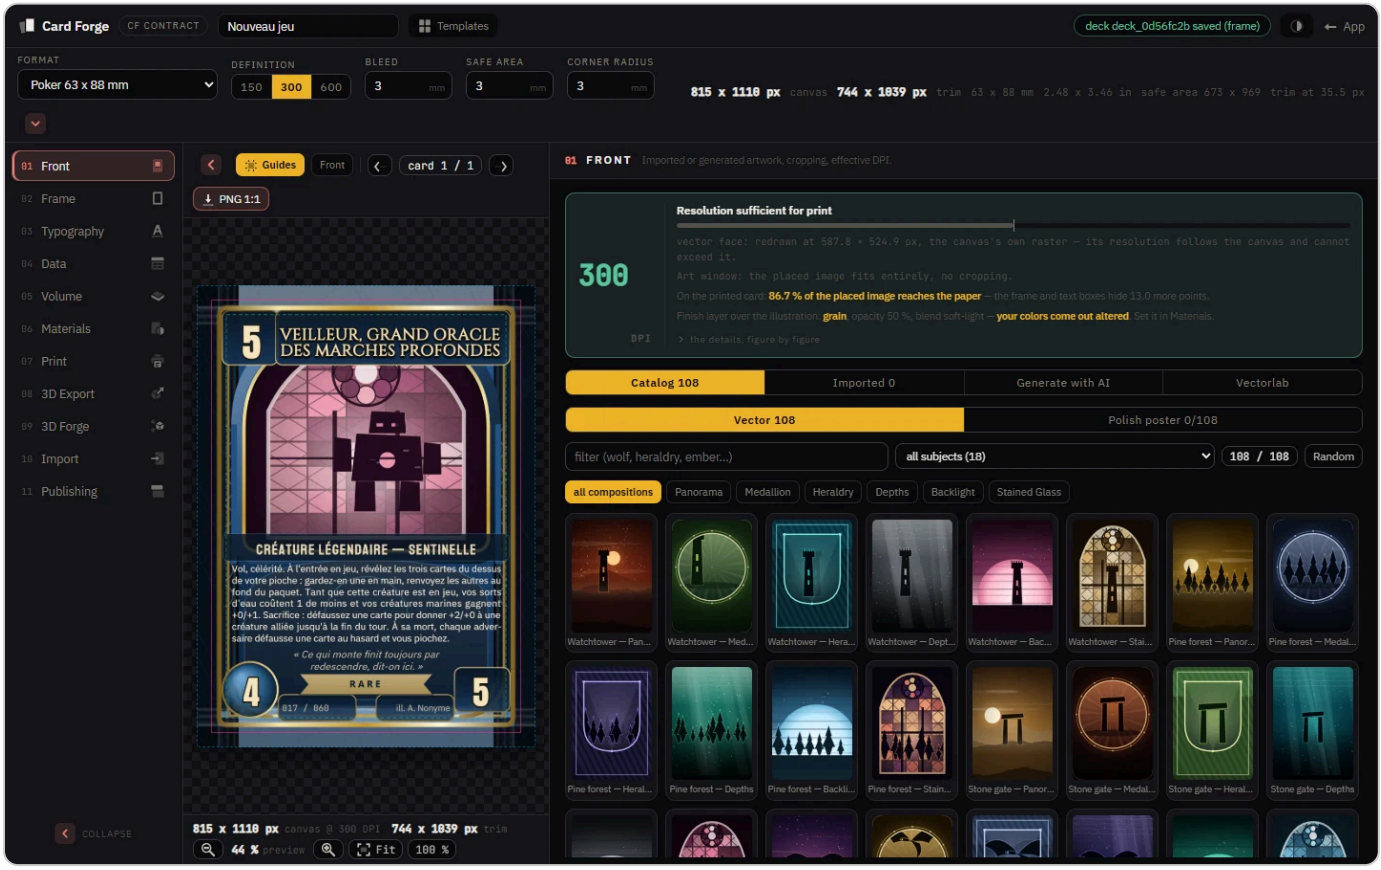

1. Eleven modules, from face to publishing 2. Format, DPI, bleed, safe zone

At the top the format, resolution, bleed and safe zone; on the left the eleven modules; in the middle the card.

1. Open the Cards tab of the Game Assets hub. **Templates** offers starter decks (**New deck from this template**).
2. At the top, set the format (poker, bridge, tarot, mini, square, domino, business card, jumbo, micro), the resolution (150, 300 or 600 DPI), bleed, safe zone and corner radius. These settings apply to every part and to the export.
3. On the left, eleven modules; in the middle, the card with its guides, front or back, card by card.

The eleven modules

01 Face · 02 Frame · 03 Typography

The illustration (imported, generated or taken from the local catalogue), borders, rarity and back, the title, rules and flavour text, with the shipped fonts.

04 Data

The deck comes from a CSV, a Notion export ZIP or a **Google Sheets link…**; each column feeds a field. The deck statistics (histograms, cost curve) and the “Back” line spot gaps.

05 Volume · 06 Materials

3D preview with real thickness and edge; eight PBR texture channels computed locally.

07 Print

Home sheets or printer templates, preflight, printer pack and duplex test chart (see below).

08 3D Export · 09 3D Forge

A GLB and a glTF with every PBR map; the 3D Forge turns the card into a printable object (see the 3D object).

10 Import · 11 Publishing

Import existing cards; deliver around the card: virtual tables, rulebook, mockup, product sheet.

Edit this face in the Vectorlab opens the rendered face in a Vectorlab document at the deck's physical size, with guides and a locked layer.

Several languages

Columns suffixed by language (name_fr, name_en…, twelve languages) are recognised: pick the active language and every part follows, with what is missing card by card. Translating a column with a language model is proposed card by card; nothing enters the table without your click, and the quote shows before any call (local Ollama: free).

Print

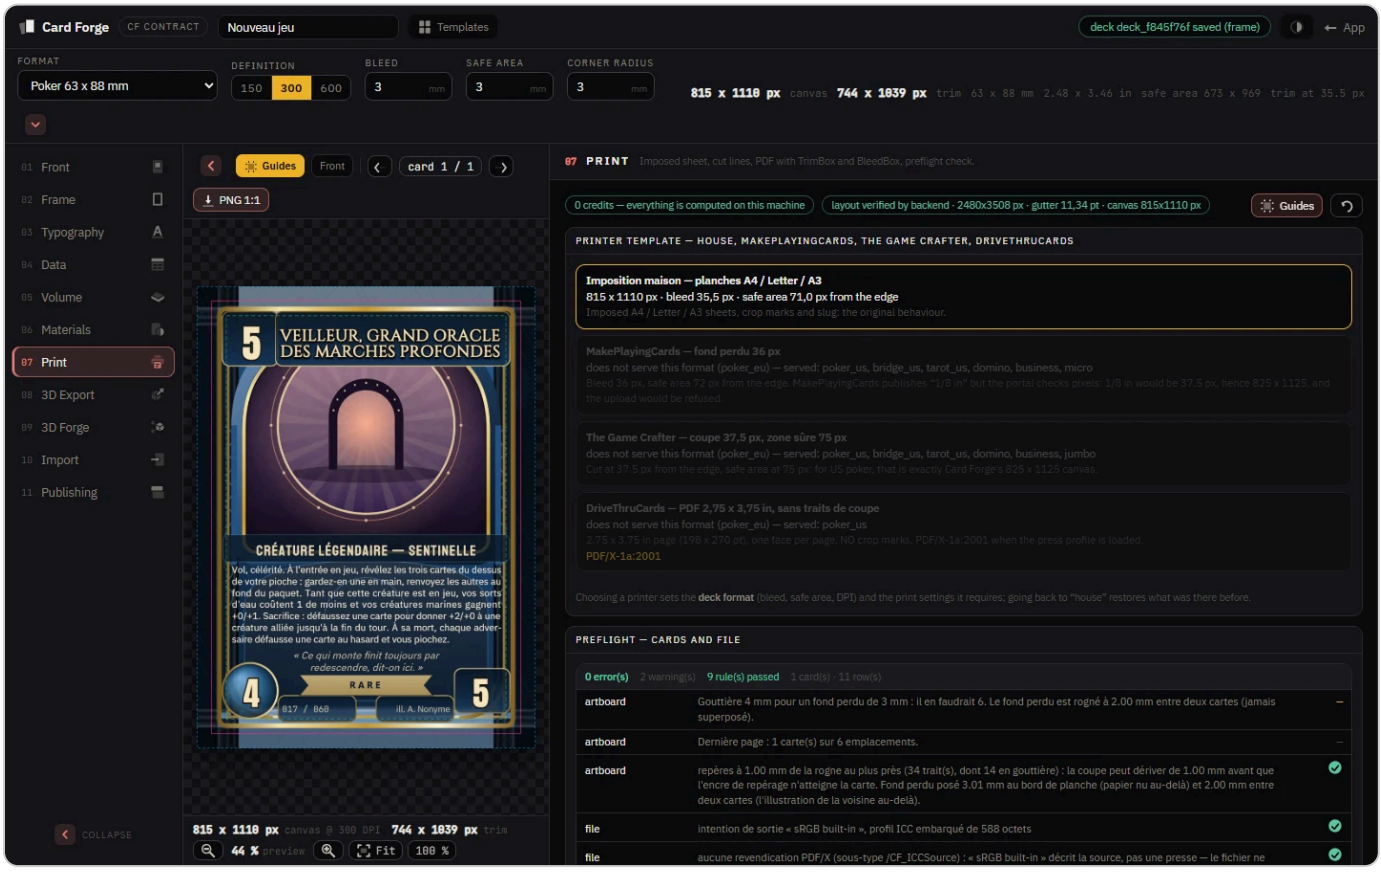

1. A printer's template in one click 2. DriveThruCards: PDF X-1a with a press profile

Module 07: home sheets, or the exact template of MakePlayingCards, The Game Crafter or DriveThruCards, preflight-checked.

1. Module 07: choose home imposition (A4 or Letter sheets, crop lines or crosses) or the exact template of MakePlayingCards, The Game Crafter or DriveThruCards (bleed, safe zone, DPI, pixels expected by the portal). An unverified format is greyed out and says so.
2. Preflight re-reads the produced file itself: bleed, crop marks, the PDF's TrimBox and BleedBox, PNG density. **Re-check** runs it again.
3. **Printer pack (.zip)** delivers named fronts and backs, padded to the deck size, with a manifest ready to upload. For DriveThruCards, the PDF becomes X-1a (CMYK) if the press profile is loaded; otherwise the gap is written on screen.
4. **Duplex test chart (PDF)** prints two pages with a vernier to a tenth of a millimetre to measure your printer's front-back offset.

Publish and deliver

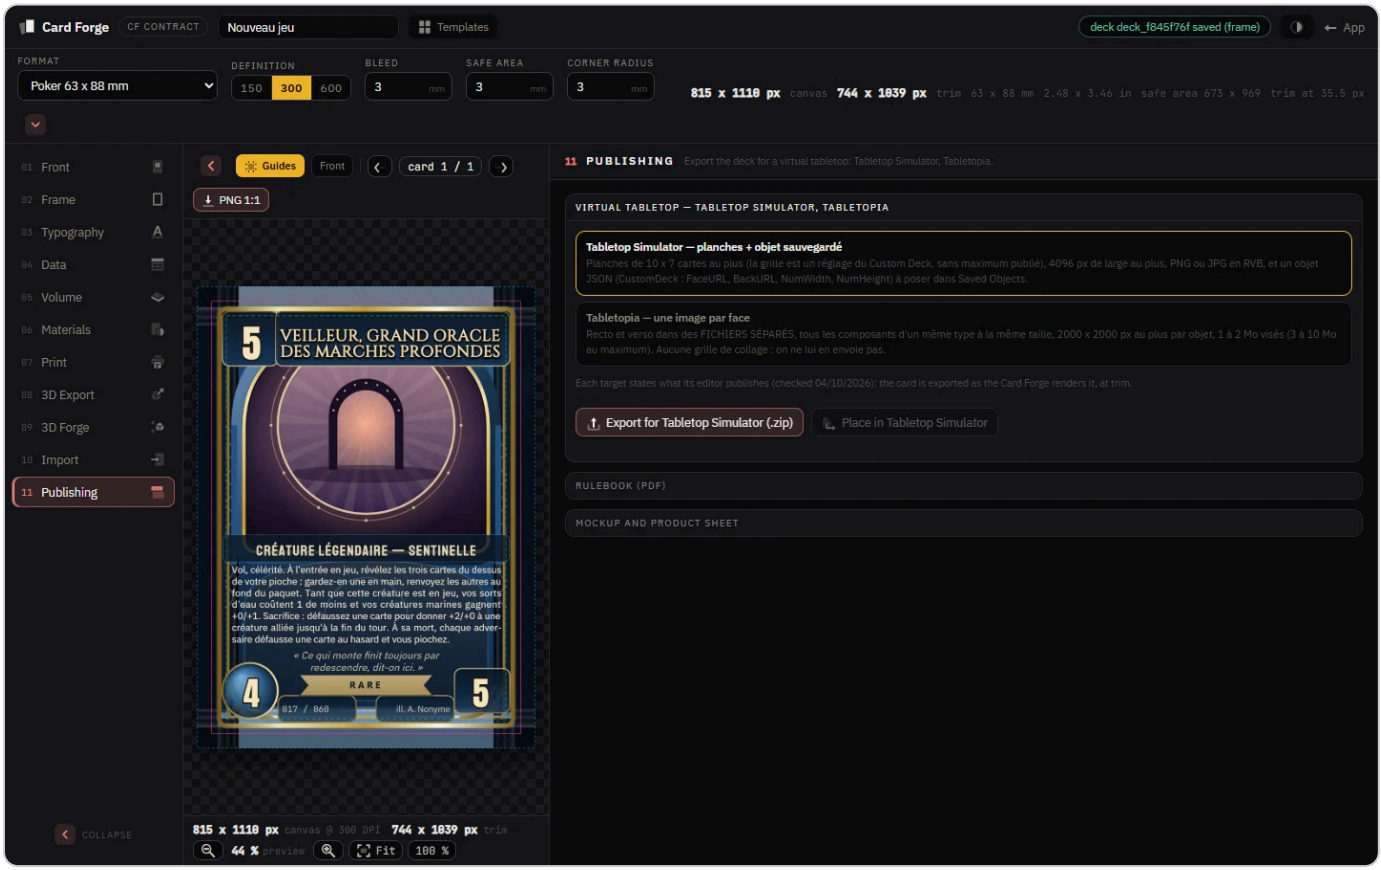

1. Export for a virtual table 2. 11 · Publishing

Module 11: virtual tables (Tabletop Simulator, Tabletopia), rulebook, fan mockup, product sheet.

Virtual tables

Export for Tabletop Simulator (.zip) (10 × 7 sheets and saved object) or **Export for Tabletopia (.zip)** (one image per face). **Place in Tabletop Simulator** copies the deck into your saved objects.

Rulebook

Composed at 300 DPI with the app's fonts (“#” headings) and a card sheet.

Mockup and sheet

Cards fanned out (five at most) in square, story or landscape format, and a product sheet with figures, ready to copy.

Deck art can also be generated in a batch: a template prompt with {columns}, a style, a bible entity; quote per row, cap per batch and your spending caps are checked before the run.

Good to know

Paid: generating a face or the deck art, translating with an online model, and the 3D Forge's image-to-3D engines. Everything else is free and runs on your PC.

GAMES AND 3D

From image to 3D object

You will be able to turn a card into a printable 3D piece with Card Forge's 3D Forge, and produce a complete 3D model (mesh, texture, skeleton, animations, formats) in the 3D Studio. Local pipelines are free; image-to-3D engines are paid, price shown before any run.

From card to object: the 3D Forge

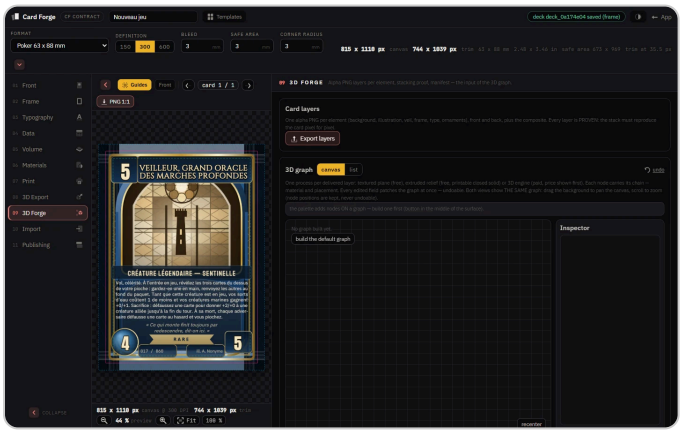

1. 1. Export the layers

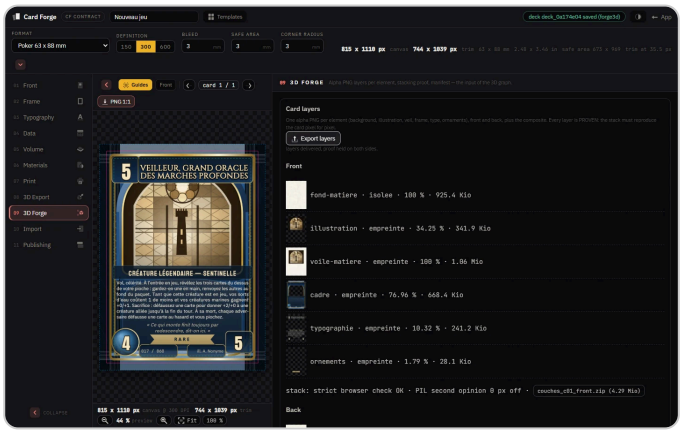

2. Six layers, stacking proven to 0 px

Module 09 Forge 3D: the card is exported in layers, each layer is processed in a graph, assembly produces a GLB and an STL.

1. In Card Forge, open module 09 3D Forge then **Export layers**: background, illustration, veil, frame, typography, ornaments and composite, front and back. The screen proves the stacking (0 px gap) before delivering, and the server checks it again.
2. Edit the graph: each layer gets a flat or relief treatment (closed solid), a Material Forge material tiled at physical scale, a finish (holographic silver or gold, lacquer, leather, brushed metal, emissive) and a placement. A relief with no depth entered takes its material's physical height, otherwise 0.6 mm.
3. The “3D mesh” treatment sends a layer to an image-to-3D engine: the price shows on the node and at the foot of the graph before any run.
4. **Build the 3D artifact** assembles the final GLB, and an STL if every element is a closed solid; otherwise the refusal says why. The manifest lists elements, engines and credits.

The “Game objects” panel also creates printable 3D pieces: a **Token**, a **Card relief token**, a pawn or a stand, plus the **Unfolded box (PDF)…**, the vector pattern of a tuck box at the deck's thickness.

The 3D Studio

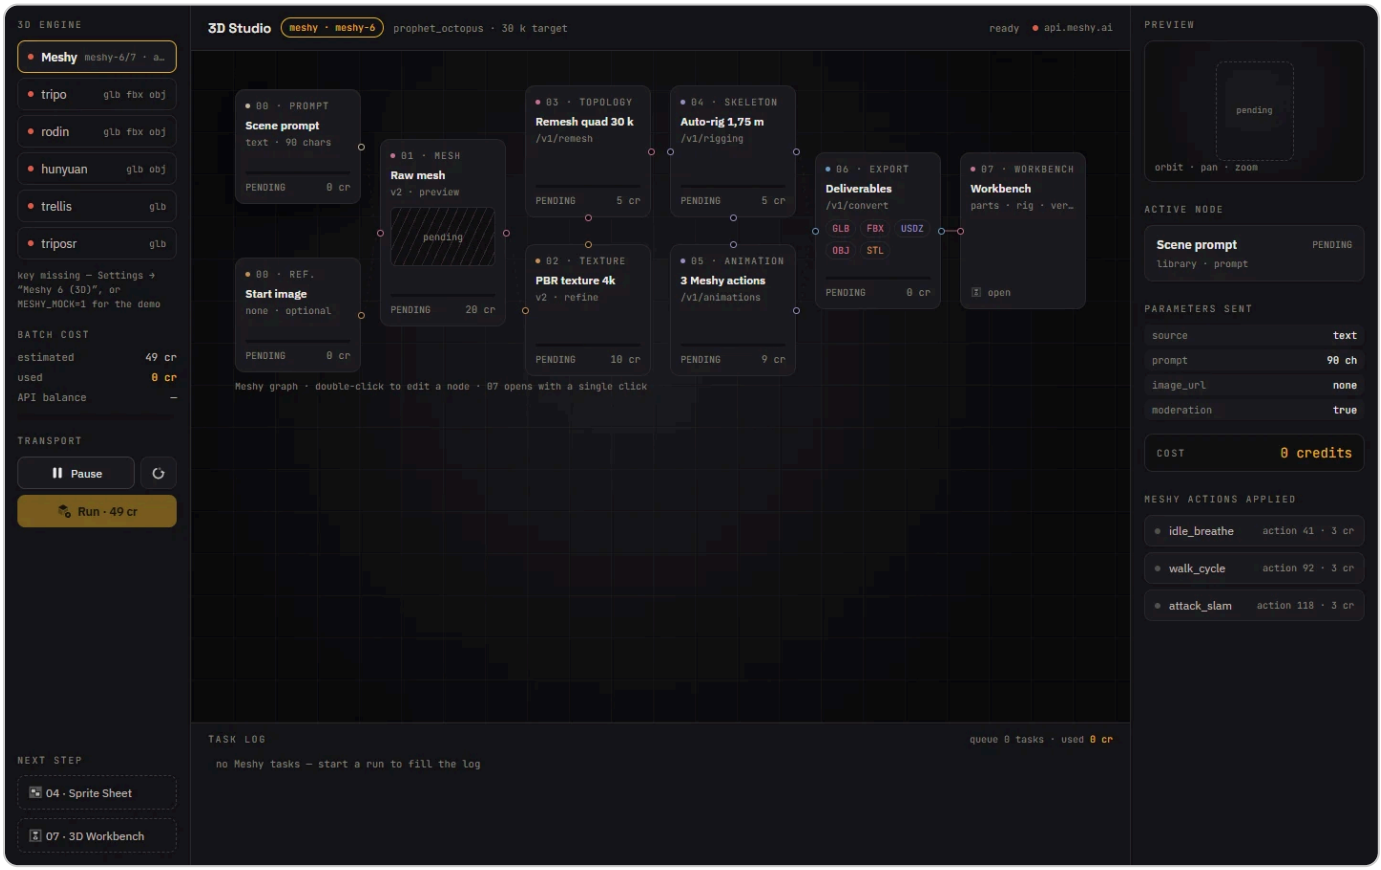

1. Each step is a priced node 2. The total price before launching

A Meshy graph from prompt to export: mesh, texture, topology, skeleton, animations; the series cost shows before Launch.

The 3D Studio (a tab of the Game Assets hub) chains a Meshy graph from prompt or start image to deliverables: mesh, PBR texture, topology, skeleton (auto-rig), animations, GLB, FBX, USDZ, OBJ or STL export, then the Workbench. Each node carries its cost; the series cost shows before **Run**, and the log follows each task. A Meshy key is needed; fal engines (tripo, rodin, hunyuan, trellis, triposr) use your fal key.

Prepare the views before paying for the mesh

1. In the hub's 3D card, the “Views first” block generates the views first (front, back, profiles): a failed view does not cost a mesh.
2. For free, reuse the views of a bible character's sheet, or build the view set from your photos (phone ones first, straightened and cut out).
3. Per thumbnail: **Remove the background** (local, free) or replay a single view with a corrected prompt (paid).
4. “Run · engine” launches the mesh on the approved views: only the engine is billed.

Rig, conversion, optimisation

fal workshop

fal workshop (rig a job) lists your 3D jobs: Meshy rig with animations (walk, run…), quote per task; greyed buttons say why (key, approval, texture).

Convert and import

Free local conversion to OBJ, STL, 3MF or glTF (STL and 3MF at a size in mm), or by Meshy to FBX, USDZ, BLEND. Import a model (OBJ, STL, glTF, GL makes it a job.

LOD and textures

The Optimise area derives levels of detail with measured loss and mobile, PC, print budgets, and exports textures to Blender, Unity, Unreal or Godot conventions.

Engine bench

The medians measured on your machine per engine (triangles, weight, watertightness, cost); add an existing job to it for free. An optional local GPU service makes Hunyuan free.

Good to know

Paid: image-to-3D engines (fal in dollars, Meshy in credits), rigging, Meshy conversion and generated views. The price is announced first, the spending table prices rig and refine, and texturing without a fal key is refused straight away instead of failing later.

GAMES AND 3D

Create materials (Material Forge)

You will be able to create seamless materials in Material Forge, from a free catalogue, a generator, a photo or a description, and export them to Blender, Unity, Unreal or Godot. The catalogue, generators and maps are free and local; only the AI-generated base image is billed.

A material from the catalogue

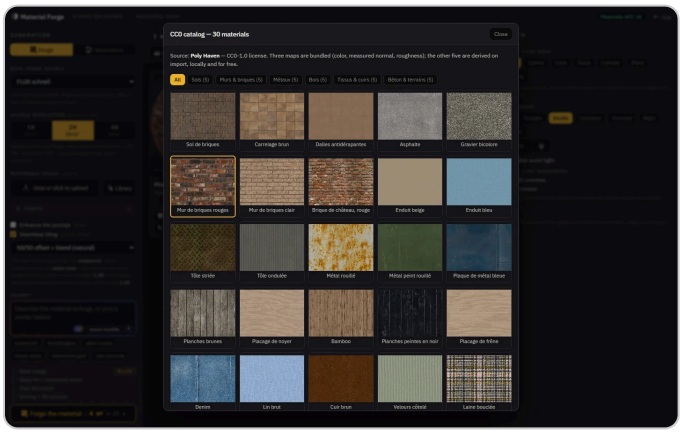

1. Pick, then Import

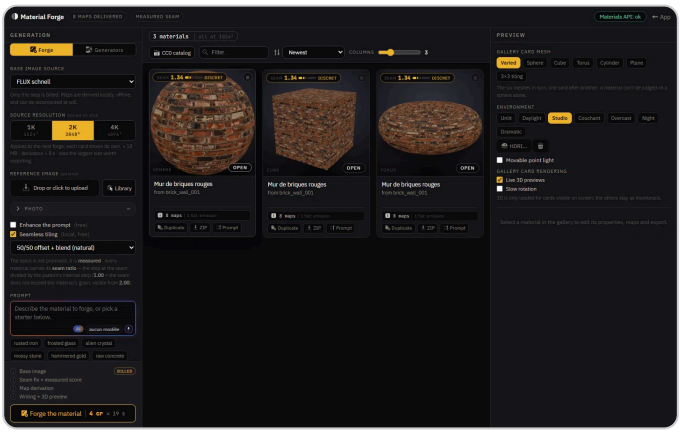

2. The material, in 3D preview, with its 8 maps

Thirty CC0 Poly Haven materials, shipped: one click imports them, their PBR maps are derived locally.

1. Open the Materials tab of the Game Assets hub, then **CC0 catalog**: thirty free Poly Haven materials in six families (floors, walls and bricks, metals, wood, fabrics and leathers, concrete and terrain).
2. Pick one or more then import: missing maps are derived locally, the source credit is kept.
3. The material lands in the gallery with a real 3D preview, its eight PBR maps and its measured seam score. One click reopens it in the editor.

The generators

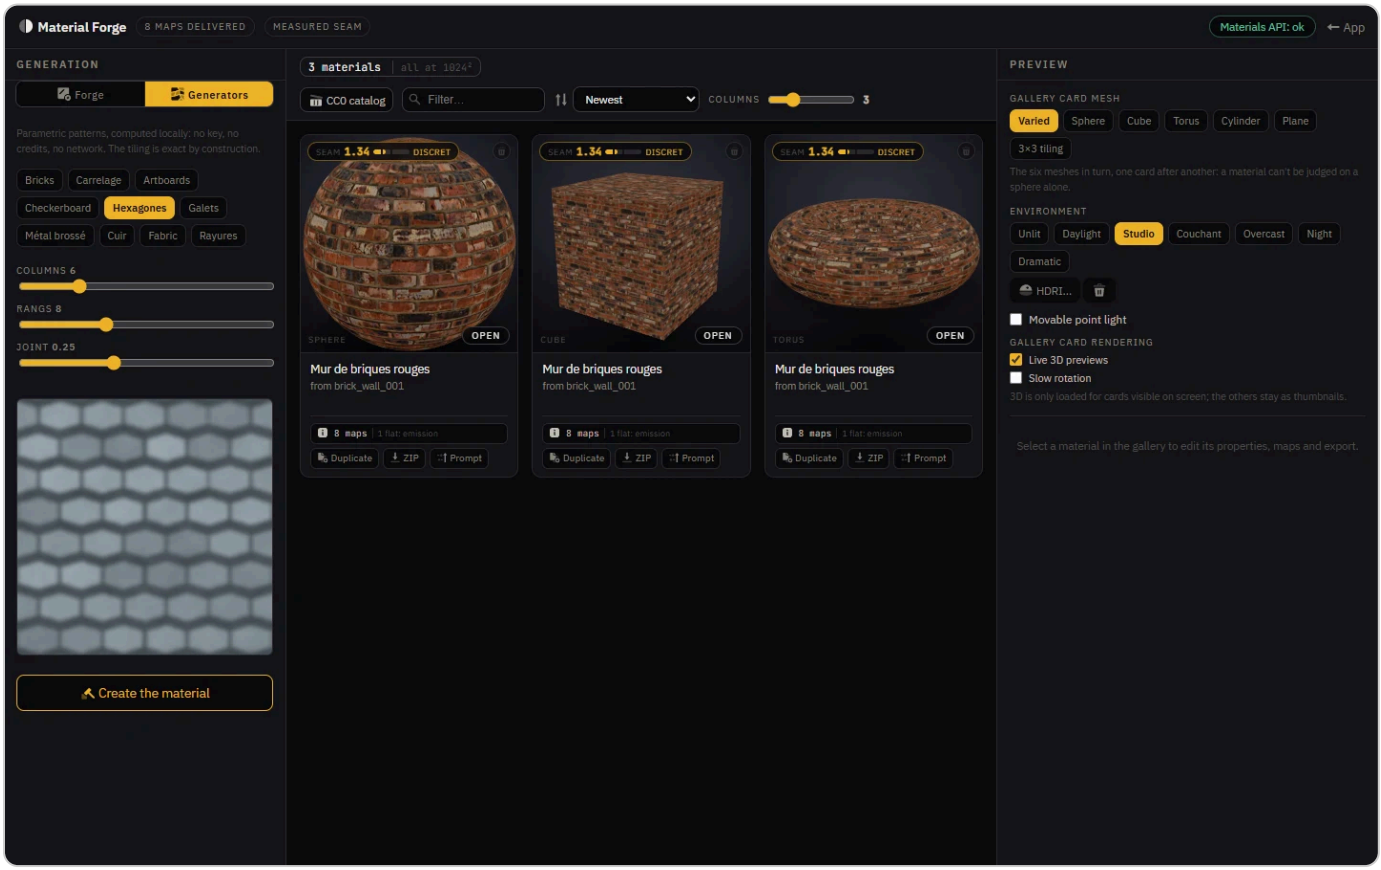

1. Bricks, tiles, hexagons, leather… 2. Create the material: free

Ten seamless generators, local and free, tuned live before creating the material.

The **Generators** tab offers ten seamless patterns: bricks, tiles, planks, checkerboard, hexagons, pebbles, brushed metal, leather, fabric, stripes. Tune them live, then **Create the material**. Local and free.

Forge from a description or a photo

1. **Forge** tab: describe the material (rusty iron, frosted glass…) or start from a reference image; pick the resolution (1K, 2K, 4K) and the seam method.
2. **Forge the material**: only the base image is billed (price shown on the button); seams, maps and preview are computed locally, recomputable at will.
3. A raw photo becomes a material: click the four corners of a surface shot at an angle to straighten it, then remove the baked lighting with an adjustable strength (the badge shows the before and after gap). On a phone, the field opens the camera.

Judge and tune

Environments

Seven HDRIs provided, and your own (.hdr or 2:1 image) to add or delete. Varied gallery meshes: sphere, cube, torus, cylinder, plane, 3×3 tiling.

Compare, My model

Compare opens a second preview, same environment and camera; “My model” puts the material on a Workbench model.

Physical height

The height map declares its millimetres (0 to 20 mm), reused by the card 3D Forge's relief.

Finishes

A preset (lacquer, leather, animated emissive, brushed metal) is tried in the preview before being applied or cancelled. Emission above 1 keeps its hue.

Export

Each material exports as a ZIP with a manifest, following a naming convention: standard, Blender (Principled BSDF slots indicated), Unity, Unreal or Godot. It also serves as a source in Tile Lab and as a material in the 3D Forge.

Good to know

Only the base image of “Forge the material” is billed, within your caps. Catalogue, generators, photo, delighting, maps and exports are free.

GAMES AND 3D

Prepare a part before slicing (Workbench)

You will be able to prepare a 3D part before sending it to the slicer: scale it, repair it, lay it down, check its overhangs, hollow it, arrange it on the bed, then export it to the printer or a game engine. Free and local. Nothing is irreversible: every gesture that changes the part writes one more version, and all of them stay on disk.

Open the Workbench

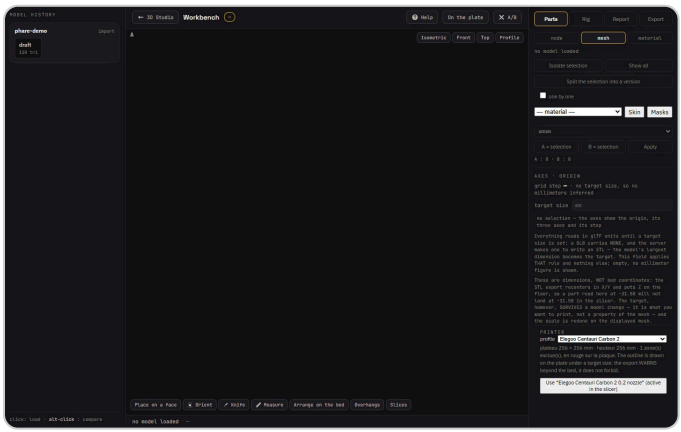

1. Load a version

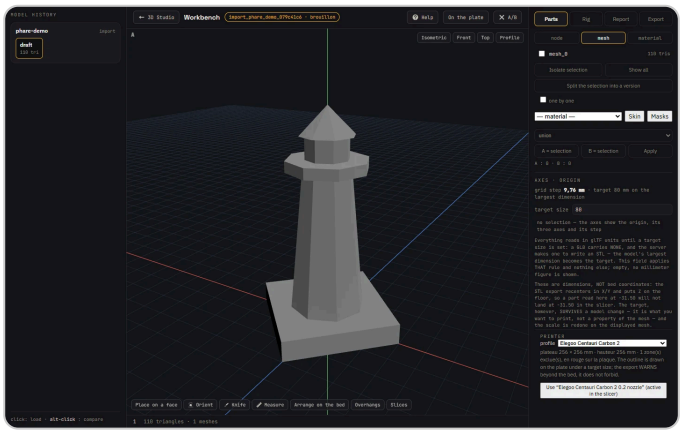

2. 80 mm: the target size 3. The printer and its real bed

On the left the model's life, in the middle the part, on the right the Parts, Rig, Sheet, Export tabs; the target size gives the millimetres.

In the 3D Studio, click the `07 · 3D Workbench` node; the Workbench takes over the page, and `3D Studio` at the top left brings you back. On the left, **the model's life** lists each model and its versions: one click loads a version, `Alt`+click compares it (A/B view). On the right, four tabs: `Parts`, `Rig`, `Report`, `Export`. `Help` opens this chapter's step-by-step and glossary.

The main gesture

1. **Set a target size.** `target size` field: the largest wanted dimension, in millimetres (80 for a figurine, 250 for a board). A 3D file carries no unit: while this field is empty the Workbench shows “glTF units” and refuses whatever needs a physical dimension (hollowing, arranging, exporting for print). A refusal beats an invented millimetre.

2. **Pick the printer.** The Elegoo Centauri Carbon 2 profile is built in (256 × 256 mm bed, 256 mm height, one excluded area); installed OrcaSlicer or ElegooSlicer profiles show up too, read, never modified, and you can enter one by hand. On the plate, the green outline is the real bed, the red rectangle the area where nothing may sit.

3. **Repair.** `Report` tab, `Repair in one click`: coincident vertices welded, duplicate faces and flat triangles dropped, normals unified. Capping holes is unchecked by default: it adds material. The detail appears in the bottom bar.

4. **Lay the part flat.** `Place on a face` (`F`), then click the face that must touch the bed. The gesture joins the queue, and `write the version` saves it. If unsure, *Orient* proposes three ranked poses (support, overhang, height) in the Report tab. A computed orientation does not always find the best pose: **always review the proposal**, no score knows which face you want to look good. If the queue holds both a placement and a recentring, the bar flags the contradiction and greys out writing.

5. **Look at the overhangs, then the slices** (see below).

6. **Hollow it, when it helps.** Report, `Hollow`, 2 mm wall to begin with: the solid inside becomes a shell. It is first a resin (SLA) gesture; with filament (FDM), the slicer already fills the inside. It needs a closed mesh, and the Workbench states the thickness that holds when the requested wall is too thick. For resin, then drill a drain hole: *Drill (drain hole)*, 3 to 4 mm, then click the face, preferably near the bed. The drill goes through the one wall under the click and stitches the hole closed.

7. **Arrange on the bed**, **measure**, then **export** (see below).

> **The order that works**: repair, lay down, inspect, hollow, export. Repair first, because every other gesture starts from sound geometry, and hollowing refuses an open mesh.

Look before slicing

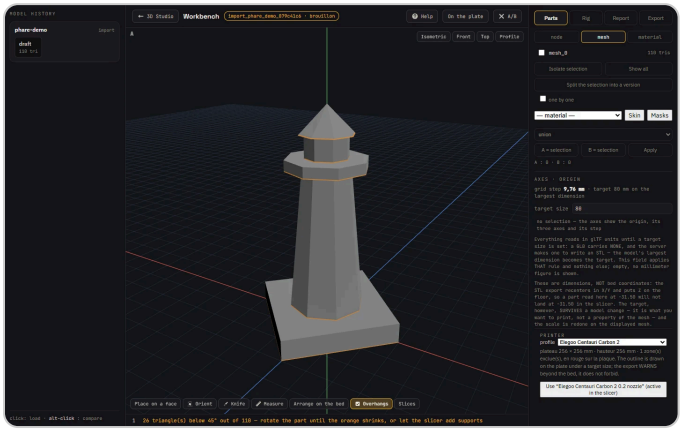

1. 26 triangles under 45°: the balcony

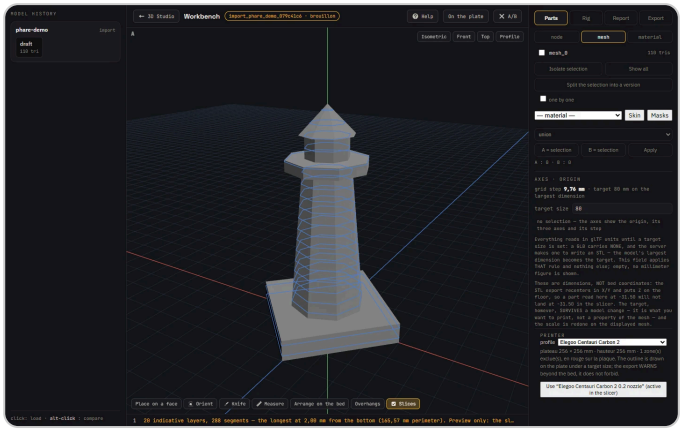

2. Twenty sections, in blue

Overhangs paints overly tilted faces orange; Slices shows twenty sections. Previews: nothing is written.

`Overhangs` paints orange the faces leaning less than 45° from horizontal: there, a layer has almost nothing under it and the slicer will add a support. Re-lay the part until the orange recedes. `Slices` draws twenty cross-sections in blue to see the real shape at each height. Both are **indicative previews**: nothing is written, and the slicer does the real slicing.

The Report tab

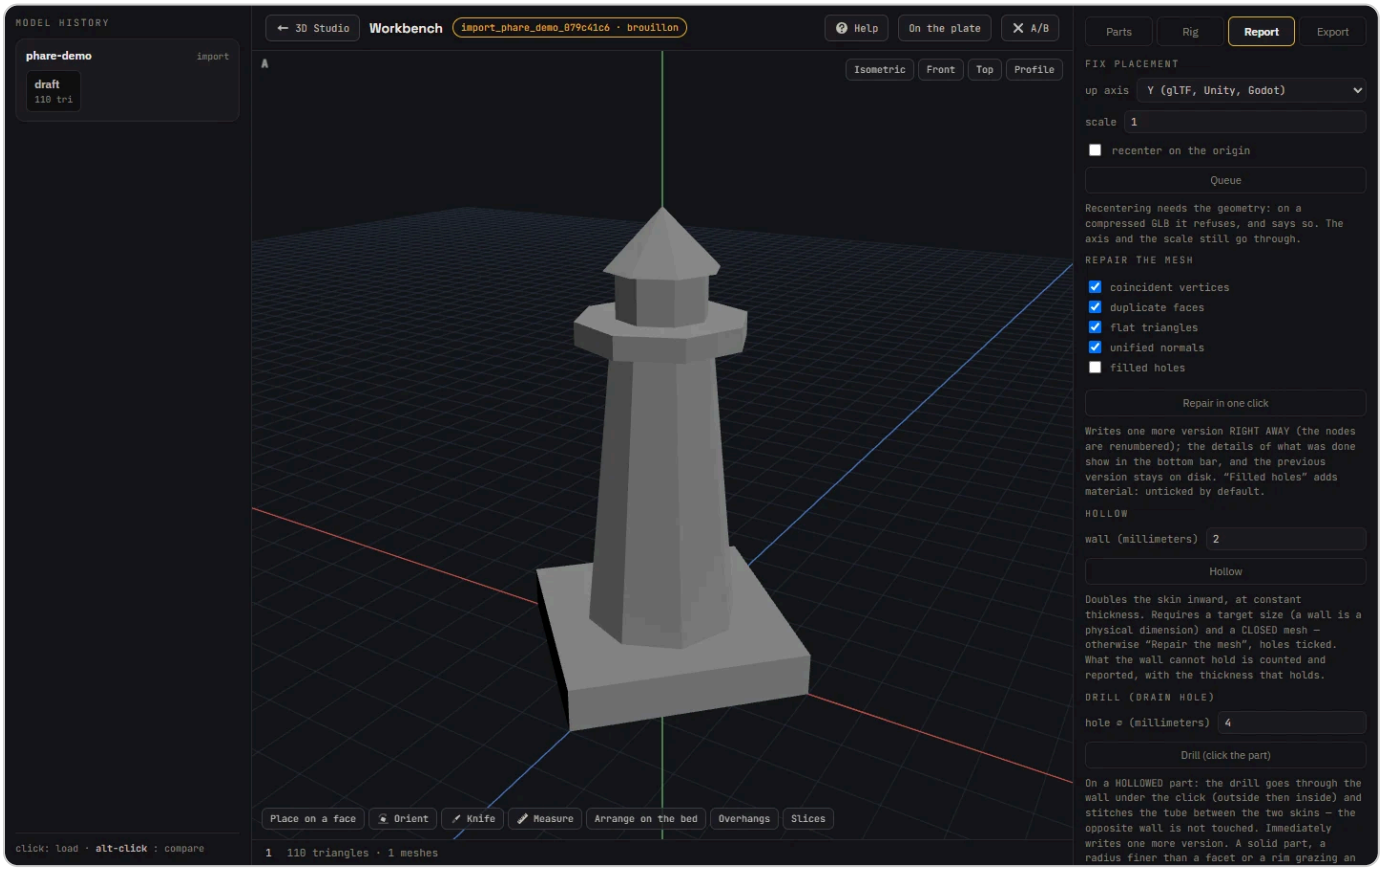

1. Each gesture writes a version

The Sheet tab: one-click repair, proposed orientations, hollow, drill, decimate, materials.

Decimate, extract

Decimate reduces triangles towards a target and says if it is not reached; extracting one by one splits each element into a sibling version.

Materials

Colour, opacity, metal, roughness, emission, alpha mode and double side, with preview; only the touched fields are written.

Measure

Measure: two clicks give a distance, its x/y/z components and the angle between the two faces, drawn in the scene.

Parts, booleans, connectors

Booleans

Parts tab: **A = selection**, **B = selection**, then union, difference (A − B) or intersection and **Apply**. The bar states the seam and lost materials.

Cut and assemble

The **Knife** cuts the part; its bar adds a connector: dowel (holes and a separate pin), plug and its socket with clearance, or a sliding dovetail.

Skin

A Material Forge material on the ticked parts: **Skin**. **Masks** shows cavities (red) and edges (green) in view B.

The **Rig** tab draws the skeleton through the skin, highlights a bone's chain, colours its influence weights, rotates a bone with three x/y/z sliders and plays the shipped animation clips.

Arrange on the bed

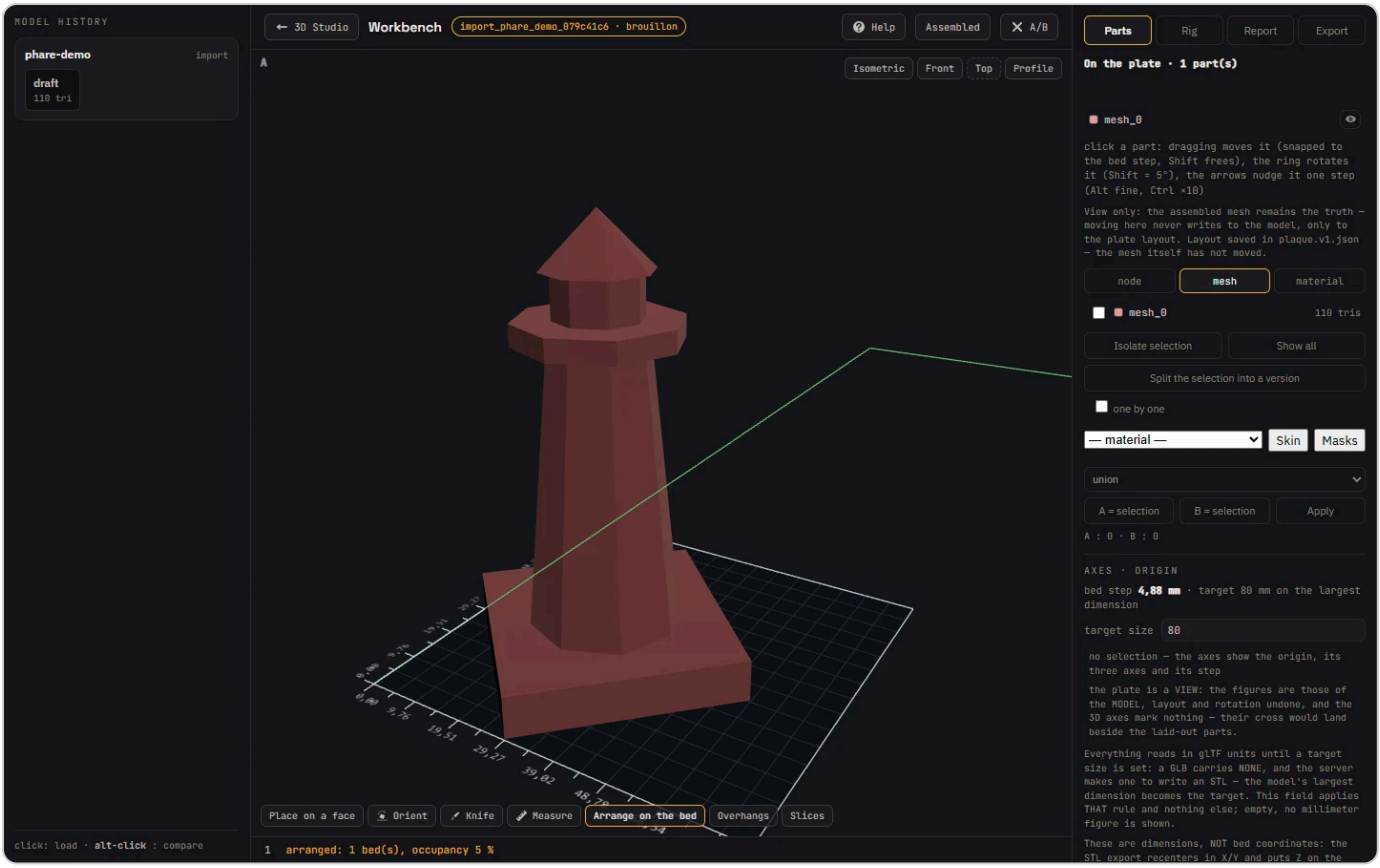

1. The green outline: the real bed 2. Arranged, rotated if it saves space

On the plate shows the printer's real bed; Arrange on the bed places the parts there with 2 mm gaps.

On the plate separates the pieces, then **Arrange on the bed** lays them inside the real bed, rotated when that saves room, 2 mm apart, on up to eight beds; what does not fit is reported. Drag a piece (snapped to

the bed grid, `Shift` frees it), nudge it with the arrow keys: its position, size and angle are shown. The plate is a **view**: it is saved, but it does not change the model.

Export

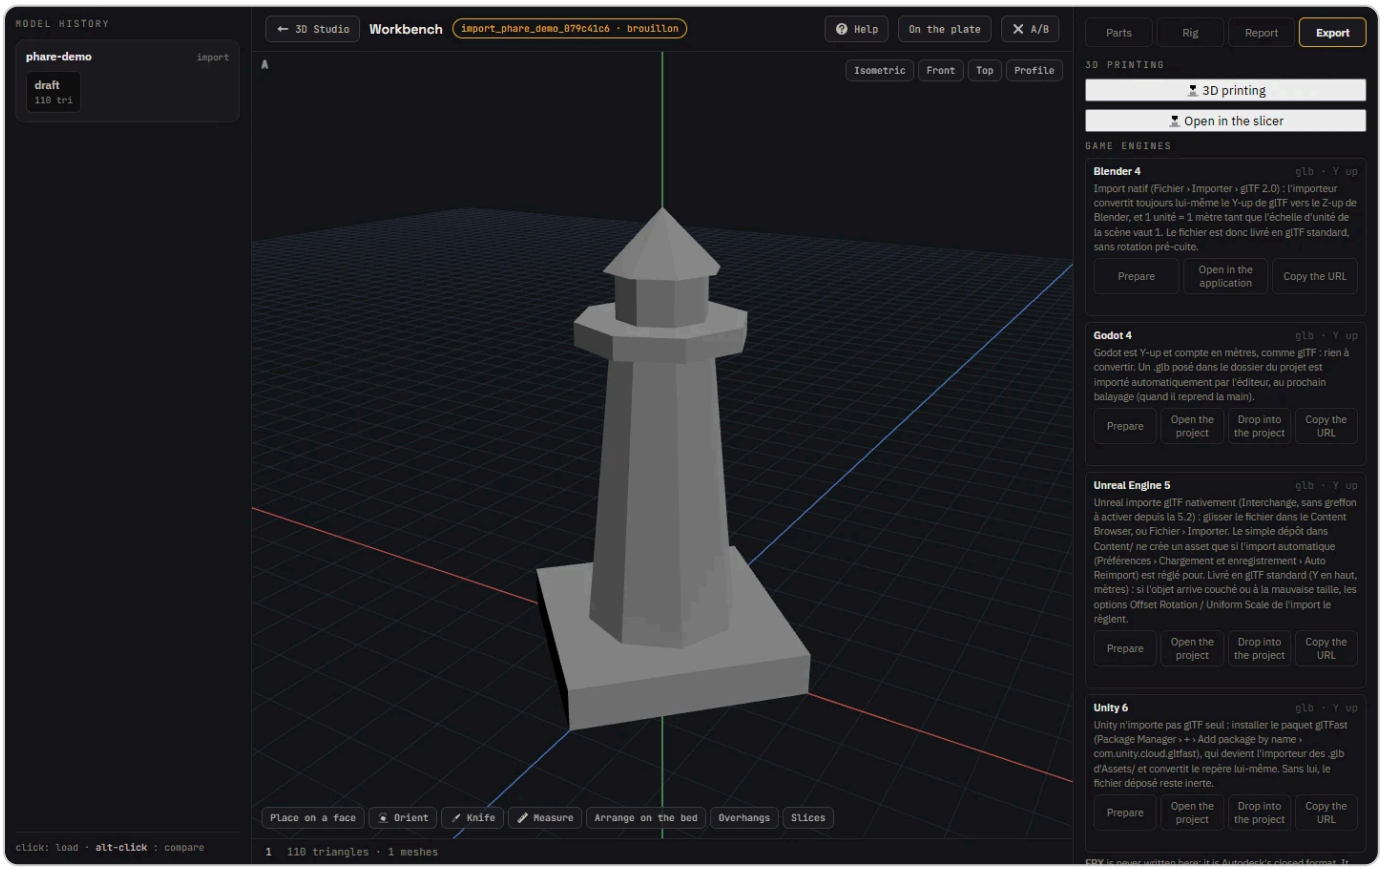

1. To printing, in millimetres

Export tab: STL and 3MF of the shown version, or glTF prepared for Blender, Godot, Unreal or Unity.

1. `Export` tab, `3D printing`: the STL and 3MF of the **displayed version**, at the target size in millimetres (see 3D-print your creations). `Open in the slicer` takes you there.
2. For an engine: `Prepare` writes one glTF per target (Blender, Godot, Unreal, Unity) with its import sheet; dropping copies it into the engine project; `Open` launches Blender or the project; `Back from Blender` re-imports a fixed GLB as one more version and says what was lost.

Glossary

The words your slicer will use, explained once.

Term	What it is
Base (bed contact)	The face that touches the bed. Wide and flat, the print holds; narrow, it comes loose.
Overhang	A wall leaning too far to rest on the layer below. Prusa puts the clean limit without support between **45 and 60°** from horizontal.
Support	A disposable structure holding overhangs up. The slicer generates it; the Workbench only helps you need less of it.
Brim	A one-layer collar around the part, to keep its corners from lifting.
Raft	A platform printed UNDER the whole part. Safer than a brim, costlier, and it leaves a rough face.
Skirt	An outline drawn next to the part, not touching it, to prime the flow before starting.
Infill	The pattern inside. *Gyroid* is about equally strong in every direction; *Lightning* only supports top surfaces. 15 % is usually enough.
Seam	Each wall loop starts and ends somewhere, leaving a visible vertical line. The slicer places it (*Nearest*, *Aligned*, *Rear*) or scatters it (*Random*).
Retraction	Filament pulled back during travel moves, to avoid strings (*stringing*).
Layer	A horizontal slice. 0.2 mm is the usual compromise; 0.12 mm for detail, 0.28 mm for speed.
Perimeter (wall)	The number of wall loops. Two perimeters at 0.45 mm make a 0.9 mm wall: the minimum to aim for when modeling.
Bridging	Printing a layer over empty space, between two supports. It works over a few centimeters, with good cooling.
Warping	Corners lifting as they cool. Remedies: clean bed, brim, closed enclosure. Common with ABS/ASA, rare with PLA.
Watertightness	A solid is watertight when its surface is closed, with no hole or free edge. That is what the Workbench bar calls "fermé" (closed).
Manifold	Every edge belongs to exactly two faces. A mesh can be closed and still be hard to slice cleanly if it is not.

Decimation	Reducing the triangle count. A 500,000-triangle mesh prints no better than a 100,000 one, but slows everything down. *Report* tab › *Decimate*.
Hollowing	Replacing the solid inside with a shell of given thickness. Mostly useful in resin; 2 mm is a good starting point.
Drain hole	The hole(s) letting uncured resin (and air) out of a hollowed part. In the Workbench: *Report* tab › *Drill (drain hole)*.

Where to learn more

Every link below was opened and checked on 05/10/2026 (5 October 2026). They are grouped by what they teach you.

Topic	Resource	Checked
Supports and overhangs	Prusa — Support material (https://help.prusa3d.com/article/support-material_1698) (the threshold is measured from horizontal, 90 = vertical)	05/10/2026
Modeling for print	Prusa — Modeling with 3D printing in mind (https://help.prusa3d.com/article/modeling-with-3d-printing-in-mind_164135) (45–60° overhangs, walls, tolerances)	05/10/2026
Bridges	Prusa — Poor bridging (https://help.prusa3d.com/article/poor-bridging_1802)	05/10/2026
Layers and perimeters	Prusa — Layers and perimeters (https://help.prusa3d.com/article/layers-and-perimeters_1748)	05/10/2026
Infill	Prusa — Infill patterns (https://help.prusa3d.com/article/infill-patterns_177130)	05/10/2026
Seam	Prusa — Seam position (https://help.prusa3d.com/article/seam-position_151069)	05/10/2026
Skirt and brim	Prusa — Skirt and brim (https://help.prusa3d.com/article/skirt-and-brim_133969)	05/10/2026
Warping	Prusa — Warping (https://help.prusa3d.com/article/warping_2011)	05/10/2026
PLA filament	Prusa — PLA (https://help.prusa3d.com/article/pla_2062) (for Prusament PLA: nozzle 215 °C first layer then 210 °C, bed 60 °C)	05/10/2026
PETG filament	Prusa — PETG (https://help.prusa3d.com/article/petg_2059) (nozzle 230 °C then 240 °C, bed 85 °C then 90 °C)	05/10/2026
Choosing a filament	Prusa — Filament material guide (https://help.prusa3d.com/filament-material-guide) (comparison table of materials)	05/10/2026
Print defects	Simplify3D — Print Quality Guide (https://www.simplify3d.com/resources/print-quality-troubleshooting/) (about 26 defects, with photos)	05/10/2026
Slicer — prepare	OrcaSlicer — Object manipulation (https://github.com/SoftFever/OrcaSlicer/wiki/prepare_object_manipulation)	05/10/2026
Slicer — orient	OrcaSlicer — Auto orientation (https://github.com/SoftFever/OrcaSlicer/wiki/prepare_auto_orient) (which warns it may not always find the best pose)	05/10/2026
Slicer — calibrate	OrcaSlicer — Wiki (Calibrations) (https://github.com/SoftFever/OrcaSlicer/wiki/Home) (temperature, flow, retraction, pressure advance…)	05/10/2026

Machine	Elegoo — Centauri Carbon 2 wiki (https://wiki.elegoo.com/centauri-carbon-2-combo) (troubleshooting guides and error codes)	05/10/2026

These links come from slicer makers and manufacturers: they describe THEIR machine and THEIR software. Numeric settings (temperatures, speeds) apply to the hardware they name; on another machine they are starting points, not truths. Two repair tools often recommended elsewhere, **Microsoft 3D Builder** and **Autodesk Meshmixer**, are no longer maintained, so they are not listed here.

GAMES AND 3D

3D-print your creations

You will be able to send everything the app produces to your 3D printer: generated models, Workbench parts, 3D Forge cards and Vectorlab drawings. Everything is local and free: no provider call, and the app never generates G-code.

Before you start

The pipeline is designed for the Elegoo Centauri Carbon 2 (256 × 256 × 256 mm bed) and its slicer ElegooSlicer, shipped with the machine; OrcaSlicer works too. Install the slicer once: Windows associates .3mf files with it, and **Open in the slicer** uses that. Without an association, the app tells you; then set `SLICER_PATH` (the slicer's path) in the data folder's .env file.

Each export writes a folder in `DeepotusVideoGenData\assets\print3d\`: the STL (the mesh), the 3MF (the one you open, in millimetres and named) and `impression.json` (provenance, scale, watertightness).

Where to export from

1. **From the Workbench**: Export tab, **3D printing** writes the displayed version (repaired, laid down, hollowed…) at the target size. It is the recommended path: see Prepare a part.
2. **From Game Assets · 3D**: each generated model has its **3D printing** button. Give the largest dimension in millimetres (80 for a figurine); the part is centred and placed on the floor. The file "optimised" for game engines is refused with a message saying what to take.
3. **From the cards' 3D Forge**: when every element is a closed solid, the manifest delivers an STL and the print button appears. Leave the size empty to keep the card's real millimetres. Watertightness is marked *guaranteed* there.
4. **From Vectorlab**: File › **3D Print…** turns the drawing into parts: one height per layer, tiles, a bevelled logo, a map relief or pixel art. See 3D printing from Vectorlab.
5. **Slice in the slicer.** The 3MF opens at the right size, laid on the bed. Supports, fine orientation, infill, filament profiles: that is the slicer's job. Above 256 mm, the export warns you without blocking you.

Good to know

> ⚠ The watertightness of meshes from image-to-3D engines is marked *unknown* in `impression.json`: ElegooSlicer and OrcaSlicer repair small defects on import and report it. Go through the Workbench first to repair in one click. Only the cards' 3D Forge guarantees its solids, through its “closed solid” check.

GAMES AND 3D

Stage shots on the 3D Set

You will be able to block a shot in 3D on the Stage before any paid run: place volumes or a real mesh, frame at the film's format, set a camera move, then send it all back to the storyboard shot. Free and local.

Compose the scene

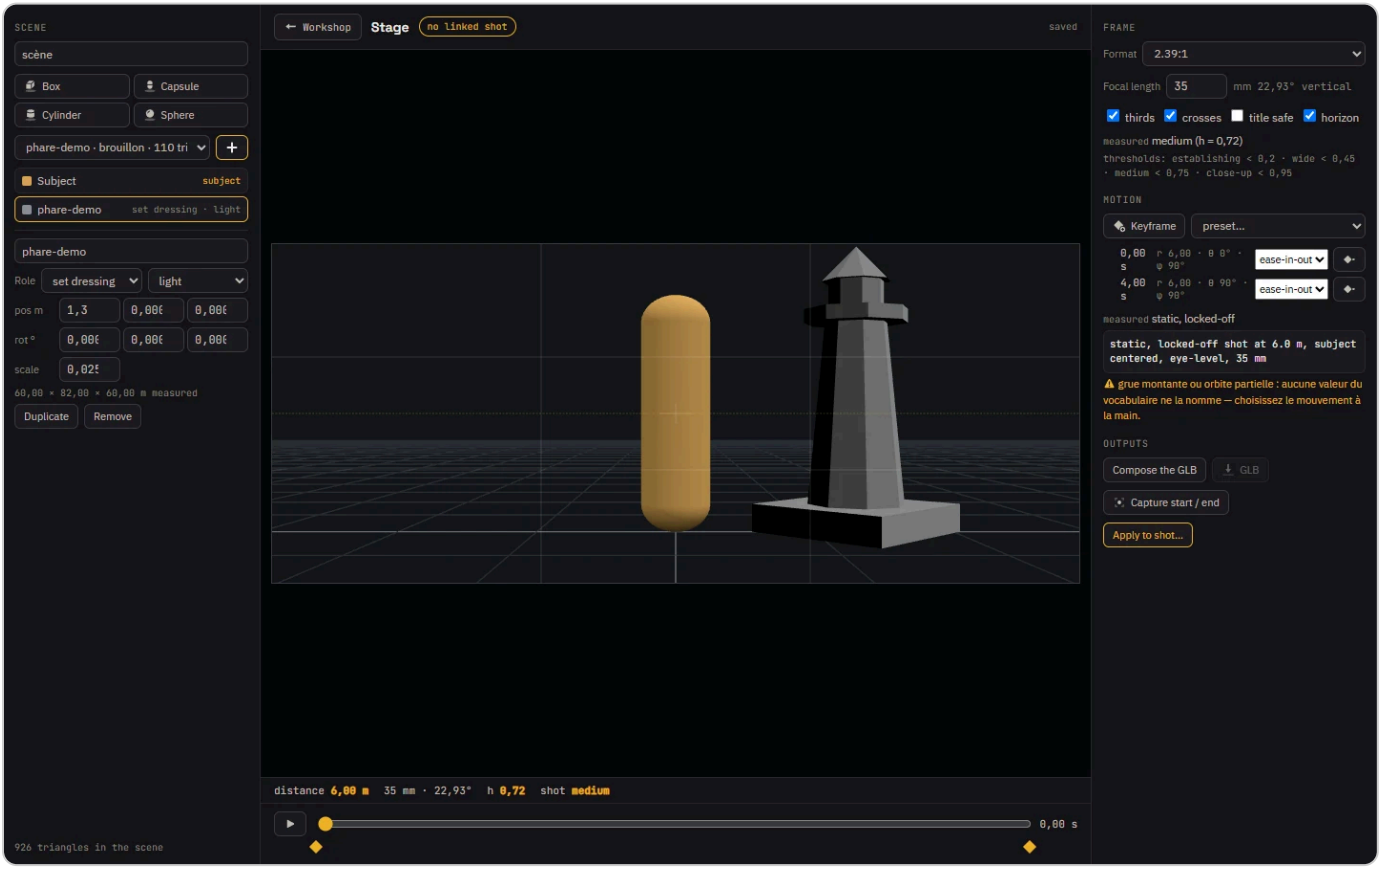

1. The lighthouse, placed from the Workbench 2. 2.39:1 frame, focal length, guides 3. Capture start and end to the Library

Simple volumes or a mesh, a frame at the film's format, then a camera move in keyframes.

1. From a shot card of the Workshop storyboard, the Stage button opens that shot's scene; **Workshop** brings you back.
2. Place simple volumes: **Box**, Capsule, **Cylinder**, **Sphere**, or a mesh from a 3D job or a bible entity, chosen in the list (grouped under “Bible entities” and “3D meshes”) then placed with the button next to it.
3. Each object has a name, a role (subject, set dressing, marker), a position, a rotation, a scale and its measured dimensions; **Duplicate** or **Remove**.

Frame and animate

Frame

16:9, 9:16, 1:1 or 2.39:1 format, focal length in millimetres, guides (thirds, crosses, title safe, horizon). The shot type (wide, medium, close-up) is measured from the subject's height in the frame.

Motion

`K` places a camera keyframe at the playhead; a preset lays a dolly-in, a 90° orbit, a crane down or a locked-off shot. `Space` plays the move in a loop.

Outputs

1. `Capture start / end`: the frame of the first and last keyframe goes to the Library, with its provenance.
2. `Compose the GLB`: the scene's GLB, versioned with its sheet.
3. `Apply to shot…`: proposes to the linked shot the shot type, the move, the motion prompt and the end images, compared before and after; nothing changes without your confirmation.

Good to know

The Stage is for deciding before paying: the captured end images become the start and end frames of a video shot, and the measured move feeds its prompt.

REFERENCE

End-to-end recipes, from idea to post

You will be able to chain screens to go from an idea to a published or printed result. Each path points to the chapters that detail every step; any cost is announced before each run.

An image, then a post

1. In the Library, describe the image then **Create image** (paid, price shown), or upload your own.
2. Retouch it if needed in Photolab: **Send to...** › Photolab, then Save.
3. **Send to...** › Scheduler creates a post draft; the Scheduler sets caption, channels and slot. Nothing goes out without your click.

A reel from the news

1. In News, let the free ranking sort today's articles.
2. **Daily pipeline**: pick the reel form (animated cards are free), then **Prepare (free)**.
3. Review the script and caption, then **Schedule batch**: the post lands in the Scheduler with its reel.

An episode of your novel

1. In Chapters, paste your chapter; the Workshop derives the breakdown, the bible and the storyboard.
2. Block the tricky shots on the 3D Stage before paying for any video.
3. Assemble the episode, finish it in the Editor (titles, mix, subtitles), then deliver in each network's formats.

A printed card game

1. In Card Forge, start from a template, import your data (CSV, Google Sheets, Notion) and set the format.
2. Retouch a face in Vectorlab if needed (“Edit this face in the Vectorlab”).

3. Module 07: pick your printer's template, let preflight re-read the file, then export the printer pack. The 3D Forge adds tokens and a box.

A 3D-printed figurine

1. Pick a cleanly cut-out image and generate the model in the 3D tab of Game Assets (paid, price shown), or prepare its views for free first.
2. Open it in the Workbench: target size, repair, lay down, check overhangs, arrange on the bed.
3. Export › 3D printing, then Open in the slicer: see 3D-print your creations.

A campaign for a client

1. Create the client's space and verify its channels: Work for clients. Link their WhatsApp in the same place: Talk to your clients on WhatsApp.
2. Start from a post that inspires you: Scan a page, an inspiration, a diagram. The draft lands in the Scheduler on the project's channels.
3. Have the client approve it, then schedule it; the WhatsApp broadcast goes out with it.

A video tutorial of the app

1. Record your gesture, dress it with a skin and edit the recording: Record the app and go live.
2. Schedule the video's draft in the Scheduler, and export the walkthrough so someone else can replay it.

Assets for your game

1. Animate a character in Sprite Lab and export the sheet for Unity, Godot or Phaser.
2. Build a tileset in Tile Lab from two materials, paint a map, export to Tiled, LDtk or Godot.

REFERENCE

Read the icons

You will recognise an icon at a glance, find its name, and print the full sheet to keep next to your screen. Free, no key needed.

One icon, one function

The whole app speaks the same language: the **Deepotus Glyph** set, 526 icons drawn with the same stroke. Each one stands for **a single function**: the same drawing means the same thing in the editor, the Photolab or the Card Forge. The old emojis and symbols are gone from buttons, menus and lists.

1. Rest the pointer on a button: its tooltip gives its name and, if it has one, its keyboard shortcut.
2. Look that name up in the icon sheet to read what it does and what not to confuse it with.
3. In the Vectorlab, the Photolab, the Sprite Lab and the Tile Lab, wait a little longer (900 ms) or click the small `?` button: an animated card SHOWS you the option.

Icon families

Icons are sorted into families. Knowing the family already tells you what to expect when you click.

Navigation

A screen, a tab or a panel: you go there, nothing is changed.

Actions

A button that does something: create, save, send, delete.

Editing

Change something already there: align, flip, group, lock.

Tools

The drawing and retouching tools of the labs.

Layers and adjustments

Layers, masks, styles and image adjustments.

Media and playback

Play, cut and edit video and sound.

States

What the app tells you, without you having to click.

Categories

Content families, each with its own colour.

3D

Objects and actions of the 3D workshops.

States worth knowing

A state icon is not clicked: it informs you. These are the ones to read before starting a paid render.

Icon	What it says
	Done: the render, the upload or the save worked.
	In progress or waiting: loading, generating, or queued.
	Failed: the tooltip or the message next to it says why.
	Warning: nothing is broken, but read before going on.
	Spending limit close or reached: beyond it, the app refuses the spend and asks you what to do.
	Generative field: AI can fill it in for you.
	Locked: nothing can move, change or delete the item while it is.
	Pinned (Studio): this node result is reused for free on the next render.

Print the icon sheet

The icon sheet shows them all, sorted by family, with their name, their function and the look-alike icons not to confuse them with.

1. Open the icon sheet.
2. To print only part of it, click a family (for example **Vectorlab tools**) or type a word in the search: only what is shown gets printed.
3. Click **Print**: each family starts on a new A4 page, black on white. You can also download the full sheet as a PDF.

> On the sheet, the red line starting with "≠" names the look-alike icons: Close hides a window, Delete destroys, Remove takes something out of a list without destroying it.

Animated cards in the labs

In the labs, some options have an animated card in three beats made of real screenshots: the starting point, the gesture, the result. It opens when you rest the pointer on the option for 900 ms, or with the `?` button next to it (in the Sprite Lab and the Tile Lab, even when the option is greyed out). Here they all are; cards of labs not yet translated still show the French interface.

Vectorlab

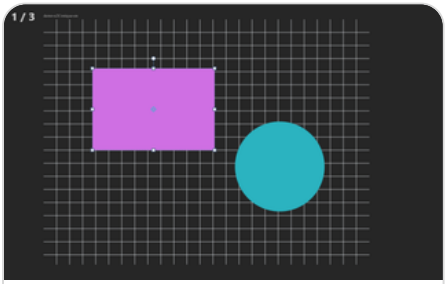

Conical
The color spins around the middle of your shape, like the spokes of a wheel.
`Appearance › Conical`

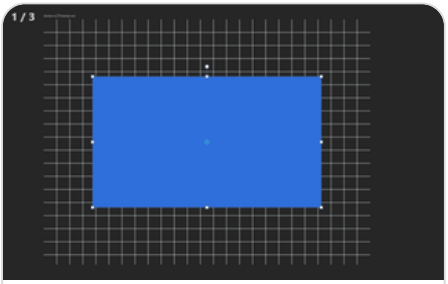

Transparency
Your shape slowly fades from one side to the other, as if it vanished into fog.
`Appearance › Transparency`

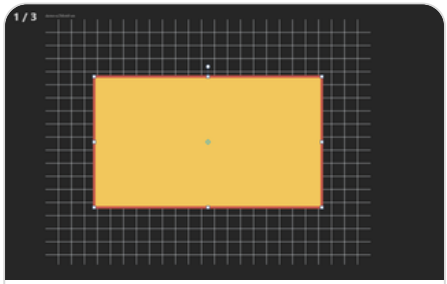

Pattern
Your shape fills up with little repeated stripes, like striped fabric.
`Appearance › Pattern`

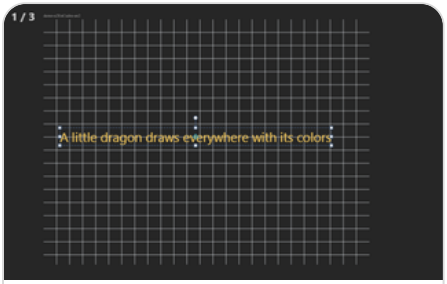

Frame
Your long line of text fits into a box and wraps by itself, like in a book.
`Appearance › Frame`

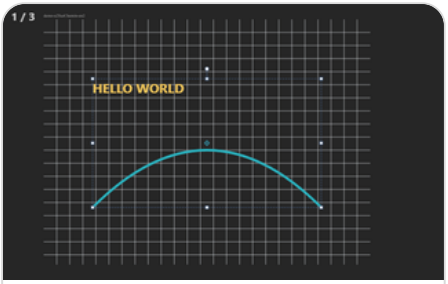

On path
Your words lie down on the line you drew and follow it wherever it goes.
`Appearance › On path`

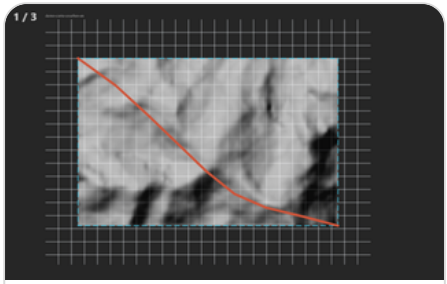

Contours
Lines join the places of the same height on your map, like the steps of a mountain.
`Real map › Contours`

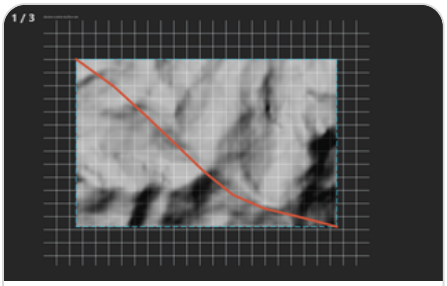

Slice

Your map turns into six-sided tiles, colored by height: water, meadow, forest, mountain.

`Real map › Slice`

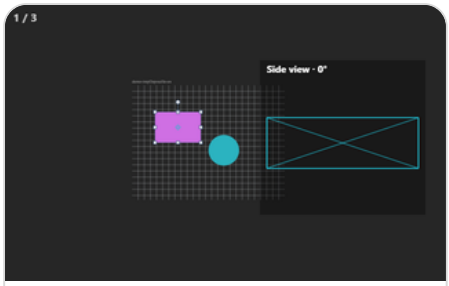

Side draft

The sides of your piece lean: the more you turn, the wider the bottom and the smaller the top.

`3D printing › Side draft`

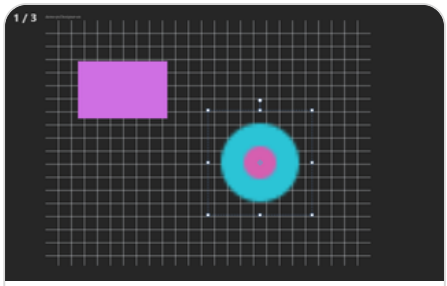

Set as reference

Your picture becomes a reference: it fades, stays still, and a grid of squares lies on it to guide you.

`Pixel persona › Set as reference`

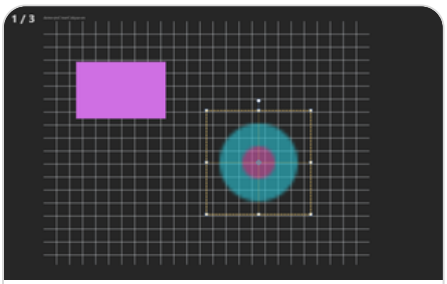

Create pixel layer

A see-through sheet lies on your reference picture: you draw on it, square by square.

`Pixel persona › Create pixel layer`

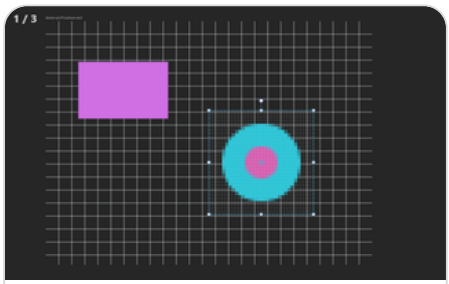

Pixelate image

Your picture is now made of big squares, like bricks: as many squares as the tile size.

`Pixel persona › Pixelate image`

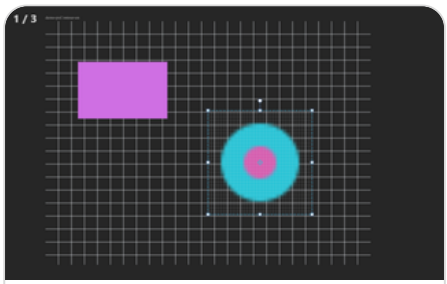

Dark outline

A dark line goes around your drawing so it stands out from the background.

`Pixel persona › Dark outline`

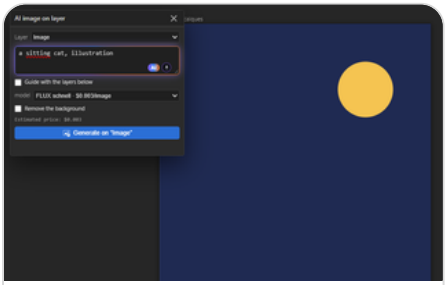

AI image on layer

You write what you want to see: the picture appears on the chosen layer, under your text.

`AI image on layer › Generate`

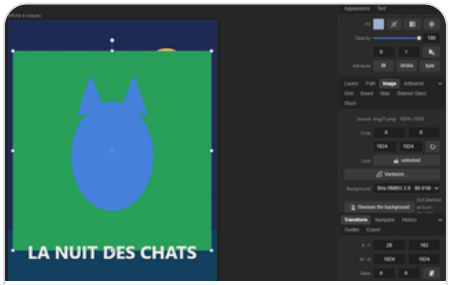

Remove the background

The background around your picture disappears: only the subject is left, sitting on your drawing.

`Image tab › Remove the background`

Photolab

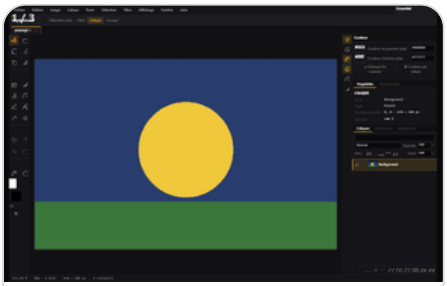

Workspace

Pick a workspace in the list: panels and tools rearrange for what you want to do.

`Workspace list`

Rectangular marquee (M)

Drag across the picture: a dotted rectangle surrounds the part you are going to change.

`Tools › Rectangular marquee (M)`

Magic wand (W)

Click a colour: the whole patch of that colour is selected at once.

`Tools › Magic wand (W)`

Crop (C)

Draw a frame on the picture then press Enter: everything outside the frame disappears.

`Tools › Crop (C)`

Brush (B)

Drag across the picture: a stroke of your colour is painted wherever you go.

`Tools › Brush (B)`

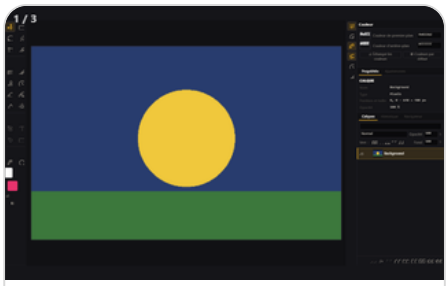

Clone stamp (S)

Alt-click what you want to copy, then paint elsewhere: the same piece of picture reappears under your brush.

`Tools › Clone stamp (S)`

Black and white

One click adds a layer that removes the colours: your picture turns grey, and you can remove it whenever you like.

`Adjustment layer › Black and white`

New AI layer...

You describe a picture: it arrives in a new layer, placed above your photo.

`Layers foot › New AI layer`

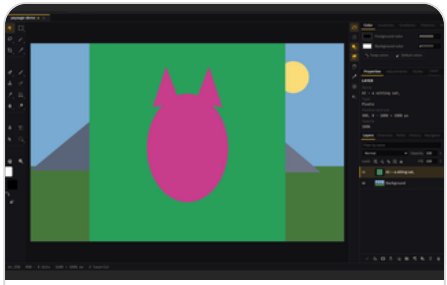

Remove layer background

Your layer's background fades away; the original stays hidden underneath, just in case.

`Layer › Remove layer background`

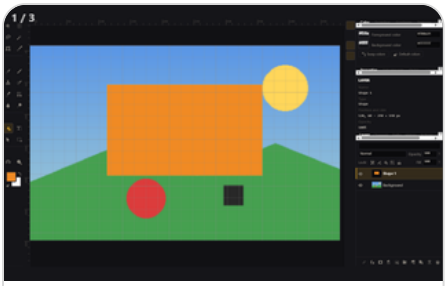

Add Anchor Point (P)

Click the edge of a shape: a new point appears there, so you can bend it at that spot.

`Tools › Add Anchor Point (P)`

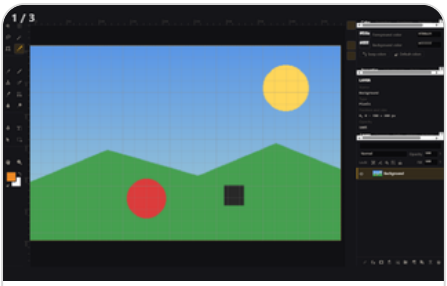

Color Sampler (I)

Click on colors: each spot keeps a number, and the Info panel tells you its color.

`Tools › Color Sampler (I)`

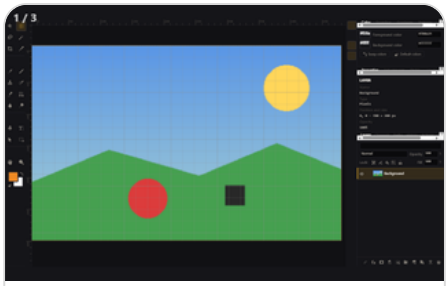

Single Column Marquee (M)

Click the picture: a single column, thin as a thread, is selected from top to bottom.

`Tools › Single Column Marquee (M)`

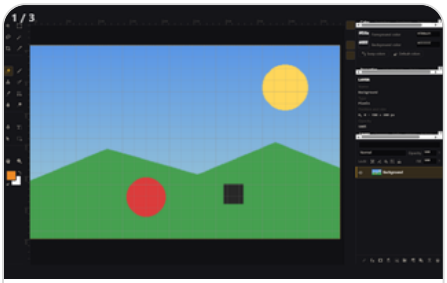

Content-Aware Move (J)

Circle an object and drag it: it moves, and the hole it leaves fills itself in.

`Tools › Content-Aware Move (J)`

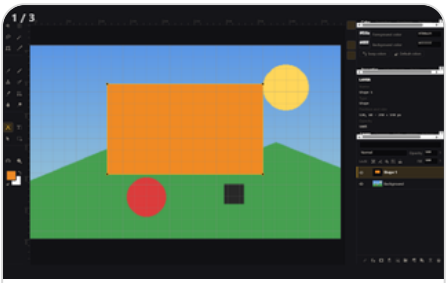

Convert Point (P)

Drag from a sharp corner: it turns round, and the edge curves.

`Tools › Convert Point (P)`

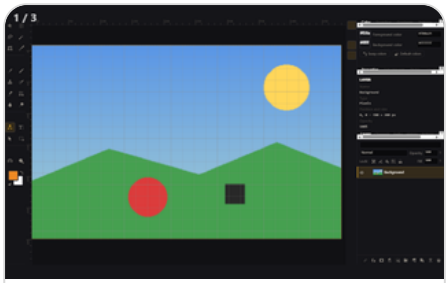

Curvature Pen (P)

Place a few points: the pen draws soft curves through each one by itself.

`Tools › Curvature Pen (P)`

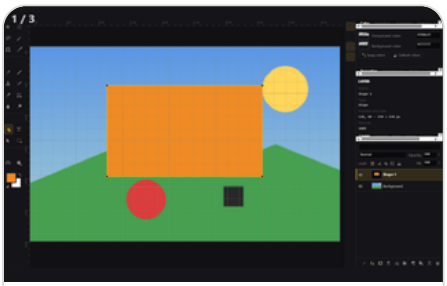

Delete Anchor Point (P)

Click a corner of a shape: that point goes away and the shape gets simpler.

`Tools › Delete Anchor Point (P)`

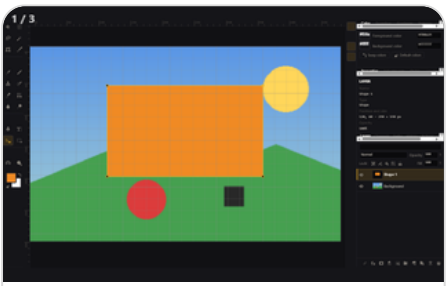

Direct Selection (A)

Grab one point of the shape and pull it: only that corner moves, the rest stays put.

`Tools › Direct Selection (A)`

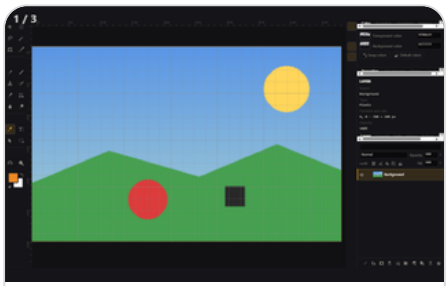

Freeform Pen (P)

Draw freehand like with a pencil: your line becomes a nice smooth shape.

`Tools › Freeform Pen (P)`

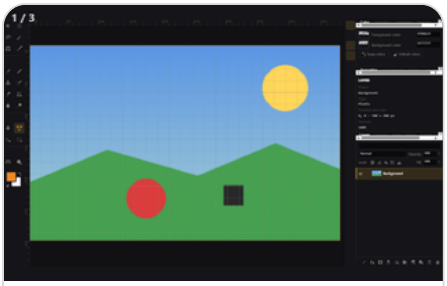

Horizontal Type Mask (T)

Type a word: its letters become a selection that you can then fill with color.

`Tools › Horizontal Type Mask (T)`

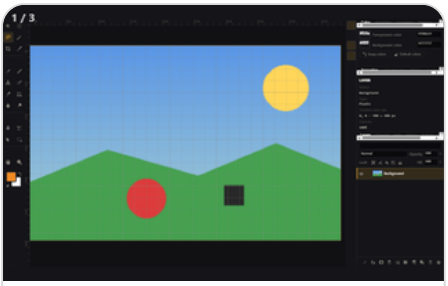

Magnetic Lasso (L)

Click around an object: the line sticks to its edge by itself, then Enter selects it.

`Tools › Magnetic Lasso (L)`

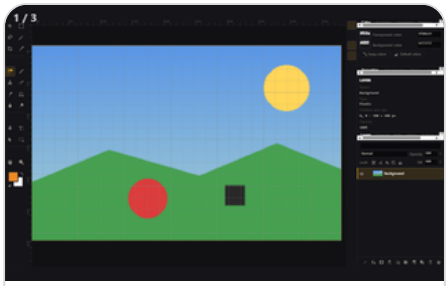

Patch (J)

Circle the spot, then drag your outline to a clean area: the spot takes on the look of that place.

`Tools › Patch (J)`

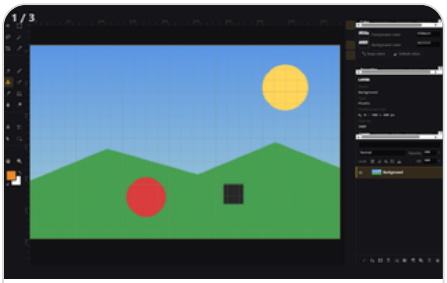

Pattern Stamp (S)

Drag across the picture: you paint with a repeating pattern, like an ink stamp.

`Tools › Pattern Stamp (S)`

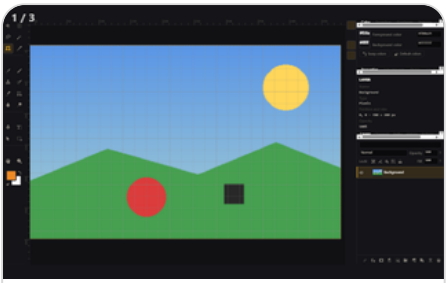

Perspective Crop (C)

Draw a frame, pull its corners onto the tilted shape, then press Enter: the picture is straightened and cut.

`Tools › Perspective Crop (C)`

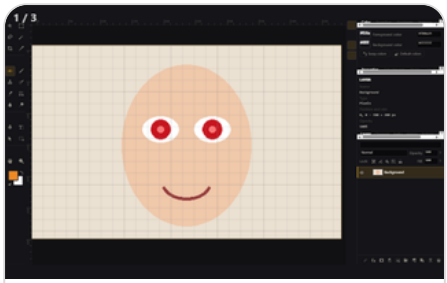

Red Eye (J)

Click a red eye: the pupil turns dark again, like in a real photo.

`Tools › Red Eye (J)`

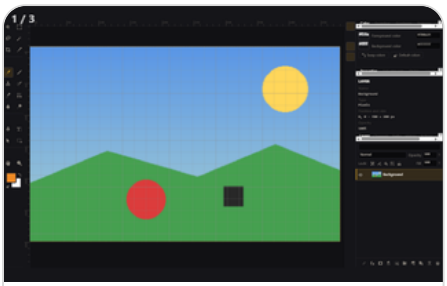

Remove (J)

Brush over what's in the way: the object vanishes and the surroundings take its place.

`Tools › Remove (J)`

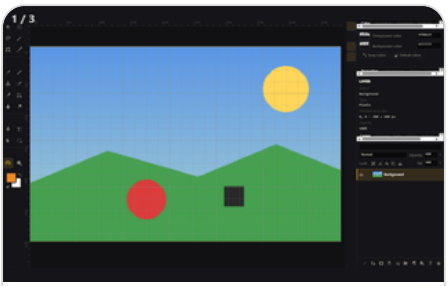

Rotate View (R)

Turn the canvas like a sheet of paper on a table: the picture doesn't change, only your view turns.

`Tools › Rotate View (R)`

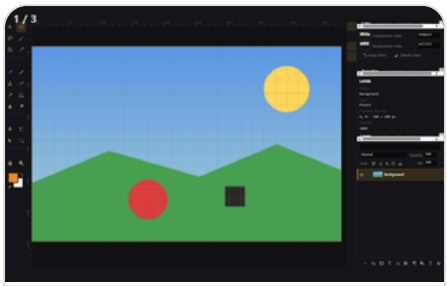

Single Row Marquee (M)

Click the picture: a single row, thin as a thread, is selected from edge to edge.

`Tools › Single Row Marquee (M)`

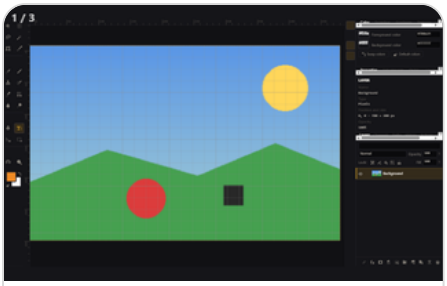

Vertical Type (T)

Click and type a word: the letters are written from top to bottom, one under another.

`Tools › Vertical Type (T)`

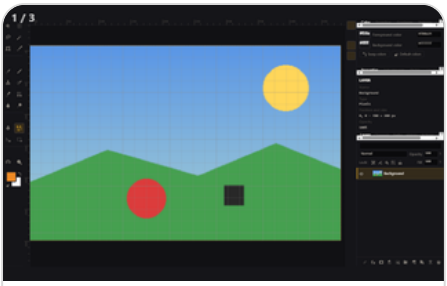

Vertical Type Mask (T)

Type a word from top to bottom: its letters become a selection, ready to be filled.

`Tools › Vertical Type Mask (T)`

Sprite Lab

Auto-detect

The tool counts the pictures on your sheet by itself and draws a box around each one.

`Sheet › Auto-detect`

Feet

All your characters put their feet at the same height, as if standing on the same ground.

`Sheet › Alignment › Feet`

Mirror

Your character turns around and looks the other way, like in a mirror.

`Preview › Mirror`

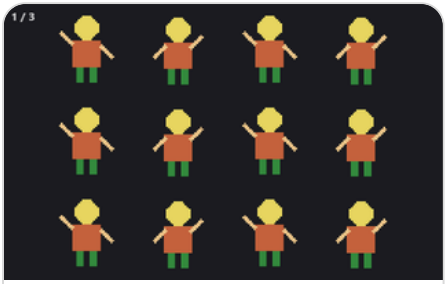

Pixel art (9b)

Your drawing becomes big squares and only a few colours, like in an old video game.

`Pixel art (9b) box`

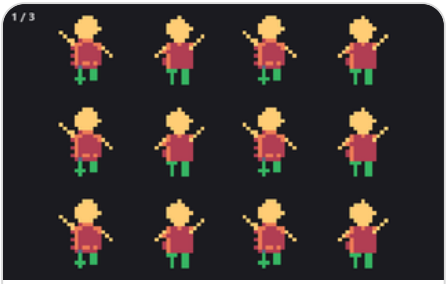

Outline

A line surrounds your character so it shows up on any background.

`Post-processing › Outline`

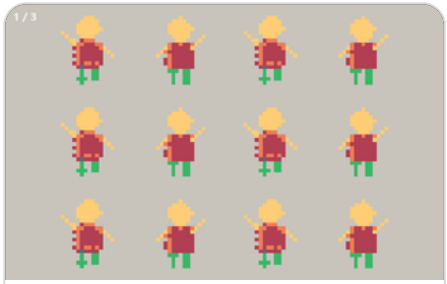

Shadow X / Y

A shadow appears behind your character, slightly offset, as if a lamp were lighting it.

`Post-processing › Shadow X / Y`

Tile Lab

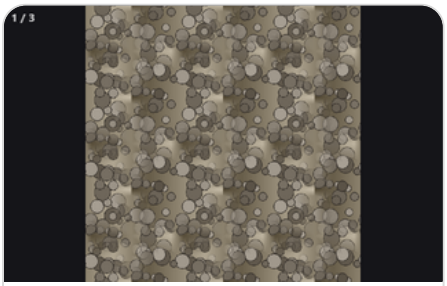

Method
Mirror folds your texture like a sheet folded in four: the edges always meet perfectly.
`Seamless mode › Method`

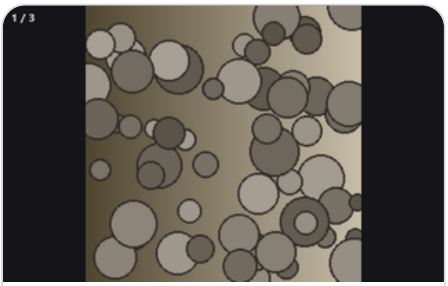

Make seamless
Your texture joins up: laid side by side many times, you can no longer see where each piece ends.
`Seamless mode › Make seamless`

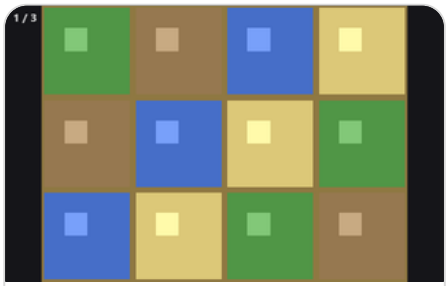

Placement grid
Your cells become diamonds, staggered like the paving of a game seen at an angle.
`Tile sheet › Placement grid`

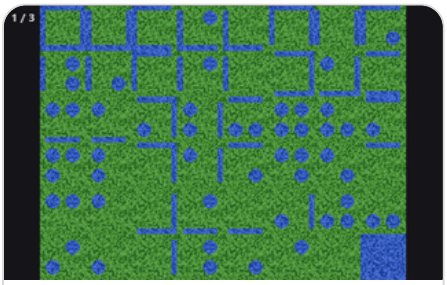

Tileset
You choose how many pieces to make: many for nicely rounded corners, or sixteen to go fast.
`Set mode › Set`

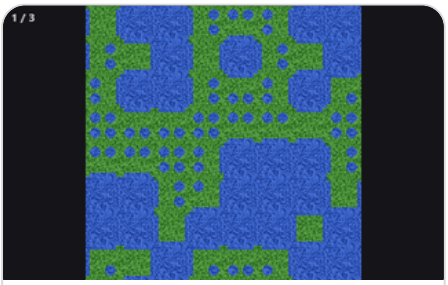

New preview
A new map draws itself at random with your tiles, to see if they go well together.
`Set mode › New preview`

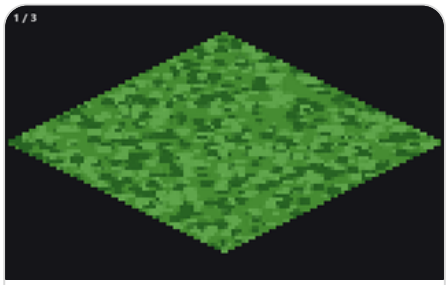

Shape
Your tile changes shape: a diamond for a game seen at an angle, or a six-sided hexagon.
`Shapes mode › Shape`

Good to know

The icon of the app, of its shortcuts and of browser tabs is the real Deepotus logo. Inside the app, the logo shown follows your active brand kit (chapter Templates and brand kits).

One single chevron is used everywhere to expand and collapse: pointing down, the section is open; pointing right, it is closed.

REFERENCE

Keyboard shortcuts

The most useful keyboard shortcuts, screen by screen. Each screen has its own; Photolab and Vectorlab also list them in their Help menu, and Photolab lets you change them.

Everywhere

Gesture	Shortcut
Quick command (go to a screen, an action)	`Ctrl` + `K`
Search this guide	`/`

Editor

Gesture	Shortcut
Cut the clip at the playhead	`Alt` + `C`
In and out range, cut the range	`I`, `U`, `Shift` + `X`
Add a marker, list them	`Shift` + `M`, `Ctrl` + `M`
Quick fade on the chosen clip	`Alt` + `T`
Undo, redo	`Ctrl` + `Z`, `Ctrl` + `Y`

Details: Edit a video.

Photolab

Gesture	Shortcut
New, open, save	`Ctrl`+`N`, `Ctrl`+`O`, `Ctrl`+`S`
Free transform	`Ctrl`+`T`
Levels, curves, hue	`Ctrl`+`L`, `Ctrl`+`M`, `Ctrl`+`U`
Search a command	`Ctrl`+`K`
Rulers, grid, guides	`Ctrl`+`R`, `Ctrl`+`'`, `Ctrl`+`;`
Change shortcuts	`Ctrl`+`Alt`+`Shift`+`K`

Details: Retouch a photo and Layers, masks and workspaces.

Vectorlab

Gesture	Shortcut
All shortcuts	`F1`
Fit page, 100%, grid	`Ctrl`+`0`, `Ctrl`+`1`, `G`
Duplicate, group, ungroup	`Ctrl`+`D`, `Ctrl`+`G`, `Ctrl`+`Shift`+`G`
Select, node, pen, text	`V`, `N`, `P`, `T`

Details: Draw vectors.

2D and 3D labs

Screen	Gesture	Shortcut
Sprite Lab	Play / pause, move the sprite	`Space`, `W` `A` `S` `D`
Workbench	Place on a face	`F`
3D Stage	Add a keyframe, play the move	`K`, `Space`

REFERENCE

Troubleshooting

You will be able to recognise the most common messages and know what to do. Always start with **Settings** › **Diagnostics**: version, disk space, latest alerts, key status (never their value) and provider balances, at a glance.

Symptoms and fixes

Symptom	What to do
A button says a key is missing, or a provider pill is red at the top right.	Add or fix the key in Settings, keys section (see Connect your keys). It works straight away, no restart.
A generation is refused: cap reached.	The window says which cap blocks and offers to adjust it. See Keep spending under control.
The HeyGen avatar list takes a while to appear the first time.	Normal on the very first load: the list is long. It is then kept on disk and shows instantly.
An ElevenLabs voice returns a plan error.	On ElevenLabs' free plan, only “premade” voices work through the API. Pick one, or move to a paid plan.
Certificate error (CERTIFICATE_VERIFY_FAILED).	An antivirus inspects HTTPS. The app uses Windows' certificates and handles it on its own; otherwise, turn off the antivirus's HTTPS inspection.
“Port 8765 already in use”.	The app is already running: open `http://deepotus.localhost:8765`.
After a restart, my keys or Library seem “gone”.	Nothing is lost: everything lives in `%LOCALAPPDATA%\DeepotusVideoGenData`. Start the app again from its Desktop icon.
Photolab returns to its welcome screen with an engine message.	The retouching engine was restarted; unsaved documents are lost. Save often (see Photolab).
A 3D file is refused at print export.	Go through the Workbench: set a target size and repair in one click before exporting.
I deleted an image by mistake.	It is in the Library Trash: Restore.

Where are my files?

In `%LOCALAPPDATA%\DeepotusVideoGenData\assets`: `images` for images, `outputs\final` for rendered videos, `outputs\uploads` for your uploads, `audio` for sounds, `print3d` for print exports. To take everything to another PC, see Move to another computer.

Good to know

> An app message always says what blocks and how to fix it. If a button is greyed out, hover it: its tooltip says why.