

Deepotus Video Gen guide

Everything you can do with the app, one task at a time: short steps, animated pictures taken from the real app, and the icons you will see on screen.



Take a tour of the app



Install, pick a language, update



Connect your keys and AI engines



Keep spending under control



Settings, personas and branding



Move to another computer



Drive it from your phone



Make a video in minutes (Quick)



Compose with the Studio



Reuse templates and brand kits



Tell a story in episodes



Plan and publish your week



Connect your accounts



Make game assets



Create playing cards (Card Forge)



From image to 3D object



Prepare a part before slicing (Workbench)



3D-print your creations



End-to-end recipes, from
idea to post



Read the icons



Troubleshooting




GET STARTED

Take a tour of the app

You will find your way around the app, move from one screen to another without searching, and read what the top bar tells you. Five minutes, free.

Deepotus Video Gen is a studio that runs **on your own computer**. It turns an idea into a video, a picture, a 3D object or a published post. You plug in your own AI provider keys: you pay those providers, not the app. Everything you create stays with you.

The screen at a glance




1. The rail: one screen per task 2. Ctrl+K: go anywhere from the keyboard 3. This month's spending and key status, always in sight
The screen rail on the left, the top bar with the quick command, the spending counter and the key status.

- 1 The rail, on the left:** one screen per task. **Quick** for a one-click video, **Studio** to compose, **Edit** to cut a video, **Library** to find your media, **Settings** for the keys. The arrow at the top of the rail folds it into icons only.
- 2 The quick command:** click **Quick command...** or press **Ctrl + K**, type a few letters, Enter.
- 3 The spending counter** shows what you spent this month with your providers. One click opens Settings.

4 The badges `fal`, `heygen` and `voice` turn green when the matching key is set. One click takes you to that key.

5 The **render queue**, top right, lists what is being computed and what just finished, in the `Running` and `Recent` tabs.

Go fast with the quick command



1. Type, then Enter

Ctrl+K opens the quick command: type a few letters to go to a screen or start an action.

The quick command goes to any screen (`Go to Studio`, `Go to News` ...), opens the last render, creates a starter graph or opens a Settings section. `Replay onboarding` restarts the first-launch tour.

The screens of the rail



Quick

A one-click video: animated clip, talking avatar, composition, voiceover.



Studio

A node graph to compose a video piece by piece.



Chapters

From a text or a novel to a narrated video, scene by scene.



Sound & VFX

Music, voices, sound effects and visual effects.



Edit

A multitrack timeline to cut, grade and deliver.



Scheduler

Plan and publish your posts for the week.



Templates

Reusable layouts and your brand kits.



News

Follow the news and turn it into short videos.



Library

All your pictures, renders, sounds and objects, with their origin.



Game Assets

3D, sprites, tiles, materials and playing cards.



Vectorlab

Vector drawing, pixel art, maps and printable objects.



Photolab

Photo retouching with layers, masks and filters.

The guide opens from the foot of the rail: [FR](#) or [EN](#) .

What the app does to protect you



Nothing leaves without you

Only this PC can change, generate or spend. Another device on the network can only read, and only once paired.



Errors that make sense

A missing key or a server refusal is shown in plain words, with what to do.



Safe imports

A picture with the same name becomes name-1, name-2...; a file that is not a real picture is refused and named.



In-app dialogs

Confirmations and text prompts use one Deepotus dialog, never a browser pop-up.

Good to know



The desktop icon opens the app at `http://deepotus.localhost:8765`. Each browser tab carries the name of the open tool, for example “Deepotus — Vectorlab”.



Only AI generations cost money, and your provider bills them. The chapter [Keep spending under control](#) shows how to set limits.

Shortcuts

Key	Effect
<code>Ctrl</code> + <code>K</code>	Opens the quick command
<code>Esc</code>	Closes the quick command or a dialog



GET STARTED

Install, pick a language, update

You will launch the app for the first time, pick its language, know what each use costs and keep it up to date. Ten minutes, free.

First launch


- 1 During installation, pick the language: **French** or **English**. It becomes the language of the whole app.
- 2 Double-click the **Deepotus Video Gen** desktop icon. The app starts without a black window, then your browser opens on the studio.
- 3 On the very first launch, a five-step welcome introduces the studio and lets you pick a persona, your generation providers and the networks you publish to.
- 4 Set at least one key, the **fal.ai** one: it makes pictures and video clips. The chapter [Connect your keys](#) shows how.



The welcome can be replayed at any time from the quick command (**Ctrl** + **K**):

Replay onboarding .

Switch language



1. French or English, for the whole app

Settings › Provider defaults › Interface language: the page reloads in the chosen language.

- 1 Open **Settings** › **Provider defaults**.
- 2 In **Interface language**, pick **Français** or **English**. The page reloads in the chosen language.

Translated: the sidebar, Settings, the Library, Quick, News, the Scheduler, the Studio, Templates, Episodes, the editor, Sound & VFX, subtitles, machine transfer and every confirmation dialog. Lists sent by the server (effects, transitions, presets) also arrive in the chosen language. The labs (Vectorlab, Photolab, Card Forge...) are being translated one by one; meanwhile their known texts are translated on display. What you type yourself is never translated.


What each use costs

The app is free. Only AI generations cost money, billed by the provider of the key. Rough figures at default rates:

Action	Provider	Indicative cost
Create a picture	fal.ai (FLUX)	≈ \$0.003
10 s video clip	fal.ai (Seedance)	≈ \$0.40
Talking avatar, 1 min	HeyGen	≈ \$0.24
Voiceover	ElevenLabs	per character, free tier available
News summary, post plan	Anthropic, OpenAI, Gemini or Ollama	a few cents, free with Ollama
Import your videos, sounds, pictures	none	free

Your real rates go into [Settings](#) > [Pricing & budget](#) ; see [Keep spending under control](#).

Keep the app up to date



1. Your version is up to date 2. Check asks for the latest published version

The version block in Diagnostics: the installed version, its status, and a button to check right now.

- 1 Once a day, the app checks whether a new version is published. It never blocks anything.
- 2 When there is one, a card announces it: **Download** fetches the installer and shows where it is saved, **Later** hides the card for that version, **Notes** tells what changes.
- 3 To check right now: **Settings** > **Diagnostics**, **Check** button.
- 4 Close the app, run the downloaded installer: nothing is installed on its own.

Good to know

i Your keys, files and database live in the `%LOCALAPPDATA%\DeepotusVideoGenData` folder, apart from the program: they survive updates and reinstalls.



Start with the fal.ai key alone, then add avatars, voices and AI as you go. Set a monthly cap from day one: see [Keep spending under control.](#)



GET STARTED


Connect your keys and AI engines

You will set an AI provider's key, choose which engine does what, and protect your keys with a password. Ten minutes, free; keys are created with each provider.

Before you start

An API key is a code the provider gives you on its website, in your account. It lets the app work on your behalf; the provider bills you. Only one is required to start: **fal.ai**.

Set a key



1. The key's status 2. Guide: where to create the key, and its price 3. Paste the key, Save: it works right away

Settings > API keys: each provider has its row, its status, its guide and its Save button.

1

Open **Settings** > **API keys**. Each provider has its row and its status: **missing** until a key is set.

2

Click **Guide**: it tells where to create the key with the provider and how it is billed.

- 3 Paste the key in the field, then **Save**. It works **right away**, no restart.
- 4 Click **Test** on the row: the provider says whether it accepts the key, without spending anything.

To set several keys at once, fill in the fields then **Save all changes**.

Which key unlocks what

Key	What it unlocks
fal.ai	Pictures (FLUX) and video clips (Seedance). The only key you need.
HeyGen	Talking avatars, alone or in a composition.
ElevenLabs	Voiceover and narration.
Anthropic, OpenAI or Gemini	News summaries, post plans, scripts and storyboards written by AI. One is enough; OpenAI also makes pictures (GPT Image).
Meshy	Premium picture-to-3D (optional, Meshy subscription).
Figma	Import Figma layers into the Library and Templates.
Ollama (no key)	An AI model installed on your PC for post plans: free, nothing leaves your machine.

Network keys (X, Telegram, YouTube, Instagram, TikTok) go into **Connected accounts**: see [Connect your accounts](#).

Choose the engine for each job

In **Settings** > **Provider defaults**, each role has its engine: picture generation, voice, picture to video, avatar, voiceover, news summary, post planning. Only keys actually set are offered.



Default video model

The model new Seedance nodes start with in the Studio. Each node can change it.



Default voice model

The ElevenLabs model of new voiceover nodes.



Local AI with Ollama


Install Ollama, download a model, then enter its name in the **Local LLM (Ollama)** card of API keys.



Gemini: pick the version

Left empty, the “Gemini — model” field always follows the latest Flash version; fill it in to pin one.

Protect your keys in the vault



1. A master password encrypts all your keys 2. The .dzk archive carries them, with its own password

Settings > Vault: your keys encrypted with a password, and the archive to carry them to another computer.

1 Open **Settings** > **Vault** .

2 Choose a master password (at least 8 characters), type it twice, then **Set up vault** . Your keys are encrypted and leave the plain-text file.

3

Close removes the keys from memory: nothing paid can go out. **Open** makes them active again with the password.

4

Remember on this PC opens the vault on its own at launch, for your Windows session;

Stop remembering asks for the password again.

On each row of **API keys**, a badge tells where the key lives: in the file, in the vault, or vault closed (one click takes you there).

Good to know



The vault password is stored nowhere: if lost, it cannot be recovered. Write it down, or keep an encrypted archive.



Export archive creates a **.dzk** file that carries your keys, caps and rates to another PC or your phone, protected by its own password. **Import an archive...** takes them back.



Settings > **Diagnostics** shows the status of all keys at once (never their value) and the providers' balances.



GET STARTED



Keep spending under control

You will know what you spend, set a cap per month and per engine, and decide yourself when a render may go over it. Five minutes, free.

Where to see what you spend

The dollar counter in the top bar shows the month's total. One click opens **Settings** > **Pricing & budget**, where everything is set.

Set your rates and caps



1. A cap for each video generation

2. Above a cap, the app asks you before spending

Settings > Pricing & budget: your rates, then one cap per engine and the bar of what you spent this month.

- 1 At the top of **Pricing & budget**, enter your providers' rates (price of a picture, of a second of video, of a voice character...). The costs shown in Quick, the Studio and the Scheduler become accurate. **Save pricing**.
- 2 **Video cap per request** limits the cost of a single video generation; **Max generated seconds per clip** limits its length.
- 3 Further down, **Monthly caps**: one overall cap and one per engine (fal, OpenAI, Anthropic, ElevenLabs, HeyGen, Meshy), in dollars. A bar shows what was spent.
- 4 Set **Alert at (% of a cap)**: a banner warns you once a month when a cap reaches that level. **Save caps**.

When a render would go over a cap

Before leaving, every paid generation is checked against your caps. If it goes over, a “Monthly spending cap” dialog opens, wherever you are in the app:

- 1 “Tirer quand même” (run anyway) sends exactly the same request, once.
- 2 “Annuler” (cancel) sends nothing and tells you why.

Actual versus estimated

Below the caps, the **Monthly spend — actual vs estimated** table gives one row per engine and per screen: the number of runs, what was estimated, what was actually billed, and the difference.



Actual

The provider gave the exact price (Meshy, HeyGen): that is what counts.



Partly actual

Some runs have their exact price, the rest is still estimated.



Estimated

The provider gives no price: the app counts the quote made at launch.

Good to know



Spending made from your phone goes into the same record and the same caps, counted once.



The Studio sends the cost it displays as the request's cap: a render cannot cost more than what you saw before clicking.




GET STARTED

Settings, personas and branding

You will find any setting, give your content a voice with personas and dress the app in your colours with brand kits. Ten minutes, free.

Find a setting



1. Type what you are looking for
 2. The matching sections show up
 3. Enter opens the first one
- Settings search finds a setting by its name and opens the right section directly.

- 1 Open **Settings** in the rail.
- 2 Type a word in **Search settings...**: accents and capitals do not matter.
- 3 Enter opens the first result, Esc clears the field.

The Settings sections


Section	What you do there
Diagnostics	Version, disk space used, latest alerts, key status and provider balances.
Vault	Encrypt your keys and carry them (chapter Connect your keys).
Devices	Pair a phone (chapter Drive it from your phone).
API keys	Set and test the AI providers' keys.
Connected accounts	Link X, Telegram, YouTube, Instagram and TikTok (chapter Connect your accounts).
Personas	The tone of your content.
Branding	Name, logo, colours: your brand kits.
Caption pack	One-click brand labels under the Scheduler's caption editor.
Provider defaults	The engine of each role, default models, language.
Paths	Where your pictures and renders are stored (button to copy the path).
News	Settings of the News screen.
Appearance	Reduced motion and visual effects.
Pricing & budget	Your rates and caps (chapter Keep spending under control).
Machine transfer	Take your work to another PC (chapter Move to another computer).

Give a voice with personas

A persona is a personality: its tone, its audience, its way of speaking. The active persona sets the News scriptwriter, the tone of the prompt generator and the default voiceover.

- 1 Open [Settings](#) > [Personas](#) .
- 2 Built-in personas can be read but not edited: duplicate one, or click [New persona](#) .
- 3 Click a persona to make it **active**.

Dress the app with a brand kit



1. Create a kit 2. The active kit dresses the whole app

Settings > Branding: several brand kits, a single active one, and its fields right below.

- 1 Open **Settings** > **Branding**. The kit list shows which one is **active**.
- 2 **New kit** creates one; **Activate** makes it the active kit and refreshes the app.
- 3 The fields below edit the active kit: **Logo...**, app name, sub-label, two tagline lines, primary and accent colour. **Save brand**.
- 4 **Rename**, **Duplicate** and **Delete** manage each kit. The active kit and the last kit cannot be deleted: the greyed-out button says why.

The active kit dresses the splash screen, the sidebar, the app's accents and Template renders that use the brand. **Reset to deepotus** restores the original look.

Good to know

i A Template render keeps the kit that was active when you sent it, even if you switch kits while it computes.

i **Settings** > **Appearance** > **Reduced motion** turns off decorative animations; it also follows the Windows setting.



GET STARTED


Move to another computer

You will take all your work to another computer, check that nothing got damaged on the way, and get your keys back there. Fifteen minutes, free; the app must be installed on both machines.


What travels

The transfer takes everything the app created: the Library and the origin of every file, renders, the bible, communication plans, vector documents, card series and 3D models. Your API keys **never** go into the package.

Export from the old PC



1. What goes, and how much it weighs



2. Extras to tick: logs and discards 3. A dated folder is written at the destination

Settings › Machine transfer: a dated package, each file with its checksum; keys stay on this PC.

- 1 Open **Settings** › **Machine transfer**: the number of files and the size to take are shown.
- 2 Click **Export...**, then pick the destination: a drive, a USB stick, or a path.
- 3 Tick the extras if you want them: **Logs** and **Discards — dated trash folders**. They are unchecked by default, with their size.
- 4 **Start export**: a dated folder is written at the destination, each file with its checksum. Do not close the app while it writes.
- 5 Click **Check integrity**: every written file is read back and compared with its checksum.

Import on the new PC

- 1 Install the app, then open **Settings** > **Machine transfer** and click **Import...**.
- 2 Give the package folder, then **Check**: the app reads the package without writing anything and tells where it comes from (date, machine, version).
- 3 **Start import**. Import **adds**: nothing this PC already has is overwritten. A damaged file is left out and named.

Take your keys too

- 1 On the old PC: **Settings** > **Vault** > **Export archive**, with a separate password.
- 2 On the new PC: **Import an archive...** and that password. Your keys, caps and rates come back.

Good to know

i Paths from the old machine are rewritten for the new one: your projects find their files without any setting.

! Stopping an export leaves an incomplete package. If it happens, run the export again, then **Check integrity** before importing.



GET STARTED


Drive it from your phone

You will link your phone to the PC, carry your posts and chapters in your pocket, and cut off a lost phone. Fifteen minutes; you need the Deepotus companion app on the phone and the same Wi-Fi.

Before you start

By default, the app only listens on this PC: no other device can reach it. To open it to your home Wi-Fi, replace `HOST=127.0.0.1` with `HOST=0.0.0.0` in the `.env` file of the data folder, then restart the app. Windows will ask to allow Python through the firewall.

Pair the phone



1. By default, only this PC can reach the app 2. Shows a QR code valid for 5 minutes

Settings > Devices: the network status, the list of devices, and pairing by QR code.

- 1 Open **Settings** > **Devices** . The words **this PC only** or **local network open** tell whether the phone will be able to reach you.
- 2 Click **Pair a device** : a QR code appears, valid for five minutes, for a single device.
- 3 In the phone app, open **Pair** and scan the QR code. It closes on its own once read.
- 4 To give your keys to the phone, export the encrypted archive from the **Vault** and open it in the phone app, with its password.

Five devices at most. A paired phone **reads** what the PC produces; it can neither write to the PC, nor spend through the PC, nor see your keys.

What the phone does



Sync

The phone fetches the Library and its pictures, and resumes an interrupted download.



Send pictures to the PC

Your photos land in the Library with the "mobile" source, without duplicates.



The week's batch

Take the approved posts and publish them on Telegram or X from the phone. A post taken by the phone is no longer published by the PC.



Inbox

What you share to Deepotus lands here; pictures go on to the PC.



Offline chapters

Take a chapter, write without a network, hand it back to the PC.



Generate pictures

The cost is shown before each generation, at your PC's rates, and counted in your caps.

A chapter taken to the phone

While a chapter is on the phone, it is read-only on the PC, with a banner. You can force it back to the PC. If the two versions disagree, the phone's text stays in the log, to copy or take back: nothing is lost.

Remove a device

- 1 In **Settings** > **Devices**, click **Revoke** on the phone's row and confirm.
- 2 Access is cut right away; the posts and chapters it had taken come back to the PC.

Good to know

 Lost your phone? Revoking is not enough: your keys travelled in the archive. Regenerate them with each provider; the Devices page gives the link to each console.

 The phone is notified when a render finishes or fails, when a post is published or awaits an action, and when a cap is near.




MAKE VIDEOS

Make a video in minutes (Quick)

Create your first images

Library tab: your library of images, video renders, sounds and captions. Everything lands here.




Library: create an image (FLUX), the **Prompt manager**, **Upload image** and **Upload video** (your own clip). Hover a render to play it.

- 1 "Describe an image to create" field → describe in one line, Enter. The FLUX image (~3 s) lands straight in the library.
- 2 Out of inspiration? Click **Prompt manager**: pick ingredients (subject, mood, light, motion...), "Random", or "Refine with AI", then *Use this prompt* — it **overwrites** the field for a more personal result.
- 3 Already have visuals? **Upload image** (or drag-and-drop anywhere). For your videos, see chapter 5; for sounds, chapter 6.

i Prompt examples: "neon-red octopus mascot, abyssal background, cinematic" · "stylized market chart, bioluminescent mood" · "underwater throne, light rays, 9:16".

Video type ① — the cinematic clip (Seedance)

Quick tab, **Seedance** mode: one image + a prompt = a vertical cinematic clip.




Quick / Seedance: sources on the left, real-format preview on the right, estimated cost before you launch.

- 1 Pick the start image. Optional: an end image for a transition.
- 2 Write the motion prompt, grab a *template*, or open the **prompt generator** (it overwrites the default prompt). Set style, duration (5–60 s), format (9:16 default).
- 3 *Generate* → follow progress in the **Job Dock** (bottom bar). 30–45 s later the render is in Library.

i **Rename a render:** in the Job Dock, click the pencil  next to the name (or play it first with ▶). The new name instantly reflects in the Library and pickers.

Video type ② — the talking avatar (HeyGen)

Still in **Quick**, **HeyGen** mode: a script read by an avatar.




Quick / HeyGen: avatar + voice selection, upload your photo for a custom avatar, script.

- 1 Pick an avatar and a voice from the lists (search field available).
- 2 **Create your own avatar:** *Upload photo* → your photo (PNG/JPG) becomes a talking avatar in ~30 s, auto-selected.
- 3 Write the script (up to 4,900 chars), *Generate*. Allow 1 to 3 minutes.

🕒 The first avatar load — why it takes ~1 minute

HeyGen gives you a **very large catalogue** — often **1,000+ avatars** and 2,000 voices. The very first time the app requests this list, HeyGen's server can take **up to a minute** to respond. **This is HeyGen's own service being slow** — not an app bug, not a key problem, and not a network or antivirus issue.

So you don't have to wait, the app **pre-loads the list in the background at startup**: by the time you reach the HeyGen tab, the avatars are usually **already there**. Once fetched, the list is **cached** — every later open is **instant**.



Left: the HeyGen tab once loaded (1,281 avatars, search field, avatar & voice selectors). Right: why the first load takes ~1 min — pre-loaded at startup, then instant.


i **If you open the tab right after launch:** the list may briefly show "loading..." (up to ~1 min, **once**). Let it fill, then everything is instant — including after you create your own avatar (the list refreshes itself). The **HeyGen status** pill (top-right and in Settings) stays **instant** at all times: it checks your key and credits without loading the whole catalogue.

Composition mode combines both: a Seedance clip + a HeyGen avatar in one video (sequential or split-screen). **Anti-cut** is automatic: the avatar always finishes its sentence.

Video type ③ — your own video (UGC)

Many creators film themselves on a phone — a selfie, a reaction, an unboxing. Deepotus welcomes these raw clips (**UGC**, user-generated content) and marries them to your animations.

- 1 **Simplest:** Library → **Upload video**. Your clip is imported, its real duration is measured, and it appears in the *Renderers* tab like any render — reusable everywhere (Scheduler, Studio).
- 2 **In the Studio:** drag a **UGC video** node onto the canvas, click it, then *Upload your video*. A thumbnail fills the node.



Studio: the "UGC + animation" starter. On the right, the **UGC video** node with the **Duration master** switch.

The "master" duration — the key idea

When you turn on **Duration master** on your clip, **its real length** drives the final video. Generated animations (Seedance) are then calibrated around your clip — no more guessing "how many seconds?". You talk for 12 seconds to camera → the composition is 12 seconds, your illustration animation follows. It's the equivalent, for your own footage, of the avatar anti-cut.

Quick recipe: Studio → "UGC + animation" starter → upload your clip in the UGC video node (master on) → connect an Image + Seedance as the animated half → *Run*. Output: your video at the bottom, the animation on top, all at your clip's length.

Audio, voiceover & music

Sound makes a video. Deepotus now has a full **audio side**: import your own sounds, turn a script into a **voiceover**, and lay a **background music** track under any render.

The audio Library

The **Library** has an **Audio** tab next to Images and Renders. It gathers every sound: imported files, generated voiceovers and music tracks.

1 **Import a sound:** the *Upload sound* button (or drag-and-drop onto the Audio tab) → pick an `.mp3` / `.wav` / `.m4a` (and more). Its duration is measured; it appears in the Audio tab, playable inline, reusable everywhere.

2 The Audio tab also lists every **voiceover you generate**, so your narrations and your imported tracks live side by side.

i The Audio tab links a few **royalty-free music banks** (Pixabay Music, Freesound, Free Music Archive, YouTube Audio Library). Only add music you have the right to use — Deepotus cannot ship licensed music for you.

Voiceover (ElevenLabs)

Turn text into a natural narration with your **ElevenLabs** key.

1 In **Quick**, open the **Voice Over** tab: a voiceover **on its own** (no image, no video), perfect to narrate over your own visuals. Write a short script (up to 2,500 characters — for longer narrations use Chapters), pick the **language**; the counter and the **estimated cost** update live.

2 Pick a voice from the **scrollable list** (▶ to preview, filter field) or keep the **app default voice**, and save your voice + language combos as reusable **castings** (★ Save). *Generate voice* → the mp3 plays inline and is **saved to Library** → **Audio**, ready to reuse (Studio, Episodes, Scheduler).

⚠ **ElevenLabs free plan:** the built-in "premade" voices (George, Sarah...) work via the API. Custom "library / community" voices require a paid ElevenLabs plan — the app tells you clearly when a chosen voice is paid-only.

Background music (BGM)

Drop a music bed under your video — the app mixes it for you.

1 **In the Studio:** add a **Music track** node, pick a track from the audio Library (with an inline preview), set its **volume** and toggle **loop to duration**. Wire it into the Render node.

2 **On the simple render too:** you can attach a background track before generating, no Studio needed.

3 The track is automatically **looped or trimmed to the exact length** of the final video, then mixed under the voice/clip. *Run* — your render comes back with music.




MAKE VIDEOS

Compose with the Studio

Compose with the Studio (node editor)

The **Studio** is the advanced workshop: you wire *nodes* (sources → generators → montage → render) like a visual recipe. Don't worry — always start from a ready-made starter graph.



Studio: the node palette on the left, the canvas in the middle, the inspector on the right. Top bar: **New**, **Preview**, **Run**.

The basic gestures

- 1 **Start from a template:** in the palette, *Starter graphs* section, click "Seedance solo", "Avatar post", "UGC + animation", "News reel" or "Timeline montage".
- 2 **Start from scratch:** **New** button (top bar) → a blank canvas with just the *Render* output node. The right panel shows the **graph name**, editable.
- 3 **Add a node — drag & drop:** grab any entry from the palette (Image, Seedance, UGC video, Music track, Text...) and **drop it on the canvas**. The node appears where you dropped it.

- 4 **Wire:** drag from a node's output port (right side) to another's input port (left side). **Move:** drag the node body. **Delete:** the **x** in the node header (it also removes its connections).
- 5 **Check then run:** *Preview* shows inputs/the render; *Run* compiles the graph and launches the real render (follow it in the Job Dock).

i When you pick an asset in a node (a Library image, an existing render, your UGC clip, a music track), a **thumbnail fills the node:** you see at a glance what the chain will produce.

Studio — Effects, masks & preview

Two additions to the node editor (Studio).

The “Effects / Mask” node

Drag the **Effects / Mask** node (*Composition* section) and wire its output to the new **fx** port on the **Render** node. The panel shows a **thumbnail grid** so you pick an effect at a glance.

Available effects: **LUT / Grade** (named grades — Teal&Orange, Cyberpunk, Deep-sea, Noir, Vintage...), **VHS** (sequenceable line displacement, intensity + speed), **Colorize** (Sepia, Duotone, Matrix, Red-alert...), parametric **Gradients** (2 colors, angle, opacity, blend) — plus modern classics: grain, vignette, chromatic aberration, glitch, bloom, halation, scanlines/CRT, cinematic letterbox, old film, sharpen, blur, dreamy, pixelate, camera shake, mirror, invert.

Under **“Apply to”**, choose *Whole render* (global post-pass) or **a specific node** (the targeted layer). Add several Effects nodes to stack effects or target different layers.

- 1 + Effect → click a thumbnail → set intensity (and preset / speed depending on the effect).
- 2 Wire the output to the Render **fx** port, choose the target, then *Preview* or *Run*.

“Preview” button — see it before you pay

The **Preview** button (top bar) produces a **free, fast** render of the final composition: not-yet-generated **Seedance/HeyGen** clips are replaced by their **source still**. You see the framing, layers, **overlays and effects** exactly as in the final render — *before* the paid **Run**.


⚠ Text overlays (Text overlay / Ticker / Separator) now apply on **every** composition at the Render node (UGC, montage, spatial compose), not only spatial compose.



MAKE VIDEOS

Reuse templates and brand kits

Templates tab: ready-made 9:16 layouts (news reel + avatar + brand strip, split-screen, PIP corner, triptych...). Each thumbnail reflects its **real region layout** — they're all different.



Templates: gallery of layouts (distinct thumbnails), spatial editor on the right (drag regions or their handles).

- 1 Click a thumbnail: the right editor shows its regions. Move them, resize them.
- 2 *Save layout* stores your version — it **appears immediately** in the gallery. *New template* creates a blank one.
- 3 *Open in Studio* switches to the node editor for deeper editing.



Tell a story in episodes

You have a story to tell — a chapter, a lore drop, a serialised novel. The **Episodes** page turns a block of text into a **narrated, illustrated vertical video**, scene by scene, ready for the Scheduler. Open it from the sidebar (**Episodes**). A clickable stepper walks you through four steps.

Step 1 — the chapter & its narration

- 1 Give the episode a **title**, then **paste your text** or *Upload* a `.txt`, `.docx` or `.pdf` — the text is extracted for you.
- 2 Pick a **voice** and a **language** (preview a voice with ►).
- 3 *Generate narration* → the whole chapter is read by ElevenLabs and saved to the audio Library. Long chapters are split and stitched back together automatically.

Step 2 — the storyboard (scenes)

A video needs shots. The app cuts your chapter into **scenes**, each one a bit of text + its own illustration.

- 1 Choose the split: **By paragraph** (verbatim, instant) or **By AI** (the AI groups the text into visual beats and writes a rich illustration prompt for each). You can **generate both and compare** before locking one in.
- 2 **Edit freely**: tweak each scene's text and its illustration prompt, reorder them (↑ ↓), add or remove scenes (+ / ×).

Step 3 — illustrations & motion

- 1 *Generate all illustrations* (or one at a time): one image per scene, with your chosen image model (FLUX or GPT Image).
- 2 Per scene, pick the **motion**: **Ken Burns** (a slow cinematic zoom/pan, default), **Seedance** (a true animated clip) or **Still** (a fixed image).

Step 4 — assemble & send

- 1 *Assemble the episode*: each scene's illustration is animated over its slice of narration, then all scenes are concatenated into one vertical video — follow it in the Job Dock.
- 2 Preview the result, then *Send to Scheduler* → a draft post on **YouTube** and **Instagram**, ready to caption and schedule.


i Serialising a novel? Keep the same persona and voice across the series for consistency, then publish one episode per day from the Scheduler.



PUBLISH

Plan and publish your week


Scheduler tab: the publishing calendar, and the heart of "human-in-the-loop" — **nothing goes out without your approval** (unless you decide otherwise).



Scheduler: week/month, AI plan, document import, the post's graph + inspector on the right.


Three ways to fill the calendar

- 1 **Generate plan:** describe your week in one sentence → the AI proposes a full plan (dated posts, hooks, captions, formats). *Review the preview* before "Add to calendar".



"Generate plan": a one-sentence brief becomes a full plan, reviewed by you (human-in-the-loop).

- 2 **Import doc:** drop your existing plan (.md, .docx, .pdf) — transcribed into dated posts, **weeks preserved**.




"Import doc": your strategy document becomes calendar posts.

- 3 **New post** or + *add post* on a day: manual creation. Episodes also land here as drafts (chapter 7).

Generate the render — before posting

Select a post: the right inspector opens the post's whole workshop. **Produce** button: the app builds the image (FLUX) then the clip (Seedance) from the plan's ideas — or reuses the source you chose. A **thumbnail** of the render appears on the post-graph's *Render* node, and a **video preview** shows in the inspector: you see the result **before** clicking Send or enabling Auto.



A post's inspector: channels, caption, and the **emoji bar** + the **Telegram Premium pack** (branded one-tap tags) — plus the render preview before publishing.

Polish the caption — emojis & Premium pack

Under the caption, two one-tap rows, ideal for Telegram and X:

- 1 **Quick emojis:** 🍷 🍷 🍷 🍷 🍷 🍷 — inserted at the cursor.
- 2 **Telegram Premium pack:** branded tags per content vertical (👤 Deepotus Protocol, 🗨️ Rippled Signal, 👁️ Prophet, 🎲 Board Game, 🎲 D&D, 🏆 \$DEEP...) to sign each post in one gesture.

Auto or assisted?


Each post has a mode. **Assisted** (default): at the set time, the post flips to READY — you click *Send*. **Auto**: it goes out by itself on connected channels (Telegram, X). Status by color: gray = draft, cyan = scheduled, amber = ready, green = posted, red = failed.



PUBLISH

Connect your accounts

Settings → **Connected accounts**: the credentials the Scheduler uses. All stored locally.



Connected accounts: real status per channel — **auto** (publishes alone) or **assisted** (you click Send).

Telegram (5 minutes, free, recommended)

- 1 On Telegram, talk to `@BotFather` → `/newbot` → copy the token.
- 2 Add the bot as an **administrator** of your channel.
- 3 Paste token + chat ID, *Save*, restart, then *Send test message*.

X (Twitter)


Create your 4 keys at `developer.x.com` (writes available pay-per-use) and paste them. *Verify credentials* button. **YouTube** and **Instagram** stay assisted.



GAMES AND 3D


Make game assets

The **Game Assets** category (sidebar) turns an image from your Library into a **textured 3D model** ready for a game engine or 3D tool (Unity, Unreal, Godot, Blender...). Only the **fal.ai** key is required. Five generation engines, from cheapest to premium: **TripoSR** ($\approx$ \$0.07), **Hunyuan3D**, **Tripo** (default — best quality/price), **TRELLIS** and **Rodin**. The estimated cost is shown live before you start.




The Game Assets form: source image, engine (with description and rate), subject, options and export formats.

- 1 Pick a **source image** (your Library images — ideally a well-isolated character/object on a neutral background), then the **engine** from the dropdown.
- 2 Options: **Multi-view** (on by default, 4 views) first generates other angles via Seedream for a more faithful mesh; **Textures**; **T-pose** for characters you plan to rig; and the **export formats** — GLB always included, plus FBX / OBJ / STL / USDZ depending on the engine. A format the chosen engine cannot emit is flagged; each extra format can cost one more generation (the estimate shows it).
- 3 **Generate** opens the **pipeline canvas**: the Input → Views → Engine → Response nodes animate and fill in live (also visible when reopening an asset via its card title).



The live pipeline: each node fills in as the generation progresses (views, engine, result).

- 4 On the result card: ► **Rotate** spins the 3D model in the built-in viewer; each **format** downloads in one click; the **Shots** strip lets you download each view or send it to *Library* → *Images* (★) to reuse as a source. Rename (✎) or delete (🗑️) the asset from its card header.
- 5 All your assets also live in **Library** → **3D tab**: preview, full-size 3D rotation, GLB download, favorites ★, rename and delete.




The Library 3D tab gathers all your generated assets.

 3D rendering runs in your browser (WebGL): keep the tab in the foreground for smooth rotation. Files live under `DeepotusVideoGenData\assets\outputs\assets3d\` and follow your migration (chapter 16).



Create playing cards (Card Forge)


The **Card Forge** sub-tab (inside **Game Assets**, next to 3D / Sprites / Tiles / Materials) is a full playing-card editor: face, frame, typography, CSV data, 3D volume, PBR materials, print export and 3D export — eight modules in one document. No key is required; only **AI face generation** uses your existing fal.ai key.



1. Pick the format and resolution: they apply to the whole deck 2. The catalog offers 108 ready-made illustrations 3. The gauge tells you if the picture is sharp enough to print

The Face module: format and resolution at the top, a 108-illustration catalog, the effective-resolution gauge.

- 1 At the top of the lab, pick the **format** (poker, bridge, tarot, domino...), the **resolution** (150 / 300 / 600 DPI), the **bleed** and the **safe zone** — these settings apply to all eight modules and to the final export.
- 2 Browse the modules in the left rail: **01 Face** (imported illustration, AI-generated, or picked from the local catalog), **02 Frame** (borders, rarity, card back), **03 Typography** (title, rules box, flavor text, 23 already-bundled fonts), **04 Data** (CSV import → columns to fields → full deck), **05 Volume** (3D preview with real thickness and edge), **06 Materials** (8 PBR texture channels computed locally).
- 3 **07 Print** runs a **preflight check** on both the sheet and the produced file: bleed, cut marks, the PDF's `TrimBox` / `BleedBox`, the PNG's `phys` density (300 DPI = 11811 px/m) — every line is verified against the delivered file's bytes, never against an on-screen setting.



The preflight check: 0 errors, TrimBox and BleedBox present in the PDF, PNG density verified byte by byte.


4 08 3D Export exports both a **.glb** and a standalone **.gltf** with the full PBR map set (base color, normal, roughness, metallic, occlusion, height, emissive, ORM), ready for Unity, Unreal, Godot or Blender.

! The **emissive** channel is now adjustable in the 3D export (an **Emission 0..1** slider — luminescent ink on matte paper, or extinguished gilding) and wired into the GLB, glTF and MTL; it stays capped at 1.0 and its mask is derived from the card's luminance (adjustable threshold). A pose hidden by the frame is fixed **in one click** from the Face module — a panel button plus a contextual offer next to the masking percentage.



From image to 3D object

Card Forge's **09 Forge 3D** module turns your card into a printable 3D piece or an "NFT-grade" artifact: the card exports **as layers** (base material, illustration, veil, frame, typography, ornaments + the composite), each layer is processed in a **graph**, and assembly produces a clean GLB — plus an STL whenever everything is watertight. The whole free circuit runs locally, no key needed.



The Forge 3D graph: a material chain (silver holographic finish) on the background, an image→3D engine on the illustration — 30 cr (~\$0.60) shown before any launch.

- 1 Export the layers:** the screen proves the stacking **before** uploading (0 px of difference between the composite and the layer pile), the server counter-proves it in PIL and seals a manifest. An empty layer is still written — absence is measured, never guessed.
- 2 Edit the graph:** each layer gets a **plane** treatment (flat card) or a **relief** (a grid displaced by luminance, a solid closed by construction), a **material** from the Material Forge shop (tiled at physical scale, in millimeters), a **silver or gilded holographic** finish (iridescence, clearcoat, anisotropy — verified in the GLB's bytes) and a **transform** (position, rotation, scale, z stacking).

3 **Image→3D engines** (optional, paid): the **mesh 3D** treatment sends the layer to one of 7 engines — tripo, hunyuan, trellis, rodin, triposr (fal, priced in \$) or **Meshy 6/7 through the direct API** (credits; ultra +5 cr on meshy-7). The price shows on the node **and** in the graph footer before any launch; the status chip follows the job (queued → running → served · real credits → or failure with the literal message); relaunching a node resets its folder.

4 **Build**: assembly merges the engines' GLBs at their layer's place (full reindexing, the engine file's identity is dropped), writes the final GLB and the **STL** — accepted only when every element is a closed solid, with a motivated refusal otherwise. The docket lists the elements, their engines and the credits consumed.

 The image→3D engines are the only paid part of Card Forge: fal bills in \$ on your existing key, Meshy in credits on yours (`MESHY_API_KEY` , the Settings field shared with the 3D Studio). To try everything without spending a credit, set `MESHY MOCK=1` in the data `.env` : the local simulator runs the full circuit, fake credits included.



Prepare a part before slicing (Workbench)

Chapter 20 showed you how to get a file *out*. This one shows what to do *before*. The **Workbench** screen (Game Assets → 3D Studio → 07 · Établi 3D →) is the workshop where you repair, lay down, inspect, hollow, arrange and measure a part, so the slicer only has to slice. Nothing below is irreversible: **every gesture that changes the part writes one more VERSION**, and all of them stay on disk — the *model's life* column on the left lists them. (The Workbench screen itself is in French, like the rest of the application; its labels are quoted below as they appear.)

- 1 Set a target size.** Right rail, *taille cible* field: type the largest wanted dimension, in millimeters (80 for a figurine, 250 for a board). A 3D file carries no unit — while this field is empty the Workbench shows "glTF units" and refuses whatever needs a physical dimension (hollowing, arranging, exporting for print). That is deliberate: a refusal beats an invented millimeter.
- 2 Pick the printer.** Same rail, *imprimante* menu: the **Elegoo Centauri Carbon 2** profile is built into the application (256 × 256 mm bed, 256 mm height, one excluded area). If OrcaSlicer or ElegooSlicer is installed, its profiles show up too — read, never modified. On the plate, the green outline IS the chosen machine's bed; the red rectangle is the area where nothing may sit.
- 3 Repair in one click.** *Fiche* tab → *Réparer en un clic*. Checked by default: coincident vertices welded, duplicate faces dropped, flat triangles dropped, normals unified. *Trous bouchés* (cap holes) is unchecked by default: capping adds material, tick it knowingly. The detail appears in the bottom bar — "8 retourné(s) — fermé" (8 flipped — closed). A mesh left open prints badly: better to know here than at the machine.
- 4 Lay the part flat.** *Poser sur une face* (key F), then click the face that must touch the bed: the part turns and lands on the floor. The gesture joins the queue — the bottom bar says "1 modification en attente" — and *écrire la version* saves it. This gesture decides almost everything else: a well-placed part needs few supports. If unsure, *Orienter* proposes three ranked poses in the *Fiche* tab — share of the surface on the bed, share overhanging, height — and one click lays the part on the one you pick. A computed orientation does not always find the best pose: **always review the proposal**, no score knows which face you want to look good.
- 5 Look at the overhangs, then the slices.** *Surplombs* button: faces leaning less than 45° from horizontal turn orange — there, a layer has almost nothing under it and the slicer will add a support. Re-lay the part until the orange recedes. *Tranches* button: twenty cross-sections of the part, in blue, to see its real shape at each height. Both are **indicative previews**: nothing is written, and the slicer does the real slicing.

6 **Hollow it, when it helps.** *Fiche* tab → *Creuser*, 2 mm wall to begin with: the solid inside is replaced by a shell. It is first a **resin** (SLA) gesture, where a solid part costs a lot of material; with filament (FDM), the slicer already fills the inside with partial infill. It needs a *closed* mesh (step 3), and the Workbench states the thickness that holds when the requested wall is too thick for the part. For resin, then drill a drain hole: same tab, *Percer (drainage)*, a 3 to 4 mm diameter, then click the face to drill — preferably near the build plate, so the resin can flow out. The drill goes through the one wall under the click and stitches the hole closed; the opposite wall is left untouched. A solid part is refused (hollow it first).

7 **Arrange on the bed.** *Sur la plaque* separates the pieces, then *Ranger sur le plateau* lays them inside the real bed, rotated when that saves room, 2 mm apart. What does not fit moves to a second plate, and what is larger than the machine is reported. Drag a piece with the mouse (snapped to the bed grid, Shift frees it), nudge it with the arrow keys; its position, size and angle show under "Sur la plaque". The plate is a **view**: it is saved, but it does not change the model.

8 **Measure, then export.** *Mesurer* → two clicks give a distance, its x/y/z components and the angle between the two faces. Finally, *Export* tab → → **Impression 3D** writes the STL and the 3MF of the **displayed version**, at the target size in millimeters (chapter 20), and *Ouvrir dans le slicer* takes you there.

i **The order that works:** repair → lay down → inspect → hollow → export. Repair first, because every other gesture starts from sound geometry — and hollowing refuses an open mesh.

Glossary

The words your slicer will use, explained once.

Term	What it is
Base (bed contact)	The face that touches the bed. Wide and flat, the print holds; narrow, it comes loose.
Overhang	A wall leaning too far to rest on the layer below. Prusa puts the clean limit without support between 45 and 60° from horizontal.
Support	A disposable structure holding overhangs up. The slicer generates it; the Workbench only helps you need less of it.
Brim	A one-layer collar around the part, to keep its corners from lifting.
Raft	A platform printed UNDER the whole part. Safer than a brim, costlier, and it leaves a rough face.
Skirt	An outline drawn next to the part, not touching it, to prime the flow before starting.
Infill	The pattern inside. <i>Gyroid</i> is about equally strong in every direction; <i>Lightning</i> only supports top surfaces. 15 % is usually enough.
Seam	Each wall loop starts and ends somewhere, leaving a visible vertical line. The slicer places it (<i>Nearest, Aligned, Rear</i>) or scatters it (<i>Random</i>).
Retraction	Filament pulled back during travel moves, to avoid strings (<i>stringing</i>).
Layer	A horizontal slice. 0.2 mm is the usual compromise; 0.12 mm for detail, 0.28 mm for speed.
Perimeter (wall)	The number of wall loops. Two perimeters at 0.45 mm make a 0.9 mm wall: the minimum to aim for when modeling.
Bridging	Printing a layer over empty space, between two supports. It works over a few centimeters, with good cooling.
Warping	Corners lifting as they cool. Remedies: clean bed, brim, closed enclosure. Common with ABS/ASA, rare with PLA.
Watertightness	A solid is watertight when its surface is closed, with no hole or free edge. That is what the Workbench bar calls "fermé" (closed).
Manifold	Every edge belongs to exactly two faces. A mesh can be closed and still be hard to slice cleanly if it is not.

Decimation	Reducing the triangle count. A 500,000-triangle mesh prints no better than a 100,000 one, but slows everything down. <i>Fiche</i> tab → <i>Décimer</i> .
Hollowing	Replacing the solid inside with a shell of given thickness. Mostly useful in resin; 2 mm is a good starting point.
Drain hole	The hole(s) letting uncured resin (and air) out of a hollowed part. In the Workbench: <i>Fiche</i> tab → <i>Percer</i> (drainage).

Where to learn more

Every link below was opened and checked on 05/10/2026 (5 October 2026). They are grouped by what they teach you.

Topic	Resource	Checked
Supports and overhangs	Prusa — Support material (https://help.prusa3d.com/article/support-material_1698) (the threshold is measured from horizontal, 90 = vertical)	05/10/2026
Modeling for print	Prusa — Modeling with 3D printing in mind (https://help.prusa3d.com/article/modeling-with-3d-printing-in-mind_164135) (45–60° overhangs, walls, tolerances)	05/10/2026
Bridges	Prusa — Poor bridging (https://help.prusa3d.com/article/poor-bridging_1802)	05/10/2026
Layers and perimeters	Prusa — Layers and perimeters (https://help.prusa3d.com/article/layers-and-perimeters_1748)	05/10/2026
Infill	Prusa — Infill patterns (https://help.prusa3d.com/article/infill-patterns_177130)	05/10/2026
Seam	Prusa — Seam position (https://help.prusa3d.com/article/seam-position_151069)	05/10/2026
Skirt and brim	Prusa — Skirt and brim (https://help.prusa3d.com/article/skirt-and-brim_133969)	05/10/2026
Warping	Prusa — Warping (https://help.prusa3d.com/article/warping_2011)	05/10/2026
PLA filament	Prusa — PLA (https://help.prusa3d.com/article/pla_2062) (for Prusament PLA: nozzle 215 °C first layer then 210 °C, bed 60 °C)	05/10/2026
PETG filament	Prusa — PETG (https://help.prusa3d.com/article/petg_2059) (nozzle 230 °C then 240 °C, bed 85 °C then 90 °C)	05/10/2026
Choosing a filament	Prusa — Filament material guide (https://help.prusa3d.com/filament-material-guide) (comparison table of materials)	05/10/2026
Print defects	Simplify3D — Print Quality Guide (https://www.simplify3d.com/resources/print-quality-troubleshooting/) (about 26 defects, with photos)	05/10/2026
Slicer — prepare	OrcaSlicer — Object manipulation (https://github.com/SoftFever/OrcaSlicer/wiki/prepare_object_manipulation)	05/10/2026
Slicer — orient	OrcaSlicer — Auto orientation (https://github.com/SoftFever/OrcaSlicer/wiki/prepare_auto_orient) (which warns it may not always find the best pose)	05/10/2026

Slicer — calibrate	OrcaSlicer — Wiki (Calibrations) (https://github.com/SoftFever/OrcaSlicer/wiki/Home) (temperature, flow, retraction, pressure advance...)	05/10/2026
Machine	Elegoo — Centauri Carbon 2 wiki (https://wiki.elegoo.com/centauri-carbon-2-combo) (troubleshooting guides and error codes)	05/10/2026

 These links come from slicer makers and manufacturers: they describe THEIR machine and THEIR software. Numeric settings (temperatures, speeds) apply to the hardware they name; on another machine they are starting points, not truths. Two repair tools often recommended elsewhere, **Microsoft 3D Builder** and **Autodesk Meshmixer**, are no longer maintained, so they are not listed here.



3D-print your creations

Everything the app produces in 3D — and even your vector illustrations — exports to your printer in one gesture. The circuit is designed for the **Elegoo Centauri Carbon 2** (256 × 256 × 256 mm build plate) and its slicer **ElegooSlicer** (bundled with the machine, an OrcaSlicer fork — OrcaSlicer works too). Every export writes a folder under `DeepotusVideoGenData\assets\print3d\` with three files: the **STL** (the mesh), the **3MF** (the one you open — it carries the millimeter unit and the name), and `impression.json` (provenance, scale, watertightness). Everything is local and free: no provider calls.

- 1 Install ElegooSlicer once.** The installer associates `.3mf` files with the slicer on Windows: the app's "Open in slicer" button uses that directly. Without an association the app tells you so, and you can set `SLICER_PATH` (the slicer executable path) in the data `.env`.
- 2 From the card Forge 3D** (Card Forge, module 09): when the build slip delivers an STL — that is, when every element is a **closed solid**, the check that guarantees it — a → **3D printing** button appears next to it. Leave the size empty to keep the card's real millimeters, or type a target (e.g. `80` for a figurine). These exports are marked *guaranteed* watertight: the producer proved it.
- 3 From Game Assets · 3D:** every generated model (card characters, props) carries its → **3D printing** button. Give the largest dimension in millimeters — 80 for a figurine, 250 for a board — and the app centers the piece and rests it on the bed. The converter reads the original `model.glb`; the "optimized" file (compressed for game engines) is refused with a message that says what to use instead.
- 4 From the Vectorlab** (Export menu → 3D printing...): your vector illustrations become **printable reliefs**. Give a global height in millimeters, then override per layer with `name=mm` — for example `2, contours=5` for a stained-glass piece: glass at 2 mm, leads at 5. Strokes without a fill count by their stroke width; text objects are skipped (and counted in the message). Flat dimensions follow the document's unit (an A4 comes out at 210 × 297 mm).
- 5 Slice in the slicer.** The `.3mf` opens at the right size, in millimeters, resting on the bed. Supports, fine orientation, infill, filament profiles: that's the slicer's job — the app never generates G-code. Beyond 256 mm the export warns you the slicer will have to cut or shrink, without blocking you.



Meshes coming from the image→3D engines are marked *unknown* watertightness in `impression.json`: ElegooSlicer and OrcaSlicer repair small defects on import and say so. Only the card Forge 3D guarantees its solids — its "closed solid" check states it, not a promise.



End-to-end recipes, from idea to post



Recipe A — a meme posted to Telegram in 5 minutes

1

Library → "Create image": describe the meme (~3 s).

2

Scheduler → *New post*: caption (+ a 🔥 emoji), Telegram channel, time in 5 min, *Auto* mode.

3

Or don't wait: **Send now**. Posted.



Recipe B — your selfie video + an animation (UGC)

1

Film 10–15 s on your phone, transfer the file to the PC.

2

Studio → "UGC + animation" starter → *UGC video* node → *Upload your video* (🕒 Duration master on).

3

Choose the image + prompt of the animated half (Seedance) → *Run*. The output lasts exactly your clip.

4

Scheduler → attach the render, polish the caption (emoji + pack), schedule.



Recipe C — today's news reel

1

News → *Refresh* → tick 2–3 articles.

2

Build script (persona voice) → *Build illustration*.

3

Compose in **Templates** → *Send to Scheduler* → schedule or publish.



Recipe D — an episode of your novel

1

Episodes → title + paste (or upload) the chapter → pick a voice & language → *Generate narration*.

2

By AI (or By paragraph) → review the scenes → *Generate all illustrations* → set each scene's motion (Ken Burns / Seedance).

3

Assemble the episode → preview → *Send to Scheduler* (YouTube + Instagram drafts). One episode a day, done.



REFERENCE

Read the icons

You will recognise an icon at a glance, find its name, and print the full sheet to keep next to your screen. Free, no key needed.

One icon, one function

The whole app speaks the same language: the **Deepotus Glyph** set, 526 icons drawn with the same stroke. Each one stands for **a single function**: the same drawing means the same thing in the editor, the Photolab or the Card Forge. The old emojis and symbols are gone from buttons, menus and lists.

- 1 Rest the pointer on a button: its tooltip gives its name and, if it has one, its keyboard shortcut.
- 2 Look that name up in the icon sheet to read what it does and what not to confuse it with.
- 3 In the Vectorlab, the Photolab, the Sprite Lab and the Tile Lab, wait a little longer (900 ms) or click the small  button: an animated card SHOWS you the option.

Icon families

Icons are sorted into families. Knowing the family already tells you what to expect when you click.



Navigation

A screen, a tab or a panel: you go there, nothing is changed.



Actions

A button that does something: create, save, send, delete.



Editing

Change something already there: align, flip, group, lock.



Tools

The drawing and retouching tools of the labs.



Layers and adjustments

Layers, masks, styles and image adjustments.



Media and playback

Play, cut and edit video and sound.



States

What the app tells you, without you having to click.



Categories

Content families, each with its own colour.



3D

Objects and actions of the 3D workshops.

States worth knowing

A state icon is not clicked: it informs you. These are the ones to read before starting a paid render.

Icon	What it says
	Done: the render, the upload or the save worked.
	In progress or waiting: loading, generating, or queued.
	Failed: the tooltip or the message next to it says why.
	Warning: nothing is broken, but read before going on.
	Spending limit close or reached: beyond it, the app refuses the spend and asks you what to do.
	Generative field: AI can fill it in for you.
	Locked: nothing can move, change or delete the item while it is.
	Pinned (Studio): this node result is reused for free on the next render.

Print the icon sheet

The [icon sheet](#) shows them all, sorted by family, with their name, their function and the look-alike icons not to confuse them with.


- 1 Open the [icon sheet](#).
- 2 To print only part of it, click a family (for example **Vectorlab tools**) or type a word in the search: only what is shown gets printed.
- 3 Click **Print**: each family starts on a new A4 page, black on white. You can also download the [full sheet as a PDF](#).

i On the sheet, the red line starting with "≠" names the look-alike icons: Close hides a window, Delete destroys, Remove takes something out of a list without destroying it.

Animated cards in the labs

In the labs, some options have an animated card in three beats made of real screenshots: the starting point, the gesture, the result. It opens when you rest the pointer on the option for 900 ms, or with the  button next to it (in the Sprite Lab and the Tile Lab, even when the option is greyed out). Here they all are; cards of labs not yet translated still show the French interface.


Vectorlab



Conical

The color spins around the middle of your shape, like the spokes of a wheel.


Appearance > Conical



Transparency

Your shape slowly fades from one side to the other, as if it vanished into fog.

Appearance > Transparency



Pattern

Your shape fills up with little repeated stripes, like striped fabric.


Appearance > Pattern



Frame

Your long line of text fits into a box and wraps by itself, like in a book.


Appearance > Frame



On path

Your words lie down on the line you drew and follow it wherever it goes.


Appearance > On path



Contours

Lines join the places of the same height on your map, like the steps of a mountain.


Real map > Contours



Slice

Your map turns into six-sided tiles, colored by height: water, meadow, forest, mountain.


Real map › Slice



Side draft

The sides of your piece lean: the more you turn, the wider the bottom and the smaller the top.


3D printing › Side draft



Set as reference

Your picture becomes a reference: it fades, stays still, and a grid of squares lies on it to guide you.


Pixel persona › Set as reference



Create pixel layer

A see-through sheet lies on your reference picture: you draw on it, square by square.


Pixel persona › Create pixel layer



Pixelate image

Your picture is now made of big squares, like bricks: as many squares as the tile size.

Pixel persona › Pixelate image




Dark outline

A dark line goes around your drawing so it stands out from the background.

Pixel persona › Dark outline

Photolab



Workspace

Pick a workspace in the list: panels and tools rearrange for what you want to do.

Workspace list



Rectangular marquee (M)

Drag across the picture: a dotted rectangle surrounds the part you are going to change.


Tools › Rectangular marquee (M)



Magic wand (W)

Click a colour: the whole patch of that colour is selected at once.

Tools › Magic wand (W)



Crop (C)

Draw a frame on the picture then press Enter: everything outside the frame disappears.


Tools › Crop (C)



Brush (B)

Drag across the picture: a stroke of your colour is painted wherever you go.


Tools › Brush (B)



Clone stamp (S)

Alt-click what you want to copy, then paint elsewhere: the same piece of picture reappears under your brush.

Tools › Clone stamp (S)




Black and white

One click adds a layer that removes the colours: your picture turns grey, and you can remove it whenever you like.

Adjustment layer › Black and white


Sprite Lab



Auto-detect

The tool counts the pictures on your sheet by itself and draws a box around each one.


Sheet › Auto-detect



Feet

All your characters put their feet at the same height, as if standing on the same ground.


Sheet › Alignment › Feet



Mirror

Your character turns around and looks the other way, like in a mirror.


Preview › Mirror



Pixel art (9b)

Your drawing becomes big squares and only a few colours, like in an old video game.


Pixel art (9b) box



Outline

A line surrounds your character so it shows up on any background.

Post-processing › Outline



Shadow X / Y

A shadow appears behind your character, slightly offset, as if a lamp were lighting it.

Post-processing › Shadow X / Y


Tile Lab



Method

Mirror folds your texture like a sheet folded in four: the edges always meet perfectly.


Seamless mode › Method



Make seamless

Your texture joins up: laid side by side many times, you can no longer see where each piece ends.


Seamless mode › Make seamless



Placement grid

Your cells become diamonds, staggered like the paving of a game seen at an angle.


Tile sheet › Placement grid



Set

You choose how many pieces to make: many for nicely rounded corners, or sixteen to go fast.


Set mode › Set



New preview

A new map draws itself at random with your tiles, to see if they go well together.

Set mode › New preview



Shape

Your tile changes shape: a diamond for a game seen at an angle, or a six-sided hexagon.

Shapes mode › Shape

Good to know

i The icon of the app, of its shortcuts and of browser tabs is the real Deepotus logo. Inside the app, the logo shown follows your active brand kit (chapter [Templates and brand kits](#)).

i One single chevron is used everywhere to expand and collapse: pointing down, the section is open; pointing right, it is closed.



Troubleshooting

Symptom	Fix
<code>down</code> pill top-right	The engine isn't running: relaunch the Desktop icon.
Red pill on <code>fal</code> / <code>heygen</code> / <code>voice</code>	Missing/invalid key: Settings → API keys, fix, restart.
HeyGen avatars are slow to appear (or the list looks empty at first)	Normal on the very first load: HeyGen returns 1,000+ avatars and takes up to ~1 min to respond (their service is slow). The app pre-loads the list at startup and caches it — wait once, then it's instant. If the HeyGen status is green in Settings, your key is fine: it's neither the key, the network, nor the antivirus.
A voice gives a "paid plan" error	On the ElevenLabs free plan, only "premade" voices work via the API; custom library/community voices need a paid ElevenLabs plan. Pick a premade voice, or upgrade ElevenLabs.
After a restart, keys/library "gone"	No data is lost (it lives in <code>%LOCALAPPDATA%\DeepotusVideoGenData</code>). Relaunch cleanly via the Desktop icon.
The plan uses the built-in planner	No Anthropic key nor Ollama model (chapters 11–12).
<code>CERTIFICATE_VERIFY_FAILED</code> error	Antivirus inspecting HTTPS — handled automatically; otherwise disable HTTPS inspection.
"Port 8765 already in use"	The app is already running — open <code>http://deepotus.localhost:8765</code> .
Where are my files?	Images: <code>...\DeepotusVideoGenData\assets\images</code> · Videos: <code>...\assets\outputs\final</code> · Your uploads: <code>...\assets\outputs\uploads</code> · Sounds: <code>...\assets\audio</code> .

Deepotus Video Gen v2.8.0 — guide built from the real app. Everything runs locally on your machine; only calls to providers (fal.ai, HeyGen, ElevenLabs...) go over the Internet, with your own keys.