



Deepotus Video Gen

Complete guide — from double-click to your first auto-published post. **v2.4.0** · [Version française](#)

Never touched an editing tool? Good — this guide is written for you. Deepotus Video Gen turns an idea into a post-ready vertical video, with no software to learn. Everything runs **on your computer**; you plug in your own API keys (you pay your providers, not the app).

Five ways to make a video, from simplest to richest:

① a still image animated into a **cinematic clip** (Seedance) · ② an **avatar that speaks** your script (HeyGen) · ③ **your own footage** shot on a phone (UGC) · ④ a **composition** mixing several sources (Templates or Studio) · ⑤ a whole **narrated novel chapter** turned into an illustrated video (Episodes). Add **voiceover and background music** on top, then the **Scheduler** plans and publishes on its own.

0 · Start & API keys

1 · Set up a persona

2 · Create your images

3 · Cinematic clip (Seedance)

4 · Talking avatar (HeyGen)

5 · Your own video (UGC)

6 · Audio, voiceover & music

7 · Episodes — narrated novel

8 · Compose — Templates

9 · Compose — Studio (nodes)

10 · Schedule & publish

11 · AI engines (cloud & Ollama)

12 · Connect your accounts

13 · White-label

14 · Troubleshooting & recipes

15 · Studio — Effects, masks & preview

16 · Moving to another PC — migration


17 · Game Assets — 3D objects

18 · Card Forge — card editor

19 · Forge 3D — layers, graph & engines


0 Start & API keys

Double-click the **Deepotus Video Gen** icon on the Desktop. The app starts silently (no black window), then your browser opens on the studio. On first launch a 5-step wizard guides you: persona, providers, channels. Replay it anytime (command palette  → "Replay onboarding").



The first-run wizard. The "Generators" step lists your providers, including the **Ollama (local LLM)** option.

- 1 The app is **BYO keys**: you use your own API accounts. Only one is **required** to start — **fal.ai** (it makes both images *and* video clips). Create it at `fal.ai/dashboard/keys`.
- 2 Paste your keys in **Settings** → **API keys**, hit *Save*, then restart the app. The pills top-right (`fal` / `heygen` / `voice`) turn **green** when a key is active.



Settings → API keys: each key has a field, a status and a Save button. At the bottom, the **Local LLM (Ollama)** card (chapter 11).

How much does it cost? (ballpark)

Action	Service	Indicative cost
Create an image	fal.ai (FLUX)	≈ \$0.003
10 s cinematic clip	fal.ai (Seedance)	≈ \$0.30–0.60
1 min talking avatar	HeyGen	≈ 6 credits
Voiceover / narration	ElevenLabs	≈ per characters · free tier available
Upload your own video / sound	—	free
News summary / marketing plan	Anthropic / OpenAI / Gemini / Ollama	≈ \$0.01 / free (Ollama)
Publish to Telegram / X	—	free / your X credits

Which key unlocks what?


API key	Required?	What it unlocks
fal.ai	Required	Images (FLUX) <i>and</i> video clips (Seedance) — the only key you need to start.
HeyGen	Optional	Talking avatars (video type ②) and avatar + clip compositions.
ElevenLabs	Optional	Voiceover / narration — for your UGC videos and the Episodes feature (chapter 7), saved straight to the audio Library.
Anthropic, OpenAI or Gemini (any one is enough)	Optional	News summaries, marketing plans and genuinely AI-written scripts (News, Studio & Episode storyboards). Without a key: deterministic brand-template version.
Ollama (local LLM)	Optional	Same AI uses as above, 100% local and free (chapter 11).
X (4 keys)	Optional	Auto-publishing to X/Twitter (chapter 12).
Telegram (bot token)	Optional	Auto-publishing to a Telegram channel.

Tip: start with **fal.ai** alone, then add keys as you explore avatars, voiceover, AI and publishing.

Costs are billed by your providers, not the app. Watch your credits on their dashboards. Your keys, files and database stay **on your machine** (`%LOCALAPPDATA%\DeepotusVideoGenData`) — they survive updates and reinstalls.

1 Set up a persona

The persona is your **brand voice**: it feeds marketing plans, avatar scripts and the prompt generator. Go to **Settings** → **Personas**.




Settings → Personas: create, edit and activate your brand voice.

- 1 Click *New persona*: name, tone (e.g. "prophetic, playful"), audience, voice mode.
- 2 Click the card to **activate** it — the active persona is echoed across the generation screens.

2 Create your first images

Library tab: your library of images, video renders, sounds and captions. Everything lands here.




Library: create an image (FLUX), the **Prompt manager**, **Upload image** and **Upload video** (your own clip). Hover a render to play it.

- 1 "Describe an image to create" field → describe in one line, Enter. The FLUX image (~3 s) lands straight in the library.
- 2 Out of inspiration? Click **Prompt manager**: pick ingredients (subject, mood, light, motion...), "Random", or "Refine with AI", then *Use this prompt* — it **overwrites** the field for a more personal result.
- 3 Already have visuals? **Upload image** (or drag-and-drop anywhere). For your videos, see chapter 5; for sounds, chapter 6.

Prompt examples: "neon-red octopus mascot, abyssal background, cinematic" · "stylized market chart, bioluminescent mood" · "underwater throne, light rays, 9:16".

3 Video type ① — the cinematic clip (Seedance)

Quick tab, **Seedance** mode: one image + a prompt = a vertical cinematic clip.




Quick / Seedance: sources on the left, real-format preview on the right, estimated cost before you launch.

- 1 Pick the start image. Optional: an end image for a transition.
- 2 Write the motion prompt, grab a *template*, or open the **prompt generator** (it overwrites the default prompt). Set style, duration (5–60 s), format (9:16 default).
- 3 *Generate* → follow progress in the **Job Dock** (bottom bar). 30–45 s later the render is in Library.

Rename a render: in the Job Dock, click the pencil  next to the name (or play it first with ►). The new name instantly reflects in the Library and pickers.

4 Video type ② — the talking avatar (HeyGen)

Still in **Quick**, **HeyGen** mode: a script read by an avatar.




Quick / HeyGen: avatar + voice selection, upload your photo for a custom avatar, script.

- 1 Pick an avatar and a voice from the lists (search field available).
- 2 **Create your own avatar:** *Upload photo* → your photo (PNG/JPG) becomes a talking avatar in ~30 s, auto-selected.
- 3 Write the script (up to 4,900 chars), *Generate*. Allow 1 to 3 minutes.

🕒 The first avatar load — why it takes ~1 minute v1.15

HeyGen gives you a **very large catalogue** — often **1,000+ avatars** and 2,000 voices. The very first time the app requests this list, HeyGen's server can take **up to a minute** to respond. **This is HeyGen's own service being slow** — not an app bug, not a key problem, and not a network or antivirus issue.

So you don't have to wait, the app **pre-loads the list in the background at startup**: by the time you reach the HeyGen tab, the avatars are usually **already there**. Once fetched, the list is **cached** — every later open is **instant**.



Left: the HeyGen tab once loaded (1,281 avatars, search field, avatar & voice selectors). Right: why the first load takes ~1 min — pre-loaded at startup, then instant.


If you open the tab right after launch: the list may briefly show "loading..." (up to ~1 min, **once**). Let it fill, then everything is instant — including after you create your own avatar (the list refreshes itself). The **HeyGen status** pill (top-right and in Settings) stays **instant** at all times: it checks your key and credits without loading the whole catalogue.

Composition mode combines both: a Seedance clip + a HeyGen avatar in one video (sequential or split-screen). **Anti-cut** is automatic: the avatar always finishes its sentence.

5 Video type ③ — your own video (UGC)

Many creators film themselves on a phone — a selfie, a reaction, an unboxing. Deepotus welcomes these raw clips (**UGC**, user-generated content) and marries them to your animations.

- 1 **Simplest:** Library → **Upload video**. Your clip is imported, its real duration is measured, and it appears in the *Renders* tab like any render — reusable everywhere (Scheduler, Studio).
- 2 **In the Studio:** drag a **UGC video** node onto the canvas, click it, then *Upload your video*. A thumbnail fills the node.



Studio: the **"UGC + animation"** starter. On the right, the **UGC video** node with the **Duration master** switch.

The "master" duration — the key idea

When you turn on **Duration master** on your clip, **its real length** drives the final video. Generated animations (Seedance) are then calibrated around your clip — no more guessing "how many seconds?". You talk for 12 seconds to camera → the composition is 12 seconds, your illustration animation follows. It's the equivalent, for your own footage, of the avatar anti-cut.

Quick recipe: Studio → "UGC + animation" starter → upload your clip in the UGC video node (master on) → connect an Image + Seedance as the animated half → **Run**. Output: your video at the bottom, the animation on top, all at your clip's length.

6 Audio, voiceover & music v1.15.6

Sound makes a video. Deepotus now has a full **audio side**: import your own sounds, turn a script into a **voiceover**, and lay a **background music** track under any render.

The audio Library

The **Library** has an **Audio** tab next to Images and Renders. It gathers every sound: imported files, generated voiceovers and music tracks.

- 1 **Import a sound:** the *Upload sound* button (or drag-and-drop onto the Audio tab) → pick an `.mp3` / `.wav` / `.m4a` (and more). Its duration is measured; it appears in the Audio tab, playable inline, reusable everywhere.
- 2 The Audio tab also lists every **voiceover you generate**, so your narrations and your imported tracks live side by side.

The Audio tab links a few **royalty-free music banks** (Pixabay Music, Freesound, Free Music Archive, YouTube Audio Library). Only add music you have the right to use — Deepotus cannot ship licensed music for you.

Voiceover (ElevenLabs) v1.17.0

Turn text into a natural narration with your **ElevenLabs** key.

- 1 In **Quick**, open the **Voice Over** tab: a voiceover **on its own** (no image, no video), perfect to narrate over your own visuals. Write a short script (up to 2,500 characters — for longer narrations use Chapters), pick the **language**; the counter and the **estimated cost** update live.
- 2 Pick a voice from the **scrollable list** (▶ to preview, filter field) or keep the **app default voice**, and save your voice + language combos as reusable **castings** (★ Save). *Generate voice* → the mp3 plays inline and is **saved to Library** → **Audio**, ready to reuse (Studio, Episodes, Scheduler).

ElevenLabs free plan: the built-in "premade" voices (George, Sarah...) work via the API. Custom "library / community" voices require a paid ElevenLabs plan — the app tells you clearly when a chosen voice is paid-only.

Background music (BGM)

Drop a music bed under your video — the app mixes it for you.

- 1 **In the Studio:** add a **Music track** node, pick a track from the audio Library (with an inline preview), set its **volume** and toggle **loop to duration**. Wire it into the Render node.
- 2 **On the simple render too:** you can attach a background track before generating, no Studio needed.
- 3 The track is automatically **looped or trimmed to the exact length** of the final video, then mixed under the voice/clip. *Run* — your render comes back with music.

7 Video type ⑤ — Episodes (a narrated novel → video) v1.15.6

You have a story to tell — a chapter, a lore drop, a serialised novel. The **Episodes** page turns a block of text into a **narrated, illustrated vertical video**, scene by scene, ready for the Scheduler. Open it from the sidebar (**Episodes**). A clickable stepper walks you through four steps.

Step 1 — the chapter & its narration

- 1 Give the episode a **title**, then **paste your text** or *Upload* a `.txt`, `.docx` or `.pdf` — the text is extracted for you.
- 2 Pick a **voice** and a **language** (preview a voice with ►).
- 3 *Generate narration* → the whole chapter is read by ElevenLabs and saved to the audio Library. Long chapters are split and stitched back together automatically.

Step 2 — the storyboard (scenes)

A video needs shots. The app cuts your chapter into **scenes**, each one a bit of text + its own illustration.

- 1 Choose the split: **By paragraph** (verbatim, instant) or **By AI** (the AI groups the text into visual beats and writes a rich illustration prompt for each). You can **generate both and compare** before locking one in.
- 2 **Edit freely**: tweak each scene's text and its illustration prompt, reorder them (↑ ↓), add or remove scenes (+ / ×).

Step 3 — illustrations & motion

- 1 *Generate all illustrations* (or one at a time): one image per scene, with your chosen image model (FLUX or GPT Image).
- 2 Per scene, pick the **motion**: **Ken Burns** (a slow cinematic zoom/pan, default), **Seedance** (a true animated clip) or **Still** (a fixed image).


Step 4 — assemble & send

- 1 *Assemble the episode*: each scene's illustration is animated over its slice of narration, then all scenes are concatenated into one vertical video — follow it in the Job Dock.
- 2 Preview the result, then *Send to Scheduler* → a draft post on **YouTube** and **Instagram**, ready to caption and schedule.

Serialising a novel? Keep the same persona and voice across the series for consistency, then publish one episode per day from the Scheduler.

8 Video type ④ — compose with Templates

Templates tab: ready-made 9:16 layouts (news reel + avatar + brand strip, split-screen, PIP corner, triptych...). Each thumbnail reflects its **real region layout** — they're all different.




Templates: gallery of layouts (distinct thumbnails), spatial editor on the right (drag regions or their handles).

- 1 Click a thumbnail: the right editor shows its regions. Move them, resize them.
- 2 *Save layout* stores your version — it **appears immediately** in the gallery. *New template* creates a blank one.
- 3 *Open in Studio* switches to the node editor for deeper editing.

9 Compose with the Studio (node editor) EXPANDED

The **Studio** is the advanced workshop: you wire *nodes* (sources → generators → montage → render) like a visual recipe. Don't worry — always start from a ready-made starter graph.



Studio: the node palette on the left, the canvas in the middle, the inspector on the right. Top bar: **New**, **Preview**, **Run**.


The basic gestures

- 1 **Start from a template:** in the palette, *Starter graphs* section, click "Seedance solo", "Avatar post", "UGC + animation", "News reel" or "Timeline montage".
- 2 **Start from scratch:** **New** button (top bar) → a blank canvas with just the *Render* output node. The right panel shows the **graph name**, editable.
- 3 **Add a node — drag & drop:** grab any entry from the palette (Image, Seedance, UGC video, Music track, Text...) and **drop it on the canvas**. The node appears where you dropped it.
- 4 **Wire:** drag from a node's output port (right side) to another's input port (left side). **Move:** drag the node body. **Delete:** the **x** in the node header (it also removes its connections).
- 5 **Check then run:** *Preview* shows inputs/the render; *Run* compiles the graph and launches the real render (follow it in the Job Dock).

When you pick an asset in a node (a Library image, an existing render, your UGC clip, a music track), a **thumbnail fills the node**: you see at a glance what the chain will produce.

10 Schedule & publish your week EXPANDED


Scheduler tab: the publishing calendar, and the heart of "human-in-the-loop" — **nothing goes out without your approval** (unless you decide otherwise).



Scheduler: week/month, AI plan, document import, the post's graph + inspector on the right.


Three ways to fill the calendar

- 1 **Generate plan:** describe your week in one sentence → the AI proposes a full plan (dated posts, hooks, captions, formats). *Review the preview* before "Add to calendar".



"Generate plan": a one-sentence brief becomes a full plan, reviewed by you (human-in-the-loop).

- 2 **Import doc:** drop your existing plan (.md, .docx, .pdf) — transcribed into dated posts, **weeks preserved**.




"Import doc": your strategy document becomes calendar posts.

3 New post or + *add post* on a day: manual creation. Episodes also land here as drafts (chapter 7).

Generate the render — before posting

Select a post: the right inspector opens the post's whole workshop. **Produce** button: the app builds the image (FLUX) then the clip (Seedance) from the plan's ideas — or reuses the source you chose. A **thumbnail** of the render appears on the post-graph's *Render* node, and a **video preview** shows in the inspector: you see the result **before** clicking Send or enabling Auto.



A post's inspector: channels, caption, and the **emoji bar** + the **Telegram Premium pack** (branded one-tap tags) — plus the render preview before publishing.

Polish the caption — emojis & Premium pack

Under the caption, two one-tap rows, ideal for Telegram and X:

- 1 Quick emojis:** 🍷 🔥 🚀 🥰 🎉 🍀 💯 — inserted at the cursor.
- 2 Telegram Premium pack:** branded tags per content vertical (👤 Deepotus Protocol, 🗨️ Rippled Signal, 👁️ Prophet, 🎲 Board Game, 🎲 D&D, 🏦 \$DEEP...) to sign each post in one gesture.

Auto or assisted?

Each post has a mode. **Assisted** (default): at the set time, the post flips to `READY` — you click *Send*. **Auto**: it goes out by itself on connected channels (Telegram, X). Status by color: gray = draft, cyan = scheduled, amber = ready, green = posted, red = failed.

11 Choose your AI engine (cloud & local) v1.15

Several features call a language model: **news summaries** (News), **marketing plan generation** (Scheduler) and **Episode storyboards** (the "By AI" split). As of v1.15, you choose the **provider** for each.

Cloud providers — Anthropic, OpenAI, Gemini

Paste whichever key(s) you want in **Settings** → **API keys**: **Anthropic** (Claude), **OpenAI** (GPT) and/or **Google Gemini**. Just **one** is enough to enable summaries and AI plans; add several if you want to switch between them.

Gemini model v1.21: by default the app uses `gemini-flash-latest`, Google's stable alias that automatically tracks the latest Flash release — nothing to do. To pin a specific version (e.g. `gemini-2.0-flash`), fill the "**Gemini — modèle**" field in **Settings** → **API keys** and restart the backend. The same Gemini key also powers the Studio's native Google video models (Veo 3.1).

- 1 In **Settings** → **Provider defaults**, pick the provider per role (summarizer, planner) from a dropdown. A role is only marked "missing" when **none** of its possible providers has a key.
- 2 **Automatic fallback**: if the chosen provider is unavailable (no key, outage), the app moves to the next available one — you're never stuck.

Engine priority: your choice in *Provider defaults* first, then the default order **Anthropic** → **OpenAI** → **Gemini** → **Ollama** → **built-in planner** (deterministic). The plan badge shows which engine actually ran.


Fully local LLM — Ollama (free, private)

Plans can also run on an **AI model installed on your machine**: it's **free** and your strategy never leaves your computer.

- 1 Install **Ollama** from `ollama.com` (Windows).
- 2 Pull a model: `ollama pull qwen2.5:14b-instruct`. *Recommended: 14B or more.*
- 3 **Settings** → **API keys** → **Local LLM (Ollama)**: leave the default URL, set the *model name*, *Save*, restart the app.

12 Connect your accounts (publish)

Settings → **Connected accounts**: the credentials the Scheduler uses. All stored locally.



Connected accounts: real status per channel — **auto** (publishes alone) or **assisted** (you click Send).

Telegram (5 minutes, free, recommended)


- 1 On Telegram, talk to `@BotFather` → `/newbot` → copy the token.
- 2 Add the bot as an **administrator** of your channel.
- 3 Paste token + chat ID, *Save*, restart, then *Send test message*.

X (Twitter)

Create your 4 keys at `developer.x.com` (writes available pay-per-use) and paste them. *Verify credentials* button. **YouTube** and **Instagram** stay assisted.

13 White-label (rebrand the studio)

Settings → **Branding**: name, subtitle, taglines, colors and logo apply to the splash, the sidebar and every accent of the app.



Branding: live wordmark preview. Stored in the data folder — survives updates.

- 1 Edit the name, taglines, primary color and accent.
- 2 *Logo...* to upload your logo (auto-normalized).
- 3 *Save brand* — applied everywhere. *Reset to deepotus* restores the default brand.

Your logo also appears on the **startup loading screen**: replace it here and your own logo greets users every launch — the studio becomes truly yours.

14 Troubleshooting

Symptom	Fix
<code>down</code> pill top-right	The engine isn't running: relaunch the Desktop icon.
Red pill on <code>fal</code> / <code>heygen</code> / <code>voice</code>	Missing/invalid key: Settings → API keys, fix, restart.
HeyGen avatars are slow to appear (or the list looks empty at first)	Normal on the very first load: HeyGen returns 1,000+ avatars and takes up to ~1 min to respond (their service is slow). The app pre-loads the list at startup and caches it — wait once, then it's instant. If the HeyGen status is green in Settings, your key is fine: it's neither the key, the network, nor the antivirus.
A voice gives a "paid plan" error	On the ElevenLabs free plan, only "premade" voices work via the API; custom library/community voices need a paid ElevenLabs plan. Pick a premade voice, or upgrade ElevenLabs.
After a restart, keys/library "gone"	No data is lost (it lives in <code>%LOCALAPPDATA%\DeepotusVideoGenData</code>). Relaunch cleanly via the Desktop icon.
The plan uses the built-in planner	No Anthropic key nor Ollama model (chapters 11–12).
<code>CERTIFICATE_VERIFY_FAILED</code> error	Antivirus inspecting HTTPS — handled automatically; otherwise disable HTTPS inspection.
"Port 8765 already in use"	The app is already running — open <code>http://127.0.0.1:8765</code> .
Where are my files?	Images: <code>...\DeepotusVideoGenData\assets\images</code> · Videos: <code>...</code> <code>\assets\outputs\final</code> · Your uploads: <code>...\assets\outputs\uploads</code> · Sounds: <code>...</code> <code>\assets\audio</code> .

-  **Recipe A — a meme posted to Telegram in 5 minutes**
- 1

Library → "Create image": describe the meme (~3 s).
- 2

Scheduler → *New post*: caption (+ a 🔥 emoji), Telegram channel, time in 5 min, *Auto* mode.
- 3

Or don't wait: **Send now**. Posted.

Recipe B — your selfie video + an animation (UGC)

- 1 Film 10–15 s on your phone, transfer the file to the PC.
- 2 **Studio** → "UGC + animation" starter → *UGC video* node → *Upload your video* (🕒 Duration master on).
- 3 Choose the image + prompt of the animated half (Seedance) → *Run*. The output lasts exactly your clip.
- 4 **Scheduler** → attach the render, polish the caption (emoji + pack), schedule.

Recipe C — today's news reel

- 1 **News** → *Refresh* → tick 2-3 articles.
- 2 *Build script* (persona voice) → *Build illustration*.
- 3 Compose in **Templates** → *Send to Scheduler* → schedule or publish.

Recipe D — an episode of your novel v1.15.6

- 1 **Episodes** → title + paste (or upload) the chapter → pick a voice & language → *Generate narration*.
- 2 **By AI** (or By paragraph) → review the scenes → *Generate all illustrations* → set each scene's motion (Ken Burns / Seedance).
- 3 *Assemble the episode* → preview → *Send to Scheduler* (YouTube + Instagram drafts). One episode a day, done.

15 Studio — Effects, masks & preview new

Two additions to the node editor (Studio).

The “Effects / Mask” node

Drag the **Effects / Mask** node (*Composition* section) and wire its output to the new **fx** port on the **Render** node. The panel shows a **thumbnail grid** so you pick an effect at a glance.

Available effects: **LUT / Grade** (named grades — Teal&Orange, Cyberpunk, Deep-sea, Noir, Vintage...), **VHS** (sequenceable line displacement, intensity + speed), **Colorize** (Sepia, Duotone, Matrix, Red-alert...), parametric **Gradients** (2 colors, angle,

opacity, blend) — plus modern classics: grain, vignette, chromatic aberration, glitch, bloom, halation, scanlines/CRT, cinematic letterbox, old film, sharpen, blur, dreamy, pixelate, camera shake, mirror, invert.

Under “**Apply to**”, choose *Whole render* (global post-pass) or a **specific node** (the targeted layer). Add several Effects nodes to stack effects or target different layers.

- 1 + Effect → click a thumbnail → set intensity (and preset / speed depending on the effect).
- 2 Wire the output to the Render **fx** port, choose the target, then *Preview* or *Run*.

“Preview” button — see it before you pay

The **Preview** button (top bar) produces a **free, fast** render of the final composition: not-yet-generated **Seedance/HeyGen** clips are replaced by their **source still**. You see the framing, layers, **overlays and effects** exactly as in the final render — *before* the paid **Run**.

Text overlays (Text overlay / Ticker / Separator) now apply on **every** composition at the Render node (UGC, montage, spatial compose), not only spatial compose.

16 Moving to another PC — migration & export new

Everything is stored **locally**: the app in `%LOCALAPPDATA%\DeepotusVideoGen`, and your data (generations, **calendar + scheduled posts**, API keys) in `%LOCALAPPDATA%\DeepotusVideoGenData`. To continue on another PC (travel) **without interrupting the schedule**, a kit lives in `Desktop\deepotus-migration`.

With the scripts (recommended)

- 1 On THIS PC, with a USB/external drive plugged in: `export-migration.ps1 -Dest "E:\deepotus-transfer"`. Copies the app + all data + a consistent DB snapshot.
- 2 Copy that folder to the other laptop, then **double-click** `IMPORT-TOUT.cmd` at its root: it restores **everything in one click** — app, generations, calendar, keys, *and* your Claude Code sessions if they are in the kit. (Manual equivalent: `import-migration.ps1 -Src "E:\deepotus-transfer"`.) The `DeepotusVideoGen-Setup-*.exe` installer is also bundled for a clean install on a third machine.
- 3 Launch **Deepotus Video Gen**: *Library* = your renders, *Scheduler* = your scheduled posts, *Settings* = your keys.


Manually

Close the app, copy the **two folders** above to the **same** `%LOCALAPPDATA%\` paths on the other PC, then run `scripts\create-desktop-shortcut.ps1` (or double-click `scripts\launch-silent.vbs`).

Do **not** run the app on both PCs at once (scheduled posts could double-publish). Best: **same Windows username** on both (paths match; otherwise the import script rewrites them). The **Scheduler runs while the app is open** — keep it running around posting times on the travel laptop, and keep the old PC closed.

17 Game Assets — image → 3D object new


The **Game Assets** category (sidebar) turns an image from your Library into a **textured 3D model** ready for a game engine or 3D tool (Unity, Unreal, Godot, Blender...). Only the **fal.ai** key is required. Five generation engines, from cheapest to premium: **TripoSR** ($\approx$ \$0.07), **Hunyuan3D**, **Tripo** (default — best quality/price), **TRELLIS** and **Rodin**. The estimated cost is shown live before you start.



The Game Assets form: source image, engine (with description and rate), subject, options and export formats.


- 1 Pick a **source image** (your Library images — ideally a well-isolated character/object on a neutral background), then the **engine** from the dropdown.
- 2 Options: **Multi-view** (on by default, 4 views) first generates other angles via Seedream for a more faithful mesh; **Textures**; **T-pose** for characters you plan to rig; and the **export formats** — GLB always included, plus FBX / OBJ / STL / USDZ depending on the engine. A format the chosen engine cannot emit is flagged; each extra format can cost one more generation (the estimate shows it).

- 3 **Generate** opens the **pipeline canvas**: the Input → Views → Engine → Response nodes animate and fill in live (also visible when reopening an asset via its card title).



The live pipeline: each node fills in as the generation progresses (views, engine, result).

- 4 On the result card: ► **Rotate** spins the 3D model in the built-in viewer; each **format** downloads in one click; the **Shots** strip lets you download each view or send it to *Library* → *Images* (★) to reuse as a source. Rename (✎) or delete (🗑️) the asset from its card header.
- 5 All your assets also live in **Library** → **3D tab**: preview, full-size 3D rotation, GLB download, favorites ★, rename and delete.




The Library 3D tab gathers all your generated assets.

3D rendering runs in your browser (WebGL): keep the tab in the foreground for smooth rotation. Files live under `DeepotusVideoGenData\assets\outputs\assets3d\` and follow your migration (chapter 16).


18 Card Forge — playing-card editor

The **Card Forge** sub-tab (inside **Game Assets**, next to 3D / Sprites / Tiles / Materials) is a full playing-card editor: face, frame, typography, CSV data, 3D volume, PBR materials, print export and 3D export — eight modules in one document. No key is required; only **AI face generation** uses your existing fal.ai key.



The Face module: a local catalog of 108 illustrations, import, AI generation, and a live effective-resolution gauge.

- 1 At the top of the lab, pick the **format** (poker, bridge, tarot, domino...), the **resolution** (150 / 300 / 600 DPI), the **bleed** and the **safe zone** — these settings apply to all eight modules and to the final export.
- 2 Browse the modules in the left rail: **01 Face** (imported illustration, AI-generated, or picked from the local catalog), **02 Frame** (borders, rarity, card back), **03 Typographie** (title, rules box, flavor text, 23 already-bundled fonts), **04 Data** (CSV import → columns to fields → full deck), **05 Volume** (3D preview with real thickness and edge), **06 Materials** (8 PBR texture channels computed locally).
- 3 **07 Print** runs a **preflight check** on both the sheet and the produced file: bleed, cut marks, the PDF's TrimBox / BleedBox, the PNG's **phys** density (300 DPI = 11811 px/m) — every line is verified against the delivered file's bytes, never against an on-screen setting.



The preflight check: 0 errors, TrimBox and BleedBox present in the PDF, PNG density verified byte by byte.

4 08 3D Export exports both a `.glb` and a standalone `.gltf` with the full PBR map set (base color, normal, roughness, metallic, occlusion, height, emissive, ORM), ready for Unity, Unreal, Godot or Blender.

The **emissive** channel is now adjustable in the 3D export (an **Emission 0..1** slider — luminescent ink on matte paper, or extinguished gilding) and wired into the GLB, glTF and MTL; it stays capped at 1.0 and its mask is derived from the card's luminance (adjustable threshold). A pose hidden by the frame is fixed **in one click** from the Face module — a panel button plus a contextual offer next to the masking percentage.

19 Forge 3D — layers, graph & image→3D engines new

Card Forge's **09 Forge 3D** module turns your card into a printable 3D piece or an "NFT-grade" artifact: the card exports **as layers** (base material, illustration, veil, frame, typography, ornaments + the composite), each layer is processed in a **graph**, and assembly produces a clean GLB — plus an STL whenever everything is watertight. The whole free circuit runs locally, no key needed.

- 1 **Export the layers:** the screen proves the stacking **before** uploading (0 px of difference between the composite and the layer pile), the server counter-proves it in PIL and seals a manifest. An empty layer is still written — absence is measured, never guessed.
- 2 **Edit the graph:** each layer gets a **plane** treatment (flat card) or a **relief** (a grid displaced by luminance, a solid closed by construction), a **material** from the Material Forge shop (tiled at physical scale, in millimeters), a **silver or gilded holographic** finish (iridescence, clearcoat, anisotropy — verified in the GLB's bytes) and a **transform** (position, rotation, scale, z stacking).
- 3 **Image→3D engines** (optional, paid): the **mesh 3D** treatment sends the layer to one of 7 engines — tripot, hunyuan, trellis, rodin, tripot (fal, priced in \$) or **Meshy 6/7 through the direct API** (credits; ultra +5 cr on meshy-7). The price shows on the node **and** in the graph footer before any launch; the status chip follows the job (queued → running → served · real credits → or failure with the literal message); relaunching a node resets its folder.
- 4 **Build:** assembly merges the engines' GLBs at their layer's place (full reindexing, the engine file's identity is dropped), writes the final GLB and the **STL** — accepted only when every element is a closed solid, with a motivated refusal otherwise. The docket lists the elements, their engines and the credits consumed.

The image→3D engines are the only paid part of Card Forge: fal bills in \$ on your existing key, Meshy in credits on yours (`MESHY_API_KEY` , the Settings field shared with the 3D Studio). To try everything without spending a credit, set `MESHY MOCK=1` in the data `.env` : the local simulator runs the full circuit, fake credits included.