Deepotus Video Gen guide

Everything you can do with the app, one task at a time: short steps, animated pictures taken from the real app, and the icons you will see on screen.

- Take a tour of the app
- Install, pick a language, update
- Connect your keys and AI engines
- Keep spending under control
- Settings, personas and branding
- Move to another computer
- Drive it from your phone
- Make a video in minutes (Quick)
- Compose with the Studio
- Reuse templates and brand kits
- Tell a story in episodes
- Sound, voices and visual effects
- Play an avatar, deferred or live
- Install and train a local RVC voice
- Edit a video
- Grade, mix and add titles
- Deliver: renders and formats
- Plan and publish your week

Connect your accounts

Follow the news and turn it into posts

Organise and find your media

Retouch a photo (Photolab)

Layers, masks and workspaces

Draw vectors (Vectorlab)

Pixel art, export, maps and 3D objects

Make game assets

Animate sprites (Sprite Lab)

Paint tiled levels (Tile Lab)

Create playing cards (Card Forge)

From image to 3D object

Prepare a part before slicing (Workbench)

3D-print your creations

End-to-end recipes, from idea to post

Read the icons

Troubleshooting

GET STARTED

Take a tour of the app

You will find your way around the app, move from one screen to another without searching, and read what the top bar tells you. Five minutes, free.

Deepotus Video Gen is a studio that runs **on your own computer**. It turns an idea into a video, a picture, a 3D object or a published post. You plug in your own AI provider keys: you pay those providers, not the app. Everything you create stays with you.

The screen at a glance

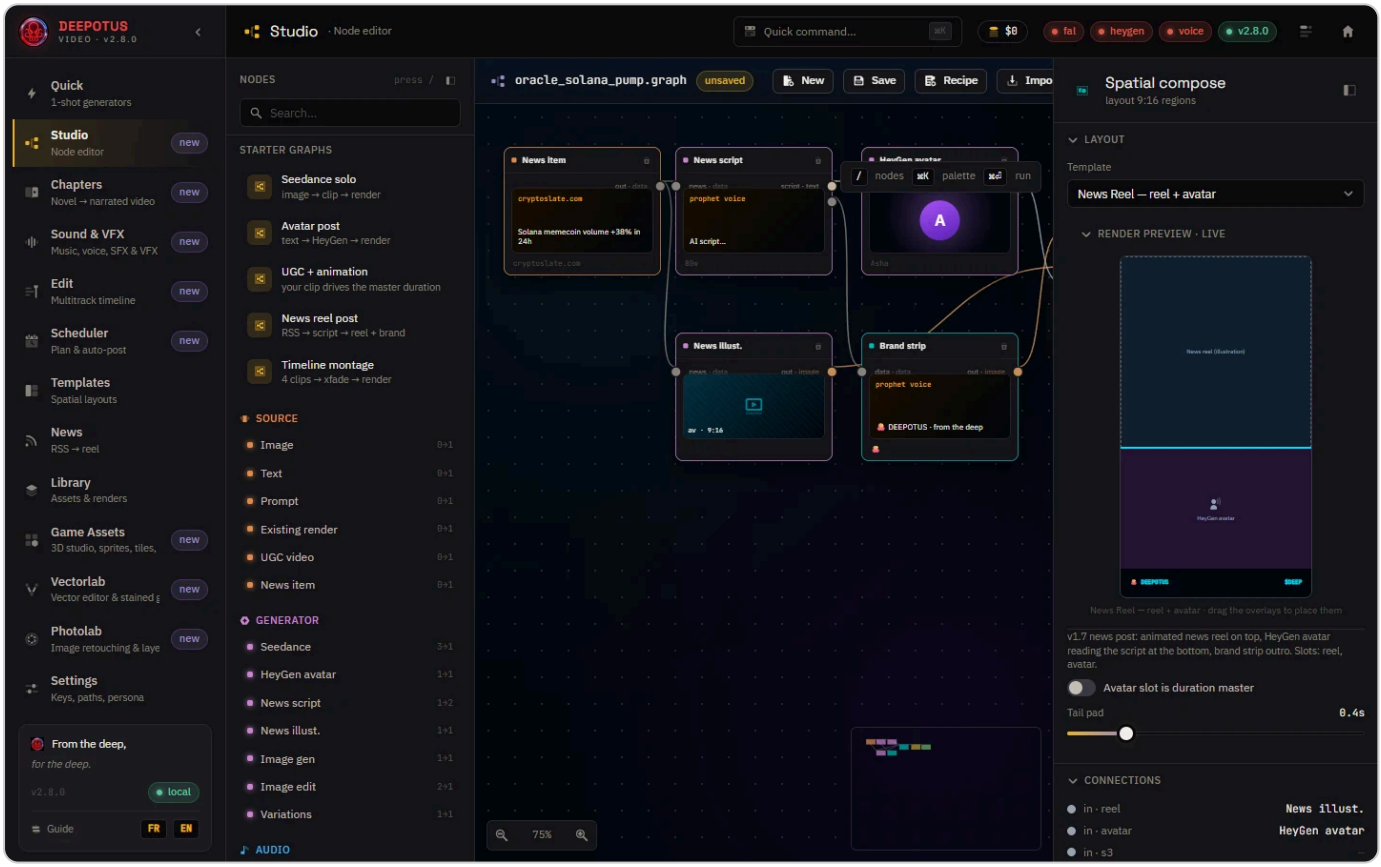

1. The rail: one screen per task. 2. Ctrl+K: go anywhere from the keyboard. 3. This month's spending and key status, always in sight.

The screen rail on the left, the top bar with the quick command, the spending counter and the key status.

1. **The rail, on the left**: one screen per task. `Quick` for a one-click video, `Studio` to compose, `Edit` to cut a video, `Library` to find your media, `Settings` for the keys. The arrow at the top of the rail folds it into icons only.
2. **The quick command**: click `Quick command...` or press `Ctrl` + `K`, type a few letters, Enter.
3. **The spending counter** shows what you spent this month with your providers. One click opens Settings.

4. **The badges** `fal`, `heygen` and `voice` turn green when the matching key is set. One click takes you to that key.

5. **The render queue**, top right, lists what is being computed and what just finished, in the `Running` and `Recent` tabs.

Go fast with the quick command

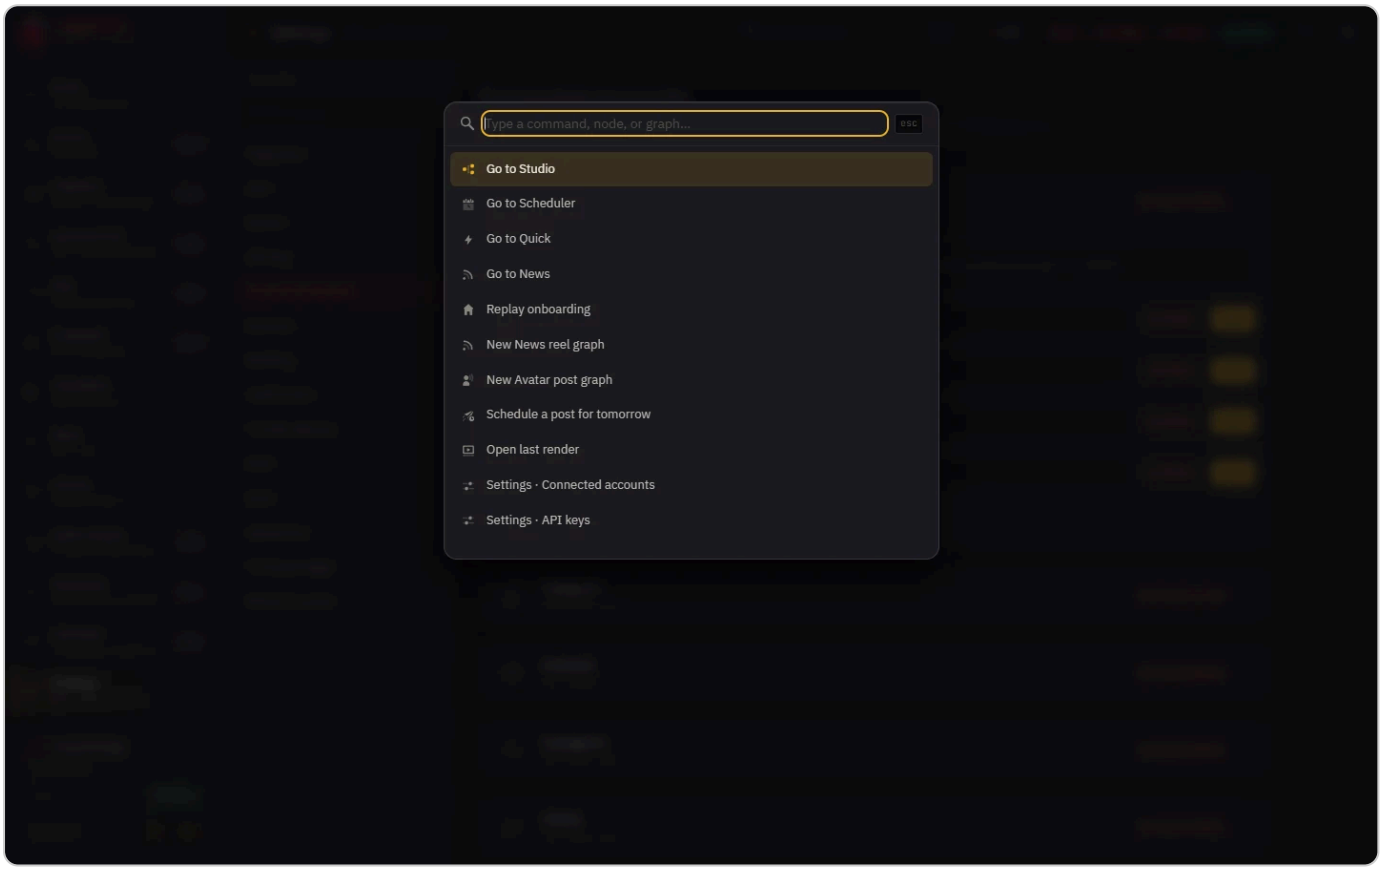

1. Type, then Enter

Ctrl+K opens the quick command: type a few letters to go to a screen or start an action.

The quick command goes to any screen (`Go to Studio`, `Go to News`...), opens the last render, creates a starter graph or opens a Settings section. `Replay onboarding` restarts the first-launch tour.

The screens of the rail

Quick
A one-click video: animated clip, talking avatar, composition, voiceover.

Studio
A node graph to compose a video piece by piece.

Chapters
From a text or a novel to a narrated video, scene by scene.

Sound & VFX
Music, voices, sound effects and visual effects.

Edit
A multitrack timeline to cut, grade and deliver.

Scheduler
Plan and publish your posts for the week.

Templates
Reusable layouts and your brand kits.

News
Follow the news and turn it into short videos.

Library
All your pictures, renders, sounds and objects, with their origin.

Game Assets
3D, sprites, tiles, materials and playing cards.

Vectorlab
Vector drawing, pixel art, maps and printable objects.

Photolab
Photo retouching with layers, masks and filters.

The guide opens from the foot of the rail: `FR` or `EN`.

What the app does to protect you

Nothing leaves without you

Only this PC can change, generate or spend. Another device on the network can only read, and only once paired.

Errors that make sense

A missing key or a server refusal is shown in plain words, with what to do.

Safe imports

A picture with the same name becomes name-1, name-2…; a file that is not a real picture is refused and named.

In-app dialogs

Confirmations and text prompts use one Deepotus dialog, never a browser pop-up.

Good to know

The desktop icon opens the app at `http://deepotus.localhost:8765`. Each browser tab carries the name of the open tool, for example “Deepotus — Vectorlab”.

Only AI generations cost money, and your provider bills them. The chapter Keep spending under control shows how to set limits.

Shortcuts

Key	Effect
`Ctrl` + `K`	Opens the quick command
`Esc`	Closes the quick command or a dialog

GET STARTED

Install, pick a language, update

You will launch the app for the first time, pick its language, know what each use costs and keep it up to date. Ten minutes, free.

First launch

1. During installation, pick the language: **French** or **English**. It becomes the language of the whole app.
2. Double-click the **Deepotus Video Gen** desktop icon. The app starts without a black window, then your browser opens on the studio.
3. On the very first launch, a five-step welcome introduces the studio and lets you pick a persona, your generation providers and the networks you publish to.
4. Set at least one key, the **fal.ai** one: it makes pictures and video clips. The chapter Connect your keys shows how.

> The welcome can be replayed at any time from the quick command (`Ctrl` + `K`): **Replay onboarding**.

Switch language

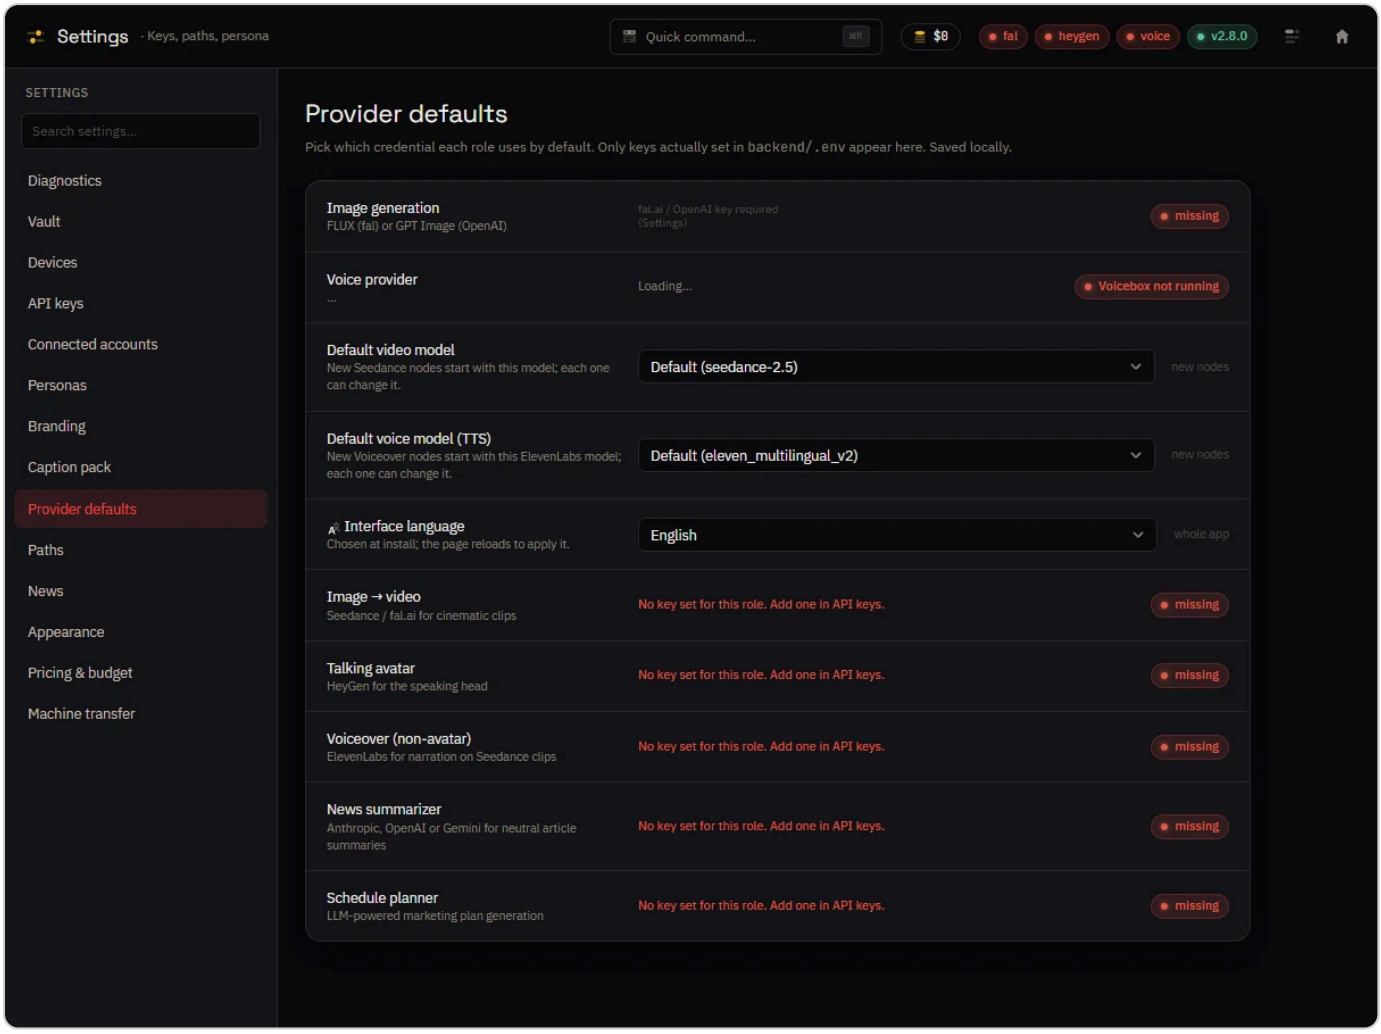

1. French or English, for the whole app

Settings › Provider defaults › Interface language: the page reloads in the chosen language.

1. Open **Settings** › **Provider defaults**.
2. In **Interface language**, pick **Français** or **English**. The page reloads in the chosen language.

Translated: the sidebar, Settings, the Library, Quick, News, the Scheduler, the Studio, Templates, Episodes, the editor, Sound & VFX, subtitles, machine transfer and every confirmation dialog. Lists sent by the server (effects, transitions, presets) also arrive in the chosen language. The labs (Vectorlab, Photolab, Card Forge…) are being translated one by one; meanwhile their known texts are translated on display. What you type yourself is never translated.

What each use costs

The app is free. Only AI generations cost money, billed by the provider of the key. Rough figures at default rates:

Action	Provider	Indicative cost
Create a picture	fal.ai (FLUX)	`≈ $0.003`
10 s video clip	fal.ai (Seedance 2.5, 720p)	`≈ $4.73`
Talking avatar, 1 min	HeyGen	`≈ $0.24`
Voiceover	ElevenLabs	per character, free tier available
News summary, post plan	Anthropic, OpenAI, Gemini or Ollama	a few cents, free with Ollama
Import your videos, sounds, pictures	none	free

Your real rates go into **Settings** › **Pricing & budget**; see Keep spending under control.

Keep the app up to date

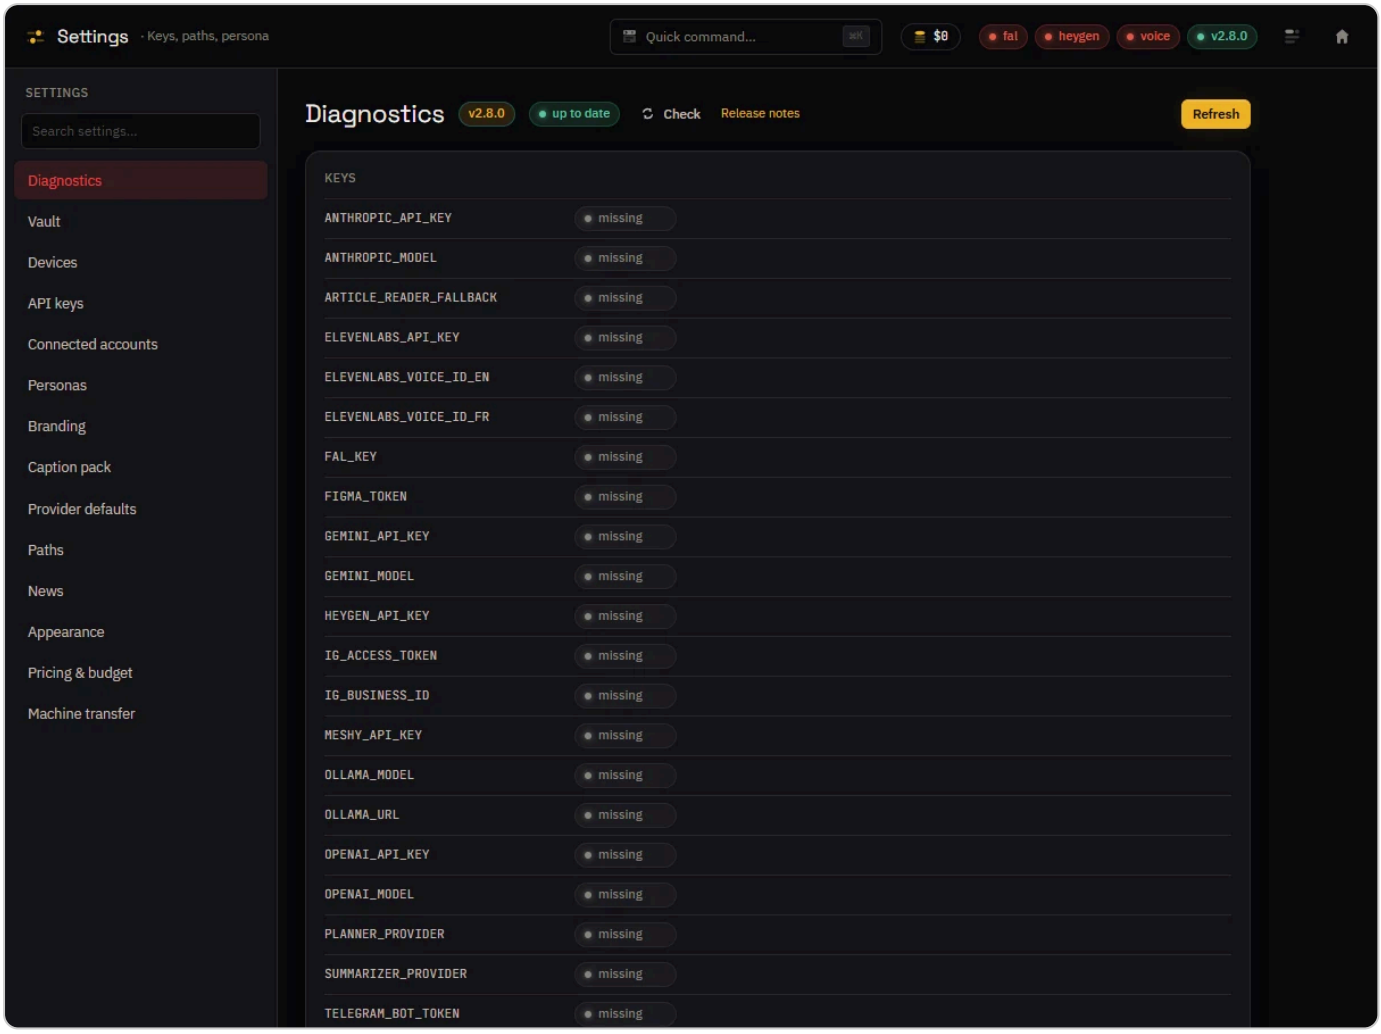

1. Your version is up to date 2. Check asks for the latest published version

The version block in Diagnostics: the installed version, its status, and a button to check right now.

1. Once a day, the app checks whether a new version is published. It never blocks anything.
2. When there is one, a card announces it: **Download** fetches the installer and shows where it is saved, **Later** hides the card for that version, **Notes** tells what changes.
3. To check right now: **Settings** › **Diagnostics**, **Check** button.
4. Close the app, run the downloaded installer: nothing is installed on its own.

Good to know

> Your keys, files and database live in the `%LOCALAPPDATA%\DeepotusVideoGenData` folder, apart from the program: they survive updates and reinstalls.

Start with the fal.ai key alone, then add avatars, voices and AI as you go. Set a monthly cap from day one: see Keep spending under control.

GET STARTED

Connect your keys and AI engines

You will set an AI provider's key, choose which engine does what, and protect your keys with a password. Ten minutes, free; keys are created with each provider.

Before you start

An API key is a code the provider gives you on its website, in your account. It lets the app work on your behalf; the provider bills you. Only one is required to start: **fal.ai**.

Set a key

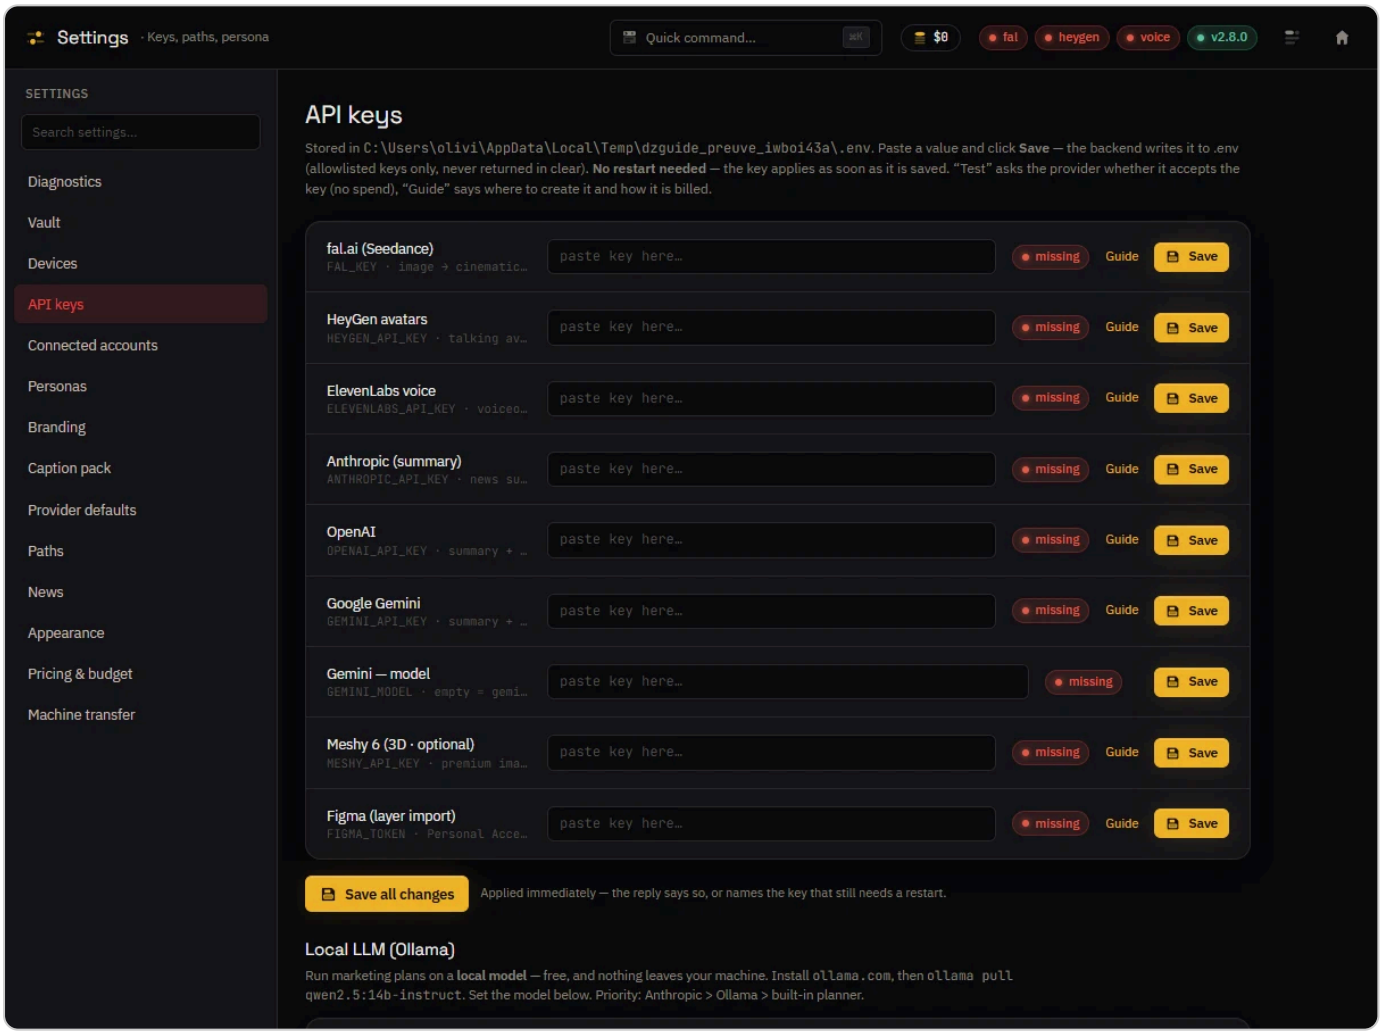

1. The key's status 2. Guide: where to create the key, and its price 3. Paste the key, Save: it works right away

Settings › API keys: each provider has its row, its status, its guide and its Save button.

1. Open **Settings** › **API keys**. Each provider has its row and its status: **missing** until a key is set.
2. Click **Guide**: it tells where to create the key with the provider and how it is billed.

3. Paste the key in the field, then **Save**. It works **right away**, no restart.
4. Click **Test** on the row: the provider says whether it accepts the key, without spending anything.

To set several keys at once, fill in the fields then **Save all changes**.

Which key unlocks what

Key	What it unlocks
fal.ai	Pictures (FLUX) and video clips (Seedance). The only key you need.
HeyGen	Talking avatars, alone or in a composition.
ElevenLabs	Voiceover and narration.
Anthropic, OpenAI or Gemini	News summaries, post plans, scripts and storyboards written by AI. One is enough; OpenAI also makes pictures (GPT Image).
Meshy	Premium picture-to-3D (optional, Meshy subscription).
Figma	Import Figma layers into the Library and Templates.
Ollama (no key)	An AI model installed on your PC for post plans: free, nothing leaves your machine.

Network keys (X, Telegram, YouTube, Instagram, TikTok) go into **Connected accounts**: see Connect your accounts.

Choose the engine for each job

In **Settings** › **Provider defaults**, each role has its engine: picture generation, voice, picture to video, avatar, voiceover, news summary, post planning. Only keys actually set are offered.

Default video model
The model new Seedance nodes start with in the Studio. Each node can change it.

Default voice model
The ElevenLabs model of new voiceover nodes.

Local AI with Ollama
Install Ollama, download a model, then enter its name in the **Local LLM (Ollama)** card of API keys.

Gemini: pick the version

Left empty, the “Gemini — model” field always follows the latest Flash version; fill it in to pin one.

Protect your keys in the vault

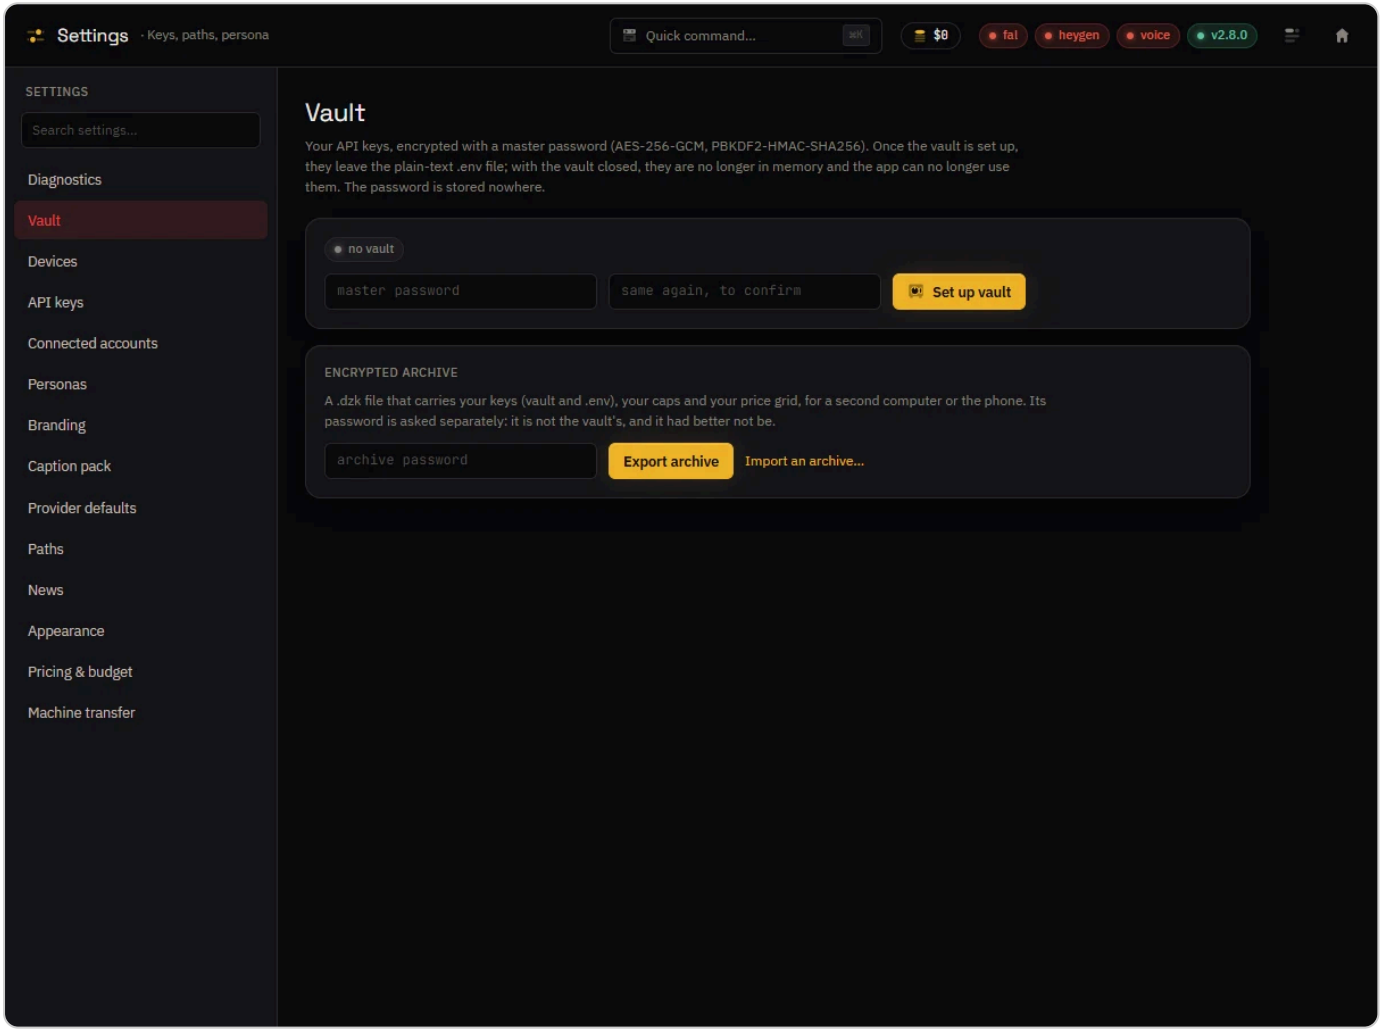

1. A master password encrypts all your keys 2. The .dzk archive carries them, with its own password

Settings › Vault: your keys encrypted with a password, and the archive to carry them to another computer.

1. Open **Settings** › **Vault**.
2. Choose a master password (at least 8 characters), type it twice, then **Set up vault**. Your keys are encrypted and leave the plain-text file.

3. **Close** removes the keys from memory: nothing paid can go out. **Open** makes them active again with the password.
4. **Remember on this PC** opens the vault on its own at launch, for your Windows session; **Stop remembering** asks for the password again.

On each row of **API keys**, a badge tells where the key lives: in the file, in the vault, or vault closed (one click takes you there).

Good to know

> ⚠ The vault password is stored nowhere: if lost, it cannot be recovered. Write it down, or keep an encrypted archive.

> ℹ **Export archive** creates a `.dzk` file that carries your keys, caps and rates to another PC or your phone, protected by its own password. **Import an archive…** takes them back.

> ℹ **Settings** › **Diagnostics** shows the status of all keys at once (never their value) and the providers' balances.

GET STARTED

Keep spending under control

You will know what you spend, set a cap per month and per engine, and decide yourself when a render may go over it. Five minutes, free.

Where to see what you spend

The dollar counter in the top bar shows the month's total. One click opens **Settings** › **Pricing & budget**, where everything is set.

Set your rates and caps

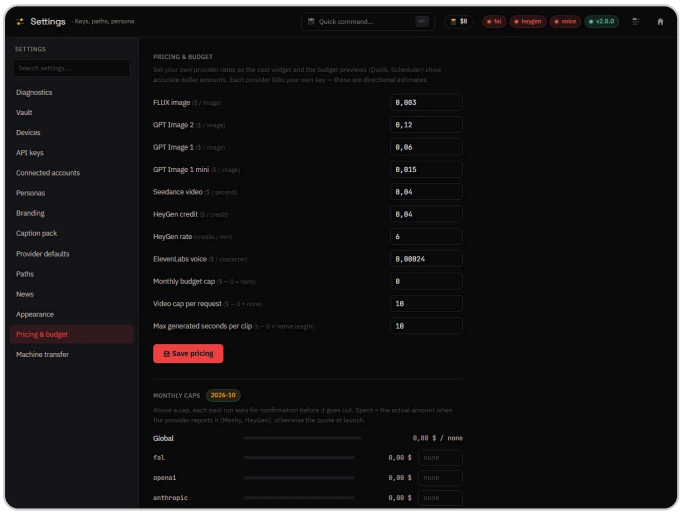

1. A cap for each video generation

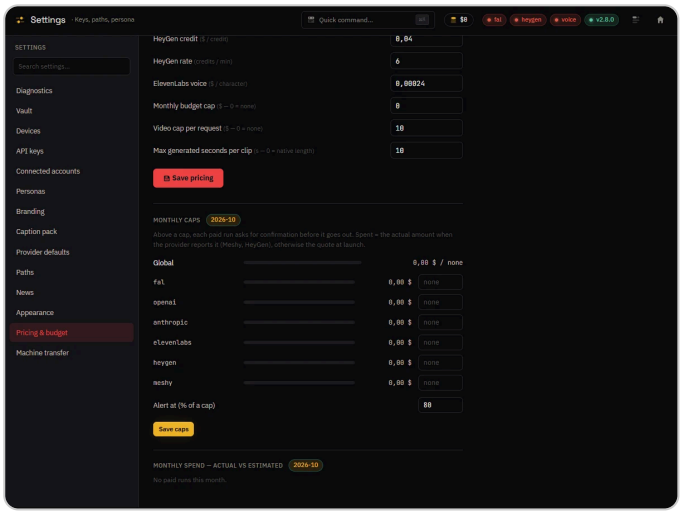

2. Above a cap, the app asks you before spending

Settings › Pricing & budget: your rates, then one cap per engine and the bar of what you spent this month.

1. At the top of **Pricing & budget**, enter your providers' rates (price of a picture, of a second of video, of a voice character…). The costs shown in Quick, the Studio and the Scheduler become accurate. **Save pricing**.
2. **Video cap per request** limits the cost of a single video generation; **Max generated seconds per clip** limits its length.
3. Further down, **Monthly caps**: one overall cap and one per engine (fal, OpenAI, Anthropic, ElevenLabs, HeyGen, Meshy), in dollars. A bar shows what was spent.
4. Set **Alert at (% of a cap)**: a banner warns you once a month when a cap reaches that level. **Save caps**.

When a render would go over a cap

Before leaving, every paid generation is checked against your caps. If it goes over, a “Monthly spending cap” dialog opens, wherever you are in the app:

1. “Tirer quand même” (run anyway) sends exactly the same request, once.
2. “Annuler” (cancel) sends nothing and tells you why.

Actual versus estimated

Below the caps, the **Monthly spend — actual vs estimated** table gives one row per engine and per screen: the number of runs, what was estimated, what was actually billed, and the difference.

Actual
The provider gave the exact price (Meshy, HeyGen): that is what counts.

Partly actual
Some runs have their exact price, the rest is still estimated.

Estimated
The provider gives no price: the app counts the quote made at launch.

Good to know

Spending made from your phone goes into the same record and the same caps, counted once.

The Studio sends the cost it displays as the request's cap: a render cannot cost more than what you saw before clicking.

GET STARTED

Settings, personas and branding

You will find any setting, give your content a voice with personas and dress the app in your colours with brand kits. Ten minutes, free.

Find a setting

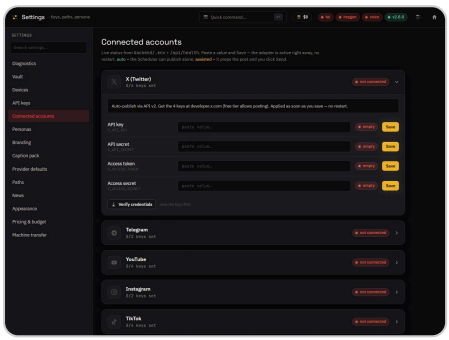

1. Type what you are looking for

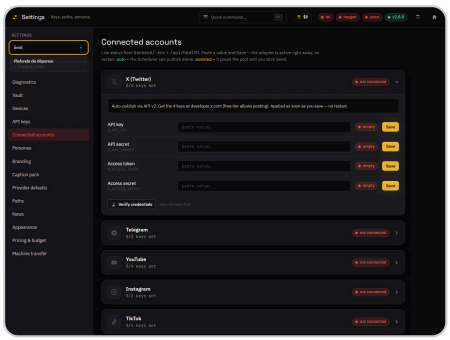

2. The matching sections show up

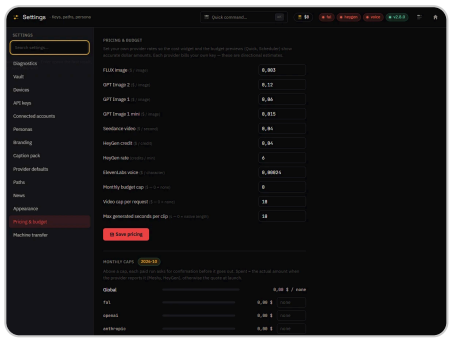

3. Enter opens the first one

Settings search finds a setting by its name and opens the right section directly.

1. Open **Settings** in the rail.
2. Type a word in **Search settings…**: accents and capitals do not matter.
3. Enter opens the first result, Esc clears the field.

The Settings sections

Section	What you do there
Diagnostics	Version, disk space used, latest alerts, key status and provider balances.
Vault	Encrypt your keys and carry them (chapter Connect your keys).
Devices	Pair a phone (chapter Drive it from your phone).
API keys	Set and test the AI providers' keys.
Connected accounts	Link X, Telegram, YouTube, Instagram and TikTok (chapter Connect your accounts).
Personas	The tone of your content.
Branding	Name, logo, colours: your brand kits.
Caption pack	One-click brand labels under the Scheduler's caption editor.
Provider defaults	The engine of each role, default models, language.
Paths	Where your pictures and renders are stored (button to copy the path).
News	Settings of the News screen.
Appearance	Reduced motion and visual effects.
Pricing & budget	Your rates and caps (chapter Keep spending under control).
Machine transfer	Take your work to another PC (chapter Move to another computer).

Give a voice with personas

A persona is a personality: its tone, its audience, its way of speaking. The active persona sets the News scriptwriter, the tone of the prompt generator and the default voiceover.

1. Open **Settings** › **Personas**.
2. Built-in personas can be read but not edited: duplicate one, or click **New persona**.
3. Click a persona to make it **active**.

Dress the app with a brand kit

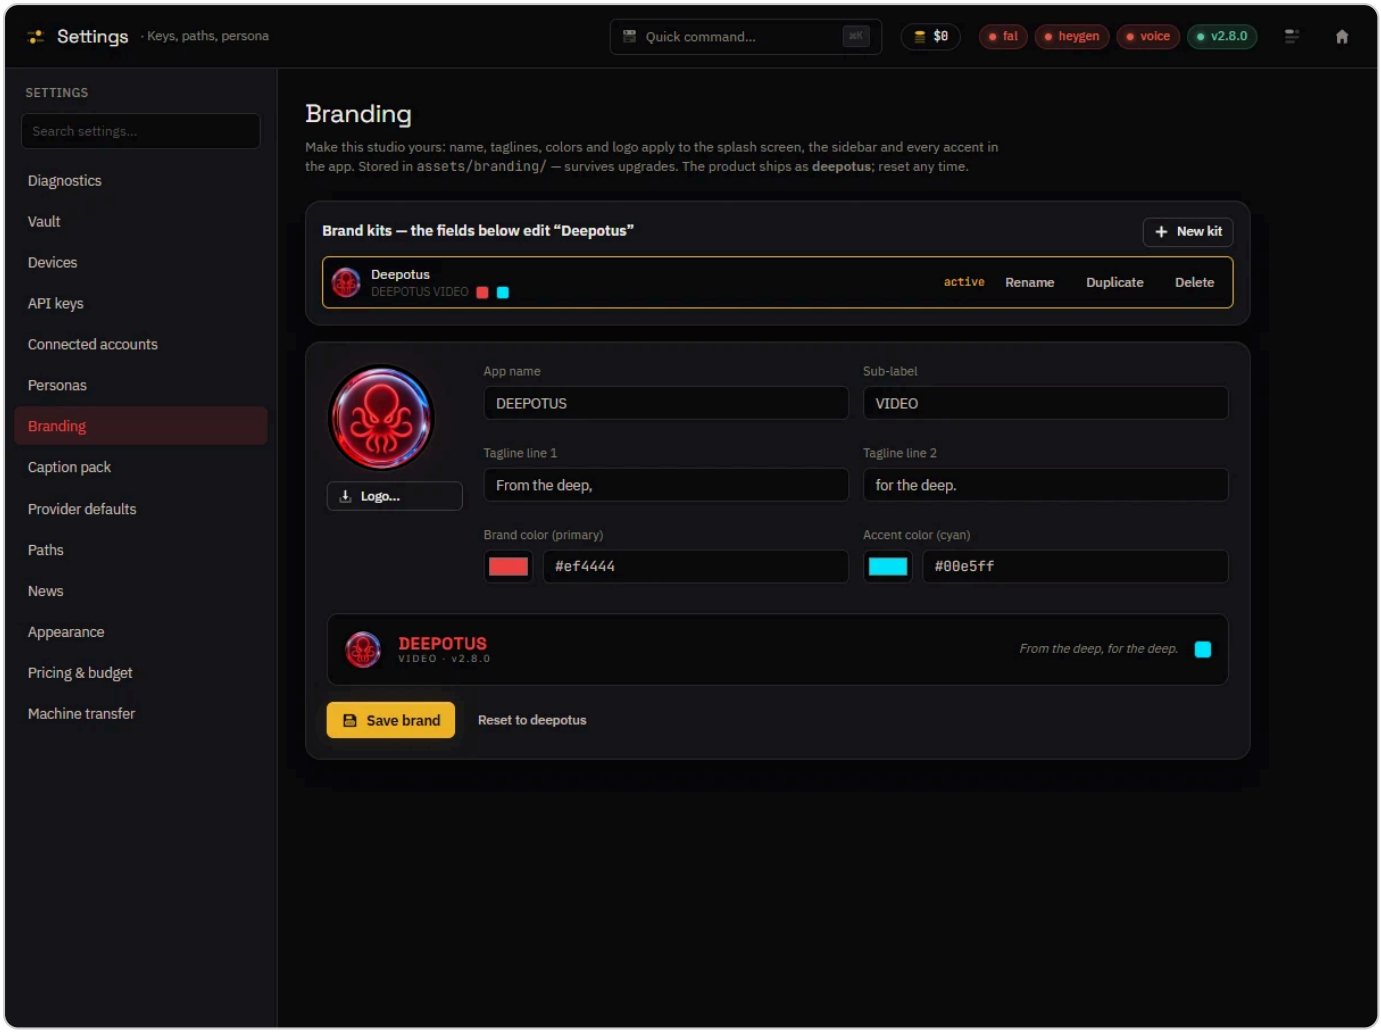

1. Create a kit 2. The active kit dresses the whole app

Settings › Branding: several brand kits, a single active one, and its fields right below.

1. Open **Settings** › **Branding**. The kit list shows which one is **active**.
2. **New kit** creates one; **Activate** makes it the active kit and refreshes the app.
3. The fields below edit the active kit: **Logo...**, app name, sub-label, two tagline lines, primary and accent colour. **Save brand**.
4. **Rename**, **Duplicate** and **Delete** manage each kit. The active kit and the last kit cannot be deleted: the greyed-out button says why.

The active kit dresses the splash screen, the sidebar, the app's accents and Template renders that use the brand. **Reset to deepotus** restores the original look.

Good to know

A Template render keeps the kit that was active when you sent it, even if you switch kits while it computes.

Settings › **Appearance** › **Reduced motion** turns off decorative animations; it also follows the Windows setting.

GET STARTED

Move to another computer

You will take all your work to another computer, check that nothing got damaged on the way, and get your keys back there. Fifteen minutes, free; the app must be installed on both machines.

What travels

The transfer takes everything the app created: the Library and the origin of every file, renders, the bible, communication plans, vector documents, card series and 3D models. Your API keys **never** go into the package.

Export from the old PC

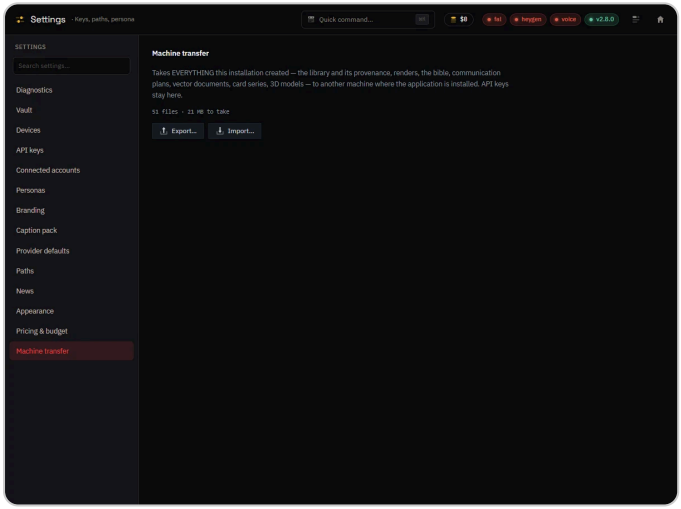

1. What goes, and how much it weighs

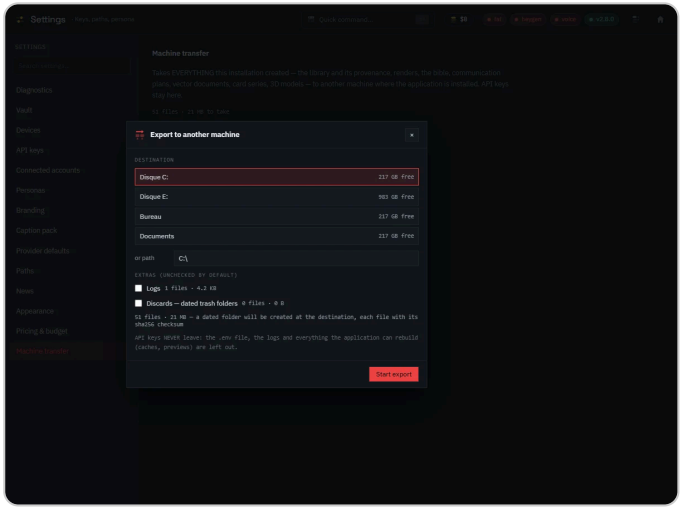

2. Extras to tick: logs and discards 3. A dated folder is written at the destination

Settings › Machine transfer: a dated package, each file with its checksum; keys stay on this PC.

1. Open **Settings** › **Machine transfer**: the number of files and the size to take are shown.
2. Click **Export...**, then pick the destination: a drive, a USB stick, or a path.
3. Tick the extras if you want them: **Logs** and **Discards — dated trash folders**. They are unchecked by default, with their size.
4. **Start export**: a dated folder is written at the destination, each file with its checksum. Do not close the app while it writes.
5. Click **Check integrity**: every written file is read back and compared with its checksum.

Import on the new PC

1. Install the app, then open **Settings** › **Machine transfer** and click **Import…**.
2. Give the package folder, then **Check**: the app reads the package without writing anything and tells where it comes from (date, machine, version).
3. **Start import**. Import **adds**: nothing this PC already has is overwritten. A damaged file is left out and named.

Take your keys too

1. On the old PC: **Settings** › **Vault** › **Export archive**, with a separate password.
2. On the new PC: **Import an archive…** and that password. Your keys, caps and rates come back.

Good to know

> Paths from the old machine are rewritten for the new one: your projects find their files without any setting.

> Stopping an export leaves an incomplete package. If it happens, run the export again, then **Check integrity** before importing.

GET STARTED

Drive it from your phone

You will link your phone to the PC, carry your posts and chapters in your pocket, and cut off a lost phone. Fifteen minutes; you need the Deepotus companion app on the phone and the same Wi-Fi.

Before you start

By default, the app only listens on this PC: no other device can reach it. To open it to your home Wi-Fi, replace `HOST=127.0.0.1` with `HOST=0.0.0.0` in the `.env` file of the data folder, then restart the app. Windows will ask to allow Python through the firewall.

Pair the phone

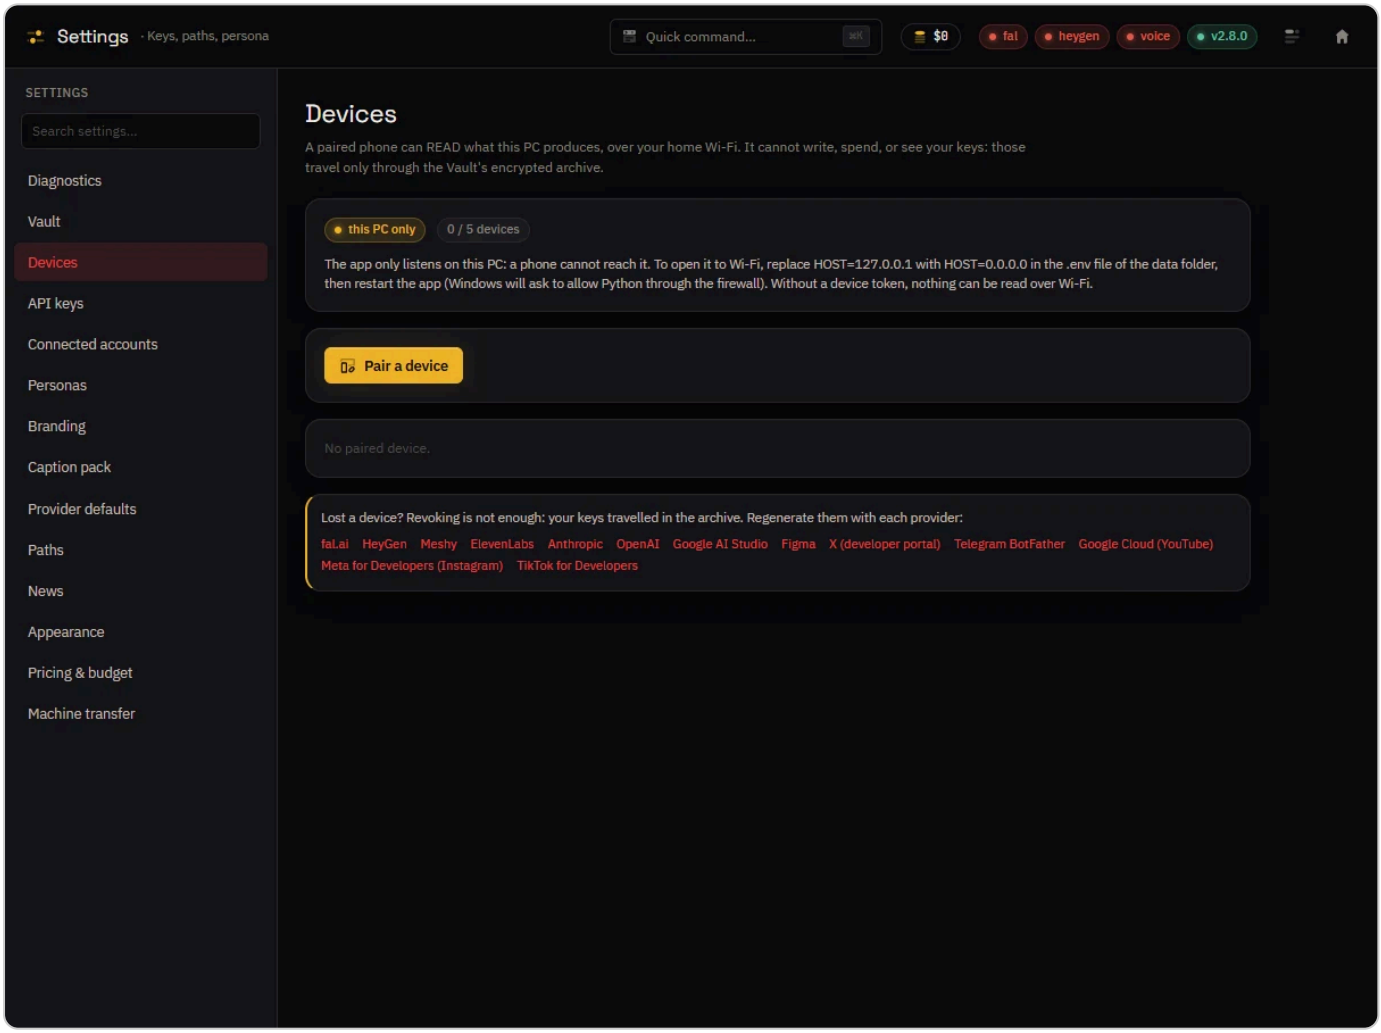

1. By default, only this PC can reach the app 2. Shows a QR code valid for 5 minutes

Settings › Devices: the network status, the list of devices, and pairing by QR code.

1. Open **Settings** › **Devices**. The words **this PC only** or **local network open** tell whether the phone will be able to reach you.
2. Click **Pair a device**: a QR code appears, valid for five minutes, for a single device.
3. In the phone app, open **Pair** and scan the QR code. It closes on its own once read.
4. To give your keys to the phone, export the encrypted archive from the **Vault** and open it in the phone app, with its password.

Five devices at most. A paired phone **reads** what the PC produces; it never sees your keys. It only writes to the PC to drop images and for Avatar live, whose spending goes through your caps.

What the phone does

Sync

The phone fetches the Library and its pictures, and resumes an interrupted download.

Send pictures to the PC

Your photos land in the Library with the "mobile" source, without duplicates.

The week's batch

Take the approved posts and publish them on Telegram or X from the phone. A post taken by the phone is no longer published by the PC.

Inbox

What you share to Deepotus lands here; pictures go on to the PC.

Offline chapters

Take a chapter, write without a network, hand it back to the PC.

Generate pictures

The cost is shown before each generation, at your PC's rates, and counted in your caps.

Avatar live

Film a take and start a Recast, or open a Live session: the PC does the work, the price is shown first (see Play an avatar).

A chapter taken to the phone

While a chapter is on the phone, it is read-only on the PC, with a banner. You can force it back to the PC. If the two versions disagree, the phone's text stays in the log, to copy or take back: nothing is lost.

Remove a device

1. In **Settings** › **Devices**, click **Revoke** on the phone's row and confirm.
2. Access is cut right away; the posts and chapters it had taken come back to the PC.

Good to know

> Lost your phone? Revoking is not enough: your keys travelled in the archive. Regenerate them with each provider; the Devices page gives the link to each console.

> The phone is notified when a render finishes or fails, when a post is published or awaits an action, and when a cap is near.

MAKE VIDEOS

Make a video in minutes (Quick)

You will make a vertical video in minutes: an animated clip from a picture, an avatar saying your text, both together, or a voiceover. Paid depending on the provider; the cost is shown before you click.

Before you start

Quick has four tabs: **Seedance** (animated clip), **HeyGen** (talking avatar), **Composition** (avatar and clip together) and **Voiceover**. A tab stays greyed out while its provider's key is missing: fal.ai for Seedance, HeyGen for the avatar (chapter Connect your keys). The voiceover can also run for free on your PC with Voicebox.

Animate a picture into a clip

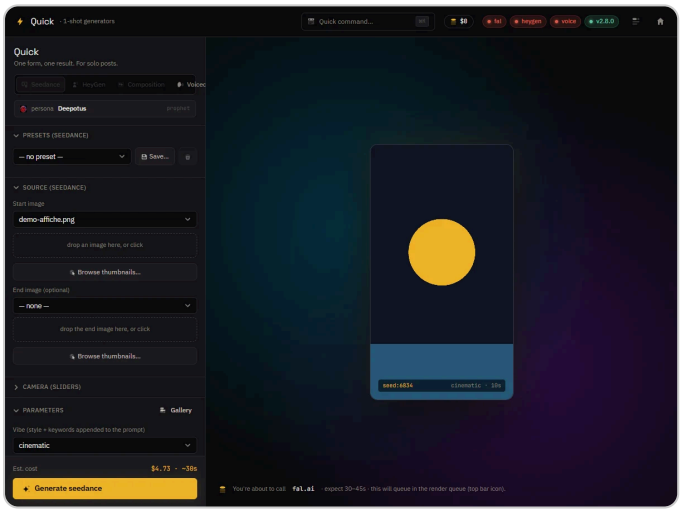

1. Drop the picture to animate

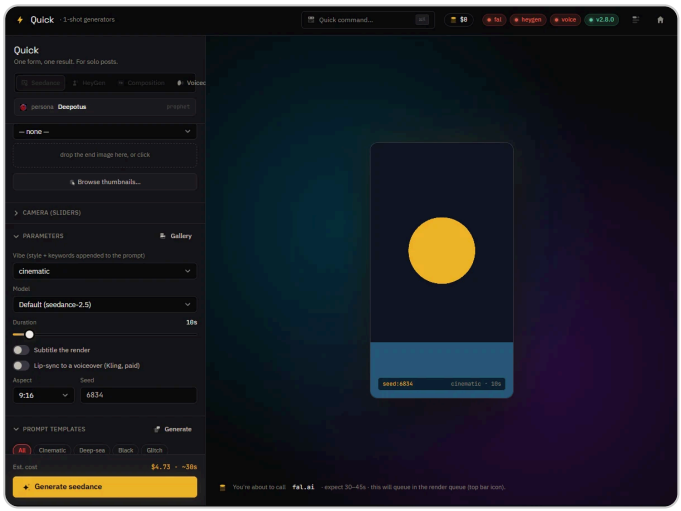

2. Model, duration and format 3. The cost is shown before you click

Quick, Seedance tab: the start image, the parameters, the estimated cost, then Generate.

1. Open **Quick**, **Seedance** tab.
2. Pick the **start image**: drop it on the area or click **Browse thumbnails...** to take it from the Library. An **end image** is optional; if the chosen model does not accept one, the field says so and names the models that do.
3. Write what should move, or start from a **prompt template**, or click **Open prompt generator**.
4. In **Parameters**, set the model, the duration and the format (9:16 by default).
5. Read the **estimated cost**, then click the **Generate** button. The render goes through the render queue, top right, and lands in the Library.

With the default model, Seedance 2.5 at 720p, one second costs about `$0.47`: a 10 s clip comes to about `$4.73`. A shorter duration or another model costs less; the estimated cost updates with every setting.

Direct the camera

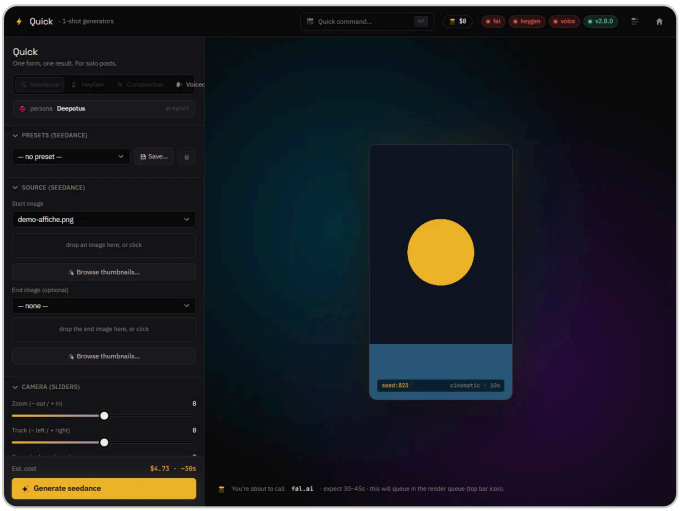

1. Six sliders, one camera sentence

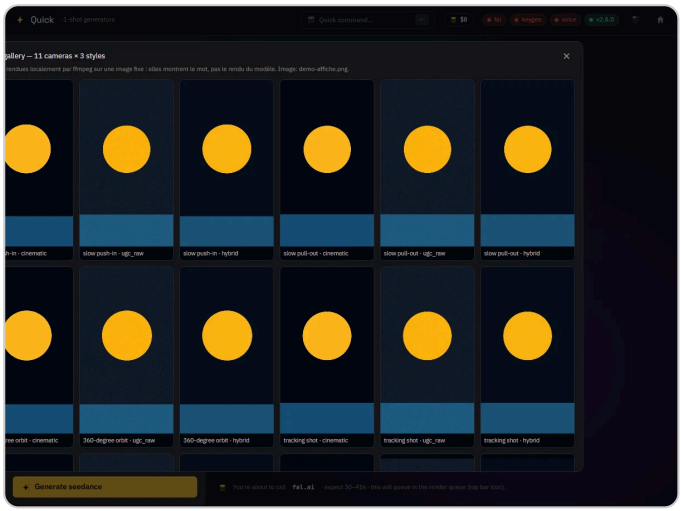

2. The Gallery shows each camera move

Camera sliders become a sentence added to the prompt; the Gallery shows the moves.

1. Unfold **Camera (sliders)**: six sliders (push in, turn, rise…) become a sentence suited to the model, added to the prompt. It is shown before the render.
2. To choose by eye, click **Gallery**: eleven camera moves in three styles, as small animations computed on your PC. It is a reminder, not a promise of what the model will do.

Make an avatar talk

1. **HeyGen** tab: pick an avatar and its voice. The lists show up right away, from a copy kept on your PC and refreshed in the background.
2. For an avatar that looks like you, import a photo: it becomes a talking avatar, selected automatically.
3. Write the script, then generate. The pauses in your text are kept.

Videos go through HeyGen and its Avatar III engine, picked according to the avatar. The HeyGen balance is shown in dollars. Count about `$0.24` a minute at default rates.

Compose and dub

Composition

An avatar and a clip in the same video, one after the other or side by side (vertical or horizontal split).

Voiceover

Choose the provider for each generation: Voicebox, on your PC and free, or ElevenLabs. A missing provider is refused, never silently replaced.

Music under the voice

When a render has both a voice and music, the music dips on its own under the voice.

Burned-in subtitles

Subtitle the render: aligning a text you provide is free; automatic transcription is paid and only runs if you tick it.

Go further

Presets

At the top of each tab, save your settings under a name and reapply them in one click.

Lip-sync

Syncs the clip's lips to a voiceover (Kling, paid); limits and price are shown beforehand.

Reopen in Quick

Every render keeps its settings: **Reopen in Quick**, from the Library or the render queue, puts them back as they were.

Extend a clip

From **Send to**, Extend adds about 7 generated seconds to a clip already rendered, with or without sound; the price is on each button.

Good to know

Leaving Quick and coming back finds your settings as you left them.

For a still picture, use the Library: **Create image**, with the **Prompt manager**. For your own filmed video, **Upload video** in the Library, then the UGC video node of the Studio.

MAKE VIDEOS

Compose with the Studio

You will compose a video piece by piece by linking nodes, see it before paying, never pay twice for what is done, and reuse a graph like a recipe. Paid depending on the nodes; the graph's cost is shown before running.

The Studio screen

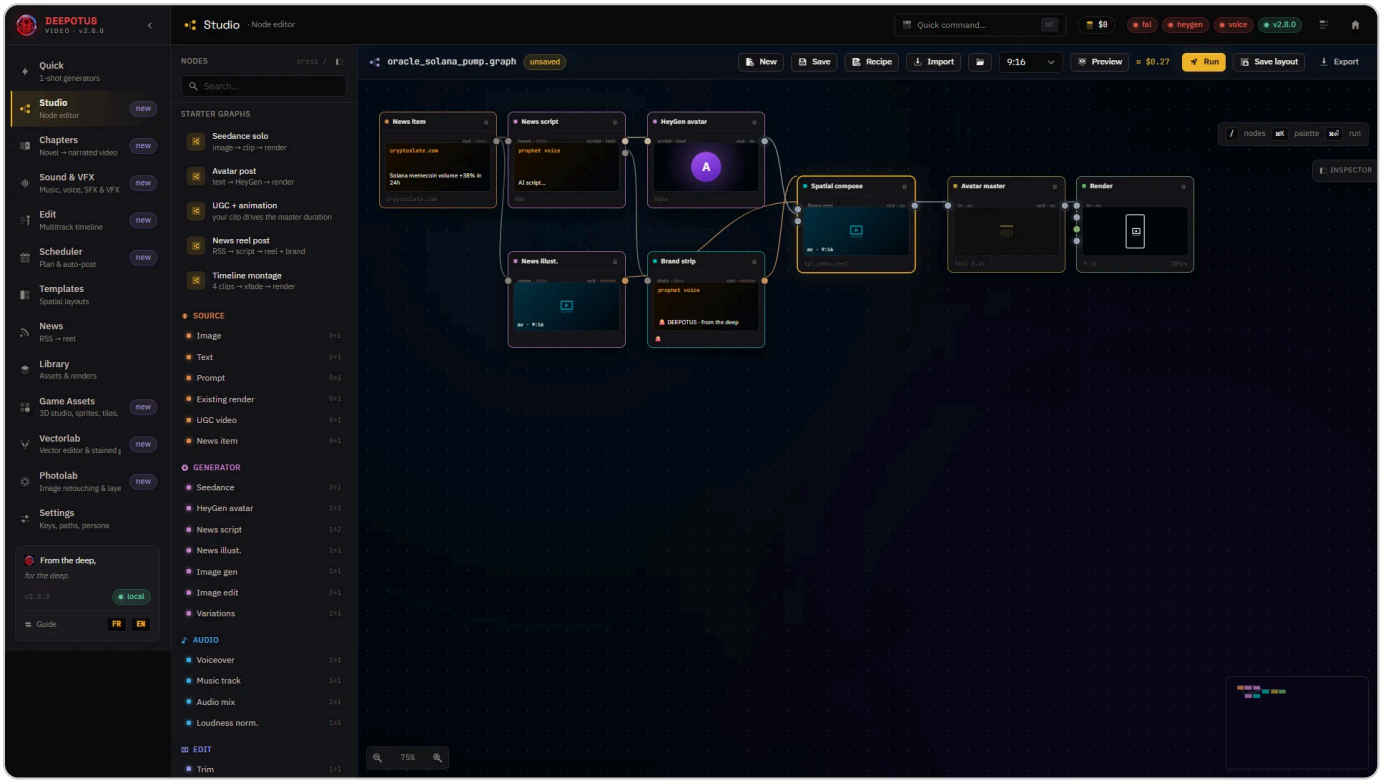

1. Start from a ready-made graph 2. Preview: see before paying 3. Run: the cost shown becomes the cap

The nodes palette, the graph, the inspector; at the top, the free preview, the cost and Run.

On the left, the **nodes** palette (`/` key) and the **starter graphs**. In the centre, the graph. On the right, the inspector of the selected node, which can be folded. At the top, the graph bar: **New**, **Save**, **Recipe**, **Import**, the format, **Preview**, the estimated cost and **Run**.

Compose a graph

1. Start from a starter graph (**Seedance solo**, **Avatar post**, **UGC + animation**, **News reel post**, **Timeline montage**) or from a blank page with **New**.
2. Drag nodes from the palette: sources (picture, text, your video), generators (Seedance, HeyGen avatar, pictures), audio, editing, composition, output.

3. Link an output, on the right of a node, to an input, on the left of another. Click a node to set it in the inspector.
4. **Preview** computes the final composition for free: clips not generated yet are replaced by their picture.
5. Read the estimated cost, then **Run**. That amount is sent as a cap: the render cannot cost more. **Save** keeps the graph under a name.

Never pay twice for what is done

The paid render of a Seedance or HeyGen node is **pinned** automatically: as long as its request does not change, it is reused for free, and the estimated cost says which nodes are.

The last eight renders
The inspector lists the node's last renders; **Use** picks the one that feeds the rest.

Regenerate
For a new paid take despite the pin, explicitly ask for the node to be regenerated.

Reuse a graph as a recipe

1. **Recipe** freezes the graph: its pictures and texts become blanks to fill in.
2. Run the recipe with new values: a quote is shown, you confirm. Unchanged nodes are reused for free.
3. From a picture in the Library, **Send to** also offers to run a recipe on it.

Go further

Engine duel

On a picture generation node, the same prompt goes to two models once you agree on the cost; you compare side by side and keep the winner.

Effects and masks

The Effects / Mask node, plugged into the Render's fx input: grades, VHS, colorisation…, on the whole render or a single layer.

AI fields

A field with a coloured edge is sent to an AI. Its badge shows the model and its price; its microphone dictates your text.

Voice and mix

The Voiceover node picks Voicebox or ElevenLabs; Audio mix dips the music under the voice (light, medium, strong).

Import a graph

Import opens a JSON graph; missing sources are listed and nothing is saved without your consent.

Frame by frame

The result drawer steps frame by frame with the , and . keys.

Good to know

A video generation above the announced cost or your per-request cap is refused before any call. The generated length is capped in Settings; the rest is extended on your PC.

New Seedance nodes start with Seedance 2.5; a saved graph keeps its model. The default model is changed in **Settings** › **Provider defaults**.

If the browser cannot dictate, the recording can be transcribed by the server: it is paid, a quote is shown and “No” is the default choice.

Shortcuts

Key	Effect
`/`	Opens the nodes palette
`Ctrl` + `Enter`	Runs the graph
`,` `.`	Previous frame, next frame in the result

MAKE VIDEOS

Reuse templates and brand kits

You will start from a ready-made layout, touch it up (slots, masks, texts, animations), adapt it to another format and dress it in your brand. Free; only renders of AI-generated slots cost money.

Pick a template

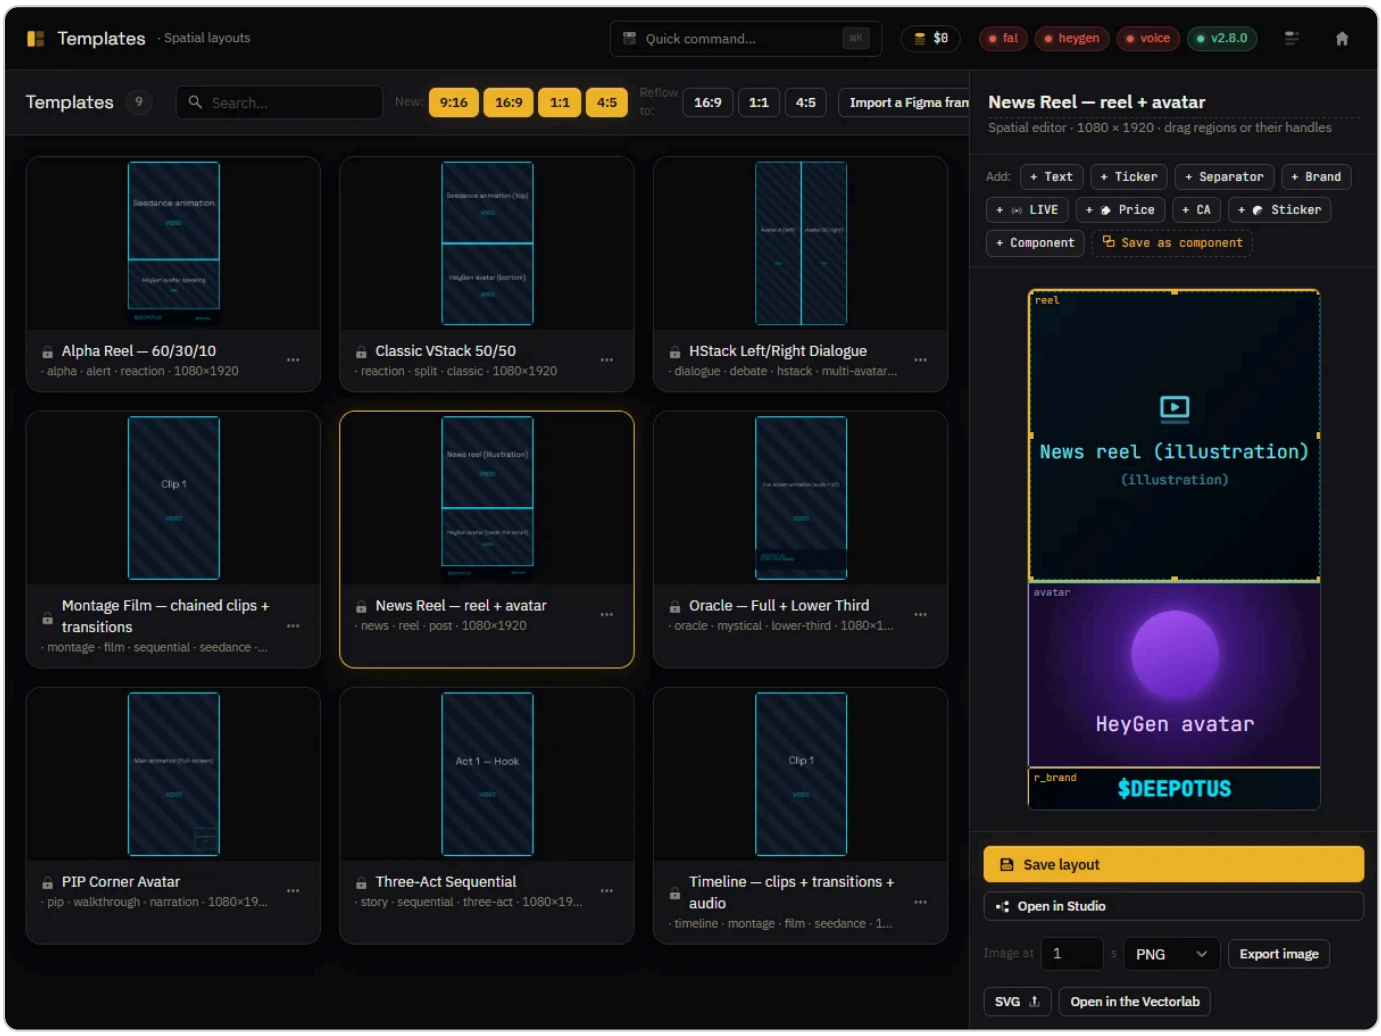

1. A blank layout in the format you want 2. The same layout, in another format

One card per layout; New in a format, or Reflow to another format.

1. Open **Templates** in the rail. The gallery shows a real picture of each layout; the search field filters by name or tag.
2. Click a card to open it in the editor, or create a blank one with **New** and a format (9:16, 16:9, 1:1, 4:5).
3. For the same layout in another format, **Reflow to** creates a copy with a before / after preview; the original does not change.

Touch up a layout

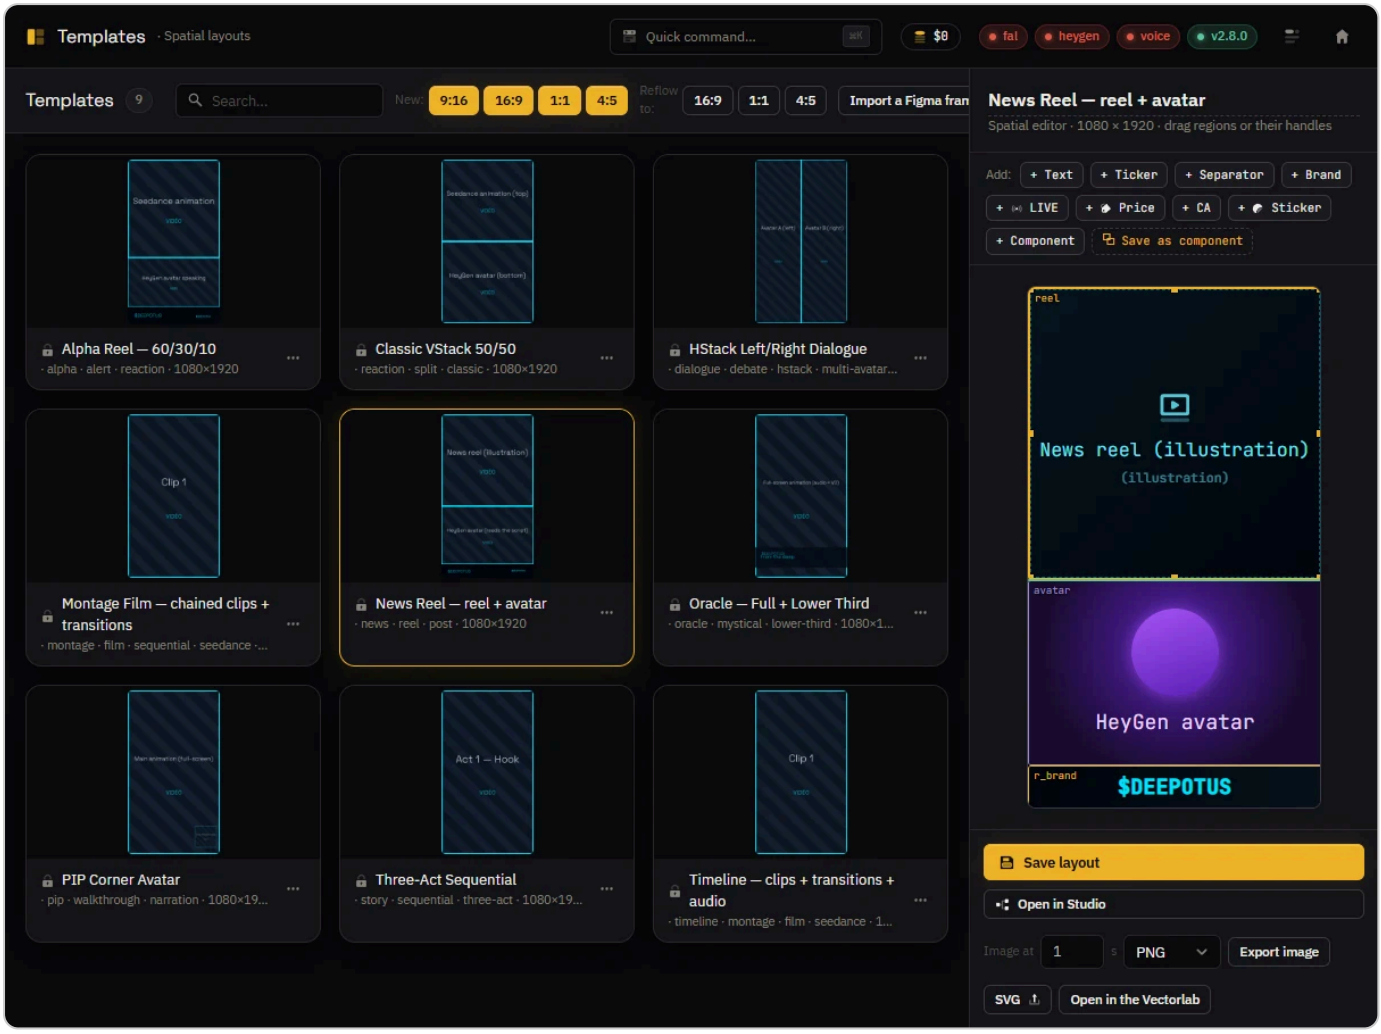

1. Add text, ticker, brand, component 2. Click a slot: drag it, resize it 3. Save, then open in the Studio

The slots on the canvas, and on the right the settings of the selected slot: mask, sample, animation.

1. Click a slot on the canvas: drag its body to move it, its handles to resize it. The inspector, on the right, shows its settings.
2. **Mask** gives a picture or video slot a shape: rounded, ellipse, polygon (triangle, diamond, hexagon, star, or your own points), cut-out windows, soft edge, coloured rim.
3. **Animation** brings the slot in and out with a fade, slide or pop, with duration, delay and curve. `Replay on canvas` shows it.
4. **Add** places a text, a ticker, a separator, the brand, a badge or a component.
5. `Save layout`, then `Open in Studio` to fill it in and render it.

Polish the texts

Ready-made styles

Social subtitle, neon, banner, gradient title; or your own settings: outline, shadow, rounded background, gradient.

Text that fits

The text stays in its slot: it wraps, then shrinks down to a minimum size.

Text on arc

Sets a title as an arch or a smile; an arc too large is reduced to fit.

Exact preview

Renders the slot with the real render engine, to see the final result.

Dress it with your brand

Templates use your **brand kit**: name, colours, logo. Kits are managed in `Settings` › `Branding` (chapter Settings, personas and branding). A render keeps the kit active when it was sent, even if you switch while it computes.

Go further

Components

`Save as component` turns slots into a reusable group; each instance can change its texts and colours, then go back to the component.

Preview sample

A Library picture fills a slot for thumbnails and picture export; never for the video render.

Export

`Export image` as PNG, JPEG or WebP (it lands in the Library), or as an SVG to drop into Figma.

To the Vectorlab

Open in the Vectorlab opens the layout as an editable vector document.

From Figma

Import a Figma frame… turns it into an editable template; the number of Figma calls is announced beforehand (Figma key required).

Avatar in the right format

A HeyGen avatar is generated in the format closest to its slot: no empty band, no cut-off head.

Good to know

In the Studio, the layout node shows the transition planned by the template as long as you do not pick another. The “cyan flash” transition is a real colour flash, of your choice.

MAKE VIDEOS

Tell a story in episodes

You will turn a text, a chapter or a novel into a narrated, illustrated video, scene by scene, and prepare its storyboard before paying for any render. Free up to narration and pictures; their cost is shown beforehand.

Two ways of working

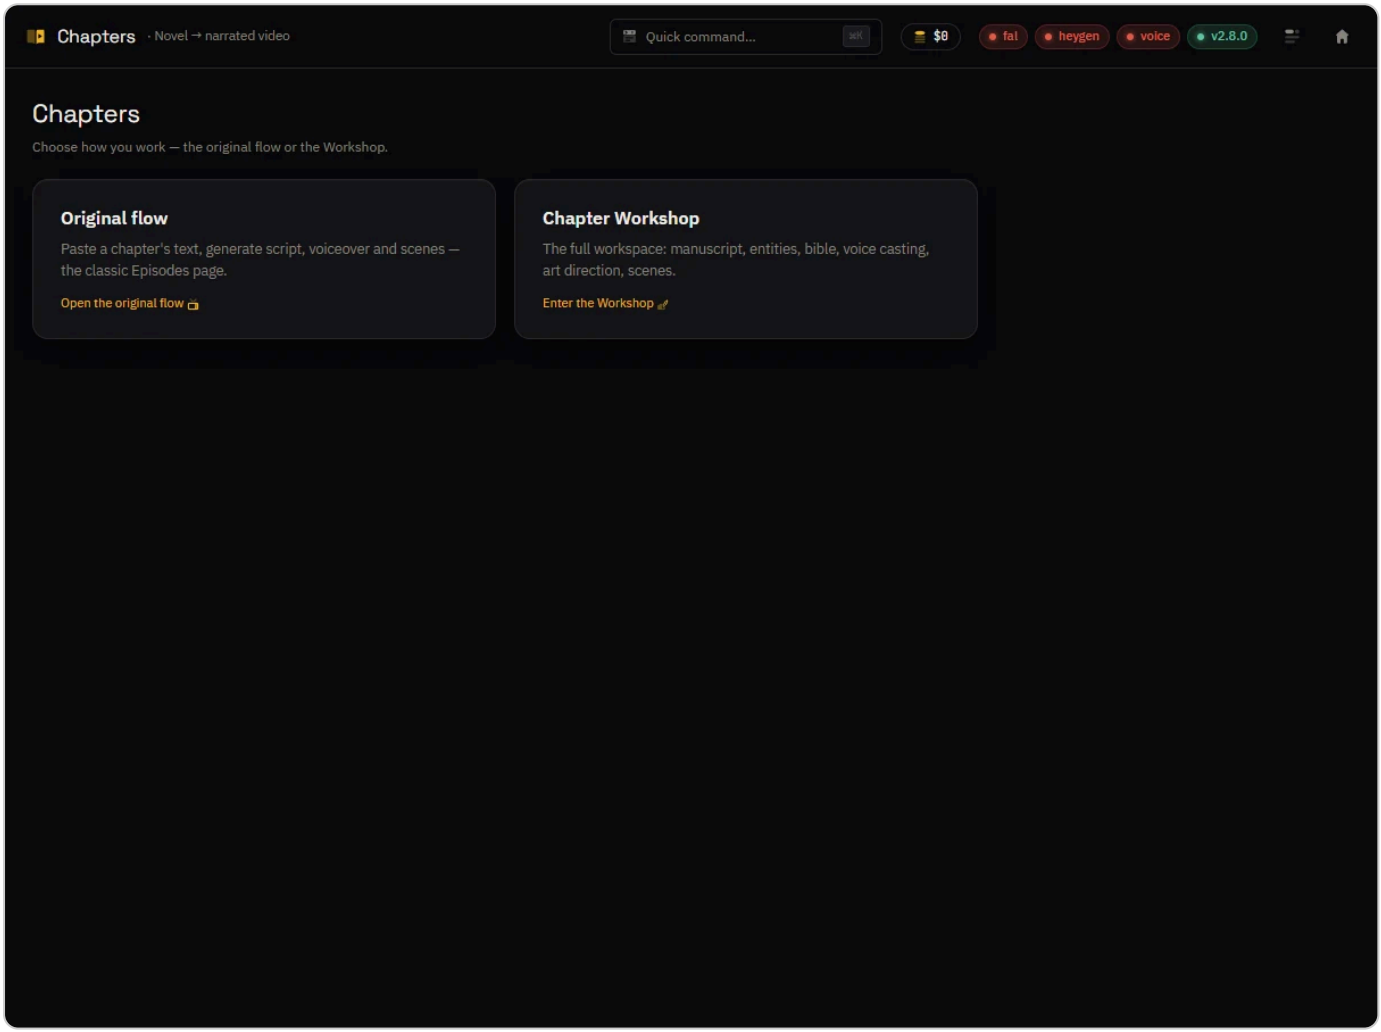

1. The original flow: from text to narrated episode 2. The Workshop: manuscript, bible, storyboard

Chapters offers the original flow (Episodes) or the full Chapter Workshop.

Open `Chapters` in the rail. The **original flow** (Episodes) goes straight from text to narrated video in four steps. The **Chapter Workshop** is the full space: manuscript, bible of characters and places, screenplay, storyboard, animatic.

An episode in four steps

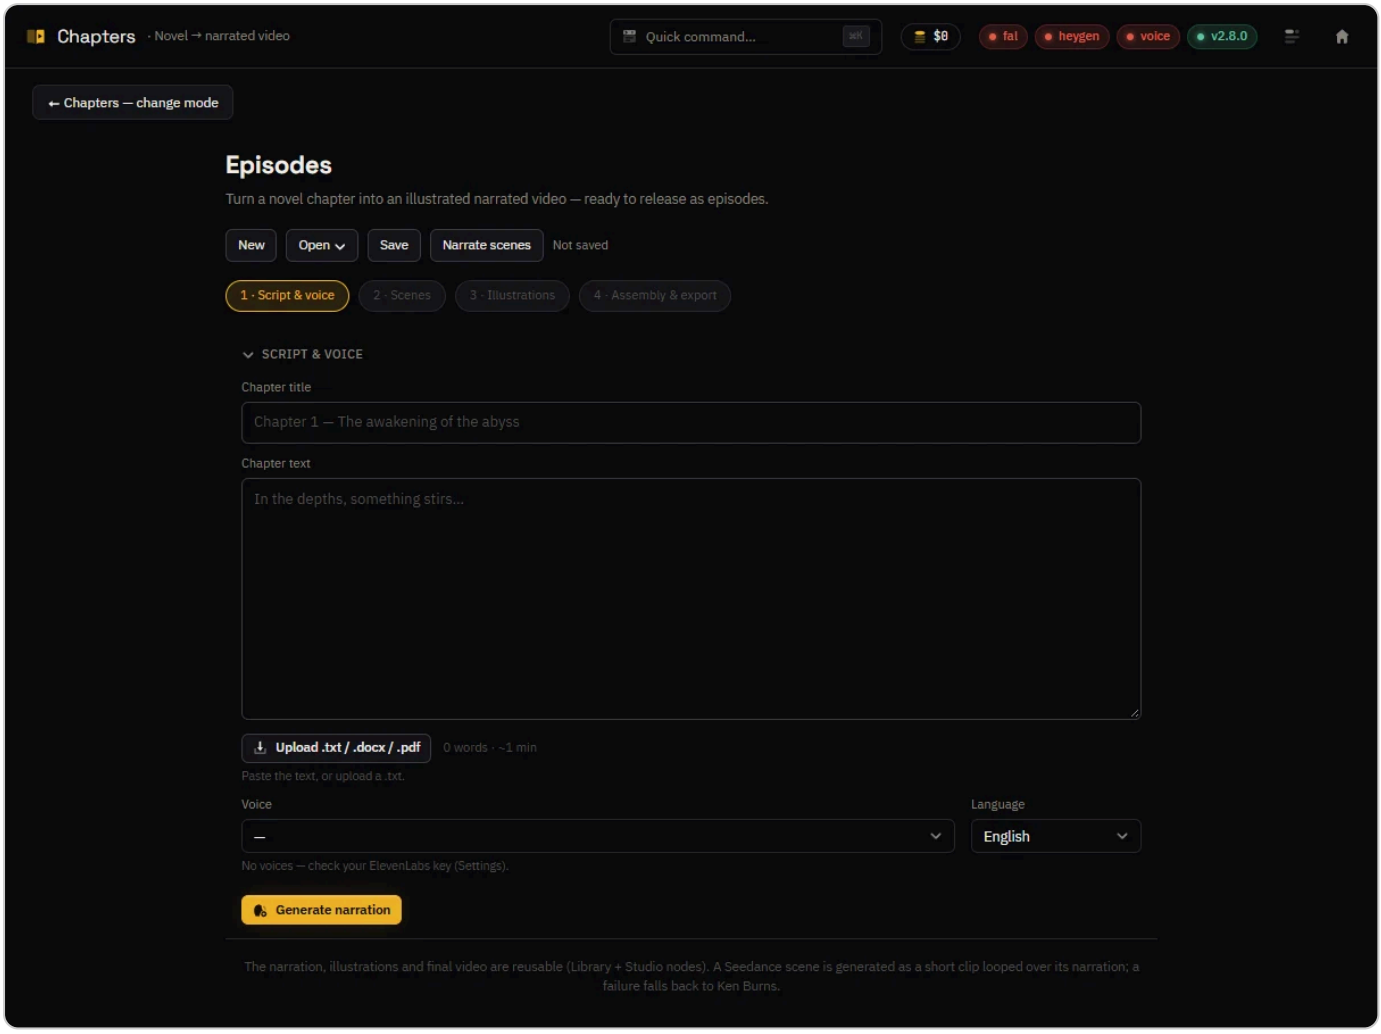

1. Four steps, left to right 2. Paste the text or upload a file 3. The voice reads the whole chapter

Episodes: script and voice, scenes, illustrations, assembly.

1. **Script and voice**: give a title, paste the text or upload a .txt, .docx or .pdf. Pick the voice and language, then `Generate narration`: the whole chapter is read and stored in the audio Library.
2. **Scenes**: split the text by paragraph (word for word) or by AI (into visual beats, with an illustration prompt). Edit, reorder, add or remove scenes.
3. **Illustrations**: one picture per scene, with the model of your choice. For each scene, pick the animation: slow camera move (Ken Burns), real Seedance clip or still picture.
4. **Assembly and export**: a quote gives the price of the Seedance scenes, then the episode is cut on its narration. If a Seedance scene fails, it falls back to a camera move and you are told.

At the top, `New`, `Open` and `Save` manage your episodes; a dot marks what is not saved. `Narrate scenes` never pays again for a scene already narrated with the same text, voice and language.

The Chapter Workshop

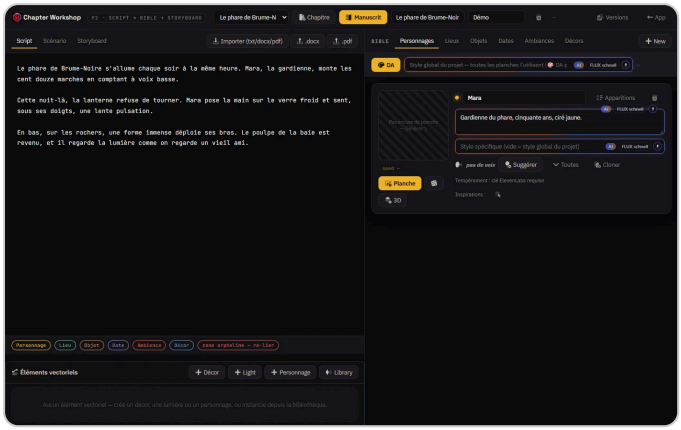

1. Three modes: script, screenplay, storyboard 2. The bible: characters, places, objects…

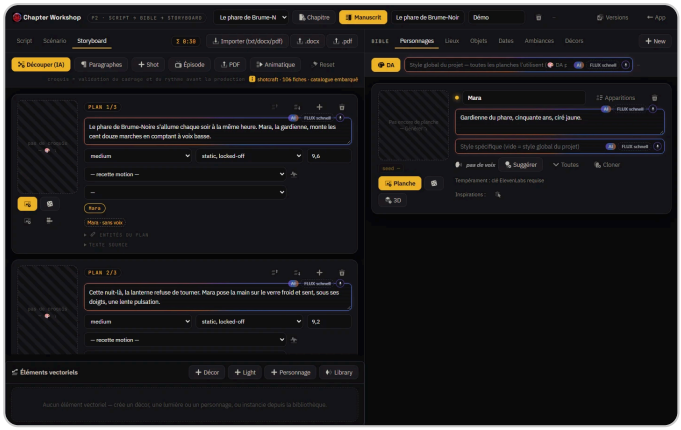

3. The storyboard, then the free animatic

The chapter text in the centre, the bible on the right; Script, Screenplay and Storyboard at the top.

The Workshop is not translated yet: its buttons below are quoted as they appear, in French.

1. **Script**: write or import the chapter text; `Manuscrit` imports a whole novel, split into chapters, and fills the bible. Recognised names are highlighted.
2. **Bible**, on the right: characters, places, objects, dates, moods, sets. Each entry has its description, its picture sheet and, for a character, its voice and temperament.
3. **Screenplay**: `Adapter (IA)` (adapt with AI) rewrites the chapter as a screenplay without touching the manuscript; `Voice-over` has each scene read by its character's voice.
4. **Storyboard**: `Découper (IA)` (split with AI) or `Paragraphes` (paragraphs) creates the shots, with their framing, camera, duration and the entities in them.
5. `Animatique` (animatic) cuts the shots into a rehearsal video, silent and free, to watch before paying for a render. It then goes into a new Edit project, as a full film or a 30 s reel.

Go further

Versions

Every time the text is overwritten, a copy is kept (ten at most): compare, restore, copy.

Exports

The manuscript as .docx or PDF, the screenplay in industry format, the storyboard as a PDF with four shots per page.

Import a screenplay

A Fountain or Final Draft file replaces the scenes (the old screenplay is kept) or is added after them.

Production picture

A shot's picture is generated with the characters' views as references; the cost is stated and confirmed.

Appearances

A bible entry lists its mentions, shots and scenes; one click opens the shot in the storyboard.

To an episode

Épisode (episode) creates a narrated episode (one shot, one scene) without paying anything; narration and render then start in Episodes.

Good to know

Rewriting a selection in the bible's tone, generating a production picture or adding the animatic's guide voice is paid: the cost is always stated and confirmed first.

A chapter taken to your phone is read-only on the PC, with a banner; see Drive it from your phone.

MAKE VIDEOS

Sound, voices and visual effects

You will generate music, sound effects and a voiceover you direct, store and find your sounds, and add free visual effects. Each generator's price is shown; visual effects and post-processing are free.

The Sound & VFX screen

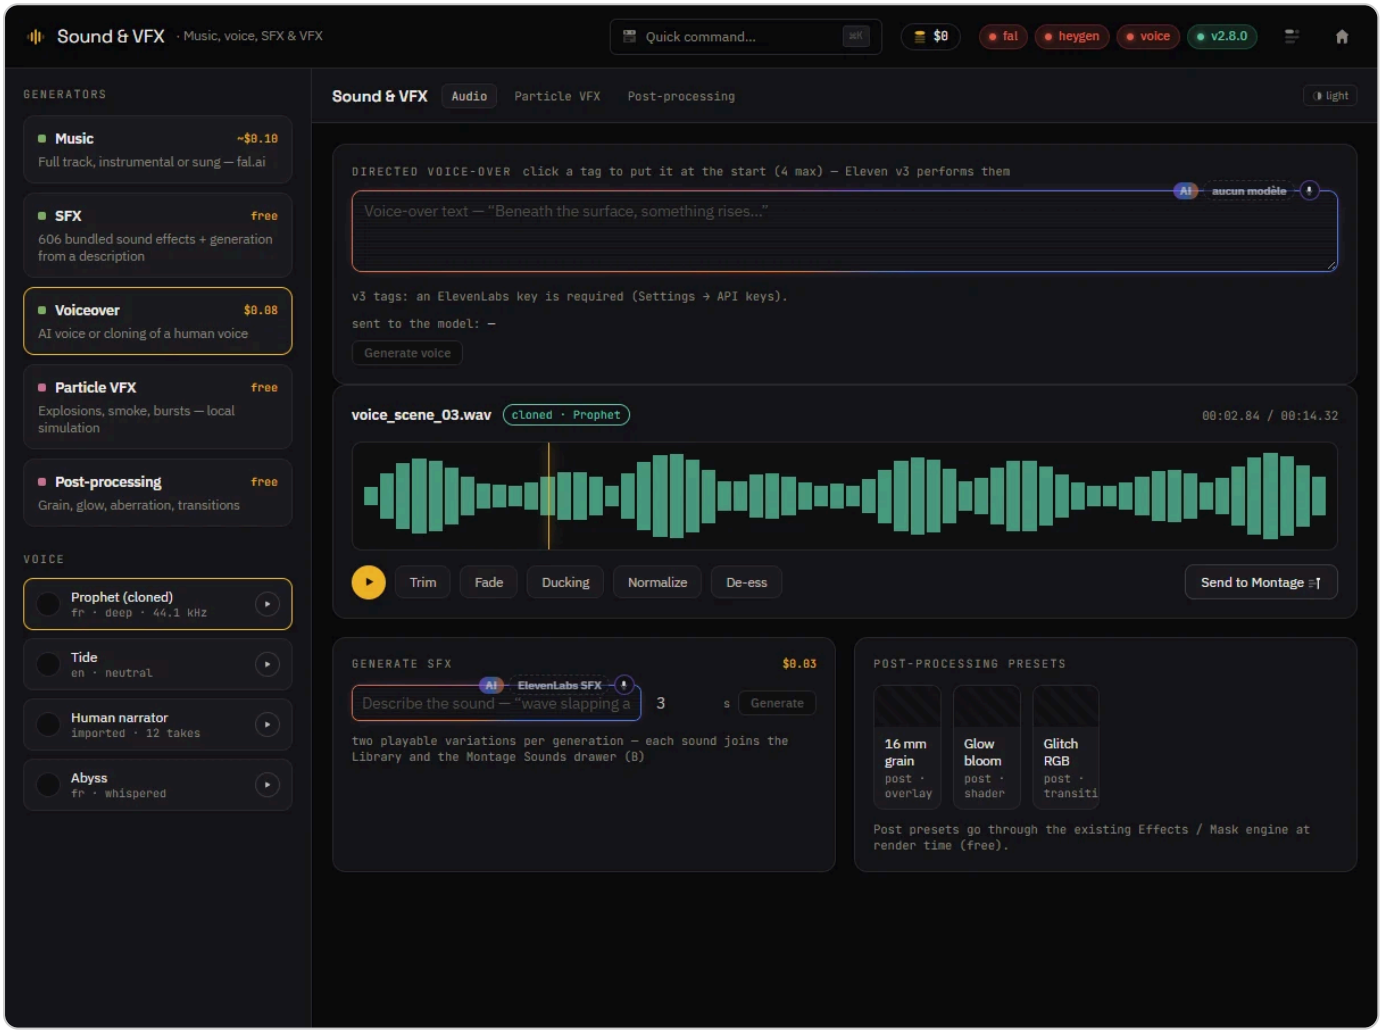

1. Music, effects, voice, particles, post: the price is shown 2. Direct the voice's acting

The generators at the top; directed voiceover, sound effects and post-processing below.

Open **Sound & VFX** in the rail. At the top, the **generators** and their price: music, sound effects (606 bundled sounds, plus generation from a description), voiceover, particles and post-processing. Below, three tabs: **Audio**, **Particle VFX** and **Post-processing**.

Direct a voiceover

1. In **Directed voice-over**, write the text.
2. Click emotion or acting tags (four at most): they go at the start and the Eleven v3 model performs them. The screen shows exactly what is sent to the model.
3. **Generate voice**. Trim, add a fade, dip the music under the voice or normalise, then **Send to Montage**.

Outside the v3 model, or with Voicebox, the tags are removed and you are told.

Music and sound effects

Music

Instrumental or sung: describe the style, write the lyrics section by section (a skeleton is suggested from the persona). Billed per second, under your caps.

Sound effects

Describe the sound and its length: two variations to listen to, which join the Library and the editor's Sounds drawer.

Voice and music mix

Mixes a voice and a piece of music, dipping the music under the voice; the mix lands in the Library.

A character's voice

In the Workshop, a character gets a cloned voice (paid) and a temperament, performed on each of its lines.

The Sounds drawer

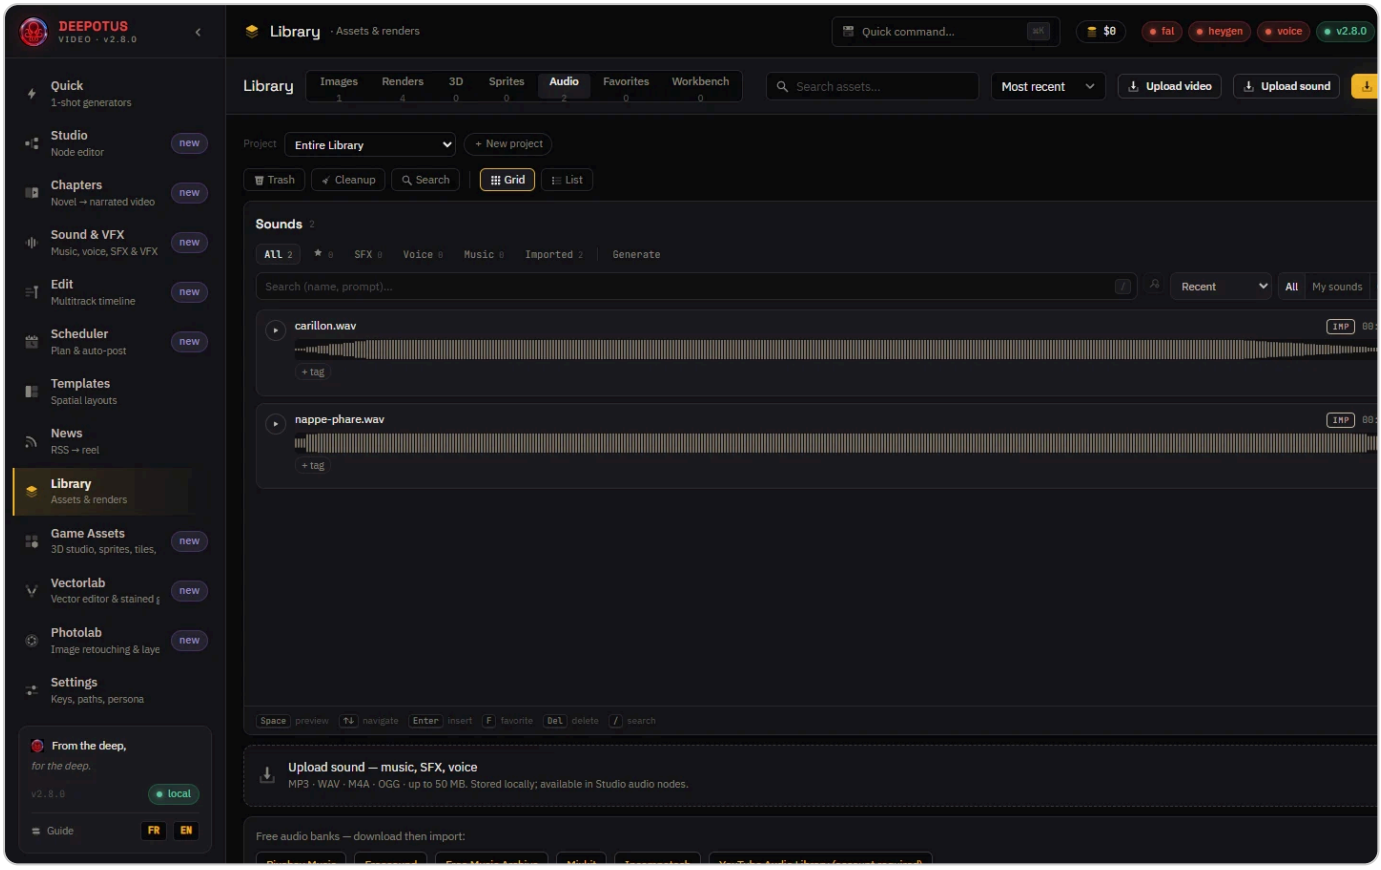

1. All sounds, then by kind: SFX, voice, music 2. Each sound: listen, tag, enhance, split

Library, Audio chip: the Sounds drawer, with its search, filters and actions.

All your sounds live in the Sounds drawer: in the Library (Audio tab) and in the editor (<kbd>B</kbd> key). Filter by kind (effects, voices, music, imported), between your sounds and the bundled catalog, search by name or tag.

Enhance

Equaliser, noise reduction, compressor and normalised loudness in one click, on your PC, no provider.

Split into stems

Splits music into vocals, drums, bass… (paid, quote on first click); each stem stays linked to its parent.

Isolate the voice

Removes the background from a voice take (paid, quote on first click).

Search by description

Describe what you want to hear; sounds similar to a sound can be listed too. Needs the local search service, otherwise the button says why it is greyed out.

Preview

Hover a sound to hear it (option to turn on); **Space** plays the selected row.

Free visual effects

1. **Particle VFX**: explosions, smoke, sparks, magic… simulated on your PC, to place on a shot.
2. **Post-processing**: 16 mm grain, glow, glitch, aberration; applied at render time by the Studio's Render node.
3. In the editor's effects rack, an effect can go **behind** the shot's cut-out subject, or in front of it. Cutting out the subject is paid; a first click gives the quote.

Good to know

Split, isolate, cut out: a first click arms the action and shows the price, a second one starts it, **Esc** disarms. Nothing goes out without that second click.

MAKE VIDEOS

Play an avatar, deferred or live

You will learn to create a Character from a few photos, have that Character replay a video in another setting, and film yourself on your webcam transformed live, with their voice. Twenty minutes. Everything is billed per second: the price shows before each launch and counts against your caps.

Before you start

The Avatar live screen opens from the rail. It needs keys depending on what you do: **fal** for Recast and the setting, **Decart** for Live, **ElevenLabs** for the Character's voice. The free local voice needs an NVIDIA graphics card and Voixbox (see Install and train a local RVC voice).

Only clone your own face and voice, or those of someone who has agreed to it. The app makes you confirm this for every Character and keeps the date.

Create a Character

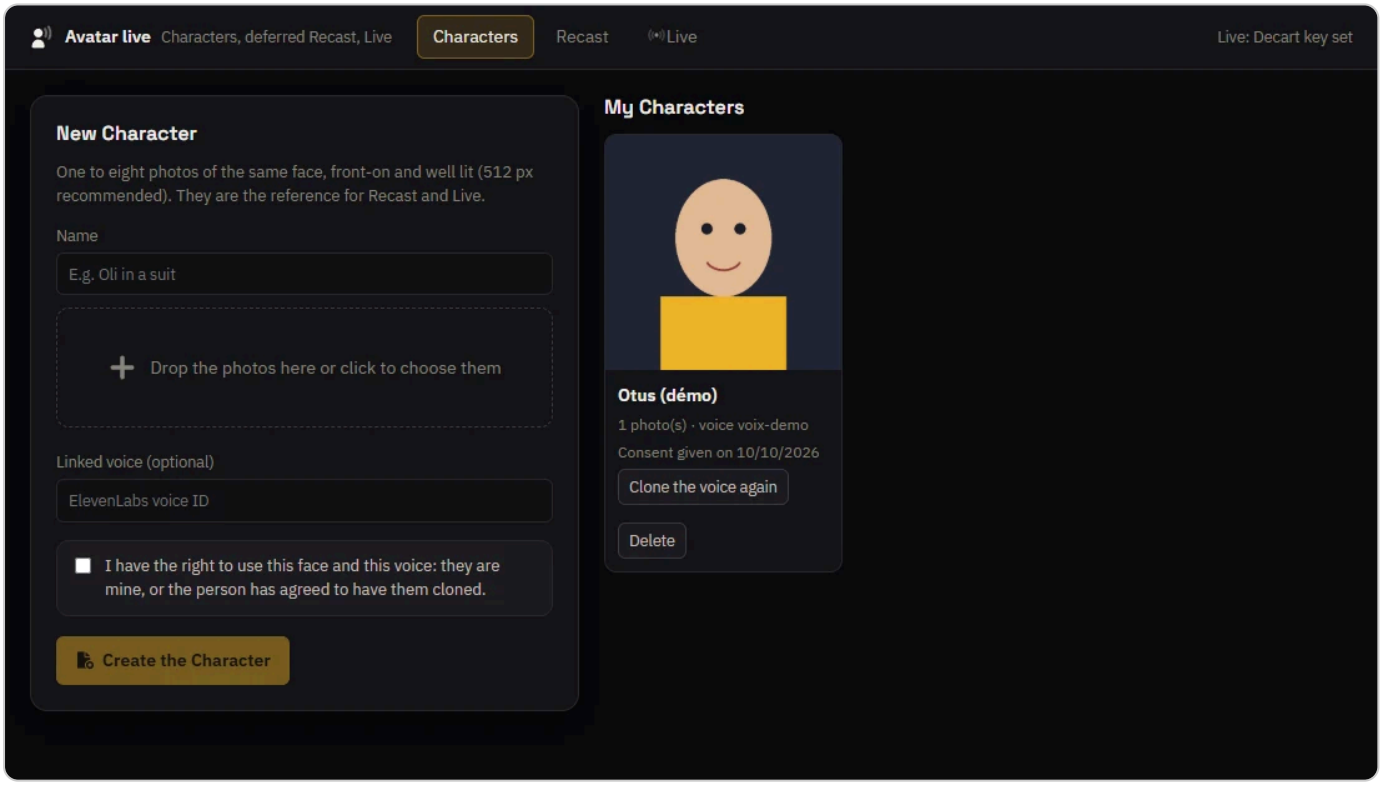

1. One to eight photos of the same face 2. Consent, dated, for every Character 3. Your Characters, ready for Recast and Live

Avatar live › Characters: reference photos, consent, then the list of your Characters.

1 Characters tab: give the Character a Name.

2. Drop one to eight photos of the same face, front-on and well lit (512 px recommended). They are the reference for Recast and Live.
3. Tick the consent box (“I have the right to use this face and this voice…”) then click **Create the Character**.
4. Optional: **Clone the voice** with one to five clips of the voice alone (30 s to 3 min in total). Cloning goes through ElevenLabs.

Recast: replay a video

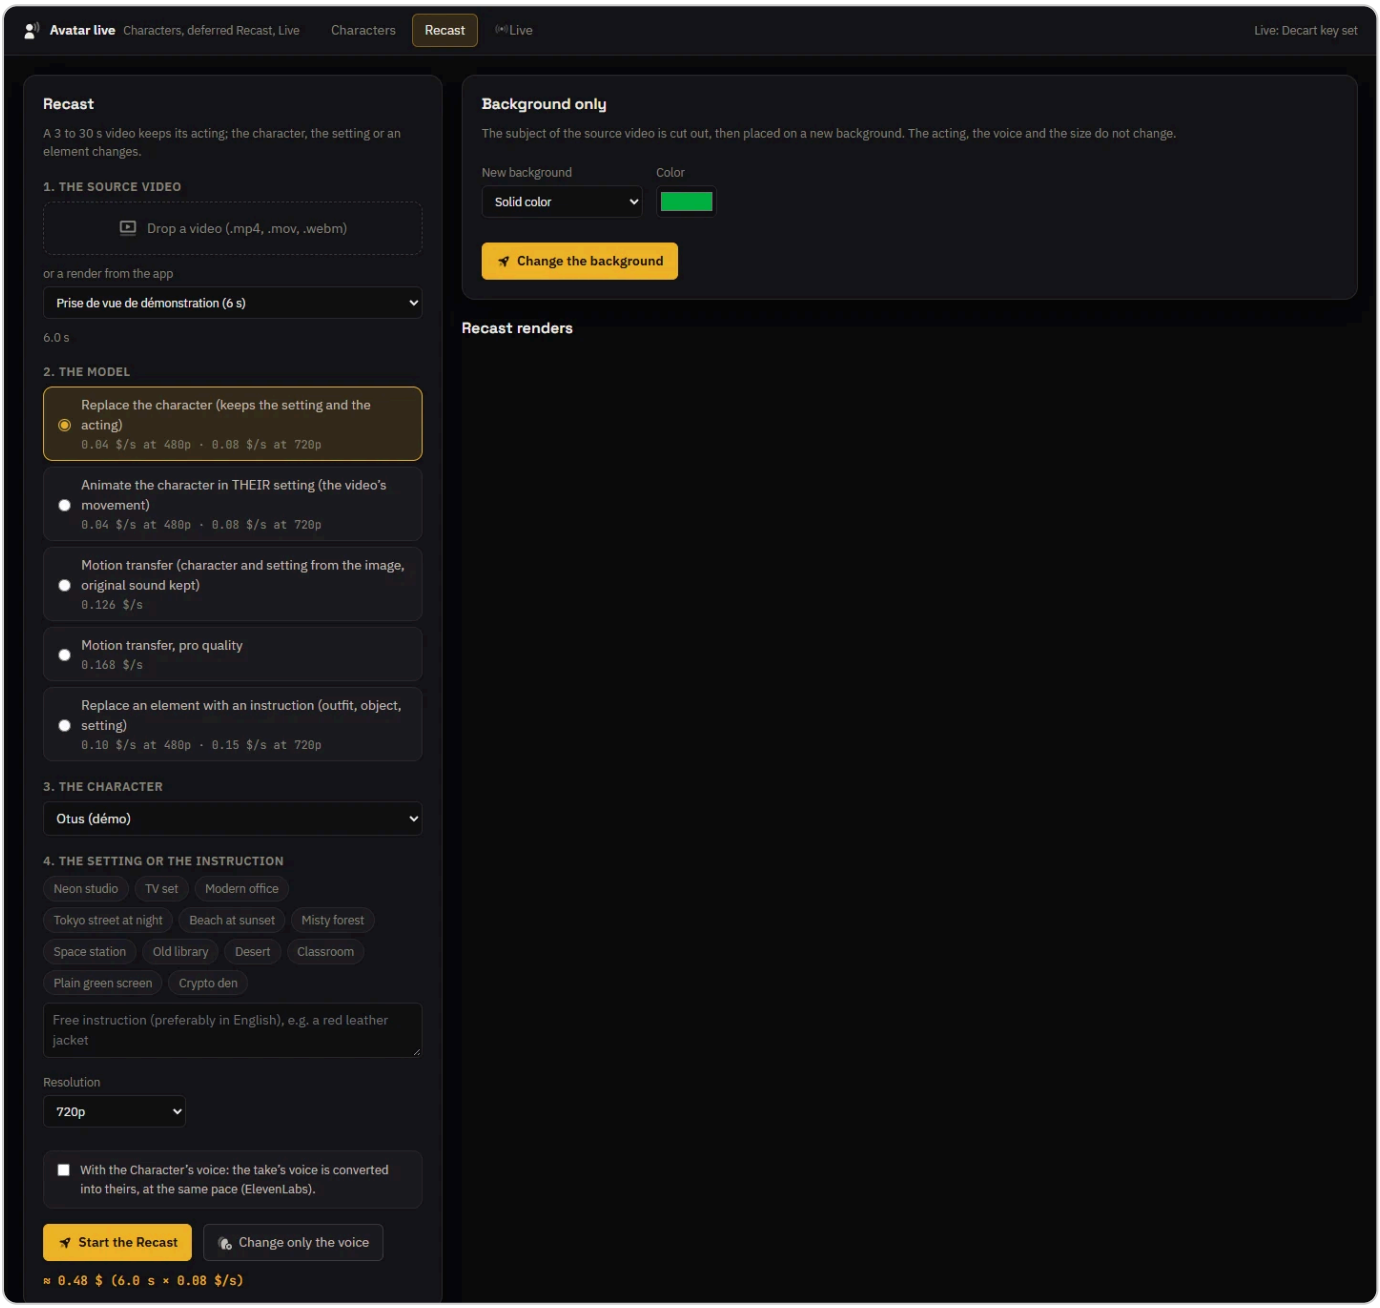

1. Each model shows its price per second 2. A ready-made setting, or your own instruction 3. The quote first, then the launch

Recast: a source video, a model and its price, the Character, the setting; the quote shows before you start.

A 3 to 30 s video keeps its acting; the Character, the setting or an element changes. The render takes a few minutes and joins your renders.

1. **Recast** tab: drop the source video, or pick a render from the app.
2. Choose the **Model**: its price per second is shown next to it.
3. Choose the Character, then a setting from the **Presets** or a free instruction (preferably in English).
4. Tick the Character's voice if you want the take's voice converted into theirs, at the same pace.
5. Read the quote under the button, then **Start the Recast**.

To try at half price, click **480p draft**: if you like it, “Finalize in 720p” on its card reruns the same Recast with the same seed. The final render is very close to the draft, without being identical. The motion transfer models have no draft.

Only one person is replaced per take. For the setting, **Suggest an instruction from the take** has the vision AI connected in your keys describe the place (never the people); read the instruction before you start.

Live: your webcam, transformed

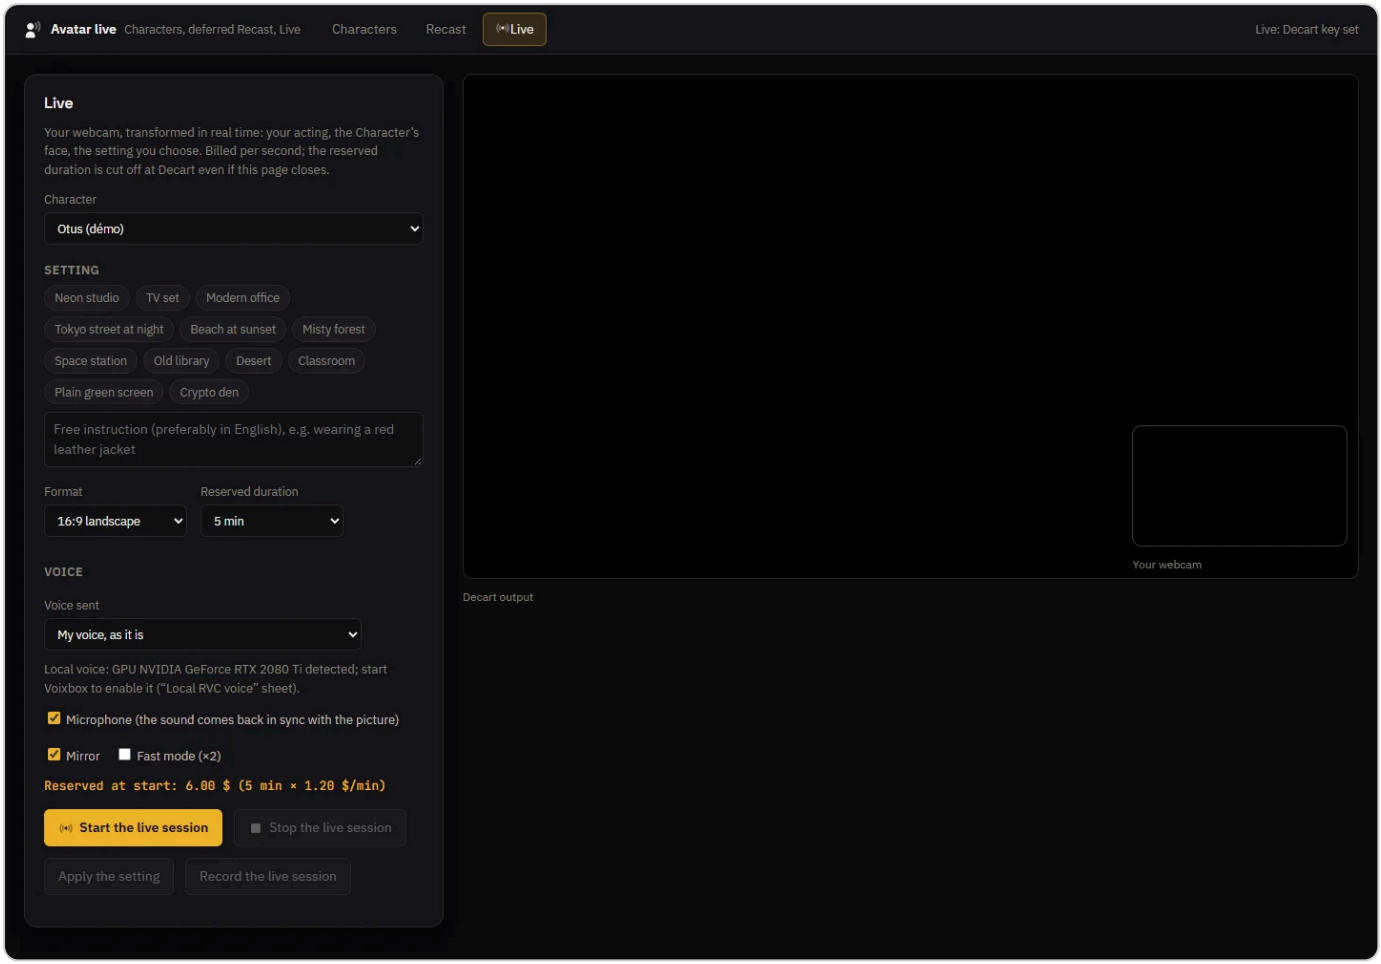

1. The setting, switchable mid-session 2. Your voice, the Character's, or a local RVC voice 3. The reserved price is shown before you start

Live: Character, setting, format, reserved duration and voice sent; the reserved price shows before you start.

1. **Live** tab: the first time, paste your Decart key and click **Save the key**. It goes to the vault; the page never sees it.
2. Choose the **Character** (or keep your face), a **Setting**, the **Format** and the **Reserved duration**.
3. Choose the **Voice sent**: your own as it is, the Character's through ElevenLabs (about 2 s delay), or a local RVC voice (under a second).
4. **Start the live session**. The transformed picture appears next to your webcam; switch settings on the fly with **Apply the setting**.
5. **Record the live session** keeps the footage in your renders, ready for the Editor or a Recast.

When a voice is converted, the picture is delayed by as much as the voice: lips and sound arrive together. The live bar shows the converted segments, the latency and any losses.

Going further

Background only

The subject of the video is cut out, then placed on a color, a Library image or a video. **Check the cut-out** first shows one picture on the new background; the video only goes after that.

Change only the voice

Keeps a video's picture and replaces its voice with the Character's. Billed at the ElevenLabs character equivalent.

In the Studio

The Recast and voice conversion nodes plug into an existing render or a UGC video. The graph quote counts their price before **Run**; the Recast node also has its 480p draft.

On the phone

The companion app starts a Recast (480p draft included, then its final version), films a take or opens a Live session, at your PC's prices and under the same caps.

Good to know

Live is reserved up front, by the minute; only the time actually played is billed. The reserved duration is cut off at Decart even if the page closes.

A tab hidden for 60 s stops the live session, so you don't pay for nothing. Prefer Edge or Chrome: they can delay the picture to match the voice.

MAKE VIDEOS

Install and train a local RVC voice

You will learn to install Voixbox, record your voice, train it with RVC and have it speak during a Live session, on your own graphics card. Free: nothing goes to the Internet. Allow thirty minutes to install, a quarter of an hour to record and about an hour of computing per voice.

Before you start

You need an **NVIDIA** graphics card with 6 GB of memory or more, and about 6 GB of disk. Without one, Live offers your voice as it is or the Character's voice through ElevenLabs (paid, nothing to install).

The engine is **RVC** (MIT license), through its official WebUI: the same install trains a voice and makes it speak. Only clone your own voice, or that of someone who has agreed to it.

Install Voixbox (once)

1. Open PowerShell in the app folder, where `tools\voixbox` lives (by default `%LOCALAPPDATA%\DeepotusVideoGen`).
2. Run `powershell -ExecutionPolicy Bypass -File tools\voixbox\installer.ps1`.
3. The script installs `uv`, downloads a verified version of the WebUI, creates its Python, adds torch for CUDA (about 3 GB) then its weights (560 MB).
4. It ends with `CUDA True | NVIDIA …`: the card is detected. On any failure, it stops and says why.

Record your voice

The quality of the cloned voice depends almost entirely on this recording.

1. Record 10 to 15 minutes of your voice alone: no music, no second voice, no echo.
2. A quiet room, the mic 15-20 cm away, always the same mic. Speak the way you do on your live sessions, with varied sentences.
3. Avoid clipping. Prefer several 30 s to 2 min files, in WAV.
4. Put them in a folder that holds nothing else, for example `C:\Voix\oli`.

Train an RVC voice

1. Run `powershell -ExecutionPolicy Bypass -File tools\voixbox\webui.ps1`. The WebUI opens in French at `http://127.0.0.1:7865`; keep the PowerShell window open.
2. « Entraîner » (Train) tab, step 1: a short **experiment name** with no spaces or accents (`oli`), rate `40k`, pitch guidance **Oui** (yes), version `v2`.
3. Step 2: your recordings folder, and the `rmvpe` pitch algorithm.
4. Step 3: save every `25` epochs, **200** epochs in total (100 for a first try), batch `8` (or `4` with 6 GB), small final model at each save: **Oui** (yes).
5. Click « Entraînement en un clic » (one-click training). Allow about 1 h 15 for 200 epochs on an RTX 2080 Ti.
6. Put the voice where Voixbox looks for it: `powershell -ExecutionPolicy Bypass -File tools\voixbox\ranger-voix.ps1 -Nom oli`.

Listen to it before using it: in the WebUI's « Inférence du modèle » (model inference) tab, pick `oli.pth` and give it a short recording. A crackling voice needs more epochs or a cleaner recording.

Speak with it during a Live session

1. Run `powershell -ExecutionPolicy Bypass -File tools\voixbox\lancer.ps1` and keep the window open: it reports how many voices are trained.
2. In **Avatar live**, **Live** tab, choose the local RVC voice under **Voice sent**.
3. Set **Transpose (semitones)** if needed: `+12` to go from a male to a female voice.
4. **Start the live session**. Each half-second of voice is converted in about 110 ms; the picture is delayed by 0.75 s to stay in sync with the lips.

Good to know

> The Live screen tells you what is missing: no NVIDIA card, Voixbox not started, or no trained voice. Another port is set with `VOIXBOX_URL` in the `.env`.

> ⚠ With Avast or an antivirus that inspects connections, always go through the scripts in `tools\voixbox`: they remove the `SSLKEYLOGFILE` variable for themselves only and read Windows certificates. Otherwise Python crashes (`no OPENSSL_Applink`).

MAKE VIDEOS

Edit a video

You will place your shots on a timeline, cut them, move them and chain them, with the gestures of a professional editor. Free: editing and previews run on your PC.

The editor screen

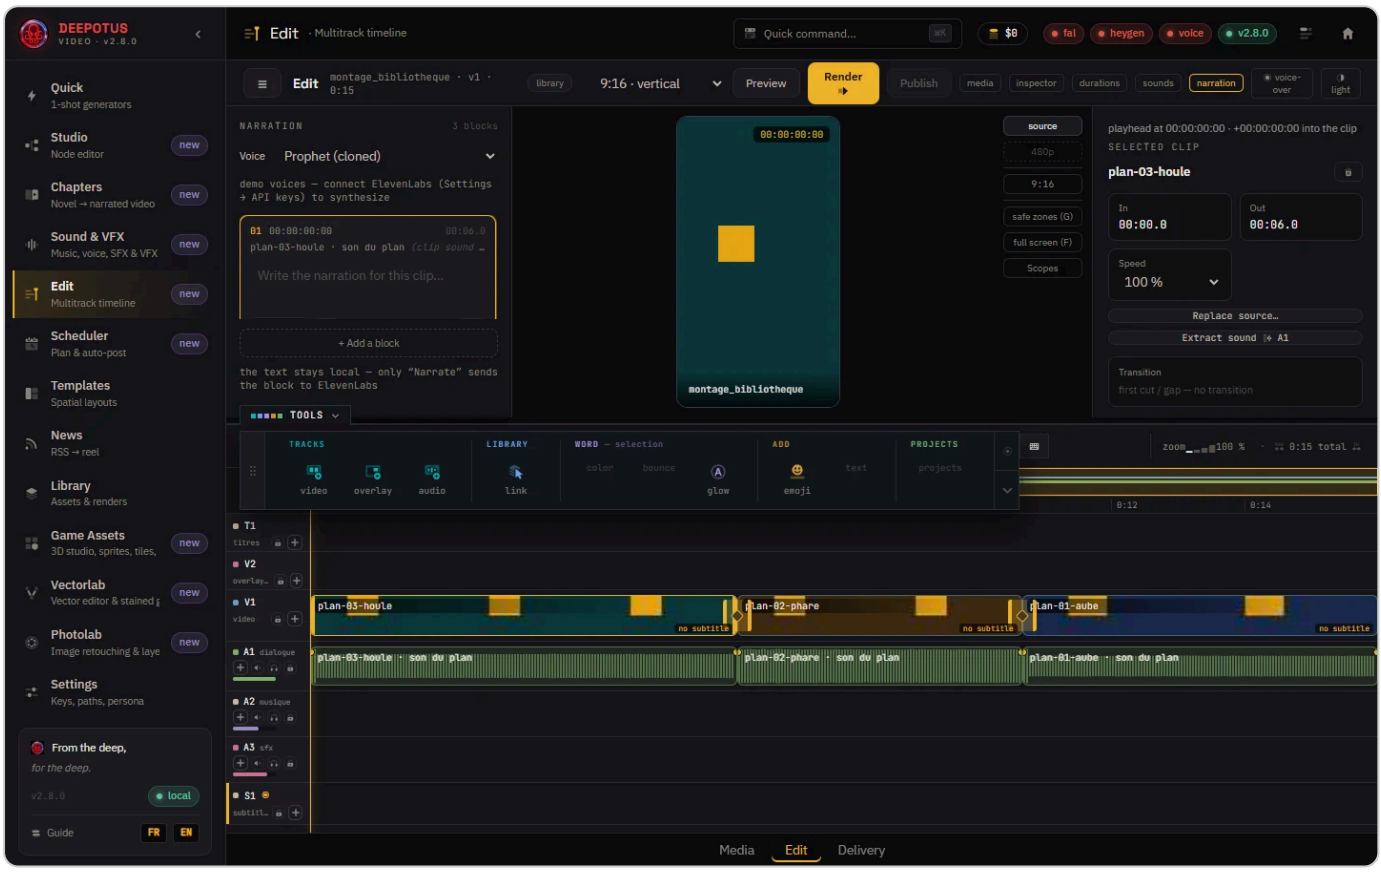

1. Shots on V1, their sound on A1 right below 2. Preview: a free 480p preview 3. Three views: Media, Edit, Delivery

The timeline at the bottom (video, sound, subtitles), the viewer in the centre, the inspector on the right; Preview, Render and Publish at the top.

At the bottom, the **timeline**: titles (T1), overlays (V2), main video (V1), then dialogue (A1), music (A2), sound effects (A3) and subtitles (S1). In the centre, the viewer; on the right, the inspector of the selected shot. At the top, **Preview** (free 480p preview), **Render** and **Publish**. The bottom bar switches between three views: **Media**, **Edit** and **Delivery**.

Place shots

1. Open the **media** drawer: all your finished video renders, with thumbnail, length and origin. Filter by origin (Studio, Chapters, Quick…), by title or by rating (three stars and up).

2. Click a media item to place it at the playhead. The **edit mode** decides how: overwrite, insert and push the rest, at the end, above, ripple overwrite, or fill the range.

3. From the Library, **Send to** › Edit also places a video on V1 with its sound on A1. A Studio render or a Chapters episode opens straight in the editor.

Cut and adjust

Cut

The blade (`Alt`+`C`) cuts the shot at the playhead. A gap is selected with one click and `Del` closes it.

Roll, slip, slide

`Alt` + drag the diamond moves the cut; `Alt` + drag the centre slips the content; `Shift` + drag moves the shot between its neighbours.

Range and markers

`I` and `U` set an in and out range; `Shift`+`X` cuts it on every track.

`Shift`+`M` sets a marker, `Ctrl`+`M` lists them.

Everything undoes

`Ctrl`+`Z` and `Ctrl`+`Y` over fifty steps: shots, mix, tracks, subtitles, range, markers.

Cut at scene changes

Right-click a clip: the app spots each scene change and cuts it for you.

Shots too long

In the ☰ menu › View, the timeline highlights shots that run too long and jump cuts from the same source.

Chain two shots

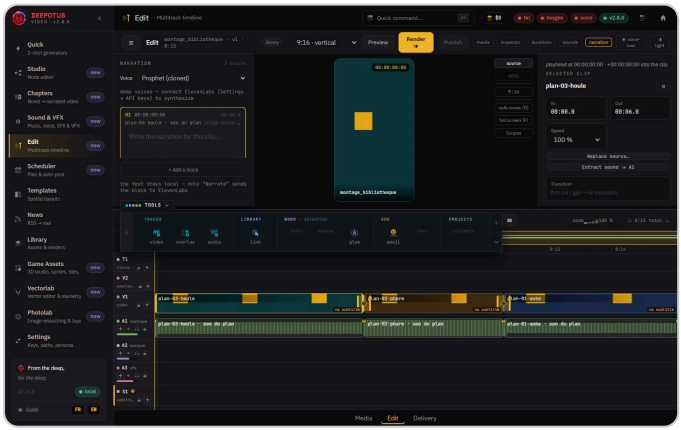

1. Click the diamond between two shots

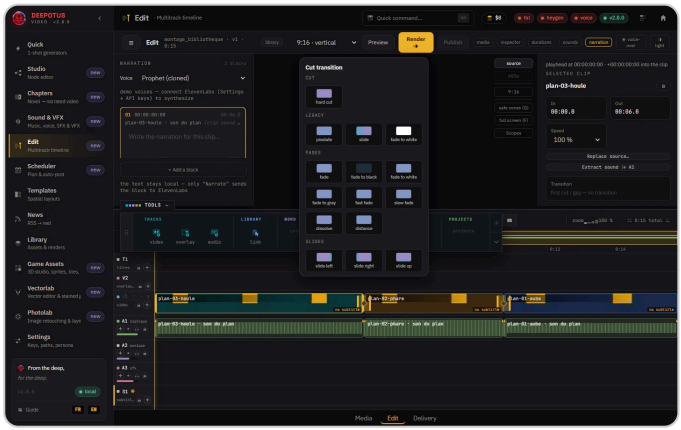

2. Pick the transition and its duration

A click on the diamond between two shots opens the transitions gallery and frame-by-frame trimming.

1. Click the **diamond** between two shots: the gallery opens, with 58 transitions sorted by family (fades, slides, wipes, zooms, pixels, shapes).
2. The same panel shows the last frame of the left shot and the first of the right one; shift the cut frame by frame.
3. For a quick fade on the selected shot: `Alt`+`T`. Simple fades play live; the others after a **Preview**.

Organise your projects

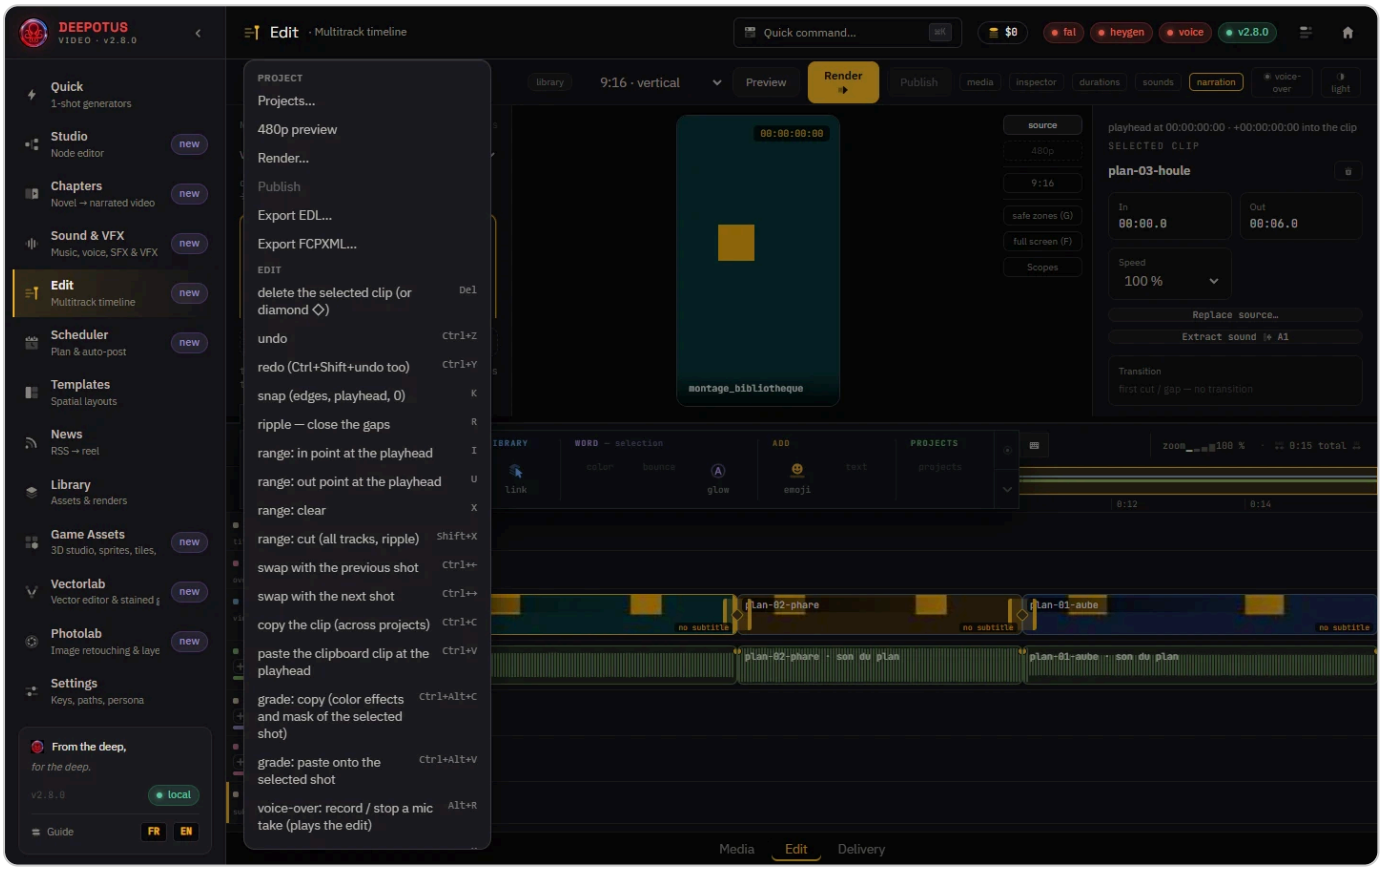

1. Project, Edit, Timeline, Markers, View, Help

The ☰ menu lists every action in six sections, each with its shortcut.

The ☰ menu lists every action in six sections, with its shortcut. **Projects…** opens the grid of your projects: open, duplicate, rename, delete, or compare two versions. Before opening another one, an unnamed edit is put away safely on its own.

New edit

An empty project, with its default tracks, and nothing else.

Copy between projects

`Ctrl` + `C` on a clip, `Ctrl` + `V` in another project.

Auto-clips

From a long video, AI suggests 1 to 8 excerpts of 15 to 60 s, then creates a project with the chosen one (paid; a free version without AI exists).

Your shortcuts

The `?` panel lists the shortcuts; a “Resolve” preset aligns them with DaVinci Resolve. Export them and import them back.

Shortcuts

Key	Effect
`Space`	Play / pause
`Alt` + `C`	Blade: cut at the playhead
`I` / `U`	Range in / out at the playhead
`Ctrl` + `Left` / `Right`	Swap the shot with its neighbour
`O`	Fold the tools bar
`B` / `T`	Sounds drawer / Narration panel
`G` / `F`	Safe zones / viewer full screen

MAKE VIDEOS

Grade, mix and add titles

You will give your shots a consistent colour, check them on the scopes, set their speed and framing, add titles and polish the sound. Free: everything is computed on your PC, except subtitle transcription and translation.

Grade a shot

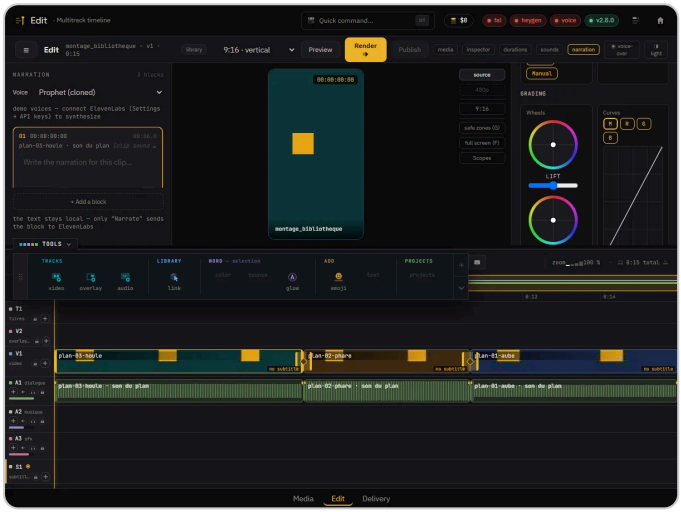

1. Shadows, midtones, highlights

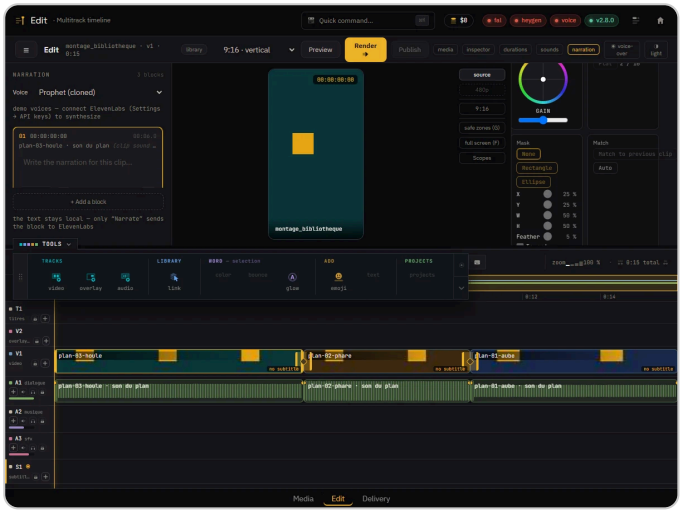

2. Match a shot to the previous one

Clip inspector, Grading panel: wheels, curves, mask, automatic match.

1. Select a shot: the inspector shows the **Grading** panel.
2. **Wheels**: shadows (lift), midtones (gamma) and highlights (gain). **Curves**: up to sixteen points, on the whole picture or each colour.
3. **Mask** limits the setting to a rectangle or an ellipse, feathered or inverted.
4. **Auto** removes a colour cast; **Match to previous clip** aligns the colours with the shot before.
5. **Copy grade** then **Paste grade** carry the setting over to another shot.

When playback stops, the viewer shows the shot's real graded picture; during playback, the picture stays raw.

Check on the scopes

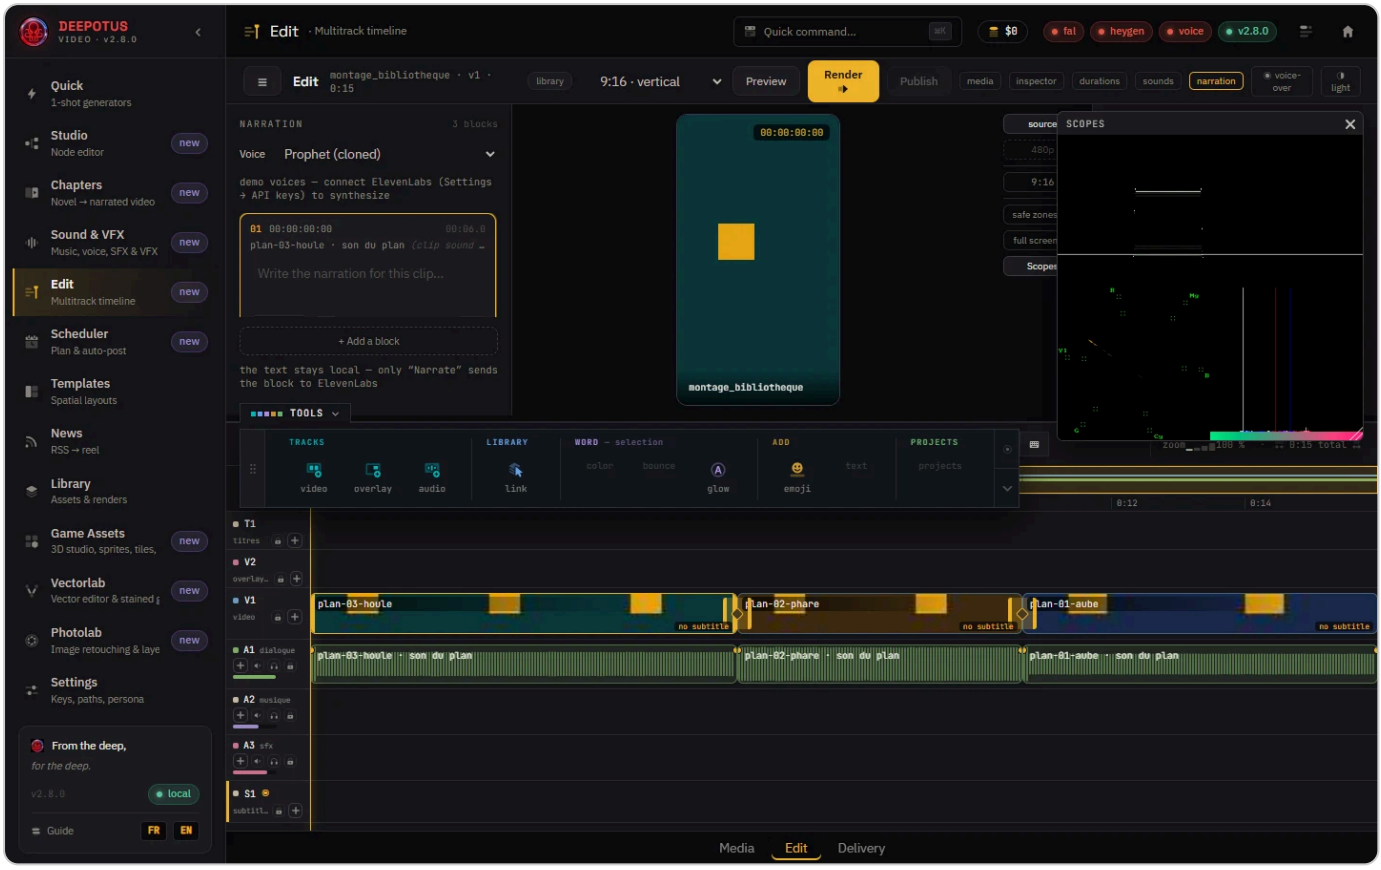

1. Scopes are computed when playback stops

Waveform, vectorscope and histogram of the graded picture, in a window you can move.

Scopes opens a movable window: waveform, vectorscope and histogram of the graded picture under the playhead, recomputed when playback stops. They account for speed, stabilisation and the render frame rate, and stay available in full screen.

Effects, speed and framing

The effects rack

Noise reduction, deflicker, debanding, sharpening, green-screen keying, monochrome, per-band hue and saturation, trails, lens distortion.

Speed and ramp

From 25 to 400 %. Interpolation smooths the motion; **Split at playhead** gives the rest of the shot another speed.

Stabilise and reframe

Stabilisation analyses a shaky shot; framing picks the part kept: centred, following the motion, or manual; dynamic zoom animates a reframe.

Adjustment clip

On track J1 (`Shift`+`J`), its effects apply to everything below, with fade in and fade out.

Overlays

Position, scale, rotation and opacity animated by keyframes, rounded corners, drop shadow.

Hand-drawn ramps

An effect ramp can follow a Bézier curve you draw; the shake's ramp raises or lowers its amplitude.

Titles and subtitles

1. `Shift`+`T` places a title card on track T1, burned in at render time.
2. Pick its template (full frame, lower third, caption, counter, chapter, quote, hashtag, call to action), its text, colour and font.
3. Subtitles have one track per language: only one is burned in, the others are exported. Translation can create a new language track.

Polish the sound

The audio rack

Six-band equaliser, noise reducer with adjustable floor, 50 / 60 Hz hum removal, stereo pan; listen to the rendered result.

Learn the noise

Set a range on a stretch of noise only: the app learns it and removes it at render time.

Voiceover from the mic

`voice-over` (`Alt`+`R`) records a take while the edit plays and places it on the dialogue track.

Narration

The `narration` panel (`T`) writes and synthesises the voice block by block; only “Narrate” sends a block to the provider.

Good to know

In the ☰ menu › View, the shot lightbox shows every V1 shot with its graded picture: one click goes to the shot.

Subtitle transcription and translation are paid; the rest of this chapter is free.

MAKE VIDEOS

Deliver: renders and formats

You will export your edit in the right format, at the right loudness, and send it to the Scheduler or to another program. Free: rendering happens on your PC.

Render the video

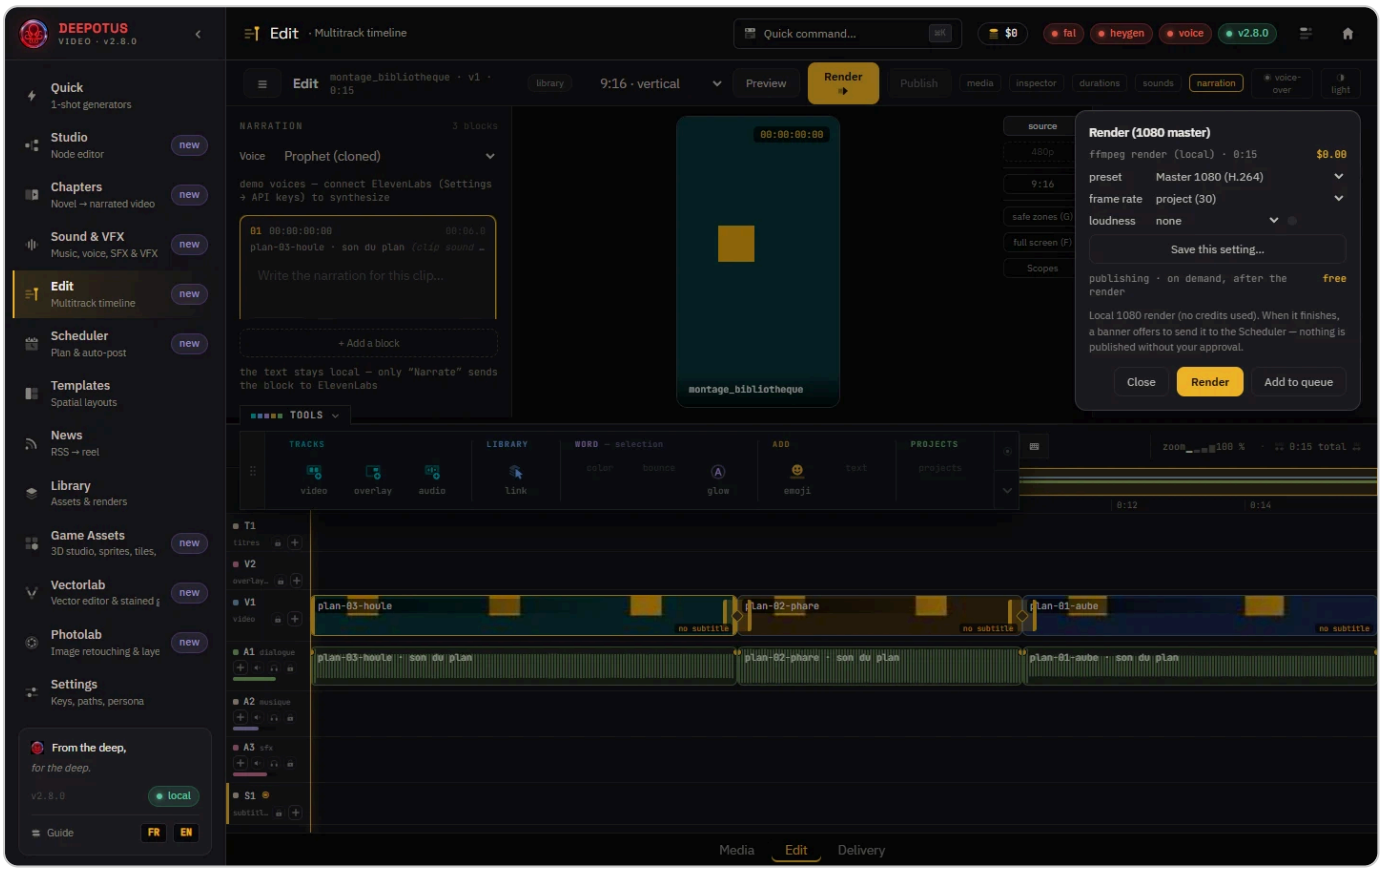

1. Pick the preset, then render or queue

Render… opens the render panel: output preset, frame rate, loudness, range, render queue.

1. Click **Render** in the top bar: the render panel opens.
2. Pick the **preset**: master 1080, web 4K, social 720, ProRes 422, H.265, WebM, audio only (AAC, MP3, WAV) or GIF 480. The project's format is kept.
3. Set the **frame rate** (24, 25, 30 or 60 fps) and the **loudness**: −14 LUFS for YouTube and TikTok, −16 for a podcast, −23 for broadcast. A badge compares the last measurement with the target.
4. **Render** starts the render; **Add to queue** puts it behind the others. To export only the I/O range, tick the box provided.

5 When it is done, a banner offers to send the video to the Scheduler, with your channels, a date and a prefilled caption. Nothing is published without you.

The Delivery view

The bottom bar › **Delivery** lists the project's renders and previews, their status (queued, running, done, failed), and leads to the Library. **Publish** sends the last final render back to the Scheduler.

Go further

Your own presets

Save this setting… keeps your output setting as a preset of your own, listed with the others.

To another program

☰ › Project › **Export EDL…** or **Export FCPXML…** to take the timeline into DaVinci Resolve or Final Cut.

True to the timeline

The render lasts exactly the timeline, frame for frame; a fade is centred on the cut with the source's real frames.

Good to know

Run a 480p **Preview** first: it is free and shows transitions, effects and mix.

PUBLISH

Plan and publish your week

You will fill your week with posts, review them, approve them in one go and let them go out on time on Telegram, X, YouTube, Instagram and TikTok. The Scheduler is free; only what you produce (image, clip, avatar) has a price, shown before the run.

Before you start

For a post to go out on its own, its network must be connected: see Connect your accounts. Without a connected account everything still works in assisted mode: the post is ready at the set time and you click `Send`.

The week at a glance

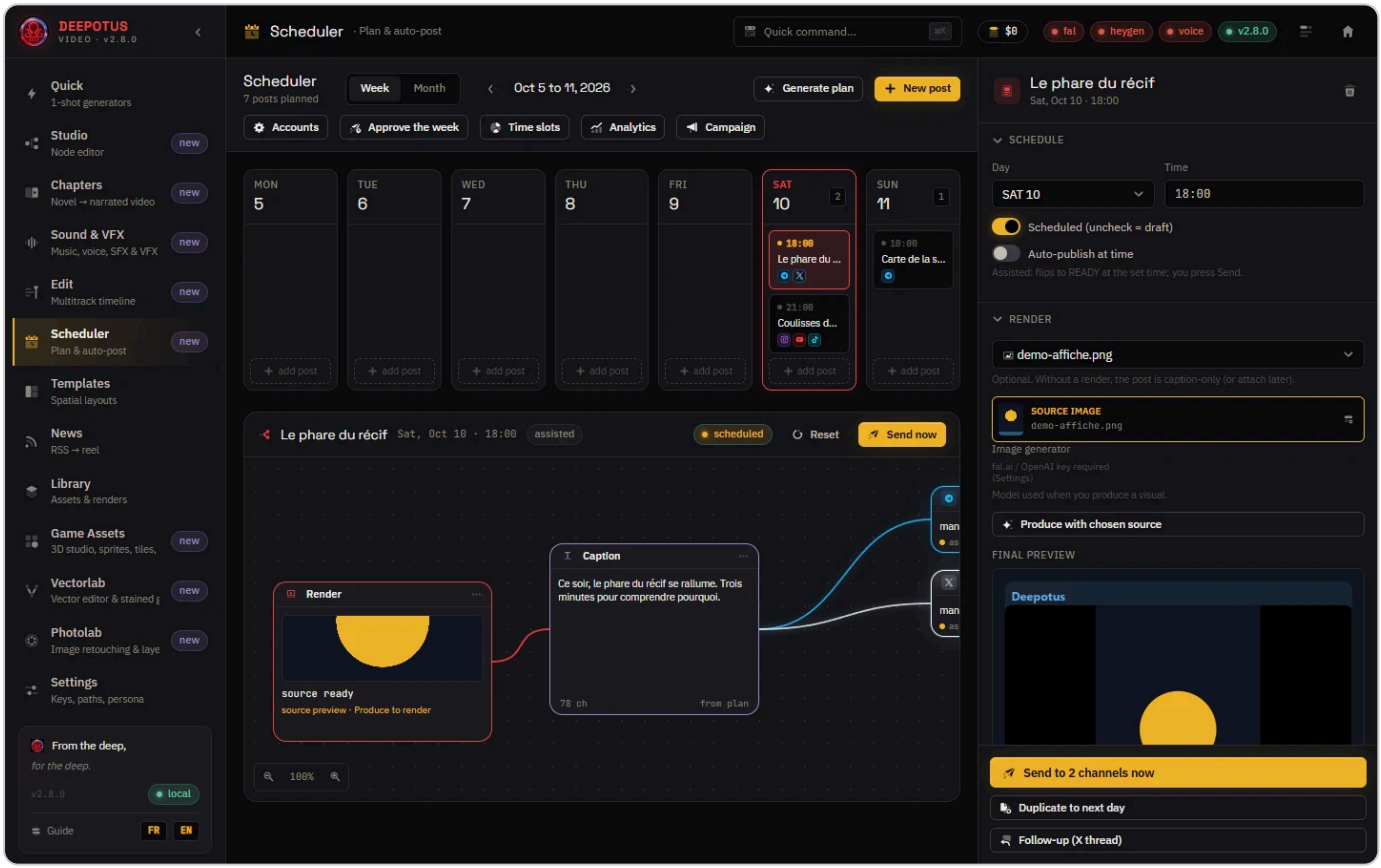

1. Fill the week with a plan 2. The graph: render, caption, channels 3. Nothing goes out without your click

The calendar at the top, the chosen post's graph in the middle, its inspector on the right.

Open `Scheduler` in the rail. At the top, the calendar in `Week` or `Month` view. Click a post: its **graph** (render, caption, one node per channel) shows in the middle and its **inspector** on the right. The screen fits the window from 1280 to 1600 px; the toolbar wraps when needed.

Fill the week

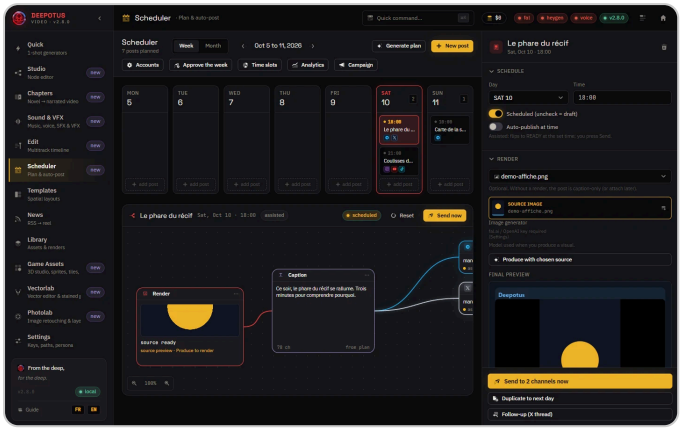

1. Generate plan

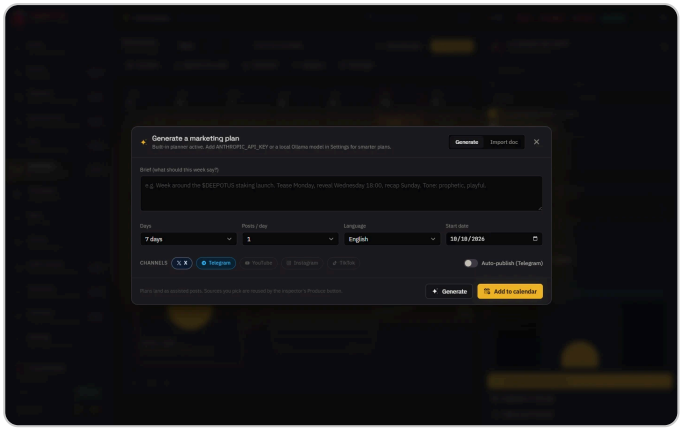

2. Describe the week in one sentence

A brief, a start date, the number of days and posts: you review the plan before it enters the calendar.

1. **Generate plan**: describe the week in one sentence, pick the number of days, posts per day, language, start date and channels, then **Generate**. Without an AI key the built-in planner does the job; with an Anthropic key or a local Ollama model the plan is finer.
2. Review the proposal, then **Add to calendar**. Times follow your per-channel time slots.
3. Already have a written schedule? The **Import doc** tab turns a .md, .txt, .docx or .pdf file into dated posts, keeping weeks and themes.
4. For a single post: **New post**, or **add post** at the bottom of a day. Episodes from Chapters land here as drafts.

Prepare a post

In the inspector, set the day and time, tick **Scheduled** (unticked, the post stays a draft), pick the channels and write the caption, the same on every channel, with its character count. **Produce with chosen source** makes the visual or clip the plan calls for; the render runs in the background and its thumbnail appears on the graph's Render node.

Coulisses du montage

Sat, Oct 10 · 21:00

FINAL PREVIEW

Reels · zones indicatives · 59 car.

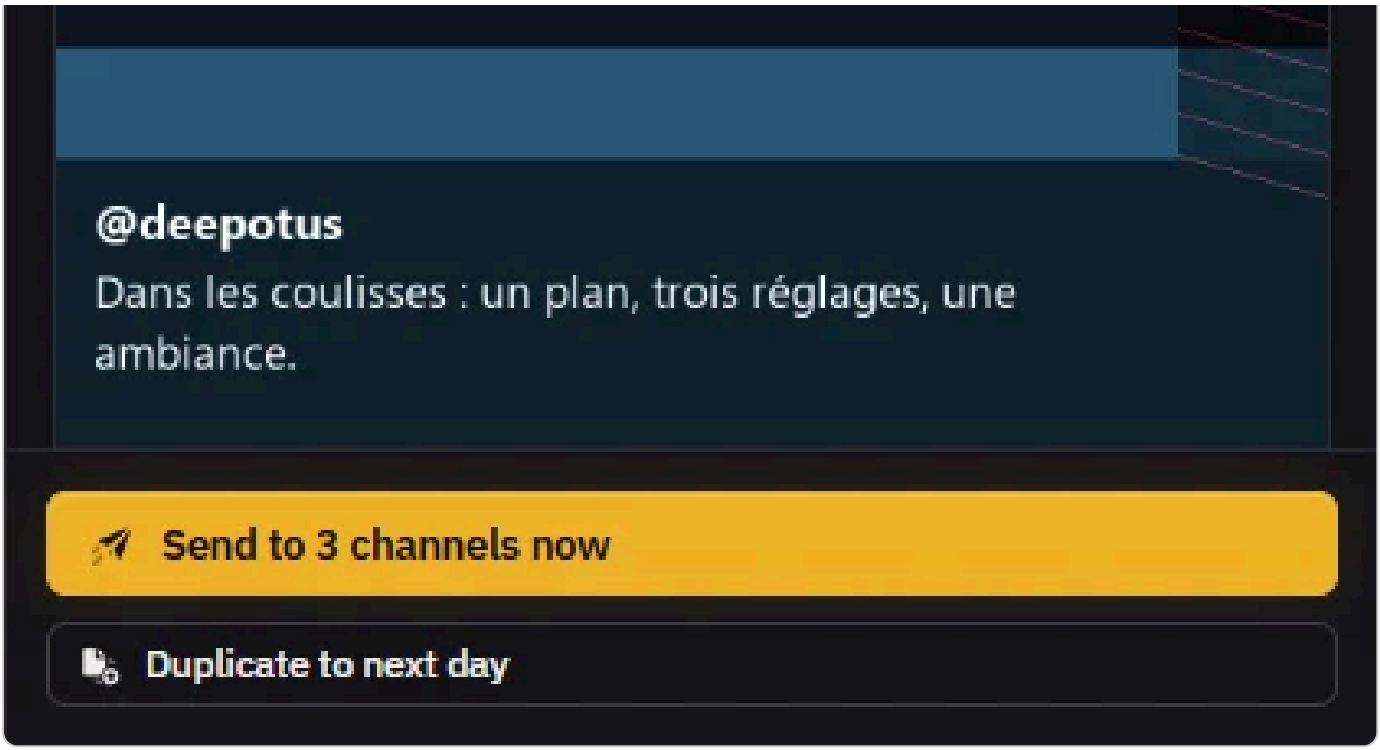

1. Put nothing important under the dark areas

For a vertical post, the preview darkens the areas the network's interface covers.

The **final preview** shows the post as it will look on the network. For a Reel, a Short or a TikTok, it darkens the top and bottom bands and the right rail that the network's interface covers: put nothing important there.

Approve and publish

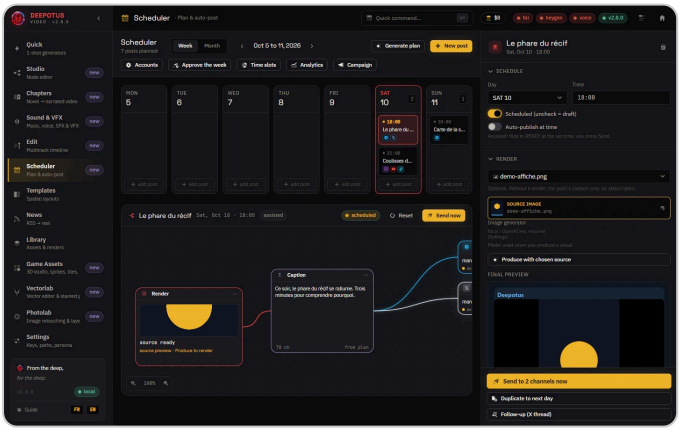

1. Approve the week

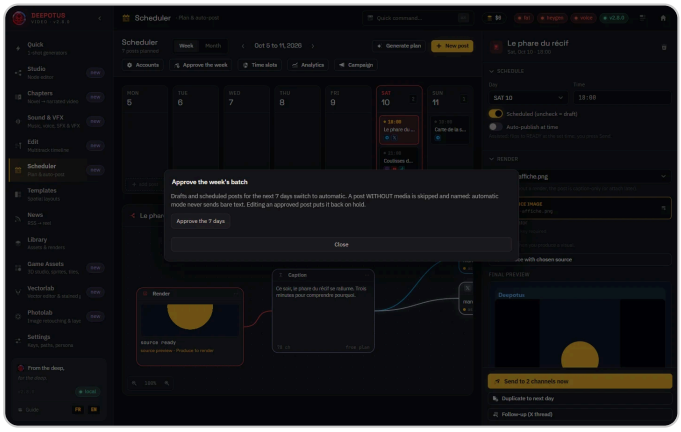

2. One click approves the whole batch

Batch approval: the next 7 days switch to automatic, a post without media is skipped.

1. Each post has a mode. **Assisted** (default): at the set time it flips to ready and you click `Send`. **Auto** (`Auto-publish at time`): it goes out on its own on its connected channels.

2. **Approve the week** then **Approve the 7 days** switches the drafts and scheduled posts of the next 7 days to automatic in one go. A post without media is skipped and named; really editing an approved post puts it back on hold.
3. To publish right now: **Send now** on the graph, or the yellow inspector button that counts the channels.

Before each send, the network's quota is checked (X 500 posts a month, Instagram 100 per rolling 24 h, YouTube 100 a day, TikTok 15 a day); a channel already published is never published again.

Going further

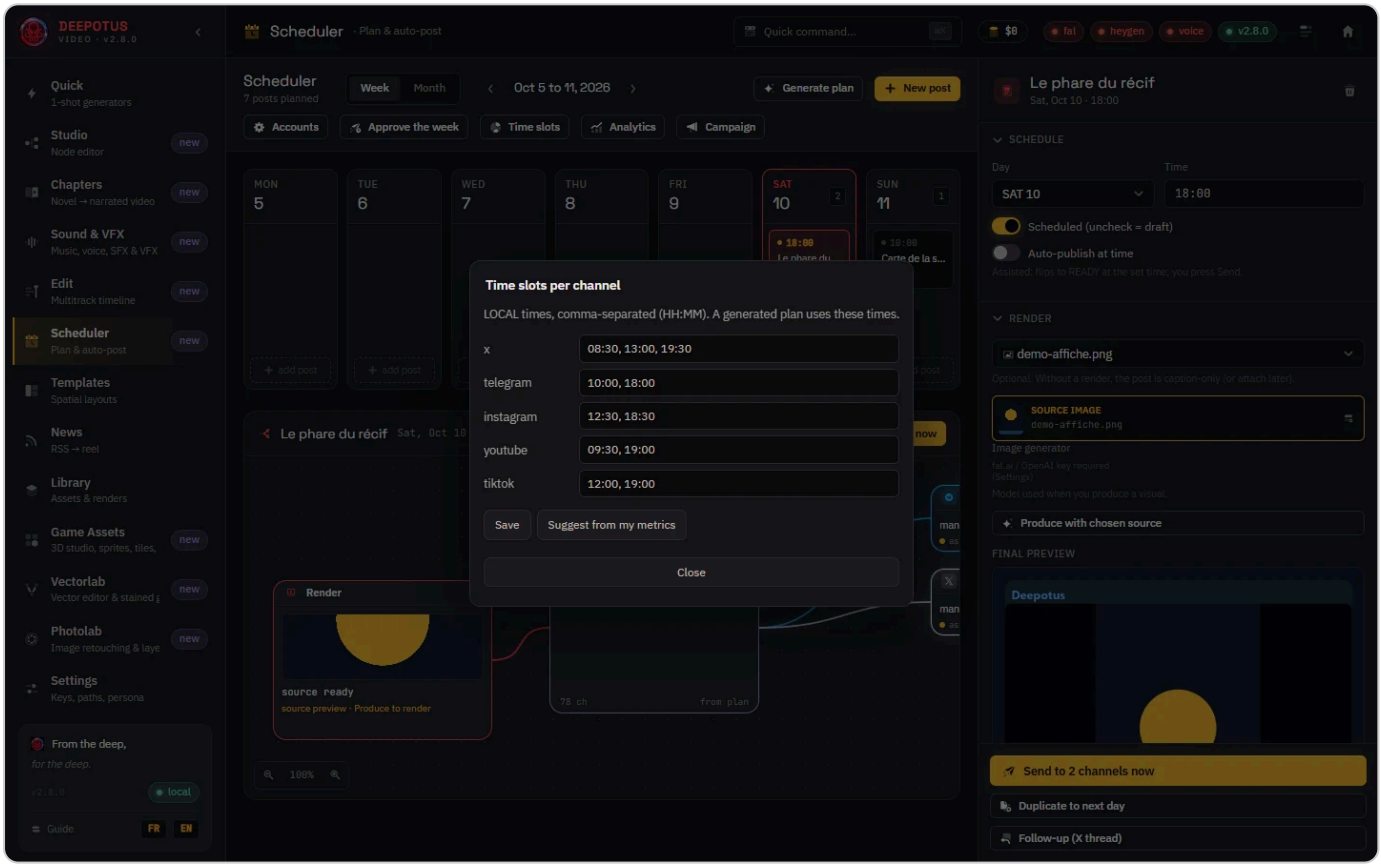

1. Times separated by commas 2. Or let your metrics suggest

Local times per channel, applied by generated plans.

Time slots

Time slots: local times per channel, separated by commas, then **Save**. **Suggest from my metrics** picks the hour with the best engagement, from 5 measured posts.

Analytics

Views, likes and engagement over the last 28 days by channel, format and week, with the top 5. Metrics are re-read every day; **Refresh now** re-reads 10 posts.

Campaign brief

Campaign › **Campaign brief**: objective, key messages, fixed sections and banned terms. The next plan follows it; banned terms are removed from captions and hashtags.

Recurring series

A rule (days, time, channels, caption template with {date}, {weekday}, {week}): **Create series**, then **Create 4 weeks**. These are drafts; creating twice duplicates nothing.

Suggested recycling

Your best-measured posts, published more than three weeks ago, come back with a varied caption. **Suggest for tomorrow** creates a draft; no paid AI is called.

X thread

Follow-up (X thread) writes the follow-up to an X post: it goes out two minutes later and replies to the previous message, never before it.

Duplicate, sign

Duplicate to next day copies the post. Under the caption, a row of emojis and the **Telegram Premium pack** insert your brand tags, set in Settings.

With the phone

Approved posts form a batch the paired phone can take along. A post handed to the phone says so in the inspector, with "Take back on the PC".

Good to know

> After an action (approve, create a series, suggest a recycle, add a follow-up), the list reloads by itself: no need to leave the screen.

The calendar, built-in plans, approval and sending are free. Producing an image, a clip or an avatar is paid, under your spending caps (see Control your spending).

PUBLISH

Connect your accounts

You will plug in Telegram, X, YouTube, Instagram and TikTok so the Scheduler can publish on its own. Free: all it takes is the access each network provides, stored on your PC.

Where it is set

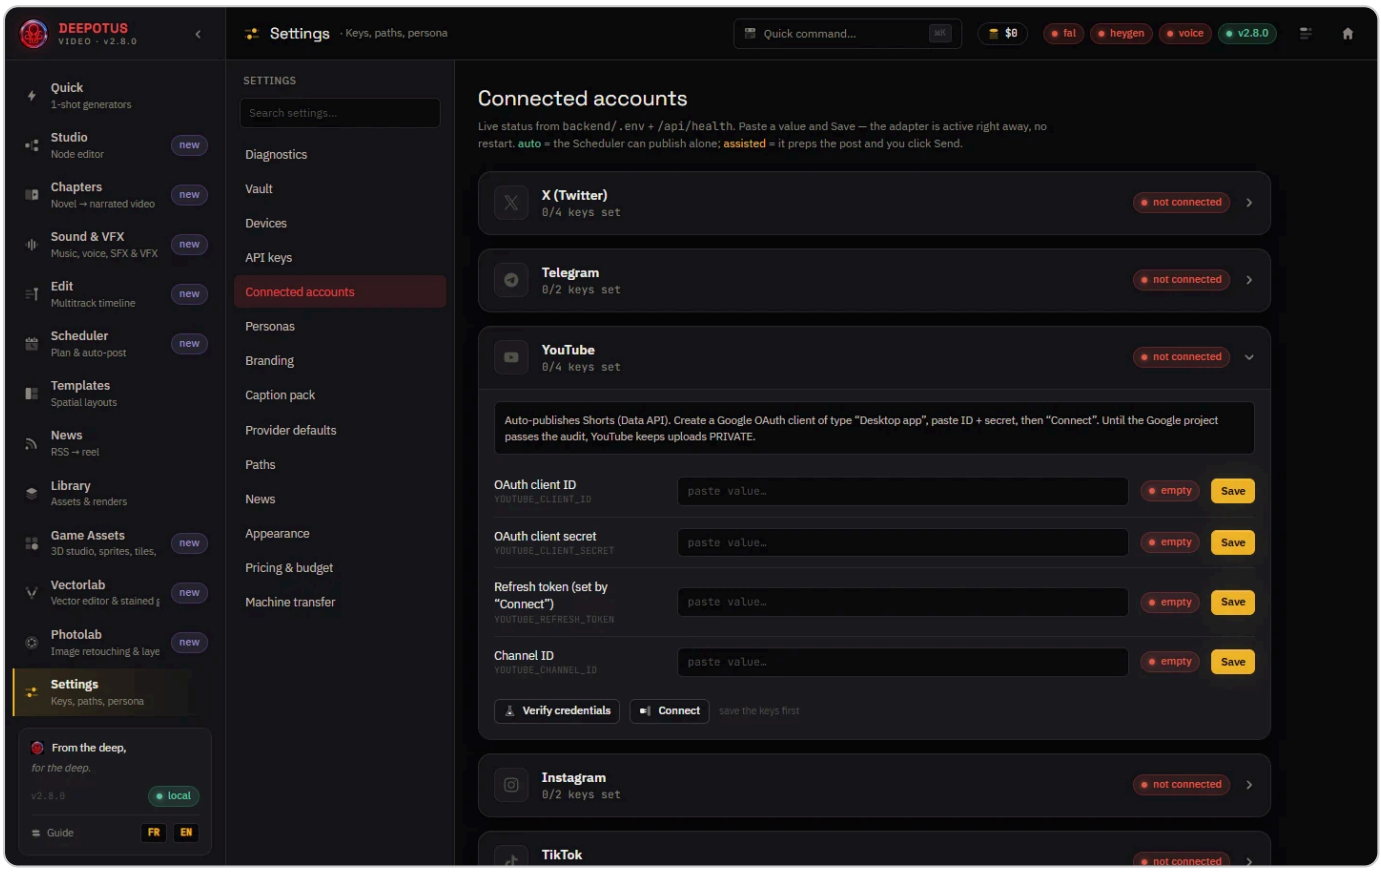

1. Paste the ID and secret, then Save 2. Connect opens the consent page

Settings, Connected accounts: one block per network, its keys, its live status.

Open `Settings` › `Connected accounts`. One block per network shows how many keys are set and its live status: **auto** (the Scheduler publishes alone), **assisted** (it preps the post and you click Send) or **not connected**.

1. Expand the network's block and read its note: it says where to get the access.
2. Paste each value and click `Save` next to it. The adapter is active right away, no restart.
3. Check with `Verify credentials` (or `Send test message` for Telegram).

Network by network

Telegram

The simplest, in five minutes. Create a bot with @BotFather, make it an administrator of the channel, paste the bot token and the channel chat ID.

X

The four keys from developer.x.com (API key and secret, access token and secret). The free tier allows posting, up to 500 posts a month.

YouTube

Create a Google OAuth client of type "Desktop app", paste its ID and secret, then **Connect**. Until the Google project passes the audit, YouTube keeps Shorts private, and the Scheduler says so.

Instagram

A business account and a Page token (EAA…) or Instagram token (IG…). Reels are uploaded from your PC, no hosting, up to 100 posts per rolling 24 h.

TikTok

A TikTok for Developers app (Login Kit, "Desktop" platform): paste the client key and secret, then **Connect**. Without an audit, uploads are private and the account must be too; set TIKTOK_AUDITED to true after the audit.

Connect YouTube and TikTok without copying a code

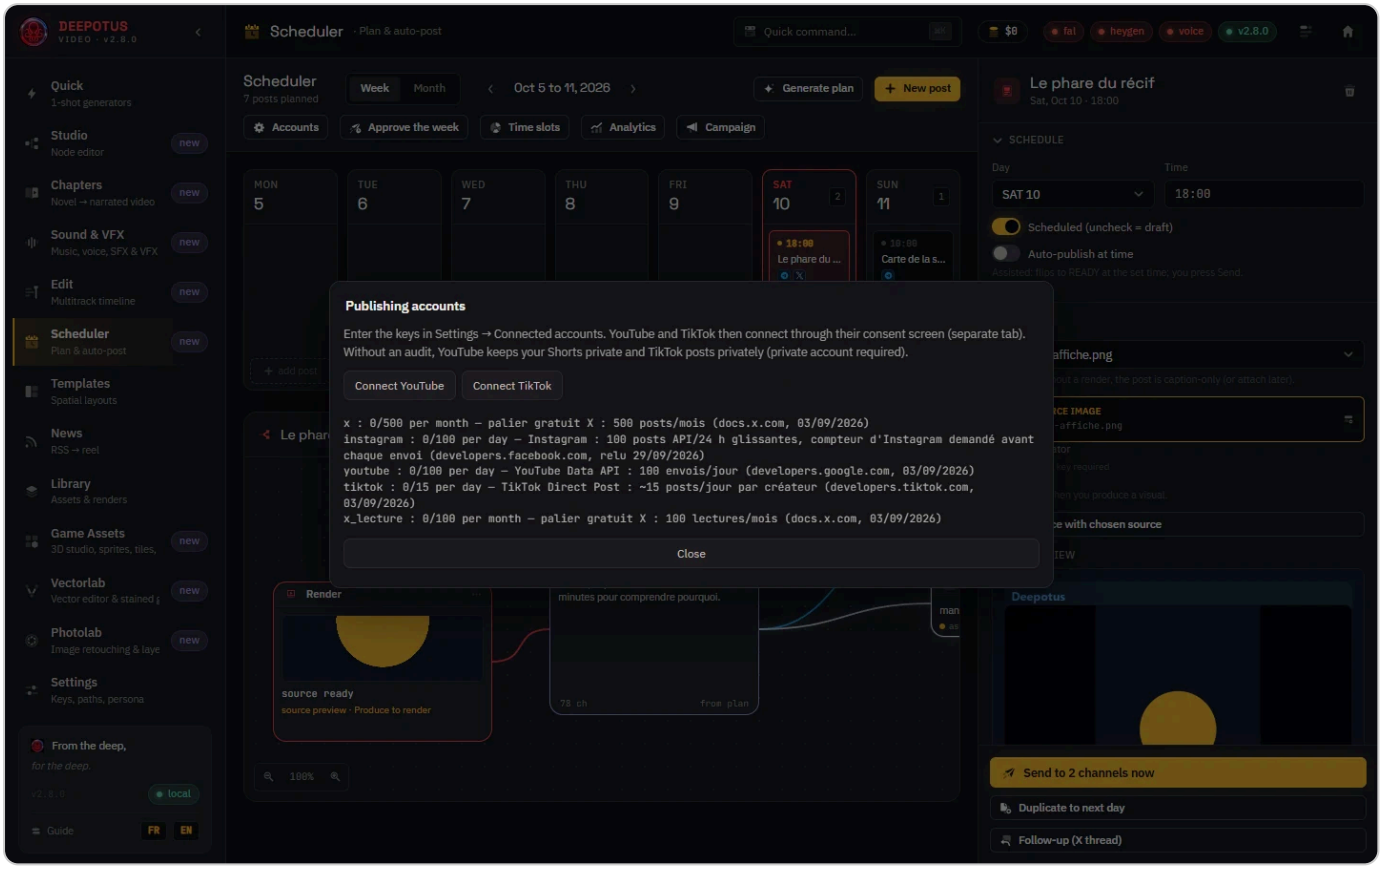

1. Connect through the consent page 2. Each network's quota, checked before sending

In the Scheduler, Accounts shows the month's quotas and connects YouTube and TikTok.

Connect opens the network's consent page in a new tab: accept, come back, the status turns to connected and the refresh token is saved on its own. The same gesture exists in the Scheduler: the **Accounts** button, then **Connect YouTube** or **Connect TikTok**. That panel also shows each network's quota for the month.

Good to know

> A network that is not connected blocks nothing: its posts switch to assisted mode and you send them in one click, or from the phone.

PUBLISH

Follow the news and turn it into posts

You will follow your news feeds, let the free ranking sort today's articles, and go from article to scheduled post with the daily pipeline. Sorting, filtering and preparing are free; AI and illustrated reels are paid, on request only.

The News screen

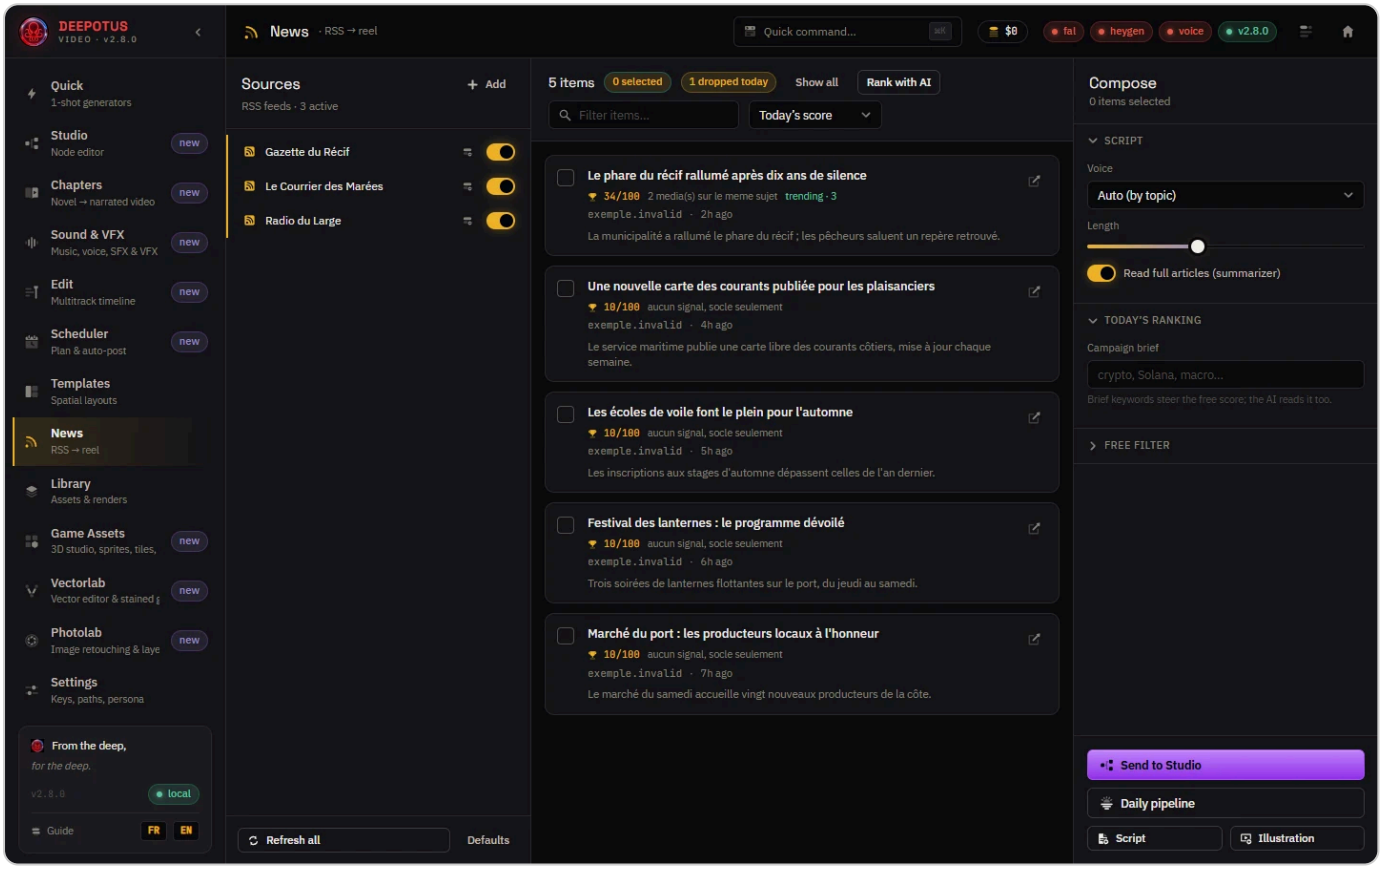

1. Picked up by three outlets: trending 2. Filtered-out items, reasons on hover 3. The score is free; AI on request

Feeds on the left, ranked articles in the middle, Compose on the right.

Open `News` in the rail. On the left, your **sources**; in the middle, the ranked **articles**; on the right, the **Compose** panel. What you ticked, the query, the sort, the voice, the length and the script are back when you return to the screen.

Choose your sources

1. `Add`: paste the address of an RSS or Atom feed. `Defaults` restores the bundled pack (crypto, geopolitics, economy, politics).
2. A source's switch pauses it without deleting it.

3 **Refresh all** re-reads every feed. The screen also does it on its own, once a day. Pickups of the same wire story are merged into a single article.

Today's ranking

By default, articles are sorted by **Today's score**, a free score: your brief's words, pickup by other outlets, a readable article. The reason and the score show on each card; **Most recent** sorts by date.

Trending

A topic picked up by three outlets the same day is marked trending, in green on the card.

Already covered

A topic close to a reel launched in the last 30 days is marked in amber, title on hover, and moves down the ranking. An overused source moves down too.

Dropped

The dropped chip counts the articles the filter removed; hover it for the reasons. **Show all** shows the full ranking, **Top only** goes back to the first five.

Rank with AI

Rank with AI re-ranks the batch with a language model: paid, under your caps, and only when you click.

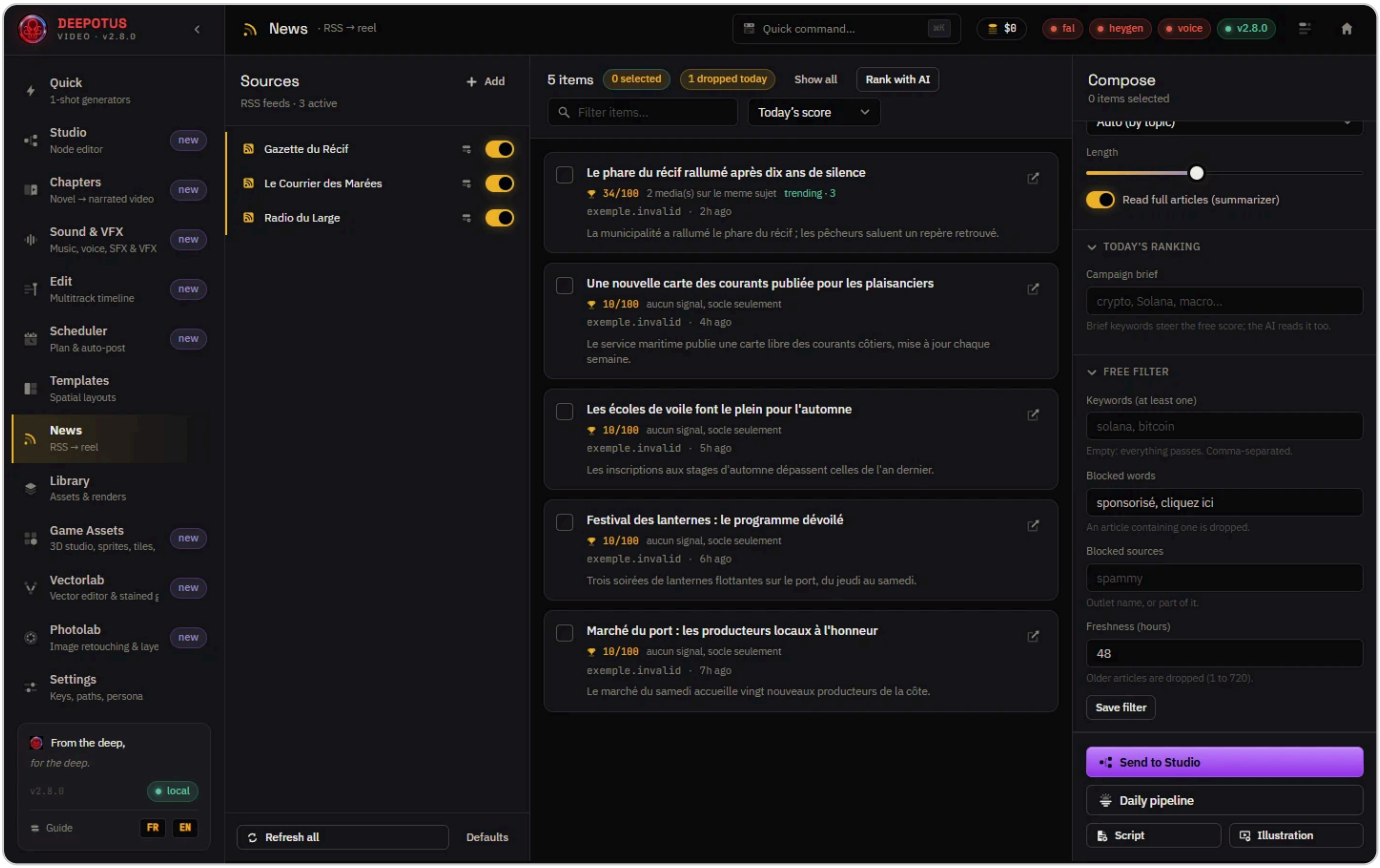

1. A blocked word drops the article 2. Saving re-ranks at once

Keywords, blocklists and freshness decide what enters the ranking.

In Compose, **Campaign brief** says what matters today and steers the score. **Free filter** sets the keywords (empty: everything passes), blocked words and sources, and freshness in hours (1 to 720); **Save filter** re-ranks at once.

From article to post: the daily pipeline

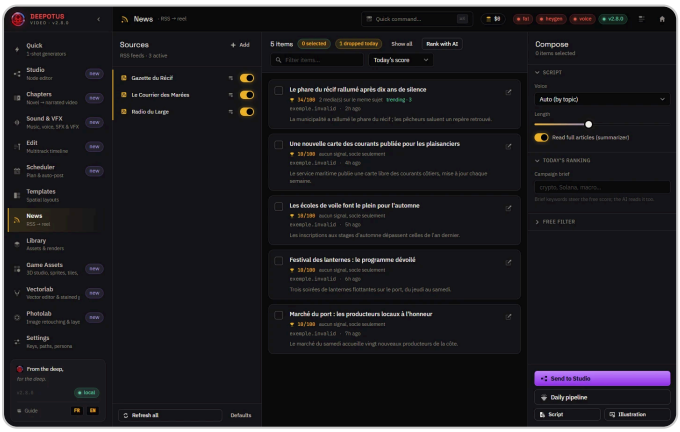

1. Daily pipeline

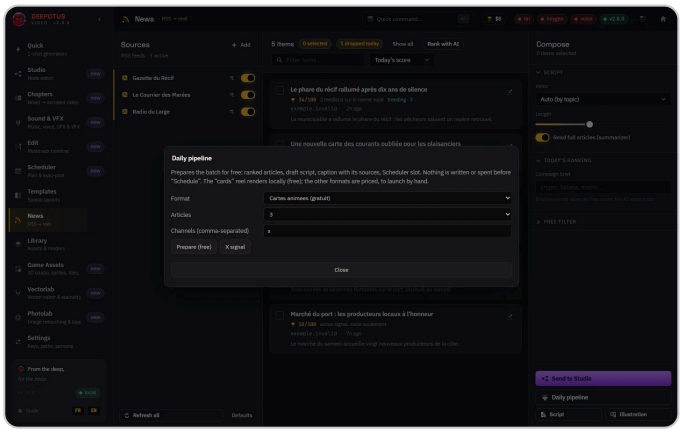

2. Preparing costs nothing

From article to scheduled post: prepare for free, polish if needed, schedule.

1. Click **Daily pipeline**. Pick the reel **format**, the number of articles and the channels. Each format's cost shows before any run: animated cards are free, rendered on your PC; AI illustration, video shots, avatar and voice-over are paid.
2. **Prepare (free)** assembles the batch without any paid call: ranked articles, draft script, caption ending with its sources (outlet, date, domain), Scheduler slot.
3. If needed, **Polish with AI** rewrites the script (paid). If the spending cap stops the polishing, the draft is kept.
4. Review the script, caption and slot, then **Schedule batch**: the post lands in the Scheduler, with its cards reel rendered for free. Nothing is written or spent before that click.

Going further

X signal

X signal reads recent X posts on your brief: on click only, three times a day at most, counted in your X read quota.

Voice by topic

With **Auto (by topic)**, the persona's tone is picked from the topic's words, for free, and the reason is shown.

To the Studio

Tick articles, then **Send to Studio** to build them yourself in a graph (see Compose with the Studio).

Good to know

Everything paid on this screen (rank with AI, polish, X signal, illustrated reels) goes through your spending caps; without a key, the button says so instead of firing.

IMAGES AND DRAWING

Organise and find your media

You will be able to find any image, sound or render, file it in a project, send it to the screen that needs it and clean up without losing anything. Everything is free and stays on your PC, except captions by a vision model (paid, quote shown).

The Library screen

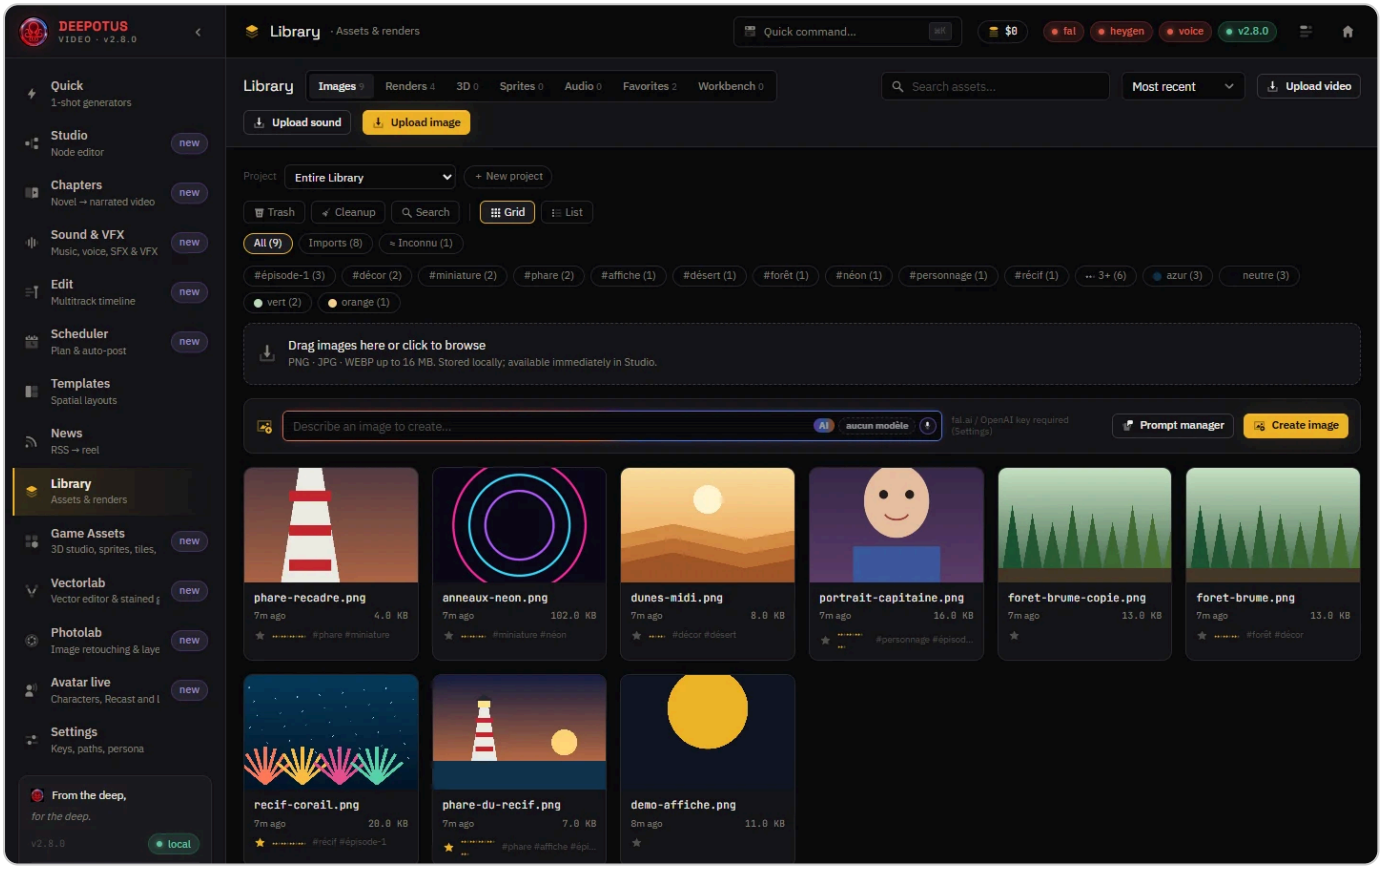

1. One tab per kind of media 2. Look at a single project 3. One click on a tag filters the grid 4. Grid or sortable list

Tabs by kind, the project you are looking at, tag and hue chips, then the grid.

Open **Library** in the rail. At the top, one tab per kind: **Images**, **Renders**, 3D, Sprites, Audio, **Favorites** and **Workbench**, where a model's versions are listed in lineage order. Everything the app produces arrives here: generators, Photolab, Vectorlab, uploads, phone.

1. **Source** chips say where each asset comes from (Generator, Materials, Uploads, mobile…). One click filters.
2. **Tag** chips are counted; “3+” keeps images rated three and up. Each image also gets its dominant colour: the **hue** chips (twelve hues and neutral) filter by colour.

3. Sort by **Most recent** or by size, then **Grid** or list: the list has ten sortable columns (name, source, size, date, dimensions, tags, rating, favourite, licence, hue) and a row opens the sheet.

4. Each thumbnail shows the file's size and age; the star sets a favourite without opening the sheet, and the dots recall the rating.

An image's sheet

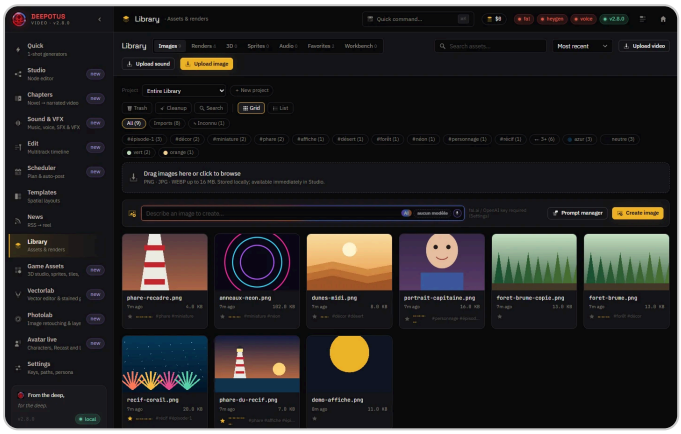

1. Open an image

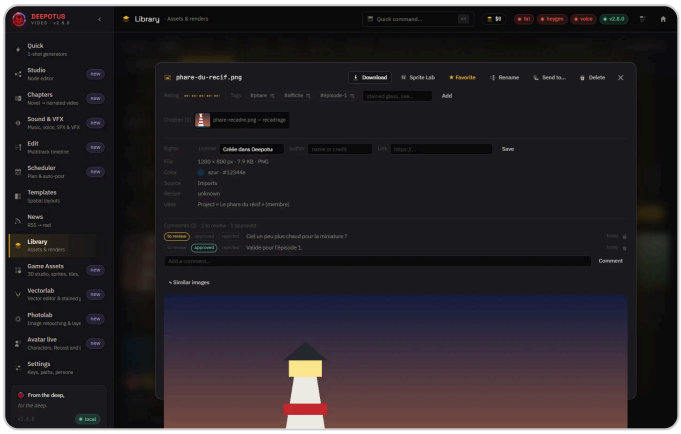

2. Its children: crop, upscale… 3. Licence, author and link 4. Review comments

One click on a thumbnail opens its sheet: rating, tags, lineage, rights, file, uses and comments.

One click on a thumbnail opens its sheet. From top to bottom you find:

Rating and tags

A rating from 0 to 5, separate from the favourite, and free tags: type, then **Add**. Search covers tags.

Parent and children

Parent shows the original image and the relation (crop, upscale, retouch…); children appear as clickable thumbnails.

Rights

Licence (suggested from the source), author and link, then **Save**. An unknown licence is flagged in amber: check it before publishing.

File, recipe, uses

Dimensions, weight, format, dominant colour, source; the generator's recipe (prompt, style, model, seed); and every use: renders, posts, bible, shots, projects.

Review comments

Type, then **Comment** (or Enter). Each comment has a clickable status: to review, approved, rejected; for a sound, the mm:ss moment. They follow the asset when renamed or restored.

Similar images

≈ Similar images shows visually close images, thanks to the CLIP model installed on your PC.

Replay recipe regenerates a generated image with its original recipe, seed included. It is paid: the price is shown and you confirm before the run; the new image arrives at the top of the list.

Send to another screen

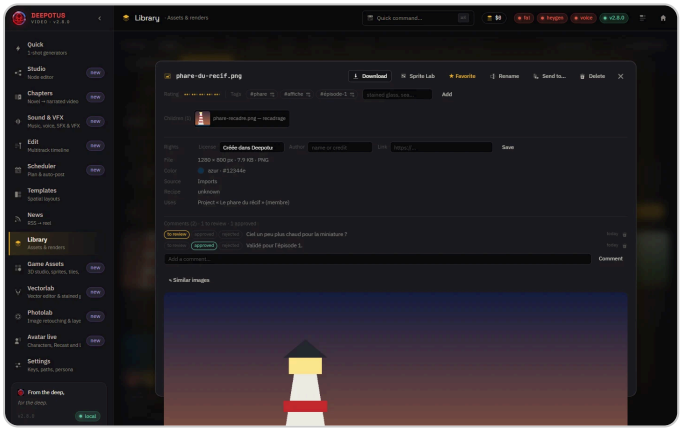

1. Send to…

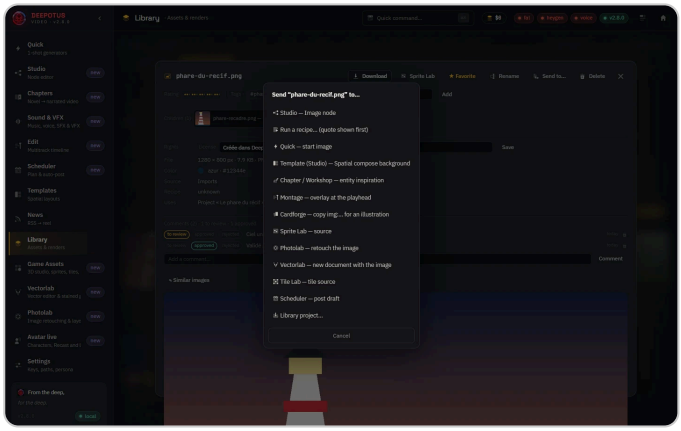

2. Retouch it in Photolab 3. Or file it in a project

From the sheet, Send to… lists every screen that can take the image, and projects.

1. In the sheet, click **Send to…**.
2. Pick the screen: Studio, Quick (start image), Template, Chapter or Workshop, Editor, Cardforge, Sprite Lab, Scheduler… Three targets are specific to images: **Photolab** to retouch it, **Vectorlab** which creates a new document at its size, and **Tile Lab** which makes it a tile's source.
3. **Library project…** files the asset in a project: the ones already holding it are ticked (one click removes it), and **+ New project…** creates one.

File by project

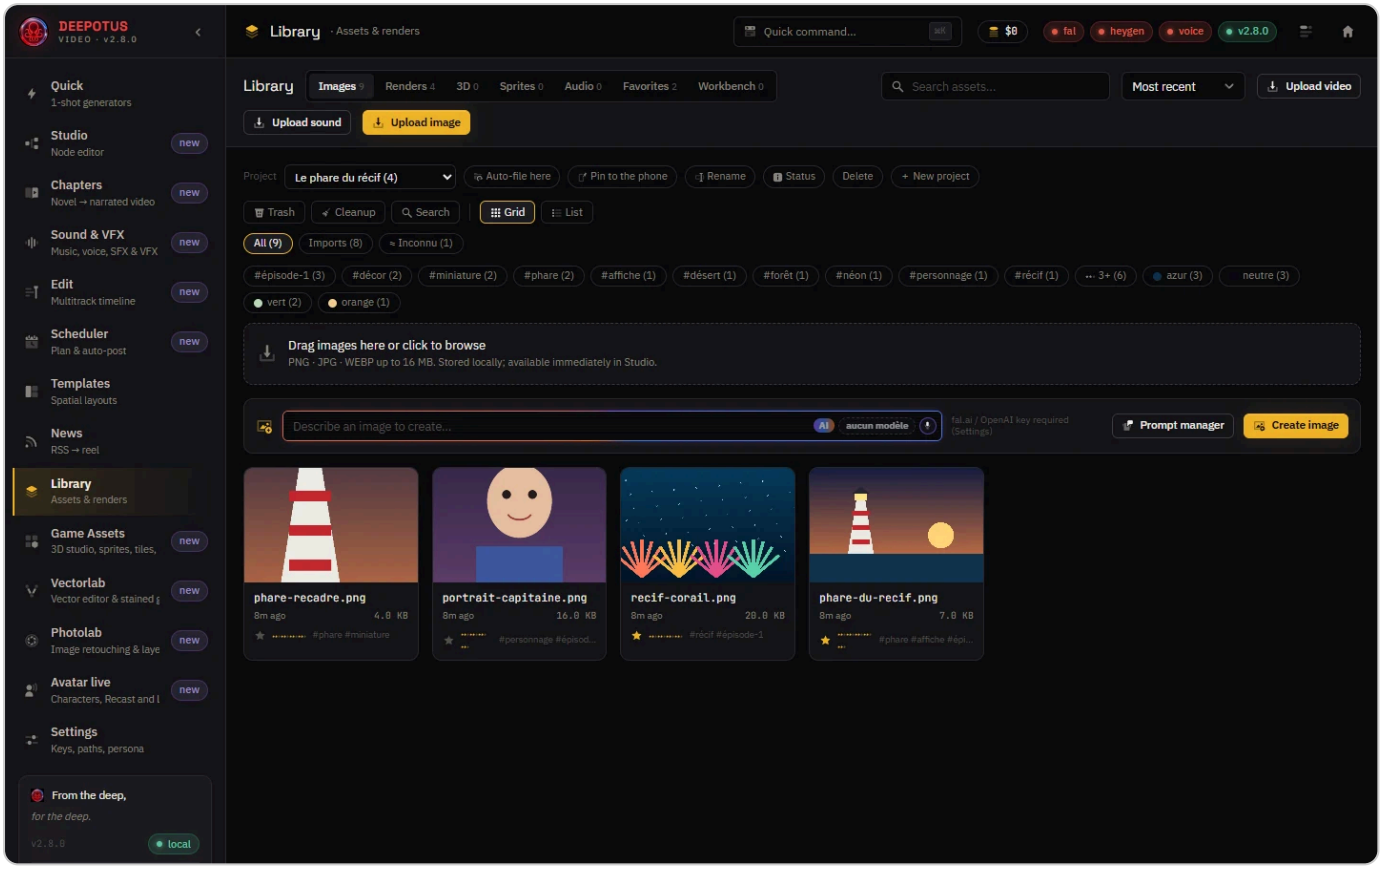

1. Le phare du récif: 4 assets 2. Everything you make is filed here 3. The mobile companion's project

Choosing a project filters every tab; its bar auto-files into it, pins it to the phone or shows its status.

The **Project** list at the top of the screen looks at a single project: every tab then only shows its assets. `+ New project` creates one; a project can be renamed or deleted without touching its assets.

Active project

`Auto-file here`: everything you produce next (images, renders, 3D, sprites) is filed in it. The reminder stays on screen while it is active, and one click stops it.

On the phone

`Pin to the phone`: a single project is pinned, and it is the one the mobile companion shows.

Project status

The `Status` button sorts its assets into edited, published, printed or unused.

Search

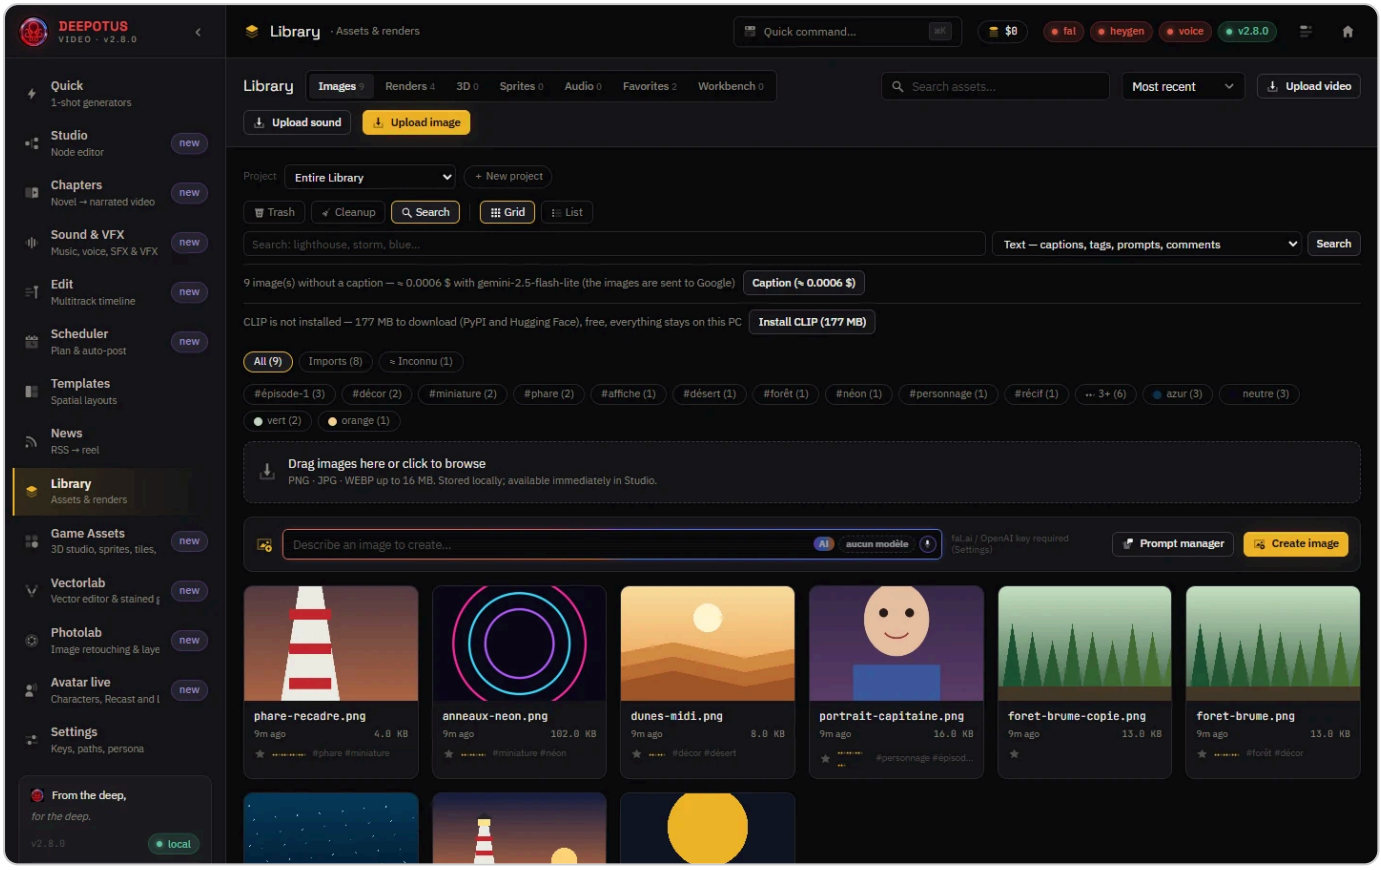

1. Captions, tags, prompts, comments 2. Paid: the quote is shown first 3. CLIP: free, on your PC

Free text search, captions by a vision model (paid, quote shown) or local CLIP, by meaning.

The `Search` button opens the search bar, with three engines:

1. **Text**: captions, tags, prompts, recipes, comments and names, ignoring accents and case. Free, instant.
2. **Captions**: a vision model describes images that have none, so you find them better. Paid: the quote is shown (images are sent to Google), the cap is checked, then progress counts the images done.
3. **CLIP**: searches by meaning (“a lighthouse at sunset”). `Install CLIP` downloads about 177 MB once, then indexing and search run on your PC, for free, sending nothing.

The `Audio` tab mounts the Sounds drawer (tags, filters, sorting, actions) above the upload: see Sound & VFX.

Clean up without losing anything

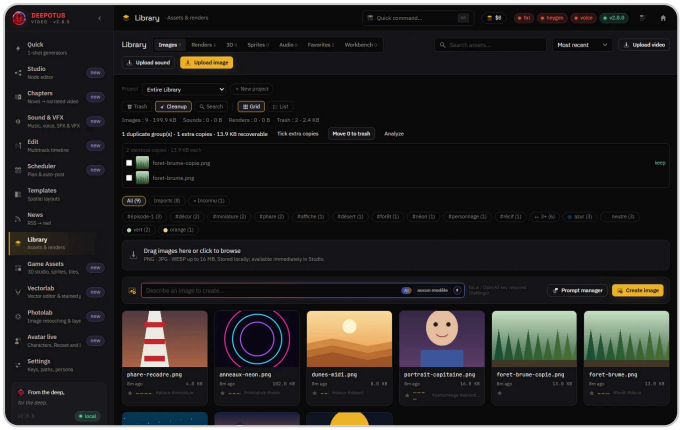

1. A group of exact duplicates

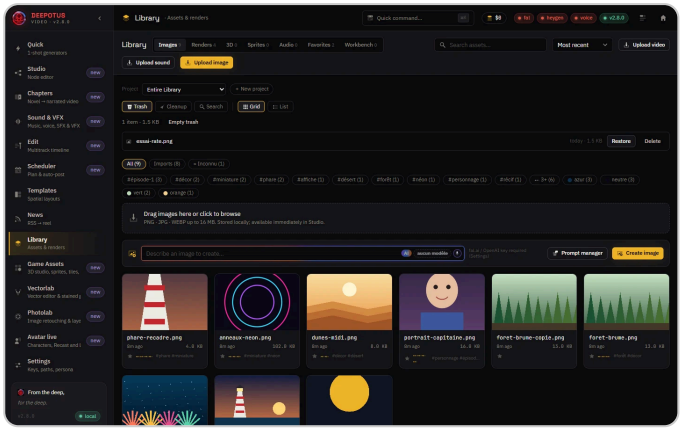

2. Nothing is lost until you empty it

Cleanup shows the weight and exact duplicates; what you discard goes to the trash, where you can restore it.

1. Deleting an image, a sound or a render moves it to the **Trash**, never further.
2. **Cleanup** shows the weight per kind and groups of exact duplicates. **Tick extra copies** keeps the used copy (or the oldest); protected copies (favourite, use, project) stay greyed out. Confirm to move them to the trash.
3. The **Trash** counts items and their weight, and flags those older than 30 days. **Restore** puts the asset back with its projects; if its name was taken, it gets a neighbouring name and says so.
4. Emptying deletes permanently: old items, everything, or a single item, always after confirmation.

Good to know

Favourites are saved by the app, no longer by the browser: they follow you from screen to screen and through a transfer. Old browser favourites were imported once.

Only **Replay recipe** and vision-model captions cost money: price shown first, spending caps checked. Sorting, text search, CLIP, projects and cleanup are free.

IMAGES AND DRAWING

Retouch a photo (Photolab)

You will be able to open an image in Photolab, select it, retouch it, adjust it and send it back to the Library or another screen. Free and offline: every gesture is computed on your PC by the photocraft engine shipped with Deepotus, with the same pixels as the original application.

Open an image

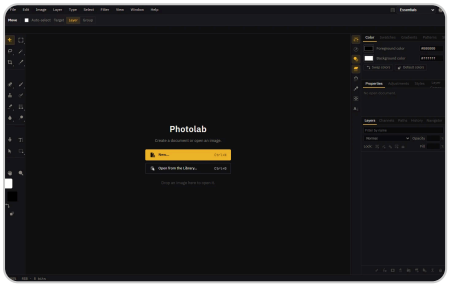

1. An image from the Library

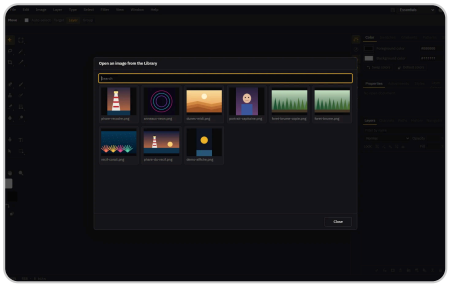

2. Pick it

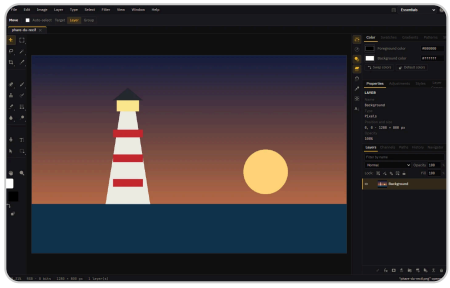

3. One tab per open document

The welcome screen offers a new document or a Library image; dropping a file works too.

1. Open **Photolab** in the app bar, right after Vectorlab. From the Library, **Send to…** › Photolab opens the image directly.
2. With no document, the welcome screen offers **New…** (`Ctrl`+`N`) and **Open from the Library…** (`Ctrl`+`O`). You can also drag a file onto the work area: it enters the Library and opens.
3. A new document is chosen by category (Photo, Film and video, Mobile, Print, Web) and preset: HDTV 1080p, 4K UHD, Instagram square, Story 9:16, A4 300 ppi…
4. Each document has its tab; a dot means it has changed. The status bar at the bottom gives the zoom, mode and depth (RGB · 8 bits), size in pixels and number of layers.

The screen follows the Deepotus language, French or English: menus, tools, panels, dialogs and history.

Pick a tool

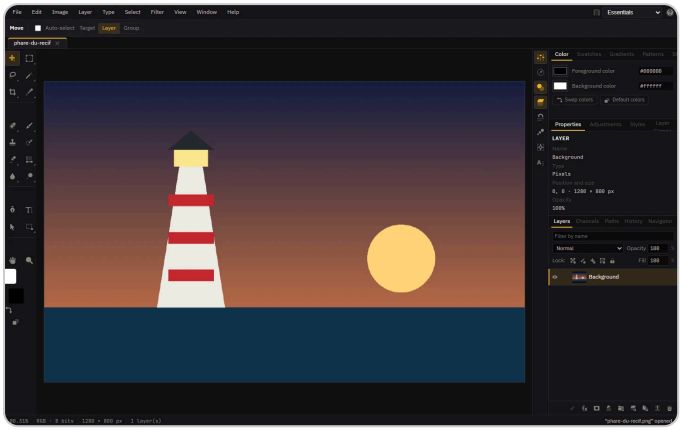

1. Magic Wand (W)

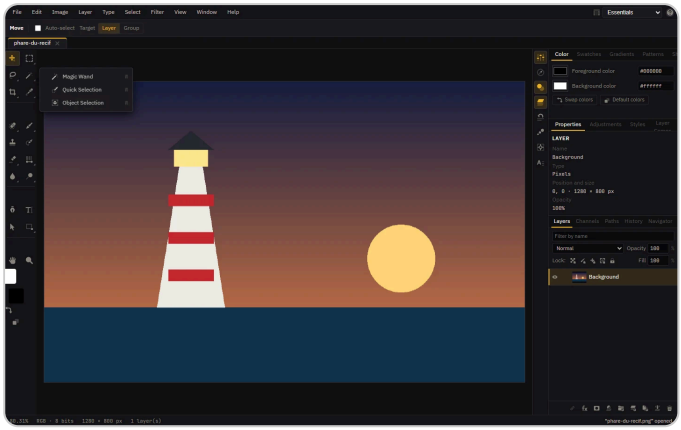

2. The other tools in the group

A slot marked with a corner holds several tools: right-click or long-press to choose.

The toolbar on the left has twenty slots. A slot marked with a corner holds several tools: right-click or long-press to choose. Each tool has its letter (V, M, L, W, C, I, J, B, S, Y, E, G, O, P, T, A, U, H, Z), and the same letter cycles to the next tool in the group. The strip at the top shows the active tool's settings.

Select

Rectangle and ellipse (M), lasso and polygonal lasso (L), magic wand, quick selection and object selection (W). `Shift` adds, `Alt` subtracts.

Crop and measure

Crop (C): draw the frame, `Enter` to cut. The eyedropper (I) picks a colour; the ruler, count and notes are in the next chapter.

Paint

Brush, pencil and mixer brush (B), eraser and background eraser (E), gradient and paint bucket (G). `[` and `]` set the size, a digit the opacity (4 then 5 = 45%).

Repair

Spot healing and healing brush (J), clone stamp (S, `Alt`+click for the source), history brush (Y), dodge, burn, sponge, blur, sharpen and smudge (O).

Text and shapes

Type (T): click for a line, drag for a paragraph. Pen (P), path selection (A), rectangle, ellipse, triangle, polygon, line and custom shapes (U).

Move around

Hand (H) or hold `Space`; zoom (Z) or `Ctrl`+wheel. `Ctrl`+`0` fits the screen, `Ctrl`+`1` shows 100%.

The swatches under the toolbar hold the foreground and background colours: `X` swaps them, `D` restores black and white.

Rectangular marquee (M)

Drag across the picture: a dotted rectangle surrounds the part you are going to change.

Magic wand (W)

Click a colour: the whole patch of that colour is selected at once.

Crop (C)

Draw a frame on the picture then press Enter: everything outside the frame disappears.

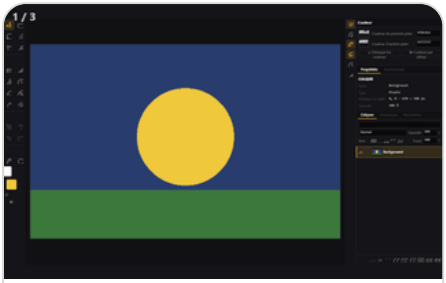

Brush (B)

Drag across the picture: a stroke of your colour is painted wherever you go.

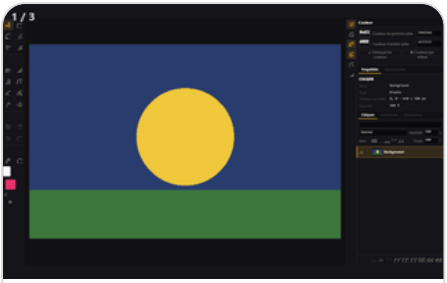

Clone stamp (S)

Alt-click what you want to copy, then paint elsewhere: the same piece of picture reappears under your brush.

Work the selection

The **Select** menu holds **All** (`Ctrl`+`A`), deselect (`Ctrl`+`D`) and inverse (`Ctrl`+`Shift`+`I`). The engine can also find the subject, the sky, a colour range or the focus area. Modify › Border, Smooth, Expand, Contract and Feather refine the outline; a selection can be saved to a channel and loaded back.

Select and Mask… (`Alt`+`Ctrl`+`R`) refines difficult edges, hair included: seven views (onion skin, on black, on white…), radius, smoothing, feather, colour decontamination, then output to a selection, a mask or a new layer.

Transform

Free Transform (`Ctrl`+`T`) sets an eight-handle frame. A corner scales and keeps proportions (`Shift` frees them), an edge stretches one axis, outside rotates (`Shift` by 15°), `Ctrl` distorts or skews. The strip gives X, Y, width, height, angle and interpolation; `Enter` confirms. Edit › Transform also offers 90° and 180° rotations and flips.

Finally, the Edit menu holds cut, copy, paste (including paste in place), fill and stroke, content-aware fill, **Fade** to soften the last filter, and the **Search** palette (`Ctrl`+`K`) that finds any command by name.

Adjust and filter

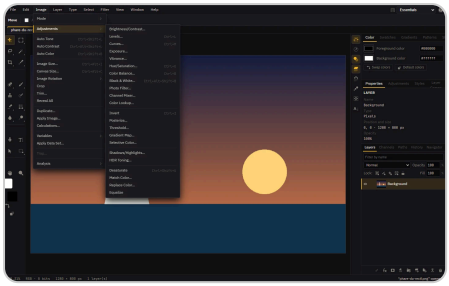

1. Image › Adjustments › Curves…

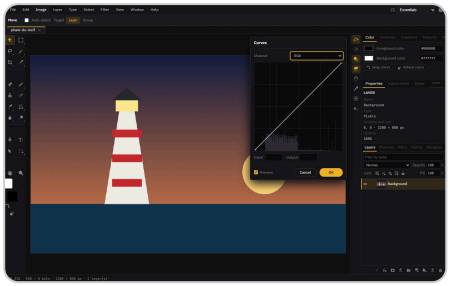

2. Click to add a point

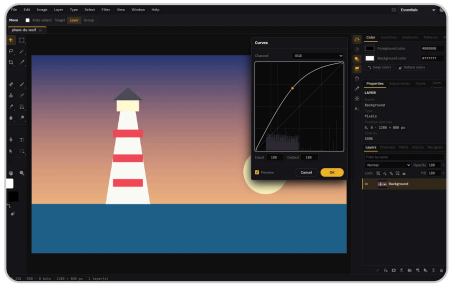

3. The preview follows every move

Image › Adjustments › Curves…: every point you move is recomputed by the engine on a copy, until OK.

1. Open **Image** › **Adjustments** and pick an adjustment: **Levels** (`Ctrl`+`L`), **Curves** (`Ctrl`+`M`), Hue/Saturation (`Ctrl`+`U`), Color Balance, **Black & White**, Exposure, Vibrance, Shadows/Highlights…
2. The dialog shows a **Preview** computed by the engine on a copy. In Curves, a click adds a point, dragging moves it, a double-click removes it, per RGB, red, green or blue channel.
3. **OK** applies and adds a history step; **Cancel** leaves the image untouched.

The **Filter** menu works the same way: blur (including the blur gallery), sharpen, noise, distort, pixelate, render, stylize, plus the **Filter Gallery…** and its 47 stackable artistic effects. Last Filter replays with `Ctrl`+`Alt`+`F`. The Image menu also gives automatic corrections, image and canvas size, rotation, and Image › Mode to switch to grayscale, CMYK, Lab, duotone, indexed colour or 16 bits.

Black and white

One click adds a layer that removes the colours: your picture turns grey, and you can remove it whenever you like.

Save and send

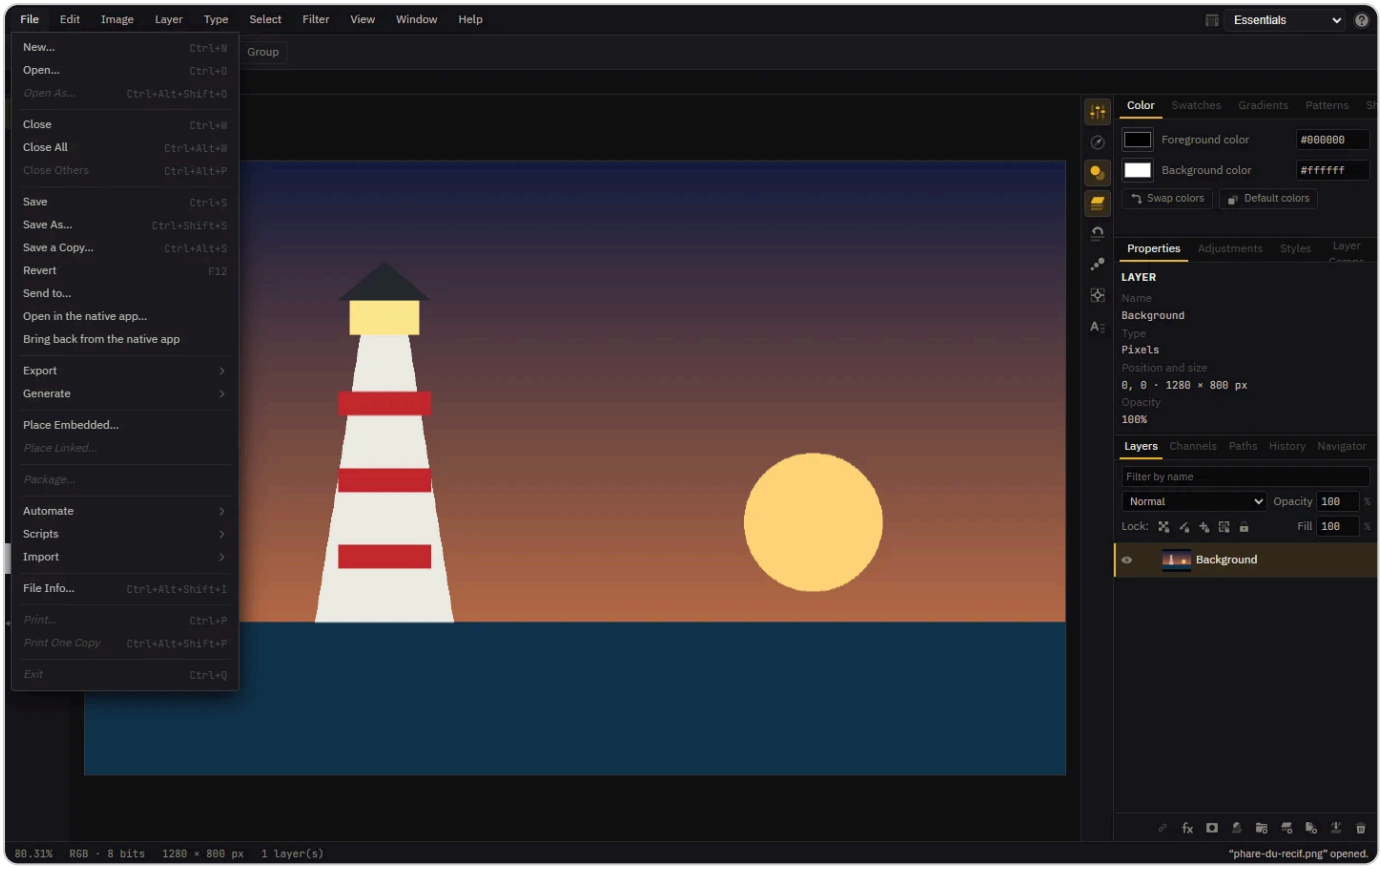

1. To the Library, with layers 2. To Vectorlab, Tile Lab… 3. The fallback: the native app

The File menu: save to the Library (with layers), export, or send to another screen.

1. **Save** (`Ctrl`+`S`) files the image as PNG in the Library, source “Photolab”, with a working file that keeps the layers. It never overwrites: each save is a new file, linked to the original image in its sheet.
2. Reopened from the Library, a Photolab image comes back with its layers.
3. **Save As…** exports to PNG, JPG, WEBP, TIFF, PSD or PCRAFT and downloads it; **Save a Copy…** leaves the document as is; Export › Quick Export as PNG outputs the image in one click, following your preferences.
4. **Send to…** passes the image to another screen, through the Library.

Close (`Ctrl`+`W`), Close All and Close Others ask for confirmation on a changed document. **Revert** (`F12`) reloads the last save. File Info… holds title, author, keywords and copyright.

Good to know

What the screen cannot do yet opens in the native application: **Open in the native app…**, then **Bring back from the native app** returns the document, layers included.

If the engine must restart, the screen says so and returns to the welcome screen: unsaved documents are lost. Save often. The installer checks every engine file; Help › **About Photolab** gives its version and licences.

Shortcuts

Gesture	Shortcut
Undo, redo	`Ctrl`+`Z`, `Ctrl`+`Shift`+`Z`
New, open, save	`Ctrl`+`N`, `Ctrl`+`O`, `Ctrl`+`S`
Free transform	`Ctrl`+`T`
Levels, curves, hue	`Ctrl`+`L`, `Ctrl`+`M`, `Ctrl`+`U`
Tip size, hardness	`[` `]`, `Shift`+`[` `Shift`+`]`
Search a command	`Ctrl`+`K`

IMAGES AND DRAWING

Layers, masks and workspaces

You will be able to work with non-destructive layers in Photolab: adjustments, masks, styles and smart objects, then arrange the screen your way with panels, workspaces, rulers and preferences. Everything is free and computed on your PC.

The Layers panel

Each layer has its eye, thumbnail and name; double-click the name to rename it. Click to pick, `Ctrl`+click to add, `Shift`+click for a range. Drag a row to place it above, below or inside a group. At the top of the panel: a field to filter by name, the blend mode (27 modes), opacity, fill and locks (transparent pixels, pixels, position, nesting, all).

The panel footer creates a layer or a group, links layers, adds a mask, duplicates, merges down and deletes. The **Layer** menu goes further: new layer via copy (`Ctrl`+`J`), group, clipping mask (`Ctrl`+`Alt`+`G`), arrange, align and distribute, merge (`Ctrl`+`E`) or flatten, rasterize, matting.

Add an adjustment layer

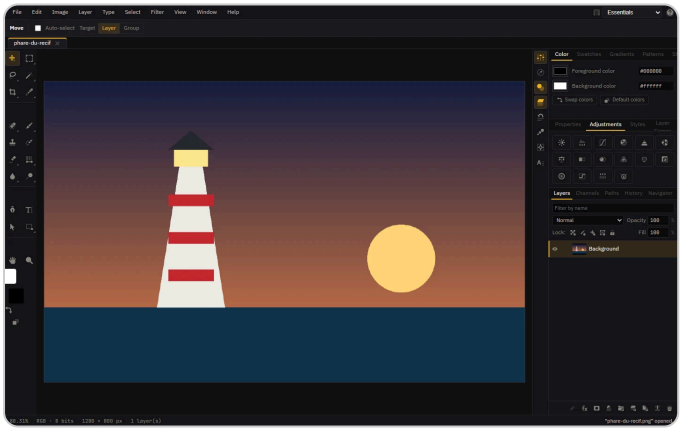

1. Black & White

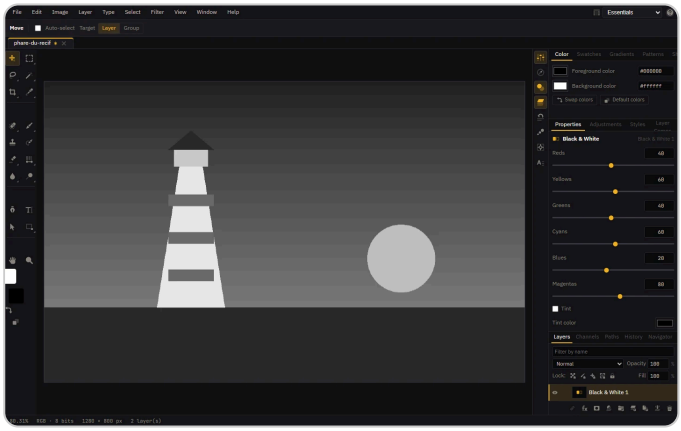

2. A non-destructive adjustment layer 3. Its sliders in Properties

Adjustments tab: one click adds an adjustment layer, editable at any time in Properties.

1. Open the **Adjustments** tab, next to **Properties**.
2. Click an adjustment, for example **Black & White**: an adjustment layer is added above the active layer. The image below is not touched.
3. Its sliders appear in **Properties**; Curves and Levels have their own editor there. Come back at any time by picking the layer.

Sixteen adjustments work this way (brightness, levels, curves, hue, colour balance, channel mixer, gradient map…), as well as fill layers: solid colour, gradient or pattern.

Masks, styles, smart objects

Layer mask

The mask button reveals the selection (or everything); `Alt` does the opposite. On the mask thumbnail: `Shift`+click disables it, `Alt`+click shows it alone, `Ctrl`+click loads it as a selection.

Vector mask

On a layer that already has a mask, the same button adds a vector mask, or Layer › Vector Mask from the active path.

Layer style

Drop shadow, glows, bevel, stroke, satin, overlays: a dialog with preview. A style can be copied, pasted or cleared; the **Styles** panel offers ready ones (neon, chrome, gold…).

Smart filters

Filter › Convert for Smart Filters: every filter applied afterwards appears under the layer, with its eye; double-click to reopen its settings, drag to change the order.

A single layer can be exported: Layer › Export As… (PNG, JPG or WEBP, at the chosen scale, to download or the Library) or quick export as PNG.

Black and white

One click adds a layer that removes the colours: your picture turns grey, and you can remove it whenever you like.

The panels

On the right, panels are arranged in tabbed groups; the rail's icon column collapses or reopens each group, and the **Window** menu ticks the visible ones.

Colour and swatches

Color (`F6`) with hex code, Swatches (personal groups kept by Deepotus), Gradients and Patterns.

History and info

History: a click on a step goes back to it. Navigator, per-channel Histogram, Info (`F8`) under the pointer, Channels and their shortcuts (`Ctrl`+`2`, `Ctrl`+`3`…).

Brushes

Tips with search, Brush Settings (`F5`), Clone Source and Tool Presets (e.g. soft eraser 60 px).

Text and paths

Character, Paragraph, Glyphs, character and paragraph styles; Paths, Shapes; Layer Comps to switch from one state to another.

Choose a workspace

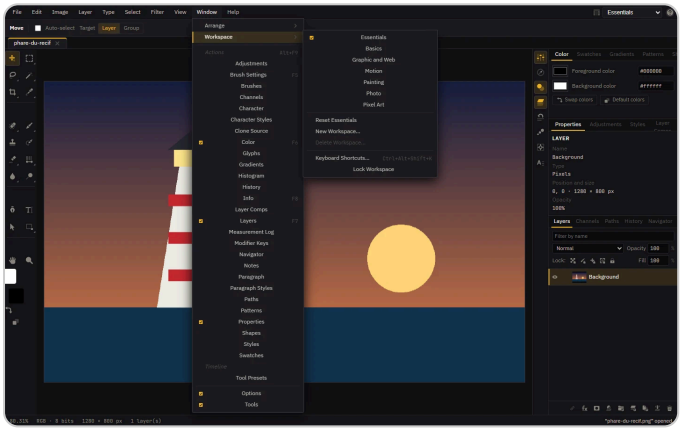

1. Seven built-in workspaces

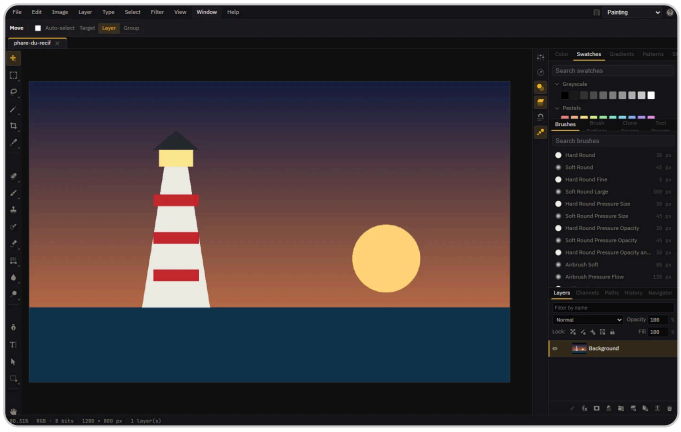

2. Painting: brushes and colours up front

Window › Workspace arranges panels and tools for a task; the layout is remembered per workspace.

1. Open **Window** › **Workspace**, or the list at the top right.
2. Pick one of seven layouts: **Essentials** (default), Basic, Graphic and Web, Motion, **Painting**, **Photo**, **Pixel Art**. Panels and tools arrange for the task.
3. What you move, collapse or hide is remembered for that workspace. Reset restores it; Lock Workspace stops remembering.
4. New Workspace… saves your layout (up to twenty), kept in the Deepotus data folder.

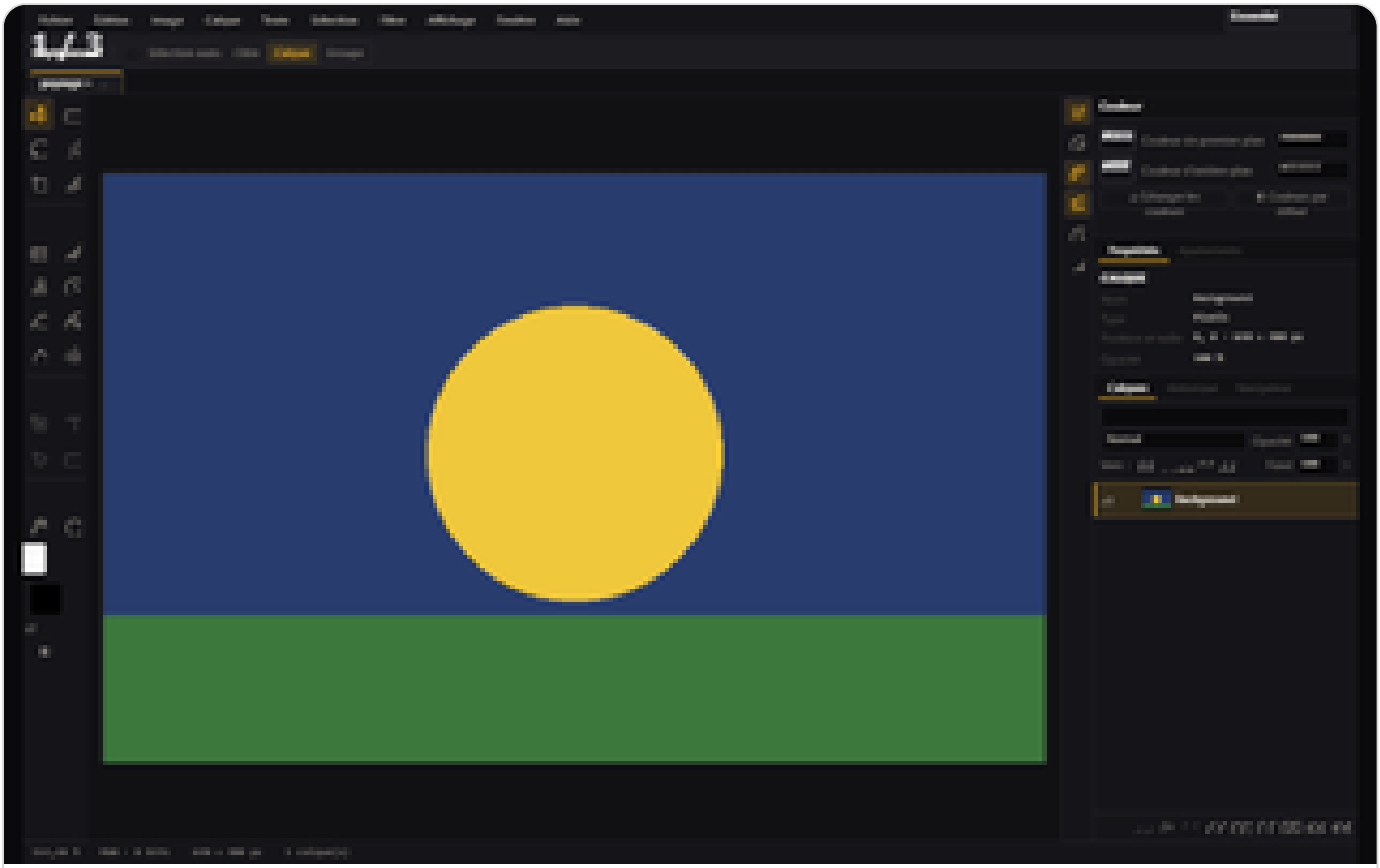

Workspace

Pick a workspace in the list: panels and tools rearrange for what you want to do.

Rulers, grid and guides

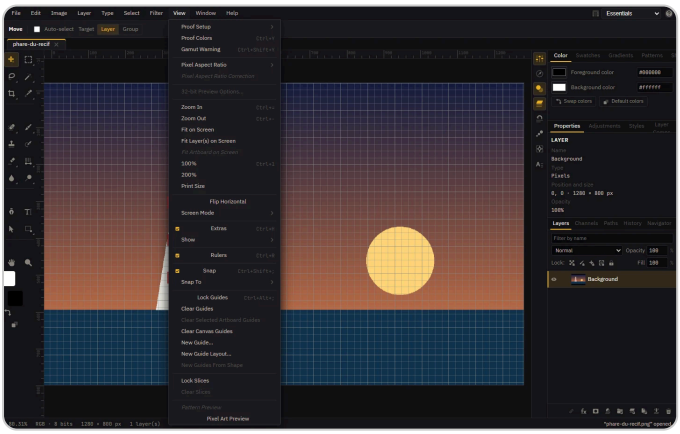

1. View › Rulers

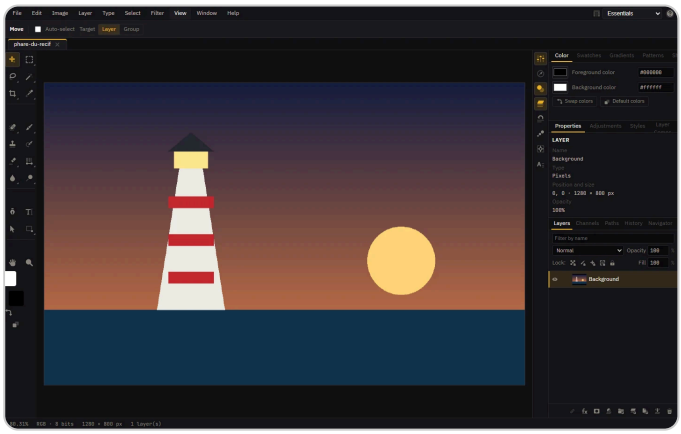

2. Graduated rulers and grid

View › Rulers (Ctrl+R) then Show › Grid: drag from a ruler to place a guide.

In **View** : **Rulers** (`Ctrl` + `R`), Show › **Grid** (`Ctrl` + `'`) and **Guides** (`Ctrl` + `;`). Drag from a ruler to place a guide; the Move tool moves it, or removes it by dragging it out. Snap (`Ctrl` + `Shift` + `;`) aligns selections and layers to guides, grid, layers or document edges. Extras (`Ctrl` + `H`) hides everything at once.

The same menu offers zooms, horizontal flip of the view (the document does not change), pixel art preview with crisp pixels and screen modes: `F` goes to the next one, `Esc` back to standard mode.

Measure, annotate, slice

Ruler (I)

Drag to measure distance and angle; `Alt` from an end makes a protractor. Straighten Layer puts the line level.

Count (I)

A click places a numbered mark, `Alt`+click removes it. Measurements are recorded in the log, exportable as CSV for a spreadsheet.

Note (I)

A note with author and colour, placed on the image; the Notes panel steps from one to the next. File › Import › Notes… copies them from another document.

Slices (C)

Cut the image for the web, or from guides; slice selection moves them, divides them and sets their options (link, alt text).

Preferences and shortcuts

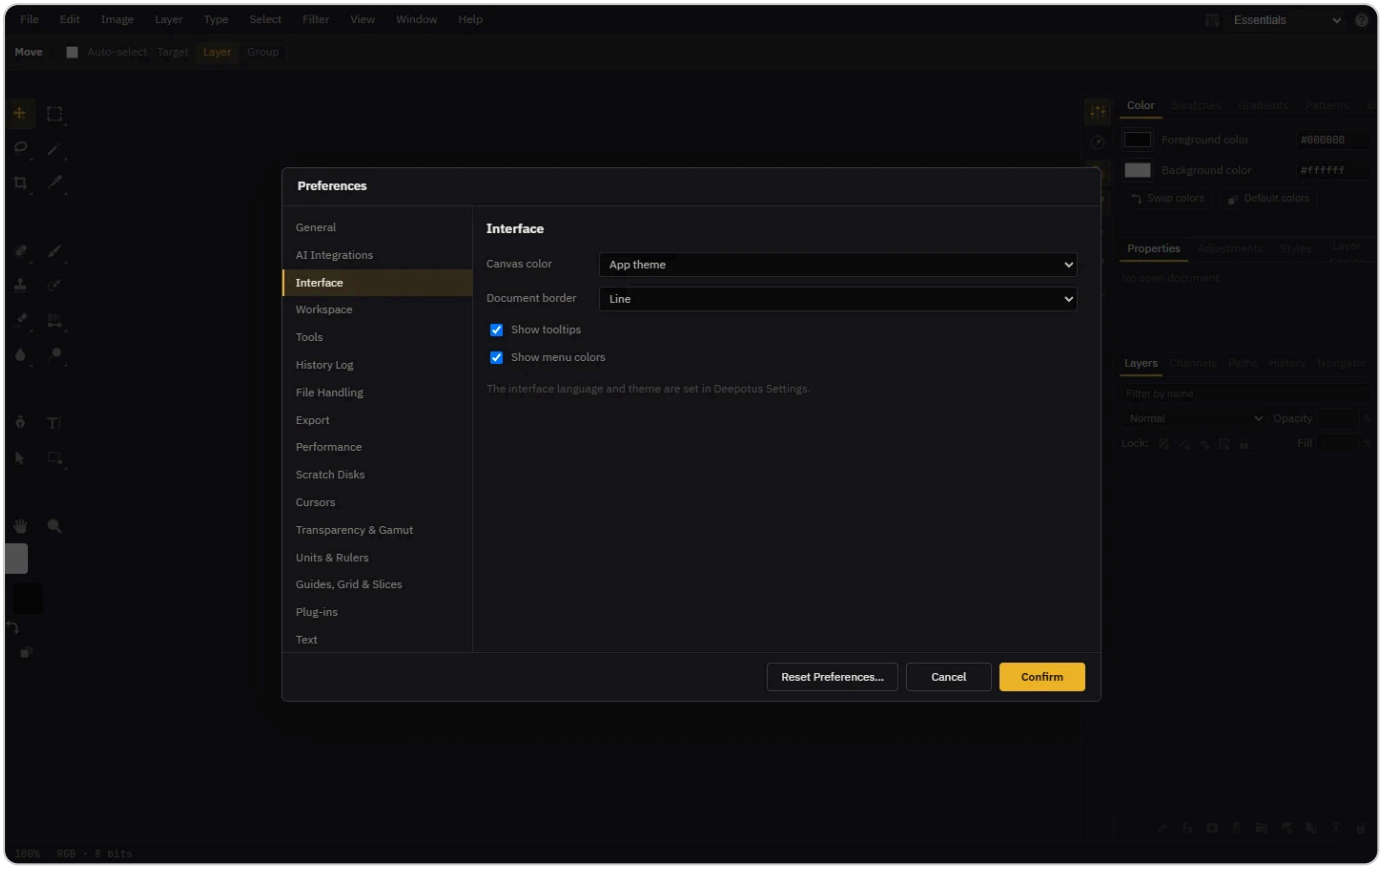

1. One section per menu entry

Edit › Preferences: eighteen sections, kept in the Deepotus data folder.

Edit › **Preferences** opens eighteen sections, saved in the Deepotus data folder: interface (canvas colour, document border, tooltips), tools, file handling and quick export, cursors, transparency, units and rulers, guides and grid, text. Some sections simply say what Deepotus handles in their place (performance, scratch disks…). The reset button restores all defaults, after confirmation.

Keyboard shortcuts

Change a command's shortcut or a tool's letter; a conflict is named and removed from the other command. Summarize… downloads the list as an HTML page.

Menus and presets

Hide or colour menu entries. The preset manager organises brushes, shapes and patterns; they export and re-import as .json.

Touch screen

Modifier Keys shows `Shift`, `Ctrl` and `Alt` on screen: a click applies to the next gesture, a double click holds the key.

Good to know

Workspaces, preferences, shortcuts and personal swatches are kept by Deepotus in its data folder: they follow you through a transfer to another PC.

IMAGES AND DRAWING

Draw vectors (Vectorlab)

You will be able to draw a poster, a logo or a sheet in vectors in Vectorlab: shapes, pen, text, appearance, layers, symbols, then send it to other screens. Free, everything is computed on your PC; only AI illustration is paid, price shown first.

Open Vectorlab

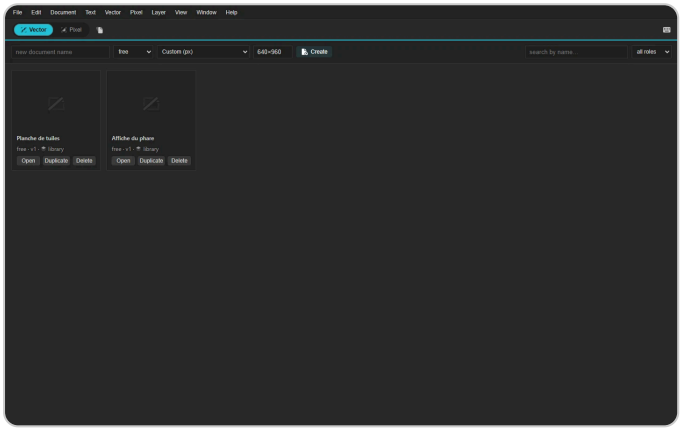

1. Open a document

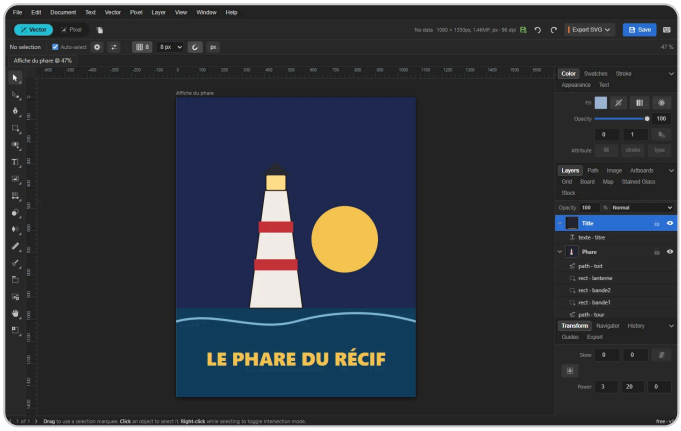

2. Two personas: Vector and Pixel 3. Tools, grouped by family

The home lists your documents: one click on a card opens it; New creates one.

1. Open **Vectorlab** in the app bar. The home lists your documents: one click on a card opens it. Pick a format (square, 16:9, A4, playing card…) then **Create** for a new document.
2. A document also comes from elsewhere: **Send to…** › Vectorlab from the Library or Photolab, “Edit this face in Vectorlab” from Card Forge, “Open in Vectorlab” from a Templates layout.
3. At the top left, two pills: **Vector** to draw, **Pixel** to retouch pixels (see the next chapter). Each has its own tools and panels.

The screen follows Affinity's conventions: ten menus at the top (**File**, **Edit**, Document, **Text**, **Vector**, Pixel, **Layer**, **View**, **Window**, **Help**), a context bar showing the selection and the current tool's settings, tools on the left, the tabbed panel stack on the right and a status bar that recalls the tool's gestures. The document tab shows “name @ zoom” and a star when it is not saved.

Pick a tool

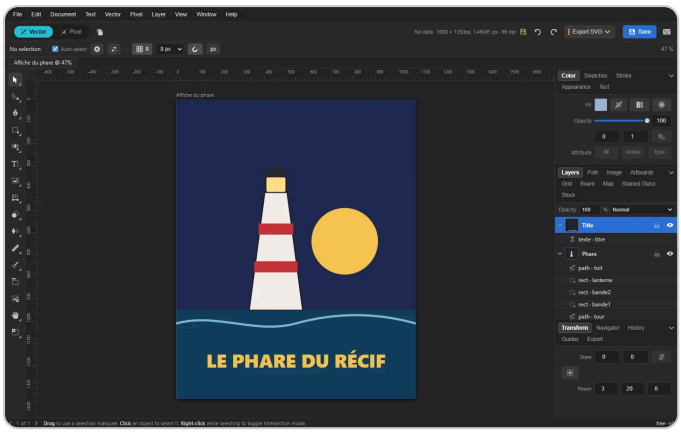

1. The Shapes family

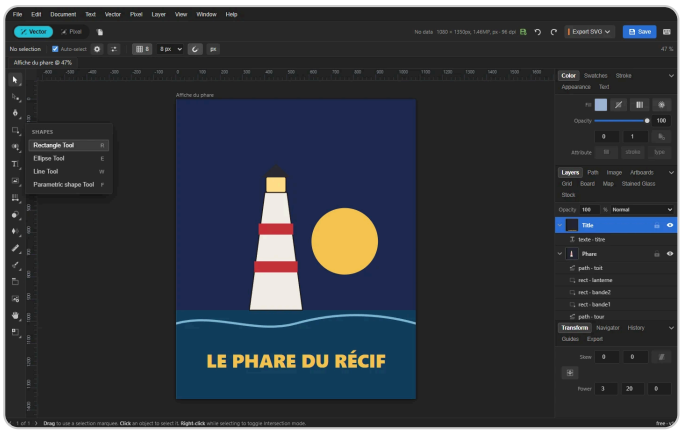

2. Rectangle, ellipse, line, parametric shape

The corner triangle or a long press opens the family; the tool's letter picks it from the keyboard.

Tools are grouped in families: only the current tool of each family is visible. The corner triangle or a long press opens the family, right-click opens the tool's settings, the letter picks it from the keyboard.

Select (V)

The marquee only takes objects fully inside (`Alt`: those it touches). `Shift` constrains, `Ctrl` resizes from the centre, `Alt`+drag copies, `Ctrl`+click reaches an object inside a group, `Tab` goes to the next one.

Pen (P) and Node (N)

Click = sharp node, drag = smooth node, click the first node to close. Four modes: Pen, Smart, Polygon, Line. The Node tool pulls handles, bends a segment, adds (double-click) or deletes nodes.

Shapes (R, E, L, F)

Rectangle, ellipse, line and parametric shapes (polygon, star, cog, arrow, donut, spiral) adjustable by handles. A single click places a shape; the Corner tool (C) rounds a corner.

Freehand

Smoothed pencil (B), vector brush (J) with flat, tapered or calligraphic profile, eraser (W), knife (X) and shape builder (S) that combines or removes areas by clicks.

Text (T)

Drag = text sized to the height you pull; click on a curve = text on that path; double-click to edit in place. The text frame flows text and flags overflow in red.

Eyedropper, measure, view

Eyedropper (I) with loupe (`Ctrl`: the whole style), measure (M), hand (H), zoom (Z, `Alt` to zoom out). `Ctrl`+`0` fits the page, `Ctrl`+`1` shows 100%.

Set the appearance

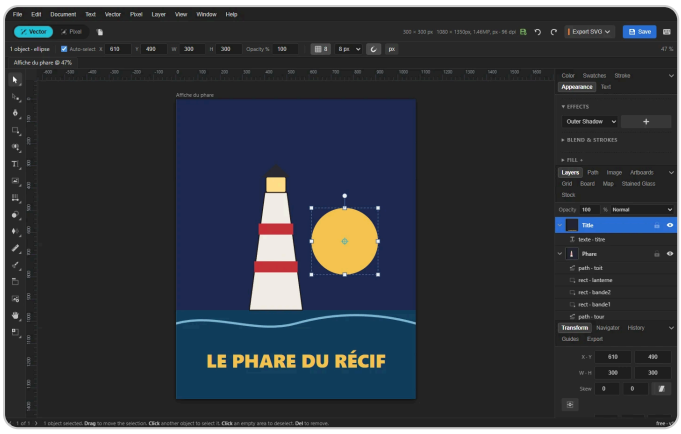

1. One click selects the sun

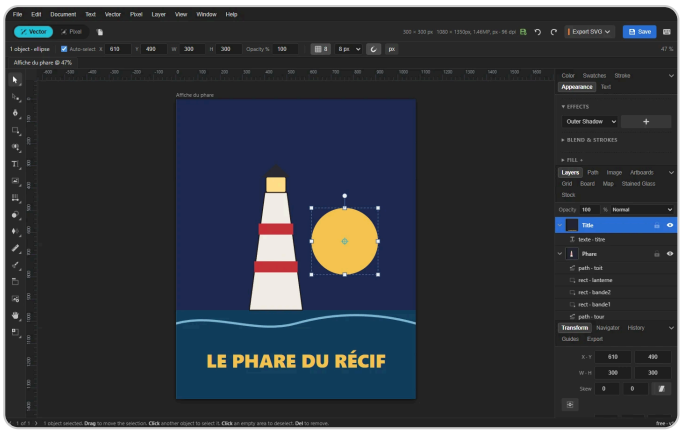

2. Fill, stroke, gradients, effects

Select an object: the context bar and the Appearance tab show its fill, stroke and opacity.

1. Select an object. The context bar gives X, Y, W, H and opacity; fields accept calculations (100/3).
2. The **Appearance** tab sets fill and stroke from the swatches, no fill or no stroke, width, opacity. **Stroke** and **Transform** have their own tabs (skew, movable pivot).
3. **Appearance +** goes further: linear, radial and conical gradients, transparency, patterns (hatching, dots, checkerboard), multiple strokes, six effects, 16 blend modes, global colours, harmonies, object styles, paste inside a shape (clipping mask), text in a frame or on a path.

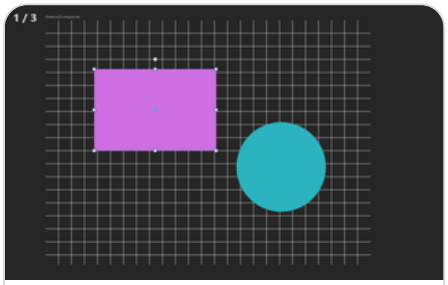

Conical
The color spins around the middle of your shape, like the spokes of a wheel.

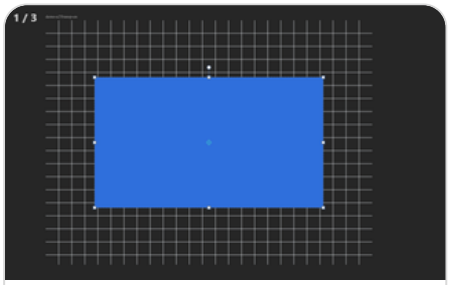

Transparency
Your shape slowly fades from one side to the other, as if it vanished into fog.

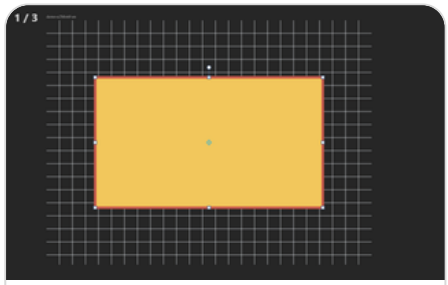

Pattern
Your shape fills up with little repeated stripes, like striped fabric.

Frame

Your long line of text fits into a box and wraps by itself, like in a book.

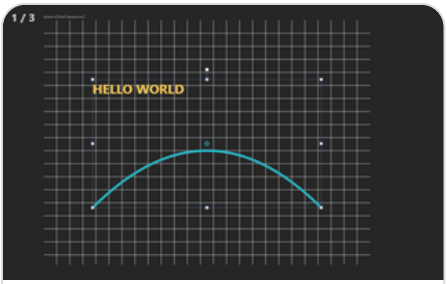

On path

Your words lie down on the line you drew and follow it wherever it goes.

An offset stroke and a text on a path stay **live**: they recompute when their source changes, until you detach them.

Layers and symbols

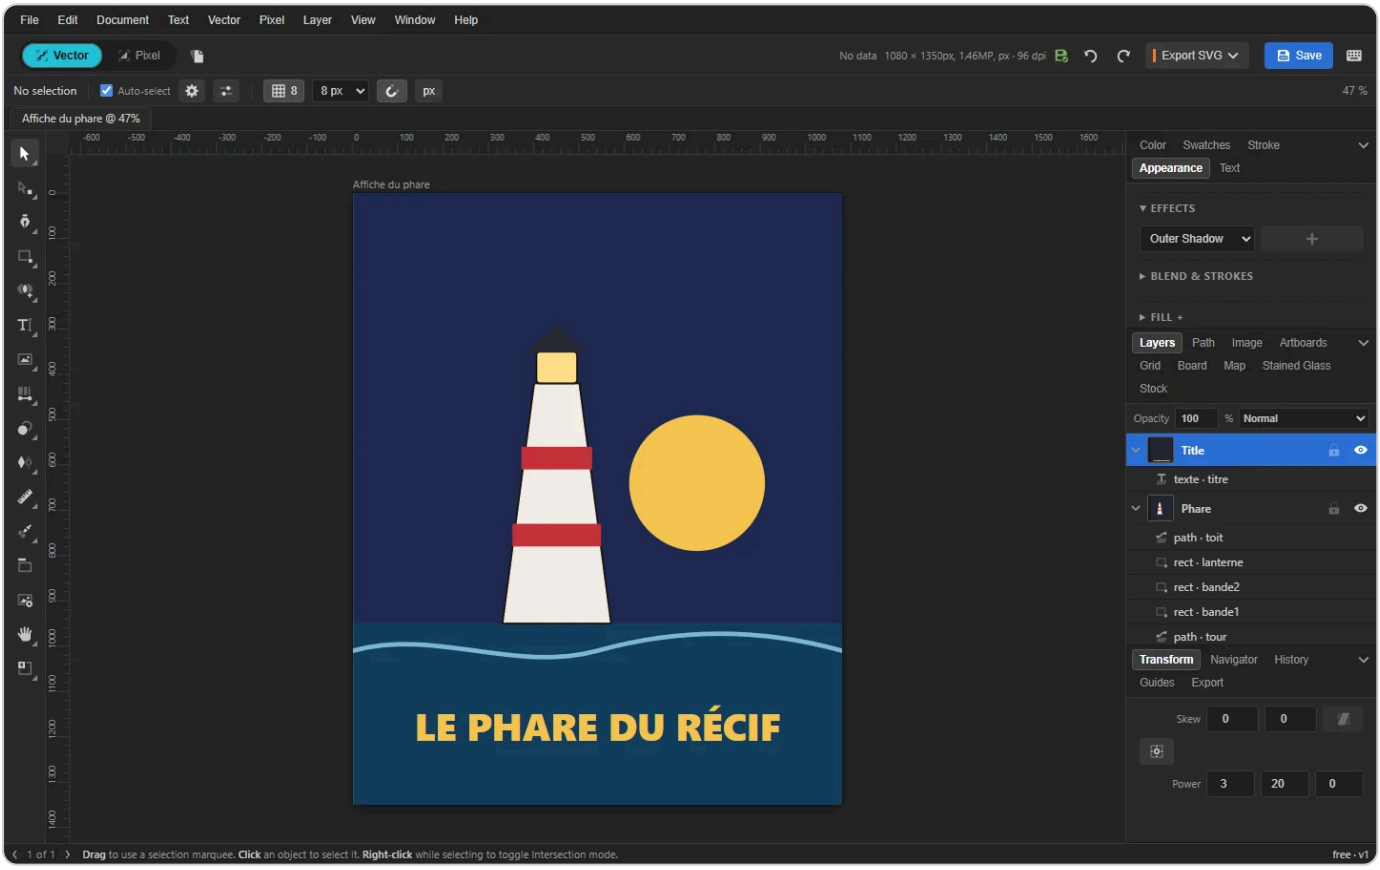

1. Background, Lighthouse, Title: three layers

Each layer unfolds into a tree of its objects, with thumbnail, lock, eye, opacity and blend mode.

The **Layers** tab shows the layer's opacity and blend mode at the top, then each layer with its thumbnail, lock and eye; a triangle unfolds it into a tree of its objects, clips, masks and effects. The **Layer** menu creates, renames, groups (`Ctrl` + `G`), locks, hides and orders.

A **symbol** is created with Vector › `Create Symbol`; its instances follow every change. Double-click an instance to edit it in place, then `Done` or `Discard`.

The Vector menu

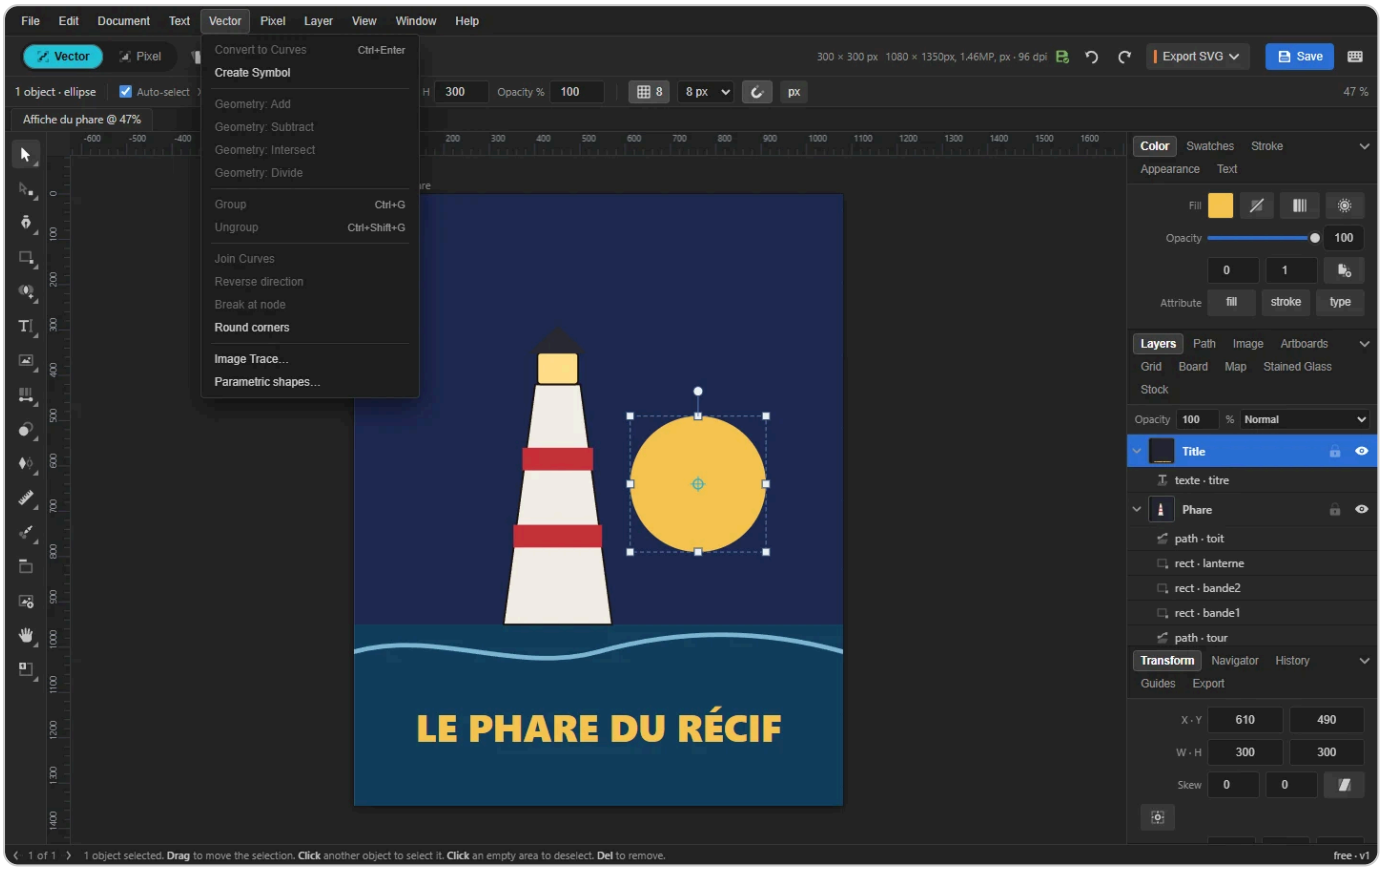

1. Combine shapes 2. Vectorise an image, locally

Geometry (union, subtract, intersect, divide), curves, symbols and image trace.

Geometry

Add, subtract, intersect and divide shapes; join, reverse, break at node; `Round corners`.

Convert to curves

A text becomes paths (`Ctrl`+`Enter`), as one or one path per glyph.

Image trace

`Image Trace...` vectorises a placed image (colours, smoothing, threshold), with preview, on your PC.

The document

The Document menu holds `Setup...` (size, resolution, unit, background), `Artboards...` exported separately, `Margins and bleed...` (guides never exported), the `Document grid...` (square, iso, triangular, hexagonal, snapping), `Board and terrains...` to paint tiles (K) and the `Real map...`. View › `Snap to objects` snaps to edges, centres and even gaps.

File › **Place image…** or **Place from Library…** places an image; the Stock tab shows the Library with search, one click places the image. An image can be cropped, locked and also receives effects.

Good to know

The document is saved as a draft every 30 seconds: when you reopen it, a window offers to restore it or to start from the server. History keeps 1000 steps; Edit › **Add Snapshot** marks a state to come back to.

Hover an option for a second, or click the "?" next to it, for its animated card. `F1` lists every shortcut; Edit › Settings (`Ctrl`+`,`) sets tooltips, grid and snapping.

AI illustration (AI family, or File › New image processing) is paid: the price is shown first, within your caps. Everything else is free.

Shortcuts

Gesture	Shortcut
Undo, redo	`Ctrl`+`Z`, `Ctrl`+`Shift`+`Z`
Duplicate, group, ungroup	`Ctrl`+`D`, `Ctrl`+`G`, `Ctrl`+`Shift`+`G`
Fit page, 100%, grid	`Ctrl`+`0`, `Ctrl`+`1`, `G`
Vectorlab home	`Ctrl`+`Alt`+`H`
All shortcuts	`F1`

IMAGES AND DRAWING

Pixel art, export, maps and 3D objects

You will be able to retouch pixels and make pixel art in Vectorlab, export your drawings in the right format, build a map from real data and turn a drawing into parts to 3D-print. Everything is free and computed on your PC.

The Pixel persona

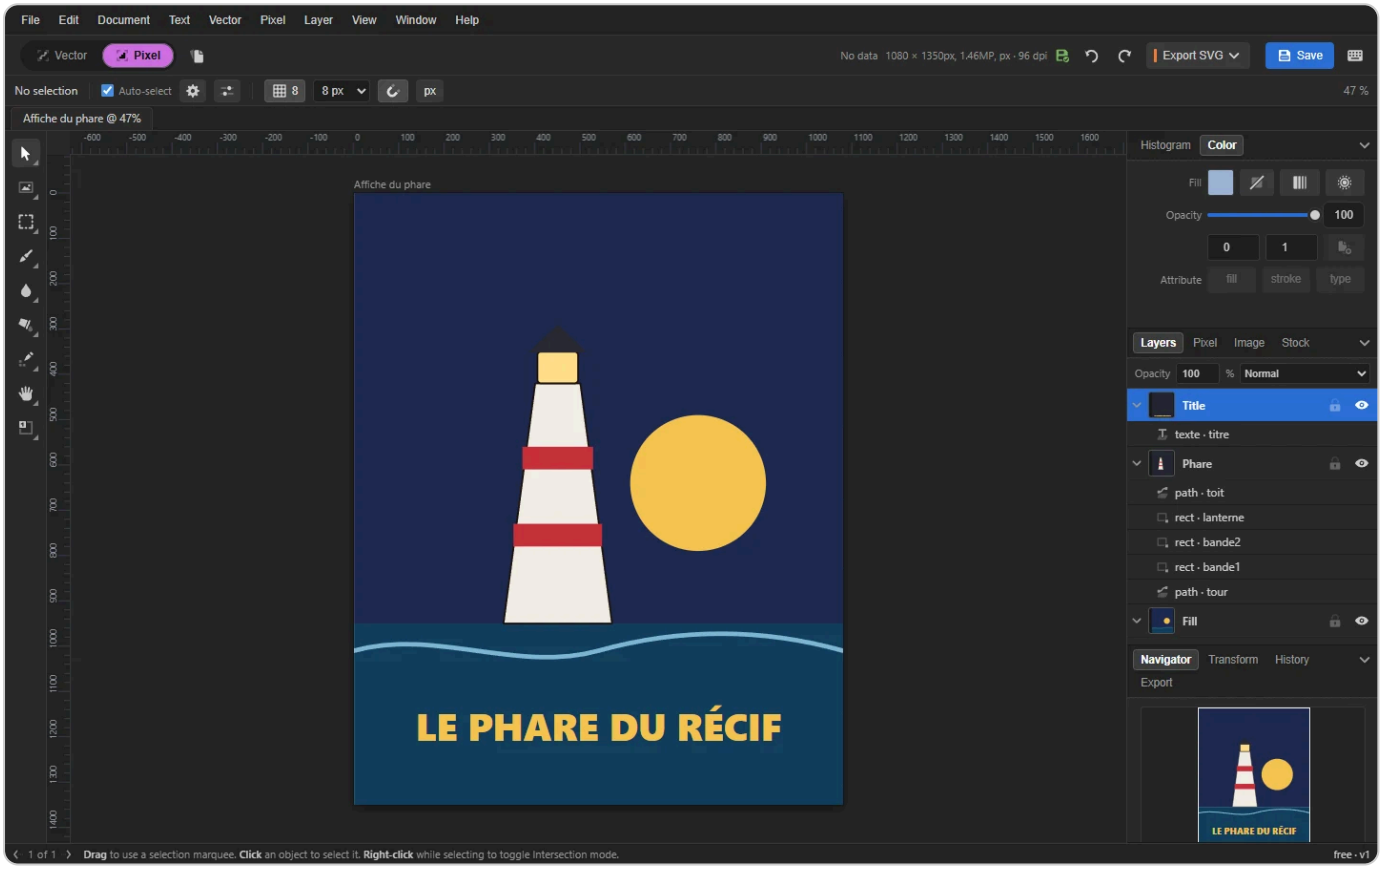

1. Pixel pill 2. The pixel tools

Same document, other tools: pixel selections, brush, eraser, bucket, pixel pencil, Pixel-art and Histogram panels.

1. Click the **Pixel** pill (or View › **Pixel Persona**). The same document keeps its layers; tools and panels change. From the Game Assets hub, the Pixel card opens this persona directly.
2. To retouch a placed image, pick it then Pixel › **Edit pixels**.
3. Select pixels: rectangle (M), lasso (L), magic wand (W) or by colour; the Pixel menu grows, contracts or inverts them (`Ctrl` + `I`).
4. Paint: brush (B) and eraser (E) with radius, hardness and round or square shape, paint bucket (G) contiguous or global, clone stamp (C), blur, dodge and burn. Right-click = secondary colour, `X` swaps them, `Alt` +click picks.

Pixel › **Adjustments and filters…** gives levels, curves, hue and saturation, black and white, threshold and blur. Every retouch undoes with `Ctrl`+`Z`, like the rest of the document. The Histogram tab shows the four channels of the edited image.

Make pixel art

The Pixel-art panel sets tile size, symmetry, palette (the same as Sprite Lab and Tile Lab), quantization and the pixel grid. The pixel pencil (K) draws without double pixels in diagonals; `Shift`+click draws a straight segment from the last point.

Trace a reference

Set as reference fades and locks an image with its grid; a pixel layer sits on top to draw cell by cell, the eyedropper takes a cell's exact, average or dominant colour.

Pixelate

Pixelate turns an image or a vector selection into big pixels. Rasterize adds palette and dithering (ordered or Floyd-Steinberg).

Tiles and seams

The 3×3 preview scores a tile's seams; the 2:1 iso tile limits painting to the diamond, with symmetry along its edges.

Animation

The timeline shows frames: play (`Enter`), loop, FPS, red (previous) and blue (next) onion skin. Export as a PNG + JSON sheet or a strip.

Finishing

Dark outline sets the drawing off the background, **Sharpen** boosts contrast; PNG export scales from ×1 to ×16 without blur.

To the other labs

Pixel › **Pixel art to Tilelab…** sends the drawing to Tile Lab; the Pixel-art panel also sends it to Sprite Lab.

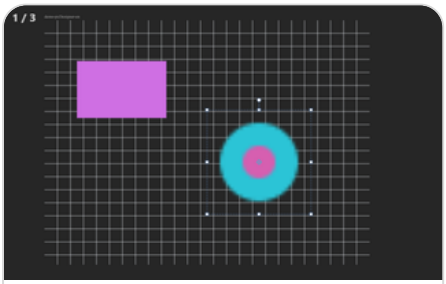

Set as reference
Your picture becomes a reference: it fades, stays still, and a grid of squares lies on it to guide you.

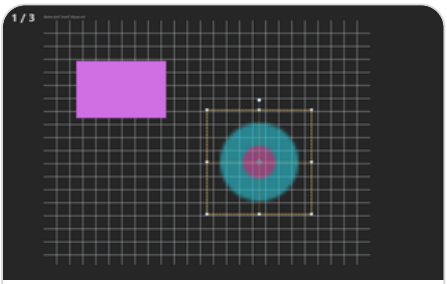

Create pixel layer
A see-through sheet lies on your reference picture: you draw on it, square by square.

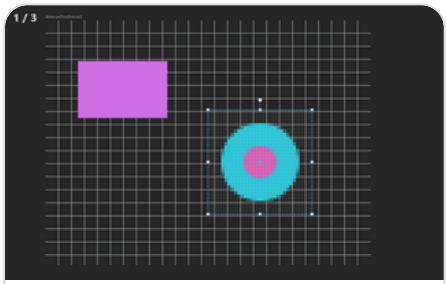

Pixelate image
Your picture is now made of big squares, like bricks: as many squares as the tile size.

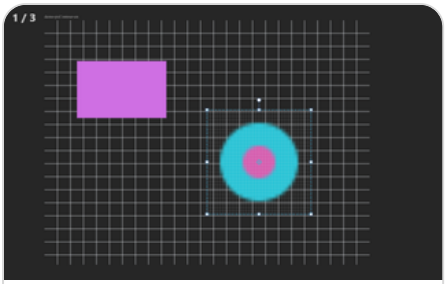

Dark outline

A dark line goes around your drawing so it stands out from the background.

Export

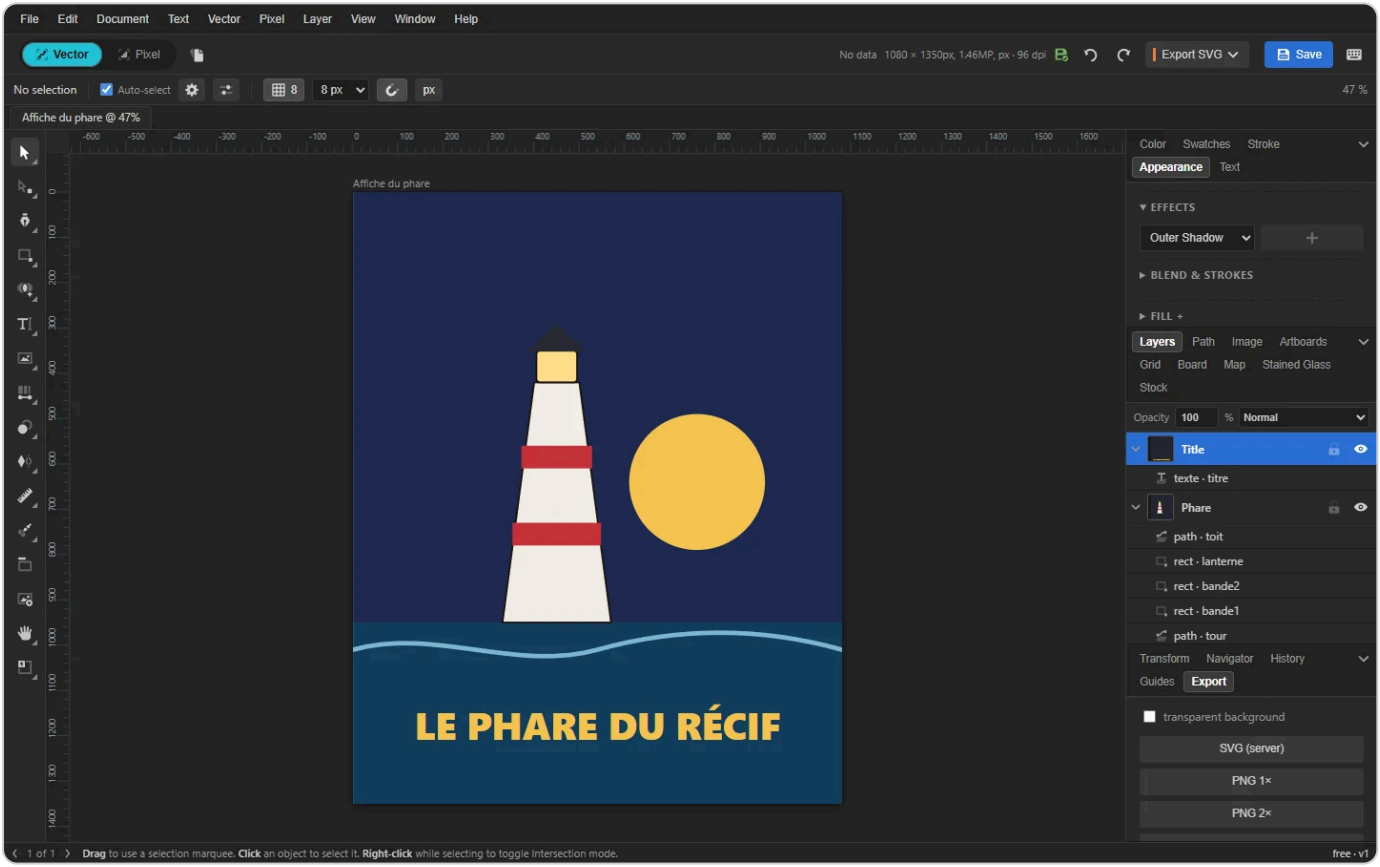

1. PNG, JPEG, WebP, SVG, PDF, DXF

Shared by both personas: formats, resolutions, slices and print presets.

1. Open the **Export** tab (File › **Export…** or Window › Export), shared by both personas.
2. Choose what goes out: the whole document, each artboard, each visible layer, each selected object, or the slices drawn with the Slice tool.

3. Choose the format (PNG, JPEG, WebP, SVG, PDF, DXF) and one or more resolutions (@1x, @2x…). File names build themselves; a batch export outputs every slice, resolution and format at once.

Vector PDF

By default: text as outlines, page at physical size, vector gradients and patterns, crop marks and bleed for the printer. A message says what had to be rasterized and why.

DXF for cutting

In millimetres, text as glyph outlines, symbols and groups included.

Quick exports

File › **Export SVG** (embedded images), **Export PNG 2×**, and **Send to Bible…** for a chapter.

A real map

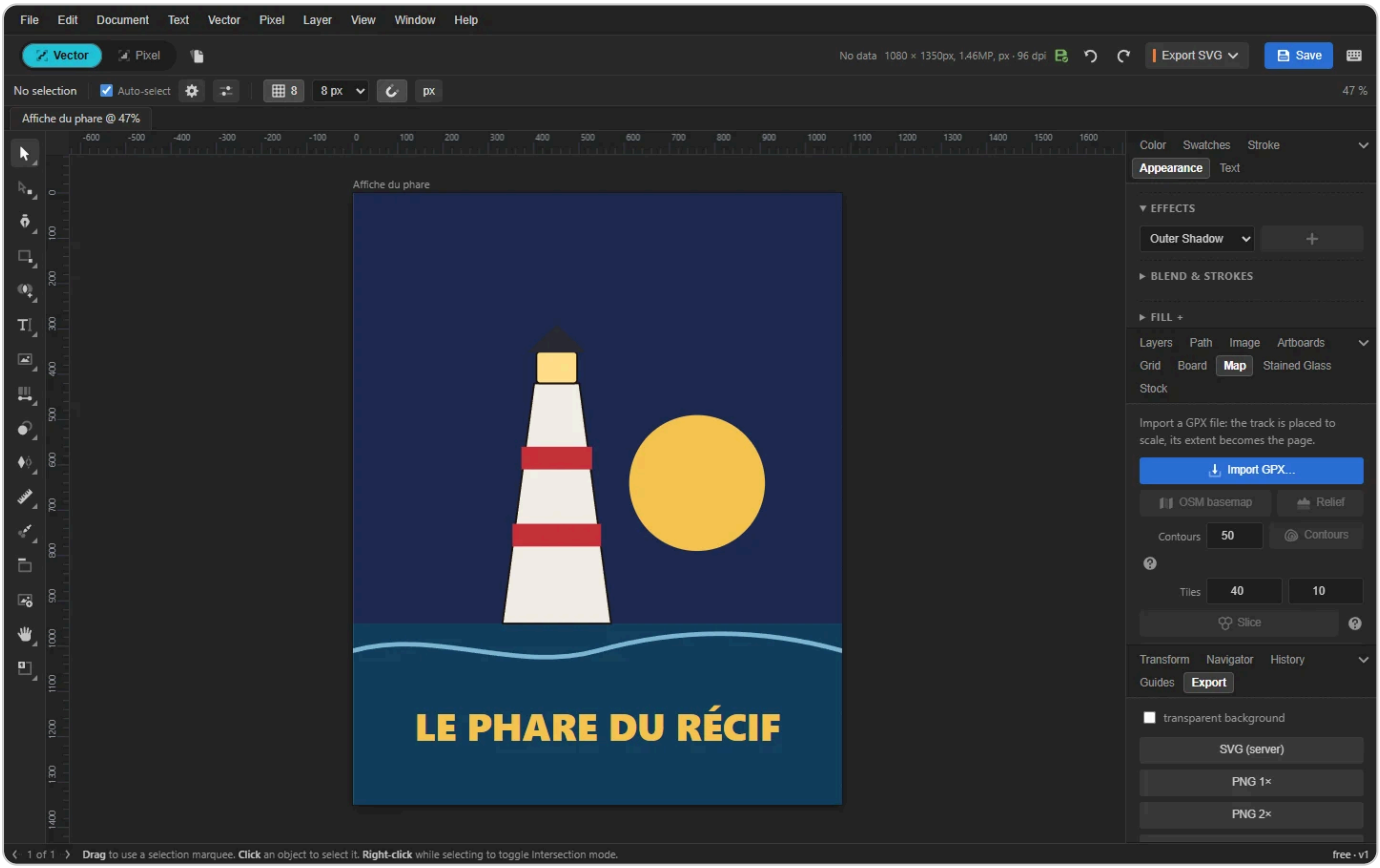

1. The Map panel

Document › Real map…: relief, contour lines, OpenStreetMap background and GPX tracks, no key, no cost.

Document › **Real map…** opens the Map panel: no key, no cost, it lays a shaded relief, draws contour lines as vector paths at the chosen interval, adds an OpenStreetMap background with its attribution, and slices

the map into hexagonal tiles coloured by altitude. File › **Import GPX…** places a track at true scale, keeps its altitude (profile, length, climbs) and can turn it into a raised ribbon.

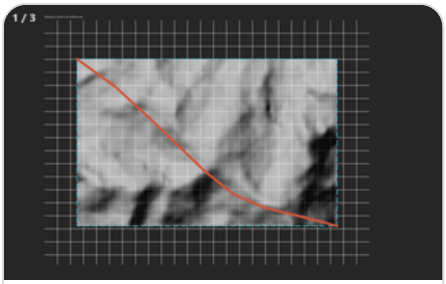

Contours
Lines join the places of the same height on your map, like the steps of a mountain.

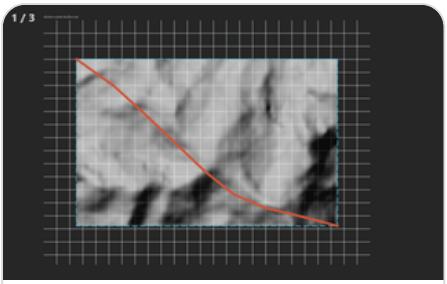

Slice
Your map turns into six-sided tiles, colored by height: water, meadow, forest, mountain.

Print in 3D

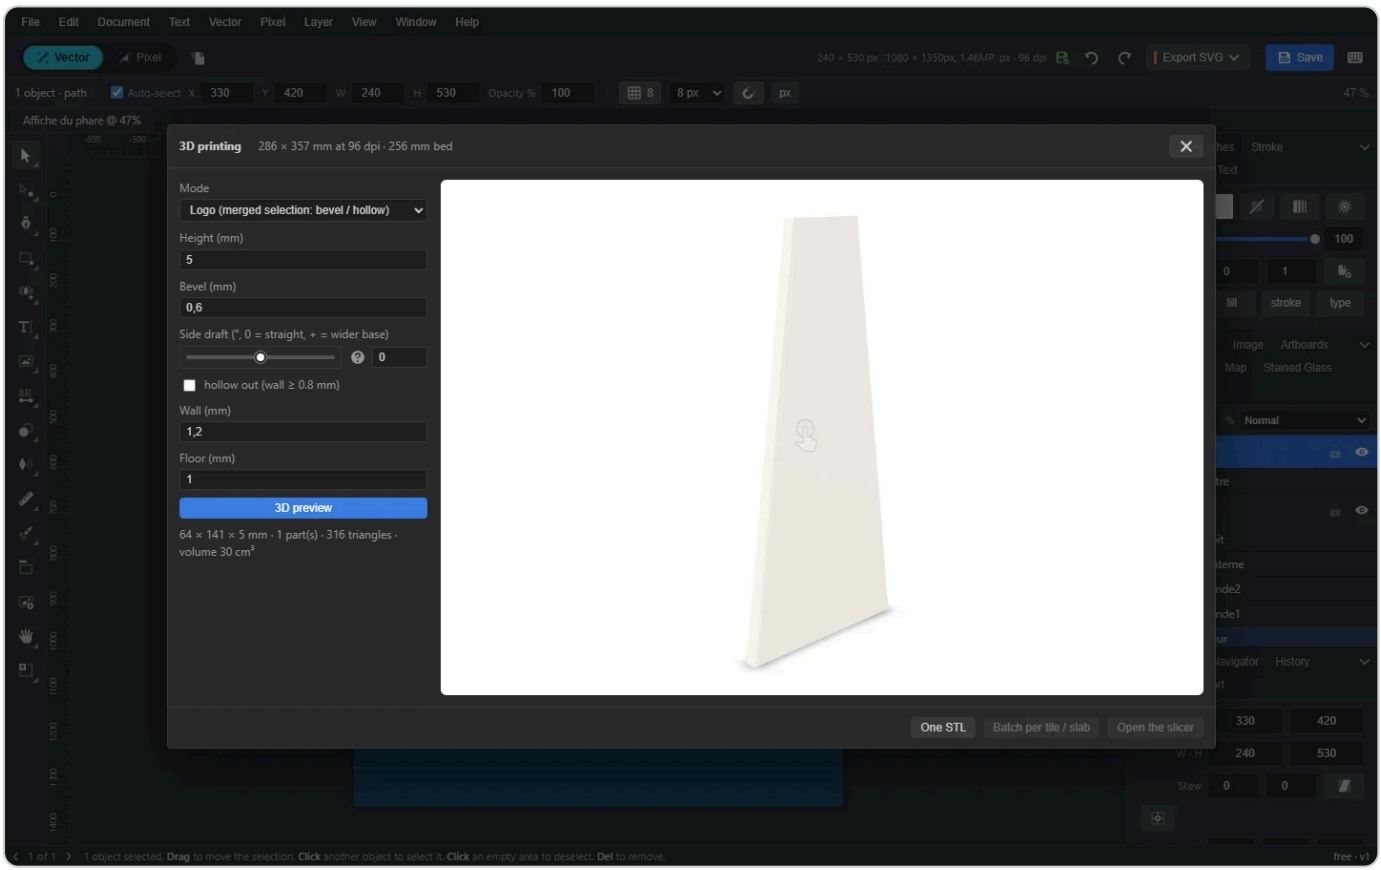

1. Layers, tiles, logo, relief or pixel art 2. The 3D preview, in part colours

File › 3D printing… turns the drawing into parts: 3D preview, then STL, 3MF or a batch per part.

1. File › **3D Print…** (or Text › 3D Logo…). Choose the mode: layers (one height per layer), board tiles with base, logo (bevel, hollowing with a minimum wall), map relief, or pixel art (one height per colour).

2. Set heights and options, then **3D preview**: the model turns in the viewer, each part in its colour, with dimensions, triangles and volume. A model wider than the bed is flagged.
3. **One STL** for a single part, or **Batch per tile / slab**: one STL per part, a coloured plateau.3mf and nomenclature.csv. **Open the slicer** hands over to your printing software.

In relief mode, numbered slabs get sockets under the base and separate keys (0.2 mm clearance) to fit together. Side draft tilts a logo's edges for demoulding.

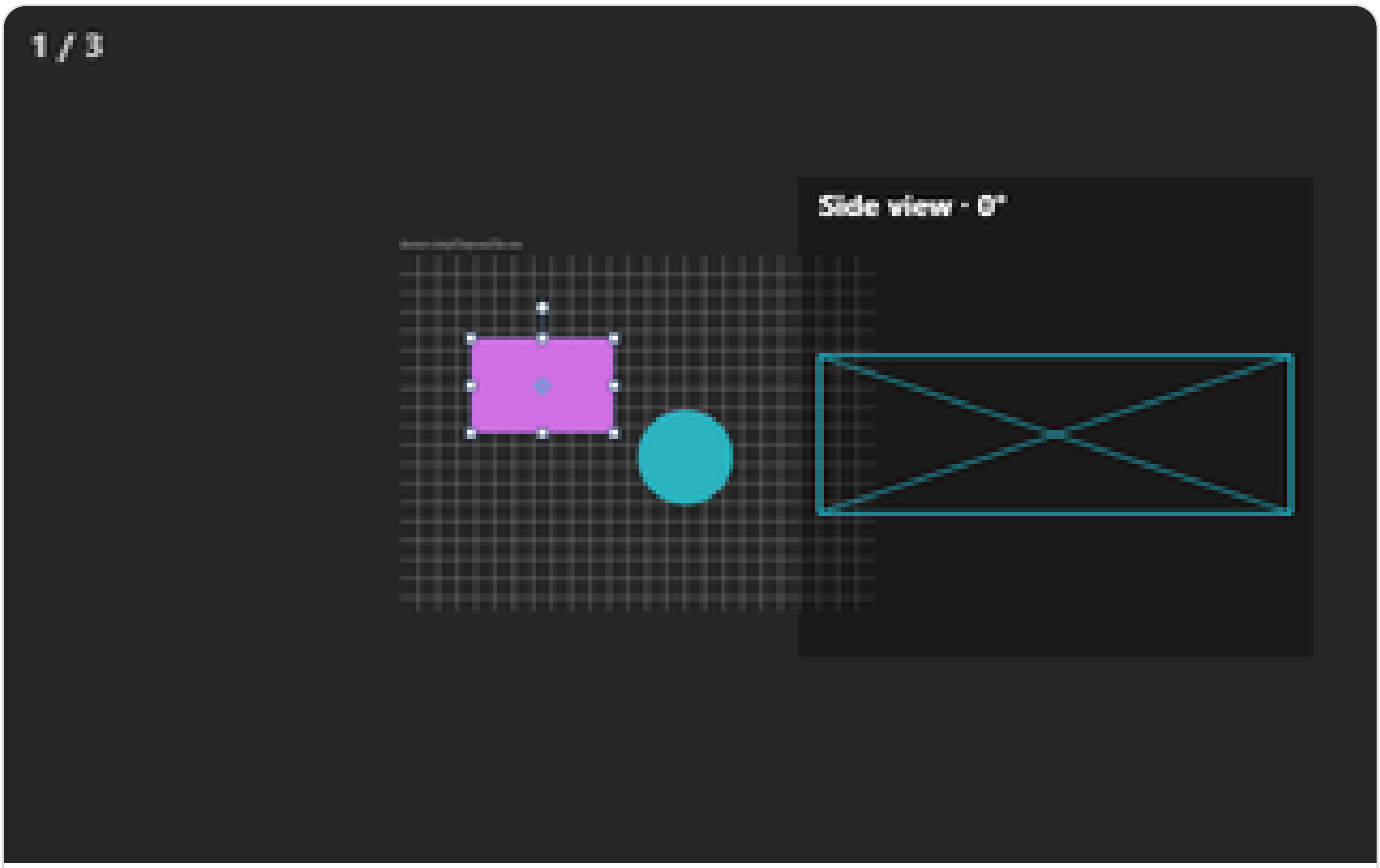

Side draft

The sides of your piece lean: the more you turn, the wider the bottom and the smaller the top.

Good to know

A text must be converted to curves to be printed: the dialog says so and skips it otherwise. To go further with a 3D model, see the Workbench and 3D printing.

GAMES AND 3D

Make game assets

You will be able to find your way around the Game Assets hub, turn an image into a 3D model and open the right 2D tool: Sprite Lab, Tile Lab, Vectorlab's Pixel persona or Card Forge. The 2D labs are free and run on your PC; 3D generation is paid, price shown first.

The Game Assets hub

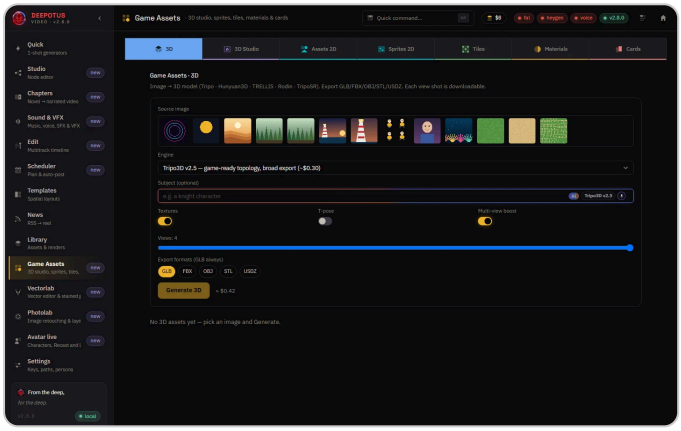

1. Image → 3D model: price shown first 2. The 2D Assets tab

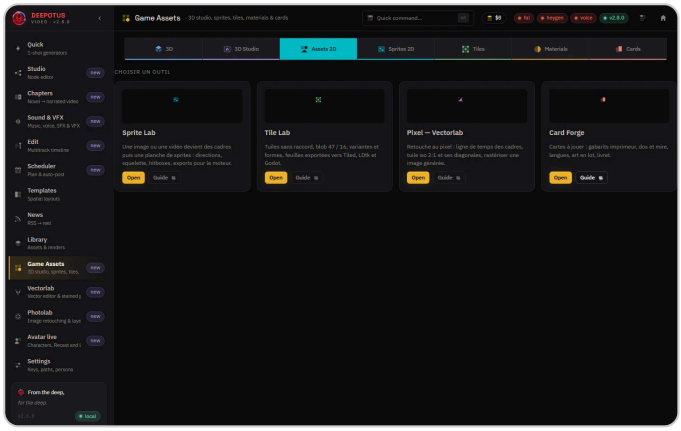

3. Four labs, each with Open and Guide

One tab per tool: 3D, 3D Studio, 2D Assets, Sprites, Tiles, Materials, Cards. 2D Assets gathers the four 2D labs.

Open `Game Assets` in the rail. One tab per tool: **3D** (image to 3D model), **3D Studio** (see Image to 3D), **Assets 2D**, **Sprites 2D** and **Tiles** (the labs opened inside the hub), **Materials** (Material Forge) and **Cards** (Card Forge).

An image becomes a 3D model

1. In the 3D tab, pick a source image from the Library: ideally a character or object cleanly cut out on a neutral background.
2. Pick the engine, from cheapest to most refined: TripoSR, Hunyuan3D, Tripo (default, good value), TRELLIS and Rodin. Only the fal.ai key is needed.
3. Options: textures, T-pose for a character to animate, multi-view (four views generated first for a more faithful model), and formats: GLB always, plus FBX, OBJ, STL or USDZ depending on the engine.
4. The total price is shown next to `Generate 3D`, before the run. The pipeline then shows each step live; the model turns in the viewer, every format and every view can be downloaded.

Models also land in the 3D tab of the Library; the Workbench then prepares them for printing.

Open a 2D tool

The `Assets 2D` tab gathers four cards, each with `Open` and `Guide`:

Sprite Lab

An image, a video or a sheet becomes an animated sprite sheet, with directions, skeleton, hitboxes and engine exports. See Animate sprites.

Tile Lab

Seamless tiles, blob 47 or 16 tilesets, variants, iso and hex shapes, painted maps, exports to Tiled, LDtk and Godot. See Paint tile scenery.

Pixel — Vectorlab

Opens Vectorlab straight in the Pixel persona: frame timeline, 2:1 iso tile, rasterize an image. See Pixel art.

Card Forge

Playing cards: printer templates, back and registration, languages, batch art, booklet. See Card Forge.

Shared palettes

Sprite Lab, Tile Lab and Vectorlab's Pixel persona share the same list of palettes: adaptive (N colours), PICO-8, Game Boy, NES, Sweetie 16, 1-bit, Grayscale 16 (sixteen greys) and Handheld 4 (four tones, handheld console style). A palette gives the same result in all three tools. Dithering (ordered or Floyd-Steinberg) softens gradients.

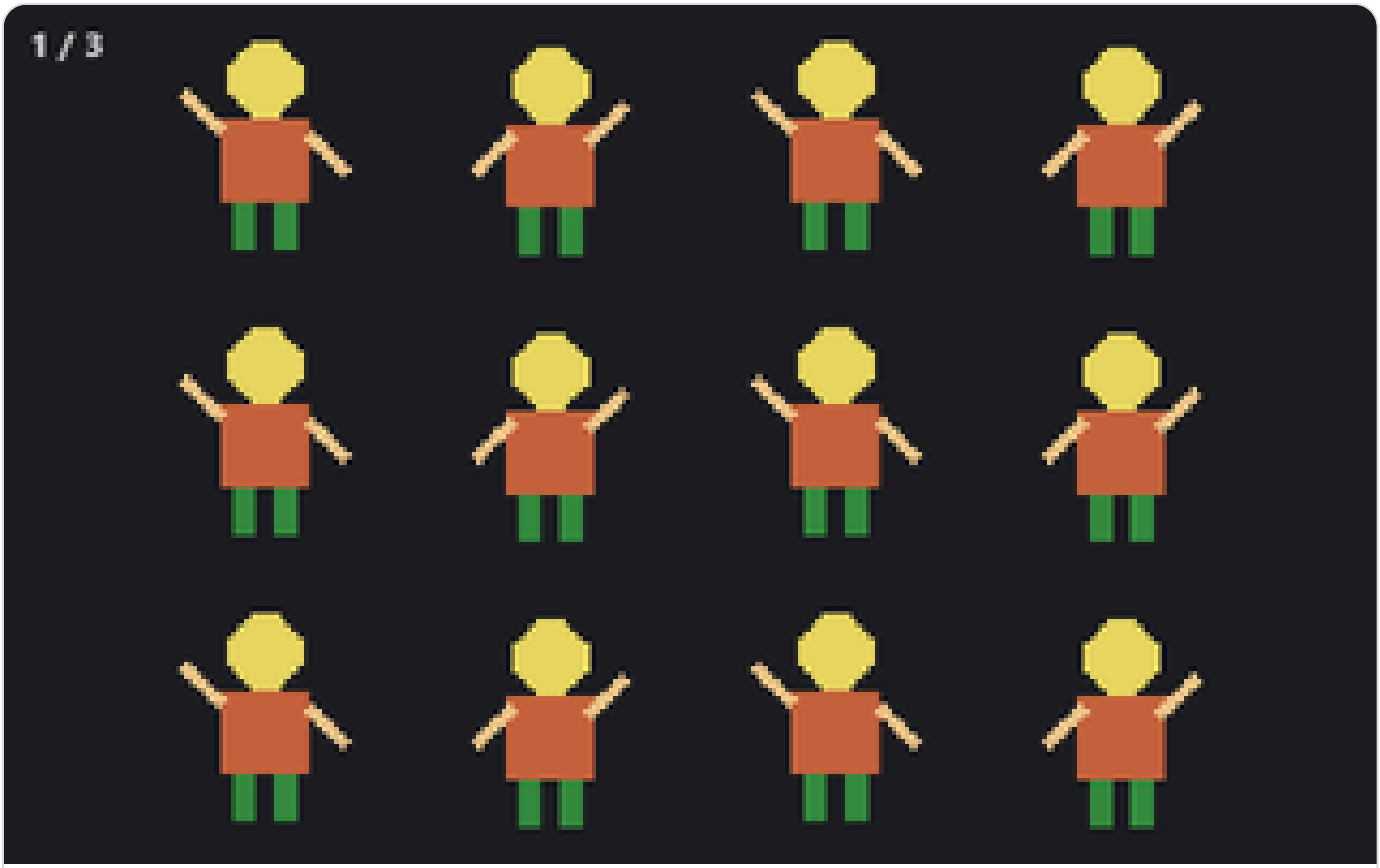

Pixel art (9b)
Your drawing becomes big squares and only a few colours, like in an old video game.

Good to know

Only 3D generation costs money: the price shown counts each extra view and format, and goes through your spending caps. The 2D labs are free.

3D rendering runs in your browser: keep the tab in front for smooth rotation. Files live in the Deepotus data folder and follow a transfer.

GAMES AND 3D

Animate sprites (Sprite Lab)

You will be able to build an animated sprite sheet in Sprite Lab: start from a shipped effect, a sheet, a video or a bible character, slice it, align it, pixelate it, then export it for your game engine. Almost everything is free and local; only animating an image and prompt generation are paid, price shown first.

The Sprite Lab screen

Open `Sprite Lab` from the Assets 2D tab of the Game Assets hub. Three columns: the **source** on the left (Get started, Image, Render, Video, Sheet, Bible, Prompt), the extracted **frames** and settings in the middle, the animated **preview** and exports on the right.

A sprite in one click

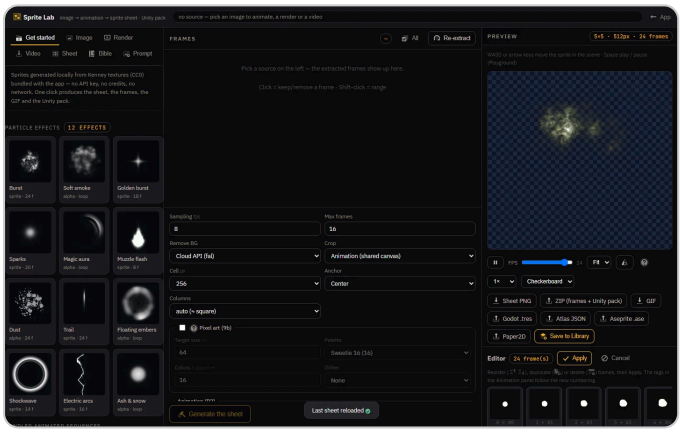

1. Pick an effect

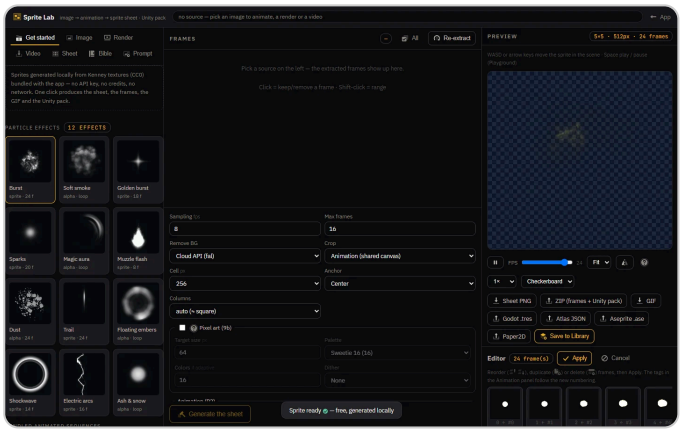

2. The preview plays the animation 3. Exports for Unity, Godot, Phaser, Aseprite, Unreal

The Start tab ships twelve particle effects and sequences: one click produces the sheet, the GIF and the exports, no key, no credit.

1. `Get started` tab: twelve particle effects (burst, smoke, sparks, aura…) and animated sequences, made from free textures shipped with the app.
2. One click produces the sheet, the frames, the GIF and the Unity pack, no key, no credit, no network.
3. The preview plays the animation: `Space` play / pause, speed 0.5×, 1× or 2×, plain backgrounds or scenes (Sky, Outdoor, Dungeon, Cave), and `W` `A` `S` `D` or the arrows walk the sprite. The mirror button flips it.

Mirror

Your character turns around and looks the other way, like in a mirror.

Start from an existing sheet

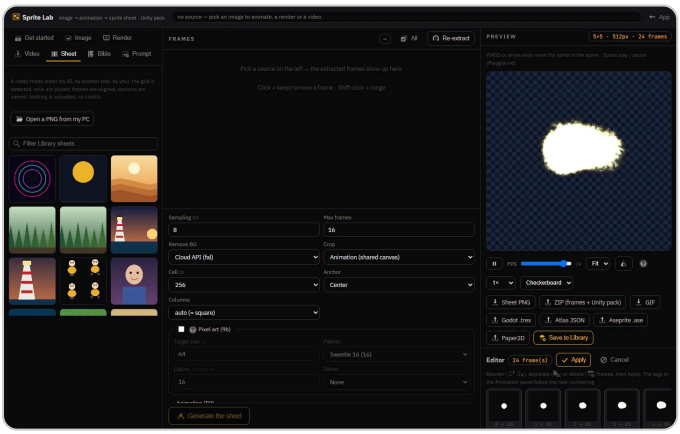

1. A sheet from the Library

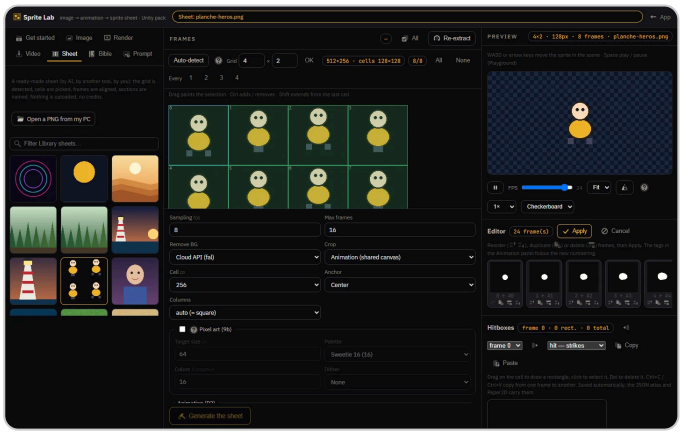

2. 4 × 2 cells detected automatically 3. The hero hops in the preview

Sheet tab: the grid is detected, cells are picked, frames are aligned; all local.

1. **Sheet** tab: pick a sheet from the Library, or **Open a PNG from my PC**. Nothing is sent, nothing is spent.

2. **Auto-detect** counts columns, rows and occupied cells from transparency; otherwise, set the columns × rows grid yourself.
3. Drag to paint the cell selection (`Ctrl` adds or removes, `Shift` extends), or **All**, **None**, one cell in 2, 3 or 4.
4. Align: **Feet** puts each character on the same floor, **Left / right** re-centres, **Both axes** does both; the arrows nudge the current frame by one pixel, **Reset** undoes.
5. Group ranges under a name (walk, jump…) as loop, ping-pong or reverse, then export the **Aligned sheet**, its JSON, or **Copy JSON**.

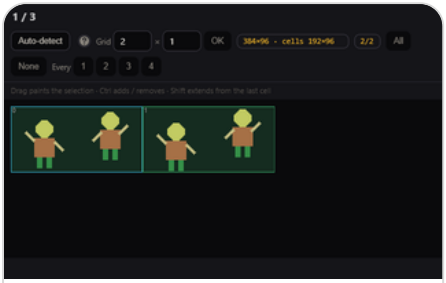

Auto-detect

The tool counts the pictures on your sheet by itself and draws a box around each one.

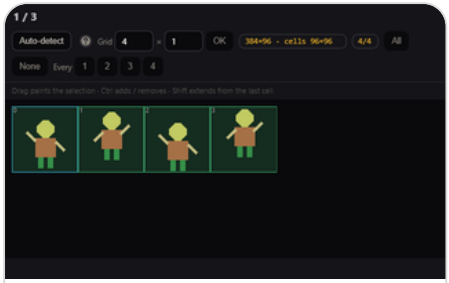

Feet

All your characters put their feet at the same height, as if standing on the same ground.

The other sources

Image, Render, Video

A Library image is animated on a plain background (paid); a render or a video from your PC is sliced into frames. Set sampling, the number of frames and background removal (cloud, local or chroma key).

Bible

4 directions (sheet) cuts a character's south, west, east, north views; **8 directions (3D model)** renders its model from eight angles. Free.

Prompt

Describe a sprite, add the persona's words in one click, then **Generate the images**: the cost is shown first, images are cut out locally and become the source.

Any source also accepts 1 to 64 Library images, taken in order as frames.

Set up the sheet

In the middle: cell (128, 256, 512 or native, no upscaling), anchor centred or at the feet, columns. **Pixel art (9b)** reduces to the target size with a shared palette and dithering. Animation: duration per frame and tags (idle, run, jump…) written into every export. Post-processing: outline, drop shadow, orphan-pixel cleanup, edge smoothing. Then **Generate the sheet**.

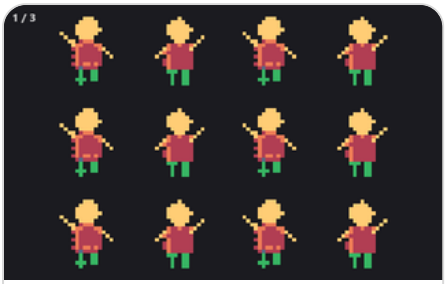

Outline
A line surrounds your character so it shows up on any background.

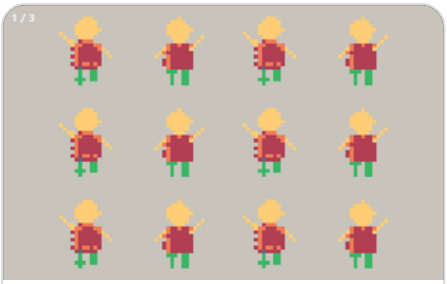

Shadow X / Y
A shadow appears behind your character, slightly offset, as if a lamp were lighting it.

Edit, hitboxes, skeleton

Editor
Reorder, duplicate or delete frames of a finished sheet, then **Apply**; tags follow. Nothing is redone or paid again.

Hitboxes
Draw per-frame hit (strike) or hurt (can be hit) rectangles, sixteen at most; adjust them, copy them from one frame to another (`Ctrl`+`C` / `Ctrl`+`V`). They go into the JSON atlas and Paper2D.

Spine skeleton
Draw bones then parts, name them, and **Write the rig** produces a Spine 3.8 skeleton with its parts.

Export

Under the preview: Sheet PNG, ZIP (frames and Unity pack, with all engine exports), GIF, Godot 4 resource (.tres), JSON atlas for Phaser and PixiJS, Aseprite file (.ase) and Unreal Paper2D sheet. Save to Library files the sheet in the Library, from where it can be sent to Vectorlab or Tile Lab.

Good to know

Paid: animating an image (Image tab), cloud background removal and prompt generation, always within your caps, price shown first. Get started, Sheet, Bible, Render, Video with local removal and every export are free.

GAMES AND 3D

Paint tiled levels (Tile Lab)

You will be able to build seamless tiles, auto-tiling tilesets, iso or hexagonal tiles in Tile Lab, and paint a small map ready for Tiled, LDtk or Godot. Everything is free and computed on your PC.

The Tile Lab screen

Open **Tile Lab** from the Assets 2D tab of the Game Assets hub. A bar of five modes: **Seamless**, **Tile sheet**, **Tileset**, **Shapes** and **Painter**. On the left, the source (Library images); on the right, the result. An image sent from the Library or Photolab (“Tile Lab — tile source”) arrives directly as the source.

A seamless tile

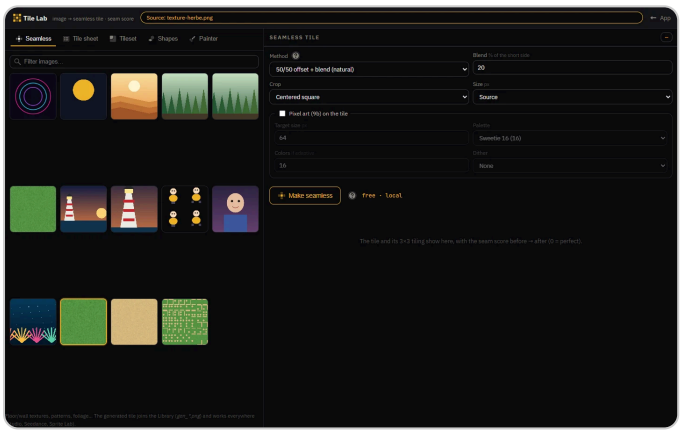

1. Make seamless: free, local

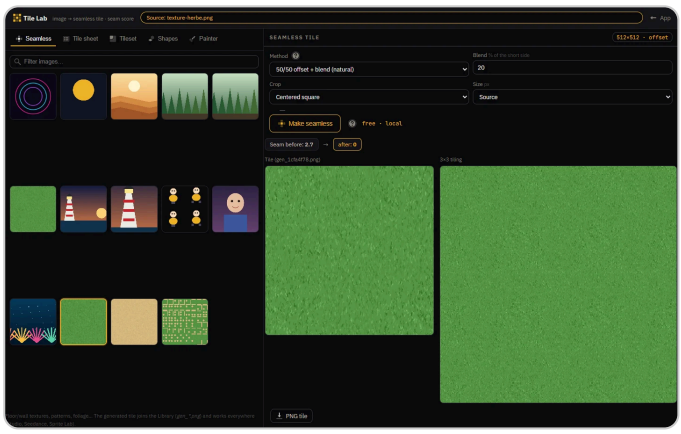

2. The tile and its 3×3 tiling

Seamless mode: an image becomes a tile that repeats without seams, with its 3×3 tiling and seam score.

1. Seamless mode: pick an image (floor, wall, foliage texture…).
2. Method: 50/50 offset with blend (natural) or 2×2 mirror (perfect, kaleidoscope effect); crop and size; optional pixel art with a shared palette.
3. **Make seamless** shows the tile, its 3×3 tiling and the seam score before and after (0 = perfect). The tile joins the Library.

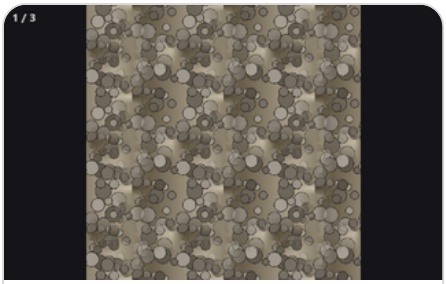

Method

Mirror folds your texture like a sheet folded in four: the edges always meet perfectly.

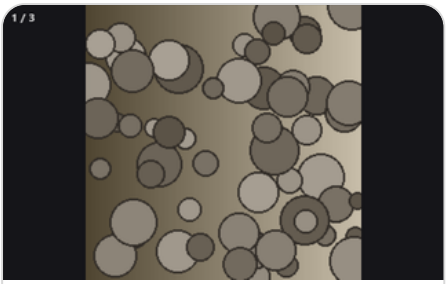

Make seamless

Your texture joins up: laid side by side many times, you can no longer see where each piece ends.

Build a tileset

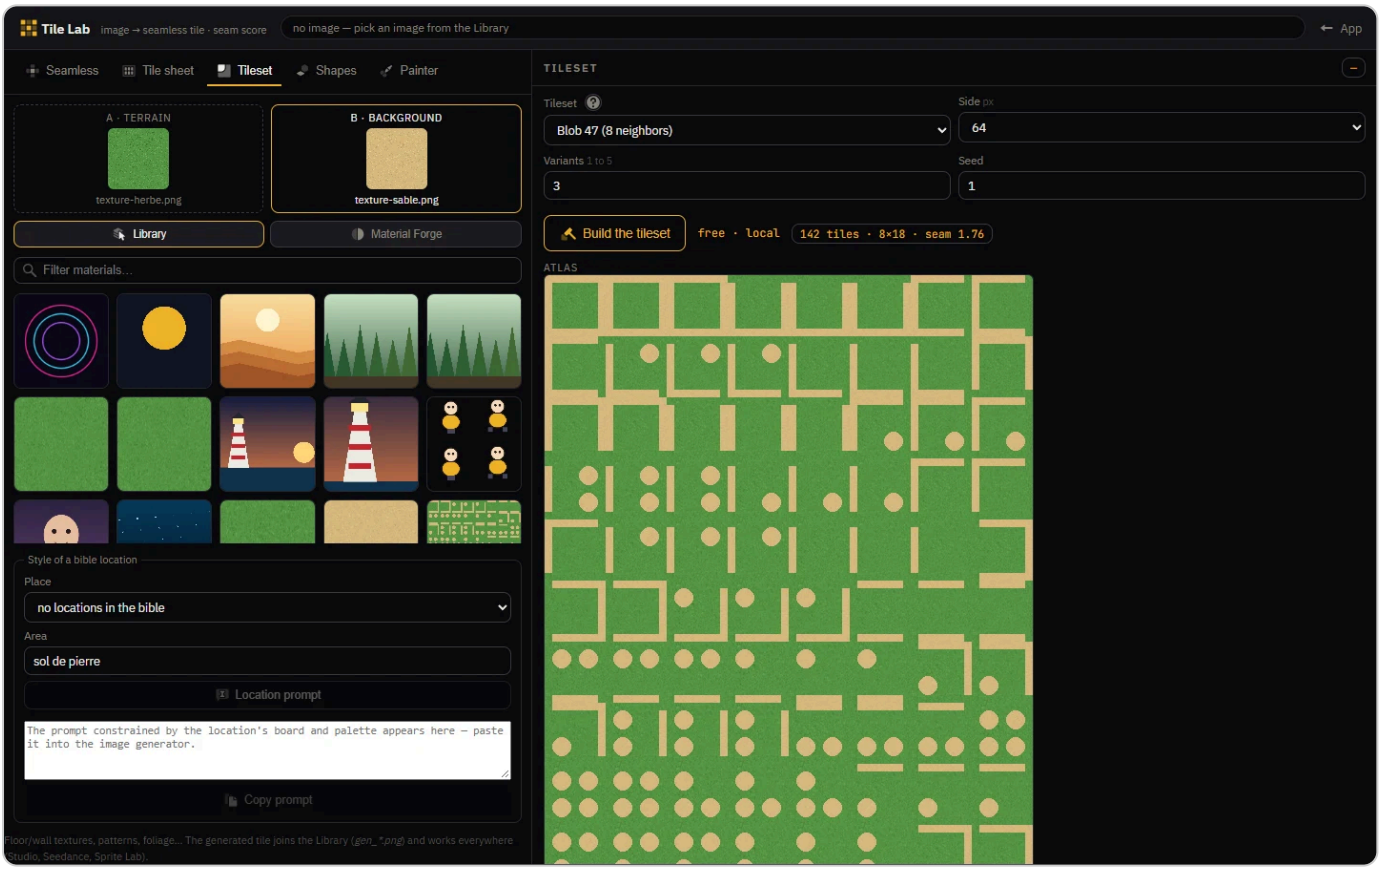

1. A: terrain, B: background 2. 142 tiles, seams measured

Tileset mode: a terrain material and a background one become a blob 47 or 16 set, auto-tiling and exportable.

1. Tileset mode: click an image for material A (terrain) then a second for B (background). The **Material Forge** toggle takes the base colour of a Material Forge material.
2. Choose the set: Blob 47 (eight neighbours, rounded corners) or Blob 16 (four edges, faster), the side in pixels, 1 to 5 variants per tile and the seed.

3. **Build the tileset** builds the atlas and measures how neighbours join; the atlas is copied to the Library.
4. **New preview** draws a random map where each cell gets the tile of its neighbours; **Measure** judges seams, visible repetition and lighting (ok / warning).
5. Export atlas.png then your engine's format: Tiled (.tsx), LDtk (.ldtk) or Godot (.tres), to be placed next to the atlas.

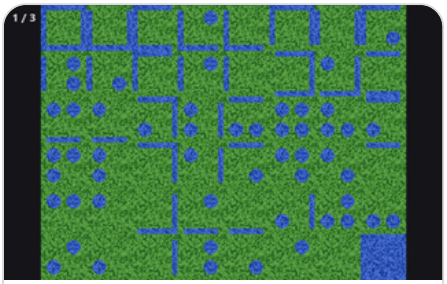

Tileset
You choose how many pieces to make: many for nicely rounded corners, or sixteen to go fast.

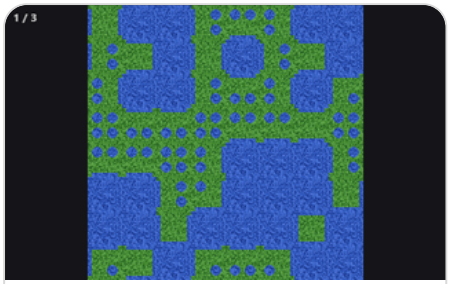

New preview
A new map draws itself at random with your tiles, to see if they go well together.

The “Style of a bible location” block composes a tile prompt bound by a location's sheet and palette: **Location prompt**, then **Copy prompt** to paste it into the image generator. Nothing is generated here.

Paint a map

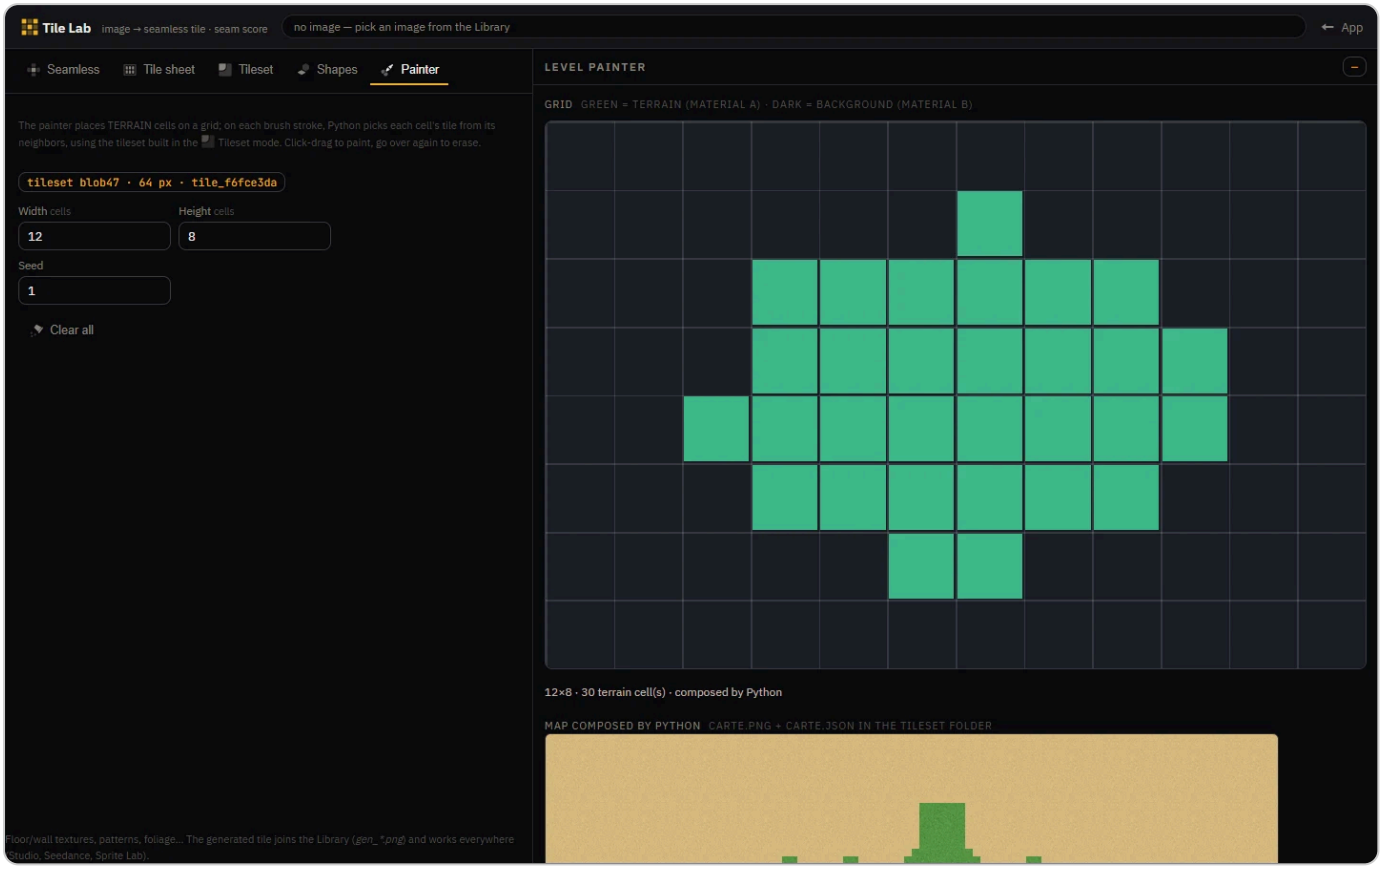

1. Click and drag to paint 2. The composed map, ready for the game

Painter mode: you paint terrain cell by cell, and the right tile of the set is placed from its neighbours.

Painter mode, once a tileset is built: set width, height and seed, then click and drag on the grid to paint terrain, go over again to erase. On each stroke, the right tile is placed from its neighbours, and the composed map shows below with carte.png and carte.json. **Clear all** empties the grid.

Shapes and tile sheets

Shapes

A material becomes a 2:1 iso diamond or flat-top hexagon tile that tiles without seams. **Build the tile** shows the tiling; export to Tiled and Godot.

Tile sheet

Open an existing sheet: **Detect** finds each tile by its transparency, the placement grid (square or 2:1 iso) tests seams, double-click = ×9 tiling.

Sheet exports

Aligned sheet and JSON, tileset and JSON, or **Separate tiles (zip)**; the aligned sheet can also join the Library.

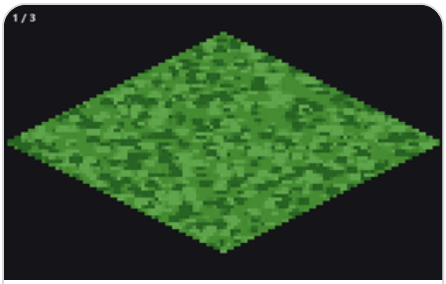

Shape

Your tile changes shape: a diamond for a game seen at an angle, or a six-sided hexagon.

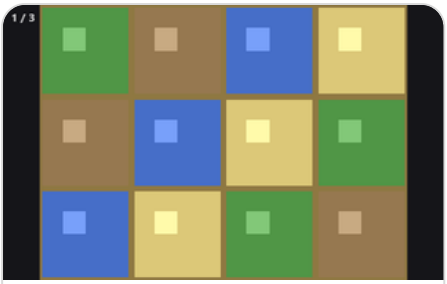

Placement grid

Your cells become diamonds, staggered like the paving of a game seen at an angle.

GAMES AND 3D

Create playing cards (Card Forge)

The **Card Forge** sub-tab (inside **Game Assets**, next to 3D / Sprites / Tiles / Materials) is a full playing-card editor: face, frame, typography, CSV data, 3D volume, PBR materials, print export and 3D export — eight modules in one document. No key is required; only **AI face generation** uses your existing fal.ai key.

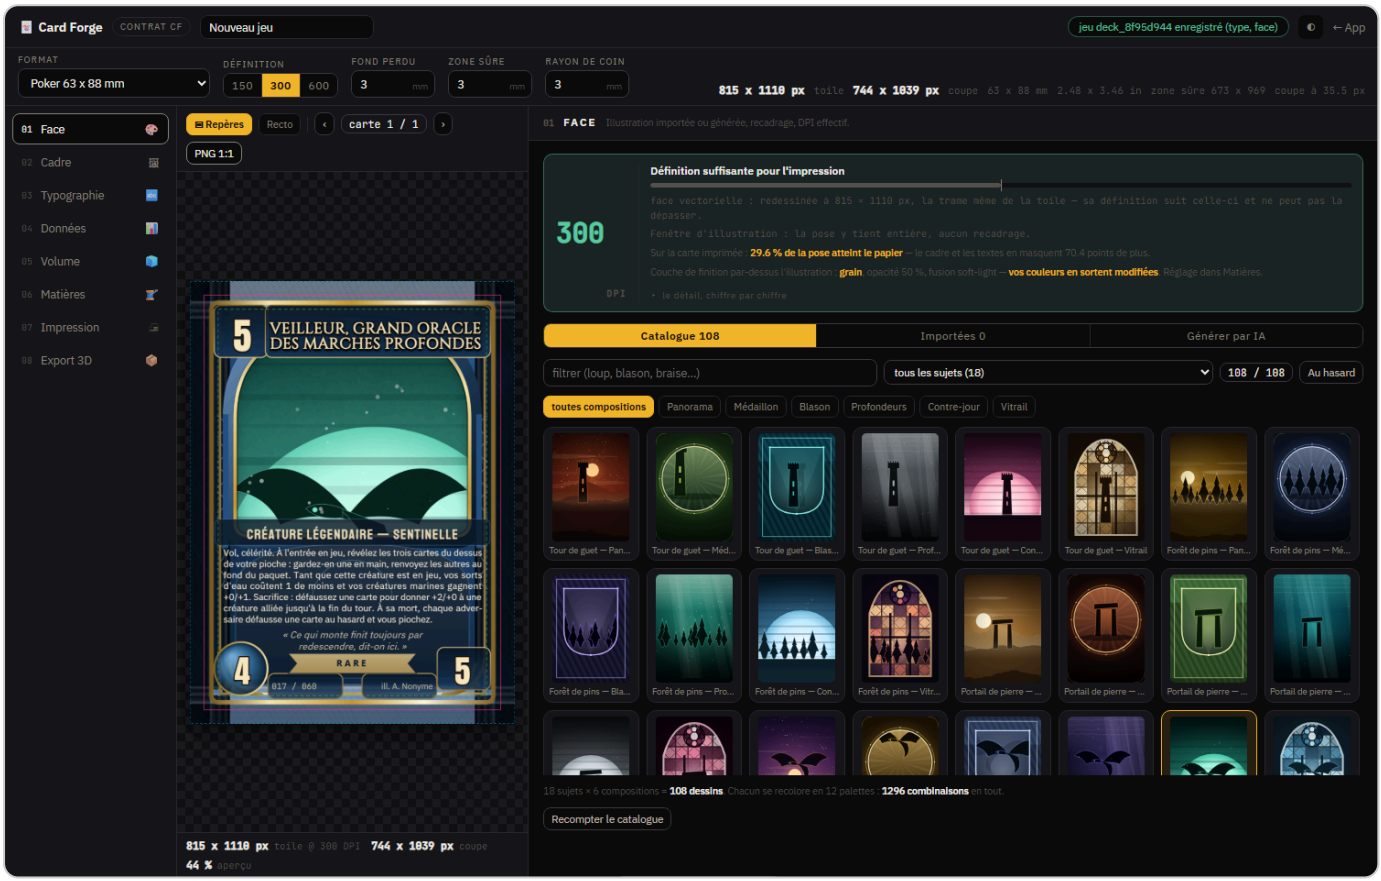

1. Pick the format and resolution: they apply to the whole deck 2. The catalog offers 108 ready-made illustrations 3. The gauge tells you if the picture is sharp enough to print

The Face module: format and resolution at the top, a 108-illustration catalog, the effective-resolution gauge.

1. At the top of the lab, pick the **format** (poker, bridge, tarot, domino…), the **resolution** (150 / 300 / 600 DPI), the **bleed** and the **safe zone** — these settings apply to all eight modules and to the final export.

2. Browse the modules in the left rail: **01 Face** (imported illustration, AI-generated, or picked from the local catalog), **02 Frame** (borders, rarity, card back), **03 Typography** (title, rules box, flavor text, 23 already-bundled fonts), **04 Data** (CSV import → columns to fields → full deck), **05 Volume** (3D preview with real thickness and edge), **06 Materials** (8 PBR texture channels computed locally).

3. **07 Print** runs a **preflight check** on both the sheet and the produced file: bleed, cut marks, the PDF's `TrimBox` / `BleedBox`, the PNG's `pHYs` density (300 DPI = 11811 px/m) — every line is verified against the delivered file's bytes, never against an on-screen setting.

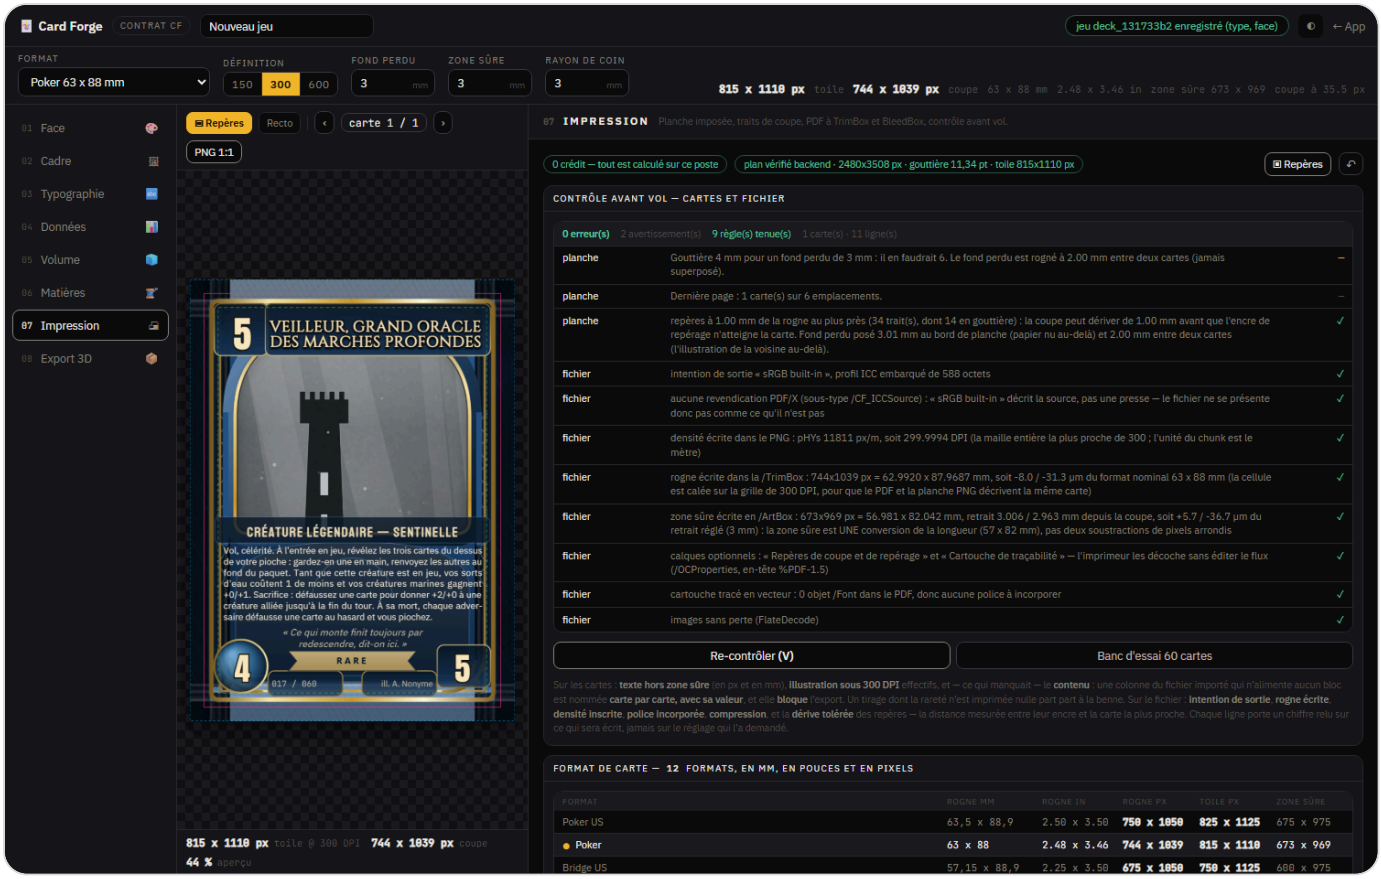

The preflight check: 0 errors, TrimBox and BleedBox present in the PDF, PNG density verified byte by byte.

4. **08 3D Export** exports both a `.glb` **and** a standalone `.gltf` with the full PBR map set (base color, normal, roughness, metallic, occlusion, height, emissive, ORM), ready for Unity, Unreal, Godot or Blender.

> ⚠ The **emissive** channel is now adjustable in the 3D export (an **Emission 0..1** slider — luminescent ink on matte paper, or extinguished gilding) and wired into the GLB, glTF and MTL; it stays capped at 1.0 and its mask is derived from the card's luminance (adjustable threshold). A pose hidden by the frame is fixed **in one click** from the Face module — a panel button plus a contextual offer next to the masking percentage.

GAMES AND 3D

From image to 3D object

Card Forge's **09 Forge 3D** module turns your card into a printable 3D piece or an "NFT-grade" artifact: the card exports **as layers** (base material, illustration, veil, frame, typography, ornaments + the composite), each layer is processed in a **graph**, and assembly produces a clean GLB — plus an STL whenever everything is watertight. The whole free circuit runs locally, no key needed.

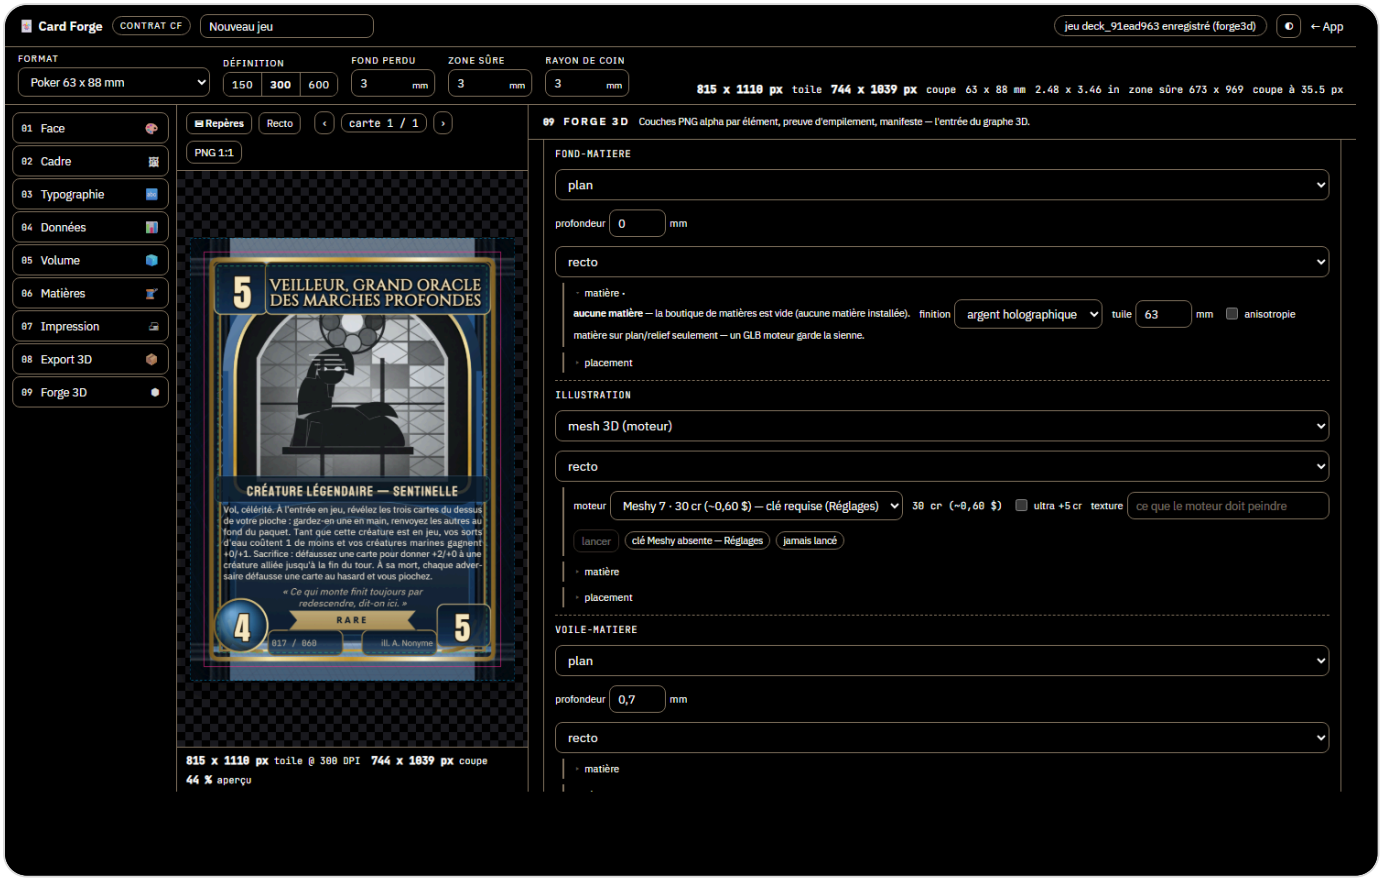

The Forge 3D graph: a material chain (silver holographic finish) on the background, an image→3D engine on the illustration — 30 cr (~$0.60) shown before any launch.

1. **Export the layers**: the screen proves the stacking **before** uploading (0 px of difference between the composite and the layer pile), the server counter-proves it in PIL and seals a manifest. An empty layer is still written — absence is measured, never guessed.

2. **Edit the graph**: each layer gets a **plane** treatment (flat card) or a **relief** (a grid displaced by luminance, a solid closed by construction), a **material** from the Material Forge shop (tiled at physical scale, in millimeters), a **silver or gilded holographic** finish (iridescence, clearcoat, anisotropy — verified in the GLB's bytes) and a **transform** (position, rotation, scale, z stacking).

3 **Image→3D engines** (optional, paid): the **mesh 3D** treatment sends the layer to one of 7 engines — tripo, hunyuan, trellis, rodin, triposr (fal, priced in $) or **Meshy 6/7 through the direct API** (credits; ultra +5 cr on meshy-7). The price shows on the node **and** in the graph footer before any launch; the status chip follows the job (queued → running → served · real credits → or failure with the literal message); relaunching a node resets its folder.

4 **Build**: assembly merges the engines' GLBs at their layer's place (full reindexing, the engine file's identity is dropped), writes the final GLB and the **STL** — accepted only when every element is a closed solid, with a motivated refusal otherwise. The docket lists the elements, their engines and the credits consumed.

⚠ The image→3D engines are the only paid part of Card Forge: fal bills in $ on your existing key, Meshy in credits on yours (`MESHY_API_KEY`, the Settings field shared with the 3D Studio). To try everything without spending a credit, set `MESHY_MOCK=1` in the data `.env`: the local simulator runs the full circuit, fake credits included.

GAMES AND 3D

Prepare a part before slicing (Workbench)

Chapter 20 showed you how to get a file *out*. This one shows what to do *before*. The **Workbench** screen (Game Assets → 3D Studio → *07 · Établi 3D* →) is the workshop where you repair, lay down, inspect, hollow, arrange and measure a part, so the slicer only has to slice. Nothing below is irreversible: **every gesture that changes the part writes one more VERSION**, and all of them stay on disk — the *model's life* column on the left lists them. (The Workbench screen itself is in French, like the rest of the application; its labels are quoted below as they appear.)

1. **Set a target size.** Right rail, *taille cible* field: type the largest wanted dimension, in millimeters (80 for a figurine, 250 for a board). A 3D file carries no unit — while this field is empty the Workbench shows "glTF units" and refuses whatever needs a physical dimension (hollowing, arranging, exporting for print). That is deliberate: a refusal beats an invented millimeter.

2. **Pick the printer.** Same rail, *imprimante* menu: the **Elegoo Centauri Carbon 2** profile is built into the application (256 × 256 mm bed, 256 mm height, one excluded area). If OrcaSlicer or ElegooSlicer is installed, its profiles show up too — read, never modified. On the plate, the green outline IS the chosen machine's bed; the red rectangle is the area where nothing may sit.

3. **Repair in one click.** *Fiche* tab → *Réparer en un clic*. Checked by default: coincident vertices welded, duplicate faces dropped, flat triangles dropped, normals unified. *Trous bouchés* (cap holes) is unchecked by default: capping adds material, tick it knowingly. The detail appears in the bottom bar — "8 retourné(s) — fermé" (8 flipped — closed). A mesh left open prints badly: better to know here than at the machine.

4. **Lay the part flat.** *Poser sur une face* (key F), then click the face that must touch the bed: the part turns and lands on the floor. The gesture joins the queue — the bottom bar says "1 modification en attente" — and *écrire la version* saves it. This gesture decides almost everything else: a well-placed part needs few supports. If unsure, *Orienter* proposes three ranked poses in the *Fiche* tab — share of the surface on the bed, share overhanging, height — and one click lays the part on the one you pick. A computed orientation does not always find the best pose: **always review the proposal**, no score knows which face you want to look good.

5. **Look at the overhangs, then the slices.** *Surplombs* button: faces leaning less than 45° from horizontal turn orange — there, a layer has almost nothing under it and the slicer will add a support. Re-lay the part until the orange recedes. *Tranches* button: twenty cross-sections of the part, in blue, to see its real shape at each height. Both are **indicative previews**: nothing is written, and the slicer does the real slicing.

6 **Hollow it, when it helps.** *Fiche* tab → *Creuser*, 2 mm wall to begin with: the solid inside is replaced by a shell. It is first a **resin** (SLA) gesture, where a solid part costs a lot of material; with filament (FDM), the slicer already fills the inside with partial infill. It needs a *closed* mesh (step 3), and the Workbench states the thickness that holds when the requested wall is too thick for the part. For resin, then drill a drain hole: same tab, *Percer (drainage)*, a 3 to 4 mm diameter, then click the face to drill — preferably near the build plate, so the resin can flow out. The drill goes through the one wall under the click and stitches the hole closed; the opposite wall is left untouched. A solid part is refused (hollow it first).

7 **Arrange on the bed.** *Sur la plaque* separates the pieces, then *Ranger sur le plateau* lays them inside the real bed, rotated when that saves room, 2 mm apart. What does not fit moves to a second plate, and what is larger than the machine is reported. Drag a piece with the mouse (snapped to the bed grid, Shift frees it), nudge it with the arrow keys; its position, size and angle show under "Sur la plaque". The plate is a **view**: it is saved, but it does not change the model.

8 **Measure, then export.** *Mesurer* → two clicks give a distance, its x/y/z components and the angle between the two faces. Finally, *Export* tab → → **Impression 3D** writes the STL and the 3MF of the **displayed version**, at the target size in millimeters (chapter 20), and *Ouvrir dans le slicer* takes you there.

> **The order that works**: repair → lay down → inspect → hollow → export. Repair first, because every other gesture starts from sound geometry — and hollowing refuses an open mesh.

Glossary

The words your slicer will use, explained once.

Term	What it is
Base (bed contact)	The face that touches the bed. Wide and flat, the print holds; narrow, it comes loose.
Overhang	A wall leaning too far to rest on the layer below. Prusa puts the clean limit without support between **45 and 60°** from horizontal.
Support	A disposable structure holding overhangs up. The slicer generates it; the Workbench only helps you need less of it.
Brim	A one-layer collar around the part, to keep its corners from lifting.
Raft	A platform printed UNDER the whole part. Safer than a brim, costlier, and it leaves a rough face.
Skirt	An outline drawn next to the part, not touching it, to prime the flow before starting.
Infill	The pattern inside. *Gyroid* is about equally strong in every direction; *Lightning* only supports top surfaces. 15 % is usually enough.
Seam	Each wall loop starts and ends somewhere, leaving a visible vertical line. The slicer places it (*Nearest*, *Aligned*, *Rear*) or scatters it (*Random*).
Retraction	Filament pulled back during travel moves, to avoid strings (*stringing*).
Layer	A horizontal slice. 0.2 mm is the usual compromise; 0.12 mm for detail, 0.28 mm for speed.
Perimeter (wall)	The number of wall loops. Two perimeters at 0.45 mm make a 0.9 mm wall: the minimum to aim for when modeling.
Bridging	Printing a layer over empty space, between two supports. It works over a few centimeters, with good cooling.
Warping	Corners lifting as they cool. Remedies: clean bed, brim, closed enclosure. Common with ABS/ASA, rare with PLA.
Watertightness	A solid is watertight when its surface is closed, with no hole or free edge. That is what the Workbench bar calls "fermé" (closed).
Manifold	Every edge belongs to exactly two faces. A mesh can be closed and still be hard to slice cleanly if it is not.

Decimation	Reducing the triangle count. A 500,000-triangle mesh prints no better than a 100,000 one, but slows everything down. *Fiche* tab → *Décimer*.
Hollowing	Replacing the solid inside with a shell of given thickness. Mostly useful in resin; 2 mm is a good starting point.
Drain hole	The hole(s) letting uncured resin (and air) out of a hollowed part. In the Workbench: *Fiche* tab → *Percer* (drainage).

Where to learn more

Every link below was opened and checked on 05/10/2026 (5 October 2026). They are grouped by what they teach you.

Topic	Resource	Checked
Supports and overhangs	Prusa — Support material (https://help.prusa3d.com/article/support-material_1698) (the threshold is measured from horizontal, 90 = vertical)	05/10/2026
Modeling for print	Prusa — Modeling with 3D printing in mind (https://help.prusa3d.com/article/modeling-with-3d-printing-in-mind_164135) (45–60° overhangs, walls, tolerances)	05/10/2026
Bridges	Prusa — Poor bridging (https://help.prusa3d.com/article/poor-bridging_1802)	05/10/2026
Layers and perimeters	Prusa — Layers and perimeters (https://help.prusa3d.com/article/layers-and-perimeters_1748)	05/10/2026
Infill	Prusa — Infill patterns (https://help.prusa3d.com/article/infill-patterns_177130)	05/10/2026
Seam	Prusa — Seam position (https://help.prusa3d.com/article/seam-position_151069)	05/10/2026
Skirt and brim	Prusa — Skirt and brim (https://help.prusa3d.com/article/skirt-and-brim_133969)	05/10/2026
Warping	Prusa — Warping (https://help.prusa3d.com/article/warping_2011)	05/10/2026
PLA filament	Prusa — PLA (https://help.prusa3d.com/article/pla_2062) (for Prusament PLA: nozzle 215 °C first layer then 210 °C, bed 60 °C)	05/10/2026
PETG filament	Prusa — PETG (https://help.prusa3d.com/article/petg_2059) (nozzle 230 °C then 240 °C, bed 85 °C then 90 °C)	05/10/2026
Choosing a filament	Prusa — Filament material guide (https://help.prusa3d.com/filament-material-guide) (comparison table of materials)	05/10/2026
Print defects	Simplify3D — Print Quality Guide (https://www.simplify3d.com/resources/print-quality-troubleshooting/) (about 26 defects, with photos)	05/10/2026
Slicer — prepare	OrcaSlicer — Object manipulation (https://github.com/SoftFever/OrcaSlicer/wiki/prepare_object_manipulation)	05/10/2026
Slicer — orient	OrcaSlicer — Auto orientation (https://github.com/SoftFever/OrcaSlicer/wiki/prepare_auto_orient) (which warns it may not always find the best pose)	05/10/2026

Slicer — calibrate	OrcaSlicer — Wiki (Calibrations) (https://github.com/SoftFever/OrcaSlicer/wiki/Home) (temperature, flow, retraction, pressure advance…)	05/10/2026
Machine	Elegoo — Centauri Carbon 2 wiki (https://wiki.elegoo.com/centauri-carbon-2-combo) (troubleshooting guides and error codes)	05/10/2026

These links come from slicer makers and manufacturers: they describe THEIR machine and THEIR software. Numeric settings (temperatures, speeds) apply to the hardware they name; on another machine they are starting points, not truths. Two repair tools often recommended elsewhere, **Microsoft 3D Builder** and **Autodesk Meshmixer**, are no longer maintained, so they are not listed here.

GAMES AND 3D

3D-print your creations

Everything the app produces in 3D — and even your vector illustrations — exports to your printer in one gesture. The circuit is designed for the **Elegoo Centauri Carbon 2** (256 × 256 × 256 mm build plate) and its slicer **ElegooSlicer** (bundled with the machine, an OrcaSlicer fork — OrcaSlicer works too). Every export writes a folder under `DeepotusVideoGenData\assets\print3d\` with three files: the **STL** (the mesh), the **3MF** (the one you open — it carries the millimeter unit and the name), and `impression.json` (provenance, scale, watertightness). Everything is local and free: no provider calls.

1. **Install ElegooSlicer once.** The installer associates `.3mf` files with the slicer on Windows: the app's "Open in slicer" button uses that directly. Without an association the app tells you so, and you can set `SLICER_PATH` (the slicer executable path) in the data `.env`.

2. **From the card Forge 3D** (Card Forge, module 09): when the build slip delivers an STL — that is, when every element is a **closed solid**, the check that guarantees it — a **→ 3D printing** button appears next to it. Leave the size empty to keep the card's real millimeters, or type a target (e.g. `80` for a figurine). These exports are marked *guaranteed* watertight: the producer proved it.

3. **From Game Assets · 3D**: every generated model (card characters, props) carries its **→ 3D printing** button. Give the largest dimension in millimeters — 80 for a figurine, 250 for a board — and the app centers the piece and rests it on the bed. The converter reads the original `model.glb`; the "optimized" file (compressed for game engines) is refused with a message that says what to use instead.

4. **From the Vectorlab** (Export menu → 3D printing…): your vector illustrations become **printable reliefs**. Give a global height in millimeters, then override per layer with `name=mm` — for example `2, contours=5` for a stained-glass piece: glass at 2 mm, leads at 5. Strokes without a fill count by their stroke width; text objects are skipped (and counted in the message). Flat dimensions follow the document's unit (an A4 comes out at 210 × 297 mm).

5. **Slice in the slicer.** The `.3mf` opens at the right size, in millimeters, resting on the bed. Supports, fine orientation, infill, filament profiles: that's the slicer's job — the app never generates G-code. Beyond 256 mm the export warns you the slicer will have to cut or shrink, without blocking you.

> ⚠ Meshes coming from the image→3D engines are marked *unknown* watertightness in `impression.json`: ElegooSlicer and OrcaSlicer repair small defects on import and say so. Only the card Forge 3D guarantees its solids — its "closed solid" check states it, not a promise.

REFERENCE

End-to-end recipes, from idea to post

🍳 Recipe A — a meme posted to Telegram in 5 minutes

1. **Library** → "Create image": describe the meme (~3 s).
2. **Scheduler** → *New post*: caption (+ a 🔥 emoji), Telegram channel, time in 5 min, *Auto* mode.
3. Or don't wait: **Send now**. Posted.

🍳 Recipe B — your selfie video + an animation (UGC)

1. Film 10–15 s on your phone, transfer the file to the PC.
2. **Studio** → "UGC + animation" starter → *UGC video* node → *Upload your video* (⏱ Duration master on).
3. Choose the image + prompt of the animated half (Seedance) → *Run*. The output lasts exactly your clip.
4. **Scheduler** → attach the render, polish the caption (emoji + pack), schedule.

🍳 Recipe C — today's news reel

1. **News** → *Refresh* → tick 2-3 articles.
2. *Build script* (persona voice) → *Build illustration*.
3. Compose in **Templates** → *Send to Scheduler* → schedule or publish.

Recipe D — an episode of your novel

1. **Episodes** → title + paste (or upload) the chapter → pick a voice & language → *Generate narration*.
2. **By AI** (or By paragraph) → review the scenes → *Generate all illustrations* → set each scene's motion (Ken Burns / Seedance).
3. *Assemble the episode* → preview → *Send to Scheduler* (YouTube + Instagram drafts). One episode a day, done.

REFERENCE

Read the icons

You will recognise an icon at a glance, find its name, and print the full sheet to keep next to your screen. Free, no key needed.

One icon, one function

The whole app speaks the same language: the **Deepotus Glyph** set, 526 icons drawn with the same stroke. Each one stands for **a single function**: the same drawing means the same thing in the editor, the Photolab or the Card Forge. The old emojis and symbols are gone from buttons, menus and lists.

1. Rest the pointer on a button: its tooltip gives its name and, if it has one, its keyboard shortcut.
2. Look that name up in the icon sheet to read what it does and what not to confuse it with.
3. In the Vectorlab, the Photolab, the Sprite Lab and the Tile Lab, wait a little longer (900 ms) or click the small `?` button: an animated card SHOWS you the option.

Icon families

Icons are sorted into families. Knowing the family already tells you what to expect when you click.

Navigation

A screen, a tab or a panel: you go there, nothing is changed.

Actions

A button that does something: create, save, send, delete.

Editing

Change something already there: align, flip, group, lock.

Tools

The drawing and retouching tools of the labs.

Layers and adjustments

Layers, masks, styles and image adjustments.

Media and playback

Play, cut and edit video and sound.

States

What the app tells you, without you having to click.

Categories

Content families, each with its own colour.

3D

Objects and actions of the 3D workshops.

States worth knowing

A state icon is not clicked: it informs you. These are the ones to read before starting a paid render.

Icon	What it says
	Done: the render, the upload or the save worked.
	In progress or waiting: loading, generating, or queued.
	Failed: the tooltip or the message next to it says why.
	Warning: nothing is broken, but read before going on.
	Spending limit close or reached: beyond it, the app refuses the spend and asks you what to do.
	Generative field: AI can fill it in for you.
	Locked: nothing can move, change or delete the item while it is.
	Pinned (Studio): this node result is reused for free on the next render.

Print the icon sheet

The icon sheet shows them all, sorted by family, with their name, their function and the look-alike icons not to confuse them with.

1. Open the icon sheet.
2. To print only part of it, click a family (for example **Vectorlab tools**) or type a word in the search: only what is shown gets printed.
3. Click **Print**: each family starts on a new A4 page, black on white. You can also download the full sheet as a PDF.

> On the sheet, the red line starting with "≠" names the look-alike icons: Close hides a window, Delete destroys, Remove takes something out of a list without destroying it.

Animated cards in the labs

In the labs, some options have an animated card in three beats made of real screenshots: the starting point, the gesture, the result. It opens when you rest the pointer on the option for 900 ms, or with the **?** button next to it (in the Sprite Lab and the Tile Lab, even when the option is greyed out). Here they all are; cards of labs not yet translated still show the French interface.

Vectorlab

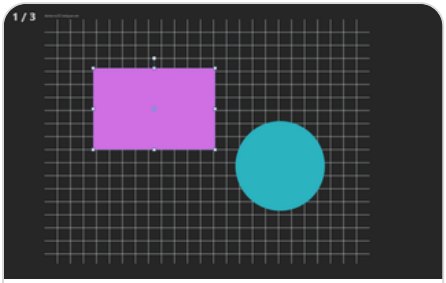

Conical

The color spins around the middle of your shape, like the spokes of a wheel.

`Appearance › Conical`

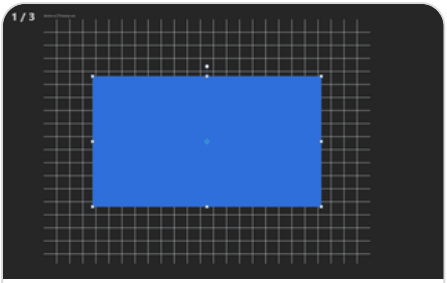

Transparency

Your shape slowly fades from one side to the other, as if it vanished into fog.

`Appearance › Transparency`

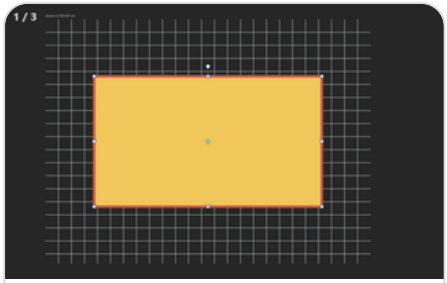

Pattern

Your shape fills up with little repeated stripes, like striped fabric.

`Appearance › Pattern`

Frame

Your long line of text fits into a box and wraps by itself, like in a book.

`Appearance › Frame`

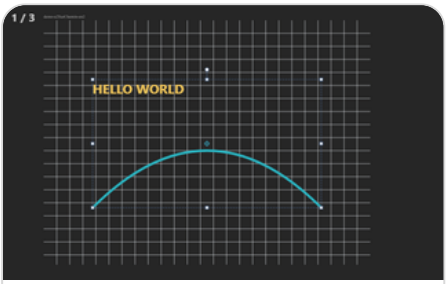

On path

Your words lie down on the line you drew and follow it wherever it goes.

`Appearance › On path`

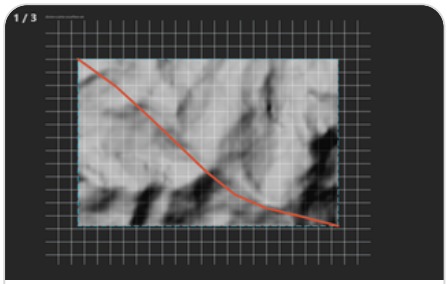

Contours

Lines join the places of the same height on your map, like the steps of a mountain.

`Real map › Contours`

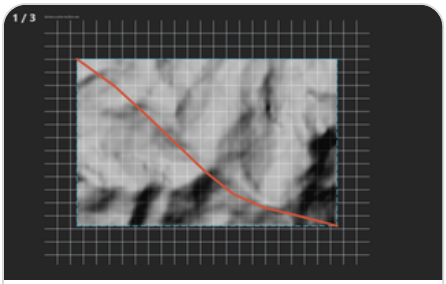

Slice

Your map turns into six-sided tiles, colored by height: water, meadow, forest, mountain.

`Real map › Slice`

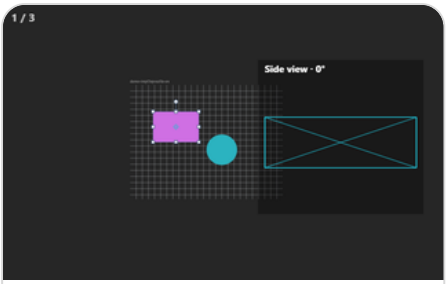

Side draft

The sides of your piece lean: the more you turn, the wider the bottom and the smaller the top.

`3D printing › Side draft`

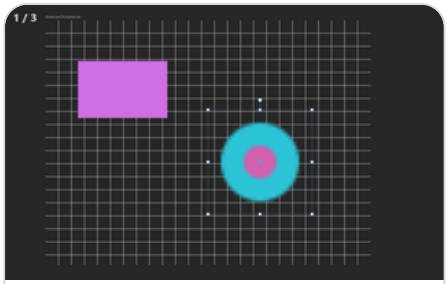

Set as reference

Your picture becomes a reference: it fades, stays still, and a grid of squares lies on it to guide you.

`Pixel persona › Set as reference`

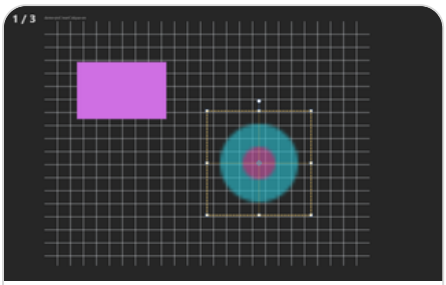

Create pixel layer

A see-through sheet lies on your reference picture: you draw on it, square by square.

`Pixel persona › Create pixel layer`

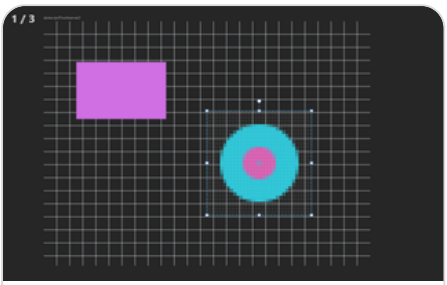

Pixelate image

Your picture is now made of big squares, like bricks: as many squares as the tile size.

`Pixel persona › Pixelate image`

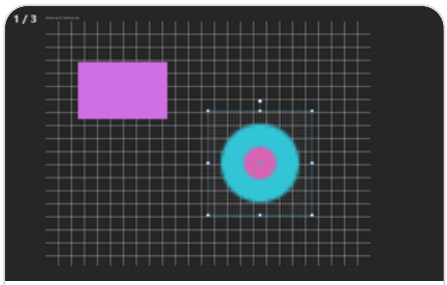

Dark outline

A dark line goes around your drawing so it stands out from the background.

`Pixel persona › Dark outline`

Photolab

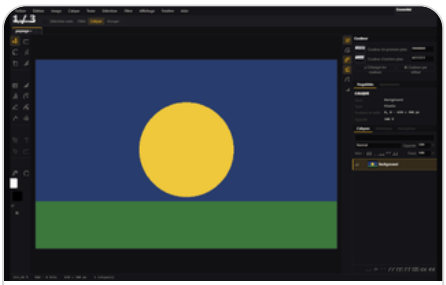

Workspace

Pick a workspace in the list: panels and tools rearrange for what you want to do.

`Workspace list`

Rectangular marquee (M)

Drag across the picture: a dotted rectangle surrounds the part you are going to change.

`Tools › Rectangular marquee (M)`

Magic wand (W)

Click a colour: the whole patch of that colour is selected at once.

`Tools › Magic wand (W)`

Crop (C)

Draw a frame on the picture then press Enter: everything outside the frame disappears.

`Tools › Crop (C)`

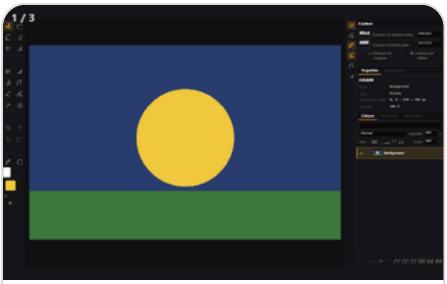

Brush (B)

Drag across the picture: a stroke of your colour is painted wherever you go.

`Tools › Brush (B)`

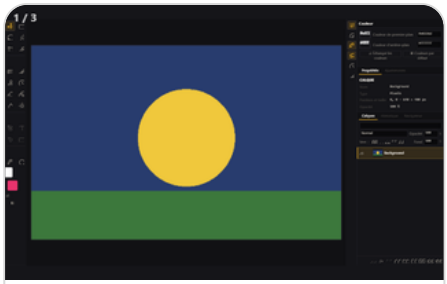

Clone stamp (S)

Alt-click what you want to copy, then paint elsewhere: the same piece of picture reappears under your brush.

`Tools › Clone stamp (S)`

Black and white

One click adds a layer that removes the colours: your picture turns grey, and you can remove it whenever you like.

`Adjustment layer › Black and white`

Sprite Lab

Auto-detect

The tool counts the pictures on your sheet by itself and draws a box around each one.

`Sheet › Auto-detect`

Feet

All your characters put their feet at the same height, as if standing on the same ground.

`Sheet › Alignment › Feet`

Mirror

Your character turns around and looks the other way, like in a mirror.

`Preview › Mirror`

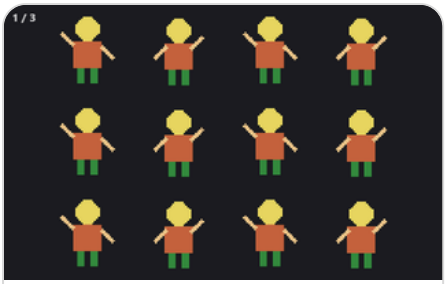

Pixel art (9b)

Your drawing becomes big squares and only a few colours, like in an old video game.

`Pixel art (9b) box`

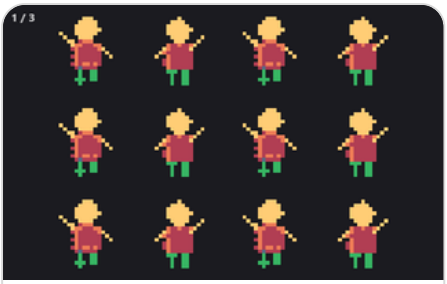

Outline

A line surrounds your character so it shows up on any background.

`Post-processing › Outline`

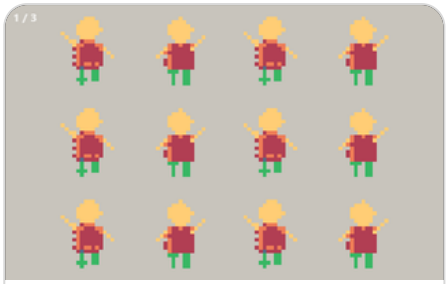

Shadow X / Y

A shadow appears behind your character, slightly offset, as if a lamp were lighting it.

`Post-processing › Shadow X / Y`

Tile Lab

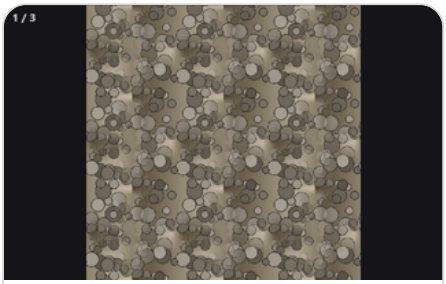

Method

Mirror folds your texture like a sheet folded in four: the edges always meet perfectly.

`Seamless mode › Method`

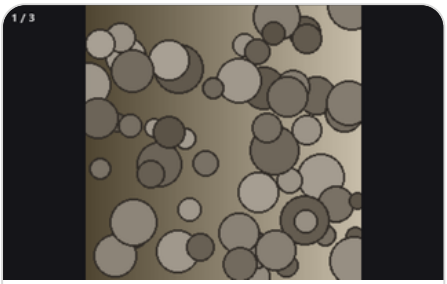

Make seamless

Your texture joins up: laid side by side many times, you can no longer see where each piece ends.

`Seamless mode › Make seamless`

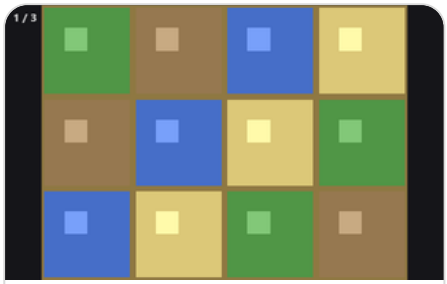

Placement grid

Your cells become diamonds, staggered like the paving of a game seen at an angle.

`Tile sheet › Placement grid`

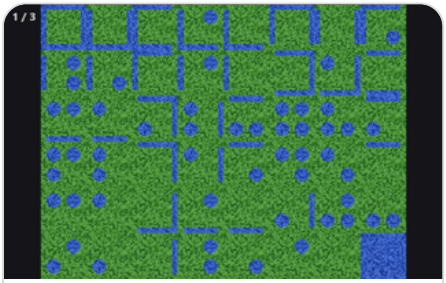

Tileset

You choose how many pieces to make: many for nicely rounded corners, or sixteen to go fast.

`Set mode › Set`

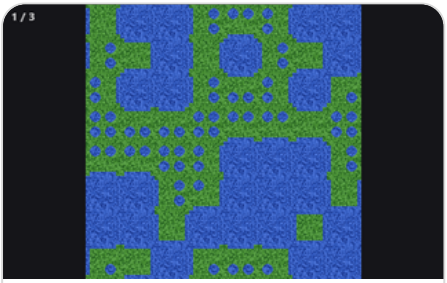

New preview

A new map draws itself at random with your tiles, to see if they go well together.

`Set mode › New preview`

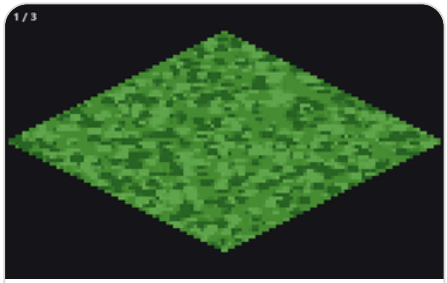

Shape

Your tile changes shape: a diamond for a game seen at an angle, or a six-sided hexagon.

`Shapes mode › Shape`

Good to know

The icon of the app, of its shortcuts and of browser tabs is the real Deepotus logo. Inside the app, the logo shown follows your active brand kit (chapter Templates and brand kits).

One single chevron is used everywhere to expand and collapse: pointing down, the section is open; pointing right, it is closed.

REFERENCE

Troubleshooting

Symptom	Fix
`down` pill top-right	The engine isn't running: relaunch the Desktop icon.
Red pill on `fal` / `heygen` / `voice`	Missing/invalid key: Settings → API keys, fix, restart.
HeyGen avatars are slow to appear (or the list looks empty at first)	**Normal on the very first load**: HeyGen returns 1,000+ avatars and takes up to ~1 min to respond (their service is slow). The app pre-loads the list at startup and caches it — wait once, then it's instant. If the **HeyGen status** is green in Settings, your key is fine: it's neither the key, the network, nor the antivirus.
A voice gives a "paid plan" error	On the ElevenLabs free plan, only "premade" voices work via the API; custom library/community voices need a paid ElevenLabs plan. Pick a premade voice, or upgrade ElevenLabs.
After a restart, keys/library "gone"	No data is lost (it lives in `%LOCALAPPDATA%\DeepotusVideoGenData`). Relaunch cleanly via the Desktop icon.
The plan uses the built-in planner	No Anthropic key nor Ollama model (chapters 11–12).
`CERTIFICATE_VERIFY_FAILED` error	Antivirus inspecting HTTPS — handled automatically; otherwise disable HTTPS inspection.
"Port 8765 already in use"	The app is already running — open `http://deepotus.localhost:8765`.
Where are my files?	Images: `…\DeepotusVideoGenData\assets\images` · Videos: `…\assets\outputs\final` · Your uploads: `…\assets\outputs\uploads` · Sounds: `…\assets\audio`.