

Deepotus Video Gen guide

Everything you can do with the app, one task at a time: short steps, animated pictures taken from the real app, and the icons you will see on screen.



Take a tour of the app



Install, pick a language, update



Connect your keys and AI engines



Keep spending under control



Settings, personas and branding



Move to another computer



Drive it from your phone



Make a video in minutes (Quick)



Compose with the Studio



Reuse templates and brand kits



Tell a story in episodes



Sound, voices and visual effects



Play an avatar, deferred or live



Install and train a local RVC voice



Edit a video



Grade, mix and add titles



Deliver: renders and formats



Plan and publish your week



Connect your accounts



Make game assets



Create playing cards (Card Forge)



From image to 3D object



Prepare a part before slicing (Workbench)



3D-print your creations



End-to-end recipes, from idea to post



Read the icons



Troubleshooting




GET STARTED

Take a tour of the app

You will find your way around the app, move from one screen to another without searching, and read what the top bar tells you. Five minutes, free.

Deepotus Video Gen is a studio that runs **on your own computer**. It turns an idea into a video, a picture, a 3D object or a published post. You plug in your own AI provider keys: you pay those providers, not the app. Everything you create stays with you.

The screen at a glance




1. The rail: one screen per task 2. Ctrl+K: go anywhere from the keyboard 3. This month's spending and key status, always in sight
The screen rail on the left, the top bar with the quick command, the spending counter and the key status.

- 1 The rail, on the left:** one screen per task. **Quick** for a one-click video, **Studio** to compose, **Edit** to cut a video, **Library** to find your media, **Settings** for the keys. The arrow at the top of the rail folds it into icons only.
- 2 The quick command:** click **Quick command...** or press **Ctrl + K**, type a few letters, Enter.
- 3 The spending counter** shows what you spent this month with your providers. One click opens Settings.

4 The badges `fal`, `heygen` and `voice` turn green when the matching key is set. One click takes you to that key.

5 The **render queue**, top right, lists what is being computed and what just finished, in the `Running` and `Recent` tabs.

Go fast with the quick command



1. Type, then Enter

Ctrl+K opens the quick command: type a few letters to go to a screen or start an action.

The quick command goes to any screen (`Go to Studio`, `Go to News` ...), opens the last render, creates a starter graph or opens a Settings section. `Replay onboarding` restarts the first-launch tour.

The screens of the rail



Quick

A one-click video: animated clip, talking avatar, composition, voiceover.



Studio

A node graph to compose a video piece by piece.



Chapters

From a text or a novel to a narrated video, scene by scene.



Sound & VFX

Music, voices, sound effects and visual effects.



Edit

A multitrack timeline to cut, grade and deliver.



Scheduler

Plan and publish your posts for the week.



Templates

Reusable layouts and your brand kits.



News

Follow the news and turn it into short videos.



Library

All your pictures, renders, sounds and objects, with their origin.



Game Assets

3D, sprites, tiles, materials and playing cards.



Vectorlab

Vector drawing, pixel art, maps and printable objects.



Photolab

Photo retouching with layers, masks and filters.

The guide opens from the foot of the rail: [FR](#) or [EN](#) .

What the app does to protect you



Nothing leaves without you

Only this PC can change, generate or spend. Another device on the network can only read, and only once paired.



Errors that make sense

A missing key or a server refusal is shown in plain words, with what to do.



Safe imports

A picture with the same name becomes name-1, name-2...; a file that is not a real picture is refused and named.



In-app dialogs

Confirmations and text prompts use one Deepotus dialog, never a browser pop-up.

Good to know



The desktop icon opens the app at `http://deepotus.localhost:8765`. Each browser tab carries the name of the open tool, for example “Deepotus — Vectorlab”.



Only AI generations cost money, and your provider bills them. The chapter [Keep spending under control](#) shows how to set limits.

Shortcuts

Key	Effect
Ctrl + K	Opens the quick command
Esc	Closes the quick command or a dialog



GET STARTED

Install, pick a language, update

You will launch the app for the first time, pick its language, know what each use costs and keep it up to date. Ten minutes, free.

First launch


- 1 During installation, pick the language: **French** or **English**. It becomes the language of the whole app.
- 2 Double-click the **Deepotus Video Gen** desktop icon. The app starts without a black window, then your browser opens on the studio.
- 3 On the very first launch, a five-step welcome introduces the studio and lets you pick a persona, your generation providers and the networks you publish to.
- 4 Set at least one key, the **fal.ai** one: it makes pictures and video clips. The chapter [Connect your keys](#) shows how.



The welcome can be replayed at any time from the quick command (**Ctrl** + **K**):

Replay onboarding .

Switch language



1. French or English, for the whole app

Settings › Provider defaults › Interface language: the page reloads in the chosen language.

1 Open **Settings** › **Provider defaults**.

2 In **Interface language**, pick **Français** or **English**. The page reloads in the chosen language.

Translated: the sidebar, Settings, the Library, Quick, News, the Scheduler, the Studio, Templates, Episodes, the editor, Sound & VFX, subtitles, machine transfer and every confirmation dialog. Lists sent by the server (effects, transitions, presets) also arrive in the chosen language. The labs (Vectorlab, Photolab, Card Forge...) are being translated one by one; meanwhile their known texts are translated on display. What you type yourself is never translated.


What each use costs

The app is free. Only AI generations cost money, billed by the provider of the key. Rough figures at default rates:

Action	Provider	Indicative cost
Create a picture	fal.ai (FLUX)	≈ \$0.003
10 s video clip	fal.ai (Seedance 2.5, 720p)	≈ \$4.73
Talking avatar, 1 min	HeyGen	≈ \$0.24
Voiceover	ElevenLabs	per character, free tier available
News summary, post plan	Anthropic, OpenAI, Gemini or Ollama	a few cents, free with Ollama
Import your videos, sounds, pictures	none	free

Your real rates go into **Settings** > **Pricing & budget** ; see [Keep spending under control](#).

Keep the app up to date



1. Your version is up to date 2. Check asks for the latest published version

The version block in Diagnostics: the installed version, its status, and a button to check right now.

- 1 Once a day, the app checks whether a new version is published. It never blocks anything.
- 2 When there is one, a card announces it: **Download** fetches the installer and shows where it is saved, **Later** hides the card for that version, **Notes** tells what changes.
- 3 To check right now: **Settings** > **Diagnostics**, **Check** button.
- 4 Close the app, run the downloaded installer: nothing is installed on its own.

Good to know

i Your keys, files and database live in the `%LOCALAPPDATA%\DeepotusVideoGenData` folder, apart from the program: they survive updates and reinstalls.



Start with the fal.ai key alone, then add avatars, voices and AI as you go. Set a monthly cap from day one: see [Keep spending under control.](#)



GET STARTED


Connect your keys and AI engines

You will set an AI provider's key, choose which engine does what, and protect your keys with a password. Ten minutes, free; keys are created with each provider.

Before you start

An API key is a code the provider gives you on its website, in your account. It lets the app work on your behalf; the provider bills you. Only one is required to start: **fal.ai**.

Set a key



1. The key's status 2. Guide: where to create the key, and its price 3. Paste the key, Save: it works right away

Settings > API keys: each provider has its row, its status, its guide and its Save button.

1

Open **Settings** > **API keys**. Each provider has its row and its status: **missing** until a key is set.

2

Click **Guide**: it tells where to create the key with the provider and how it is billed.

- 3 Paste the key in the field, then **Save**. It works **right away**, no restart.
- 4 Click **Test** on the row: the provider says whether it accepts the key, without spending anything.

To set several keys at once, fill in the fields then **Save all changes**.

Which key unlocks what

Key	What it unlocks
fal.ai	Pictures (FLUX) and video clips (Seedance). The only key you need.
HeyGen	Talking avatars, alone or in a composition.
ElevenLabs	Voiceover and narration.
Anthropic, OpenAI or Gemini	News summaries, post plans, scripts and storyboards written by AI. One is enough; OpenAI also makes pictures (GPT Image).
Meshy	Premium picture-to-3D (optional, Meshy subscription).
Figma	Import Figma layers into the Library and Templates.
Ollama (no key)	An AI model installed on your PC for post plans: free, nothing leaves your machine.

Network keys (X, Telegram, YouTube, Instagram, TikTok) go into **Connected accounts**: see [Connect your accounts](#).

Choose the engine for each job

In **Settings** > **Provider defaults**, each role has its engine: picture generation, voice, picture to video, avatar, voiceover, news summary, post planning. Only keys actually set are offered.



Default video model

The model new Seedance nodes start with in the Studio. Each node can change it.



Default voice model

The ElevenLabs model of new voiceover nodes.



Local AI with Ollama


Install Ollama, download a model, then enter its name in the **Local LLM (Ollama)** card of API keys.



Gemini: pick the version

Left empty, the “Gemini — model” field always follows the latest Flash version; fill it in to pin one.

Protect your keys in the vault



1. A master password encrypts all your keys 2. The .dzk archive carries them, with its own password

Settings > Vault: your keys encrypted with a password, and the archive to carry them to another computer.

1 Open **Settings** > **Vault** .

2 Choose a master password (at least 8 characters), type it twice, then **Set up vault** . Your keys are encrypted and leave the plain-text file.

3

Close removes the keys from memory: nothing paid can go out. **Open** makes them active again with the password.

4

Remember on this PC opens the vault on its own at launch, for your Windows session;

Stop remembering asks for the password again.

On each row of **API keys**, a badge tells where the key lives: in the file, in the vault, or vault closed (one click takes you there).

Good to know



The vault password is stored nowhere: if lost, it cannot be recovered. Write it down, or keep an encrypted archive.



Export archive creates a **.dzk** file that carries your keys, caps and rates to another PC or your phone, protected by its own password. **Import an archive...** takes them back.



Settings > **Diagnostics** shows the status of all keys at once (never their value) and the providers' balances.



GET STARTED


Keep spending under control

You will know what you spend, set a cap per month and per engine, and decide yourself when a render may go over it. Five minutes, free.


Where to see what you spend

The dollar counter in the top bar shows the month's total. One click opens **Settings** > **Pricing & budget**, where everything is set.

Set your rates and caps



1. A cap for each video generation



2. Above a cap, the app asks you before spending

Settings > Pricing & budget: your rates, then one cap per engine and the bar of what you spent this month.

- 1 At the top of **Pricing & budget**, enter your providers' rates (price of a picture, of a second of video, of a voice character...). The costs shown in Quick, the Studio and the Scheduler become accurate. **Save pricing**.
- 2 **Video cap per request** limits the cost of a single video generation; **Max generated seconds per clip** limits its length.
- 3 Further down, **Monthly caps**: one overall cap and one per engine (fal, OpenAI, Anthropic, ElevenLabs, HeyGen, Meshy), in dollars. A bar shows what was spent.
- 4 Set **Alert at (% of a cap)**: a banner warns you once a month when a cap reaches that level. **Save caps**.

When a render would go over a cap

Before leaving, every paid generation is checked against your caps. If it goes over, a “Monthly spending cap” dialog opens, wherever you are in the app:

- 1 “Tirer quand même” (run anyway) sends exactly the same request, once.
- 2 “Annuler” (cancel) sends nothing and tells you why.

Actual versus estimated

Below the caps, the **Monthly spend — actual vs estimated** table gives one row per engine and per screen: the number of runs, what was estimated, what was actually billed, and the difference.



Actual

The provider gave the exact price (Meshy, HeyGen): that is what counts.



Partly actual

Some runs have their exact price, the rest is still estimated.



Estimated

The provider gives no price: the app counts the quote made at launch.

Good to know



Spending made from your phone goes into the same record and the same caps, counted once.



The Studio sends the cost it displays as the request's cap: a render cannot cost more than what you saw before clicking.




GET STARTED

Settings, personas and branding

You will find any setting, give your content a voice with personas and dress the app in your colours with brand kits. Ten minutes, free.

Find a setting



1. Type what you are looking for
 2. The matching sections show up
 3. Enter opens the first one
- Settings search finds a setting by its name and opens the right section directly.

- 1 Open **Settings** in the rail.
- 2 Type a word in **Search settings...**: accents and capitals do not matter.
- 3 Enter opens the first result, Esc clears the field.

The Settings sections


Section	What you do there
Diagnostics	Version, disk space used, latest alerts, key status and provider balances.
Vault	Encrypt your keys and carry them (chapter Connect your keys).
Devices	Pair a phone (chapter Drive it from your phone).
API keys	Set and test the AI providers' keys.
Connected accounts	Link X, Telegram, YouTube, Instagram and TikTok (chapter Connect your accounts).
Personas	The tone of your content.
Branding	Name, logo, colours: your brand kits.
Caption pack	One-click brand labels under the Scheduler's caption editor.
Provider defaults	The engine of each role, default models, language.
Paths	Where your pictures and renders are stored (button to copy the path).
News	Settings of the News screen.
Appearance	Reduced motion and visual effects.
Pricing & budget	Your rates and caps (chapter Keep spending under control).
Machine transfer	Take your work to another PC (chapter Move to another computer).

Give a voice with personas

A persona is a personality: its tone, its audience, its way of speaking. The active persona sets the News scriptwriter, the tone of the prompt generator and the default voiceover.

- 1 Open [Settings](#) > [Personas](#) .
- 2 Built-in personas can be read but not edited: duplicate one, or click [New persona](#) .
- 3 Click a persona to make it **active**.

Dress the app with a brand kit



1. Create a kit 2. The active kit dresses the whole app

Settings > Branding: several brand kits, a single active one, and its fields right below.

- 1 Open **Settings** > **Branding**. The kit list shows which one is **active**.
- 2 **New kit** creates one; **Activate** makes it the active kit and refreshes the app.
- 3 The fields below edit the active kit: **Logo...**, app name, sub-label, two tagline lines, primary and accent colour. **Save brand**.
- 4 **Rename**, **Duplicate** and **Delete** manage each kit. The active kit and the last kit cannot be deleted: the greyed-out button says why.

The active kit dresses the splash screen, the sidebar, the app's accents and Template renders that use the brand. **Reset to deepotus** restores the original look.

Good to know

i A Template render keeps the kit that was active when you sent it, even if you switch kits while it computes.

i **Settings** > **Appearance** > **Reduced motion** turns off decorative animations; it also follows the Windows setting.



GET STARTED


Move to another computer

You will take all your work to another computer, check that nothing got damaged on the way, and get your keys back there. Fifteen minutes, free; the app must be installed on both machines.


What travels

The transfer takes everything the app created: the Library and the origin of every file, renders, the bible, communication plans, vector documents, card series and 3D models. Your API keys **never** go into the package.

Export from the old PC



1. What goes, and how much it weighs



2. Extras to tick: logs and discards 3. A dated folder is written at the destination

Settings › Machine transfer: a dated package, each file with its checksum; keys stay on this PC.

- 1 Open **Settings** › **Machine transfer**: the number of files and the size to take are shown.
- 2 Click **Export...**, then pick the destination: a drive, a USB stick, or a path.
- 3 Tick the extras if you want them: **Logs** and **Discards — dated trash folders**. They are unchecked by default, with their size.
- 4 **Start export**: a dated folder is written at the destination, each file with its checksum. Do not close the app while it writes.
- 5 Click **Check integrity**: every written file is read back and compared with its checksum.

Import on the new PC

- 1 Install the app, then open **Settings** > **Machine transfer** and click **Import...**.
- 2 Give the package folder, then **Check**: the app reads the package without writing anything and tells where it comes from (date, machine, version).
- 3 **Start import**. Import **adds**: nothing this PC already has is overwritten. A damaged file is left out and named.

Take your keys too

- 1 On the old PC: **Settings** > **Vault** > **Export archive**, with a separate password.
- 2 On the new PC: **Import an archive...** and that password. Your keys, caps and rates come back.

Good to know

i Paths from the old machine are rewritten for the new one: your projects find their files without any setting.

! Stopping an export leaves an incomplete package. If it happens, run the export again, then **Check integrity** before importing.



GET STARTED


Drive it from your phone

You will link your phone to the PC, carry your posts and chapters in your pocket, and cut off a lost phone. Fifteen minutes; you need the Deepotus companion app on the phone and the same Wi-Fi.

Before you start

By default, the app only listens on this PC: no other device can reach it. To open it to your home Wi-Fi, replace `HOST=127.0.0.1` with `HOST=0.0.0.0` in the `.env` file of the data folder, then restart the app. Windows will ask to allow Python through the firewall.

Pair the phone



1. By default, only this PC can reach the app 2. Shows a QR code valid for 5 minutes

Settings > Devices: the network status, the list of devices, and pairing by QR code.

- 1 Open **Settings** > **Devices**. The words **this PC only** or **local network open** tell whether the phone will be able to reach you.
- 2 Click **Pair a device**: a QR code appears, valid for five minutes, for a single device.
- 3 In the phone app, open **Pair** and scan the QR code. It closes on its own once read.
- 4 To give your keys to the phone, export the encrypted archive from the **Vault** and open it in the phone app, with its password.

Five devices at most. A paired phone **reads** what the PC produces; it never sees your keys. It only writes to the PC to drop images and for Avatar live, whose spending goes through your caps.

What the phone does



Sync

The phone fetches the Library and its pictures, and resumes an interrupted download.



Send pictures to the PC

Your photos land in the Library with the "mobile" source, without duplicates.



The week's batch

Take the approved posts and publish them on Telegram or X from the phone. A post taken by the phone is no longer published by the PC.



Inbox

What you share to Deepotus lands here; pictures go on to the PC.



Offline chapters

Take a chapter, write without a network, hand it back to the PC.



Generate pictures

The cost is shown before each generation, at your PC's rates, and counted in your caps.



Avatar live

Film a take and start a Recast, or open a Live session: the PC does the work, the price is shown first (see [Play an avatar](#)).

A chapter taken to the phone

While a chapter is on the phone, it is read-only on the PC, with a banner. You can force it back to the PC. If the two versions disagree, the phone's text stays in the log, to copy or take back: nothing is lost.

Remove a device

- 1 In **Settings** > **Devices**, click **Revoke** on the phone's row and confirm.
- 2 Access is cut right away; the posts and chapters it had taken come back to the PC.

Good to know

 Lost your phone? Revoking is not enough: your keys travelled in the archive. Regenerate them with each provider; the Devices page gives the link to each console.

 The phone is notified when a render finishes or fails, when a post is published or awaits an action, and when a cap is near.



MAKE VIDEOS



Make a video in minutes (Quick)

You will make a vertical video in minutes: an animated clip from a picture, an avatar saying your text, both together, or a voiceover. Paid depending on the provider; the cost is shown before you click.

Before you start

Quick has four tabs: **Seedance** (animated clip), **HeyGen** (talking avatar), **Composition** (avatar and clip together) and **Voiceover**. A tab stays greyed out while its provider's key is missing: fal.ai for Seedance, HeyGen for the avatar (chapter [Connect your keys](#)). The voiceover can also run for free on your PC with Voicebox.

Animate a picture into a clip




1. Drop the picture to animate
 2. Model, duration and format
 3. The cost is shown before you click
- Quick, Seedance tab: the start image, the parameters, the estimated cost, then Generate.

1. Open **Quick**, **Seedance** tab.
2. Pick the **start image**: drop it on the area or click **Browse thumbnails...** to take it from the Library. An **end image** is optional; if the chosen model does not accept one, the field says so and names the models that do.
3. Write what should move, or start from a **prompt template**, or click **Open prompt generator**.
4. In **Parameters**, set the model, the duration and the format (9:16 by default).
5. Read the **estimated cost**, then click the **Generate** button. The render goes through the render queue, top right, and lands in the Library.

🔄 With the default model, Seedance 2.5 at 720p, one second costs about **\$0.47** : a 10 s clip comes to about **\$4.73** . A shorter duration or another model costs less; the estimated cost updates with every setting.

Direct the camera



1. Six sliders, one camera sentence



2. The Gallery shows each camera move

Camera sliders become a sentence added to the prompt; the Gallery shows the moves.

1. Unfold **Camera (sliders)**: six sliders (push in, turn, rise...) become a sentence suited to the model, added to the prompt. It is shown before the render.
2. To choose by eye, click **Gallery** : eleven camera moves in three styles, as small animations computed on your PC. It is a reminder, not a promise of what the model will do.

Make an avatar talk

1. **HeyGen** tab: pick an avatar and its voice. The lists show up right away, from a copy kept on your PC and refreshed in the background.
2. For an avatar that looks like you, import a photo: it becomes a talking avatar, selected automatically.
3. Write the script, then generate. The pauses in your text are kept.

Videos go through HeyGen and its Avatar III engine, picked according to the avatar. The HeyGen balance is shown in dollars. Count about **\$0.24** a minute at default rates.

Compose and dub



Composition

An avatar and a clip in the same video, one after the other or side by side (vertical or horizontal split).



Voiceover

Choose the provider for each generation: Voicebox, on your PC and free, or ElevenLabs. A missing provider is refused, never silently replaced.



Music under the voice

When a render has both a voice and music, the music dips on its own under the voice.



Burned-in subtitles

Subtitle the render :

aligning a text you provide is free; automatic transcription is paid and only runs if you tick it.

Go further



Presets

At the top of each tab, save your settings under a name and reapply them in one click.



Lip-sync

Syncs the clip's lips to a voiceover (Kling, paid); limits and price are shown beforehand.



Reopen in Quick

Every render keeps its settings: **Reopen in Quick**, from the Library or the render queue, puts them back as they were.



Extend a clip

From **Send to**, **Extend** adds about 7 generated seconds to a clip already rendered, with or without sound; the price is on each button.

Good to know



Leaving Quick and coming back finds your settings as you left them.



For a still picture, use the Library: **Create image**, with the **Prompt manager**. For your own filmed video, **Upload video** in the Library, then the UGC video node of the [Studio](#).




MAKE VIDEOS

Compose with the Studio

You will compose a video piece by piece by linking nodes, see it before paying, never pay twice for what is done, and reuse a graph like a recipe. Paid depending on the nodes; the graph's cost is shown before running.

The Studio screen



1. Start from a ready-made graph 2. Preview: see before paying 3. Run: the cost shown becomes the cap

The nodes palette, the graph, the inspector; at the top, the free preview, the cost and Run.

On the left, the **nodes** palette (**/** key) and the **starter graphs**. In the centre, the graph. On the right, the inspector of the selected node, which can be folded. At the top, the graph bar: **New**, **Save**, **Recipe**, **Import**, the format, **Preview**, the estimated cost and **Run**.

Compose a graph

- 1 Start from a starter graph (**Seedance solo**, **Avatar post**, **UGC + animation**, **News reel post**, **Timeline montage**) or from a blank page with **New**.
- 2 Drag nodes from the palette: sources (picture, text, your video), generators (Seedance, HeyGen avatar, pictures), audio, editing, composition, output.

- 3 Link an output, on the right of a node, to an input, on the left of another. Click a node to set it in the inspector.
- 4 **Preview** computes the final composition for free: clips not generated yet are replaced by their picture.
- 5 Read the estimated cost, then **Run**. That amount is sent as a cap: the render cannot cost more. **Save** keeps the graph under a name.

Never pay twice for what is done

The paid render of a Seedance or HeyGen node is **pinned** automatically: as long as its request does not change, it is reused for free, and the estimated cost says which nodes are.



The last eight renders

The inspector lists the node's last renders; **Use** picks the one that feeds the rest.



Regenerate

For a new paid take despite the pin, explicitly ask for the node to be regenerated.

Reuse a graph as a recipe

- 1 **Recipe** freezes the graph: its pictures and texts become blanks to fill in.
- 2 Run the recipe with new values: a quote is shown, you confirm. Unchanged nodes are reused for free.
- 3 From a picture in the Library, **Send to** also offers to run a recipe on it.

Go further



Engine duel

On a picture generation node, the same prompt goes to two models once you agree on the cost; you compare side by side and keep the winner.



Effects and masks

The Effects / Mask node, plugged into the Render's fx input: grades, VHS, colorisation..., on the whole render or a single layer.



AI fields

A field with a coloured edge is sent to an AI. Its badge shows the model and its price; its microphone dictates your text.



Voice and mix

The Voiceover node picks Voicebox or ElevenLabs; Audio mix dips the music under the voice (light, medium, strong).



Import a graph

Import opens a JSON graph; missing sources are listed and nothing is saved without your consent.



Frame by frame

The result drawer steps frame by frame with the **,** and **.** keys.

Good to know



A video generation above the announced cost or your per-request cap is refused before any call. The generated length is capped in Settings; the rest is extended on your PC.



New Seedance nodes start with Seedance 2.5; a saved graph keeps its model. The default model is changed in **Settings** > **Provider defaults**.



If the browser cannot dictate, the recording can be transcribed by the server: it is paid, a quote is shown and "No" is the default choice.

Shortcuts

Key	Effect
/	Opens the nodes palette
Ctrl + Enter	Runs the graph
, .	Previous frame, next frame in the result




MAKE VIDEOS

Reuse templates and brand kits

You will start from a ready-made layout, touch it up (slots, masks, texts, animations), adapt it to another format and dress it in your brand. Free; only renders of AI-generated slots cost money.

Pick a template




1. A blank layout in the format you want 2. The same layout, in another format

One card per layout; New in a format, or Reflow to another format.

- 1 Open **Templates** in the rail. The gallery shows a real picture of each layout; the search field filters by name or tag.
- 2 Click a card to open it in the editor, or create a blank one with **New** and a format (9:16, 16:9, 1:1, 4:5).
- 3 For the same layout in another format, **Reflow** to creates a copy with a before / after preview; the original does not change.

Touch up a layout



1. Add text, ticker, brand, component 2. Click a slot: drag it, resize it 3. Save, then open in the Studio

The slots on the canvas, and on the right the settings of the selected slot: mask, sample, animation.

- 1 Click a slot on the canvas: drag its body to move it, its handles to resize it. The inspector, on the right, shows its settings.
- 2 **Mask** gives a picture or video slot a shape: rounded, ellipse, polygon (triangle, diamond, hexagon, star, or your own points), cut-out windows, soft edge, coloured rim.
- 3 **Animation** brings the slot in and out with a fade, slide or pop, with duration, delay and curve. **Replay on canvas** shows it.
- 4 **Add** places a text, a ticker, a separator, the brand, a badge or a component.
- 5 **Save layout**, then **Open in Studio** to fill it in and render it.

Polish the texts



Ready-made styles

Social subtitle, neon, banner, gradient title; or your own settings: outline, shadow, rounded background, gradient.



Text that fits

The text stays in its slot: it wraps, then shrinks down to a minimum size.



Text on arc

Sets a title as an arch or a smile; an arc too large is reduced to fit.



Exact preview

Renders the slot with the real render engine, to see the final result.

Dress it with your brand

Templates use your **brand kit**: name, colours, logo. Kits are managed in **Settings** > **Branding** (chapter [Settings, personas and branding](#)). A render keeps the kit active when it was sent, even if you switch while it computes.

Go further



Components

Save as component turns slots into a reusable group; each instance can change its texts and colours, then go back to the component.



Preview sample

A Library picture fills a slot for thumbnails and picture export; never for the video render.



Export

Export image as PNG, JPEG or WebP (it lands in the Library), or as an SVG to drop into Figma.



To the Vectorlab

[Open in the Vectorlab](#)

opens the layout as an editable vector document.



From Figma

[Import a Figma frame...](#)

turns it into an editable template; the number of Figma calls is announced beforehand (Figma key required).



Avatar in the right format

A HeyGen avatar is generated in the format closest to its slot: no empty band, no cut-off head.

Good to know



In the Studio, the layout node shows the transition planned by the template as long as you do not pick another. The "cyan flash" transition is a real colour flash, of your choice.




MAKE VIDEOS

Tell a story in episodes

You will turn a text, a chapter or a novel into a narrated, illustrated video, scene by scene, and prepare its storyboard before paying for any render. Free up to narration and pictures; their cost is shown beforehand.

Two ways of working



1. The original flow: from text to narrated episode 2. The Workshop: manuscript, bible, storyboard

Chapters offers the original flow (Episodes) or the full Chapter Workshop.

Open **Chapters** in the rail. The **original flow** (Episodes) goes straight from text to narrated video in four steps. The **Chapter Workshop** is the full space: manuscript, bible of characters and places, screenplay, storyboard, animatic.

An episode in four steps

1. Four steps, left to right 2. Paste the text or upload a file 3. The voice reads the whole chapter


Episodes: script and voice, scenes, illustrations, assembly.

- 1 Script and voice:** give a title, paste the text or upload a .txt, .docx or .pdf. Pick the voice and language, then **Generate narration**: the whole chapter is read and stored in the audio Library.
- 2 Scenes:** split the text by paragraph (word for word) or by AI (into visual beats, with an illustration prompt). Edit, reorder, add or remove scenes.
- 3 Illustrations:** one picture per scene, with the model of your choice. For each scene, pick the animation: slow camera move (Ken Burns), real Seedance clip or still picture.
- 4 Assembly and export:** a quote gives the price of the Seedance scenes, then the episode is cut on its narration. If a Seedance scene fails, it falls back to a camera move and you are told.

At the top, **New**, **Open** and **Save** manage your episodes; a dot marks what is not saved.

Narrate scenes never pays again for a scene already narrated with the same text, voice and language.

The Chapter Workshop



1. Three modes: script, screenplay, storyboard 2. The bible: characters, places, objects...

3. The storyboard, then the free animatic

The chapter text in the centre, the bible on the right; Script, Screenplay and Storyboard at the top.

The Workshop is not translated yet: its buttons below are quoted as they appear, in French.

- 1 Script:** write or import the chapter text; **Manuscrit** imports a whole novel, split into chapters, and fills the bible. Recognised names are highlighted.
- 2 Bible,** on the right: characters, places, objects, dates, moods, sets. Each entry has its description, its picture sheet and, for a character, its voice and temperament.
- 3 Screenplay:** **Adapter (IA)** (adapt with AI) rewrites the chapter as a screenplay without touching the manuscript; **Voice-over** has each scene read by its character's voice.
- 4 Storyboard:** **Découper (IA)** (split with AI) or **Paragraphes** (paragraphs) creates the shots, with their framing, camera, duration and the entities in them.
- 5 Animatique** (animatic) cuts the shots into a rehearsal video, silent and free, to watch before paying for a render. It then goes into a new Edit project, as a full film or a 30 s reel.

Go further



Versions

Every time the text is overwritten, a copy is kept (ten at most): compare, restore, copy.



Exports

The manuscript as .docx or PDF, the screenplay in industry format, the storyboard as a PDF with four shots per page.



Import a screenplay

A Fountain or Final Draft file replaces the scenes (the old screenplay is kept) or is added after them.



Production picture

A shot's picture is generated with the characters' views as references; the cost is stated and confirmed.



Appearances

A bible entry lists its mentions, shots and scenes; one click opens the shot in the storyboard.



To an episode

Épisode (episode) creates a narrated episode (one shot, one scene) without paying anything; narration and render then start in Episodes.

Good to know



Rewriting a selection in the bible's tone, generating a production picture or adding the animatic's guide voice is paid: the cost is always stated and confirmed first.



A chapter taken to your phone is read-only on the PC, with a banner; see [Drive it from your phone](#).




MAKE VIDEOS

Sound, voices and visual effects

You will generate music, sound effects and a voiceover you direct, store and find your sounds, and add free visual effects. Each generator's price is shown; visual effects and post-processing are free.

The Sound & VFX screen



1. Music, effects, voice, particles, post: the price is shown 2. Direct the voice's acting

The generators at the top; directed voiceover, sound effects and post-processing below.

Open **Sound & VFX** in the rail. At the top, the **generators** and their price: music, sound effects (606 bundled sounds, plus generation from a description), voiceover, particles and post-processing. Below, three tabs: **Audio**, **Particle VFX** and **Post-processing**.

Direct a voiceover

- 1 In **Directed voice-over**, write the text.
- 2 Click emotion or acting tags (four at most): they go at the start and the Eleven v3 model performs them. The screen shows exactly what is sent to the model.
- 3 **Generate voice**. Trim, add a fade, dip the music under the voice or normalise, then **Send to Montage**.

Outside the v3 model, or with Voicebox, the tags are removed and you are told.

Music and sound effects



Music

Instrumental or sung: describe the style, write the lyrics section by section (a skeleton is suggested from the persona). Billed per second, under your caps.



Sound effects

Describe the sound and its length: two variations to listen to, which join the Library and the editor's Sounds drawer.



Voice and music mix


Mixes a voice and a piece of music, dipping the music under the voice; the mix lands in the Library.



A character's voice

In the Workshop, a character gets a cloned voice (paid) and a temperament, performed on each of its lines.

The Sounds drawer



1. All sounds, then by kind: SFX, voice, music 2. Each sound: listen, tag, enhance, split
Library, Audio chip: the Sounds drawer, with its search, filters and actions.

All your sounds live in the Sounds drawer: in the Library (Audio tab) and in the editor (**B** key). Filter by kind (effects, voices, music, imported), between your sounds and the bundled catalog, search by name or tag.



Enhance

Equaliser, noise reduction, compressor and normalised loudness in one click, on your PC, no provider.



Split into stems

Splits music into vocals, drums, bass... (paid, quote on first click); each stem stays linked to its parent.



Isolate the voice

Removes the background from a voice take (paid, quote on first click).



Search by description

Describe what you want to hear; sounds similar to a sound can be listed too. Needs the local search service, otherwise the button says why it is greyed out.



Preview

Hover a sound to hear it (option to turn on); **Space** plays the selected row.

Free visual effects

- 1 **Particle VFX** : explosions, smoke, sparks, magic... simulated on your PC, to place on a shot.
- 2 **Post-processing** : 16 mm grain, glow, glitch, aberration; applied at render time by the Studio's Render node.
- 3 In the editor's effects rack, an effect can go **behind** the shot's cut-out subject, or in front of it. Cutting out the subject is paid; a first click gives the quote.

Good to know



Split, isolate, cut out: a first click arms the action and shows the price, a second one starts it, **Esc** disarms. Nothing goes out without that second click.



Play an avatar, deferred or live

You will learn to create a Character from a few photos, have that Character replay a video in another setting, and film yourself on your webcam transformed live, with their voice. Twenty minutes. Everything is billed per second: the price shows before each launch and counts against your caps.

Before you start

The **Avatar live** screen opens from the rail. It needs keys depending on what you do: **fal** for Recast and the setting, **Decart** for Live, **ElevenLabs** for the Character's voice. The free local voice needs an NVIDIA graphics card and Voibox (see [Install and train a local RVC voice](#)).

Only clone your own face and voice, or those of someone who has agreed to it. The app makes you confirm this for every Character and keeps the date.

Create a Character

The screenshot shows the 'Avatar live' application interface. At the top, there's a navigation bar with tabs: 'Avatar live', 'Characters, deferred Recast, Live', 'Characters' (selected), 'Recast', and 'Live'. On the right of the bar, it says 'Live: Decart key set'. The main area is split into two panels. The left panel, titled 'New Character', contains instructions: 'One to eight photos of the same face, front-on and well lit (512 px recommended). They are the reference for Recast and Live.' Below this is a 'Name' input field with the placeholder 'E.g. Oli in a suit'. A large dashed box with a plus icon and text 'Drop the photos here or click to choose them' is for image uploads. Below that is a 'Linked voice (optional)' input field with 'ElevenLabs voice ID'. A checkbox with the text 'I have the right to use this face and this voice: they are mine, or the person has agreed to have them cloned.' is present. At the bottom is a 'Create the Character' button. The right panel, titled 'My Characters', shows a list of characters. The first entry is 'Otus (d  mo)' with a placeholder image, '1 photo(s) - voice voix-demo', and 'Consent given on 10/10/2026'. It has buttons for 'Clone the voice again' and 'Delete'.

1. One to eight photos of the same face 2. Consent, dated, for every Character 3. Your Characters, ready for Recast and Live

Avatar live > Characters: reference photos, consent, then the list of your Characters.

1

Characters tab: give the Character a **Name**.

- 2 Drop one to eight photos of the same face, front-on and well lit (512 px recommended). They are the reference for Recast and Live.
- 3 Tick the consent box ("I have the right to use this face and this voice...") then click **Create the Character**.
- 4 Optional: **Clone the voice** with one to five clips of the voice alone (30 s to 3 min in total). Cloning goes through ElevenLabs.

Recast: replay a video

Avatar live

Characters, deferred Recast, Live

Characters

Recast

Live

Live: Decart key set

Recast

A 3 to 30 s video keeps its acting; the character, the setting or an element changes.

1. THE SOURCE VIDEO

Drop a video (.mp4, .mov, .webm)

or a render from the app

Prise de vue de démonstration (6 s)

6.0 s

2. THE MODEL

Replace the character (keeps the setting and the acting)

0.04 \$/s at 480p - 0.08 \$/s at 720p

Animate the character in THEIR setting (the video's movement)

0.04 \$/s at 480p - 0.08 \$/s at 720p

Motion transfer (character and setting from the image, original sound kept)

0.126 \$/s

Motion transfer, pro quality

0.168 \$/s

Replace an element with an instruction (outfit, object, setting)

0.10 \$/s at 480p - 0.15 \$/s at 720p

3. THE CHARACTER

Otus (d  mo)

4. THE SETTING OR THE INSTRUCTION

Neon studioTV setModern officeTokyo street at nightBeach at sunsetMisty forestSpace stationOld libraryDesertClassroomPlain green screenCrypto den

Free instruction (preferably in English), e.g. a red leather jacket

Resolution

720p

☐ With the Character's voice: the take's voice is converted into theirs, at the same pace (ElevenLabs).

Start the Recast

Change only the voice

   0.48 \$ (6.0 s    0.08 \$/s)

Background only

The subject of the source video is cut out, then placed on a new background. The acting, the voice and the size do not change.

New background

Solid color

Color

Change the background


Recast renders

1. Each model shows its price per second 2. A ready-made setting, or your own instruction 3. The quote first, then the launch Recast: a source video, a model and its price, the Character, the setting; the quote shows before you start.

A 3 to 30 s video keeps its acting; the Character, the setting or an element changes. The render takes a few minutes and joins your renders.

- 1 **Recast** tab: drop the source video, or pick a render from the app.
- 2 Choose the **Model**: its price per second is shown next to it.
- 3 Choose the Character, then a setting from the **Presets** or a free instruction (preferably in English).
- 4 Tick the Character's voice if you want the take's voice converted into theirs, at the same pace.
- 5 Read the quote under the button, then **Start the Recast**.

Live: your webcam, transformed



1. The setting, switchable mid-session 2. Your voice, the Character's, or a local RVC voice 3. The reserved price is shown before you start Live: Character, setting, format, reserved duration and voice sent; the reserved price shows before you start.

- 1 **Live** tab: the first time, paste your Decart key and click **Save the key**. It goes to the vault; the page never sees it.

- 2 Choose the **Character** (or keep your face), a **Setting**, the **Format** and the **Reserved duration**.
- 3 Choose the **Voice sent**: your own as it is, the Character's through ElevenLabs (about 2 s delay), or a local RVC voice (under a second).
- 4 **Start the live session**. The transformed picture appears next to your webcam; switch settings on the fly with **Apply the setting**.
- 5 **Record the live session** keeps the footage in your renders, ready for the Editor or a Recast.

When a voice is converted, the picture is delayed by as much as the voice: lips and sound arrive together. The live bar shows the converted segments, the latency and any losses.

Going further



Background only

The subject of the video is cut out, then placed on a color, a Library image or a video. The acting, the voice and the size do not change.



Change only the voice

Keeps a video's picture and replaces its voice with the Character's. Billed at the ElevenLabs character equivalent.



On the phone

The companion app starts a Recast, films a take or opens a Live session, at your PC's prices and under the same caps.

Good to know



Live is reserved up front, by the minute; only the time actually played is billed. The reserved duration is cut off at Decart even if the page closes.



A tab hidden for 60 s stops the live session, so you don't pay for nothing. Prefer Edge or Chrome: they can delay the picture to match the voice.



Install and train a local RVC voice

You will learn to install Voixbox, record your voice, train it with RVC and have it speak during a Live session, on your own graphics card. Free: nothing goes to the Internet. Allow thirty minutes to install, a quarter of an hour to record and about an hour of computing per voice.

Before you start

You need an **NVIDIA** graphics card with 6 GB of memory or more, and about 6 GB of disk. Without one, Live offers your voice as it is or the Character's voice through ElevenLabs (paid, nothing to install).

The engine is **RVC** (MIT license), through its official WebUI: the same install trains a voice and makes it speak. Only clone your own voice, or that of someone who has agreed to it.

Install Voixbox (once)

- 1 Open PowerShell in the app folder, where `tools\voixbox` lives (by default `%LOCALAPPDATA%\DeepotusVideoGen`).
- 2 Run `powershell -ExecutionPolicy Bypass -File tools\voixbox\installer.ps1`.
- 3 The script installs `uv`, downloads a verified version of the WebUI, creates its Python, adds torch for CUDA (about 3 GB) then its weights (560 MB).
- 4 It ends with `CUDA True | NVIDIA ...`: the card is detected. On any failure, it stops and says why.

Record your voice

The quality of the cloned voice depends almost entirely on this recording.

- 1 Record 10 to 15 minutes of your voice alone: no music, no second voice, no echo.
- 2 A quiet room, the mic 15-20 cm away, always the same mic. Speak the way you do on your live sessions, with varied sentences.
- 3 Avoid clipping. Prefer several 30 s to 2 min files, in WAV.
- 4 Put them in a folder that holds nothing else, for example `C:\Voix\oli`.

Train an RVC voice

- 1 Run `powershell -ExecutionPolicy Bypass -File tools\voixbox\webui.ps1`. The WebUI opens in French at `http://127.0.0.1:7865`; keep the PowerShell window open.
- 2 « Entraîner » (Train) tab, step 1: a short **experiment name** with no spaces or accents (`oli`), rate `40k`, pitch guidance **Oui** (yes), version `v2`.
- 3 Step 2: your recordings folder, and the `rmvpe` pitch algorithm.
- 4 Step 3: save every `25` epochs, **200** epochs in total (100 for a first try), batch `8` (or `4` with 6 GB), small final model at each save: **Oui** (yes).
- 5 Click « Entraînement en un clic » (one-click training). Allow about 1 h 15 for 200 epochs on an RTX 2080 Ti.
- 6 Put the voice where Voixbox looks for it: `powershell -ExecutionPolicy Bypass -File tools\voixbox\ranger-voix.ps1 -Nom oli`.

Listen to it before using it: in the WebUI's « Inférence du modèle » (model inference) tab, pick `oli.pth` and give it a short recording. A crackling voice needs more epochs or a cleaner recording.

Speak with it during a Live session

- 1 Run `powershell -ExecutionPolicy Bypass -File tools\voixbox\lancer.ps1` and keep the window open: it reports how many voices are trained.
- 2 In `Avatar live`, `Live` tab, choose the local RVC voice under `Voice sent`.
- 3 Set `Transpose (semitones)` if needed: `+12` to go from a male to a female voice.
- 4 `Start the live session`. Each half-second of voice is converted in about 110 ms; the picture is delayed by 0.75 s to stay in sync with the lips.

Good to know

i The Live screen tells you what is missing: no NVIDIA card, Voixbox not started, or no trained voice. Another port is set with `VOIXBOX_URL` in the `.env`.



With Avast or an antivirus that inspects connections, always go through the scripts in `tools\voixbox` : they remove the `SSLKEYLOGFILE` variable for themselves only and read Windows certificates. Otherwise Python crashes (`no OPENSSL_AppLink`).




MAKE VIDEOS

Edit a video

You will place your shots on a timeline, cut them, move them and chain them, with the gestures of a professional editor. Free: editing and previews run on your PC.

The editor screen



1. Shots on V1, their sound on A1 right below 2. Preview: a free 480p preview 3. Three views: Media, Edit, Delivery

The timeline at the bottom (video, sound, subtitles), the viewer in the centre, the inspector on the right; Preview, Render and Publish at the top.

At the bottom, the **timeline**: titles (T1), overlays (V2), main video (V1), then dialogue (A1), music (A2), sound effects (A3) and subtitles (S1). In the centre, the viewer; on the right, the inspector of the selected shot. At the top, **Preview** (free 480p preview), **Render** and **Publish**. The bottom bar switches between three views: **Media**, **Edit** and **Delivery**.

Place shots

- 1 Open the **media** drawer: all your finished video renders, with thumbnail, length and origin. Filter by origin (Studio, Chapters, Quick...), by title or by rating (three stars and up).

- 2 Click a media item to place it at the playhead. The **edit mode** decides how: overwrite, insert and push the rest, at the end, above, ripple overwrite, or fill the range.
- 3 From the Library, **Send to** > Edit also places a video on V1 with its sound on A1. A Studio render or a Chapters episode opens straight in the editor.

Cut and adjust



Cut

The blade (**Alt** + **C**) cuts the shot at the playhead. A gap is selected with one click and **Del** closes it.



Roll, slip, slide

Alt + drag the diamond moves the cut; **Alt** + drag the centre slips the content; **Shift** + drag moves the shot between its neighbours.



Range and markers

I and **U** set an in and out range; **Shift** + **X** cuts it on every track. **Shift** + **M** sets a marker, **Ctrl** + **M** lists them.



Everything undoes

Ctrl + **Z** and **Ctrl** + **Y** over fifty steps: shots, mix, tracks, subtitles, range, markers.



Cut at scene changes


Right-click a clip: the app spots each scene change and cuts it for you.




Shots too long

In the menu > View, the timeline highlights shots that run too long and jump cuts from the same source.

Chain two shots



1. Click the diamond between two shots




2. Pick the transition and its duration

A click on the diamond between two shots opens the transitions gallery and frame-by-frame trimming.

- 1 Click the **diamond** between two shots: the gallery opens, with 58 transitions sorted by family (fades, slides, wipes, zooms, pixels, shapes).
- 2 The same panel shows the last frame of the left shot and the first of the right one; shift the cut frame by frame.
- 3 For a quick fade on the selected shot: **Alt** + **T**. Simple fades play live; the others after a **Preview**.

Organise your projects



1. Project, Edit, Timeline, Markers, View, Help

The ≡ menu lists every action in six sections, each with its shortcut.

The ≡ menu lists every action in six sections, with its shortcut. **Projects...** opens the grid of your projects: open, duplicate, rename, delete, or compare two versions. Before opening another one, an unnamed edit is put away safely on its own.



New edit

An empty project, with its default tracks, and nothing else.



Copy between projects

Ctrl + **C** on a clip,
Ctrl + **V** in another project.



Auto-clips

From a long video, AI suggests 1 to 8 excerpts of 15 to 60 s, then creates a project with the chosen one (paid; a free version without AI exists).



Your shortcuts

The **?** panel lists the shortcuts; a “Resolve” preset aligns them with DaVinci Resolve. Export them and import them back.

Shortcuts

Key	Effect
Space	Play / pause
Alt + C	Blade: cut at the playhead
I / U	Range in / out at the playhead
Ctrl + Left / Right	Swap the shot with its neighbour
O	Fold the tools bar
B / T	Sounds drawer / Narration panel
G / F	Safe zones / viewer full screen




MAKE VIDEOS


Grade, mix and add titles

You will give your shots a consistent colour, check them on the scopes, set their speed and framing, add titles and polish the sound. Free: everything is computed on your PC, except subtitle transcription and translation.

Grade a shot



1. Shadows, midtones, highlights




2. Match a shot to the previous one

Clip inspector, Grading panel: wheels, curves, mask, automatic match.

- 1 Select a shot: the inspector shows the **Grading** panel.
- 2 **Wheels**: shadows (lift), midtones (gamma) and highlights (gain). **Curves**: up to sixteen points, on the whole picture or each colour.
- 3 **Mask** limits the setting to a rectangle or an ellipse, feathered or inverted.
- 4 **Auto** removes a colour cast; **Match to previous clip** aligns the colours with the shot before.
- 5 **Copy grade** then **Paste grade** carry the setting over to another shot.

When playback stops, the viewer shows the shot's real graded picture; during playback, the picture stays raw.

Check on the scopes



1. Scopes are computed when playback stops

Waveform, vectorscope and histogram of the graded picture, in a window you can move.

Scopes opens a movable window: waveform, vectorscope and histogram of the graded picture under the playhead, recomputed when playback stops. They account for speed, stabilisation and the render frame rate, and stay available in full screen.

Effects, speed and framing



The effects rack

Noise reduction, deflicker, debanding, sharpening, green-screen keying, monochrome, per-band hue and saturation, trails, lens distortion.



Speed and ramp

From 25 to 400 %. Interpolation smooths the motion; **Split at playhead** gives the rest of the shot another speed.



Stabilise and reframe

Stabilisation analyses a shaky shot; framing picks the part kept: centred, following the motion, or manual; dynamic zoom animates a reframe.



Adjustment clip

On track J1 (**Shift** + **J**), its effects apply to everything below, with fade in and fade out.



Overlays

Position, scale, rotation and opacity animated by keyframes, rounded corners, drop shadow.



Hand-drawn ramps

An effect ramp can follow a Bézier curve you draw; the shake's ramp raises or lowers its amplitude.

Titles and subtitles

- 1 **Shift** + **T** places a title card on track T1, burned in at render time.
- 2 Pick its template (full frame, lower third, caption, counter, chapter, quote, hashtag, call to action), its text, colour and font.
- 3 Subtitles have one track per language: only one is burned in, the others are exported. Translation can create a new language track.

Polish the sound



The audio rack

Six-band equaliser, noise reducer with adjustable floor, 50 / 60 Hz hum removal, stereo pan; listen to the rendered result.



Learn the noise

Set a range on a stretch of noise only: the app learns it and removes it at render time.



Voiceover from the mic

voice-over (**Alt** + **R**) records a take while the edit plays and places it on the dialogue track.



Narration

The **narration** panel (**T**) writes and synthesises the voice block by block; only "Narrate" sends a block to the provider.

Good to know



In the ☰ menu › View, the shot lightbox shows every V1 shot with its graded picture: one click goes to the shot.



Subtitle transcription and translation are paid; the rest of this chapter is free.




MAKE VIDEOS

Deliver: renders and formats

You will export your edit in the right format, at the right loudness, and send it to the Scheduler or to another program. Free: rendering happens on your PC.

Render the video



1. Pick the preset, then render or queue

Render... opens the render panel: output preset, frame rate, loudness, range, render queue.

- 1 Click **Render** in the top bar: the render panel opens.
- 2 Pick the **preset**: master 1080, web 4K, social 720, ProRes 422, H.265, WebM, audio only (AAC, MP3, WAV) or GIF 480. The project's format is kept.
- 3 Set the **frame rate** (24, 25, 30 or 60 fps) and the **loudness**: -14 LUFS for YouTube and TikTok, -16 for a podcast, -23 for broadcast. A badge compares the last measurement with the target.
- 4 **Render** starts the render; **Add to queue** puts it behind the others. To export only the I/O range, tick the box provided.

- 5 When it is done, a banner offers to send the video to the Scheduler, with your channels, a date and a prefilled caption. Nothing is published without you.

The Delivery view

The bottom bar > **Delivery** lists the project's renders and previews, their status (queued, running, done, failed), and leads to the Library. **Publish** sends the last final render back to the Scheduler.

Go further



Your own presets

Save this setting... keeps your output setting as a preset of your own, listed with the others.



To another program

 > Project >

Export EDL... or

Export FCPXML... to take the timeline into DaVinci Resolve or Final Cut.



True to the timeline

The render lasts exactly the timeline, frame for frame; a fade is centred on the cut with the source's real frames.

Good to know




Run a 480p **Preview** first: it is free and shows transitions, effects and mix.



PUBLISH

Plan and publish your week


Scheduler tab: the publishing calendar, and the heart of "human-in-the-loop" — **nothing goes out without your approval** (unless you decide otherwise).



Scheduler: week/month, AI plan, document import, the post's graph + inspector on the right.


Three ways to fill the calendar

- 1 **Generate plan:** describe your week in one sentence → the AI proposes a full plan (dated posts, hooks, captions, formats). *Review the preview* before "Add to calendar".



"Generate plan": a one-sentence brief becomes a full plan, reviewed by you (human-in-the-loop).

- 2 **Import doc:** drop your existing plan (.md, .docx, .pdf) — transcribed into dated posts, **weeks preserved**.




"Import doc": your strategy document becomes calendar posts.

- 3 **New post** or + *add post* on a day: manual creation. Episodes also land here as drafts (chapter 7).

Generate the render — before posting

Select a post: the right inspector opens the post's whole workshop. **Produce** button: the app builds the image (FLUX) then the clip (Seedance) from the plan's ideas — or reuses the source you chose. A **thumbnail** of the render appears on the post-graph's *Render* node, and a **video preview** shows in the inspector: you see the result **before** clicking Send or enabling Auto.



A post's inspector: channels, caption, and the **emoji bar** + the **Telegram Premium pack** (branded one-tap tags) — plus the render preview before publishing.

Polish the caption — emojis & Premium pack

Under the caption, two one-tap rows, ideal for Telegram and X:

- 1 Quick emojis:** 🍷 🔥 🚀 🥰 🎉 🍀 💪 💰 — inserted at the cursor.
- 2 Telegram Premium pack:** branded tags per content vertical (🧙 Deepotus Protocol, 🗨️ Rippled Signal, 👁️ Prophet, 🎲 Board Game, 🎲 D&D, 🏆 \$DEEP...) to sign each post in one gesture.

Auto or assisted?


Each post has a mode. **Assisted** (default): at the set time, the post flips to READY — you click *Send*. **Auto**: it goes out by itself on connected channels (Telegram, X). Status by color: gray = draft, cyan = scheduled, amber = ready, green = posted, red = failed.



PUBLISH

Connect your accounts

Settings → **Connected accounts**: the credentials the Scheduler uses. All stored locally.



Connected accounts: real status per channel — **auto** (publishes alone) or **assisted** (you click Send).

Telegram (5 minutes, free, recommended)

- 1 On Telegram, talk to `@BotFather` → `/newbot` → copy the token.
- 2 Add the bot as an **administrator** of your channel.
- 3 Paste token + chat ID, *Save*, restart, then *Send test message*.

X (Twitter)


Create your 4 keys at `developer.x.com` (writes available pay-per-use) and paste them. *Verify credentials* button. **YouTube** and **Instagram** stay assisted.



GAMES AND 3D


Make game assets

The **Game Assets** category (sidebar) turns an image from your Library into a **textured 3D model** ready for a game engine or 3D tool (Unity, Unreal, Godot, Blender...). Only the **fal.ai** key is required. Five generation engines, from cheapest to premium: **TripoSR** ($\approx \$0.07$), **Hunyuan3D**, **Tripo** (default — best quality/price), **TRELLIS** and **Rodin**. The estimated cost is shown live before you start.




The Game Assets form: source image, engine (with description and rate), subject, options and export formats.

- 1 Pick a **source image** (your Library images — ideally a well-isolated character/object on a neutral background), then the **engine** from the dropdown.
- 2 Options: **Multi-view** (on by default, 4 views) first generates other angles via Seedream for a more faithful mesh; **Textures**; **T-pose** for characters you plan to rig; and the **export formats** — GLB always included, plus FBX / OBJ / STL / USDZ depending on the engine. A format the chosen engine cannot emit is flagged; each extra format can cost one more generation (the estimate shows it).
- 3 **Generate** opens the **pipeline canvas**: the Input → Views → Engine → Response nodes animate and fill in live (also visible when reopening an asset via its card title).



The live pipeline: each node fills in as the generation progresses (views, engine, result).

- 4 On the result card: ► **Rotate** spins the 3D model in the built-in viewer; each **format** downloads in one click; the **Shots** strip lets you download each view or send it to *Library* → *Images* (★) to reuse as a source. Rename (✎) or delete (🗑️) the asset from its card header.
- 5 All your assets also live in **Library** → **3D tab**: preview, full-size 3D rotation, GLB download, favorites ★, rename and delete.




The Library 3D tab gathers all your generated assets.

⚠ 3D rendering runs in your browser (WebGL): keep the tab in the foreground for smooth rotation. Files live under `DeepotusVideoGenData\assets\outputs\assets3d\` and follow your migration (chapter 16).



Create playing cards (Card Forge)


The **Card Forge** sub-tab (inside **Game Assets**, next to 3D / Sprites / Tiles / Materials) is a full playing-card editor: face, frame, typography, CSV data, 3D volume, PBR materials, print export and 3D export — eight modules in one document. No key is required; only **AI face generation** uses your existing fal.ai key.



1. Pick the format and resolution: they apply to the whole deck 2. The catalog offers 108 ready-made illustrations 3. The gauge tells you if the picture is sharp enough to print

The Face module: format and resolution at the top, a 108-illustration catalog, the effective-resolution gauge.

- 1 At the top of the lab, pick the **format** (poker, bridge, tarot, domino...), the **resolution** (150 / 300 / 600 DPI), the **bleed** and the **safe zone** — these settings apply to all eight modules and to the final export.
- 2 Browse the modules in the left rail: **01 Face** (imported illustration, AI-generated, or picked from the local catalog), **02 Frame** (borders, rarity, card back), **03 Typography** (title, rules box, flavor text, 23 already-bundled fonts), **04 Data** (CSV import → columns to fields → full deck), **05 Volume** (3D preview with real thickness and edge), **06 Materials** (8 PBR texture channels computed locally).
- 3 **07 Print** runs a **preflight check** on both the sheet and the produced file: bleed, cut marks, the PDF's **TrimBox** / **BleedBox**, the PNG's **phys** density (300 DPI = 11811 px/m) — every line is verified against the delivered file's bytes, never against an on-screen setting.



The preflight check: 0 errors, TrimBox and BleedBox present in the PDF, PNG density verified byte by byte.

4

08 3D Export exports both a **.glb** and a standalone **.gltf** with the full PBR map set (base color, normal, roughness, metallic, occlusion, height, emissive, ORM), ready for Unity, Unreal, Godot or Blender.




The **emissive** channel is now adjustable in the 3D export (an **Emission 0..1** slider — luminescent ink on matte paper, or extinguished gilding) and wired into the GLB, glTF and MTL; it stays capped at 1.0 and its mask is derived from the card's luminance (adjustable threshold). A pose hidden by the frame is fixed **in one click** from the Face module — a panel button plus a contextual offer next to the masking percentage.



From image to 3D object

Card Forge's **09 Forge 3D** module turns your card into a printable 3D piece or an "NFT-grade" artifact: the card exports **as layers** (base material, illustration, veil, frame, typography, ornaments + the composite), each layer is processed in a **graph**, and assembly produces a clean GLB — plus an STL whenever everything is watertight. The whole free circuit runs locally, no key needed.



The Forge 3D graph: a material chain (silver holographic finish) on the background, an image→3D engine on the illustration — 30 cr (~\$0.60) shown before any launch.

- 1 Export the layers:** the screen proves the stacking **before** uploading (0 px of difference between the composite and the layer pile), the server counter-proves it in PIL and seals a manifest. An empty layer is still written — absence is measured, never guessed.
- 2 Edit the graph:** each layer gets a **plane** treatment (flat card) or a **relief** (a grid displaced by luminance, a solid closed by construction), a **material** from the Material Forge shop (tiled at physical scale, in millimeters), a **silver or gilded holographic** finish (iridescence, clearcoat, anisotropy — verified in the GLB's bytes) and a **transform** (position, rotation, scale, z stacking).

3 **Image→3D engines** (optional, paid): the **mesh 3D** treatment sends the layer to one of 7 engines — tripot, hunyuan, trellis, rodin, tripot (fal, priced in \$) or **Mesby 6/7 through the direct API** (credits; ultra +5 cr on mesby-7). The price shows on the node **and** in the graph footer before any launch; the status chip follows the job (queued → running → served · real credits → or failure with the literal message); relaunching a node resets its folder.

4 **Build**: assembly merges the engines' GLBs at their layer's place (full reindexing, the engine file's identity is dropped), writes the final GLB and the **STL** — accepted only when every element is a closed solid, with a motivated refusal otherwise. The docket lists the elements, their engines and the credits consumed.

 The image→3D engines are the only paid part of Card Forge: fal bills in \$ on your existing key, Mesby in credits on yours (`MESBY_API_KEY` , the Settings field shared with the 3D Studio). To try everything without spending a credit, set `MESBY_MOCK=1` in the data `.env` : the local simulator runs the full circuit, fake credits included.



Prepare a part before slicing (Workbench)

Chapter 20 showed you how to get a file *out*. This one shows what to do *before*. The **Workbench** screen (Game Assets → 3D Studio → 07 · Établi 3D →) is the workshop where you repair, lay down, inspect, hollow, arrange and measure a part, so the slicer only has to slice. Nothing below is irreversible: **every gesture that changes the part writes one more VERSION**, and all of them stay on disk — the *model's life* column on the left lists them. (The Workbench screen itself is in French, like the rest of the application; its labels are quoted below as they appear.)

- 1 Set a target size.** Right rail, *taille cible* field: type the largest wanted dimension, in millimeters (80 for a figurine, 250 for a board). A 3D file carries no unit — while this field is empty the Workbench shows "glTF units" and refuses whatever needs a physical dimension (hollowing, arranging, exporting for print). That is deliberate: a refusal beats an invented millimeter.
- 2 Pick the printer.** Same rail, *imprimante* menu: the **Elegoo Centauri Carbon 2** profile is built into the application (256 × 256 mm bed, 256 mm height, one excluded area). If OrcaSlicer or ElegooSlicer is installed, its profiles show up too — read, never modified. On the plate, the green outline IS the chosen machine's bed; the red rectangle is the area where nothing may sit.
- 3 Repair in one click.** *Fiche* tab → *Réparer en un clic*. Checked by default: coincident vertices welded, duplicate faces dropped, flat triangles dropped, normals unified. *Trous bouchés* (cap holes) is unchecked by default: capping adds material, tick it knowingly. The detail appears in the bottom bar — "8 retourné(s) — fermé" (8 flipped — closed). A mesh left open prints badly: better to know here than at the machine.
- 4 Lay the part flat.** *Poser sur une face* (key F), then click the face that must touch the bed: the part turns and lands on the floor. The gesture joins the queue — the bottom bar says "1 modification en attente" — and *écrire la version* saves it. This gesture decides almost everything else: a well-placed part needs few supports. If unsure, *Orienter* proposes three ranked poses in the *Fiche* tab — share of the surface on the bed, share overhanging, height — and one click lays the part on the one you pick. A computed orientation does not always find the best pose: **always review the proposal**, no score knows which face you want to look good.
- 5 Look at the overhangs, then the slices.** *Surplombs* button: faces leaning less than 45° from horizontal turn orange — there, a layer has almost nothing under it and the slicer will add a support. Re-lay the part until the orange recedes. *Tranches* button: twenty cross-sections of the part, in blue, to see its real shape at each height. Both are **indicative previews**: nothing is written, and the slicer does the real slicing.

6 **Hollow it, when it helps.** *Fiche* tab → *Creuser*, 2 mm wall to begin with: the solid inside is replaced by a shell. It is first a **resin** (SLA) gesture, where a solid part costs a lot of material; with filament (FDM), the slicer already fills the inside with partial infill. It needs a *closed* mesh (step 3), and the Workbench states the thickness that holds when the requested wall is too thick for the part. For resin, then drill a drain hole: same tab, *Percer (drainage)*, a 3 to 4 mm diameter, then click the face to drill — preferably near the build plate, so the resin can flow out. The drill goes through the one wall under the click and stitches the hole closed; the opposite wall is left untouched. A solid part is refused (hollow it first).

7 **Arrange on the bed.** *Sur la plaque* separates the pieces, then *Ranger sur le plateau* lays them inside the real bed, rotated when that saves room, 2 mm apart. What does not fit moves to a second plate, and what is larger than the machine is reported. Drag a piece with the mouse (snapped to the bed grid, Shift frees it), nudge it with the arrow keys; its position, size and angle show under "Sur la plaque". The plate is a **view**: it is saved, but it does not change the model.

8 **Measure, then export.** *Mesurer* → two clicks give a distance, its x/y/z components and the angle between the two faces. Finally, *Export* tab → → **Impression 3D** writes the STL and the 3MF of the **displayed version**, at the target size in millimeters (chapter 20), and *Ouvrir dans le slicer* takes you there.

i **The order that works:** repair → lay down → inspect → hollow → export. Repair first, because every other gesture starts from sound geometry — and hollowing refuses an open mesh.

Glossary

The words your slicer will use, explained once.

Term	What it is
Base (bed contact)	The face that touches the bed. Wide and flat, the print holds; narrow, it comes loose.
Overhang	A wall leaning too far to rest on the layer below. Prusa puts the clean limit without support between 45 and 60° from horizontal.
Support	A disposable structure holding overhangs up. The slicer generates it; the Workbench only helps you need less of it.
Brim	A one-layer collar around the part, to keep its corners from lifting.
Raft	A platform printed UNDER the whole part. Safer than a brim, costlier, and it leaves a rough face.
Skirt	An outline drawn next to the part, not touching it, to prime the flow before starting.
Infill	The pattern inside. <i>Gyroid</i> is about equally strong in every direction; <i>Lightning</i> only supports top surfaces. 15 % is usually enough.
Seam	Each wall loop starts and ends somewhere, leaving a visible vertical line. The slicer places it (<i>Nearest, Aligned, Rear</i>) or scatters it (<i>Random</i>).
Retraction	Filament pulled back during travel moves, to avoid strings (<i>stringing</i>).
Layer	A horizontal slice. 0.2 mm is the usual compromise; 0.12 mm for detail, 0.28 mm for speed.
Perimeter (wall)	The number of wall loops. Two perimeters at 0.45 mm make a 0.9 mm wall: the minimum to aim for when modeling.
Bridging	Printing a layer over empty space, between two supports. It works over a few centimeters, with good cooling.
Warping	Corners lifting as they cool. Remedies: clean bed, brim, closed enclosure. Common with ABS/ASA, rare with PLA.
Watertightness	A solid is watertight when its surface is closed, with no hole or free edge. That is what the Workbench bar calls "fermé" (closed).
Manifold	Every edge belongs to exactly two faces. A mesh can be closed and still be hard to slice cleanly if it is not.

Decimation	Reducing the triangle count. A 500,000-triangle mesh prints no better than a 100,000 one, but slows everything down. <i>Fiche</i> tab → <i>Décimer</i> .
Hollowing	Replacing the solid inside with a shell of given thickness. Mostly useful in resin; 2 mm is a good starting point.
Drain hole	The hole(s) letting uncured resin (and air) out of a hollowed part. In the Workbench: <i>Fiche</i> tab → <i>Percer</i> (drainage).

Where to learn more

Every link below was opened and checked on 05/10/2026 (5 October 2026). They are grouped by what they teach you.

Topic	Resource	Checked
Supports and overhangs	Prusa — Support material (https://help.prusa3d.com/article/support-material_1698) (the threshold is measured from horizontal, 90 = vertical)	05/10/2026
Modeling for print	Prusa — Modeling with 3D printing in mind (https://help.prusa3d.com/article/modeling-with-3d-printing-in-mind_164135) (45–60° overhangs, walls, tolerances)	05/10/2026
Bridges	Prusa — Poor bridging (https://help.prusa3d.com/article/poor-bridging_1802)	05/10/2026
Layers and perimeters	Prusa — Layers and perimeters (https://help.prusa3d.com/article/layers-and-perimeters_1748)	05/10/2026
Infill	Prusa — Infill patterns (https://help.prusa3d.com/article/infill-patterns_177130)	05/10/2026
Seam	Prusa — Seam position (https://help.prusa3d.com/article/seam-position_151069)	05/10/2026
Skirt and brim	Prusa — Skirt and brim (https://help.prusa3d.com/article/skirt-and-brim_133969)	05/10/2026
Warping	Prusa — Warping (https://help.prusa3d.com/article/warping_2011)	05/10/2026
PLA filament	Prusa — PLA (https://help.prusa3d.com/article/pla_2062) (for Prusament PLA: nozzle 215 °C first layer then 210 °C, bed 60 °C)	05/10/2026
PETG filament	Prusa — PETG (https://help.prusa3d.com/article/petg_2059) (nozzle 230 °C then 240 °C, bed 85 °C then 90 °C)	05/10/2026
Choosing a filament	Prusa — Filament material guide (https://help.prusa3d.com/filament-material-guide) (comparison table of materials)	05/10/2026
Print defects	Simplify3D — Print Quality Guide (https://www.simplify3d.com/resources/print-quality-troubleshooting/) (about 26 defects, with photos)	05/10/2026
Slicer — prepare	OrcaSlicer — Object manipulation (https://github.com/SoftFever/OrcaSlicer/wiki/prepare_object_manipulation)	05/10/2026
Slicer — orient	OrcaSlicer — Auto orientation (https://github.com/SoftFever/OrcaSlicer/wiki/prepare_auto_orient) (which warns it may not always find the best pose)	05/10/2026

Slicer — calibrate	OrcaSlicer — Wiki (Calibrations) (https://github.com/SoftFever/OrcaSlicer/wiki/Home) (temperature, flow, retraction, pressure advance...)	05/10/2026
Machine	Elegoo — Centauri Carbon 2 wiki (https://wiki.elegoo.com/centauri-carbon-2-combo) (troubleshooting guides and error codes)	05/10/2026

 These links come from slicer makers and manufacturers: they describe THEIR machine and THEIR software. Numeric settings (temperatures, speeds) apply to the hardware they name; on another machine they are starting points, not truths. Two repair tools often recommended elsewhere, **Microsoft 3D Builder** and **Autodesk Meshmixer**, are no longer maintained, so they are not listed here.



3D-print your creations

Everything the app produces in 3D — and even your vector illustrations — exports to your printer in one gesture. The circuit is designed for the **Elegoo Centauri Carbon 2** (256 × 256 × 256 mm build plate) and its slicer **ElegooSlicer** (bundled with the machine, an OrcaSlicer fork — OrcaSlicer works too). Every export writes a folder under `DeepotusVideoGenData\assets\print3d\` with three files: the **STL** (the mesh), the **3MF** (the one you open — it carries the millimeter unit and the name), and `impression.json` (provenance, scale, watertightness). Everything is local and free: no provider calls.

- 1 Install ElegooSlicer once.** The installer associates `.3mf` files with the slicer on Windows: the app's "Open in slicer" button uses that directly. Without an association the app tells you so, and you can set `SLICER_PATH` (the slicer executable path) in the data `.env`.
- 2 From the card Forge 3D** (Card Forge, module 09): when the build slip delivers an STL — that is, when every element is a **closed solid**, the check that guarantees it — a → **3D printing** button appears next to it. Leave the size empty to keep the card's real millimeters, or type a target (e.g. `80` for a figurine). These exports are marked *guaranteed* watertight: the producer proved it.
- 3 From Game Assets · 3D:** every generated model (card characters, props) carries its → **3D printing** button. Give the largest dimension in millimeters — 80 for a figurine, 250 for a board — and the app centers the piece and rests it on the bed. The converter reads the original `model.glb`; the "optimized" file (compressed for game engines) is refused with a message that says what to use instead.
- 4 From the Vectorlab** (Export menu → 3D printing...): your vector illustrations become **printable reliefs**. Give a global height in millimeters, then override per layer with `name=mm` — for example `2, contours=5` for a stained-glass piece: glass at 2 mm, leads at 5. Strokes without a fill count by their stroke width; text objects are skipped (and counted in the message). Flat dimensions follow the document's unit (an A4 comes out at 210 × 297 mm).
- 5 Slice in the slicer.** The `.3mf` opens at the right size, in millimeters, resting on the bed. Supports, fine orientation, infill, filament profiles: that's the slicer's job — the app never generates G-code. Beyond 256 mm the export warns you the slicer will have to cut or shrink, without blocking you.



Meshes coming from the image→3D engines are marked *unknown* watertightness in `impression.json`: ElegooSlicer and OrcaSlicer repair small defects on import and say so. Only the card Forge 3D guarantees its solids — its "closed solid" check states it, not a promise.



End-to-end recipes, from idea to post



Recipe A — a meme posted to Telegram in 5 minutes

1

Library → "Create image": describe the meme (~3 s).

2

Scheduler → *New post*: caption (+ a 🔥 emoji), Telegram channel, time in 5 min, *Auto* mode.

3

Or don't wait: **Send now**. Posted.



Recipe B — your selfie video + an animation (UGC)

1

Film 10–15 s on your phone, transfer the file to the PC.

2

Studio → "UGC + animation" starter → *UGC video* node → *Upload your video* (🕒 Duration master on).

3

Choose the image + prompt of the animated half (Seedance) → *Run*. The output lasts exactly your clip.

4

Scheduler → attach the render, polish the caption (emoji + pack), schedule.



Recipe C — today's news reel

1

News → *Refresh* → tick 2–3 articles.

2

Build script (persona voice) → *Build illustration*.

3

Compose in **Templates** → *Send to Scheduler* → schedule or publish.



Recipe D — an episode of your novel

1

Episodes → title + paste (or upload) the chapter → pick a voice & language → *Generate narration*.

2

By AI (or By paragraph) → review the scenes → *Generate all illustrations* → set each scene's motion (Ken Burns / Seedance).

3

Assemble the episode → preview → *Send to Scheduler* (YouTube + Instagram drafts). One episode a day, done.



Read the icons

You will recognise an icon at a glance, find its name, and print the full sheet to keep next to your screen. Free, no key needed.

One icon, one function

The whole app speaks the same language: the **Deepotus Glyph** set, 526 icons drawn with the same stroke. Each one stands for **a single function**: the same drawing means the same thing in the editor, the Photolab or the Card Forge. The old emojis and symbols are gone from buttons, menus and lists.

- 1 Rest the pointer on a button: its tooltip gives its name and, if it has one, its keyboard shortcut.
- 2 Look that name up in the icon sheet to read what it does and what not to confuse it with.
- 3 In the Vectorlab, the Photolab, the Sprite Lab and the Tile Lab, wait a little longer (900 ms) or click the small  button: an animated card SHOWS you the option.

Icon families

Icons are sorted into families. Knowing the family already tells you what to expect when you click.



Navigation

A screen, a tab or a panel: you go there, nothing is changed.



Actions

A button that does something: create, save, send, delete.



Editing

Change something already there: align, flip, group, lock.



Tools

The drawing and retouching tools of the labs.



Layers and adjustments

Layers, masks, styles and image adjustments.



Media and playback

Play, cut and edit video and sound.



States

What the app tells you, without you having to click.



Categories

Content families, each with its own colour.



3D

Objects and actions of the 3D workshops.

States worth knowing

A state icon is not clicked: it informs you. These are the ones to read before starting a paid render.

Icon	What it says
✓	Done: the render, the upload or the save worked.
⌚	In progress or waiting: loading, generating, or queued.
✖	Failed: the tooltip or the message next to it says why.
⚠	Warning: nothing is broken, but read before going on.
🔄	Spending limit close or reached: beyond it, the app refuses the spend and asks you what to do.
🧠	Generative field: AI can fill it in for you.
🔒	Locked: nothing can move, change or delete the item while it is.
📌	Pinned (Studio): this node result is reused for free on the next render.

Print the icon sheet

The [icon sheet](#) shows them all, sorted by family, with their name, their function and the look-alike icons not to confuse them with.


- 1 Open the [icon sheet](#).
- 2 To print only part of it, click a family (for example **Vectorlab tools**) or type a word in the search: only what is shown gets printed.
- 3 Click **Print**: each family starts on a new A4 page, black on white. You can also download the [full sheet as a PDF](#).

i On the sheet, the red line starting with "≠" names the look-alike icons: Close hides a window, Delete destroys, Remove takes something out of a list without destroying it.

Animated cards in the labs

In the labs, some options have an animated card in three beats made of real screenshots: the starting point, the gesture, the result. It opens when you rest the pointer on the option for 900 ms, or with the  button next to it (in the Sprite Lab and the Tile Lab, even when the option is greyed out). Here they all are; cards of labs not yet translated still show the French interface.


Vectorlab



Conical

The color spins around the middle of your shape, like the spokes of a wheel.


Appearance > Conical



Transparency

Your shape slowly fades from one side to the other, as if it vanished into fog.

Appearance > Transparency



Pattern

Your shape fills up with little repeated stripes, like striped fabric.


Appearance > Pattern



Frame

Your long line of text fits into a box and wraps by itself, like in a book.


Appearance > Frame



On path

Your words lie down on the line you drew and follow it wherever it goes.


Appearance > On path



Contours

Lines join the places of the same height on your map, like the steps of a mountain.


Real map > Contours



Slice

Your map turns into six-sided tiles, colored by height: water, meadow, forest, mountain.


Real map › Slice



Side draft

The sides of your piece lean: the more you turn, the wider the bottom and the smaller the top.


3D printing › Side draft



Set as reference

Your picture becomes a reference: it fades, stays still, and a grid of squares lies on it to guide you.


Pixel persona › Set as reference



Create pixel layer

A see-through sheet lies on your reference picture: you draw on it, square by square.


Pixel persona › Create pixel layer



Pixelate image

Your picture is now made of big squares, like bricks: as many squares as the tile size.

Pixel persona › Pixelate image




Dark outline

A dark line goes around your drawing so it stands out from the background.

Pixel persona › Dark outline

Photolab



Workspace

Pick a workspace in the list: panels and tools rearrange for what you want to do.


Workspace list



Rectangular marquee (M)

Drag across the picture: a dotted rectangle surrounds the part you are going to change.


Tools › Rectangular marquee (M)



Magic wand (W)

Click a colour: the whole patch of that colour is selected at once.

Tools › Magic wand (W)



Crop (C)

Draw a frame on the picture then press Enter: everything outside the frame disappears.


Tools › Crop (C)



Brush (B)

Drag across the picture: a stroke of your colour is painted wherever you go.

Tools › Brush (B)



Clone stamp (S)

Alt-click what you want to copy, then paint elsewhere: the same piece of picture reappears under your brush.

Tools › Clone stamp (S)




Black and white

One click adds a layer that removes the colours: your picture turns grey, and you can remove it whenever you like.

Adjustment layer › Black and white


Sprite Lab



Auto-detect

The tool counts the pictures on your sheet by itself and draws a box around each one.


Sheet › Auto-detect



Feet

All your characters put their feet at the same height, as if standing on the same ground.


Sheet › Alignment › Feet



Mirror

Your character turns around and looks the other way, like in a mirror.


Preview › Mirror



Pixel art (9b)

Your drawing becomes big squares and only a few colours, like in an old video game.


Pixel art (9b) box



Outline

A line surrounds your character so it shows up on any background.

Post-processing › Outline




Shadow X / Y

A shadow appears behind your character, slightly offset, as if a lamp were lighting it.

Post-processing › Shadow X / Y


Tile Lab



Method

Mirror folds your texture like a sheet folded in four: the edges always meet perfectly.


Seamless mode › Method



Make seamless

Your texture joins up: laid side by side many times, you can no longer see where each piece ends.


Seamless mode › Make seamless



Placement grid

Your cells become diamonds, staggered like the paving of a game seen at an angle.


Tile sheet › Placement grid



Set

You choose how many pieces to make: many for nicely rounded corners, or sixteen to go fast.


Set mode › Set



New preview

A new map draws itself at random with your tiles, to see if they go well together.

Set mode › New preview



Shape

Your tile changes shape: a diamond for a game seen at an angle, or a six-sided hexagon.

Shapes mode › Shape

Good to know

i The icon of the app, of its shortcuts and of browser tabs is the real Deepotus logo. Inside the app, the logo shown follows your active brand kit (chapter [Templates and brand kits](#)).

i One single chevron is used everywhere to expand and collapse: pointing down, the section is open; pointing right, it is closed.



Troubleshooting

Symptom	Fix
<code>down</code> pill top-right	The engine isn't running: relaunch the Desktop icon.
Red pill on <code>fal</code> / <code>heygen</code> / <code>voice</code>	Missing/invalid key: Settings → API keys, fix, restart.
HeyGen avatars are slow to appear (or the list looks empty at first)	Normal on the very first load: HeyGen returns 1,000+ avatars and takes up to ~1 min to respond (their service is slow). The app pre-loads the list at startup and caches it — wait once, then it's instant. If the HeyGen status is green in Settings, your key is fine: it's neither the key, the network, nor the antivirus.
A voice gives a "paid plan" error	On the ElevenLabs free plan, only "premade" voices work via the API; custom library/community voices need a paid ElevenLabs plan. Pick a premade voice, or upgrade ElevenLabs.
After a restart, keys/library "gone"	No data is lost (it lives in <code>%LOCALAPPDATA%\DeepotusVideoGenData</code>). Relaunch cleanly via the Desktop icon.
The plan uses the built-in planner	No Anthropic key nor Ollama model (chapters 11–12).
<code>CERTIFICATE_VERIFY_FAILED</code> error	Antivirus inspecting HTTPS — handled automatically; otherwise disable HTTPS inspection.
"Port 8765 already in use"	The app is already running — open <code>http://deepotus.localhost:8765</code> .
Where are my files?	Images: <code>...\DeepotusVideoGenData\assets\images</code> · Videos: <code>...\assets\outputs\final</code> · Your uploads: <code>...\assets\outputs\uploads</code> · Sounds: <code>...\assets\audio</code> .

Deepotus Video Gen v2.8.0 — guide built from the real app. Everything runs locally on your machine; only calls to providers (fal.ai, HeyGen, ElevenLabs...) go over the Internet, with your own keys.